



RTL Design Sherpa

The Monitor System as a Design Surface AMBA Monitor System White Paper — Rev 1.1

September 2026

Table of Contents

1 The Monitor System as a Design Surface.....	2
2 The spine.....	2
3 Which monitor: the lite, unless you can say why not.....	3
4 Identity space allocation.....	5
5 Where to insert monitoring.....	6
6 Timestamp policy.....	7
7 Drain path selection.....	8
8 Packet-type filtering.....	9
9 Aggregation topology.....	10
10 What it costs.....	11
11 Validating a tweak in simulation.....	12
12 Out of scope.....	12

1 The Monitor System as a Design Surface


Audience: SoC integrators deciding how to instrument a design with the RTL Design Sherpa monitor system. **Not** a status snapshot of what is in the tree (that is [overview.md](#)) and not a port list (those are the per-module pages under [monitor/](#)). This paper sits one level up: here is the spine, here are the six axes the integrator owns, and here is what each choice costs.

Version: 1.1, 2026-09-28. Numbers are from the tree at that date; each is cited to the page or report it comes from. Revision 1.1 makes the recommendation explicit: the lite monitor is the realistic choice on every port, and the full monitor is the exception you justify.

2 The spine

Every monitored port emits 128-bit packets onto a monitor bus (monbus). An arbiter tree merges the ports of a block into one stream, and a monbus group sinks that stream into two paths: an

error FIFO the host reads over AXI-Lite behind an interrupt, and a write FIFO the group flushes as AXI bursts into a memory ring. Filtering happens at the group, per protocol and per packet type. The packet layout is fixed ([monitor_package_spec.md](#)); everything else on this page is a knob.



Monitor system spine: ports, arbiter tree, group, two drains

3 Which monitor: the lite, unless you can say why not

Two monitors implement the port end of the spine. `axi_monitor_base`, the full monitor, tracks transactions in a CAM and drives six reporter cones (error, timeout, completion, threshold, performance, debug) with ID and address filtering. `axi_monitor_lite` tracks the same transactions in per-ID linked lists, emits errors, timeouts, completions with latency, thresholds and address-range events, and drops the performance and debug classes and the ID filter. Measured at their default parameters, out of context, on the Kintex-7 the boards ship on ([monitor/monitor_characterization.md](#)):

	Full monitor	Lite	Ratio
AXI4 read monitor, LUTs	7,006	1,115	6.3x
AXI4 read monitor, flops	5,219	1,027	5.1x
AXI-Lite read monitor, LUTs	2,891	929	3.1x
6.667 ns on Kintex-7 325T -2, register to register	misses (CAM at 16 slots)	meets, 0.8 to 1.5 ns spare	
10 ns on Artix-7 100T -1	misses by 1.45 ns	meets by 0.6 ns	

That is why the lite is the default: the bridge generator instantiates it on every monitored port, STREAM's three in-core monitors and RAPIDS' descriptor and AXIS monitors are lite, and every board build since 2026-09-26 carries it. It is also the more honest of the two under load. A full monitor that runs out of table entries stalls the command channel (`block_ready`); the lite never stalls the bus, counts the command it could not track (`refused_count`), and when its output queue is full it counts the events it could not deliver and reports the count in one `Error/EVENT_DROPPED` packet as soon as the bus is free. Its contract is exact and is proved in the tree: every command is completed, refused or live, and every generated event is delivered, reported dropped or pending, with no fourth outcome (`val/amba/test_axi_mon_block_ready.py`, `test_axi_monitor_soak.py`, `formal/amba/axi_monitor_lite`). A stalled bus is the one effect an instrument must never have on the design it watches.

What you give up, and when it matters:

Full-monitor feature	Lite	Reach for the full monitor if
performance windows (<code>cfg_perf_enable</code> , byte/beat/idle counters)	none; the in-datapath <code>axi_bus_meter</code> and the pass-through observer measure throughput without a monitor	you need windowed bandwidth numbers from the monitor itself
debug trace packets	none	you need per-beat trace, not events
ID-range filter (<code>ID_FILTER_ENABLE</code>)	none; one lite per port, sliced by <code>unit_id/agent_id</code> instead	several instances must share one port by ID
16 slots	8 by default, parameter	more than eight transactions in flight per port and every one must be tracked, not counted as refused
address filter	address ranges kept (<code>N_ADDR_RANGES</code> , error or match per range)	

The rest of this paper assumes the lite at the ports. Every axis below is the same for both monitors; the numbers are the lite's.



4 Identity space allocation

Three fields in every packet are designer-owned and fixed at build time:

Field	Bits	Set by	Meaning
unit_id	8	wrapper parameter UNIT_ID	which subsystem
agent_id	16	wrapper parameter AGENT_ID	which block inside it
channel_id	9	the monitor, from the transaction	which channel or AXI ID

The tools decode nothing from these fields; a host-side table maps the (unit_id, agent_id) tuple to a name. That makes the allocation yours. STREAM uses unit_id = 0x01 for the whole DMA and spends agent_id on blocks: 0x10 descriptor engine, 0x30 scheduler, 0x40 the widened scheduler set (projects/components/dma-ip/stream/rtl/macro/scheduler_group.sv).

Worked example, one level down. A subsystem with several internal buses can split unit_id into a subsystem nibble and a sub-bus nibble: unit_id = {sub_system[3:0], sub_bus[3:0]} gives sixteen subsystems each with up to sixteen tracked internal buses, and agent_id is still free to name the block on each bus. A host filter on unit_id[7:4] then isolates a subsystem, and on the full byte a single bus, with no change to the RTL. The cost is nothing: the fields exist in every packet regardless.




Identity allocation: one unit per subsystem, or a subsystem nibble and a sub-bus nibble

5 Where to insert monitoring

Placement	What it sees	Cost	Use when
Per port (the default; the bridge generator puts a lite monitor on every monitored port)	every transaction on that port, attributed by ID, with completion latency	one lite per port: about 1,100 LUTs / 1,030 FFs at defaults, 677 / 831 with the bridge's narrower parameters	you need to know WHICH port misbehaved

Placement	What it sees	Cost	Use when
		(monitor_characterization.md , axi_monitor_lite.md)	
Mid-fabric	the traffic crossing one internal link, through a pass-through observer with its own APB configuration (projects/components/misc/rtl/axi4_intf_master_observer.sv)	one observer per link	localizing violations the fabric itself introduces (ordering, ID collisions between ports)
Root of tree	the aggregate of everything below, once	one monitor	area-constrained; you only need “something is wrong”, not where

Per port and root of tree are the ends of one trade: resolution against area. A three-port read bridge with per-port lite monitors spends about 2,000 LUTs on monitors and about 1,400 on the shared group; the same bridge monitored once at its root spends 677 and 1,400, and can no longer tell the ports apart. With the full monitor the same per-port choice would cost about 21,000 LUTs, which is why per-port monitoring was not realistic before the lite and is the default with it.



Insertion points: per port, mid-fabric, root of tree

6 Timestamp policy

A 64-bit timestamp rides beside every packet (monbus_timestamp_t). Today it is the monbus group family’s local free-running counter (mon_time_out), distributed to the wrappers as i_mon_time, and the group writes its low 60 bits into the first beat of every three-beat record. Within one group every packet is on one clock; across groups there is no shared time base.

Two directions are open to the integrator:

- **Hybrid** {global_us[47:0], local_cyc[15:0]}: a system-wide microsecond count in the upper bits, the wrapper's own cycle count in the lower. Cross-subsystem correlation and per-wrapper cycle resolution then share the one field, and nothing in the packet layout moves.
- **External time source** (PTP or a chip-level counter): drive i_mon_time from it instead of the group's counter. The wrappers do not care where the value comes from.

Neither is prototyped; the hybrid form is the one this paper recommends when it is, because it needs no new field and no host-side re-basing.



Timestamp policy: the group's local counter today, the hybrid time base as the tweak


7 Drain path selection

The group has two sinks and a per-type steering mask ([monbus_group.md](#)):

Sink	Record	Reached by	Best for
Error FIFO	192-bit {timestamp, packet}, read back as three 64-bit beats over the AXI-Lite slave; irq_out while non-empty	cfg_<proto>_err_select[type] = 1	the few packets a handler must act on now: errors, timeouts
Write FIFO	the same record as three beats (raw) or one beat (compressed), flushed as AXI bursts into the ring	everything not dropped and not selected for the error FIFO	bulk trace: completions, latencies, thresholds, for offline analysis

Sink	Record	Reached by	Best for
	<code>[cfg_base_addr, cfg_limit_addr]</code> when <code>cfg_flush_watermark</code> beats are queued or the flush timeout elapses		

The choice is per packet type per protocol, at runtime. The common shape is errors and timeouts to the interrupt path and everything else to the trace ring; a debug session flips completions over to the error FIFO for a while and reads them one at a time.



Drain paths: error FIFO with interrupt, write FIFO into the memory ring


8 Packet-type filtering

Three layers, all runtime-programmable over the control APB:

1. **Type mask** at the group, `cfg_<proto>_pkt_mask[type]`: 1 drops the type. Also on every lite wrapper (`cfg_axi_pkt_mask`) so a dropped type never even leaves the port.
2. **Event mask** at the group, `cfg_<proto>_<class>_mask[event_code[3:0]]`: 1 drops one event within a type, for example `RESP_SLVERR` but not `RESP_DECERR`.
3. **Enable pins** on the wrappers, `cfg_error_enable`, `cfg_compl_enable`, `cfg_timeout_enable`, `cfg_threshold_enable`: a class that is off is not generated at all, so it costs no monbus bandwidth.

The congestion pitfall, and how the lite retires it. On the full monitor, completion and performance packets are the high-rate classes; with both enabled on a busy port the monbus

saturates and lower-priority packets are lost silently. The rule from [AXI_Monitor_Configuration_Guide.md](#) stands there: never enable `cfg_compl_enable` and `cfg_perf_enable` together. The lite has no performance class to collide with, holds a timeout or a latency event that loses the pick until it can go, and when its queue is genuinely full it counts the loss and reports the count in one `Error/EVENT_DROPPED` packet as soon as the bus is free. Nothing is lost silently; the consumer always knows how many events it did not see, which is the property a filtering strategy can be built on.




Packet-type filtering: enable pins, type mask, event mask, then drain steering; the lite counts what its queue could not take

9 Aggregation topology

Topology	Block	When
Tree of round-robin arbiters (default)	monbus_arbiter, any fan-in, ACK-mode grants	ports of similar rate; simplest, fair
Weighted	arbiter_wrr_pwm_monbus (weights per client, PWM flow control); arbiter_rr_pwm_monbus is the unweighted sibling	one port far busier than the others and you want its packets to win proportionally, not equally
Protocol-partitioned groups	one monbus_group_core per protocol family instead of the shared AXI / AXIS / CORE slots	the families need different rings, different interrupt handlers, or different flush policies

The group core is shared by every wrapper family (AXI4, AXI5, AXI-Lite, AXIS, Wishbone), so partitioning is a topology decision, not a new block.



Aggregation topologies: round-robin tree, weighted, protocol-partitioned

10 What it costs

Standalone, at default parameters, from the characterization sweep ([monitor/monitor_characterization.md](#)), and in place, from the bridge fixture with the bridge's parameters ([axi_monitor_lite.md](#)):

	Full monitor	Lite
One AXI4 read monitor at defaults (16 / 8 slots), LUTs / FFs	7,006 / 5,219	1,115 / 1,027
One read monitor in the bridge, LUTs / FFs	3,249 / 1,628	677 / 831
Three-port read bridge with group, LUTs	12,625	5,339
The unmonitored bridge, LUTs	826	826
Kintex-7 325T -2 at 6.667 ns, bridge WNS	+0.212 ns	+1.092 ns
Kintex-7 at 6.667 ns, standalone at defaults, register to register	-0.166 ns	+1.005 ns
A group (error FIFO 64 records, write	1,700 to 2,000 / 1,150 to 1,390	same, shared

Full monitor	Lite
--------------	------

FIFO 96 beats), LUTs
/ FFs

The lite bridge fixture meets 10 ns on the Artix-7 100T -1 since 2026-09-28 (+0.366 ns), after the group planner and the lite's event stage were re-pipelined (amba ISSUE-001, monitor-lite ISSUE-002). The per-variant characterization – every wrapper family, with and without clock gating, full against lite, on both parts, from one repeatable flow – is

[monitor/monitor_characterization.md](#). Its headline: at default parameters the full AXI monitor is about 7,000 LUTs and the lite about 1,100, and the lite meets 150 MHz on the Kintex-7 with 0.8 ns or more to spare where the 16-slot full monitor does not.

11 Validating a tweak in simulation

Three templates in the tree exercise exactly the knobs above:

- `val/amba/test_axi_mon_block_ready.py`: saturate a port and prove the monitor's admission contract, `block_ready` on the full monitor and the `refuse-count` identity on the lite.
- `val/amba/test_axi_monitor_soak.py::monitor_soak_monlite`: random traffic with backpressure, asserting that every generated event is delivered or counted as dropped. Change a mask or a drain and this is the identity that must still close.
- `formal/amba/monbus_group_core`: the routing rule, `monbus_ready`, the FIFO accounting and the legality of the flush burst, proved at the ports. A change to the group's filtering or drain logic re-runs against it.

The error-injection template the original outline named (`test_bridge_1x2_rd_monitor_error_inject.py`) is not in the tree; the bridge's generated monitor stress tests cover that ground.

12 Out of scope

Packet bit layout ([monitor_package_spec.md](#)), per-module ports and timing ([monitor/](#)), and test recipes (the test sources).