



RTL Design Sherpa

RTL Clock Domain Crossing Module Reference — Rev 0.90

October 2026

Table of Contents

1 Clock Domain Crossing (CDC).....	22
1.1 Choosing a technique.....	22
1.1.1 Quick reference.....	23
1.2 Reset Considerations.....	24
1.2.1 Why encoding decides reset behavior.....	24
1.2.2 Async FIFO reset behavior.....	28
1.2.3 Silicon evidence.....	29
1.2.4 Safe usage checklist for 2-phase.....	29
1.2.5 Recovering an existing design.....	29
1.3 cdc_synchronizer.....	30
1.4 cdc_open_loop.....	31
1.5 cdc_2_phase_handshake.....	32
1.5.1 Shared handshake interface.....	33
1.5.2 Theory of operation.....	33
1.5.3 Swapping between variants.....	34
1.6 cdc_4_phase_handshake.....	34
1.6.1 Theory of operation.....	35
1.7 Async FIFO Pointers (Gray and Johnson).....	35
1.7.1 Gray vs Johnson encoding.....	35
1.7.2 Sizing the depth.....	37
1.7.3 Walkthrough: is a non-power-of-2 FIFO worth it at 512 bits?.....	39
1.8 SDC constraints.....	43

1.9 Common mistakes.....	44
1.10 Verification status.....	45
1.11 Related Modules.....	46
1.12 Testing.....	47
1.13 Navigation.....	47
2 Binary-to-Gray converter (bin2gray.sv).....	48
2.1 Overview.....	48
2.2 Parameters.....	48
2.3 Ports.....	49
2.4 Functional Description.....	49
2.4.1 Mathematical definition.....	49
2.4.2 Why Gray code?.....	49
2.4.3 Gray code sequences.....	49
2.4.4 Core logic.....	50
2.4.5 Companion converter: gray2bin.....	51
2.5 Timing Characteristics.....	52
2.6 Usage Examples.....	53
2.7 Design Notes.....	53
2.8 Related Modules.....	53
2.9 Testing.....	54
2.9.1 Comprehensive Test Bench.....	54
2.9.2 Property-Based Verification.....	55
2.9.3 Coverage Model.....	57
2.9.4 Test files.....	57

2.10 Navigation.....	58
3 Gray-to-binary converter (gray2bin.sv).....	59
3.1 Overview.....	59
3.2 Parameters.....	59
3.3 Ports.....	59
3.4 Functional Description.....	59
3.4.1 Conversion algorithm.....	60
3.4.2 Worked conversions.....	60
3.4.3 Implementation.....	61
3.5 Timing Characteristics.....	62
3.6 Usage Examples.....	62
3.6.1 Asynchronous FIFO pointers.....	62
3.6.2 Clock domain crossing counter.....	62
3.6.3 Position encoding.....	63
3.7 Design Notes.....	63
3.7.1 Common mistakes.....	63
3.8 Related Modules.....	64
3.9 Testing.....	64
3.9.1 Exhaustive Testing.....	64
3.9.2 Property-Based Verification.....	64
3.9.3 Random Testing.....	65
3.9.4 Test files.....	65
3.10 Navigation.....	66
4 Johnson-to-binary converter (johnson2bin.sv).....	67

4.1 Overview.....	67
4.2 Parameters.....	67
4.3 Ports.....	67
4.4 Functional Description.....	68
4.4.1 The Johnson sequence.....	68
4.4.2 The algorithm.....	68
4.4.3 Three-part binary construction.....	69
4.4.4 Width calculations.....	69
4.4.5 Why this algorithm works.....	69
4.4.6 Worked conversions.....	70
4.4.7 Comparison with standard gray2bin.....	70
4.5 Timing Characteristics.....	71
4.6 Usage Examples.....	71
4.7 Design Notes.....	71
4.7.1 The clock and reset ports are unused.....	71
4.7.2 When to use Johnson vs. Gray.....	72
4.7.3 Resource planning.....	72
4.7.4 Verification challenges.....	72
4.7.5 Invalid Johnson sequences.....	72
4.7.6 Debug visibility.....	73
4.7.7 Common issues.....	73
4.7.8 Advanced topics.....	73
4.8 Related Modules.....	73
4.9 Testing.....	73

4.10 Navigation.....	74
5 Binary-Gray counter (counter_bingray.sv).....	75
5.1 Overview.....	75
5.2 Parameters.....	75
5.3 Ports.....	75
5.4 Functional Description.....	76
5.4.1 Gray code sequence (4-bit).....	76
5.4.2 Internal logic.....	77
5.5 Timing Characteristics.....	77
5.6 Waveforms.....	78
5.6.1 Scenario 1: Binary vs Gray Code Comparison.....	78
5.6.2 Scenario 2: Single-Bit Transitions (CDC Safety) KEY FEATURE.....	78
5.6.3 Scenario 3: Lookahead Signal (counter_bin_next).....	78
5.6.4 Scenario 4: Enable and Reset Control.....	79
5.7 Usage Examples.....	80
5.8 Design Notes.....	80
5.8.1 Common mistakes.....	81
5.9 Related Modules.....	81
5.10 Testing.....	81
5.10.1 Test scenarios.....	81
5.10.2 Coverage.....	81
5.10.3 Assertions.....	82
5.10.4 Test files.....	82
5.11 Navigation.....	82

6 Johnson counter (counter_johnson.sv).....	83
6.1 Overview.....	83
6.2 Parameters.....	83
6.3 Ports.....	83
6.4 Functional Description.....	84
6.4.1 State sequences.....	84
6.4.2 NOT self-starting – reset is mandatory.....	85
6.4.3 Symmetry.....	86
6.4.4 Single-bit transitions – but Johnson is not Gray.....	86
6.5 Timing Characteristics.....	86
6.6 Waveforms.....	87
6.6.1 Scenario 1: Walking Ones and Walking Zeros Pattern.....	87
6.6.2 Scenario 2: Single-Bit Transitions (CDC Safety) KEY FEATURE.....	87
6.6.3 Scenario 3: Enable Control.....	87
6.6.4 Scenario 4: Reset Behavior.....	88
6.7 Usage Examples.....	88
6.7.1 Example 1: 4-phase clock generator.....	88
6.7.2 Example 2: LED chaser.....	89
6.7.3 Example 3: state machine controller.....	89
6.7.4 Example 4: quadrature signal generation.....	90
6.7.5 Multi-phase clock for pipeline stages.....	90
6.7.6 Sequential control (memory refresh).....	90
6.7.7 Pattern generation for BIST.....	91
6.8 Design Notes.....	91

6.9 Related Modules.....	91
6.10 Testing.....	92
6.10.1 Functional tests.....	92
6.10.2 Coverage.....	92
6.10.3 Assertions.....	92
6.10.4 Test files.....	93
6.11 Navigation.....	93
7 Asynchronous FIFO (fifo_async.sv).....	94
7.1 Overview.....	94
7.2 Parameters.....	94
7.3 Ports.....	95
7.4 Functional Description.....	96
7.4.1 Clock domain crossing strategy.....	96
7.4.2 Core components.....	96
7.4.3 Gray code pointer system.....	96
7.4.4 Clock domain crossing.....	97
7.4.5 Memory organization.....	97
7.4.6 Full/empty detection.....	97
7.4.7 The power-of-2 requirement.....	98
7.5 Timing Characteristics.....	98
7.6 Waveforms.....	99
7.7 Usage Examples.....	100
7.8 Design Notes.....	100
7.8.1 Depth sizing.....	100

7.8.2 Clock relationship.....	100
7.8.3 Almost full/empty settings.....	100
7.8.4 Error checking.....	101
7.9 Related Modules.....	101
7.10 Testing.....	101
7.11 Navigation.....	101
8 GAXI asynchronous FIFO (gaxi_fifo_async.sv).....	102
8.1 Overview.....	102
8.2 Parameters.....	103
8.3 Ports.....	104
8.4 Functional Description.....	104
8.4.1 Why Gray code?.....	104
8.4.2 Dependencies.....	104
8.5 Timing Characteristics.....	107
8.6 Usage Examples.....	107
8.6.1 Example 1: basic CDC FIFO.....	107
8.6.2 Example 2: high-speed to low-speed CDC.....	108
8.7 Design Notes.....	108
8.7.1 Depth sizing for clock ratio.....	108
8.7.2 Reset synchronization.....	108
8.7.3 Metastability protection.....	109
8.7.4 Error checking.....	109
8.7.5 Common mistakes.....	110
8.8 Related Modules.....	110

8.9 Testing.....	110
8.10 References.....	110
8.11 Navigation.....	111
9 gaxi_skid_buffer_async.....	112
9.1 Overview.....	112
9.1.1 Key Features.....	112
9.2 Parameters.....	112
9.3 Ports.....	113
9.4 Functional Description.....	114
9.4.1 Internal Components.....	114
9.5 Timing Characteristics.....	114
9.5.1 Latency Components.....	114
9.5.2 Throughput.....	115
9.6 Usage Examples.....	115
9.6.1 Example 1: Basic Async Skid Buffer.....	115
9.6.2 Example 2: Pipeline Stage with CDC.....	115
9.7 Design Notes.....	116
9.7.1 When to Use This Module.....	116
9.7.2 Depth Selection.....	116
9.7.3 Reset Synchronization.....	116
9.7.4 Resource Utilization.....	116
9.7.5 Common Issues.....	117
9.8 Related Modules.....	117
9.9 Testing.....	118

9.10 References.....	118
9.11 Navigation.....	118
10 glitch_free_n_dff_arn.....	119
10.1 Overview.....	119
10.2 Parameters.....	119
10.3 Ports.....	119
10.4 Functional Description.....	119
10.4.1 The Metastability Problem.....	119
10.4.2 Packed Array Structure.....	120
10.4.3 Shift Register Behavior.....	121
10.4.4 Reset Behavior.....	121
10.4.5 MTBF (Mean Time Between Failures) Analysis.....	121
10.5 Timing Characteristics.....	122
10.5.1 Synchronizer Latency.....	122
10.5.2 Performance Impact.....	122
10.6 Usage Examples.....	123
10.6.1 Reset Synchronizer Pattern.....	123
10.6.2 Enable Signal Synchronization.....	123
10.6.3 Status Flag Collection.....	123
10.7 Design Notes.....	123
10.7.1 Appropriate Use Cases.....	123
10.7.2 Inappropriate Use Cases.....	123
10.7.3 Multi-Bit Considerations.....	124
10.7.4 Synthesis Considerations.....	124

10.7.5 Debug Features.....	124
10.8 Related Modules.....	125
10.9 Testing.....	125
10.9.1 Verification Strategies.....	125
10.9.2 Test Suite.....	126
10.10 Navigation.....	126
11 reset_sync.....	127
11.1 Overview.....	127
11.2 Parameters.....	127
11.3 Ports.....	128
11.3.1 Inputs.....	128
11.3.2 Outputs.....	128
11.4 Functional Description.....	129
11.4.1 The Reset Problem.....	129
11.4.2 Asynchronous Assert, Synchronous Deassert.....	129
11.4.3 Synchronization Register Chain.....	129
11.4.4 Reset Logic.....	129
11.4.5 Output Assignment.....	129
11.5 Timing Characteristics.....	130
11.5.1 Reset Assertion (Asynchronous).....	130
11.5.2 Reset Deassertion (Synchronous, N=3).....	130
11.5.3 State Sequence for N=3.....	130
11.6 Usage Examples.....	131
11.6.1 Choosing N Value.....	131

11.6.2 Global Reset Distribution.....	132
11.6.3 Clock Domain Crossing.....	132
11.6.4 Power-On Reset.....	133
11.6.5 Reset Controller Integration.....	133
11.7 Design Notes.....	133
11.7.1 1. Why the chain looks inverted.....	133
11.7.2 2. Metastability Resolution.....	134
11.7.3 3. Initialization.....	134
11.7.4 4. Vendor attributes.....	134
11.7.5 5. Parameterization Benefits.....	134
11.7.6 Reset Tree Planning.....	134
11.7.7 Clock Quality Requirements.....	134
11.7.8 Reset Source Management.....	135
11.7.9 Incorrect Sensitivity List.....	135
11.7.10 Wrong Reset Polarity.....	135
11.7.11 Clock Domain Issues.....	135
11.8 Related Modules.....	135
11.9 Testing.....	136
11.9.1 Test Coverage.....	136
11.9.2 Test Scenarios.....	136
11.9.3 Testbench Scenarios.....	136
11.9.4 Assertions.....	136
11.9.5 Coverage Points.....	136
11.9.6 Running Tests.....	136

11.9.7 Test Results.....	137
11.10 Navigation.....	137
12 clock_pulse.....	138
12.1 Overview.....	138
12.2 Parameters.....	138
12.3 Ports.....	138
12.4 Functional Description.....	139
12.4.1 Internal Counter.....	139
12.4.2 Core Logic.....	139
12.4.3 Operation Principles.....	139
12.5 Related Modules.....	140
12.6 Timing Characteristics.....	140
12.6.1 Timing Characteristics.....	140
12.6.2 Basic Operation (WIDTH=4).....	140
12.6.3 Reset Behavior.....	140
12.6.4 Different WIDTH Values.....	141
12.7 Usage Examples.....	141
12.8 Design Notes.....	141
12.8.1 Resource Utilization.....	141
12.8.2 Synthesis Considerations.....	141
12.8.3 Design Considerations.....	142
12.8.4 Common Applications.....	143
12.9 Testing.....	143
12.9.1 Comprehensive Test Bench.....	143

12.9.2 Coverage Model.....	147
12.9.3 Formal Properties.....	148
12.10 Navigation.....	149
13 apb4_slave_cdc.....	150
13.1 Overview.....	150
13.1.1 Key Features.....	150
13.2 Parameters.....	150
13.3 Ports.....	151
13.4 Functional Description.....	152
13.4.1 Clock Domains.....	152
13.4.2 CDC Structure.....	153
13.4.3 Reset Behavior.....	154
13.4.4 Why Not the Previous 2-Phase Handshake.....	154
13.5 Design Notes.....	155
13.6 Timing Characteristics.....	155
13.6.1 Maximum Clock Ratio.....	155
13.6.2 Timing Constraints.....	155
13.7 Waveforms.....	155
13.7.1 Scenario 1: Write Transaction with CDC.....	156
13.7.2 Scenario 2: Read Transaction with CDC.....	156
13.7.3 Scenario 3: Back-to-Back Writes with CDC.....	157
13.7.4 Scenario 4: Back-to-Back Reads with CDC.....	157
13.7.5 Scenario 5: Write-to-Read Transition with CDC.....	157
13.7.6 Scenario 6: Read-to-Write Transition with CDC.....	158

13.7.7 Scenario 7: Error Response with CDC.....	158
13.8 Usage Examples.....	158
13.9 Related Modules.....	159
13.10 Testing.....	160
13.11 References.....	160
13.12 Navigation.....	160
14 apb4_slave_cdc_cg.....	161
14.1 Overview.....	161
14.2 Parameters.....	162
14.2.1 Derived Parameters (do not override).....	162
14.3 Ports.....	162
14.3.1 Clock-gating status outputs.....	163
14.4 Functional Description.....	164
14.5 Timing Characteristics.....	165
14.6 Usage Examples.....	165
14.7 Design Notes.....	166
14.8 Related Modules.....	166
14.9 Testing.....	167
14.10 Navigation.....	167
15 APB5 Slave CDC.....	168
15.1 Overview.....	168
15.1.1 Key Features.....	168
15.2 Parameters.....	168
15.2.1 Derived Parameters (do not override).....	169

15.3 Ports.....	170
15.3.1 Clock and Reset.....	170
15.3.2 APB5 Slave Interface.....	170
15.3.3 Backend Interface.....	170
15.4 Functional Description.....	171
15.4.1 CDC FIFO Pointer Encoding.....	172
15.4.2 Reset Synchronization.....	173
15.5 Timing Characteristics.....	174
15.5.1 Latency.....	174
15.5.2 Timing Considerations.....	174
15.6 Usage Examples.....	174
15.7 Design Notes.....	175
15.7.1 FIFO Depth Sizing.....	175
15.8 Related Modules.....	175
15.9 Testing.....	175
15.10 Navigation.....	176
16 APB5 Slave CDC (Clock-Gated).....	177
16.1 Overview.....	177
16.1.1 Key Features.....	177
16.2 Parameters.....	177
16.2.1 Derived Parameters (do not override).....	178
16.3 Ports.....	179
16.3.1 Clock and Reset.....	179
16.3.2 Clock Gating Configuration.....	179

16.3.3 APB5 Slave Interface.....	179
16.3.4 Backend Interface.....	179
16.3.5 Status Outputs.....	179
16.4 Functional Description.....	180
16.5 Timing Characteristics.....	181
16.6 Usage Examples.....	182
16.7 Design Notes.....	183
16.7.1 Power and Latency Trade-offs.....	183
16.7.2 Reset Considerations.....	184
16.8 Related Modules.....	185
16.9 Testing.....	185
16.10 Navigation.....	185

List of Tables

Table 1: CDC building blocks at a glance.....	22
Table 2: CDC technique quick reference by scenario.....	23
Table 3: How pointer encoding determines reset behavior.....	24
Table 4: 2-phase handshake reset domains per flop.....	25
Table 5: Options for recovering a design using 2-phase across independent resets	29
Table 6: cdc_synchronizer parameters.....	30
Table 7: cdc_open_loop parameters.....	31
Table 8: cdc_open_loop ports.....	32
Table 9: cdc_2_phase_handshake parameters.....	33
Table 10: Shared handshake interface ports.....	33
Table 11: Consequences of swapping handshake variants.....	34
Table 12: cdc_4_phase_handshake parameters.....	34
Table 13: Async FIFO pointer encodings.....	36
Table 14: 3-bit Gray vs the equivalent 4-bit Johnson counter, full cycle. “flips bit#” is a bit POSITION; “Ham. dist” is the number of bits that changed, which is 1 everywhere including the wrap.....	36
Table 15: Pointer width by encoding – Gray is logarithmic, Johnson is linear.....	37
Table 16: Worked depth-sizing example.....	38
Table 17: Nominal storage overhead of Gray depth rounding by data width.....	39
Table 18: 512-bit walkthrough: depth after applying the encoding constraint.....	39
Table 19: 512-bit walkthrough: nominal storage delta.....	39
Table 20: 512-bit walkthrough: pointer flop cost.....	40
Table 21: 512-bit walkthrough: whether the saving is real, by memory style.....	40
Table 22: Depth-36 case: storage delta when Gray must round 36 up to 64.....	41
Table 23: Depth-36 case: pointer flop cost.....	42
Table 24: Depth-36 case: ASIC memory styles, all of which take the word count as an argument.....	42
Table 25: CDC module reference.....	46
Table 26: bin2gray parameters.....	48
Table 27: bin2gray ports.....	49
Table 28: bin2gray resource usage by width (typical FPGA).....	52
Table 29: gray2bin parameters.....	59
Table 30: gray2bin ports.....	59
Table 31: gray2bin delay by width.....	62
Table 32: johnson2bin parameters.....	67

Table 33: johnson2bin ports.....	67
Table 34: counter_bingray parameters.....	75
Table 35: counter_bingray ports.....	75
Table 36: counter_johnson parameters.....	83
Table 37: counter_johnson ports.....	83
Table 38: fifo_async parameters.....	94
Table 39: fifo_async ports.....	95
Table 40: gaxi_fifo_async parameters.....	103
Table 41: gaxi_fifo_async ports.....	104
Table 42: gaxi_fifo_async latency by clock ratio.....	107
Table 43: Synchronizer stage guidance.....	109
Table 44: reset_sync parameters.....	127
Table 45: reset_sync state sequence for N=3 at default polarities.....	130

List of Waveforms

Waveform 1.1: 2-Phase Normal Transfer.....	25
Waveform 1.2: 2-Phase Hazard - Destination Reset Alone.....	26
Waveform 1.3: 4-Phase Normal Transfer.....	27
Waveform 1.4: 4-Phase One-Sided Reset Is Safe.....	27

1 Clock Domain Crossing (CDC)

RTL: rtl/cdc/ **Filelists:** rtl/amba/filelists/cdc_*.f **Status:** Production Ready **for common-clock, common-reset verification.** Reset behavior across independent domains is argued from the RTL and, for the 2-phase hazard, confirmed on silicon – but it is **not** covered by the formal proofs, which drive both domains from one clock and one reset. See [Verification status](#) before relying on that badge.

This page is the single reference for clock domain crossing in this repository: every CDC building block, how to choose between them, and the reset behavior that distinguishes them. That last part is the property most likely to bite you, and the least likely to show up in simulation.

CDC building blocks at a glance

Module	Category	One-line
cdc_synchronizer	Multi-flop sync	N-stage synchronizer for quasi-static signals
cdc_open_loop	Open loop	Source holds data + valid, no acknowledge
cdc_2_phase_handshake	Closed loop	Toggle (NRZ) valid/ready handshake, faster
cdc_4_phase_handshake	Closed loop	Level (RZ) valid/ready handshake, classic
fifo_async / gaxi_fifo_async	Pointer sync	Streaming data, buffered

1.1 Choosing a technique

Start with the shape of the traffic, not the protocol:

Need to cross clock domains with multi-bit data?

```
|
+-- Is data streaming / continuous?
|   |
|   +-- YES --> Async FIFO (fifo_async or gaxi_fifo_async)
|               Sustained 1-per-clock throughput, buffered
|   |
|   +-- NO, occasional transfers -->
|       |
|       +-- Does the source need backpressure (flow control)?
|           |
|           +-- YES --> Closed-loop handshake
```

```

|
|      (cdc_2_phase_handshake or
cdc_4_phase_handshake)
|
|      +-- NO, source can guarantee hold time -->
|      |
|      +-- Open-loop (cdc_open_loop)
|          Simplest, lowest area, no ack needed
|
Single-bit signal?
|
+-- Level / quasi-static --> cdc_synchronizer or glitch_free_n_dff_arn
+-- Single-cycle pulse --> sync_pulse
+-- Reset signal --> reset_sync

```

Then check the **reset rule** before committing to a handshake. It eliminates one of the two handshake variants outright in a common situation, and it is not visible from the port list.

1.1.1 Quick reference

CDC technique quick reference by scenario

Scenario	Module	Why
Config register update (CPU -> peripheral)	cdc_open_loop	Infrequent, no backpressure needed
Interrupt signal (1 bit, pulse)	sync_pulse	Single-cycle event
Status register read (slow -> fast)	cdc_synchronizer	Quasi-static level
APB slave in different clock domain	apb4_slave_cdc	Full protocol CDC
Monitor packet crossing	cdc_4_phase_handshake	Occasional, needs flow control
DMA data stream	fifo_async	High throughput, continuous
AXI channel crossing	gaxi_skid_buffer_async	Pipelined, GAXI protocol
Reset distribution	reset_sync	Async assert, sync deassert

1.2 Reset Considerations

DESIGN RULE: If the two clock domains can be reset **independently**, do not use `cdc_2_phase_handshake`. Use `cdc_4_phase_handshake` or an async FIFO. A toggle handshake stores transfer state as **parity**, and parity cannot survive a one-sided reset.

“Independently” covers more cases than people expect: a soft reset that clears a datapath but not a register block, a per-block reset, separate power domains, or a reset synchronizer whose deassertion is gated differently on each side. If any of those can reach one domain alone, the rule applies.

1.2.1 Why encoding decides reset behavior

How pointer encoding determines reset behavior

	2-phase	4-phase	Async FIFO
Transfer state	Parity (toggle) – relative	Level – absolute	Pointer position – absolute
Idle has a value?	No	Yes (req=0)	Yes (pointers equal)
One-sided reset from idle	Fabricates a transfer	Nothing happens	Read-side reset: replays consumed entries (benign only if pointers never moved). Write-side reset: fabricates occupancy
One-sided reset mid-transfer	Duplicate or lost transfer	Duplicate or dropped transfer	Entry may be re-read or dropped

The 2-phase protocol signals “a transfer happened” with a *transition*. The receiver has no absolute reference – it only knows whether the synchronized toggle differs from its previous sample:

```
assign w_req_event = w_req_sync ^ r_req_sync_d;    //  
cdc_2_phase_handshake.sv:261
```

That XOR is the entire protocol state. Each side resets its own flops:


2-phase handshake reset domains per flop

Flop	Domain	Reset by	After reset
r_req_tog	source	rst_src_n	1'b0 (:190)
r_ack_syn c, r_ack_syn c_d	source	rst_src_n	1'b0 (:173-174)
r_ack_tog	destination	rst_dst_n	1'b0
r_req_syn c, r_req_syn c_d	destination	rst_dst_n	1'b0 (:252-253)

Reset **both** domains together and everything lands at 0 – parity agrees, no event. Reset **one** and the two ends land at opposite parity, and the XOR reports an edge no source transfer produced.

Because parity is relative, there is no toggle value that means “idle”. Whatever r_req_tog happens to hold, a freshly cleared destination disagrees with it half the time. In 4-phase, req = 0 is idle, absolutely: a reset destination observing req = 0 correctly concludes nothing is pending.

1.2.1.1 Waveform 1.1: 2-Phase Normal Transfer




2-phase normal transfer

Source: [cdc2_normal_transfer.json](#)

One src_valid toggles r_req_tog 0 -> 1. Three synchronizer stages later the destination sees the new value, w_req_event pulses for one cycle, dst_valid asserts. Parity now agrees at 1 on both ends.

1.2.1.2 Waveform 1.2: 2-Phase Hazard - Destination Reset Alone



2-phase asymmetric reset hazard

Source: [cdc2_asymmetric_reset_hazard.json](#)


`rst_dst_n` asserts while `rst_src_n` stays high:

1. The destination clears `r_req_sync` and `r_req_sync_d` to 0.
2. The source is untouched – `r_req_tog` is still 1.
3. On release the sync chain refills with 1 while `r_req_sync_d` is the freshly cleared 0.
4. $w_req_event = 1 \wedge 0 = 1$ – a **phantom transfer**. `dst_valid` asserts with stale `r_src_data_hold`, for a transfer the source never sent.

Resetting the source alone abandons an in-flight transfer too: the source FSM returns to `S_IDLE` from `S_WAIT_ACK` (: 188 - 190). It does **not** clear `r_ack_tog` – that flop lives in the destination domain and is reset only by `rst_dst_n`, as the flop table above records. The spurious-ack mechanism belongs to the destination-reset case, where `r_ack_tog` clears while the source is still waiting. Either way a one-sided reset can **lose** a transfer as well as fabricate one.

Scope of the damage. The handshake itself re-synchronizes: once `r_req_sync_d` catches up to `r_req_tog`, parity agrees again. What is permanent is the **extra item already delivered downstream**. A consumer that samples the latest value self-heals on the next real transfer; a consumer that counts transfers, pushes into a FIFO, or pairs responses to requests positionally is corrupted for the rest of the session.

1.2.1.3 Waveform 1.3: 4-Phase Normal Transfer




4-phase normal transfer

Source: [cdc4_normal_transfer.json](#)

Four crossings: `r_req_src` rises, `r_ack_dst` rises, `r_req_src` falls, `r_ack_dst` falls. Both ends finish at the 0 idle level.

1.2.1.4 Waveform 1.4: 4-Phase One-Sided Reset Is Safe



4-phase one-sided reset is safe

Source: [cdc4_asymmetric_reset_safe.json](#)

`rst_dst_n` asserts mid-transfer while the source runs:

1. The destination FSM returns to D_IDLE and `r_ack_dst` clears to 0.
2. The source is unaffected, still holding `r_req_src` = 1 with data stable.
3. On release the destination re-observes `w_req_sync` = 1, treats it as a fresh request, and redoes the transfer.
4. The source sees the ack and completes normally.

Worst case is one transfer **repeated or dropped**, and only for a transfer genuinely in flight. If the link is idle when the reset lands, the destination sees `req` = 0, reads it as “nothing pending”, and does nothing.

Caveat – duplicate delivery. Step 3 re-delivers a transfer the destination may already have accepted. If the consumer is not idempotent (a FIFO push, a counter increment, a command queue), hold it in reset alongside the destination FSM or drain it afterwards. Level encoding guarantees you never get a transfer the source did not send; it does not guarantee exactly-once delivery.

1.2.2 Async FIFO reset behavior

Gray/Johnson pointers are absolute positions, and each domain resets its own pointer **and its crossed copy of the remote pointer** from the local reset:

```
rd_ptr_gray_cross_inst -> .rst_n (axi_wr_aresetn) // synchronizer,
local reset
wr_ptr_gray_cross_inst -> .rst_n (axi_rd_aresetn) // synchronizer,
local reset
```

Only the synchronizers take a reset. The decode stage beside each one has none to take: with the default `USE_JOHNSON=0` it is a gray2bin, whose entire port list is `.gray` and `.binary`. (With `USE_JOHNSON=1` it is a johnson2bin, which does carry `clk/rst_n` ports – but leaves them unused.) The local-reset property below rests on the synchronizers alone; the decode is combinational and simply follows whatever the synchronizer holds.

There is no parity state to invert, so the FIFO cannot fabricate the way a 2-phase handshake does – but “one-sided reset is safe” is too strong, and the asymmetry matters before you rely on it.

A read-side reset is benign only when the pointers have never moved. `rd_ptr` clears to 0, and the read domain’s copy of the write pointer (`wr_ptr_gray_cross_inst`) is cleared with it – but that copy is a LIVE synchronizer, and within `N_FLOP_CROSS` rd clocks it re-samples the write side’s untouched pointer. If the link has seen traffic – even fully consumed traffic, both pointers resting at some `K` – the settled state is `rd_ptr = 0` against `wr_ptr_sync = K`: the read side believes the FIFO holds `K` entries, and the `K` words it already consumed are **replayed**. Empty-flag “idle” is not enough; the pointers must be at zero.

A write-side reset from idle is not. The read domain’s copy of the write pointer (`wr_ptr_gray_cross_inst` in `fifo_async.sv`) is reset by **`rd_rst_n`**, not **`wr_rst_n`** – a write-side reset does not clear it. So with both pointers idle at some value `K`:

```
wr_rst_n pulsed:   wr_ptr      K -> 0
                   rd_ptr      K      (untouched)
                   wr_ptr_sync K -> 0 over N_FLOP_CROSS rd clocks

read side sees:   rd_ptr (K) != wr_ptr_sync (0)  => NOT EMPTY
                  count = 0 - K mod 2^(AW+1)    => 27 at K=5, AW=4
```

The read side now believes the FIFO holds more entries than it has slots, and pops words that were never written. That is fabrication, and it is worse than losing the writes.

The rule is the same either way – **quiesce the link before resetting one side** – but the reason is not symmetric, and resetting only the write domain is the dangerous direction.

This covers reset ASSERTION, not reset RELEASE. Each domain's reset must still be released synchronously *within that domain* – which is what `reset_sync` (async assert, sync deassert) is for. A raw asynchronous release can violate recovery/removal timing on the pointer flops, and if the two domains release at sufficiently different times the FIFO can begin operating while one side's synchronized copy of the remote pointer is still catching up. That window is benign for an idle FIFO (it reads empty) but not if traffic starts immediately. Hold traffic off until both domains are out of reset, or gate the producer on a released-and-settled indication.

1.2.3 Silicon evidence

This is not theoretical. On the Nexys A7 DDR2 characterization board, an APB CDC using the 2-phase handshake had its core-side reset pulsed by `CTRL.soft_reset` while the APB side stayed up. One phantom transfer entered a slave FSM that pairs responses to requests positionally (`apb4_slave.sv` returns whatever response sits at the head of its skid buffer), and **every register read from then on returned the previous register's value** – lagged by ~3 transactions, permanently, until reprogramming. Writes were unaffected, so it presented as a readback bug rather than a CDC bug. I've debugged enough of these to know that disguise is what makes them expensive.

Replacing that handshake with `gaxi_fifo_async` fixed it: reads became stable across `soft_reset`, verified on silicon.

1.2.4 Safe usage checklist for 2-phase

This is a hard stop, not a scorecard. Every condition must hold. Failing any one of them disqualifies the module – there is no partial credit and no mitigation short of changing the crossing or the reset topology.

Use `cdc_2_phase_handshake` only when **all** hold:

- ☐ `rst_src_n` and `rst_dst_n` assert and release together, **or** both derive from one reset no subsystem can pulse independently
- ☐ No soft reset, per-block reset, or power-domain reset touches one side alone
- ☐ The consumer correlates responses to requests by tag, not by position

If any box is unchecked, choose a level-encoded or pointer-encoded crossing.

1.2.5 Recovering an existing design

Options for recovering a design using 2-phase across independent resets

Option	Effort	Notes
Move both domains onto a common reset	Low	Best when the split was accidental
Swap to	Low	Pin-compatible; level

Option	Effort	Notes
cdc_4_phase_handshake		encoding self-recovers
Swap to gaxi_fifo_async	Medium	Absolute pointers fail less catastrophically than a stuck toggle, adds buffering – the quiesce-before-reset rule above still applies

1.3 cdc_synchronizer

Multi-flop synchronizer for **quasi-static** signals. Not a handshake – there is no flow control and no guarantee the destination sees every value.

```

module cdc_synchronizer #(
    parameter int WIDTH      = 1,    // Bus width to synchronize
    parameter int FLOP_COUNT = 3    // Number of synchronizer stages
(2-5)
) (
    input   logic          clk,      // Destination clock
    input   logic          rst_n,    // Asynchronous active-low
    reset
    input   logic [WIDTH-1:0] async_in, // Asynchronous input from
    source domain
    output logic [WIDTH-1:0] sync_out // Synchronized output in
    destination domain
);

```

cdc_synchronizer parameters

Parameter	Default	Description
WIDTH	1	Number of bits to synchronize (1+)
FLOP_COUNT	3	Synchronizer chain depth (2-5)

Usage rules

- Input must be **quasi-static**: it changes rarely enough, relative to the destination clock, that missing a transition is acceptable for the application. This is a statistical property, not a hard timing check the synchronizer enforces – it will sample whatever is present. As a sizing rule of thumb,

allow at least $\text{FLOP_COUNT} + 2$ destination clocks between transitions; change faster than that and the destination can miss transitions entirely or sample mid-flight. ($\text{FLOP_COUNT} + 2$ is the figure glitch_free_n_dff_arn's own header states, and `cdc_synchronizer` is a thin wrapper over it. This page previously said $+ 1$; neither is enforced by logic, but the two should not disagree, and the safer bound wins.)

- For multi-bit buses, all bits must change simultaneously **or** be Gray coded.
- $\text{FLOP_COUNT}=2$ for relaxed MTBF, 3 for production (default)
- Do **not** use for pulses (`sync_pulse`) or streaming data (`fifo_async`)

Gray-coded input must be REGISTERED in the source domain before it crosses. The single-bit-change guarantee is a property of a registered counter sequence, not of the encoding in the abstract. Feeding this module a Gray value straight out of combinational logic re-introduces exactly the hazard Gray coding exists to remove: the bits arrive with different delays, and the synchronizer can latch a transient that was never a real state. Register the Gray value, then cross it. The same applies to a Gray value that can *jump* rather than increment by one – multiple bits change and no encoding saves you.

1.4 `cdc_open_loop`

Source holds `valid + data` stable for a stretch window long enough that the destination cannot miss it. No acknowledge, so the lowest area and latency of the data-carrying crossings – at the cost of no backpressure.

cdc_open_loop parameters

Parameter	Default	Description
<code>DATA_WIDTH</code>	8	Width of the data bus
<code>STRETCH_CYCLES</code>	8	Source clocks to hold <code>valid+data</code> stable (<code>AUTO_STRETCH=0</code>)
<code>SYNC_STAGES</code>	2	Destination synchronizer depth (2-4)
<code>AUTO_STRETCH</code>	0	0 = use <code>STRETCH_CYCLES</code> ; 1 = compute from clock frequencies

Parameter	Default	Description
SRC_CLK_HZ	25_000_000	Source clock frequency (AUTO_STRETCH=1)
DST_CLK_HZ	100_000_000	Destination clock frequency (AUTO_STRETCH=1)

cdc_open_loop ports

Port	Direction	Width	Description
clk_src / rst_src_n	input	1	Source domain clock / async reset
src_valid	input	1	Single-cycle pulse: data is valid
src_data	input	DATA_WIDT H	Data to transfer
src_busy	output	1	High during stretch countdown
clk_dst / rst_dst_n	input	1	Destination domain clock / async reset
dst_valid	output	1	Single-cycle pulse: data latched
dst_data	output	DATA_WIDT H	Latched data, stable until next dst_valid

The failure mode to watch: sending a new transfer before the previous one is sampled. Without an acknowledge the source cannot know the destination latched the data. Respect `src_busy`, or space transfers by at least `STRETCH_CYCLES` source clocks (the busy window; the sizing rule makes it $\geq \text{SYNC_STAGES} + 1$ destination clocks).

1.5 `cdc_2_phase_handshake`

Toggle (NRZ) closed-loop handshake. Two synchronizer crossings per transfer instead of four.

Read [Reset Considerations](#) before choosing this module. It is pin-compatible with the 4-phase variant but not behaviorally equivalent under independent resets.

cdc_2_phase_handshake parameters

Parameter	Default	Description
DATA_WIDTH	8	Width of the data bus (1+)
SYNC_STAGES	3	Synchronizer depth for req/ack (2 or 3)
TIMEOUT_CYCLES	0	0 = disabled; >0 asserts src_timeout after stall

1.5.1 Shared handshake interface

Both handshake modules expose an identical port list, which is why swapping them is a one-word edit:

Shared handshake interface ports

Port	Direction	Width	Description
clk_src / rst_src_n	input	1	Source domain clock / async reset
src_valid	input	1	Source indicates data valid
src_ready	output	1	Handshake ready (asserted when idle)
src_data	input	DATA_WIDTH	Data from source domain
src_timeout	output	1	Stall timeout (when TIMEOUT_CYCLES > 0)
clk_dst / rst_dst_n	input	1	Destination domain clock / async reset
dst_valid	output	1	Data valid to receiver
dst_ready	input	1	Receiver ready
dst_data	output	DATA_WIDTH	Data transferred to destination domain

1.5.2 Theory of operation

Source FSM

S_IDLE src_ready=1. On src_valid: latch data, toggle r_req_tog, drop src_ready, go to S_WAIT_ACK.
S_WAIT_ACK Wait for ack edge (w_ack_event). On edge: src_ready=1, S_IDLE.

Destination FSM

D_IDLE On req edge (w_req_event): copy data into r_dst_data, raise dst_valid, go to D_WAIT_READY.
D_WAIT_READY Hold dst_valid. On dst_ready: drop dst_valid, toggle r_ack_tog, return to D_IDLE.

Edge detection

event = current_sync_output ^ previous_sync_output

1.5.3 Swapping between variants

Consequences of swapping handshake variants

Direction	Consequence
4-phase -> 2-phase	Gains throughput. Silently forfeits independent-reset tolerance. Only safe when both domains share one reset.
2-phase -> 4-phase	Costs roughly two extra synchronizer crossings (~2 x SYNC_STAGES clocks, so ~6 at the default of 3), split between the two domains. Always safe.

1.6 cdc_4_phase_handshake

Level (RZ) closed-loop handshake: req rises, ack rises, req falls, ack falls. Slower than 2-phase, and tolerant of independent domain resets.

cdc_4_phase_handshake parameters

Parameter	Default	Description
DATA_WIDTH	8	Width of the data bus (1 to 1024+)
SYNC_STAGES	3	Synchronizer depth for req/ack (2 or 3)
TIMEOUT_CYCLES	0	0 = disabled; >0 asserts src_timeout after stall

Ports are identical to the [shared handshake interface](#).

1.6.1 Theory of operation

Source FSM

```
S_IDLE          On src_valid: capture data, r_req_src=1, S_WAIT_ACK.
S_WAIT_ACK      On synchronized ack: r_req_src=0, S_WAIT_ACK_CLR.
S_WAIT_ACK_CLR  On ack cleared: src_ready=1, S_IDLE.
```

Destination FSM

```
D_IDLE          On synchronized req: latch data, dst_valid=1,
D_WAIT_READY.
D_WAIT_READY    On dst_ready: r_ack_dst=1, dst_valid=0,
D_WAIT_REQ_CLR.
D_WAIT_REQ_CLR  On req cleared: r_ack_dst=0, D_IDLE.
```

Data stability guarantee: data is captured at REQ assertion and held until ACK completion, so it is stable for the entire crossing regardless of clock ratio.

Latency: four synchronizer crossings per transfer against two for 2-phase. At the default SYNC_STAGES=3 that is roughly 12 clocks of round trip against 6 – counted in the clock domain each crossing lands in, so the wall-clock cost depends on the src:dst frequency ratio. Treat these as order-of-magnitude, not budgetable: measure your own configuration.

1.7 Async FIFO Pointers (Gray and Johnson)

For streaming data, synchronize the read and write **pointers** of a dual-port memory rather than the data. Pointers use an encoding where only one bit changes per increment, so any sample during a transition yields either the old or the new value – never a corrupted intermediate. The memory itself needs no CDC.

```
Write Domain:                                     Read Domain:
  wr_ptr (binary) ---> bin2gray --->             gray2bin ---> compare with
rd_ptr                                             [sync N stages]             for empty
detection
```

```
gray2bin <--- [sync N stages] <--- bin2gray <--- rd_ptr (binary)
compare with wr_ptr for full detection
```

1.7.1 Gray vs Johnson encoding

Both `fifo_async` (rtl/cdc) and `gaxi_fifo_async` (rtl/cdc) select the encoding with a `USE_JOHNSON` parameter:

Async FIFO pointer encodings

USE_JOHNSON	Encoding	Pointer width	Converter	Legal DEPTH
0 (default)	Gray	$\log_2(\text{DEPTH}) + 1$	gray2bin (combinational)	power of 2 only
1	Johnson	DEPTH	johnson2bin (combinational)	any depth, odd included

An illegal combination (Gray with a non-power-of-2 depth) fails at **elaboration** with an explicit \$error, not at runtime.

1.7.1.1 State walk: why the pointers are different widths

Both encodings guarantee the same CDC-critical property – exactly one bit changes per increment – but they buy it very differently. Walking a 3-bit Gray counter through its full cycle and putting the equivalent Johnson counter beside it shows why.

A 3-bit Gray counter has $2^{*3} = 8$ states. To reach the same 8 states a Johnson counter needs $2N = 8$, i.e. **4 bits** – it only ever produces $2 \times \text{WIDTH}$ states, not 2^{*}WIDTH . The Johnson sequence fills with ones from the LSB, then flushes them back out with zeros; that is the whole counter (counter_johnson.sv:32: {counter_gray[WIDTH-2:0], ~counter_gray[WIDTH-1]}).

The two “which bit” columns give the **bit POSITION that flipped** (0 = LSB). They are not counts. The Hamming-distance columns state the count explicitly, and it is 1 on every row – that is the property being demonstrated.

3-bit Gray vs the equivalent 4-bit Johnson counter, full cycle. “flips bit#” is a bit POSITION; “Ham. dist” is the number of bits that changed, which is 1 everywhere including the wrap.

Step	Gray (3-bit)	flips bit#	Ham. dist	Johnson (4-bit)	flips bit#	Ham. dist	Johnson phase
0	000	0	1	0000	0	1	ones in
1	001	1	1	0001	1	1	ones in
2	011	0	1	0011	2	1	ones in
3	010	2	1	0111	3	1	ones in
4	110	0	1	1111	0	1	zeros in
5	111	1	1	1110	1	1	zeros in
6	101	0	1	1100	2	1	zeros in

Step	Gray (3-bit)	flips bit#	Ham. dist	Johnson (4-bit)	flips bit#	Ham. dist	Johnson phase
7	100	2	1	1000	3	1	zeros in
wrap	-> 000	2	1	-> 0000	3	1	–

Every step – **including the wrap back to state 0** – flips exactly one bit in both encodings. Hamming distance 1 is the only property the synchronizer needs, which is why both are safe to cross.

The Johnson wrap is single-bit, and this is easy to get wrong. A reviewer of this document read the earlier “bit changed” column as a bit *count* and concluded the Johnson wrap 1000 -> 0000 changed four bits. It changes one (bit 3). A twisted-ring counter is specifically constructed so that the wrap is a single-bit transition like every other step – that is what makes the 2N cycle usable as a CDC pointer at all.

The cost difference is in the width, and it compounds:

Pointer width by encoding – Gray is logarithmic, Johnson is linear

States needed (2 x DEPTH)	DEPTH		
	H	Gray pointer bits	Johnson pointer bits
8	4	3	4
16	8	4	8
32	16	5	16
64	32	6	32
128	64	7	64

Gray is log2 in the state count; Johnson is linear. At DEPTH=4 the difference is one flop; at DEPTH=64 it is 57 flops per pointer, and each pointer is duplicated per domain and again per synchronizer stage. That is the entire argument for Gray being the default – and equally, why Johnson stays viable only at the modest depths where non-power-of-2 sizing actually matters.

What Johnson buys for that width is the freedom to stop at any count at all: its 2N progression lands on 6, 10, 14 – and on 21 or 36 – just as naturally as on 8 or 16, whereas Gray’s 2**N can only land on powers of two. Odd depths are legal too; the “even only” restriction is stale language from the retired `fifo_async_div2`, and the only elaboration check in the RTL fires on Gray.

1.7.2 Sizing the depth

`depth >= burst_length * (1 - read_rate / write_rate)`

where $\text{rate} = \text{freq} * \text{duty_cycle}$. Round up, then apply the encoding's constraint: Gray rounds to the next power of two; Johnson takes the depth you ask for. The only elaboration check is $(\text{USE_JOHNSON} == 0) \ \&\& \ ((\text{DEPTH} \ \& \ (\text{DEPTH}-1)) \neq 0)$, so Johnson imposes no restriction at all.

This formula assumes the reader keeps draining throughout the burst. It sizes for a *sustained rate mismatch*, not for burst isolation. If the reader can stall completely – arbitration loss, a blocked downstream, a paused consumer – then nothing drains during the stall and you need $\text{depth} \geq \text{burst_length}$ to avoid dropping writes. Between those extremes, size for the worst-case drain rate the reader can guarantee, not its average.

One term the formula above leaves out: the full flag does not track the true fill level instantly. It is derived from a pointer that has been synchronised into the write domain AND then registered, so it lags real reads by $N_FLOP_CROSS + 1$ write clocks – the synchroniser stages plus the flop on `wr_full` itself in `fifo_control`. Production sizing therefore adds $N_FLOP_CROSS + 1$ slots on top of the raw depth – three for the default two-stage synchroniser, four if you have configured a third stage. The worked example below shows the raw depth only, so that it lines up with the storage-overhead tables that follow; add the margin before committing to a number.

Spreadsheet. `docs/fifo_depth_calculator_v2.xlsx` does this whole calculation – rate-from-duty, raw depth, synchroniser margin, then the Gray power-of-two rounding and Johnson's absence of rounding side by side, with the slot and percentage saving between them. It carries worked example cases and a one-page cheat sheet. Note that it deliberately does *not* model clock drift: crystal tolerance is ppm-level and moves the answer by well under one slot over any bounded burst, and a divergence large enough to matter is a steady-state violation whose fix is back-pressure, not a deeper FIFO. To size a worst case, enter the worst-case frequencies directly – fastest writer, slowest reader.

Worked example. Writer at 100 MHz bursts 100 back-to-back writes; reader pulls 1 word/clock at 80 MHz:

Worked depth-sizing example

Step	Value
Burst duration	$100 / 100 \text{ MHz} = 1 \mu\text{s}$
Reads during burst	$1 \mu\text{s} \times 80 \text{ MHz} = 80$
Net build-up	$100 - 80 = \mathbf{20 \text{ words}}$
Gray FIFO (power-of-2)	round up to 32
Johnson FIFO (any depth)	20 as-is

Gray adds 12 unused slots, and the nominal overhead scales with data width:

Nominal storage overhead of Gray depth rounding by data width

DATA_WIDTH	Gray depth 32	Johnson depth 20	Nominal saving
32 b	1024 b	640 b	384 b (38%)
64 b	2048 b	1280 b	768 b (38%)
256 b	8192 b	5120 b	3072 b (38%)
1024 b	32 768 b	20 480 b	12 288 b (38%)

Those are *nominal* bit counts. Whether they translate into real area depends entirely on how the array is implemented – see the walkthrough below, which is the case most people actually hit.

Add margin for synchronizer latency if you are near the safe depth – the backpressure to the writer lags real fill level by $N_FLOP_CROSS + 1$ write clocks (the synchroniser stages, plus the registered flag).

1.7.3 Walkthrough: is a non-power-of-2 FIFO worth it at 512 bits?

This is the question that motivates `USE_JOHNSON`, and the honest answer is “sometimes” – the naive 38% figure above is frequently zero in practice. Work it through.

Scenario. A 512-bit streaming path. Writer at 250 MHz bursts 96 words back-to-back; reader drains 1 word/clock at 200 MHz.

Step 1 – required depth.

$\text{depth} \geq 96 * (1 - 200/250) = 96 * 0.2 = 19.2 \rightarrow 20 \text{ words}$

Step 2 – apply the encoding constraint.

512-bit walkthrough: depth after applying the encoding constraint

Encoding	Legal depth	Chosen	Entries wasted
Gray (<code>USE_JOHNSON=0</code>)	power of 2	32	12
Johnson (<code>USE_JOHNSON=1</code>)	any depth	20	0

Step 3 – nominal storage delta.

512-bit walkthrough: nominal storage delta

	Depth	Bits
Gray	32 x 512	16 384 b (2 KiB)

	Depth	Bits
Johnson	20 x 512	10 240 b (1.25 KiB)
Delta	12 entries	6 144 b (768 B)

Step 4 – what Johnson costs. The pointer is DEPTH bits wide instead of $\log_2(\text{DEPTH})+1$, and it is duplicated per domain and again per synchronizer stage. With $N_FLOP_CROSS=2$:

512-bit walkthrough: pointer flop cost

	Gray (depth 32)	Johnson (depth 20)
Pointer width	6 b	20 b
Per domain: own binary + encoded pointer	$6 + 6 = 12$	$6 + 20 = 26$
Per domain: synchronized remote pointer	$2 \times 6 = 12$	$2 \times 20 = 40$
Per domain: gray->bin converter	0 (combinational)	0 (combinational)
Both domains	~48 flops	~132 flops

So Johnson costs roughly **+84 flops** to save **6 144 memory bits**. If the memory is flops, that is a ~73:1 win. If it is not, the trade can invert completely.

Step 5 – the part that decides it: memory granularity. MEM_STYLE (FIFO_AUTO / FIFO_SRL / FIFO_BRAM) determines whether 12 fewer entries costs 12 fewer entries of area:

512-bit walkthrough: whether the saving is real, by memory style

MEM_STYLE	Granularity	Depth 32 vs 20	Verdict
Flop array	1 entry	16 384 vs 10 240 flops	Johnson wins big – saving is exactly proportional, and dwarfs the +84 pointer flops
FIFO_SRL	32 deep per LUT (SRL32)	Both ≤ 32 , so 512 LUTs either way	No saving. Johnson costs +84

MEM_STYLE	Granularity	Depth 32 vs 20	Verdict
			flops for nothing
FIFO_BRAM	512 deep x 72 b per BRAM	512-bit width needs $\text{ceil}(512/72) = 8$ BRAMs, and both depths are far under 512	No saving. Same 8 BRAMs either way

Conclusion for this example. At 512 bits and depth ~20, Johnson is a clear win *only* if the array is implemented as flops. Under SRL or BRAM the depth rounding is absorbed by the primitive's own granularity, and choosing Johnson is a pure loss of +84 flops plus a wider pointer comparator on the critical path.

When a non-power-of-2 depth genuinely pays.

1. **Flop-based arrays** – no granularity, saving is proportional. The most common real win, and it grows with DATA_WIDTH.
2. **The depth is externally constrained, not free to round up.** If you need exactly 20 entries to cover a 20-beat burst, a credit limit, or a descriptor count, rounding to 32 is not “spare capacity”, it is 12 entries of dead silicon you provisioned to satisfy an encoding.
3. **Crossing a primitive boundary.** If the rate analysis lands at 33 and Gray forces 64, an SRL32-based FIFO goes from 2 primitives per bit to 2 – no help – but a flop array pays the full 31 entries. Check where your required depth sits relative to the primitive's step size, not relative to the next power of two.

1.7.3.1 On an ASIC the granularity argument disappears

Everything in Step 5 above is an FPGA argument. SRL32 and BRAM have a fixed step size, and that step size is what absorbs the rounding and erases Johnson's advantage. On an ASIC there is no such step: both **register files** and **compiled SRAM** are generated to a requested word count, and every memory compiler in common use takes the word count as an argument. So the depth you compute is the depth you instantiate, and an odd rounding to the next power of two is pure wasted silicon rather than something a primitive was going to charge you for anyway.

Worked case: a required depth of 36. Take the same 512-bit path, but with a rate analysis that lands on 36 entries – above 32, so Gray must round to 64.

Depth-36 case: storage delta when Gray must round 36 up to 64

	Depth	Storage at 512 b
Gray (power of 2)	64	32 768 b (4 KiB)

	Depth	Storage at 512 b
Johnson (any depth)	36	18 432 b (2.25 KiB)
Delta	28 entries	14 336 b (1.75 KiB)

Gray pays for 28 entries it will never use – 44 percent of the array – because 36 is not a power of two. The pointer cost runs the other way, by the same accounting as Step 4:

Depth-36 case: pointer flop cost

	Gray (depth 64)	Johnson (depth 36)
Pointer width	7 b	36 b
Per domain: own binary + encoded pointer	$7 + 7 = 14$	$7 + 36 = 43$
Per domain: synchronized remote pointer	$2 \times 7 = 14$	$2 \times 36 = 72$
Per domain: bin converter	0 (combinational)	0 (combinational)
Both domains	~56 flops	~230 flops

Johnson costs about **+174 flops** to save **14 336 memory bits** – roughly an 82:1 return, better than the depth-20 case because the rounding gap is wider.

Which ASIC memory style. At depth 36 both are available, and the choice is about aspect ratio rather than about the encoding:

Depth-36 case: ASIC memory styles, all of which take the word count as an argument

Style	Arbitrary depth?	Fit at 36 x 512
Flop array	Yes, no granularity at all	Works, but 18 432 flops is a lot of area and a synthesis burden
Register file	Yes, word count is a compiler argument	Good fit. Shallow and wide is what register files are for
Compiled SRAM	Yes, word count is a compiler argument	Works, but 36 words is shallow for SRAM; check the compiler's minimum depth and its area per bit at this aspect ratio

The practical answer at this shape is usually a register file: it takes depth 36 directly, it is denser than a flop array, and it avoids the periphery overhead that makes a 36-word SRAM inefficient. Confirm against your own compiler – minimum depth, word-count granularity (some compilers step in multiples of 4 or 8 words, which would reintroduce a smaller version of the FPGA problem), and area per bit at a shallow, wide aspect ratio.

Rule of thumb. Compute the required depth first. Then ask what the memory maps to and what its step size is. Only if the rounding actually crosses a step boundary – or there is no step, as with flops, register files, and compiled SRAM – does `USE_JOHNSON=1` buy you anything. On an ASIC that condition is usually met, so Johnson is worth pricing whenever the required depth is not already a power of two. On an FPGA it usually is not met, and Gray is the default: narrower pointers, and one less thing to reason about. (Both converters are combinational – `gray2bin` and `johnson2bin` alike – so the difference is pointer width, not a pipeline stage.)

1.8 SDC constraints

Both handshakes need the control toggles constrained as bounded crossings, and the quasi-static data bus constrained by `set_max_delay`, **not** `set_false_path`.

The source flop names differ between the two modules, so the constraints are not copy-paste interchangeable. The synchronizer registers (`r_req_sync`, `r_ack_sync`) and the data registers (`r_src_data_hold`, `r_dst_data`) are named the same in both; only the req/ack source flops differ:

	2-phase	4-phase
request source flop	<code>r_req_tog</code>	<code>r_req_src</code>
acknowledge source flop	<code>r_ack_tog</code>	<code>r_ack_dst</code>

The example below is written for `cdc_2_phase_handshake`. For `cdc_4_phase_handshake`, substitute per the table (its own RTL header prints the constraints with the right names).

```
# 1. Req / Ack single-bit crossings [2-phase names -- see table
above]
set_max_delay -datapath_only \
  -from [get_pins u_cdc/r_req_tog_reg/C] \
  -to   [get_pins u_cdc/r_req_sync_reg[0]/D] \
  <dst_period_ns>

set_max_delay -datapath_only \
  -from [get_pins u_cdc/r_ack_tog_reg/C] \
  -to   [get_pins u_cdc/r_ack_sync_reg[0]/D] \
  <src_period_ns>
```

2. Data bus (quasi-static, protected by the handshake)

```
set_max_delay -datapath_only \  
  -from [get_pins u_cdc/r_src_data_hold_reg[*]/C] \  
  -to   [get_pins u_cdc/r_dst_data_reg[*]/D] \  
  <dst_period_ns>
```

Why -datapath_only, and why not set_false_path:

- set_false_path tells the tool to ignore the path entirely. The data bus is then free to take arbitrarily long, but the handshake only holds it stable for a bounded window – so the destination can latch a value that has partially updated. The crossing is unconstrained, not safe.
- set_max_delay **without** -datapath_only still includes clock skew and clock-uncertainty terms in the calculation. Across asynchronous domains those terms are meaningless – the tool has no defined phase relationship to work from – so you get pessimistic, unfixable violations that tempt people back to set_false_path.
- set_max_delay -datapath_only constrains **only the combinational datapath delay** from source register to destination register, which is exactly the quantity that must fit inside the handshake’s stable window. This is the one that expresses the real requirement.

Set the value to the destination clock period (or the source period, for the ack/return direction) unless you have a specific reason to tighten it.

1.9 Common mistakes

1. Multi-flop synchronizer on a multi-bit bus that changes simultaneously. Safe only for single-bit or Gray-coded values. If multiple bits change at once, different bits may be sampled from different time points, producing a value that never existed. Fix: handshake or async FIFO.

2. set_false_path on CDC data buses. It tells the tool to ignore timing entirely, but the data bus must arrive within a bounded window. Use set_max_delay -datapath_only.

3. Assuming async FIFO depth 2 is enough. Both entries are usable – fifo_control uses a wrap-bit full test, not a reserved slot – but the synchronized pointers lag by N_FLOP_CROSS cycles, so a depth-2 FIFO spends most of its time reporting stale fullness and stalls a writer that could have proceeded. Use depth >= 4.

4. Open-loop transfer before the previous one is sampled. Without an ack the NEW transfer is silently dropped – the capture guard (src_valid && !src_busy) ignores a pulse that arrives while the previous one is still crossing, and the first transfer completes. Minimum spacing

STRETCH_CYCLES source clocks (the busy window; sized $\geq \text{SYNC_STAGES} + 1$ destination clocks).

5. Choosing 2-phase for speed without checking the reset domains. See [Reset Considerations](#). This one has cost real silicon debug time.

1.10 Verification status

The formal harness at `formal/cdc/cdc_handshake/formal_cdc_handshake.sv` wires the DUT as:

```
.clk_src (clk), .rst_src_n (rst_n),  
.clk_dst (clk), .rst_dst_n (rst_n),
```

Single clock, single reset for both domains. The proof therefore **cannot express asymmetric reset or asynchronous clocks**. Treat “formally verified” for the handshakes as scoped to protocol correctness under a common clock and reset.

And the DUT it proves is stale. The harness compiles `cdc_handshake_formal.sv`, a Yosys-compatible copy taken before the module was renamed to `cdc_4_phase_handshake` and grew parameters. The copy carries `DATA_WIDTH` alone; the live `rtl/cdc/cdc_4_phase_handshake.sv` carries `DATA_WIDTH`, `SYNC_STAGES`, `TIMEOUT_CYCLES` and `FAST_PATH`.

Updated 2026-09-16 (CDC-FORMAL-STALE closed). The fork is gone: the proof now flattens and reads the shipped module, and drives each parameter from its own sbv task. The timeout path is proven – it never fires before a transfer is sent and always fires once one stalls past the programmed count.

`FAST_PATH` no longer exists (cdc BUG-001 (was CDC-002), 2026-09-16). Formal proved it unsound: `D_IDLE` sampled `dst_ready` and then raised `dst_valid` and the ack together on the NEXT cycle, so a receiver that dropped ready in between was acknowledged for a beat it never took – silent data loss. Made correct it saved nothing, so the parameter was removed rather than left as a knob that did nothing. The destination now always waits for an observed `dst_valid && dst_ready`.

The reset behavior described in this document is argued from the RTL encoding and – for the 2-phase hazard – confirmed on silicon. It is **not** currently covered by the formal proof or by a directed test. Extending the harness to independent clocks and resets would cover it, and would be expected to fail on the 2-phase variant.

1.11 Related Modules

CDC module reference

Module	RTL	Filelist	Test
bin2gray	rtl/cdc/ bin2gray.s v	bin2gray.f	val/cdc/ test_bin2gra y.py
cdc_2_phase_hands hake	rtl/cdc/ cdc_2_phas e_handshak e.sv	cdc_2_phase_handshake. f	val/cdc/ test_cdc_2_p hase_handsha ke.py
cdc_4_phase_hands hake	rtl/cdc/ cdc_4_phas e_handshak e.sv	cdc_4_phase_handshake. f	val/cdc/ test_cdc_4_p hase_handsha ke.py
cdc_open_loop	rtl/cdc/ cdc_open_l oop.sv	cdc_open_loop.f	val/cdc/ test_cdc_ope n_loop.py
cdc_synchronizer	rtl/cdc/ cdc_synchr onizer.sv	cdc_synchronizer.f	–
counter_bingray	rtl/cdc/ counter_bi ngray.sv	counter_bingray.f	val/cdc/ test_counter _bingray.py, val/cdc/test _counter_bin gray_wavedro m.py
counter_johnson	rtl/cdc/ counter_jo hnson.sv	counter_johnson.f	val/cdc/ test_counter _johnson.py, val/cdc/test _counter_joh nson_wavedro m.py
fifo_async	rtl/cdc/ fifo_async .sv	fifo_async.f	val/cdc/ test_fifo_bu ffer_async.p y, val/cdc/test _fifo_async_ wavedrom.py
gaxi_fifo_async	rtl/cdc/ gaxi_fifo_	gaxi_fifo_async.f	val/cdc/ test_gaxi_bu

Module	RTL	Filelist	Test
	async.sv		ffer_async.py
gaxi_skid_buffer_async	rtl/cdc/gaxi_skid_buffer_async.sv	gaxi_skid_buffer_async.f	val/cdc/test_gaxi_buffer_async.py
gray2bin	rtl/cdc/gray2bin.sv	gray2bin.f	val/cdc/test_gray2bin.py
johnson2bin	rtl/cdc/johnson2bin.sv	johnson2bin.f	val/cdc/test_johnson2bin.py

Protocol-level CDC built on these blocks: apb4_slave_cdc, apb4_slave_cdc_cg, apb5_slave_cdc, apb5_slave_cdc_cg, gaxi_skid_buffer_async, axi4_to_apb4_shim.

Supporting primitives (rtl/cdc): bin2gray, gray2bin, johnson2bin, counter_bingray, counter_johnson, glitch_free_n_dff_arn, sync_pulse, reset_sync.

1.12 Testing

Each mechanism documented here has its own module page with a Testing section and suite. The whole area runs with `make -C val/cdc clean-all && make -C val/cdc run-all-{gate,func,full}-parallel`, and every rtl/cdc module carries a passing formal harness under formal/cdc/.

1.13 Navigation

- [← Back to CDC Index](#)
- [← Back to Main Documentation Index](#)

2 Binary-to-Gray converter (bin2gray.sv)

2.1 Overview

bin2gray is a purely combinational binary-to-Gray converter. Gray code – reflected binary code, unit-distance code, same thing – is a binary numeral system where two successive values differ in exactly one bit. That one-bit property is what reduces glitches and metastability exposure when a value moves between asynchronous domains, which is why Gray code turns up everywhere in clock domain crossing work.

Where it earns its keep:

- **Asynchronous FIFOs** – pointer comparisons that can't afford metastability
- **Clock domain crossing** – safe multi-bit signal transfer
- **Position encoders** – mechanical/optical encoder interfaces
- **Memory address generation** – reduces EMI and power spikes
- **Test pattern generation** – controlled single-bit transitions
- **Error detection** – single-bit error detection schemes
- **ADC/DAC interfaces** – fewer glitches in conversion systems

The whole module is one parameter and two ports:

```
module bin2gray #(
    parameter int WIDTH = 4
) (
    input wire [WIDTH-1:0] binary,
    output wire [WIDTH-1:0] gray
);
```

2.2 Parameters

bin2gray parameters

Parameter	Default	Description
WIDTH	4	Bit width of both input and output (int, any positive integer ≥ 1 ; common values 4, 8, 16, 32 for address counters and data buses). Determines the number of XOR gates required.

2.3 Ports

bin2gray ports

Port	Direction	Width	Description
binary	input	WIDTH	Input binary value
gray	output	WIDTH	Output Gray code value

2.4 Functional Description

2.4.1 Mathematical definition

The conversion follows one simple relationship:

- **MSB:** $\text{gray}[\text{WIDTH}-1] = \text{binary}[\text{WIDTH}-1]$ (MSB unchanged)
- **Other bits:** $\text{gray}[i] = \text{binary}[i] \oplus \text{binary}[i+1]$ for $i = 0$ to $\text{WIDTH}-2$

2.4.2 Why Gray code?

That one-bit property buys you four things:

1. **Single Bit Transitions:** Adjacent values differ by exactly one bit
2. **Glitch Reduction:** Eliminates intermediate states during transitions
3. **Metastability Prevention:** Safer for asynchronous clock domain crossings
4. **Mechanical Encoders:** Natural for optical/magnetic position encoders

2.4.3 Gray code sequences

Walk the tables and watch the “changed bit” column – one bit, every time.

2.4.3.1 2-bit Gray Code

Decimal	Binary	Gray	Transition
0	00	00	-
1	01	01	bit 0
2	10	11	bit 1
3	11	10	bit 0

2.4.3.2 3-bit Gray Code

Decimal	Binary	Gray	Changed Bit
0	000	000	-
1	001	001	0
2	010	011	1
3	011	010	0
4	100	110	2
5	101	111	0
6	110	101	1
7	111	100	0

2.4.3.3 4-bit Gray Code (Complete Table)

Dec	Binary	Gray	Dec	Binary	Gray
0	0000	0000	8	1000	1100
1	0001	0001	9	1001	1101
2	0010	0011	10	1010	1111
3	0011	0010	11	1011	1110
4	0100	0110	12	1100	1010
5	0101	0111	13	1101	1011
6	0110	0101	14	1110	1001
7	0111	0100	15	1111	1000

2.4.4 Core logic

```
genvar i;
generate
    for (i = 0; i < WIDTH - 1; i++) begin : gen_gray
        assign gray[i] = binary[i] ^ binary[i+1];
    end
endgenerate

assign gray[WIDTH-1] = binary[WIDTH-1];
```

What the RTL builds is a bank of XOR gates: each lower Gray bit is the XOR of the current and next binary bits, the MSB comes straight across, and everything evaluates in parallel – a single level of logic.

For WIDTH = 4, gate level:




Diagram 1

2.4.5 Companion converter: gray2bin

You'll almost always need the trip back. The inverse conversion:

```

module gray2bin #(
    parameter int WIDTH = 4
) (
    input wire [WIDTH-1:0] gray,
    output wire [WIDTH-1:0] binary
);

    genvar i;

    // MSB is unchanged
    assign binary[WIDTH-1] = gray[WIDTH-1];
  
```

```

// Other bits: XOR accumulation from MSB down
generate
    for (i = WIDTH-2; i >= 0; i--) begin : gen_binary
        assign binary[i] = binary[i+1] ^ gray[i];
    end
endgenerate

```

endmodule

2.4.5.1 Gray-to-Binary Truth Table (4-bit)

Gray	Binary	Conversion Process
0000	0000	bin[3]=0, bin[2]=0 $\oplus$ 0=0, bin[1]=0 $\oplus$ 0=0, bin[0]=0 $\oplus$ 0=0
0001	0001	bin[3]=0, bin[2]=0 $\oplus$ 0=0, bin[1]=0 $\oplus$ 0=0, bin[0]=0 $\oplus$ 1=1
0011	0010	bin[3]=0, bin[2]=0 $\oplus$ 0=0, bin[1]=0 $\oplus$ 1=1, bin[0]=1 $\oplus$ 1=0
0010	0011	bin[3]=0, bin[2]=0 $\oplus$ 0=0, bin[1]=0 $\oplus$ 1=1, bin[0]=1 $\oplus$ 0=1

2.5 Timing Characteristics

One XOR gate of delay – typically 0.1-0.3ns in modern processes – and that’s the whole critical path. Each binary bit drives at most 2 XOR gates, and the delay is constant regardless of WIDTH. This module will never be your timing problem.

Typical FPGA resource usage, by width:

bin2gray resource usage by width (typical FPGA)

WIDTH	LUTs	Delay (ns)	Max Freq (MHz)
4	3	0.2	800+
8	7	0.2	800+

WIDTH	LUTs	Delay (ns)	Max Freq (MHz)
16	15	0.3	600+
32	31	0.4	500+

2.6 Usage Examples

Every parameter and port below is read from the module declaration.

```
bin2gray #(
    .WIDTH                (4)
) u_bin2gray (
    .binary                (binary),
    .gray                  (gray)
);
```

2.7 Design Notes

- **Always use Gray for async boundaries** – it’s the baseline tool for safe asynchronous transfer of a multi-bit count.
- **Timing is a non-issue:** purely combinational, no setup/hold concerns of its own.
- **Verify the single-bit property:** adjacent values must differ by exactly one bit. Check it in simulation, not by eye.
- **Plan for the trip back:** you usually need gray2bin on the other side.
- **Check the width:** enough bits for the application range, no more.

The binary-to-Gray converter is one of those fundamental building blocks – simple, combinational, dependable. For asynchronous interfaces and glitch-free operation it’s close to indispensable. Respect the two rules from the multi-bit synchronizer example above – register the Gray value in the source domain, and make sure it only ever increments by one – and this module will never be the source of your bug.

2.8 Related Modules

- **gray2bin:** performs the inverse conversion (Gray to binary)
- **counter_bingray:** combined binary/Gray counter; what the FIFOs instantiate
- **glitch_free_n_dff_arn / cdc_synchronizer:** multi-flop synchronizers that carry the registered Gray value across
- **fifo_async / gaxi_fifo_async:** the primary users of this conversion

2.9 Testing

2.9.1 Comprehensive Test Bench

```
module tb_bin2gray;

    parameter int WIDTH = 4;
    parameter int MAX_VAL = (1 << WIDTH) - 1;

    logic [WIDTH-1:0] binary, gray;
    logic [WIDTH-1:0] expected_gray;
    logic [WIDTH-1:0] binary_check;

    // DUT
    bin2gray #(.WIDTH(WIDTH)) dut (
        .binary(binary),
        .gray(gray)
    );

    // Reference converter for checking
    gray2bin #(.WIDTH(WIDTH)) check_conv (
        .gray(gray),
        .binary(binary_check)
    );

    // Test sequence
    initial begin
        $display("Testing %d-bit Binary to Gray converter", WIDTH);

        // Test all possible values
        for (int i = 0; i <= MAX_VAL; i++) begin
            binary = i;
            expected_gray = compute_gray_reference(i);

            #1; // Allow propagation

            // Check conversion correctness
            if (gray !== expected_gray) begin
                $error("Mismatch: binary=%b, expected_gray=%b,
actual_gray=%b",
                    binary, expected_gray, gray);
            end

            // Check round-trip conversion
            if (binary_check !== binary) begin
                $error("Round-trip failed: original=%b, recovered=%b",
```

```

        binary, binary_check);
    end

    // Check single-bit change property
    if (i > 0) begin
        logic [WIDTH-1:0] prev_gray =
compute_gray_reference(i-1);
        int bit_changes = $countones(gray ^ prev_gray);
        if (bit_changes != 1) begin
            $error("Multiple bit change: %d→%d, gray %b→%b,
changes=%d",
                    i-1, i, prev_gray, gray, bit_changes);
        end
    end
    end

    $display("PASS: %d → binary:%b → gray:%b", i, binary,
gray);
    end

    $display("All tests passed!");
    $finish;
end

// Reference Gray code computation
function [WIDTH-1:0] compute_gray_reference(input [WIDTH-1:0]
bin);
    compute_gray_reference[WIDTH-1] = bin[WIDTH-1];
    for (int i = WIDTH-2; i >= 0; i--) begin
        compute_gray_reference[i] = bin[i] ^ bin[i+1];
    end
endfunction

endmodule

```

2.9.2 Property-Based Verification

```

// bin2gray is purely combinational -- it has no clock, so there is no
// clocking
// event to write concurrent assertions against. Use IMMEDIATE
// assertions in an
// always_comb block: they re-evaluate whenever an input settles, need
// no clock,
// and work under both simulation and formal.
//
// The checker declares PORTS. A checker with no ports bound with
// `(.*)` would
// connect nothing, and every assertion would silently evaluate on X.

```

```

module bin2gray_properties #(
    parameter int WIDTH = 4
) (
    input logic [WIDTH-1:0] binary,
    input logic [WIDTH-1:0] gray
);

    // Reference encode/decode as functions -- a property body cannot
    contain
    // variable declarations, continuous assigns, or genvar loops.
    function automatic logic [WIDTH-1:0] gray_encode(input logic
[WIDTH-1:0] bin);
        gray_encode = bin ^ (bin >> 1);
    endfunction

    function automatic logic [WIDTH-1:0] gray_decode(input logic
[WIDTH-1:0] gry);
        gray_decode[WIDTH-1] = gry[WIDTH-1];
        for (int i = WIDTH-2; i >= 0; i--) begin
            gray_decode[i] = gray_decode[i+1] ^ gry[i];
        end
    endfunction

    always_comb begin
        // MSB is passed through unchanged
        a_msb_unchanged: assert (gray[WIDTH-1] == binary[WIDTH-1]);

        // Output matches the reference encoding
        a_matches_reference: assert (gray == gray_encode(binary));

        // Adjacent codes differ in exactly one bit (skip the wrap at
        the top,
        // which is also single-bit but is worth asserting separately)
        if (binary < WIDTH'((2**WIDTH) - 1)) begin
            a_single_bit_change: assert (
                $countones(gray ^ gray_encode(binary + WIDTH'(1))) ==
1
            );
        end

        // Decoding the output recovers the input
        a_round_trip: assert (gray_decode(gray) == binary);
    end

endmodule

```

```

// Bind at the point of use, OUTSIDE the checker -- a module cannot
bind itself
// into another module from within its own body. Ports connect by name
via (.*).
// because the checker's port names match the DUT's.
bind bin2gray bin2gray_properties #(.WIDTH(WIDTH)) props_inst (.*);

```

The bind line belongs in a file compiled alongside the DUT (typically the testbench), not inside the checker module. Under Verilator these compile with `--assert`; under SymbiYosys the same immediate assertions are proven exhaustively rather than sampled.

2.9.3 Coverage Model

```

covergroup bin2gray_cg;

  cp_binary_values: coverpoint binary {
    bins zero = {0};
    bins powers_of_two[] = {1, 2, 4, 8, 16}; // For appropriate
WIDTH
    bins max_value = {2**WIDTH - 1};
    bins mid_range[] = {[1:2**(WIDTH-1)-1]};
    bins upper_range[] = {[2**(WIDTH-1):2**WIDTH-2]};
  }

  cp_gray_values: coverpoint gray {
    bins all_values[] = {[0:2**WIDTH-1]};
  }

  cp_bit_patterns: coverpoint binary {
    bins alternating_01 = {8'b01010101}; // For WIDTH=8
    bins alternating_10 = {8'b10101010};
    bins all_ones = {'1'};
    bins all_zeros = {'0'};
  }

  // Cross coverage between input patterns and outputs
  cross_binary_gray: cross cp_binary_values, cp_gray_values;

endcovergroup

```

2.9.4 Test files

- `val/cdc/test_bin2gray.py` – functional verification
- `pytest val/cdc/test_bin2gray.py -v`

2.10 Navigation

- [← Back to CDC Index](#)
- [← Back to Main Documentation Index](#)

3 Gray-to-binary converter (gray2bin.sv)

3.1 Overview

gray2bin converts Gray code (reflected binary code) back to standard binary using XOR reduction. When a Gray pointer crosses into a domain that needs to do arithmetic on it, this is the module that turns it back into a number – asynchronous FIFO pointer comparison is the classic case, and it shows up in plenty of other CDC paths too.

3.2 Parameters

gray2bin parameters

Parameter	Default	Description
WIDTH	4	Data width in bits

3.3 Ports

gray2bin ports

Port	Direction	Width	Description
gray	input	WIDTH	Gray code input
binary	output	WIDTH	Binary code output

3.4 Functional Description

Gray code is a binary numeral system where **only one bit changes** between consecutive values:

Binary	Gray	Transitions
000	→ 000	
001	→ 001	1 bit change
010	→ 011	1 bit change
011	→ 010	1 bit change
100	→ 110	1 bit change
101	→ 111	1 bit change
110	→ 101	1 bit change
111	→ 100	1 bit change

That property is exactly what CDC needs:

- **Metastability protection:** only one bit in transition means no multi-bit race conditions

- **Safe sampling:** intermediate values during a transition are still valid Gray codes
- **FIFO pointers:** essential for async FIFO full/empty detection

3.4.1 Conversion algorithm

Each binary bit is the XOR of all Gray bits from that position up to the MSB:

$$\text{binary}[i] = \text{gray}[\text{MSB}] \oplus \text{gray}[\text{MSB}-1] \oplus \dots \oplus \text{gray}[i]$$

The RTL says the same thing with a shift and a reduction:

```
genvar i;
generate
  for (i = 0; i < WIDTH; i++) begin : gen_gray_to_bin
    assign binary[i] = ^(gray >> i);
  end
endgenerate
```

Bit by bit, for a 4-bit converter:

```
// For 4-bit example:
assign binary[3] = gray[3];           // MSB unchanged
assign binary[2] = gray[3] ^ gray[2]; // XOR from MSB
down
assign binary[1] = gray[3] ^ gray[2] ^ gray[1]; // XOR from MSB
down
assign binary[0] = gray[3] ^ gray[2] ^ gray[1] ^ gray[0]; // XOR all
bits
```

3.4.2 Worked conversions

3.4.2.1 4-Bit Conversion Table

Gray[3:0]	Binary[3:0]	Calculation
0000	0000	0⊕0⊕0⊕0=0, 0⊕0⊕0=0, 0⊕0=0, 0=0
0001	0001	0⊕0⊕0⊕1=1, 0⊕0⊕0=0, 0⊕0=0, 0=0
0011	0010	0⊕0⊕1⊕1=0, 0⊕0⊕1=1, 0⊕0=0, 0=0
0010	0011	0⊕0⊕1⊕0=1,

Gray[3:0]	Binary[3:0]	Calculation
		$0 \oplus 0 \oplus 1 = 1, 0 \oplus 0 = 0,$ $0 = 0$
0110	0100	$0 \oplus 1 \oplus 1 \oplus 0 = 0,$ $0 \oplus 1 \oplus 1 = 0, 0 \oplus 1 = 1,$ $0 = 0$
0111	0101	$0 \oplus 1 \oplus 1 \oplus 1 = 1,$ $0 \oplus 1 \oplus 1 = 0, 0 \oplus 1 = 1,$ $0 = 0$
0101	0110	$0 \oplus 1 \oplus 0 \oplus 1 = 0,$ $0 \oplus 1 \oplus 0 = 1, 0 \oplus 1 = 1,$ $0 = 0$
0100	0111	$0 \oplus 1 \oplus 0 \oplus 0 = 1,$ $0 \oplus 1 \oplus 0 = 1, 0 \oplus 1 = 1,$ $0 = 0$

3.4.2.2 Step-by-Step Example (Gray 0110 → Binary)

Gray input: 0110

```

binary[3] = gray[3] = 0
binary[2] = gray[3] ^ gray[2] = 0 ^ 1 = 1
binary[1] = gray[3] ^ gray[2] ^ gray[1] = 0 ^ 1 ^ 1 = 0
binary[0] = gray[3] ^ gray[2] ^ gray[1] ^ gray[0] = 0 ^ 1 ^ 1 ^ 0 = 0

```

Result: binary = 0100 (decimal 4)

3.4.3 Implementation

The whole module is one line per bit:

```
assign binary[i] = ^(gray >> i);
```

Read it as: shift gray right by i, then XOR-reduce whatever is left – which XORs together every bit from position i to the MSB. The generate loop makes it work for any WIDTH, synthesis tools recognize the pattern and optimize it, all bits compute in parallel, and the result maps onto XOR tree structures.

3.5 Timing Characteristics

The delay is an XOR tree of depth $\log_2(\text{WIDTH})$ – typically 1-2 LUT delays for most widths, with the critical path running from the MSB input to the LSB output.

gray2bin delay by width

WIDTH	XOR Levels	Typical Delay
4	2	1 LUT delay
8	3	2 LUT delays
16	4	2 LUT delays
32	5	3 LUT delays

Modern synthesis tools need no help here: they recognize the XOR-reduction pattern, build balanced trees, and share XOR gates where possible.

3.6 Usage Examples

3.6.1 Asynchronous FIFO pointers

Convert the synchronized Gray pointers back to binary so you can do arithmetic on them:

```
// Convert synchronized Gray pointers back to binary for comparison
gray2bin #(.WIDTH(5)) wr_ptr_conv (
    .gray(sync_wr_ptr_gray),
    .binary(sync_wr_ptr_bin)
);

gray2bin #(.WIDTH(5)) rd_ptr_conv (
    .gray(sync_rd_ptr_gray),
    .binary(sync_rd_ptr_bin)
);

// Now can perform binary arithmetic for occupancy calculation
assign occupancy = sync_wr_ptr_bin - sync_rd_ptr_bin;
```

3.6.2 Clock domain crossing counter

```
// Convert Gray counter back to binary for address generation
gray2bin #(.WIDTH(AW)) addr_converter (
    .gray(gray_counter),
    .binary(memory_address)
);
```

3.6.3 Position encoding

```
// Convert Gray-encoded position to binary for processing
gray2bin #(.WIDTH(8)) position_decoder (
    .gray(gray_position),
    .binary(bin_position)
);
```

3.7 Design Notes

Width scaling. `binary[i]` is the XOR of Gray bits `i..WIDTH-1`, so the unshared cost is $WIDTH * (WIDTH - 1) / 2$ XOR gates – quadratic in `WIDTH`, as the module header notes (“`WIDTH * (WIDTH/2)` average”). Synthesis shares the common prefixes, so realised area lands between linear and quadratic; do not budget it as linear for wide converters. Delay grows with $\log(WIDTH)$, and the module works well up to 64+ bits.

Input validation. Gray codes have no invalid states – every input decodes to something. If you want to gate on the *source* of the Gray code instead:

```
// Gray codes have no invalid states - all inputs are valid
// However, may want to validate source of Gray code
always_comb begin
    if (gray_code_valid) begin
        binary_out = converted_binary;
    end else begin
        binary_out = '0; // Default when invalid
    end
end
```

Pipelining. If you’re closing timing at very high speed, a pipeline stage is cheap insurance:

```
// Optional pipeline stage for timing closure
always_ff @(posedge clk) begin
    if (!rst_n) begin
        binary_reg <= '0;
    end else begin
        binary_reg <= binary_comb;
    end
end
```

3.7.1 Common mistakes

Bit order confusion. The shift direction matters:

```
// WRONG: Bit order matters in Gray codes
assign binary[i] = ^(gray << i); // Left shift instead of right

// CORRECT:
assign binary[i] = ^(gray >> i); // Right shift
```

Width mismatches. Match the parameter to the actual bus width:

```
// WRONG: Input/output width mismatch
gray2bin #(.WIDTH(4)) converter (
    .gray(gray_5bit), // 5-bit input
    .binary(binary_4bit) // 4-bit output
);
```

```
// CORRECT: Match widths
gray2bin #(.WIDTH(5)) converter (
    .gray(gray_5bit),
    .binary(binary_5bit)
);
```

Timing assumptions. The output is combinational, not instantaneous:

```
// WRONG: Assuming zero delay
gray_in <= new_value;
binary_out <= converted_value; // May not be updated yet
```

```
// CORRECT: Account for combinational delay
gray_in <= new_value;
#1; // Wait for conversion
binary_out <= converted_value;
```

3.8 Related Modules

- **bin2gray**: Performs inverse conversion (binary to Gray)
- **counter_bingray**: Combined binary/Gray counter
- **fifo_async**: Uses Gray codes for CDC
- **johnson2bin**: Johnson counter to binary converter (different algorithm)

3.9 Testing

3.9.1 Exhaustive Testing

```
// For reasonable widths, test all possible inputs
for (int i = 0; i < (1 << WIDTH); i++) begin
    gray_input = int_to_gray(i); // Convert integer to Gray
    expected_binary = i; // Expected result
    #1;
    assert(binary_output == expected_binary);
end
```

3.9.2 Property-Based Verification

```
// gray2bin is purely combinational -- no clock port, so there is no
clocking
// event for a concurrent assertion to sample. Use IMMEDIATE
assertions in an
// always_comb block, the same shape bin2gray.md uses: they re-
evaluate whenever
// an input settles, need no clock, and hold under simulation and
formal alike.
```

```

//
// The checker declares PORTS. Binding one with no ports via `(*)`
// connects
// nothing, and every assertion then evaluates silently on X.

module gray2bin_properties #(
    parameter int WIDTH = 4
) (
    input logic [WIDTH-1:0] gray,
    input logic [WIDTH-1:0] binary
);

    function automatic logic [WIDTH-1:0] gray_encode(input logic
[WIDTH-1:0] bin);
        gray_encode = bin ^ (bin >> 1);
    endfunction

    always_comb begin
        // MSB passes through unchanged
        a_msb_unchanged: assert (binary[WIDTH-1] == gray[WIDTH-1]);

        // Round trip: re-encoding the output reproduces the input
        a_round_trip: assert (gray_encode(binary) == gray);
    end

endmodule

// Bind at the point of use, OUTSIDE the checker.
bind gray2bin gray2bin_properties #(.WIDTH(WIDTH)) u_props (*);

```

3.9.3 Random Testing

```

repeat (10000) begin
    random_gray = $random();
    expected_result = reference_gray2bin(random_gray);
    #1;
    assert(binary_output == expected_result);
end

```

3.9.4 Test files

- val/cdc/test_gray2bin.py – functional verification
- ```
pytest val/cdc/test_gray2bin.py -v
```

## 3.10 Navigation

---

- [← Back to CDC Index](#)

- [← Back to Main Documentation Index](#)

## 4 Johnson-to-binary converter (johnson2bin.sv)

### 4.1 Overview

johnson2bin converts Johnson counter codes to binary – the piece that makes asynchronous FIFOs with non-power-of-2 depths possible. Standard Gray-to-binary conversion won't do the job here; the Johnson sequence has its own structure, and decoding it takes position detection rather than XOR reduction.

### 4.2 Parameters

*johnson2bin parameters*

| Parameter | Default | Description                                                               |
|-----------|---------|---------------------------------------------------------------------------|
| JCW       | 10      | Johnson Counter Width (equals FIFO DEPTH)                                 |
| WIDTH     | 4       | Binary output width (typically $\lceil \log_2(\text{DEPTH}) \rceil + 1$ ) |

### 4.3 Ports

*johnson2bin ports*

| Port   | Direction | Width | Description                                                         |
|--------|-----------|-------|---------------------------------------------------------------------|
| clk    | input     | 1     | Clock input. Declared but unused; the module is combinational.      |
| rst_n  | input     | 1     | Active-low reset. Declared but unused; the module is combinational. |
| gray   | input     | JCW   | Johnson counter input (misleadingly named “gray”)                   |
| binary | output    | WIDTH | Binary output representing position in sequence                     |

## 4.4 Functional Description

---

### 4.4.1 The Johnson sequence

Johnson counters are **shift registers with inverted feedback**:

For JCW=6 (DEPTH=6):

```
State 0: 000000 ← Initial state
State 1: 000001 ← Shift left, insert ~MSB (1)
State 2: 000011
State 3: 000111
State 4: 001111
State 5: 011111
State 6: 111111 ← All 1s (fill complete)
State 7: 111110 ← Shift left, insert ~MSB (0)
State 8: 111100
State 9: 111000
State 10: 110000
State 11: 100000
State 0: 000000 ← Cycle complete (12 states total)
```

**Fill direction.** This library shifts in at the LSB. `counter_johnson.sv` implements `counter_gray <= {counter_gray[WIDTH-2:0], ~counter_gray[WIDTH-1]}`, so ones enter at bit 0 and march upward. An earlier revision of this page showed the mirror image (ones entering at the MSB), which is also a valid Johnson counter but is *not* the one this library builds – and both worked conversion examples inherited the error. Note that external references frequently use the MSB-fill convention, so check against `counter_johnson.sv` rather than against a textbook.

The properties that matter for decoding:

- **Single bit transitions:** only one bit changes per state
- **Two phases:**
  - **First half** (0 to DEPTH-1): filling with 1s from the **right** (ones enter at bit 0 and march upward: 000000 → 000001 → 000011 → ...)
  - **Second half** (DEPTH to 2×DEPTH-1): emptying 1s from the **right** (111111 → 111110 → 111100 → ...)
- **Wrap indicator:** the MSB says which half of the cycle you're in

### 4.4.2 The algorithm

The conversion works by **position detection** – finding the transition between the 1s and the 0s:

```

if (w_all_zeroes || w_all_ones) begin
 w_binary = {WIDTH{1'b0}};
end else if (gray[JCW-1]) begin
 // Second half: use trailing one position directly
 w_binary = {{(WIDTH-N){1'b0}}, w_trailing_one};
end else begin
 // First half: use leading one + 1
 w_binary = {{(WIDTH-N){1'b0}}, (w_leading_one + 1'b1)};
end

```

Position detection itself is farmed out to a helper:

```

leading_one_trailing_one #(
 .WIDTH(JCW)
) u_leading_one_trailing_one (
 .data(gray),
 .leadingone(w_leading_one),
 .trailingone(w_trailing_one),
 .all_zeroes(w_all_zeroes),
 .all_ones(w_all_ones),
 .valid(w_valid)
);

```

#### 4.4.3 Three-part binary construction

```

assign binary[WIDTH-1] = gray[JCW-1]; // MSB = wrap
indicator
assign binary[WIDTH-2:0] = w_binary[WIDTH-2:0]; // Lower bits
= position

```

#### 4.4.4 Width calculations

```

localparam int N = $clog2(JCW); // Address
bits needed
localparam int PAD_WIDTH = (WIDTH > N+1) ? WIDTH-N-1 : 0; // Padding
if needed

```

#### 4.4.5 Why this algorithm works

**First half (MSB = 0).** The pattern is 000...0111...1 – zeros followed by ones. The leading-one position tells you how far the fill has progressed: position = leading\_one + 1, because we've filled (leading\_one + 1) positions with 1s.

**Second half (MSB = 1).** The pattern flips to 111...1000...0 – ones followed by zeros. We're emptying from the right, and the trailing-one position shows how far: position = trailing\_one.

## 4.4.6 Worked conversions

### 4.4.6.1 First Half Conversion (MSB = 0)

JCW=6, WIDTH= $\text{clog2}(6)+1=4$  -- the legal pairing, and what both FIFOs instantiate.

Johnson: 001111 (set bits {0,1,2,3}: w\_leading\_one = 3, w\_trailing\_one = 0)

Logic: gray[5]=0 → first half, position = w\_leading\_one + 1 = 3 + 1 = 4

Binary: 4'b0100 = {wrap=0, addr=4} -- Johnson state 4

### 4.4.6.2 Second Half Conversion (MSB = 1)

Johnson: 111000 (set bits {3,4,5}: w\_leading\_one = 5, w\_trailing\_one = 3)

Logic: gray[5]=1 → second half, position = w\_trailing\_one = 3

Binary: 4'b1011 = {wrap=1, addr=3} -- Johnson state 9 of 12

The output is not the number 9. It is the FIFO pointer {wrap, addr}: the ninth state of the sequence is address 3 on the second lap. Reading `4'b1011` as the decimal 11 and expecting 9 is the usual first confusion with this module.

### 4.4.6.3 Special Cases

Johnson: 000000 → Binary: 4'b0000 (all zeros case)

Johnson: 111111 → Binary: 4'b1000 (all ones: lower bits forced to 0, but the RTL

unconditionally sets binary[WIDTH-1] = gray[JCW-1] = 1, so the wrap

bit is SET -- this is {wrap=1, addr=0}, NOT the same as all-zeros.

FIFO full detection depends on this.)

## 4.4.7 Comparison with standard gray2bin

| Aspect            | Standard Gray2Bin     | Johnson2Bin (johnson2bin) |
|-------------------|-----------------------|---------------------------|
| <b>Input type</b> | Traditional Gray code | Johnson counter sequence  |
| <b>Algorithm</b>  | XOR reduction         | Position detection        |
| <b>Complexity</b> | Simple combinational  | Complex position logic    |

| Aspect                   | Standard Gray2Bin    | Johnson2Bin (johnson2bin)                   |
|--------------------------|----------------------|---------------------------------------------|
| <b>Clock requirement</b> | None                 | None (clk/rst_n ports exist but are unused) |
| <b>Width scaling</b>     | Logarithmic          | Linear with JCW                             |
| <b>Use case</b>          | Power-of-2 sequences | Any sequence length                         |

## 4.5 Timing Characteristics

The critical path runs through the position detection logic: leading/trailing one detection, a binary addition, and the output mux. Synthesis complexity is higher than a standard Gray conversion, and resources scale with JCW – the Johnson width, not its log:

Resources scale with JCW (Johnson Counter Width):

- Small FIFOs ( $JCW \leq 16$ ): Reasonable overhead
- Medium FIFOs ( $JCW = 32-64$ ): Significant resources
- Large FIFOs ( $JCW > 64$ ): May become limiting factor

## 4.6 Usage Examples

This is where the module actually lives – decoding the crossed pointer in `fifo_async.sv` (USE\_JOHNSON=1):

```
// fifo_async.sv (USE_JOHNSON=1) usage:
johnson2bin #(
 .JCW(JCW), // = DEPTH
 .WIDTH(AW + 1) // Address width + wrap bit
) rd_ptr_gray2bin_inst(
 .binary(w_wdom_rd_ptr_bin), // Binary for arithmetic
 .gray(r_wdom_rd_ptr_gray), // Johnson counter from CDC
 .clk(wr_clk),
 .rst_n(wr_rst_n)
);
```

## 4.7 Design Notes

### 4.7.1 The clock and reset ports are unused

johnson2bin declares `clk` and `rst_n` but **never uses them**. The conversion is entirely combinational: a single `always_comb` block driven by `leading_one_trailing_one`, which is itself combinational (one `always_comb` plus continuous assignments, no `always_ff` anywhere).

The ports are retained so the module can be dropped into a clocked context without an interface change, and so a future revision could register the output without breaking callers. Tie them off to

whatever the surrounding domain uses; nothing in this module samples them. Do not expect a cycle of latency – the output follows gray combinationally, and long conversion paths must be registered by the caller.

#### 4.7.2 When to use Johnson vs. Gray

```
// Use Johnson counter approach when:
parameter int DEPTH = 10; // Non-power-of-2 required
parameter int DEPTH = 6; // Specific size needed
parameter int DEPTH = 14; // Standard Gray won't work

// Use standard Gray approach when:
parameter int DEPTH = 8; // Power-of-2 is acceptable
parameter int DEPTH = 16; // Efficiency more important
parameter int DEPTH = 64; // Large depth, resource conscious
```

#### 4.7.3 Resource planning

```
// Resource overhead estimation:
// Standard Gray: ~log2(DEPTH) XOR gates
// Johnson: ~DEPTH comparators + position logic

// Break-even point typically around DEPTH = 16-32
```

#### 4.7.4 Verification challenges

```
// More complex verification due to:
// 1. Two-phase sequence behavior
// 2. Position detection correctness
// 3. Special case handling (all 0s, all 1s)
// 4. Wraparound boundary conditions
```

#### 4.7.5 Invalid Johnson sequences

**Nothing checks them.** It is worth being blunt, because the `.valid(w_valid)` port on the helper reads like it does:

```
.valid(w_valid) // NOT a Johnson-pattern check -- see below
```

`leading_one_trailing_one` computes `assign valid = |data`, so `valid` means “input is nonzero” and nothing more. A pattern such as `6'b010101`, which is not a Johnson sequence at all, drives it high. And in `johnson2bin.sv` `w_valid` is declared, connected, and then never read – no output carries it, no assertion consumes it. A malformed sequence is decoded into a wrong pointer silently.

A valid Johnson pattern is one with a single 1-to-0 transition (or its mirror). If you need that enforced, it has to come from somewhere else: the upstream `counter_johnson` produces only legal patterns by construction, which is why the decoder can assume them.

#### 4.7.6 Debug visibility

```
// Monitor internal state for debugging:
// - w_leading_one: Position of leftmost 1
// - w_trailing_one: Position of rightmost 1
// - w_all_zeroes: Special case flag
// - w_all_ones: Special case flag
// - w_valid: nonzero-input flag ONLY -- not a Johnson-pattern check
```

#### 4.7.7 Common issues

1. **Invalid sequences:** non-Johnson input decodes silently to a wrong pointer
2. **Width mismatches:** JCW vs. actual Johnson counter width
3. **Phase confusion:** misunderstanding first vs. second half logic
4. **Boundary conditions:** all-zeros and all-ones handling

#### 4.7.8 Advanced topics

**Hierarchical Johnson counters.** For very large depths, a hierarchical approach is worth considering:

```
// Break large Johnson counter into smaller segments
// Use multiple johnson2bin instances with higher-level arbitration
```

**Alternative position detection.** Other implementations detect position differently:

- **Priority encoders:** find first/last set bit
- **Thermometer decoders:** convert 1-hot positions
- **LUT-based:** for small, fixed widths

### 4.8 Related Modules

---

- **counter\_johnson:** Generates Johnson counter sequences
- **leading\_one\_trailing\_one:** Position detection helper
- **fifo\_async / gaxi\_fifo\_async (USE\_JOHNSON=1):** Primary users of this conversion
- **gray2bin:** Standard Gray-to-binary for comparison
- **fifo\_control:** Uses converted binary values for status generation

### 4.9 Testing

---

From the test suite (`val/cdc/test_johnson2bin.py`):

- **Key test scenarios:**
  - Gray Johnson Counter to Binary Converter Test

- This test verifies the Gray Johnson counter to binary conversion functionality:
- CONFIGURATION:
- JCW: Johnson Counter Width (10, 12, 16, 20)
- WIDTH: Binary output width (4, 5, 6, 8)

Run levels come from the standard grid: REG\_LEVEL=GATE|FUNC|FULL selects the parameter set, TEST\_LEVEL the per-test depth. Run the whole area with `make -C val/cdc run-all-func-parallel`, never bare `pytest` for suites.

## 4.10 Navigation

---

- [← Back to CDC Index](#)
- [← Back to Main Documentation Index](#)

## 5 Binary-Gray counter (counter\_bingray.sv)

### 5.1 Overview

counter\_bingray is a dual-output counter: you get the binary and the Gray representation of the same count, registered in parallel. Asynchronous FIFOs are the reason it exists – the Gray value is what crosses the clock boundary safely (one bit changes per increment, so metastability never gets a foothold), while the binary value stays home for the arithmetic.

You'll also see it used for clock-domain-crossing counters, position encoder interfaces, state machine counters that need glitch-free outputs, and memory address generation in dual-port systems.

```
module counter_bingray #(
 parameter int WIDTH = 4
) (
 input logic clk,
 input logic rst_n,
 input logic enable,
 output logic [WIDTH-1:0] counter_bin,
 output logic [WIDTH-1:0] counter_bin_next,
 output logic [WIDTH-1:0] counter_gray
);
```

### 5.2 Parameters

*counter\_bingray parameters*

| Parameter | Default | Description                                                                                                                                         |
|-----------|---------|-----------------------------------------------------------------------------------------------------------------------------------------------------|
| WIDTH     | 4       | Bit width of both binary and Gray code outputs (int, any positive integer $\geq 1$ ).<br>Determines maximum count value ( $2^{\text{WIDTH}} - 1$ ). |

### 5.3 Ports

*counter\_bingray ports*

| Port   | Direction | Width | Description                   |
|--------|-----------|-------|-------------------------------|
| clk    | input     | 1     | System clock input            |
| rst_n  | input     | 1     | Active-low asynchronous reset |
| enable | input     | 1     | Counter enable control        |

| Port             | Direction | Width | Description                        |
|------------------|-----------|-------|------------------------------------|
| counter_bin      | output    | WIDTH | Binary counter output (registered) |
| counter_bin_next | output    | WIDTH | Next binary value (combinational)  |
| counter_gray     | output    | WIDTH | Gray code output (registered)      |

## 5.4 Functional Description

Gray code (reflected binary code) guarantees adjacent values differ by exactly one bit. That buys you:

1. **Metastability Prevention:** when crossing clock domains, only one bit changes at a time
2. **Glitch Elimination:** reduces intermediate states during transitions
3. **Asynchronous Safety:** safe for use in asynchronous circuits

The conversion is the standard one:

- **MSB:**  $\text{gray}[\text{WIDTH}-1] = \text{binary}[\text{WIDTH}-1]$
- **Other bits:**  $\text{gray}[i] = \text{binary}[i] \oplus \text{binary}[i+1]$  for  $i = 0$  to  $\text{WIDTH}-2$

### 5.4.1 Gray code sequence (4-bit)

| Decimal | Binary | Gray | Changes |
|---------|--------|------|---------|
| 0       | 0000   | 0000 | -       |
| 1       | 0001   | 0001 | bit 0   |
| 2       | 0010   | 0011 | bit 1   |
| 3       | 0011   | 0010 | bit 0   |
| 4       | 0100   | 0110 | bit 2   |
| 5       | 0101   | 0111 | bit 0   |
| 6       | 0110   | 0101 | bit 1   |
| 7       | 0111   | 0100 | bit 0   |
| 8       | 1000   | 1100 | bit 3   |
| 9       | 1001   | 1101 | bit 0   |

| Decimal | Binary | Gray | Changes |
|---------|--------|------|---------|
| 10      | 1010   | 1111 | bit 1   |
| 11      | 1011   | 1110 | bit 0   |
| 12      | 1100   | 1010 | bit 2   |
| 13      | 1101   | 1011 | bit 0   |
| 14      | 1110   | 1001 | bit 1   |
| 15      | 1111   | 1000 | bit 0   |

### 5.4.2 Internal logic

```
logic [WIDTH-1:0] w_counter_bin, w_counter_gray;
```

```
assign w_counter_bin = enable ? (counter_bin + 1) : counter_bin;
assign w_counter_gray = w_counter_bin ^ (w_counter_bin >> 1);
assign counter_bin_next = w_counter_bin;
```

The next binary value is a conditional increment; the Gray value is that next binary value XORed with itself shifted right by one. Both registers update from the same `always_ff`:

```
always_ff @(posedge clk, negedge rst_n) begin
 if (!rst_n) begin
 counter_bin <= 'b0;
 counter_gray <= 'b0;
 end else begin
 counter_bin <= w_counter_bin;
 counter_gray <= w_counter_gray;
 end
end
```

Note what this buys you over a `bin2gray` plus a separate flop: the Gray output is registered from the *next-state* value in the same process, so the encoding transient never leaves the module.

## 5.5 Timing Characteristics

Three paths matter:

1. **Binary increment** – standard binary addition timing
2. **Gray conversion** – a single XOR level: `gray = bin ^ (bin >> 1)` gives every output bit one 2-input XOR, so the delay is constant in WIDTH. (The  $\log(\text{WIDTH})$  XOR depth belongs to the *decode* direction, `gray2bin`.)
3. **Combined path** – increment + conversion in the same cycle

Clock-to-Q is the standard flip-flop delay, the combinational piece depends on the adder and XOR depth, and your setup check has to absorb the longest of them. Expect 200-400 MHz in a modern FPGA as-is; pipeline the Gray conversion if you need more.

## 5.6 Waveforms

WaveDrom timing diagrams for the Binary-Gray counter are available.

Four scenarios walk through the dual-output counter design:

### 5.6.1 Scenario 1: Binary vs Gray Code Comparison




#### *BinGray Comparison*

WaveJSON: [bingray\\_counter\\_binary\\_vs\\_gray.json](#)

Side-by-side comparison of both outputs: - Binary: Normal sequential counting (0→1→2→3→...) - Gray: Single-bit transitions between values - Shows full cycle demonstrating encoding differences - Illustrates why Gray code is CDC-safe

### 5.6.2 Scenario 2: Single-Bit Transitions (CDC Safety) KEY FEATURE




#### *BinGray Single-Bit Transitions*

WaveJSON: [bingray\\_counter\\_single\\_bit\\_transitions.json](#)

Gray code CDC safety property: - Each Gray transition changes EXACTLY one bit - Hamming distance = 1 between adjacent values - **Critical for fifo\_async CDC mechanism** - Prevents metastability in clock domain crossing

### 5.6.3 Scenario 3: Lookahead Signal (counter\_bin\_next)




#### *BinGray Lookahead*

WaveJSON: [bingray\\_counter\\_lookahead.json](#)

Combinational lookahead feature: - Predicts next binary value one cycle ahead - Useful for FIFO full/empty prediction - Shows enable gating behavior - Enables early decision-making

#### 5.6.4 Scenario 4: Enable and Reset Control



##### *BinGray Enable and Reset*

WaveJSON: [bingray\\_counter\\_enable\\_reset.json](#)

Control signal behavior: - Both outputs hold when enable=0 - Asynchronous reset to zero - Clean state transitions - Reset during counting demonstration

To regenerate these waveforms:

```
pytest val/cdc/test_counter_bingray_wavedrom.py -v
Then convert JSON to PNG:
cd docs/markdown/assets/WAVES/counter_bingray
for f in *.json; do wavedrom-cli -i "$f" -p "${f%.json}.png"; done
```

##### What Makes Binary-Gray Counters Special:

The waveforms put the unique dual-output design on display: - **Dual Encoding:** Binary for arithmetic, Gray for CDC safety - **Single-Bit Transitions:** Gray code changes one bit at a time - **Lookahead:** counter\_bin\_next provides one-cycle prediction - **CDC-Safe:** Safe for asynchronous clock domain crossing

##### Relationship to fifo\_async:

Binary-Gray counters are the foundation of the standard fifo\_async module: - **fifo\_async** uses this counter for read/write pointers - Gray outputs cross clock domains safely - Binary outputs used for arithmetic (occupancy, address generation) - Logarithmic width ( $\lceil \log_2(\text{DEPTH}) \rceil + 1$ ) for resource efficiency

##### Comparison Table:

| Feature                    | BinGray Counter                               | Johnson Counter         |
|----------------------------|-----------------------------------------------|-------------------------|
| <b>Output Width</b>        | $\lceil \log_2(\text{DEPTH}) \rceil + 1$ bits | DEPTH bits              |
| <b>CDC Method</b>          | Standard Gray code                            | Johnson code            |
| <b>Depth Support</b>       | Power-of-2 only                               | Any depth, odd included |
| <b>Resource Efficiency</b> | Logarithmic (better for large depths)         | Linear                  |

| Feature           | BinGray Counter          | Johnson Counter              |
|-------------------|--------------------------|------------------------------|
| <b>Conversion</b> | XOR tree (simple)        | Position detection (complex) |
| <b>Used In</b>    | fifo_async USE_JOHNSON=0 | fifo_async USE_JOHNSON=1     |

#### Comparison with Other Modules:

- test\_counter\_johnson\_wavedrom.py - Johnson counter (any depth, linear width)
- test\_fifo\_async\_wavedrom.py - BinGray counter in action (async FIFO, power-of-2)

## 5.7 Usage Examples

Every parameter and port below is read from the module declaration.

```
counter_bingray #(
 .WIDTH (4)
) u_counter_bingray (
 .clk (clk),
 .rst_n (rst_n),
 .enable (enable),
 .counter_bin (counter_bin),
 .counter_bin_next (counter_bin_next),
 .counter_gray (counter_gray)
);
```

## 5.8 Design Notes

Resource cost:  $2 \times \text{WIDTH}$  flip-flops (binary and Gray registers are separate), plus the increment logic and the single-level XOR for the Gray conversion. The critical path runs through the binary increment and the Gray conversion in series.

```
// Optional: Pipeline Gray conversion for high speed
logic [WIDTH-1:0] counter_gray_pipe;
always_ff @(posedge clk) begin
 counter_gray_pipe <= w_counter_gray;
end
```

Mark the Gray register as an async-crossing register so the placer keeps the sync chain tight:

```
(* ASYNC_REG = "TRUE" *) logic [WIDTH-1:0] counter_gray; // Xilinx
// synthesis attribute ASYNC_REG of counter_gray is "TRUE" // Quartus
```

Dynamic power scales with switching activity; static is minimal. Gate the clock with enable if the counter sits idle for long stretches.

### 5.8.1 Common mistakes

1. **Metastability** – the Gray value must go through a proper synchronizer in the destination domain; the encoding alone doesn't sample for you.
2. **Timing violations** – pipeline the Gray conversion if the increment+XOR path doesn't close.
3. **Incorrect FIFO status** – check the Gray-to-binary decode on the synchronized pointer before blaming the flags.
4. **Reset skew** – use proper reset synchronization in each domain.

For debug: assert the Gray properties in simulation, capture real transitions on a logic analyzer, run timing analysis on the crossing paths, and verify the synchronizer stage count against your MTBF target.

## 5.9 Related Modules

---

- **counter\_johnson**: the non-power-of-2 alternative – any depth, linear width
- **bin2gray** / **gray2bin**: the standalone converters, if you ever need the encode or decode alone
- **glitch\_free\_n\_dff\_arn**: the multi-flop synchronizer the Gray output crosses through
- **fifo\_async** / **gaxi\_fifo\_async**: the modules this counter exists to serve

## 5.10 Testing

---

### 5.10.1 Test scenarios

1. **Sequential counting**: both outputs increment correctly
2. **Gray code properties**: single-bit changes between adjacent values
3. **Reset behavior**: both outputs reset to zero
4. **Enable control**: hold behavior when disabled
5. **Rollover**: clean wrap-around from maximum value

### 5.10.2 Coverage

```
covergroup counter_bingray_cg @(posedge clk);
 cp_binary: coverpoint counter_bin {
 bins all_values[] = {[0:2**WIDTH-1]};
 }

 cp_gray: coverpoint counter_gray {
 bins all_values[] = {[0:2**WIDTH-1]};
 }
}
```

```

cp_enable: coverpoint enable {
 bins enabled = {1};
 bins disabled = {0};
}

cp_transitions: coverpoint counter_bin {
 bins transitions[] = ([0:2**WIDTH-2] ==> [1:2**WIDTH-1]);
 bins rollover = (2**WIDTH-1 ==> 0);
}
endcovergroup

```

### 5.10.3 Assertions

```

// Verify Gray code has single bit changes
property gray_single_bit_change;
 @(posedge clk) disable iff (!rst_n)
 enable && !$isunknown(counter_gray) |->
 $countones(counter_gray ^ $past(counter_gray)) <= 1;
endproperty

```

```

assert property (gray_single_bit_change);

```

```

// Verify binary and Gray relationship
property bin_gray_relationship;
 @(posedge clk) disable iff (!rst_n)
 counter_gray == (counter_bin ^ (counter_bin >> 1));
endproperty

```

```

assert property (bin_gray_relationship);

```

### 5.10.4 Test files

- val/cdc/test\_counter\_bingray.py – full functional verification
- val/cdc/test\_counter\_bingray\_wavedrom.py – WaveDrom timing diagrams

```

Full functional test (basic/medium/full levels)

```

```

pytest val/cdc/test_counter_bingray.py -v

```

```

WaveDrom waveform generation

```

```

pytest val/cdc/test_counter_bingray_wavedrom.py -v

```

## 5.11 Navigation

---

- [← Back to CDC Index](#)
- [← Back to Main Documentation Index](#)

## 6 Johnson counter (counter\_johnson.sv)

### 6.1 Overview

counter\_johnson implements a Johnson counter – also called a twisted-ring counter or a switch-tail counter. It's a shift register with a twist: the inverted output of the last stage feeds back into the first, which gives you a sequence with  $2 \times \text{WIDTH}$  unique states. You'll see these generating multi-phase clock signals, driving state machine control, and running sequential timing applications – and, in this library, serving as the non-power-of-2 FIFO pointer behind `USE_JOHNSON=1`.

```
module counter_johnson #(
 parameter int WIDTH = 4
) (
 input logic clk,
 input logic rst_n,
 input logic enable,
 output logic [WIDTH - 1:0] counter_gray
);
```

### 6.2 Parameters

*counter\_johnson parameters*

| Parameter | Default | Description                                                                                                                                                        |
|-----------|---------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| WIDTH     | 4       | Number of stages in the Johnson counter (int, $\geq 2$ ). Generates $2 \times \text{WIDTH}$ unique states; determines sequence length and number of output phases. |

The  $\geq 2$  bound is real, not a style preference. The shift expression is `counter_gray <= {counter_gray[WIDTH-2:0], ~counter_gray[WIDTH-1]}`, whose `[WIDTH-2:0]` part-select is reversed at `WIDTH=1` and fails elaboration. A 1-stage Johnson counter has no use anyway.

### 6.3 Ports

*counter\_johnson ports*

| Port  | Direction | Width | Description        |
|-------|-----------|-------|--------------------|
| clk   | input     | 1     | System clock input |
| rst_n | input     | 1     | Active-low         |

| Port         | Direction | Width | Description                     |
|--------------|-----------|-------|---------------------------------|
|              |           |       | asynchronous reset              |
| enable       | input     | 1     | Counter enable control          |
| counter_gray | output    | WIDTH | Johnson counter output register |

## 6.4 Functional Description

A Johnson counter is a shift register with inverted feedback:

- **Normal operation:** bits shift left (or right) each clock cycle
- **Feedback:** the complement of the MSB feeds back to the LSB
- **Sequence length:**  $2 \times \text{WIDTH}$  states before repeating

Mathematically, for a WIDTH-bit counter:

```
Next_State[i] = Current_State[i-1] for i = 1 to WIDTH-1
Next_State[0] = ~Current_State[WIDTH-1]
```

And here is the entire implementation:

```
always_ff @(posedge clk or negedge rst_n) begin
 if (!rst_n)
 counter_gray <= {WIDTH{1'b0}}; // Reset to all zeros
 else if (enable) begin
 counter_gray <= {counter_gray[WIDTH-2:0], ~counter_gray[WIDTH-1]};
 end
end
```

Four lines of actual logic: shift the low bits up, feed the inverted MSB into the LSB, hold when enable is low, clear on reset.

### 6.4.1 State sequences

#### 6.4.1.1 4-bit Johnson Counter (WIDTH=4)

| Step | counter_gray | Decimal | Description      |
|------|--------------|---------|------------------|
| 0    | 0000         | 0       | Reset state      |
| 1    | 0001         | 1       | First '1' enters |
| 2    | 0011         | 3       | '1' propagates   |

| Step | counter_gray | Decimal | Description      |
|------|--------------|---------|------------------|
| 3    | 0111         | 7       | '1' propagates   |
| 4    | 1111         | 15      | All '1's         |
| 5    | 1110         | 14      | First '0' enters |
| 6    | 1100         | 12      | '0' propagates   |
| 7    | 1000         | 8       | '0' propagates   |
| 8    | 0000         | 0       | Back to start    |

#### 6.4.1.2 3-bit Johnson Counter (WIDTH=3)

| Step | counter_gray | Binary Pattern |
|------|--------------|----------------|
| 0    | 000          | All zeros      |
| 1    | 001          | One bit set    |
| 2    | 011          | Two bits set   |
| 3    | 111          | All ones       |
| 4    | 110          | Two bits set   |
| 5    | 100          | One bit set    |
| 6    | 000          | Cycle repeats  |

#### 6.4.1.3 2-bit Johnson Counter (WIDTH=2)

| Step | counter_gray | State Name |
|------|--------------|------------|
| 0    | 00           | State A    |
| 1    | 01           | State B    |
| 2    | 11           | State C    |
| 3    | 10           | State D    |
| 4    | 00           | Back to A  |

### 6.4.2 NOT self-starting – reset is mandatory

This RTL is a bare twisted-ring shift register ( $\text{counter\_gray} \leftarrow \{\text{counter\_gray}[\text{WIDTH}-2:0], \sim\text{counter\_gray}[\text{WIDTH}-1]\}$ ) with **no state-correction logic**, so it is **not** self-starting:

- **Valid states:**  $2 \times \text{WIDTH}$  states in the normal sequence
- **Invalid states:**  $(2^{\text{WIDTH}} - 2 \times \text{WIDTH})$  states not in sequence
- **Recovery: none.** The invalid states form their own closed cycle and never converge to the valid sequence. For WIDTH=4 the 8 invalid states cycle 0010→0101→1011→0110→1101→1010→0100→1001→0010 forever; for WIDTH=3

the pair 010↔101 is a parasitic 2-cycle. **You must guarantee a clean reset** – after a glitch or SEU the counter can lock permanently into the parasitic cycle. If self-correction is required, add explicit decode/recovery logic (not present here).

### 6.4.3 Symmetry

The sequence has inherent symmetry: the first half fills with '1's (0→1→11→111...), the second half flushes them back out with '0's (111→110→100→000...). Each bit effectively represents a different phase of the cycle.

### 6.4.4 Single-bit transitions – but Johnson is not Gray

Johnson counter outputs share the one-bit-change property that makes Gray code safe to synchronize – consecutive states differ by exactly **one** bit, including the wrap from the last state back to state 0 – with predictable progression and easy state decode. But they are **not** Gray code:

**Johnson is not Gray code.** Johnson achieves single-bit transitions through sequential shift-register operation; Gray code achieves them through a mathematical encoding ( $\text{gray} = \text{bin} \oplus (\text{bin} \gg 1)$ ). Consequences: - Johnson yields  $2N$  states from  $N$  bits; Gray yields  $2^N$ . - Johnson is safe only for a counter stepping through its own sequence. It cannot encode an arbitrary value, and there is no closed-form binary conversion – decoding needs `johnson2bin`, which detects position. - Do not substitute one for the other in general-purpose CDC. For an arbitrary binary value crossing domains, use Gray (`bin2gray/gray2bin`).

## 6.5 Timing Characteristics

---

The logic is minimal – a shift register and an inverter – so there are no complex timing requirements. Typical maximum frequency is 400-600 MHz in modern FPGAs, limited by the shift-register timing. For extreme speeds, pipeline:

```
// For high-speed applications, consider pipeline stages
logic [WIDTH-1:0] johnson_pipe;
always_ff @(posedge clk) begin
 johnson_pipe <= {counter_gray[WIDTH-2:0], ~counter_gray[WIDTH-1]};
end
```


## 6.6 Waveforms

---

WaveDrom timing diagrams for the Johnson counter are available.

Four scenarios demonstrate the unique properties of Johnson counters:

### 6.6.1 Scenario 1: Walking Ones and Walking Zeros Pattern




#### Johnson Walking Pattern

WaveJSON: [johnson\\_counter\\_walking\\_pattern.json](#)

Complete  $2 \times \text{WIDTH}$  state cycle (8 states for  $\text{WIDTH}=4$ ): - Walking ones: 0000 → 0001 → 0011 → 0111 → 1111 - Walking zeros: 1111 → 1110 → 1100 → 1000 → 0000 - Demonstrates predictable sequential progression

### 6.6.2 Scenario 2: Single-Bit Transitions (CDC Safety) KEY FEATURE




#### Johnson Single-Bit Transitions

WaveJSON: [johnson\\_counter\\_single\\_bit\\_transitions.json](#)

Each transition changes only ONE bit: - CDC-safe like Gray codes - **Critical for the fifo\_async USE\_JOHNSON=1 CDC mechanism** - Prevents metastability in clock domain crossing - Hamming distance = 1 between all adjacent states

### 6.6.3 Scenario 3: Enable Control




#### Johnson Enable Control

WaveJSON: [johnson\\_counter\\_enable\\_control.json](#)

Enable control and state holding: - Counter advances when enable=1 - Counter holds state when enable=0 - Clean enable/disable transitions - Demonstrates conditional counting

### 6.6.4 Scenario 4: Reset Behavior



#### Johnson Reset Behavior

## WaveJSON: [johnson\\_counter\\_reset\\_behavior.json](#)

Reset and initialization: - Asynchronous reset to all zeros (0000) - Immediate reset effect - Clean restart from reset state - Reset during counting operation

---

### To regenerate these waveforms:

```
pytest val/cdc/test_counter_johnson_wavedrom.py -v
Then convert JSON to PNG:
cd docs/markdown/assets/WAVES/counter_johnson
for f in *.json; do wavedrom-cli -i "$f" -p "${f%.json}.png"; done
```

### What Makes Johnson Counters Special:

The waveforms highlight the unique properties that make Johnson counters useful: - **Single-Bit Transitions**: Only one bit changes per state, making them CDC-safe - **2×WIDTH States**: More efficient than one-hot (N states with N/2 flip-flops) - **Walking Pattern**: Natural “fill then empty” sequence useful for visual effects - **Reset-dependent**: NOT self-starting – a valid state is only guaranteed by reset

### Relationship to fifo\_async (USE\_JOHNSON=1):

Johnson counters are the foundation of the fifo\_async USE\_JOHNSON=1 CDC mechanism: - **fifo\_async with USE\_JOHNSON=1** uses Johnson counters for pointer synchronization - Single-bit transitions enable safe clock domain crossing - Linear width scaling (DEPTH bits) allows any depth, odd included

### Comparison with Other Counters:

- test\_counter\_bingray\_wavedrom.py - Binary-Gray counter (power-of-2 depths, logarithmic width)
- test\_fifo\_async\_wavedrom.py - Gray code in action (async FIFO)

## 6.7 Usage Examples

---

### 6.7.1 Example 1: 4-phase clock generator

```
counter_johnson #(
 .WIDTH(4)
) phase_gen (
 .clk(master_clk),
 .rst_n(rst_n),
 .enable(1'b1),
 .counter_gray(phases)
);
```

```

// Decode individual phases
assign phase_0 = (phases == 4'b0001); // Step 1
assign phase_1 = (phases == 4'b0011); // Step 2
assign phase_2 = (phases == 4'b0111); // Step 3
assign phase_3 = (phases == 4'b1111); // Step 4

```

### 6.7.2 Example 2: LED chaser

```

counter_johnson #(
 .WIDTH(8)
) led_chaser (
 .clk(slow_clk), // ~1Hz clock
 .rst_n(rst_n),
 .enable(chase_enable),
 .counter_gray(led_pattern)
);

// LEDs show walking pattern
assign leds[7:0] = led_pattern;

```

### 6.7.3 Example 3: state machine controller

```

counter_johnson #(
 .WIDTH(3)
) state_ctrl (
 .clk(clk),
 .rst_n(rst_n),
 .enable(state_advance),
 .counter_gray(state_code)
);

// Decode states for control logic
always_comb begin
 case (state_code)
 3'b001: current_state = INIT;
 3'b011: current_state = LOAD;
 3'b111: current_state = PROCESS;
 3'b110: current_state = STORE;
 3'b100: current_state = CLEANUP;
 3'b000: current_state = IDLE;
 default: current_state = ERROR;
 endcase
end

```

### 6.7.4 Example 4: quadrature signal generation

```

// Generate quadrature encoder-like signals
counter_johnson #(
 .WIDTH(4)
) quad_gen (

```

```

 .clk(position_clk),
 .rst_n(rst_n),
 .enable(motor_moving),
 .counter_gray(quad_state)
);

 // Extract A and B channels with 90° phase shift
 assign quad_a = quad_state[0] ^ quad_state[1];
 assign quad_b = quad_state[1] ^ quad_state[2];

```

### 6.7.5 Multi-phase clock for pipeline stages

```

 // Multi-phase clock for pipeline stages
 counter_johnson #(.WIDTH(4)) pipeline_phases (
 .clk(master_clk),
 .enable(pipeline_enable),
 .counter_gray(phase_vector)
);

 // Generate phase-shifted clocks
 assign clk_ph0 = master_clk & phase_vector[0];
 assign clk_ph1 = master_clk & phase_vector[1];
 assign clk_ph2 = master_clk & phase_vector[2];
 assign clk_ph3 = master_clk & phase_vector[3];

```

### 6.7.6 Sequential control (memory refresh)

```

 // Memory refresh controller
 counter_johnson #(.WIDTH(3)) refresh_ctrl (
 .clk(clk),
 .enable(refresh_request),
 .counter_gray(refresh_phase)
);

 // Control refresh operations
 assign precharge = (refresh_phase == 3'b001);
 assign activate = (refresh_phase == 3'b011);
 assign refresh = (refresh_phase == 3'b111);
 assign restore = (refresh_phase == 3'b110);

```

### 6.7.7 Pattern generation for BIST

```

 // Test pattern generator
 counter_johnson #(.WIDTH(8)) pattern_gen (
 .clk(test_clk),
 .enable(pattern_enable),
 .counter_gray(test_pattern)
);

```

```
// Use for built-in self-test (BIST)
assign test_data = test_pattern;
```

## 6.8 Design Notes

---

**What it does well.** The logic is minimal – a shift register and an inverter – so states can be used directly with no decode logic and no complex timing requirements. It runs fast (400-600 MHz typical in modern FPGAs, limited by the shift-register timing), it's power-friendly because only one bit changes per cycle, and the progression is completely predictable. Any width works, every bit gives you a different phase, and the parameterization is trivial.

**What it costs.** Count efficiency:  $2 \times \text{WIDTH}$  states versus  $2^{\text{WIDTH}}$  for a binary counter of the same flops. The progression is a fixed sequence – no random access, you step through it. Decoding specific states takes comparators, the output isn't binary-compatible without `johnson2bin`, and you need care to avoid state conflicts when overlapping decoded phases. And the big one, from the theory section: it is **not** self-starting, so a clean reset is a hard requirement and glitch-free operation is only guaranteed inside the valid sequence.

**Synthesis.** Resources: `WIDTH` flip-flops, almost no LUTs (shift and invert only), and simple routing between adjacent stages. Keep the tool from swallowing the register into an SRL – the feedback path needs real flops:

```
(* SHREG_EXTRACT = "NO" *) logic [WIDTH-1:0] counter_gray; // Prevent
SRL inference
```

Dynamic power is low – one bit changes per cycle – and `enable` gives you free clock gating.

## 6.9 Related Modules

---

- `counter_ring`: different shift register pattern
- `counter_bingray`: binary and Gray code counter
- `counter`: simple binary counter
- Standard shift register implementations

## 6.10 Testing

---

### 6.10.1 Functional tests

1. **Sequence verification:** check the complete  $2 \times \text{WIDTH}$  state cycle
2. **Reset behavior:** verify initialization to all zeros
3. **Enable control:** hold behavior when disabled
4. **Invalid-state containment:** force an invalid state and confirm it stays trapped in the parasitic cycle (this counter does NOT auto-recover), so the design's reset strategy is the only guarantee of a valid state

## 6.10.2 Coverage

```
covergroup johnson_cg @(posedge clk);
 cp_states: coverpoint counter_gray {
 // Valid Johnson counter states for WIDTH=4
 bins valid_states[] = {4'b0000, 4'b0001, 4'b0011, 4'b0111,
 4'b1111, 4'b1110, 4'b1100, 4'b1000};
 bins invalid_states[] = default;
 }

 cp_transitions: coverpoint counter_gray {
 bins valid_trans[] = (4'b0000 => 4'b0001),
 (4'b0001 => 4'b0011),
 (4'b0011 => 4'b0111),
 (4'b0111 => 4'b1111),
 (4'b1111 => 4'b1110),
 (4'b1110 => 4'b1100),
 (4'b1100 => 4'b1000),
 (4'b1000 => 4'b0000);
 }

 cp_enable: coverpoint enable {
 bins enabled = {1};
 bins disabled = {0};
 }
endcovergroup
```

## 6.10.3 Assertions

```
// Verify valid state progression
property johnson_sequence;
 @(posedge clk) disable iff (!rst_n)
 enable && (counter_gray == 4'b0000) | =>
 counter_gray == 4'b0001;
endproperty

assert property (johnson_sequence);

// Check for invalid states
property no_invalid_states;
 @(posedge clk) disable iff (!rst_n)
 counter_gray inside {4'b0000, 4'b0001, 4'b0011, 4'b0111,
 4'b1111, 4'b1110, 4'b1100, 4'b1000};
endproperty

assert property (no_invalid_states);
```

#### 6.10.4 Test files

- `val/cdc/test_counter_johnson.py` – full functional verification
- `val/cdc/test_counter_johnson_wavedrom.py` – WaveDrom timing diagrams

*# Full functional test (basic/medium/full levels)*

```
pytest val/cdc/test_counter_johnson.py -v
```

*# WaveDrom waveform generation*

```
pytest val/cdc/test_counter_johnson_wavedrom.py -v
```

### 6.11 Navigation

---

- [← Back to CDC Index](#)
- [← Back to Main Documentation Index](#)

## 7 Asynchronous FIFO (fifo\_async.sv)

### 7.1 Overview

An asynchronous FIFO for crossing data safely between different clock domains. **Restricted to power-of-2 depths only** – that restriction comes from the Gray code pointer implementation, and USE\_JOHNSON=1 is the supported way around it.

### 7.2 Parameters

*fifo\_async parameters*

| Parameter        | Default | Description                                                                                                                                                                       |
|------------------|---------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| MEM_STYLE        | —       | Memory implementation (FIFO_AUTO/SRL/BRAM). The BRAM branch registers the read path (registered read even when REGISTERED=0).                                                     |
| REGISTERED       | —       | 0 = mux mode, 1 = flop mode for read output                                                                                                                                       |
| DATA_WIDTH       | 8       | Width of data                                                                                                                                                                     |
| DEPTH            | 16      | FIFO depth. Power-of-2 with the default binary Gray pointers; set <b>USE_JOHNSON=1</b> for <b>non-power-of-2</b> depths (Johnson-coded pointers via counter_johnson/johnson2bin). |
| USE_JOHNSON      | 0       | 0 = Gray pointers (power-of-2 depth), 1 = Johnson pointers (arbitrary depth). This is the supported route to non-power-of-2 depth (the old fifo_async_div2 module was retired).   |
| N_FLOP_CROSS     | 2       | Number of synchronizer stages                                                                                                                                                     |
| ALMOST_WR_MARGIN | 1       | Almost full threshold                                                                                                                                                             |

| Parameter        | Default | Description            |
|------------------|---------|------------------------|
| ALMOST_RD_MARGIN | 1       | Almost empty threshold |

## 7.3 Ports

### *fifo\_async* ports

| Port            | Direction | Width      | Description                   |
|-----------------|-----------|------------|-------------------------------|
| wr_clk          | input     | 1          | Write domain clock            |
| wr_rst_n        | input     | 1          | Write domain active-low reset |
| write           | input     | 1          | Write enable signal           |
| wr_data         | input     | DATA_WIDTH | Data to write                 |
| wr_full         | output    | 1          | Write domain full flag        |
| wr_almost_full  | output    | 1          | Write domain almost full flag |
| rd_clk          | input     | 1          | Read domain clock             |
| rd_rst_n        | input     | 1          | Read domain active-low reset  |
| read            | input     | 1          | Read enable signal            |
| rd_data         | output    | DATA_WIDTH | Data read from FIFO           |
| rd_empty        | output    | 1          | Read domain empty flag        |
| rd_almost_empty | output    | 1          | Read domain almost empty flag |

## 7.4 Functional Description

### 7.4.1 Clock domain crossing strategy

The FIFO uses **Gray code pointers** for safe clock domain crossing:




Diagram 2

### 7.4.2 Core components

1. **Binary-Gray counters** (counter\_bingray) for pointer generation
2. **Multi-stage synchronizers** (glitch\_free\_n\_dff\_arn) for CDC
3. **Gray-to-binary converters** (gray2bin) for pointer comparison
4. **Shared memory array** accessible from both domains
5. **FIFO control logic** for status flag generation

### 7.4.3 Gray code pointer system

Gray codes ensure **only one bit changes** per increment:

- **Binary:** 011 → 100 (3 bits change simultaneously)
- **Gray:** 010 → 110 (only 1 bit changes)
- **Benefit:** eliminates metastability from multi-bit transitions

```
// Write domain Gray counter
counter_bingray #(.WIDTH(AW + 1)) wr_ptr_counter_gray (
 .clk(wr_clk),
 .rst_n(wr_rst_n),
 .enable(write && !wr_full),
 .counter_bin(r_wr_ptr_bin), // Binary for addressing
 .counter_bin_next(w_wr_ptr_bin_next), // Next binary value
 .counter_gray(r_wr_ptr_gray) // Gray for CDC
);
```

Pointer width:  $AW = \lceil \log_2(DEPTH) \rceil$  address bits, plus one extra bit for wrap detection – so  $DEPTH=16$  gives  $AW=4$  and a 5-bit pointer.

#### 7.4.4 Clock domain crossing

```
// Cross read pointer to write domain
glitch_free_n_dff_arn #(
 .FLOP_COUNT(N_FLOP_CROSS),
 .WIDTH(AW + 1)
) rd_ptr_gray_cross_inst (
 .q(r_wdom_rd_ptr_gray), // Synchronized output
 .d(r_rd_ptr_gray), // Gray pointer input
 .clk(wr_clk), // Destination clock
 .rst_n(wr_rst_n) // Destination reset
);
```

The default is 2 flip-flops per crossing (N\_FLOP\_CROSS=2). Each additional stage improves MTBF at the cost of latency – the classic trade.

#### 7.4.5 Memory organization

```
logic [DW-1:0] mem[DEPTH]; // Memory array -- sized by DEPTH, not by
 // 1<=AW: with USE_JOHNSON=1 and a
 // non-power-of-2 DEPTH those differ

// Write port (write domain)
always_ff @(posedge wr_clk) begin
 if (write && !wr_full) begin // !wr_full guard REQUIRED (prevents
 mem[r_wr_addr] <= wr_data;
 end
end

// Read port (read domain) -- dual mode, like fifo_sync:
if (REGISTERED != 0) begin : g_flop
 always_ff ... r_rd_data <= mem[r_rd_addr];
 assign rd_data = r_rd_data;
end else begin : g_mux
 assign rd_data = mem[r_rd_addr];
end
```

Addresses come from the binary pointers, truncated to the low bits: r\_wr\_addr = r\_wr\_ptr\_bin[AW-1:0] and r\_rd\_addr = r\_rd\_ptr\_bin[AW-1:0].

#### 7.4.6 Full/empty detection

```
// In write domain (fifo_control): computed combinatorially,
REGISTERED out
assign w_wr_full_d = (w_wdom_ptr_xor &&
 (wr_ptr_bin[AW-1:0] == wdom_rd_ptr_bin[AW-1:0]));

always_ff @(posedge wr_clk, negedge wr_rst_n) begin
```

```

 if (!wr_rst_n) wr_full <= 'b0;
 else wr_full <= w_wr_full_d;
end

// Where:
assign w_wdom_ptr_xor = wr_ptr_bin[AW] ^ wdom_rd_ptr_bin[AW];

// In read domain: same shape, and rd_empty RESETS TO 1 (empty out of reset)
assign w_rd_empty_d = (!w_rdom_ptr_xor_for_empty &&
 (rd_ptr_bin[AW:0] == w_wr_ptr_for_empty[AW:0]));

always_ff @(posedge rd_clk, negedge rd_rst_n) begin
 if (!rd_rst_n) rd_empty <= 'b1;
 else rd_empty <= w_rd_empty_d;
end

```

Both flags are registered, so each lags its pointer comparison by one cycle of its own clock – conservative in the safe direction (a write that just made the FIFO non-empty shows up at the reader one `rd_clk` later).

The algorithm in words: **full** is MSBs differ AND LSBs equal (the write pointer has wrapped and caught the read pointer); **empty** is all bits equal (pointers at the same location). The MSB is the wrap bit – it records which pointer has lapped.

### 7.4.7 The power-of-2 requirement

Why power-of-2 only, with Gray pointers:

1. **Gray code properties:** natural binary-Gray relationship
2. **Wraparound behavior:** clean modulo- $2^n$  arithmetic
3. **Address truncation:** simple bit slicing for memory addressing
4. **Pointer comparison:** efficient full/empty detection

*// Valid depths: 2, 4, 8, 16, 32, 64, 128, 256, ...*  
*// Invalid depths: 3, 5, 6, 7, 9, 10, 12, 15, ...*

## 7.5 Timing Characteristics

- **Pointer propagation:** 2-3 clock cycles (depending on `N_FLOP_CROSS`)
- **Status flag delay:** flags reflect state with synchronizer latency
- **Conservative design:** prevents overflow/underflow despite the delay

Metastability protection comes from the usual pair: Gray code transitions (single bit changes only) and multi-stage synchronization, which reduces metastability failure probability exponentially (raises MTBF). Proper timing constraints on the crossing paths are essential – see the SDC section of the [CDC reference](#).

Throughput and latency at a glance:

- **Throughput:** up to 1 operation per clock per domain
- **Latency:** 0-1 cycles (depending on REGISTERED mode)
- **CDC latency:** 2-3 cycles for status propagation
- **Efficiency:** ~100% utilization possible

## 7.6 Waveforms

---

WaveDrom timing diagrams of the Gray code CDC mechanism are available.

Run the WaveDrom test to generate detailed timing diagrams:

```
Generate Gray code waveforms (standard async FIFO implementation)
pytest val/cdc/test_fifo_async_wavedrom.py -v
```

Waveform Scenarios Generated:

1. **Write-Fill-Read-Empty Cycle**
  - Standard async FIFO operation
  - Gray code pointer progression
  - Power-of-2 depth utilization
2. **Gray Code Synchronization**
  - Efficient Gray code CDC
  - Logarithmic pointer width (vs. linear for Johnson)
  - Cross-domain pointer transitions
3. **Power-of-2 Depth Utilization**
  - Efficient addressing with power-of-2 depths
  - Full depth wraparound behavior
  - Resource-efficient pointer management

Key Differences from Johnson Counter Variant:

- **Pointer Width:** Logarithmic ( $\lceil \log_2(\text{DEPTH}) \rceil + 1$ ) vs. linear (DEPTH) for Johnson
- **Depth Restriction:** Power-of-2 only vs. **any** depth for Johnson – odd depths included. The elaboration check only fires for `USE_JOHNSON == 0` (if `((USE_JOHNSON == 0) && ((DEPTH & (DEPTH - 1)) != 0)) $error(...)`), so Johnson mode has no restriction at all. (“Even only” is stale language from the retired `fifo_async_div2`.)

- **Resource Efficiency:** Better for large depths (>32) vs. Johnson's flexibility for small depths

#### Comparison Tests:

- `test_fifo_sync_wavedrom.py` - Synchronous FIFO (single clock, no CDC)

## 7.7 Usage Examples

---

Typical applications:

- **Video processing:** different pixel and memory clock domains
- **Networking:** packet buffers between different rate domains
- **Audio systems:** sample rate conversion buffers
- **Microprocessor interfaces:** CPU and peripheral clock domains

When to use what:

- **Async FIFO** when the clock domains are truly independent
- **Sync FIFO** when a single clock domain is sufficient
- **Non-power-of-2 depth:** use `fifo_async #( .USE_JOHNSON(1) )` – the standalone `fifo_async_div2` module was retired

## 7.8 Design Notes

---

### 7.8.1 Depth sizing

```
// Calculate minimum depth for burst handling
// DEPTH ≥ burst_size + synchronizer_latency + margin
parameter int MIN_DEPTH = 16; // Typical minimum for async FIFOs
```

### 7.8.2 Clock relationship

- **Asynchronous clocks:** no phase relationship assumed
- **Clock gating:** avoid gating clocks used by the FIFO
- **Reset deassertion:** ensure proper reset release sequencing

### 7.8.3 Almost full/empty settings

- **Almost full:** `DEPTH - ALMOST_WR_MARGIN`
- **Almost empty:** `ALMOST_RD_MARGIN`
- **Guideline:** set margins > synchronizer latency

### 7.8.4 Error checking

The current RTL has **no** runtime `$display` overflow/underflow checks. The `!wr_full` write guard is the only overflow protection. The one elaboration-time check that does exist is an `$error` for a non-power-of-2 DEPTH when Gray pointers are selected (`USE_JOHNSON=0`). Add assertions in your testbench if you need overflow telemetry.

## 7.9 Related Modules

---

- **USE\_JOHNSON=1**: for non-power-of-2 depths, using Johnson counters (replaces the retired `fifo_async_div2`)
- **fifo\_sync**: for single clock domain applications
- **counter\_bingray**: binary-Gray counter implementation
- **glitch\_free\_n\_dff\_arn**: multi-stage synchronizer

## 7.10 Testing

---

- `val/cdc/test_fifo_buffer_async.py` – full functional verification
- `val/cdc/test_fifo_async_wavedrom.py` – WaveDrom timing diagrams

*# Full functional test (basic/medium/full levels)*

```
pytest val/cdc/test_fifo_buffer_async.py -v
```

*# WaveDrom waveform generation*

```
pytest val/cdc/test_fifo_async_wavedrom.py -v
```

## 7.11 Navigation

---

- [← Back to CDC Index](#)
- [← Back to Main Documentation Index](#)

## 8 GAXI asynchronous FIFO (gaxi\_fifo\_async.sv)

### 8.1 Overview

gaxi\_fifo\_async (in rtl/cdc/, production ready) moves GAXI data safely between independent clock domains. Gray code pointers and multi-flop synchronizers keep metastability out of the datapath and get your data across intact.

What you get:

- **Clock domain crossing:** safe transfer between independent clocks
- **Gray code pointers:** no multi-bit synchronization hazards
- **Configurable CDC stages:** 2-4 flop synchronizers (3 recommended; 4 for ultra-critical systems)
- **Arbitrary clock ratios:** works with any write:read clock ratio
- **Two read modes:** mux or flop mode

```
module gaxi_fifo_async #(
 parameter fifo_mem_t MEM_STYLE = FIFO_AUTO, // FIFO_AUTO /
 FIFO_SRL / FIFO_BRAM
 parameter int REGISTERED = 0, // 0 = mux
 mode, 1 = flop mode
 parameter int DATA_WIDTH = 8,
 parameter int DEPTH = 16,
 parameter int USE_JOHNSON = 0, // 0 = Gray
 (power-of-2 depth), 1 = Johnson (any depth)
 parameter int N_FLOP_CROSS = 2,
 parameter int ALMOST_WR_MARGIN = 1,
 parameter int ALMOST_RD_MARGIN = 1
) (
 // Write Domain
 input logic axi_wr_aclk,
 input logic axi_wr_aresetn,
 input logic wr_valid,
 output logic wr_ready, // not full
 input logic [DW-1:0] wr_data,

 // Read Domain
 input logic axi_rd_aclk,
 input logic axi_rd_aresetn,
 input logic rd_ready,
 output logic rd_valid, // not empty
 output logic [DW-1:0] rd_data
);
```

## 8.2 Parameters

*gaxi\_fifo\_async* parameters

| Parameter        | Default   | Description                                                                                                                                                                                                              |
|------------------|-----------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| MEM_STYLE        | FIFO_AUTO | Memory implementation: FIFO_AUTO / FIFO_SRL / FIFO_BRAM. The BRAM branch registers the read path – a registered read even when REGISTERED=0                                                                              |
| REGISTERED       | 0         | 0=mux mode, 1=flop mode (read path)                                                                                                                                                                                      |
| DATA_WIDTH       | 8         | Data bus width                                                                                                                                                                                                           |
| DEPTH            | 16        | FIFO depth. Power of 2 with Gray pointers; <b>any</b> depth – odd included – with USE_JOHNSON=1                                                                                                                          |
| USE_JOHNSON      | 0         | Pointer CDC encoding: 0 = Gray ( $\log_2(\text{DEPTH})+1$ bits, power-of-2 depth only), 1 = Johnson (DEPTH bits, <b>any</b> depth – odd included). An illegal combination fails at elaboration with an explicit \$error. |
| N_FLOP_CROSS     | 2         | Synchronizer stages (3 recommended for safety)                                                                                                                                                                           |
| ALMOST_WR_MARGIN | 1         | Almost full threshold                                                                                                                                                                                                    |
| ALMOST_RD_MARGIN | 1         | Almost empty threshold                                                                                                                                                                                                   |

**Important:** Set N\_FLOP\_CROSS=3 for production designs to ensure metastability protection.

The RTL also declares five derived parameters after these – DW, D, AW, JCW, N. They are aliases (DW = DATA\_WIDTH, AW =  $\lceil \log_2(\text{DEPTH}) \rceil$ , JCW = D, N = N\_FLOP\_CROSS), not independent knobs. SystemVerilog lets you override them, but doing so decouples a port width from the storage that backs it: override DW alone and the ports widen while the memory stays narrow, and the FIFO silently drops the upper bits of every entry. Set DATA\_WIDTH, not DW.

## 8.3 Ports

---

*gaxi\_fifo\_async* ports

| Port           | Direction | Width | Description                   |
|----------------|-----------|-------|-------------------------------|
| axi_wr_aclk    | input     | 1     | Write domain clock            |
| axi_wr_aresetn | input     | 1     | Write domain active-low reset |
| wr_valid       | input     | 1     | Write valid                   |
| wr_ready       | output    | 1     | Write ready (not full)        |
| wr_data        | input     | DW    | Write data                    |
| axi_rd_aclk    | input     | 1     | Read domain clock             |
| axi_rd_aresetn | input     | 1     | Read domain active-low reset  |
| rd_ready       | input     | 1     | Read ready                    |
| rd_valid       | output    | 1     | Read valid (not empty)        |
| rd_data        | output    | DW    | Read data                     |

## 8.4 Functional Description

---

### 8.4.1 Why Gray code?

Gray code ensures only one bit changes at a time during pointer updates:

Binary: 011 → 100 (3 bits change - hazard!)

Gray: 010 → 110 (1 bit changes - safe!)


This prevents glitches during synchronization across clock domains.

### 8.4.2 Dependencies

- `counter_bin.sv` - Binary counters
- `counter_johnson.sv` - Johnson (twisted-ring) counters, used when `USE_JOHNSON=1`. Johnson code is NOT Gray code – it is a distinct encoding

that happens to share the single-bit-change property needed for CDC. Both change exactly one bit per increment, including the wrap.

- `glitch_free_n_dff_arn.sv` - Multi-flop synchronizers
- `johnson2bin.sv` - Johnson-to-binary conversion (**combinational**; its `clk/rst_n` ports are declared but unused), used when `USE_JOHNSON=1`
- `counter_bingray.sv` / `gray2bin.sv` - Gray pointer path (combinational decode), used when `USE_JOHNSON=0`
- `fifo_control.sv` - Full/empty flag generation



## Implementation

Each domain keeps its own binary pointer for arithmetic, converts it to Gray for the crossing, and decodes the synchronized remote pointer for flag generation.

## 8.5 Timing Characteristics

*gaxi\_fifo\_async latency by clock ratio*

| Clock Ratio (wr:rd) | Latency    | Notes                                                        |
|---------------------|------------|--------------------------------------------------------------|
| 1:1 (same freq)     | 3-5 cycles | CDC<br>synchronizat<br>ion overhead                          |
| 2:1 (fast→slow)     | 4-6 cycles | Additional<br>write-side<br>delay                            |
| 1:2 (slow→fast)     | 3-5 cycles | Read-side<br>samples<br>faster                               |
| Any ratio           | 3-7 cycles | Depends on<br>synchronizer<br>stages + clock<br>relationship |

**Latency formula:**  $\sim (2 \times N\_FLOP\_CROSS) + 1$  in slower clock domain cycles.

## 8.6 Usage Examples

### 8.6.1 Example 1: basic CDC FIFO

```
gaxi_fifo_async #(
 .DATA_WIDTH(32),
 .DEPTH(16),
 .N_FLOP_CROSS(3), // 3-stage synchronizer
 .REGISTERED(0)
) u_cdc_fifo (
 // Write domain @ 100 MHz
 .axi_wr_aclk (clk_100mhz),
 .axi_wr_aresetn (rst_100_n),
 .wr_valid (domain_a_valid),
 .wr_ready (domain_a_ready),
 .wr_data (domain_a_data),

 // Read domain @ 50 MHz
 .axi_rd_aclk (clk_50mhz),
```

```

 .axi_rd_aresetn (rst_50_n),
 .rd_ready (domain_b_ready),
 .rd_valid (domain_b_valid),
 .rd_data (domain_b_data)
);

```

### 8.6.2 Example 2: high-speed to low-speed CDC

```

// Fast writer (250 MHz) → Slow reader (62.5 MHz)
// Needs deeper FIFO to handle burst traffic
gaxi_fifo_async #(
 .DATA_WIDTH(64),
 .DEPTH(32), // Deeper for rate mismatch
 .N_FLOP_CROSS(3),
 .REGISTERED(1) // Flop mode for timing
) u_rate_converter (
 .axi_wr_aclk (clk_250mhz),
 .axi_wr_aresetn (rst_fast_n),
 .wr_valid (fast_valid),
 .wr_ready (fast_ready),
 .wr_data (fast_data),

 .axi_rd_aclk (clk_62p5mhz),
 .axi_rd_aresetn (rst_slow_n),
 .rd_ready (slow_ready),
 .rd_valid (slow_valid),
 .rd_data (slow_data)
);

```

## 8.7 Design Notes

---

### 8.7.1 Depth sizing for clock ratio

When write clock >> read clock, size the FIFO to handle burst accumulation:

Required Depth  $\geq$  Burst Size  $\times$  (1 - Read Freq / Write Freq)  $\times$  Safety Margin

Example:

- Write: 100 MHz, Read: 25 MHz
  - Burst: 16 transfers
  - Safety: 1.5x
- Depth  $\geq 16 \times (1 - 25/100) \times 1.5 = 18$  entries

### 8.7.2 Reset synchronization

**Critical:** both clock domains must have properly synchronized resets!

```
// Separate reset synchronizers for each domain
```

```
reset_sync u_wr_rst_sync (
 .clk(axi_wr_aclk),
 .rst_n(global_rst_n),
 .sync_rst_n(axi_wr_aresetn)
);
```

```
reset_sync u_rd_rst_sync (
 .clk(axi_rd_aclk),
 .rst_n(global_rst_n),
 .sync_rst_n(axi_rd_aresetn)
);
```

### 8.7.3 Metastability protection

N\_FLOP\_CROSS sets the synchronizer depth, and MTBF rises steeply with it.

#### *Synchronizer stage guidance*

| Stages | Use Case                    |
|--------|-----------------------------|
| 2      | Short-term prototyping only |
| 3      | <b>Production standard</b>  |
| 4      | Ultra-critical systems      |

This page deliberately gives no MTBF figures. Real MTBF depends on the flop’s metastability time constant and resolution window, the two clock frequencies and the data toggle rate – none of which a module page can know. A second table here also drifted from the one in [glitch\\_free\\_n\\_dff\\_arn](#), which put 2 stages at “hours” where this page said “years”. One source per fact: that page owns the discussion.

**Recommendation:** always use N\_FLOP\_CROSS=3 in production.

### 8.7.4 Error checking

The RTL contains **no** assertions and no runtime \$display checks. What is there is a single empty read-domain block:

```
always_ff @(posedge axi_rd_aclk) begin
 if (w_read && r_rd_empty) begin end // empty -- reports nothing
end
```

There is no write-domain equivalent. Overflow is prevented structurally by the !wr\_full write guard rather than reported; underflow is not detected at all. If you need overflow/underflow telemetry in simulation, add it in the testbench. fifo\_async.md states the same for its module.

### 8.7.5 Common mistakes

**1. Metastability.** Symptom: random data corruption, simulation/hardware mismatch. Fix: increase `N_FLOP_CROSS` to 3 or 4.

**2. Pointer synchronization failure.** Symptom: FIFO full/empty signals incorrect. Debug:

1. Verify the clocks are truly independent (no PLL relationship violations)
2. Check reset synchronization – both domains must reset properly
3. Verify the Gray code conversion logic

**3. Underflow/overflow.** Symptom: data loss or corruption. Debug:

1. Check `DEPTH` is sufficient for the clock ratio
2. Verify flow control is respected in both domains
3. Monitor pointer values in both domains

## 8.8 Related Modules

---

- [gaxi\\_fifo\\_sync](#) - Single clock domain version
- [gaxi\\_skid\\_buffer\\_async](#) - Async skid buffer
- [GAXI Index](#) - Overview

## 8.9 Testing

---

*# Async FIFO tests with various clock ratios*

```
pytest val/cdc/test_gaxi_buffer_async.py -v
```

*# Test specific clock ratio (wr:rd)*

```
pytest val/cdc/test_gaxi_buffer_async.py -k "wr10_rd12" -v # 10ns : 12ns
```

The test matrix covers:

- Same clocks (1:1)
- 1.2x ratio (10ns : 12ns)
- 2x ratio (10ns : 20ns)
- 2.5x ratio (8ns : 20ns)

## 8.10 References

---

- **Clifford Cummings:** “Simulation and Synthesis Techniques for Asynchronous FIFO Design” (Sunburst Design)
- **Source:** `rtl/cdc/gaxi_fifo_async.sv`

- **Tests:** `val/cdc/test_gaxi_buffer_async.py`

## 8.11 Navigation

---

- [← Back to CDC Index](#)
  - [← Back to Main Documentation Index](#)
- 

**Version:** 1.0 **Last Updated:** 2025-10-06

## 9 gaxi\_skid\_buffer\_async

**Module:** gaxi\_skid\_buffer\_async.sv **Location:** rtl/cdc/ **Status:** Production Ready

---

### 9.1 Overview

---

This module is a wrapper, and it owns that: a synchronous skid buffer on the write side, an asynchronous FIFO behind it. One instantiation buys you the skid buffer's elastic buffering plus a clock domain crossing. One thing to know before you build around it — the skid buffer registers its handshake outputs, so there is no zero-latency bypass path; see Latency Components under Timing below.

#### 9.1.1 Key Features

- **CDC Capability:** Safe clock domain crossing
  - **Elastic Buffering:** Combines skid buffer + async FIFO
  - **Write-Side Buffering:** Skid buffer in write domain prevents stalls
  - **Configurable CDC:** Adjustable synchronizer stages
  - **Dual Clock Domains:** Independent write/read clocks
- 

### 9.2 Parameters

---

| Parameter    | Type           | Default   | Description                                                                         |
|--------------|----------------|-----------|-------------------------------------------------------------------------------------|
| MEM_STYLE    | fifo_mem_<br>t | FIFO_AUTO | Async FIFO storage:<br>FIFO_AUTO / FIFO_SRL /<br>FIFO_BRAM                          |
| REGISTERED   | int            | 0         | Async FIFO read mode<br>(0=mux, 1=flop)                                             |
| DATA_WIDTH   | int            | 32        | Data bus width                                                                      |
| DEPTH        | int            | 2         | Async FIFO depth (CDC<br>stage)                                                     |
| USE_JOHNSON  | int            | 0         | 0 = Gray pointers, power-<br>of-2 DEPTH only; 1 =<br>Johnson pointers, any<br>DEPTH |
| N_FLOP_CROSS | int            | 2         | Synchronizer stages (3<br>recommended)                                              |

USE\_JOHNSON reaches the async FIFO underneath and carries the same meaning it has there: it selects the pointer encoding, and with it the legal DEPTH values. Leave it at 0 and DEPTH must be a power of 2 or elaboration fails; set it to 1 and odd depths become legal at the cost of a wider pointer. See [gaxi\\_fifo\\_async](#) for the trade-off in full.

**Note:** the skid half's depth is NOT DEPTH. The wrapper instantiates `gaxi_skid_buffer` with `DATA_WIDTH` only, so it takes that module's own default of 2. DEPTH sizes the async FIFO behind it.

---

## 9.3 Ports

---


The port list splits cleanly into a write domain and a read domain:

```
module gaxi_skid_buffer_async #(
 parameter fifo_mem_t MEM_STYLE = FIFO_AUTO, // FIFO_AUTO /
 FIFO_SRL / FIFO_BRAM
 parameter int REGISTERED = 0, // 0 = mux mode, 1
 = flop mode
 parameter int DATA_WIDTH = 32,
 parameter int DEPTH = 2, // async FIFO depth
 parameter int USE_JOHNSON = 0, // 0 = Gray (power-
 of-2 depth), 1 = Johnson (any depth)
 parameter int N_FLOP_CROSS = 2, // CDC synchronizer
 stages
 parameter int DW = DATA_WIDTH // port-width alias
 used below
) (
 // Write Domain
 input logic axi_wr_aclk,
 input logic axi_wr_aresetn,
 input logic wr_valid,
 output logic wr_ready,
 input logic [DW-1:0] wr_data,

 // Read Domain
 input logic axi_rd_aclk,
 input logic axi_rd_aresetn,
 output logic rd_valid,
 input logic rd_ready,
 output logic [DW-1:0] rd_data
);
```

---

## 9.4 Functional Description



### Architecture

#### 9.4.1 Internal Components

- gaxi\_skid\_buffer (Write Domain):**
  - Provides elastic buffering
  - Prevents write-side stalls
  - Fixed depth (typically 2-4 entries)
- gaxi\_fifo\_async (CDC):**
  - Handles clock domain crossing
  - Configurable depth (DEPTH parameter)
  - Gray code pointer synchronization

## 9.5 Timing Characteristics

#### 9.5.1 Latency Components

| Component                 | Latency           | Domain     |
|---------------------------|-------------------|------------|
| Skid buffer (empty)       | 1 cycle           | Write      |
| Skid buffer (buffered)    | 1 cycle           | Write      |
| Async FIFO CDC            | 3-5 cycles        | Both       |
| <b>Total (empty path)</b> | <b>4-6 cycles</b> | End-to-end |
| <b>Total (buffered)</b>   | <b>4-6 cycles</b> | End-to-end |

`gaxi_skid_buffer` registers its handshake outputs, so even a write into an empty buffer costs one write clock before `rd_valid` rises and the async FIFO can take the data. There is no zero-latency bypass path.

### 9.5.2 Throughput

- **Write Side:** 1 transfer/cycle (when not full)
  - **Read Side:** 1 transfer/cycle (when not empty)
  - **Sustained:** Limited by slower clock domain
- 

## 9.6 Usage Examples

---

### 9.6.1 Example 1: Basic Async Skid Buffer

```
gaxi_skid_buffer_async #(
 .DATA_WIDTH(64),
 .DEPTH(4), // Async FIFO depth
 .N_FLOP_CROSS(3), // CDC stages
 .REGISTERED(0)
) u_async_skid (
 // Write domain @ 100 MHz
 .axi_wr_aclk (clk_100),
 .axi_wr_aresetn (rst_100_n),
 .wr_valid (src_valid),
 .wr_ready (src_ready),
 .wr_data (src_data),

 // Read domain @ 75 MHz
 .axi_rd_aclk (clk_75),
 .axi_rd_aresetn (rst_75_n),
 .rd_valid (dst_valid),
 .rd_ready (dst_ready),
 .rd_data (dst_data)
);
```

### 9.6.2 Example 2: Pipeline Stage with CDC

```
// Break long path AND cross clock domains
gaxi_skid_buffer_async #(
 .DATA_WIDTH(128),
 .DEPTH(8), // Larger for rate buffering
 .N_FLOP_CROSS(3),
 .REGISTERED(1) // Registered output for timing
) u_cdc_pipeline (
 .axi_wr_aclk (source_clk),
 .axi_wr_aresetn (source_rst_n),
 .wr_valid (stage_in_valid),
 .wr_ready (stage_in_ready),
 .wr_data (stage_in_data),
```

```

 .axi_rd_aclk (dest_clk),
 .axi_rd_aresetn (dest_rst_n),
 .rd_valid (stage_out_valid),
 .rd_ready (stage_out_ready),
 .rd_data (stage_out_data)
);

```

---

## 9.7 Design Notes

### 9.7.1 When to Use This Module

**Use async skid buffer when:** - Need CDC between clock domains - Want elastic buffering on write side - Write-side needs low-latency response

**Use plain async FIFO instead when:** - Only need CDC, no special write-side buffering - Want deeper buffering (larger DEPTH) - Don't need skid buffer complexity

### 9.7.2 Depth Selection

Total buffering = Skid buffer entries + DEPTH

**Recommendations:** - **DEPTH=4:** Minimal CDC buffering, light traffic - **DEPTH=8:** Moderate buffering, balanced - **DEPTH=16:** Heavy buffering, burst traffic

### 9.7.3 Reset Synchronization

**Critical:** Both domains need synchronized resets!

*// Example reset synchronization*

```

reset_sync #(.N(3)) u_wr_sync (
 .clk(axi_wr_aclk),
 .rst_n(global_rst_n),
 .sync_rst_n(axi_wr_aresetn)
);

```

```

reset_sync #(.N(3)) u_rd_sync (
 .clk(axi_rd_aclk),
 .rst_n(global_rst_n),
 .sync_rst_n(axi_rd_aresetn)
);

```

### 9.7.4 Resource Utilization

Combines resources of both sub-modules:

| Component                    | Flops                                      | LUTs                         |
|------------------------------|--------------------------------------------|------------------------------|
| Skid buffer (fixed, DEPTH=2) | $2 \times DW + \sim 4$                     | $\sim 50$                    |
| Async FIFO (DEPTH=8)         | $8 \times DW + \sim 40$                    | $\sim 120$                   |
| <b>Total (DEPTH=8)</b>       | <b><math>\sim 10 \times DW + 44</math></b> | <b><math>\sim 170</math></b> |

The skid row does not scale with DEPTH. The wrapper never overrides the skid buffer's depth, so its storage is `r_data[2]` — two words — whatever DEPTH you pass. Only the FIFO row moves.

### 9.7.5 Common Issues

#### Issue 1: CDC Violations

Symptom: metastability, data corruption.

Solution: - Set `N_FLOP_CROSS=3` (minimum for production) - Verify clocks are truly independent - Check reset synchronization

#### Issue 2: Write-Side Stalls

Symptom: `wr_ready` deasserts unexpectedly.

Debug: 1. Check async FIFO depth (DEPTH parameter) 2. Verify read-side is consuming data 3. Monitor clock ratio - may need deeper FIFO

#### Issue 3: Unexpected Latency

Symptom: data takes longer than expected.

Explanation: - Total latency = skid buffer (1 write clock) + async FIFO CDC (3-5 cycles) - Minimum 4-6 cycles even when empty — see the latency table above, which is where this number comes from - There is no zero-latency bypass: `gaxi_skid_buffer` registers its handshake outputs, so the skid hop always costs a write clock - This is inherent to CDC design

---

## 9.8 Related Modules

---

- [gaxi\\_skid\\_buffer](#) - Synchronous version
  - [gaxi\\_fifo\\_async](#) - Async FIFO without skid buffer
  - [GAXI Index](#) - Overview
-

## 9.9 Testing

---

*# Async skid buffer tests*

```
pytest val/cdc/test_gaxi_buffer_async.py -k "skid" -v
```

*# Test specific clock ratio*

```
pytest val/cdc/test_gaxi_buffer_async.py -k "skid" -k "wr10_rd20" -v
```

---

## 9.10 References

---

- **Source:** rtl/cdc/gaxi\_skid\_buffer\_async.sv
  - **Dependencies:**
    - gaxi\_skid\_buffer.sv
    - gaxi\_fifo\_async.sv
  - **Tests:** val/cdc/test\_gaxi\_buffer\_async.py
- 

**Version:** 1.0 **Last Updated:** 2025-10-06

---

## 9.11 Navigation

---

- [← Back to GAXI Index](#)
- [← Back to Main Documentation Index](#)

## 10 glitch\_free\_n\_dff\_arn

### 10.1 Overview

This is the synchronizer the rest of the CDC toolbox is built on: a parameterized multi-stage flip-flop chain for safe clock domain crossing (CDC). You set the stage count and the width; the configurable flip-flop stages knock down metastability risk while your data makes the jump between asynchronous clock domains intact.

### 10.2 Parameters

| Parameter  | Type | Default | Description                             |
|------------|------|---------|-----------------------------------------|
| FLOP_COUNT | int  | 3       | Number of synchronizer flip-flop stages |
| WIDTH      | int  | 4       | Data width in bits                      |


### 10.3 Ports

| Port  | Direction | Width | Description                                    |
|-------|-----------|-------|------------------------------------------------|
| clk   | Input     | 1     | Destination domain clock                       |
| rst_n | Input     | 1     | Destination domain active-low reset            |
| d     | Input     | WIDTH | Input data from source clock domain            |
| q     | Output    | WIDTH | Synchronized output data in destination domain |

### 10.4 Functional Description


#### 10.4.1 The Metastability Problem

When signals cross between asynchronous clock domains:



Metastability effects:

- **Undefined output:** Neither 0 nor 1, but an intermediate voltage
- **Extended resolution:** Takes longer than a normal clock period to settle
- **Propagation:** Can corrupt downstream logic
- **System failure:** Can take the entire design down



### *Multi-Stage Synchronizer Solution*

Stage by stage:

1. **Stage 1 (FF1):** Captures the source data — this is the one that may go metastable
2. **Stage 2 (FF2):** Gives metastability time to resolve, low probability of propagation
3. **Stage 3+ (FFN):** Further **increases** MTBF (Mean Time Between Failures) — each added stage gives metastability more time to resolve (see the ~1000×-per-stage table below)

#### 10.4.2 Packed Array Structure

```
logic [FC-1:0][WIDTH-1:0] r_q_array;
```

- **Efficient storage:** Single declaration for all stages
- **Clear indexing:** `r_q_array[0]` is first stage, `r_q_array[FC-1]` is last
- **Synthesis friendly:** Maps well to FPGA/ASIC resources

### 10.4.3 Shift Register Behavior

```
always_ff @(posedge clk or negedge rst_n) begin
 if (!rst_n) begin
 // Reset all flip-flops to zero
 r_q_array <= {FC{{DW{1'b0}}}};
 end else begin
 // Load new data into first stage
 r_q_array[0] <= d;
 // Shift existing data through remaining stages
 for (int i = 1; i < FC; i++) begin
 r_q_array[i] <= r_q_array[i-1];
 end
 end
end
```

### 10.4.4 Reset Behavior

- **Asynchronous reset** (active-low `rst_n`): the RTL uses `always_ff` `@(posedge clk or negedge rst_n)` (the “arn” suffix = async reset, active-low), driven by the destination-domain reset
- **Complete reset**: All stages reset to zero
- **Nested replication**: `{FC{{DW{1'b0}}}}` creates FC copies of DW zeros

### 10.4.5 MTBF (Mean Time Between Failures) Analysis

The MTBF calculation:

$$\text{MTBF} = (e^{(t_{\text{res}}/\tau)}) / (T_0 \times f_{\text{clk}} \times f_{\text{data}})$$

Where:

- $t_{\text{res}}$  = Resolution time available (clock period - setup - routing)
- $\tau$  = Flip-flop metastability resolution time constant (~200ps typical)
- $T_0$  = Metastability aperture / window constant (device-specific, ~1e-12 s)
- $f_{\text{clk}}$  = Destination clock frequency
- $f_{\text{data}}$  = Data transition frequency

**$\tau$  appears ONLY in the exponent.** The denominator constant is  $T_0$ , not  $\tau$  — they are different device parameters and are not interchangeable. Both are process/library specific: take them from the vendor’s metastability characterization data, not from a datasheet timing table.

The stage-count table below is a **rule of thumb** for relative improvement, not a result derived from this formula. Use the formula only with real  $\tau/T_0$  values for your target device; otherwise rely on the table for sizing.

Stage count impact:

| Stages | Relative MTBF   |
|--------|-----------------|
| 1      | 1x (baseline)   |
| 2      | ~1000x          |
| 3      | ~1,000,000x     |
| 4      | ~1,000,000,000x |

**Relative only, and deliberately so.** This table used to carry an absolute column — Minutes / Hours / Years / Millennia — with no stated basis, and it disagreed with this module’s own RTL header (which claims  $\sim 10^6$  hours at FLOP\_COUNT=2, about 114 years, where the table said “Hours”) by roughly eight orders of magnitude. Both figures were unsourced, so neither could be trusted and the column is gone.

An absolute MTBF is a property of your device and your clock, not of this module. If you need one, put your vendor’s tau and T0 into the formula above with your real f\_clk and f\_data. If you do not have those numbers, use the relative column for sizing and do not quote an absolute figure to anyone.

Practical guidelines:

- **2 stages:** Minimum for most applications
- **3 stages:** Recommended for high-reliability systems
- **4+ stages:** Extreme reliability requirements (aerospace, medical)

## 10.5 Timing Characteristics

### 10.5.1 Synchronizer Latency

- **Minimum latency:** 2 clock cycles (2-stage synchronizer)
- **General latency:** FLOP\_COUNT clock cycles
- **Fixed delay:** Independent of data pattern

### 10.5.2 Performance Impact

*// 2-stage: 2 cycle latency, good MTBF*

```
glitch_free_n_dff_arn #(.FLOP_COUNT(2)) fast_sync (...);
```

*// 3-stage: 3 cycle latency, excellent MTBF*

```
glitch_free_n_dff_arn #(.FLOP_COUNT(3)) safe_sync (...);
```

*// 4-stage: 4 cycle latency, extreme MTBF*

```
glitch_free_n_dff_arn #(.FLOP_COUNT(4)) ultra_safe_sync (...);
```

## 10.6 Usage Examples

---

### 10.6.1 Reset Synchronizer Pattern

```
// Asynchronous assert, synchronous deassert
logic reset_sync;
glitch_free_n_dff_arn #(.FLOP_COUNT(2), .WIDTH(1)) reset_sync_inst (
 .clk(dest_clk),
 .rst_n(1'b1), // No reset feedback
 .d(~async_reset), // Invert for active-low
 .q(reset_sync)
);
```

```
assign sync_reset_n = reset_sync & ~async_reset;
```

### 10.6.2 Enable Signal Synchronization

```
glitch_free_n_dff_arn #(.FLOP_COUNT(3), .WIDTH(1)) enable_sync (
 .clk(fast_clk),
 .rst_n(rst_n),
 .d(slow_enable),
 .q(fast_enable_sync)
);
```

### 10.6.3 Status Flag Collection

```
// Synchronize multiple independent status signals
glitch_free_n_dff_arn #(.FLOP_COUNT(2), .WIDTH(4)) status_sync (
 .clk(dest_clk),
 .rst_n(rst_n),
 .d({error_flag, done_flag, busy_flag, ready_flag}),
 .q({sync_error, sync_done, sync_busy, sync_ready})
);
```

## 10.7 Design Notes

---

### 10.7.1 Appropriate Use Cases

- **Control signals:** Reset, enable, mode switches
- **Slow status signals:** Error flags, configuration bits
- **Gray code pointers:** FIFO read/write pointers
- **Single-bit signals:** Generally safe for multi-bit if Gray coded

### 10.7.2 Inappropriate Use Cases

- **High-speed data buses:** Use proper CDC techniques (handshaking, FIFOs)
- **Binary counters:** Multiple bits can transition simultaneously

- **Address buses:** Binary addresses are not CDC-safe
- **Clock signals:** Never synchronize clocks this way

### 10.7.3 Multi-Bit Considerations

```
// SAFE: Gray-coded FIFO pointers
glitch_free_n_dff_arn #(.WIDTH(5)) gray_sync (
 .d(gray_pointer), // Gray code - only 1 bit changes
 .q(sync_gray_pointer)
);

// UNSAFE: Binary counter synchronization
// Don't do this - multiple bits can change simultaneously
glitch_free_n_dff_arn #(.WIDTH(8)) bad_sync (
 .d(binary_counter), // Binary - multiple bit transitions
 .q(sync_counter) // May capture inconsistent values
);
```

### 10.7.4 Synthesis Considerations

On FPGA:

- **Dedicated resources:** Maps to flip-flop primitives
- **Placement:** Tools may place stages in different locations
- **Timing constraints:** Require proper CDC constraints

Timing constraints:

```
Constrain the synchronizer input as asynchronous
set_false_path -to [get_pins sync_inst/r_q_array[0]/D]

Or set maximum delay if there's some relationship
set_max_delay -to [get_pins sync_inst/r_q_array[0]/D] 2.0
```

Optimization prevention:

```
// Synthesis attributes to prevent optimization (one declaration, both
attributes --
// two separate declarations of r_q_array do not compile)
(* ASYNC_REG = "TRUE", DONT_TOUCH = "TRUE" *) logic [FC-1:0][WIDTH-
1:0] r_q_array;
```

### 10.7.5 Debug Features

Flatten the array for waveform viewing:

```
// Flatten array for waveform viewing
wire [(DW*FC)-1:0] flat_r_q;
genvar i;
```

```

generate
 for (i = 0; i < FC; i++) begin : gen_flatten_memory
 assign flat_r_q[i*DW+:DW] = r_q_array[i];
 end
endgenerate

```

Simulation visibility:

- **All stages visible:** Can observe data propagation through stages
- **Metastability injection:** Can inject X values for testing
- **Timing analysis:** Observe setup/hold violations

## 10.8 Related Modules

---

- **FIFO CDC modules:** Use this for pointer synchronization
- **Reset synchronizers:** Specialized CDC for reset signals
- **Handshake synchronizers:** For complex data CDC
- **Pulse synchronizers:** For single-cycle pulse CDC

## 10.9 Testing

---

### 10.9.1 Verification Strategies

Functional verification:

- **Data integrity:** Verify output eventually matches input
- **Reset behavior:** Check proper reset operation
- **Latency verification:** Confirm expected delay

Metastability testing:

```

// Stress test with rapid data changes
initial begin
 forever begin
 @(posedge source_clk);
 d <= $random();
 #($random() % 100); // Random timing
 end
end

```

Formal verification:

- **Eventually properties:** Data will eventually propagate
- **Reset properties:** Reset behavior verification
- **No X propagation:** Ensure no unknown values propagate

## 10.9.2 Test Suite

From the test suite (`val/cdc/test_glitch_free_n_dff_arn.py`):

- **Key test scenarios:**
  - Test runner for `glitch_free_n_dff_arn` module
  - The `glitch_free_n_dff_arn` is a parameterized N-stage synchronizer for safe
  - clock domain crossing. It reduces metastability risk by passing signals
  - through multiple flip-flop stages.
  - Test Configurations:

Run levels come from the standard grid: `REG_LEVEL=GATE|FUNC|FULL` selects the parameter set, `TEST_LEVEL` the per-test depth. Run the whole area with `make -C val/cdc run-all-func-parallel`, never bare `pytest` for suites.

## 10.10 Navigation

---

- [← Back to rtl-common Index](#)
- [← Back to Main Documentation Index](#)

## 11 reset\_sync

### 11.1 Overview

The `reset_sync` module is a parameterized reset synchronizer with asynchronous assertion and synchronous deassertion. That combination is the whole point: reset needs to hit immediately when it fires, but it has to come *off* on a clock edge. Release it asynchronously and every flop downstream gets a setup/hold gamble. This is the standard circuit for proper reset distribution, and it exists to keep reset release from going metastable or violating timing.

- Asynchronous reset assertion (immediate)
- Synchronous reset deassertion (clocked)
- Parameterizable synchronization depth
- Metastability resolution
- Standard reset synchronizer topology
- Clean reset release timing

### 11.2 Parameters

*reset\_sync parameters*

| Parameter      | Type | Default | Description                                                                                                                                                   |
|----------------|------|---------|---------------------------------------------------------------------------------------------------------------------------------------------------------------|
| N              | int  | 3       | Number of synchronization stages. Must be $\geq 2$ ; an elaboration \$error rejects anything smaller.                                                         |
| KEEP_ATTRS     | bit  | 1'b1    | Attach vendor CDC attributes (ASYNC_REG, SHREG_EXTRACT, altera_attribute, syn_preserve) to the chain so tools recognize it and do not extract it into an SRL. |
| IN_ACTIVE_LOW  | bit  | 1'b1    | Polarity of the <code>rst_n</code> input. 1 = active-low (default), 0 = active-high.                                                                          |
| OUT_ACTIVE_LOW | bit  | 1'b1    | Polarity of the                                                                                                                                               |

| Parameter    | Type | Default | Description                                                                                                                                                                                                |
|--------------|------|---------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| ASYNC_ASSERT | bit  | 1'b1    | sync_rst_n output. 1 = active-low (default), 0 = active-high.<br><br>1 = asynchronous assert with synchronous deassert (FPGA best practice). 0 = fully synchronous style, where assertion is also clocked. |

**Polarity without renaming ports.** IN\_ACTIVE\_LOW and OUT\_ACTIVE\_LOW let the module sit in an active-high reset domain while the ports stay named rst\_n and sync\_rst\_n. The input gets normalized to an internal active-high form on the way in, and the requested polarity gets applied on the way out:

```
wire rst_in_h = IN_ACTIVE_LOW ? ~rst_n : rst_n;
...
sync_rst_n = OUT_ACTIVE_LOW ? ~sync_rst_h : sync_rst_h;
```

The names stick around for compatibility, which means a reset\_sync with OUT\_ACTIVE\_LOW = 0 drives an active-**high** reset out of a port still called sync\_rst\_n. Check the parameter, not the name.

## 11.3 Ports

### 11.3.1 Inputs

| Port  | Width | Description                           |
|-------|-------|---------------------------------------|
| clk   | 1     | System clock for synchronization      |
| rst_n | 1     | Asynchronous input reset (active-low) |

### 11.3.2 Outputs

| Port       | Width | Description                            |
|------------|-------|----------------------------------------|
| sync_rst_n | 1     | Synchronized reset output (active-low) |

## 11.4 Functional Description

---

### 11.4.1 The Reset Problem

Reset release is where synchronous designs get hurt. Mishandle it and you're looking at: -

**Metastability:** When reset is released asynchronously relative to clock - **Timing Violations:**

Setup/hold violations during reset release - **Race Conditions:** Different flip-flops releasing from reset at different times

### 11.4.2 Asynchronous Assert, Synchronous Deassert

The behavior you actually want: - **Assert:** Immediate (asynchronous) when reset input becomes active - **Deassert:** Synchronized to clock edge to ensure clean timing

### 11.4.3 Synchronization Register Chain

```
logic [N-1:0] r_sync_reg = '0;
```

An N-stage shift register provides the synchronization delay. Here's the part that trips everyone up: the chain carries the **active-high** internal form of the reset — a 1 in the chain means “reset asserted”. That's the opposite of what the port names suggest, and it's the single most common source of confusion when people read this module.

### 11.4.4 Reset Logic

```
// ASYNC_ASSERT = 1 (default): async assert, sync deassert
always_ff @(posedge clk or posedge rst_in_h) begin
 if (rst_in_h) r_sync_reg <= '1; // hold asserted
 through chain
 else r_sync_reg <= {r_sync_reg[N-2:0], 1'b0};
end
```

While reset is asserted, the chain is held at all ones. Once it releases, zeros shift in from the LSB, and after N clock edges the MSB has cleared and the reset is considered released.

With ASYNC\_ASSERT = 0 the sensitivity list drops `rst_in_h` and assertion gets sampled on the clock edge like any other synchronous load:

```
always_ff @(posedge clk) begin
 if (rst_in_h) r_sync_reg <= '1;
 else r_sync_reg <= {r_sync_reg[N-2:0], 1'b0};
end
```

That style can't assert reset while the clock is stopped — which is exactly why the asynchronous-assert variant is the default.

### 11.4.5 Output Assignment

```
wire sync_rst_h = r_sync_reg[N-1];
always_comb begin
```

```

 sync_rst_n = OUT_ACTIVE_LOW ? ~sync_rst_h : sync_rst_h;
end

```

The MSB of the chain is the active-high synchronized reset; the output stage applies the requested polarity.

The RTL body is written out four times inside a generate, once per combination of ASYNC\_ASSERT and KEEP\_ATTRS, because the vendor attributes have to be attached at the declaration. All four variants are logically identical apart from the attributes and the sensitivity list. Yes, it's repetitive. That's the attribute syntax's fault, not the author's.

## 11.5 Timing Characteristics

### 11.5.1 Reset Assertion (Asynchronous)

```

rst_n: _____
sync_rst_n: _____
 ^
 Immediate assertion

```

### 11.5.2 Reset Deassertion (Synchronous, N=3)

```

clk: - - - - -
rst_n: _____
r_sync_reg: 111 111 111 110 100 000 000 000 000
sync_rst_n: _____
 ^
 Synchronized deassertion
 (N clock cycles after rst_n release)

```

### 11.5.3 State Sequence for N=3

Default polarities (IN\_ACTIVE\_LOW = 1, OUT\_ACTIVE\_LOW = 1), so `rst_n = 0` means reset asserted and `sync_rst_n = 0` means reset asserted downstream. Remember the chain holds the active-high form, so 111 is “asserted”.

*reset\_sync state sequence for N=3 at default polarities*

| Clock | rst_n | rst_in_h | r_sync_reg | sync_rst_n | Description                              |
|-------|-------|----------|------------|------------|------------------------------------------|
| -     | 0     | 1        | 111        | 0          | Reset asserted, chain forced to all ones |
| 1     | 1     | 0        | 110        | 0          | Reset released;                          |

| Clock | rst_n | rst_in_h | r_sync_reg | sync_rst_n | Description                           |
|-------|-------|----------|------------|------------|---------------------------------------|
|       |       |          |            |            | first 0 shifted in                    |
| 2     | 1     | 0        | 100        | 0          | Second 0 shifted in                   |
| 3     | 1     | 0        | 000        | 1          | MSB clears; reset released downstream |
| 4+    | 1     | 0        | 000        | 1          | Stable, out of reset                  |

The output releases on the third clock edge after rst\_n goes high — those N cycles are the metastability margin the chain exists to buy you.

## 11.6 Usage Examples

### 11.6.1 Choosing N Value

#### 11.6.1.1 *N=2 (Minimum)*

```
reset_sync #(.N(2)) rst_sync_inst (
 .clk(sys_clk),
 .rst_n(external_rst_n),
 .sync_rst_n(internal_rst_n)
);
```

- **Use when:** Fast reset release required
- **Risk:** Minimal metastability margin
- **Applications:** High-speed designs with good clock quality

#### 11.6.1.2 *N=3 (Recommended)*

```
reset_sync #(.N(3)) rst_sync_inst (
 .clk(sys_clk),
 .rst_n(external_rst_n),
 .sync_rst_n(internal_rst_n)
);
```

- **Use when:** Standard applications
- **Benefit:** Good balance of speed vs. metastability margin
- **Applications:** Most digital designs

#### 11.6.1.3 $N \geq 4$ (Conservative)

```
reset_sync #(.N(4)) rst_sync_inst (
 .clk(sys_clk),
 .rst_n(external_rst_n),
 .sync_rst_n(internal_rst_n)
);
```

- **Use when:** Critical applications, poor clock quality
- **Benefit:** Maximum metastability protection
- **Trade-off:** Slower reset release

#### 11.6.2 Global Reset Distribution

*// Synchronize external reset for internal use*

```
reset_sync #(.N(3)) global_rst_sync (
 .clk(system_clock),
 .rst_n(external_reset_n),
 .sync_rst_n(global_reset_n)
);
```

*// Use synchronized reset throughout design*

```
always_ff @(posedge system_clock or negedge global_reset_n) begin
 if (!global_reset_n) begin
 // Reset state
 end else begin
 // Normal operation
 end
end
```

#### 11.6.3 Clock Domain Crossing

*// Separate reset synchronizers for each clock domain*

```
reset_sync #(.N(3)) clk_a_rst_sync (
 .clk(clk_a),
 .rst_n(system_rst_n),
 .sync_rst_n(clk_a_rst_n)
);
```

```
reset_sync #(.N(3)) clk_b_rst_sync (
 .clk(clk_b),
 .rst_n(system_rst_n),
 .sync_rst_n(clk_b_rst_n)
);
```

#### 11.6.4 Power-On Reset

```
// Synchronize power-on reset
reset_sync #(.N(3)) por_sync (
 .clk(main_clk),
 .rst_n(power_on_rst_n),
 .sync_rst_n(system_rst_n)
);
```

#### 11.6.5 Reset Controller Integration

```
module reset_controller (
 input logic sys_clk,
 input logic external_rst_n,
 input logic watchdog_rst_n,
 input logic sw_rst,
 output logic sync_rst_n
);

 logic combined_rst_n;

 // Combine all reset sources
 assign combined_rst_n = external_rst_n & watchdog_rst_n & ~sw_rst;

 // Synchronize the combined reset
 reset_sync #(.N(3)) rst_sync_inst (
 .clk(sys_clk),
 .rst_n(combined_rst_n),
 .sync_rst_n(sync_rst_n)
);

endmodule
```

### 11.7 Design Notes

---

#### 11.7.1 1. Why the chain looks inverted

The chain stores the **active-high** form of the reset, while the ports are named for the active-low one. Read the dataflow in order and the confusion evaporates:

- `rst_in_h` is the input normalized to active-high
- while `rst_in_h` is 1, the chain is forced to all ones (reset asserted)
- when `rst_in_h` falls, zeros shift in from the LSB
- after `N` edges the MSB is 0, and the output stage inverts it back to the requested polarity

So during reset the register holds 1s and after release it drains to 0s — the opposite of what the port names alone would tell you.

### 11.7.22. Metastability Resolution

The N stages buy you multiple clock cycles for metastability to resolve: - Each flip-flop stage gets a chance to resolve potential metastability - Probability of metastability propagation:  $(MTBF)^{(-N)}$  - N=2 typically sufficient, N=3 provides extra margin

### 11.7.33. Initialization

```
logic [N-1:0] r_sync_reg = '0;
```

The declaration initializer gives FPGA configuration a defined starting value. Note which value: 0, the *released* state. So on an FPGA the design comes out of configuration with reset already deasserted unless `rst_n` is asserted. On an ASIC the initializer is ignored, and the first asynchronous assertion establishes the state.

### 11.7.44. Vendor attributes

With `KEEP_ATTRS = 1` (the default) the chain carries `ASYNC_REG = "TRUE"` and `SHREG_EXTRACT = "N0"` for Xilinx, `altera_attribute` forcing synchronizer identification for Quartus, and `syn_preserve` for Synplify. These do two jobs: they stop the tool from packing the chain into an SRL primitive — which would quietly destroy your metastability margin — and they let CDC reports recognize it as a synchronizer. Set `KEEP_ATTRS = 0` only if your flow objects to the attributes.

### 11.7.55. Parameterization Benefits

Configurable N lets you tune for different requirements: - N=2: Minimum permitted, fast reset release - N=3: Standard choice, good metastability margin - N>=4: Conservative, slower release in exchange for maximum metastability protection

### 11.7.6 Reset Tree Planning

```
// Plan reset distribution hierarchy
// Level 1: Global reset synchronizer
// Level 2: Domain-specific reset synchronizers
// Level 3: Module-level reset distribution
```

### 11.7.7 Clock Quality Requirements

- **Clean Clock:** Stable, low-jitter clock required
- **Clock Frequency:** Must be running during reset release
- **Clock Gating:** Avoid gating clock used for reset synchronization

### 11.7.8 Reset Source Management

```
// Multiple reset sources need careful combination
wire combined_reset = por_n & external_rst_n & watchdog_rst_n &
~soft_reset;
```

### 11.7.9 Incorrect Sensitivity List

```
// WRONG: Missing async reset in sensitivity list
always_ff @(posedge clk) begin
 if (!rst_n) r_sync_reg <= '0;
 else r_sync_reg <= {r_sync_reg[N-2:0], 1'b1};
end
```

```
// CORRECT: Include async reset
always_ff @(posedge clk or negedge rst_n) begin
 // ... correct implementation
end
```

### 11.7.10 Wrong Reset Polarity

```
// Ensure consistent reset polarity throughout design
// This module uses active-low reset (rst_n)
```

### 11.7.11 Clock Domain Issues

```
// Each clock domain needs its own reset synchronizer
// Don't share synchronized reset across clock domains
```

## 11.8 Related Modules

---

- **Positive-edge reset:** For active-high reset systems
- **Multi-domain reset:** For complex clock domain systems
- **Reset pulse generator:** For generating reset pulses
- **Reset debouncer:** For mechanical switch inputs

Integration points worth knowing about:

- Clock domain crossing modules
- PLL/clock management units
- Power management controllers
- System controllers and reset managers

## 11.9 Testing

---

### 11.9.1 Test Coverage

Test coverage for the reset\_sync module lives in: - **Testbench Class:**

bin/TBClasses/cdc/reset\_sync\_tb.py - **Test Runner:** val/cdc/test\_reset\_sync.py

### 11.9.2 Test Scenarios

- Reset assertion during various clock phases
- Reset deassertion timing verification
- Multiple reset assertion/deassertion cycles
- Clock jitter during reset release
- Power-on behavior

### 11.9.3 Testbench Scenarios

**Basic Reset Synchronization** - Verifies N-cycle synchronization delay **Immediate Reset**

**Assertion** - Verifies asynchronous assertion behavior **Multiple Reset Cycles** - Tests repeated reset/release sequences **Reset Glitch Filtering** - Validates recovery from short reset pulses

### 11.9.4 Assertions

```
// Verify synchronous deassertion
property sync_deassert;
 @(posedge clk)
 $rose(rst_n) |-> ##N $rose(sync_rst_n);
endproperty
```

```
// Verify asynchronous assertion
property async_assert;
 $fell(rst_n) |-> $fell(sync_rst_n);
endproperty
```

### 11.9.5 Coverage Points

- All reset assertion scenarios
- Reset release with different clock phases
- Multiple consecutive reset cycles
- Parameter variation coverage (different N values)

### 11.9.6 Running Tests

```
Run all reset_sync tests (4 configurations: N=2,3,4,5)
pytest val/cdc/test_reset_sync.py -v
```

```
Run specific configuration
```

```
pytest val/cdc/test_reset_sync.py::test_reset_sync[2-min] -v
```

```
Run with waveform generation
```

```
pytest val/cdc/test_reset_sync.py -v -s
```

### 11.9.7 Test Results

All 4 parameter configurations passing (100% success rate): - N=2 (min) - PASSED - N=3 (typical) - PASSED - N=4 (safe) - PASSED - N=5 (max) - PASSED

**Bug Discovery:** the test suite caught a critical RTL bug on the initial run — an inverted reset polarity. Exactly the kind of bug that sails through a smoke test and ruins someone's bring-up. Thorough verification earns its keep.

## 11.10 Navigation

---

- [← Back to rtl-cdc Index](#)
- [← Back to Main Documentation Index](#)

## 12 clock\_pulse

### 12.1 Overview

The `clock_pulse` module is a periodic pulse generator with a configurable period: it emits a single-cycle pulse every `WIDTH` clock cycles. It's the part you reach for whenever something in the system needs to happen on a schedule — timing generation, heartbeat signals, sampling triggers, periodic events of any kind.

### 12.2 Parameters

| Parameter          | Type             | Default         | Description                                    |
|--------------------|------------------|-----------------|------------------------------------------------|
| <code>WIDTH</code> | <code>int</code> | <code>10</code> | Period of the pulse generation in clock cycles |

`WIDTH` can range from 2 to  $2^{31}-1$ . The parameter is declared `int`, which is 32-bit SIGNED, so  $2^{32}-1$  is not representable. It sets the pulse frequency ( $f_{\text{clk}} / \text{WIDTH}$ ). The pulse itself is always exactly 1 clock cycle wide, which makes the duty cycle  $1/\text{WIDTH}$  — 10% for `WIDTH=10`.

### 12.3 Ports

```
module clock_pulse #(
 parameter int WIDTH = 10 // PERIOD in clock cycles (the pulse
 // always exactly 1 cycle wide)
 itself is
) (
 input logic clk, // Input clock signal
 input logic rst_n, // Input reset signal
 output logic pulse // Output pulse signal
);
```

| Port               | Direction | Width | Description                   |
|--------------------|-----------|-------|-------------------------------|
| <code>clk</code>   | Input     | 1     | System clock input            |
| <code>rst_n</code> | Input     | 1     | Active-low asynchronous reset |
| <code>pulse</code> | Output    | 1     | Periodic pulse output         |

## 12.4 Functional Description

---

### 12.4.1 Internal Counter

```
// WIDTH is the PERIOD; the counter only needs to hold 0..WIDTH-1,
// i.e.
// $clog2(WIDTH) bits (NOT WIDTH bits).
localparam int CW = (WIDTH < 2) ? 1 : $clog2(WIDTH);

logic [CW-1:0] r_counter;
logic [CW-1:0] w_width_minus_one;

// Properly sized period-1 constant
assign w_width_minus_one = CW'(WIDTH - 1);
```

### 12.4.2 Core Logic

```
always_ff @(posedge clk or negedge rst_n) begin
 if (!rst_n) begin
 r_counter <= 'b0;
 pulse <= 'b0;
 end else begin
 if (r_counter < w_width_minus_one)
 r_counter <= r_counter + 1'b1;
 else
 r_counter <= 'b0;

 pulse <= (r_counter == w_width_minus_one);
 end
end
```

### 12.4.3 Operation Principles

1. **Counter:** Free-running counter from 0 to WIDTH-1
2. **Pulse Generation:** pulse is **registered** — pulse <= (r\_counter == WIDTH-1) — so it asserts on the cycle **after** the counter reaches WIDTH-1, i.e. during the cycle when the counter has wrapped back to 0.
3. **Auto-Reset:** Counter wraps to 0 after reaching maximum
4. **Synchronous:** All operations synchronized to input clock
5. **Single Cycle:** Pulse duration is exactly one clock cycle

That second point is the one people get wrong, so let's be blunt about it: the registered comparison means pulse never coincides with the terminal count. It lands one cycle later, on the wrapped-to-zero count. Every timing diagram below follows from that.

## 12.5 Related Modules

Nothing in the tree instantiates this module and it instantiates nothing: it is a leaf, used directly by whatever design needs it. Its nearest neighbours in `rtl/common/` are:

- `clock_divider`
- `clock_gate_ctrl`


## 12.6 Timing Characteristics

### 12.6.1 Timing Characteristics

- **Period:** WIDTH clock cycles
- **Frequency:**  $f_{\text{clk}} / \text{WIDTH}$
- **Pulse Width:** 1 clock cycle
- **Duty Cycle:**  $1/\text{WIDTH}$
- **Phase:** Because the comparison is registered, pulse is high one cycle **after** `r_counter == WIDTH-1` — that is, during the `r_counter == 0` cycle of each period, not on the WIDTH-1 count itself.




### 12.6.2 Basic Operation (WIDTH=4)

Watch where pulse actually lands — high during the `< 0 >` cell that immediately follows `< 3 >`, one cycle after the counter hit WIDTH-1. That's the registered comparison at work:

Clock:   
Counter: `< 0 > < 1 > < 2 > < 3 > < 0 > < 1 > < 2 > < 3 > < 0 >`  
Pulse: 

(The first `< 0 >` is the post-reset state, where pulse is still low.)

### 12.6.3 Reset Behavior

Clock:   
Reset\_n:   
Counter: `< 0 > < 1 > < 2 > < 3 > < 0 > < 0 > < 1 > < 2 > < 3 > < 0 >`  
Pulse: 

Two things worth seeing here, both consequences of the pulse being registered (`pulse <= (r_counter == w_width_minus_one)`):

- **The pulse that reset swallows.** The counter reaches 3 in the fourth cell, so a pulse would normally appear in the fifth — but reset is asserted exactly there, and reset forces `pulse <= 0`. That pulse never happens.

- **Recovery costs a full period.** After reset releases, the counter restarts from 0 and the first pulse appears in the cell *after* it next reaches 3 — the final cell above, not the one where the counter reads 3.

#### 12.6.4 Different WIDTH Values

| WIDTH | Period      | Duty Cycle | Use Case         |
|-------|-------------|------------|------------------|
| 2     | 2 cycles    | 50%        | Clock divider    |
| 4     | 4 cycles    | 25%        | Sampling         |
| 10    | 10 cycles   | 10%        | Timing reference |
| 100   | 100 cycles  | 1%         | Heartbeat        |
| 1000  | 1000 cycles | 0.1%       | Slow events      |

## 12.7 Usage Examples

Every parameter and port below is read from the module declaration.

```
clock_pulse #(
 .WIDTH (10)
) u_clock_pulse (
 .clk (clk),
 .rst_n (rst_n),
 .pulse (pulse)
);
```

## 12.8 Design Notes

### 12.8.1 Resource Utilization

| WIDTH | Counter Bits | LUTs | FFs | Max Freq |
|-------|--------------|------|-----|----------|
| 8     | 3            | 8    | 4   | 500MHz   |
| 16    | 4            | 12   | 5   | 450MHz   |
| 32    | 5            | 18   | 6   | 400MHz   |
| 1024  | 10           | 28   | 11  | 350MHz   |

### 12.8.2 Synthesis Considerations

For high-frequency applications, pipeline the comparison:

```
// For high-frequency applications, pipeline the comparison (WIDTH >= 2).
```

```

// IMPORTANT: because the registered compare_result ALSO gates the
// counter
// wrap, it must be produced from the WIDTH-2 threshold, not WIDTH-1.
Comparing
// against WIDTH-1 here makes the counter sequence 0..WIDTH,0 and
yields a
// period of WIDTH+1, not WIDTH. With WIDTH-2 the period is exactly
WIDTH and
// the pulse still lands on the r_counter==0 cycle, matching the base
module.
module clock_pulse_pipelined #(
 parameter int WIDTH = 1000
) (
 input logic clk,
 input logic rst_n,
 output logic pulse
);

 localparam int CW = (WIDTH < 2) ? 1 : $clog2(WIDTH);
 logic [CW-1:0] r_counter;
 logic compare_result;

 // Register the comparison one count early so the pipelined wrap
 keeps a
 // WIDTH-cycle period.
 always_ff @(posedge clk or negedge rst_n) begin
 if (!rst_n) compare_result <= 1'b0;
 else compare_result <= (r_counter == CW'(WIDTH - 2));
 end

 always_ff @(posedge clk or negedge rst_n) begin
 if (!rst_n) begin
 r_counter <= 'b0;
 pulse <= 1'b0;
 end else begin
 r_counter <= compare_result ? '0 : r_counter + 1'b1;
 pulse <= compare_result;
 end
 end

endmodule

```

### 12.8.3 Design Considerations

1. **Choose Appropriate WIDTH:** Balance between resolution and resource usage
2. **Consider Reset Timing:** Ensure clean startup behavior

3. **Validate Timing:** Verify pulse frequency matches requirements
4. **Plan for Variation:** Consider process, voltage, temperature effects
5. **Monitor Resource Usage:** Large counters can impact timing and area
6. **Use Enables:** Add enable signals for conditional operation
7. **Document Timing:** Clearly specify pulse rates and relationships

#### 12.8.4 Common Applications

1. **Timing References:** System heartbeats, watchdog timers
2. **Sampling Systems:** ADC triggers, data acquisition timing
3. **Communication:** Baud rate generation, protocol timing
4. **Memory Systems:** Refresh timing, access scheduling
5. **Test Equipment:** Pattern generation, stimulus timing
6. **Power Management:** Activity monitoring, timeout generation
7. **Display Systems:** Refresh rates, sync generation

That's the whole module. It's a small thing, but precise periodic timing shows up in nearly every design you'll ever ship — get this one right once, parameterize it well, and reuse it forever.

## 12.9 Testing

---

### 12.9.1 Comprehensive Test Bench

```
module tb_clock_pulse;

 parameter WIDTH = 10;

 logic clk, rst_n, pulse;

 // Clock generation
 initial clk = 0;
 always #5ns clk = ~clk; // 100MHz

 // DUT instantiation
 clock_pulse #(.WIDTH(WIDTH)) dut (.*);

 // Test sequence
 initial begin
 rst_n = 0;
 #100ns rst_n = 1;

 // Test basic pulse generation
 test_basic_pulse_generation();
 end
endmodule
```

```

 // Test reset behavior
 test_reset_behavior();

 // Test timing accuracy
 test_timing_accuracy();

 $display("All tests completed!");
 $finish;
end

// Test basic pulse generation
task test_basic_pulse_generation();
 int pulse_count = 0;
 int cycle_count = 0;
 begin
 $display("Testing basic pulse generation (WIDTH=%0d)...",
WIDTH);

 // Monitor for several pulse periods. WIDTH*3 + 1 edges:
 `@(posedge clk)
 // resumes BEFORE the NBA update, so pulses are first
 visible at edges
 // WIDTH+1, 2*WIDTH+1, 3*WIDTH+1 (see test_reset_behavior
 for the trace).
 repeat (WIDTH * 3 + 1) begin
 @(posedge clk);
 cycle_count++;
 if (pulse) pulse_count++;
 end

 // Verify pulse count
 if (pulse_count == 3) begin
 $display("PASS: Correct number of pulses generated");
 end else begin
 $error("FAIL: Expected 3 pulses, got %0d",
pulse_count);
 end
 end
endtask

// Test reset behavior -- the two consequences of the registered
pulse
task test_reset_behavior();
 int pulses_during_reset = 0;
 int cycles_to_first_pulse = 0;

```

```

begin
 $display("Testing reset behavior...");

 // Line the counter up so a pulse is due, then assert
 reset in // exactly that cycle. Reset forces pulse <= 0, so that
 pulse is // swallowed and never appears.
 repeat (WIDTH - 1) @(posedge clk);
 rst_n = 0;
 repeat (WIDTH * 2) begin
 @(posedge clk);
 if (pulse) pulses_during_reset++;
 end
 if (pulses_during_reset == 0) begin
 $display("PASS: no pulse while reset is asserted");
 end else begin
 $error("FAIL: %0d pulse(s) during reset",
pulses_during_reset);
 end

 // Recovery costs a full period: the counter restarts from
 0, so the // first pulse lands one cycle AFTER it next reaches
 WIDTH-1.
 @(negedge clk);
 rst_n = 1;
 // Expect the first pulse on the (WIDTH+1)th sampled edge,
 not the // WIDTHth. Three separate cycles are involved and it is
 easy to // lose one: the counter reaches WIDTH-1 at edge WIDTH-1,
 `pulse` // is NBA-assigned at edge WIDTH, and `@(posedge clk)`
 resumes in // the Active region -- BEFORE that edge's NBA update --
 so the // high value is first visible at edge WIDTH+1.
 forever begin
 @(posedge clk);
 cycles_to_first_pulse++;
 if (pulse) break;
 if (cycles_to_first_pulse > WIDTH * 2) begin
 $error("FAIL: no pulse within two periods of reset
release");
 end
 end
end

```

```

 if (cycles_to_first_pulse == WIDTH + 1) begin
 $display("PASS: first pulse %0d sampled edges after
release",
 WIDTH + 1);
 end else begin
 $error("FAIL: expected first pulse at cycle %0d, got
%0d",
 WIDTH + 1, cycles_to_first_pulse);
 end
 end
endtask

// Test timing accuracy
task test_timing_accuracy();
 time pulse_times[10];
 time pulse_periods[9];
 int i;
 real avg_period, expected_period;
 begin
 $display("Testing timing accuracy...");

 expected_period = WIDTH * 10.0; // 10ns clock period

 // Capture pulse timing
 for (i = 0; i < 10; i++) begin
 @(posedge pulse);
 pulse_times[i] = $time;
 end

 // Calculate periods
 for (i = 0; i < 9; i++) begin
 pulse_periods[i] = pulse_times[i+1] - pulse_times[i];
 end

 // Calculate average period
 avg_period = 0;
 for (i = 0; i < 9; i++) begin
 avg_period += pulse_periods[i];
 end
 avg_period = avg_period / 9.0;

 // Check accuracy (within 1% tolerance)
 if (abs(avg_period - expected_period) < expected_period *
0.01) begin
 $display("PASS: Average period = %.1f ns (expected
%.1f ns)",
 avg_period, expected_period);
 end
 end
endtask

```

```

 end else begin
 $error("FAIL: Average period = %.1f ns (expected %.1f
ns)",
 avg_period, expected_period);
 end
 end
endtask

function real abs(real value);
 abs = (value >= 0) ? value : -value;
endfunction

endmodule

```

## 12.9.2 Coverage Model

```

covergroup clock_pulse_cg @(posedge clk);

 cp_counter_values: coverpoint dut.r_counter {
 bins zero = {0};
 bins low[] = {[1:WIDTH/4]};
 bins mid[] = {[WIDTH/4+1:3*WIDTH/4]};
 bins high[] = {[3*WIDTH/4+1:WIDTH-2]};
 bins max = {WIDTH-1};
 }

 cp_pulse: coverpoint pulse {
 bins asserted = {1};
 bins deasserted = {0};
 }

 cp_reset: coverpoint rst_n {
 bins reset = {0};
 bins normal = {1};
 }

 // Transition coverage
 cp_pulse_edges: coverpoint pulse {
 bins rising = (0 => 1);
 bins falling = (1 => 0);
 bins stable_low = (0 => 0);
 bins stable_high = (1 => 1);
 }

 // Cross coverage
 cross_counter_pulse: cross cp_counter_values, cp_pulse;

```

endcovergroup

### 12.9.3 Formal Properties

**module** clock\_pulse\_properties;

```
// Bind to DUT
bind clock_pulse clock_pulse_properties props_inst (.*);

// Property: reaching max count produces a pulse on the NEXT
cycle.
// The comparison is registered (pulse <= r_counter == WIDTH-1),
so the
// implication is |=> (next cycle), NOT |-> (same cycle).
property pulse_after_max_count;
 @(posedge clk) disable iff (!rst_n)
 (dut.r_counter == WIDTH-1) |=> pulse;
endproperty

// Property: a pulse implies the PREVIOUS count was max
property pulse_implies_prev_max;
 @(posedge clk) disable iff (!rst_n)
 pulse |-> $past(dut.r_counter == WIDTH-1);
endproperty

// Property: Counter wraps correctly
property counter_wrap;
 @(posedge clk) disable iff (!rst_n)
 (dut.r_counter == WIDTH-1) |=> (dut.r_counter == 0);
endproperty

// Property: Counter increments
property counter_increment;
 @(posedge clk) disable iff (!rst_n)
 (dut.r_counter < WIDTH-1) |=>
 (dut.r_counter == $past(dut.r_counter) + 1);
endproperty

// Property: Reset behavior
property reset_behavior;
 @(posedge clk)
 !rst_n |=> (dut.r_counter == 0) && !pulse;
endproperty

// Assertions
assert property (pulse_after_max_count);
```

```
 assert property (pulse_implies_prev_max);
 assert property (counter_wrap);
 assert property (counter_increment);
 assert property (reset_behavior);
```

```
endmodule
```

## 12.10 Navigation

---

- [← Back to rtl-common Index](#)
- [← Back to Main Documentation Index](#)

## 13 apb4\_slave\_cdc

**Module:** apb4\_slave\_cdc.sv **Location:** rtl/amba/apb4/ **Status:** Production Ready

---

### 13.1 Overview

---

The APB Slave CDC (Clock Domain Crossing) module is a complete APB slave interface with the crossing built in: full APB on the pclk side, an AXI/GAXI-style backend on the aclk side. It's what lets you integrate APB peripherals running at different clock frequencies without metastability becoming your problem.

#### 13.1.1 Key Features

- **Safe CDC:** Gray-pointer asynchronous FIFOs, one per direction
- **Dual Clock Domains:** APB (pclk) and AXI (aclk) operate independently
- **Caveat: reset both domains together** — a one-sided reset is NOT safe; see [Reset Behavior](#)
- **Full APB4 Support:** Complete AMBA 4 APB protocol compliance
- **Command/Response Interface:** Clean GAXI-style backend interface
- **Buffered Operation:** Integrated skid buffers for elastic storage

**Protocol scope:** APB4 only. For APB5 signalling use apb5\_slave\_cdc from rtl/amba/apb5/ — see the [APB5 book](#).

---

### 13.2 Parameters

---

| Parameter  | Type | Default      | Description                                                                                   |
|------------|------|--------------|-----------------------------------------------------------------------------------------------|
| ADDR_WIDTH | int  | 32           | APB address bus width                                                                         |
| DATA_WIDTH | int  | 32           | APB data bus width                                                                            |
| STRB_WIDTH | int  | DATA_WIDTH/8 | Write strobe width (calculated)                                                               |
| PROT_WIDTH | int  | 3            | APB protection signal width                                                                   |
| DEPTH      | int  | 2            | Skid-buffer depth <b>in entries</b> inside the wrapped apb4_slave; also the floor for the CDC |

| Parameter       | Type | Default  | Description                                                                                                                                                                                                                                                                                                                             |
|-----------------|------|----------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| USE_JOHNSON     | int  | 0 (Gray) | FIFO depth<br>CDC-FIFO pointer encoding: 0 Gray (power-of-2 derived depth only), 1 Johnson (any FIFO depth, max(DEPTH, 4)-bit pointers; module-level DEPTH stays limited to 2..8 inclusive by the skid buffers). Gray by default — Johnson is opt-in because its pointers cost DEPTH bits in both domains and every synchronizer stage. |
| USE_2_PHASE_CDC | bit  | 1        | <b>Deprecated and ignored.</b> Has no effect on the generated logic                                                                                                                                                                                                                                                                     |

### 13.3 Ports

```

module apb4_slave_cdc #(
 parameter int ADDR_WIDTH = 32,
 parameter int DATA_WIDTH = 32,
 parameter int STRB_WIDTH = DATA_WIDTH / 8,
 parameter int PROT_WIDTH = 3,
 parameter int DEPTH = 2,
 // Async-FIFO pointer encoding: 0 = Gray (power-of-2 derived depth
 only),
 // 1 = Johnson (any depth, DEPTH-bit pointers). Forwarded to both
 CDC FIFOs.
 parameter int USE_JOHNSON = 0,
 // DEPRECATED / NO EFFECT -- retained only so existing
 instantiations
 // still elaborate. See "CDC Implementation" below.
 parameter bit USE_2_PHASE_CDC = 1'b1,
 // Derived width aliases the port list references -- not
 independent knobs
 parameter int DW = DATA_WIDTH,
 parameter int AW = ADDR_WIDTH,

```

```

parameter int SW = STRB_WIDTH,
parameter int PW = PROT_WIDTH
) (
 // Clock and Reset
 input logic aclk,
 input logic aresetn,
 input logic pclk,
 input logic presetn,

 // APB Slave Interface (pclk domain)
 input logic s_apb_PSEL,
 input logic s_apb_PENABLE,
 output logic s_apb_PREADY,
 input logic [AW-1:0] s_apb_PADDR,
 input logic s_apb_PWRITE,
 input logic [DW-1:0] s_apb_PWDATA,
 input logic [SW-1:0] s_apb_PSTRB,
 input logic [PW-1:0] s_apb_PPROT,
 output logic [DW-1:0] s_apb_PRDATA,
 output logic s_apb_PSLVERR,

 // Command Interface (aclk domain)
 output logic cmd_valid,
 input logic cmd_ready,
 output logic cmd_pwrite,
 output logic [AW-1:0] cmd_paddr,
 output logic [DW-1:0] cmd_pwdata,
 output logic [SW-1:0] cmd_pstrb,
 output logic [PW-1:0] cmd_pprot,

 // Response Interface (aclk domain)
 input logic rsp_valid,
 output logic rsp_ready,
 input logic [DW-1:0] rsp_prdata,
 input logic rsp_pslverr
);

```

---

## 13.4 Functional Description

### 13.4.1 Clock Domains

APB domain (pclk):

- APB slave interface signals
- Typical frequency: 50-200 MHz

- Used by APB master/interconnect

#### AXI domain (aclk):

- Command and response interfaces
- Can be faster or slower than pclk
- Used by backend processing logic

### 13.4.2 CDC Structure

The module is an apb4\_slave in the pclk domain plus two independent asynchronous FIFOs:

| FIFO           | Direction | Write domain   | Read domain    | Payload width               |
|----------------|-----------|----------------|----------------|-----------------------------|
| u_cmd_cdc_fifo | Command   | pclk / presetn | aclk / aresetn | CPW = AW + DW + SW + PW + 1 |
| u_rsp_cdc_fifo | Response  | aclk / aresetn | pclk / presetn | RPW = DW + 1                |

Both are gaxi\_fifo\_async instances with:

- **Pointer encoding:** absolute read/write pointers, Gray by default (USE\_JOHNSON = 0); pass USE\_JOHNSON = 1 for Johnson
- **Synchronizer depth:** N\_FLOP\_CROSS = 2 (two-flop synchronizer per crossed pointer)
- **FIFO depth:** CDC\_FIFO\_DEPTH = (DEPTH < 4) ? 4 : DEPTH — a floor of 4 entries regardless of the DEPTH used for the internal skid buffers. **Under the default USE\_JOHNSON = 0 a power of two is required**, because Gray carries a generate-scope elaboration check ((USE\_JOHNSON == 0) && ((DEPTH & (DEPTH-1)) != 0) -> \$error). The check sees the DERIVED depth max(DEPTH, 4), while the same DEPTH also reaches the wrapped apb4\_slave's two gaxi\_skid\_buffer instances UNFLOORED, whose guard accepts 2..8 inclusive (any integer). Under the default Gray encoding the floored depth must be a power of two, so the legal DEPTH set is **{2, 3, 4, 8}** (2 and 3 floor to 4); 5, 6 and 7 floor to themselves, fail the power-of-2 check, and elaborate only with USE\_JOHNSON = 1 – Johnson buys exactly those three depths. It pays for that with pointers as wide as the derived FIFO depth max(DEPTH, 4) against Gray's \$clog2(max(DEPTH, 4))+1 – at the default DEPTH = 2 that is 4 bits vs 3, not 2 vs 2 – in both domains and every synchronizer stage; Gray is the default so that cost is never paid by accident.

There is no separate metastability-hardening option — two-flop synchronization is fixed at instantiation.

### 13.4.3 Reset Behavior

`presetn` and `aresetn` are separate reset domains, but they must be asserted together. **A one-sided reset is not safe. Quiesce the bus first.**

The local reset clears that domain's own pointer, but the crossed copy of the *remote* pointer is a live synchronizer (`glitch_free_n_dff_arb`,  $N=2$ ) that keeps sampling the non-reset domain the moment reset deasserts — it does not hold at zero. The reset side comes back with its own pointer at zero and the remote pointer at whatever the other side had reached: not an empty FIFO, a mismatched one.

Neither consequence is a clean discard:

- **Consumed commands are re-presented.** Pulsing `aresetn` with an **already-consumed** command in the cmd FIFO rewinds the read pointer behind the write pointer, so the backend sees that command again after reset. It does not time out — it re-executes. An *unread* command is not at risk: its read pointer is already behind the write pointer, so resetting it to 0 rewinds nothing and the entry is delivered exactly once.
- **The response FIFO can fabricate responses.** The same rewind on the response path presents entries the APB side never queued, so the APB master can complete a transfer the backend never answered.

Pointers being absolute positions rather than toggle parity is what makes the *steady-state* crossing reliable; it does not make a one-sided reset safe.

### 13.4.4 Why Not the Previous 2-Phase Handshake

`USE_2_PHASE_CDC` selected a toggle-based handshake in an earlier revision. It is now deprecated and ignored; the parameter survives only for source-compatibility with existing instantiations.

The failure mode is worth understanding, because it's the kind that ships. A 2-phase handshake encodes each transfer as a **toggle**. Reset the two domains independently and the toggle parity desynchronizes — the link fabricates or drops exactly one transfer. Permanently, because nothing ever re-syncs it. Pair that with the `apb4_slave` FSM, which pairs commands and responses by position rather than by tag, and a single phantom transfer offsets the response stream by one forever: every read returns the previous read's data.

This was observed on the Nexys A7 ddr2-char board on 2026-07-19. Reading a single CSR eight times returned the previous register's value about three times before settling, while the non-CDC harness window was stable. The two mitigations now in place are the gray-pointer FIFOs described above and the orphan-response guard in `apb4_slave` — see [apb4\\_slave.md](#).

---

## 13.5 Design Notes

---

**This is a clock-domain crossing; the usual CDC rules bind.** Pointers cross Gray-coded through `glitch_free_n_dff_arn`, and the reset must be asserted in both domains – releasing one before the other leaves the pointer comparison meaningless until the second releases.

**Depth buys latency tolerance, not bandwidth.** A CDC FIFO's throughput is set by the slower domain; a deeper buffer absorbs a longer burst without stalling but does not raise the sustained rate. See `vault/handbook/design/cdc.md`.

---

## 13.6 Timing Characteristics

---

### 13.6.1 Maximum Clock Ratio

There is no maximum ratio between `pclk` and `acclk`. Gray-pointer FIFOs impose no relationship between the two clocks — either may be arbitrarily faster, slower, or phase-unrelated, and either may be stopped indefinitely. Stopping `acclk` simply stalls the command FIFO's read side; the APB side backpressures via `PREADY` held low and no data is lost.

What the ratio affects is throughput, not correctness. Each transfer pays the usual two-flop synchronizer latency in each direction (roughly 2-3 destination clock edges per crossing), so a very slow `acclk` directly lengthens APB wait states.

### 13.6.2 Timing Constraints

The crossed signals are the gray-coded pointers and the FIFO memory read path. Standard practice applies:

- Declare `pclk` and `acclk` asynchronous to each other (`set_clock_groups - asynchronous`).
  - Do not over-constrain the pointer synchronizer paths; the gray encoding tolerates a one-bit-at-a-time skew by construction.
- 

## 13.7 Waveforms


---

The following timing diagrams show CDC behavior with **both clock domains visible**:

**Clock Configuration:** - `apb_pclk`: 100MHz (10ns period) - `cmd_acclk`: 500MHz (2ns period), displayed with `period=0.4` for visual compactness

### 13.7.1 Scenario 1: Write Transaction with CDC

Shows APB write crossing from pclk domain to aclk domain:




*Write CDC*

WaveJSON: [apb\\_write\\_sequence\\_001.json](#)

**Key Observations:** - APB transaction in pclk domain - CMD transaction crosses to aclk domain - Note the CDC latency between domains

### 13.7.2 Scenario 2: Read Transaction with CDC

Shows APB read with response crossing back from aclk to pclk domain:




*Read CDC*

WaveJSON: [apb\\_read\\_sequence\\_001.json](#)

**Key Observations:** - APB read request in pclk domain - CMD crosses to aclk domain - RSP crosses back to pclk domain - Complete round-trip CDC visible


### 13.7.3 Scenario 3: Back-to-Back Writes with CDC



*B2B Writes CDC*

WaveJSON: [apb\\_back\\_to\\_back\\_writes\\_001.json](#)


### 13.7.4 Scenario 4: Back-to-Back Reads with CDC



*B2B Reads CDC*

WaveJSON: [apb\\_back\\_to\\_back\\_reads\\_001.json](#)


### 13.7.5 Scenario 5: Write-to-Read Transition with CDC



*Write-to-Read CDC*

WaveJSON: [apb\\_write\\_to\\_read\\_001.json](#)


### 13.7.6 Scenario 6: Read-to-Write Transition with CDC



#### Read-to-Write CDC

WaveJSON: [apb\\_read\\_to\\_write\\_001.json](#)

### 13.7.7 Scenario 7: Error Response with CDC



#### Error CDC

WaveJSON: [apb\\_error\\_001.json](#)

## 13.8 Usage Examples

```
apb4_slave_cdc #(
 .ADDR_WIDTH(32),
```

```

 .DATA_WIDTH(32),
 .DEPTH(2)
) u_apb_cdc (
 // APB clock domain
 .pclk (apb_clk),
 .presetn (apb_resetsn),

 // AXI clock domain
 .aclk (axi_clk),
 .aresetsn (axi_resetsn),

 // APB slave interface (pclk domain)
 .s_apb_PSEL (apb_psel),
 .s_apb_PENABLE (apb_penable),
 .s_apb_PREADY (apb_pready),
 .s_apb_PADDR (apb_paddr),
 .s_apb_PWRITE (apb_pwrite),
 .s_apb_PWDATA (apb_pwdata),
 .s_apb_PSTRB (apb_pstrb),
 .s_apb_PPROT (apb_pprot),
 .s_apb_PRDATA (apb_prdata),
 .s_apb_PSLVERR (apb_pslverr),

 // Command interface (aclk domain)
 .cmd_valid (cmd_valid),
 .cmd_ready (cmd_ready),
 .cmd_pwrite (cmd_pwrite),
 .cmd_paddr (cmd_paddr),
 .cmd_pwdata (cmd_pwdata),
 .cmd_pstrb (cmd_pstrb),
 .cmd_pprot (cmd_pprot),

 // Response interface (aclk domain)
 .rsp_valid (rsp_valid),
 .rsp_ready (rsp_ready),
 .rsp_prdata (rsp_prdata),
 .rsp_pslverr (rsp_pslverr)
);

```

---

## 13.9 Related Modules

---

- **APB Slave:** [apb4\\_slave.md](#)
- **Clock-Gated Variant:** [apb4\\_slave\\_cdc\\_cg.md](#)
- **APB5 Equivalent:** [apb5\\_slave\\_cdc.md](#)

- **CDC FIFO:** rtl/cdc/gaxi\_fifo\_async.sv
- 

## 13.10 Testing

---

val/amba/test\_apb4\_slave\_cdc.py exercises this module. It collects 2 parameter cases at the default REG\_LEVEL.

```
source env_python
pytest val/amba/test_apb4_slave_cdc.py -v
```

---

## 13.11 References

---

- **Source:** rtl/amba/apb4/apb4\_slave\_cdc.sv
  - **Tests:** val/amba/test\_apb4\_slave\_cdc.py
  - **WaveDrom Test:**  
val/amba/test\_apb4\_slave\_cdc.py::test\_apb4\_slave\_cdc\_wavedrom
- 

**Last Updated:** 2026-07-19

---

## 13.12 Navigation

---

- [← Back to rtl-amba Index](#)
- [← Back to Main Documentation Index](#)

## 14 apb4\_slave\_cdc\_cg

**Module:** apb4\_slave\_cdc\_cg.sv **Base Module:** [apb4\\_slave\\_cdc](#) **Location:** rtl/amba/apb4/  
**Status:** Production Ready

---

### 14.1 Overview

---

This is the **clock-gated variant** of [apb4\\_slave\\_cdc](#).

**For the clock-gating architecture and the underlying gate cell, see:**

→ [Clock-Gated Variants Guide](#) and → [amba\\_clock\\_gate\\_ctrl](#)

Both live in the Shared book, which is published as a separate PDF. Take the parameter and port names for this module from the tables below, not from the generic guide.

The apb4\_slave\_cdc\_cg module adds power optimization to apb4\_slave\_cdc through activity-based clock gating. **Both clock domains are gated independently** — the module instantiates two amba\_clock\_gate\_ctrl blocks, one on pclk and one on aclk, sharing the same configuration inputs.

Structurally it's a sibling of apb4\_slave\_cdc, not a wrapper around it: it re-instantiates the same apb4\_slave plus the same pair of gaxi\_fifo\_async CDC FIFOs, but drives them from the gated clocks. The CDC behaviour described in [apb4\\_slave\\_cdc.md](#) — gray pointers, N\_FLOP\_CROSS=2, no maximum clock ratio — applies unchanged. That includes the reset rule: a one-sided reset is NOT safe here either, for the same reason (the crossed pointer copy is a live synchronizer, not a snapshot).

While a domain is gated, that domain's ready outputs are forced low, so no handshake can complete against a stopped clock.

- **Same Functionality:** 100% equivalent to base module
  - **Dual-Domain Gating:** Separate gate cell per clock domain
  - **Runtime Control:** Gating is enabled and tuned by input signals, not parameters
  - **Bypassable:** Tie cfg\_cg\_enable = 0 for behaviour identical to the base module
- 

### 14.2 Parameters

---

In addition to all [apb4\\_slave\\_cdc](#) parameters:

| Parameter           | Type | Default | Description                                                 |
|---------------------|------|---------|-------------------------------------------------------------|
| CG_IDLE_COUNT_WIDTH | int  | 4       | Width of the idle countdown counter, shared by both domains |

USE\_2\_PHASE\_CDC is inherited from the base module and is likewise **deprecated and ignored**. There is no ENABLE\_CLOCK\_GATING parameter and no per-domain CG\_GATE\_\* parameters.

### 14.2.1 Derived Parameters (do not override)

These are declared as parameter so the elaborator can compute them, not so callers can set them. Each defaults to an expression over the parameters above; overriding one desynchronises it from its source and the design fails to elaborate or silently mis-sizes a bus. Set the parameters they are derived FROM and leave these alone.

| Derived parameter | Default expression    |
|-------------------|-----------------------|
| STRB_WIDTH        | DATA_WIDTH / 8        |
| DW                | DATA_WIDTH            |
| AW                | ADDR_WIDTH            |
| SW                | STRB_WIDTH            |
| PW                | PROT_WIDTH            |
| CPW               | AW + DW + SW + PW + 1 |
| RPW               | DW + 1                |

## 14.3 Ports

In addition to all [apb4\\_slave\\_cdc](#) ports:

| Port              | Width               | Direction | Description                                                                                        |
|-------------------|---------------------|-----------|----------------------------------------------------------------------------------------------------|
| cfg_cg_enable     | 1                   | Input     | Global clock-gate enable for both domains. 0 = never gate                                          |
| cfg_cg_idle_count | CG_IDLE_COUNT_WIDTH | Input     | Idle cycles to count down before gating                                                            |
| cg_gating         | 1                   | Output    | Asserted while the pclk domain is gated                                                            |
| cg_idle           | 1                   | Output    | Asserted one cycle after the pclk-domain activity terms go low. <b>Not an occupancy flag</b> — see |

| Port      | Width | Direction | Description                                                                                              |
|-----------|-------|-----------|----------------------------------------------------------------------------------------------------------|
|           |       |           | below                                                                                                    |
| cg_gating | 1     | Output    | Asserted while the aclk domain is gated                                                                  |
| cg_idle   | 1     | Output    | Asserted one cycle after the aclk-domain activity terms go low. <b>Not an occupancy flag</b> — see below |

**\*\_cg\_idle does not know whether the FIFOs are empty.**

amba\_clock\_gate\_ctrl has no occupancy input at all — its whole idle logic is `r_wakeup <= user_valid || axi_valid`; and assign `idle = ~r_wakeup`. So idle asserts one cycle after the activity terms drop, whatever the CDC FIFOs hold. On this wrapper the activity terms mostly hide occupancy anyway: APB holds `s_apb_PSEL` until `PREADY`, and this slave asserts `PREADY` only when the response returns, so a stalled backend keeps `pclk_user_valid` high and `cg_idle` low for the whole stall; an unconsumed response likewise holds `w_rsp_valid/rsp_valid` high. Do not read that as “idle implies empty” — it does not, here or anywhere: idle is “nothing has been handed to me recently”, the input to the gating countdown.

**Do not use it as a safe-to-reset or safe-to-power-down qualifier;** that requires an emptiness check this signal does not perform.

### 14.3.1 Clock-gating status outputs

This wrapper spans TWO clock domains and gates each independently, so it carries three pairs rather than one. They are not interchangeable.

| Port           | Dir | Width | Description                                                                         |
|----------------|-----|-------|-------------------------------------------------------------------------------------|
| pclk_cg_gating | Out | 1     | APB-domain clock gated. Masks <code>s_apb_PREADY</code> and the command-side ready. |
| pclk_cg_idle   | Out | 1     | APB-domain                                                                          |

| Port           | Dir | Width | Description                                                                                               |
|----------------|-----|-------|-----------------------------------------------------------------------------------------------------------|
| aclk_cg_gating | Out | 1     | activity terms quiet.<br>AXI-domain clock gated. Masks the response-side ready.                           |
| aclk_cg_idle   | Out | 1     | AXI-domain activity terms quiet.                                                                          |
| cg_gating      | Out | 1     | Whole-module aggregate:<br>pclk_cg_gating &&<br>aclk_cg_gating.<br>High only when BOTH domains are gated. |
| cg_idle        | Out | 1     | Whole-module aggregate:<br>pclk_cg_idle &&<br>aclk_cg_idle.                                               |

Use the per-domain signals to reason about one side's readys; use `cg_gating` only as the "is this block asleep" summary. Treating `cg_gating` as equivalent to `pclk_cg_gating` is wrong: with one domain busy the aggregate stays low while that domain's readys are still masked.

## 14.4 Functional Description

This wrapper is the base module plus one `amba_clock_gate_ctrl` instance. The datapath is untouched – every channel signal is forwarded verbatim – so functional behaviour is identical to the base module and the wrapper adds no latency of its own.

What it adds is a gated clock. `amba_clock_gate_ctrl` watches two activity terms, `user_valid` (this side's valids plus the base module's busy) and `axi_valid` (the far side's valids), registers their OR into `r_wakeup`, and stops the inner module's clock once both have been quiet for `cfg_cg_idle_count` cycles. The clock restarts on the next activity, one cycle later.

While the clock is stopped the wrapper masks its outward-facing READY signals with !  
cg\_gating, so a peer sees no acceptance until the clock runs again and no handshake is lost  
across the wake boundary.

cfg\_cg\_enable arms this behaviour; with it low the clock free-runs and the module is  
indistinguishable from its base.

---

## 14.5 Timing Characteristics

---

This module is **sequential**: it contains 1 always\_ff block(s), clocked on aclk with active-low  
asynchronous reset aresetn. Outputs derived in those blocks are registered and therefore appear  
one clock after the inputs that produced them.

Per-path cycle counts are not enumerated here; read the block that drives the signal you care  
about. No synthesis frequency or area figures are quoted – none have been measured against a  
target device.

---

## 14.6 Usage Examples

---

```
apb4_slave_cdc_cg #(
 // Base module parameters (see apb4_slave_cdc.md)
 .ADDR_WIDTH (32),
 .DATA_WIDTH (32),
 .DEPTH (2),
 .USE_JOHNSON (0), // 0 = Gray: DEPTH in {2,3,4,8}
 (floored_max(DEPTH,4) must be pow2); 1 = Johnson: any DEPTH 2..8

 // Clock gating
 .CG_IDLE_COUNT_WIDTH (4)
) u_cg (
 .pclk (apb_clk),
 .presetn (apb_resetn),
 .aclk (core_clk),
 .aresetn (core_resetn),

 // Clock gating control (applies to both domains)
 .cfg_cg_enable (1'b1),
 .cfg_cg_idle_count (4'd8),

 // Per-domain gating status
 .cg_gating (pclk_gated),
 .cg_idle (pclk_idle),
 .cg_gating (aclk_gated),
```

```
.cg_idle (aclk_idle),

 // ... all other ports same as apb4_slave_cdc
);
```

**Simulation note:** set `cfg_cg_enable = 0` when debugging CDC waveforms. Two independently gated clocks make cross-domain timing very hard to read.

---

## 14.7 Design Notes

---

**A peer's READY must never enter the activity term.** A consumer that parks its response-ready high while idle is behaving correctly; folding that signal into `user_valid` pins this block permanently awake and defeats gating entirely – silently, because function is unaffected. Ten wrappers in this repository shipped that way and nothing failed until someone measured. `val/amba/test_cg_peer_ready.py` parks the peer READY high, holds every VALID low, and requires `cg_gating`. Canonical rule: `vault/handbook/design/clock-gating-activity-terms.md`.

**cfg\_cg\_enable is not a kill switch.** It arms gating and reaches `amba_clock_gate_ctrl` only; the datapath and any monitor enables are forwarded untouched. With it low the clock free-runs and this module behaves exactly like its base.

**Gating latency.** The clock stops `cfg_cg_idle_count + 2` cycles after the last bus activity – the idle counter, plus one for the `r_wakeup` flop. Size the idle count against your traffic's inter-burst gap: too small and the block wakes constantly, too large and it never gates.

**Cost.** Five flops: `r_wakeup` plus `r_idle_counter` at `IDLE_CNTR_WIDTH`, scaling as `1 + CG_IDLE_COUNT_WIDTH`. The ICG itself is a latch or `BUFGCE`, not fabric flops.

---

## 14.8 Related Modules

---

- **Base Module Functionality:** [apb4\\_slave\\_cdc.md](#)
  - **Clock Gating Guide:** [clock\\_gated\\_variants.md](#)
  - **Detailed CG Examples:**
    - [axi4\\_master\\_rd\\_mon\\_cg.md](#) (AXI4 monitor)
    - [axil4\\_master\\_rd\\_mon\\_cg.md](#) (AXIL4 monitor)
    - [apb4\\_slave\\_cg.md](#) (APB interface)
-

## 14.9 Testing

---

`val/amba/test_apb4_slave_cdc_cg.py` exercises this module. It collects 2 parameter cases at the default `REG_LEVEL`.

```
source env_python
pytest val/amba/test_apb4_slave_cdc_cg.py -v
```

---

## 14.10 Navigation

---

- [← Back to APB Index](#)
- [← Back to rtl-amba Index](#)
- [← Back to Main Documentation Index](#)

## 15 APB5 Slave CDC

**Module:** apb5\_slave\_cdc.sv **Location:** rtl/amba/apb5/ **Status:** Production Ready

---

### 15.1 Overview

---

The APB5 Slave CDC module provides clock domain crossing between an APB5 bus clock domain and a backend clock domain: full APB5 slave on the pclk side, a command/response backend interface on the aclk side, and transactions carried safely across the asynchronous boundary in between.

#### 15.1.1 Key Features

- Full APB5 protocol support with CDC
  - Asynchronous FIFO-based clock domain crossing (Gray pointers by default; Johnson available via USE\_JOHNSON, and opt-in by design)
  - All APB5 user signals (PAUSER, PWUSER, PRUSER, PBUSER)
  - Wake-up request crossing via a dedicated level synchronizer
  - Single DEPTH parameter sizes both directions
  - Metastability protection with 2-flop pointer synchronizers
- 

### 15.2 Parameters

---

| Parameter   | Type | Default      | Description                     |
|-------------|------|--------------|---------------------------------|
| ADDR_WIDTH  | int  | 32           | APB address bus width           |
| DATA_WIDTH  | int  | 32           | APB data bus width              |
| STRB_WIDTH  | int  | DATA_WIDTH/8 | Write strobe width (calculated) |
| PROT_WIDTH  | int  | 3            | Protection signal width         |
| AUSER_WIDTH | int  | 4            | Address user signal width       |
| WUSER_WIDTH | int  | 4            | Write user signal width         |
| RUSER_WIDTH | int  | 4            | Read user signal width          |
| BUSER_WIDTH | int  | 4            | Response user signal width      |
| DEPTH       | int  | 2            | Skid-buffer depth of the        |

| Parameter       | Type | Default  | Description                                                                                                                     |
|-----------------|------|----------|---------------------------------------------------------------------------------------------------------------------------------|
| USE_JOHNSON     | int  | 0 (Gray) | wrapped apb5_slave;<br>2..8 inclusive (any integer)<br>CDC-FIFO pointer encoding: 0 Gray, 1 Johnson, -1 auto-select. See below. |
| ENABLE_PARITY   | bit  | 0        | Enable parity generation and checking                                                                                           |
| USE_2_PHASE_CDC | bit  | 1        | Deprecated and ignored — retained for source compatibility                                                                      |

There is no CMD\_DEPTH, RSP\_DEPTH or SYNC\_STAGES parameter on this module. The two asynchronous FIFOs are sized internally from DEPTH:

```
localparam int CDC_FIFO_DEPTH = (DEPTH < 4) ? 4 : DEPTH;
```

so the CDC FIFOs are never shallower than 4 entries regardless of DEPTH. The synchronizer depth is fixed at 2 flops (N\_FLOP\_CROSS(2) on both gaxi\_fifo\_async instances) and is not exposed as a parameter. If your design needs 3-flop synchronization for an extreme clock ratio, you're editing the instantiation in the RTL.

### 15.2.1 Derived Parameters (do not override)

These are declared as parameter so the elaborator can compute them, not so callers can set them. Each defaults to an expression over the parameters above; overriding one desynchronises it from its source and the design fails to elaborate or silently mis-sizes a bus. Set the parameters they are derived FROM and leave these alone.

| Derived parameter | Default expression        |
|-------------------|---------------------------|
| DW                | DATA_WIDTH                |
| AW                | ADDR_WIDTH                |
| SW                | STRB_WIDTH                |
| PW                | PROT_WIDTH                |
| AUW               | AUSER_WIDTH               |
| WUW               | WUSER_WIDTH               |
| RUW               | RUSER_WIDTH               |
| BUW               | BUSER_WIDTH               |
| CPW               | AW + DW + SW + PW + AUW + |

| Derived parameter | Default expression                |
|-------------------|-----------------------------------|
| RPW               | $WUW + 1$<br>$DW + RUW + BUW + 1$ |

## 15.3 Ports

### 15.3.1 Clock and Reset

| Port    | Width | Direction | Description                           |
|---------|-------|-----------|---------------------------------------|
| pclk    | 1     | Input     | APB bus clock                         |
| presetn | 1     | Input     | APB reset<br>(active low)             |
| acclk   | 1     | Input     | Backend/user<br>clock                 |
| aresetn | 1     | Input     | Backend/user<br>reset (active<br>low) |

The backend clock and reset are named `acclk` and `aresetn`, matching the rest of the AMBA library. There are no `bclk/bresetn` ports.

### 15.3.2 APB5 Slave Interface


Same as [apb5\\_slave](#) - operates in `pclk` domain, including the optional parity signals.

### 15.3.3 Backend Interface


Same command/response interface as [apb5\\_slave](#) - operates in the `acclk` domain. `wakeup_request` is also an `acclk`-domain input.

`parity_error_wdata` and `parity_error_ctrl` are **not** `acclk`-domain signals, despite sitting next to these ports. `apb5_slave` drives them combinationally from the APB inputs (`s_apb_PSEL` && `s_apb_PENABLE ? ... : 1'b0`), so they are `pclk`-domain (gated-`pclk` in the CG variant) pulses valid only during the APB access phase. They cross no synchronizer. Sample them in `pclk`, or synchronize them yourself before using them in `acclk`.

## 15.4 Functional Description



### Architecture



### CDC Mechanism

Both directions cross through `gaxi_fifo_async` instances with a fixed 2-flop synchronizer (`N_FLOP_CROSS(2)`) on each pointer crossing. The command FIFO is written in `pclk` and read in `aclk`; the response FIFO is written in `aclk` and read in `pclk`.

The `wakeup_request` input is the one signal that does not go through a FIFO: it crosses from `aclk` into `pclk` through a `cdc_synchronizer` before reaching the wrapped `apb5_slave`. Because it is

a level, not a pulse, this is safe — but it means `wakeup_request` must be held asserted long enough to be sampled in the `pc1k` domain (at least two `pc1k` periods).

### 15.4.1 CDC FIFO Pointer Encoding

The wrapper derives  $\text{CDC\_FIFO\_DEPTH} = (\text{DEPTH} < 4) ? 4 : \text{DEPTH}$  and hands that to two `gaxi_fifo_async` instances (command and response). Gray pointers only close on a power-of-2 depth, so a FIFO depth of 6 cannot use Gray — `gaxi_fifo_async` carries an elaboration-time `$error` for exactly that case.

`USE_JOHNSON` selects the encoding:

| USE_JOHNSON | Encoding | Pointer width                                     | Depth constraint                                                                     |
|-------------|----------|---------------------------------------------------|--------------------------------------------------------------------------------------|
| 0 (default) | Gray     | $\lceil \log_2(\max(\text{DEPTH}, 4)) \rceil + 1$ | floored $\max(\text{DEPTH}, 4)$ must be a power of 2: $\text{DEPTH}$ in {2, 3, 4, 8} |
| 1           | Johnson  | $\max(\text{DEPTH}, 4)$ bits                      | any depth                                                                            |
| -1          | auto     | per depth                                         | none — Gray when the derived FIFO depth is a power of 2, Johnson otherwise           |

**The default is Gray, not auto.** With defaults,  $\text{DEPTH}$  2, 3, 4 and 8 elaborate (2, 3 and 4 all derive a depth-4 FIFO) and **DEPTH 5, 6 and 7 fail the build**. That is intentional. Johnson costs  $\max(\text{DEPTH}, 4)$ -bit pointers against Gray's  $\lceil \log_2(\max(\text{DEPTH}, 4)) \rceil + 1$  (widths follow the floored FIFO depth, so the default  $\text{DEPTH} = 2$  pays 4 bits vs 3) — at depth 6, 6 bits against 4 — duplicated in both domains and again in every synchronizer stage. Nobody should pay that because a default quietly decided for them.

If you want a non-power-of-2 depth, say so: pass `USE_JOHNSON=1` for Johnson, or `USE_JOHNSON=-1` to restore per-depth auto-selection. The capability is intact; only the default changed, so that the choice is visible in the instantiation.

**USE\_JOHNSON=0 with DEPTH=6 is an elaboration error, deliberately.** Asking for Gray at a non-power-of-2 depth is a real configuration mistake and should fail the build.

**The pointer encoding is selectable, and defaults to Gray.** `USE_JOHNSON` feeds a localparam that resolves against the `DERIVED` depth:

```

localparam int CDC_FIFO_DEPTH = (DEPTH < 4) ? 4 : DEPTH;
localparam bit CDC_DEPTH_POW2 = ((CDC_FIFO_DEPTH & (CDC_FIFO_DEPTH - 1)) == 0);
localparam int CDC_USE_JOHNSON = (USE_JOHNSON >= 0) ? USE_JOHNSON
:
(CDC_DEPTH_POW2 ? 0 : 1);

```

With the default `USE_JOHNSON = 0`, `CDC_USE_JOHNSON` is 0 and the crossing is Gray — the default `DEPTH = 2` derives 4, a power of 2, so it elaborates. `DEPTH = 6` would not; pass `USE_JOHNSON = 1` for 6-bit Johnson pointers, or `-1` to let the derived depth decide. See the CDC FIFO pointer encoding section above for the full table.

### 15.4.2 Reset Synchronization

`presetn` and `aresetn` are independent reset domains and either may be asserted alone. **A one-sided reset is not safe. Quiesce the bus first.**

The local reset clears that domain's own pointer, but the crossed copy of the *remote* pointer is a live synchronizer (`glitch_free_n_dff_arn`, `N=2`) that keeps sampling the non-reset domain the moment reset deasserts — it does not hold at zero. The reset side therefore comes back with its own pointer at zero and the remote pointer at whatever the other side had reached, which is not an empty FIFO. It is a mismatched one.

Two concrete consequences, neither of which is a clean discard:

- **Consumed commands are re-presented.** Pulsing `aresetn` with an **already-consumed** command in the cmd FIFO rewinds the read pointer behind the write pointer, so the backend sees that command again after reset. It does not time out — it re-executes. An *unread* command is not at risk: its read pointer is already behind the write pointer, so resetting it to 0 rewinds nothing and the entry is delivered exactly once.
- **The response FIFO can fabricate responses.** The same rewind on the response path presents entries the APB side never queued, so the APB master can complete a transfer that the backend did not answer.

Pointers being absolute positions rather than toggle parity is what makes the *steady-state* crossing reliable; it does not make a one-sided reset safe.

Neither reset is internally synchronized to the other domain's clock; each is expected to be already synchronized (or asynchronously asserted and synchronously deasserted) in its own domain by the integrator.

## 15.5 Timing Characteristics

---

### 15.5.1 Latency

- Crossing latency is the async-FIFO write-to-read pointer synchronization: on the order of 2-3 destination-clock cycles per direction, plus the source-clock cycle that performs the write
- A full APB transfer therefore costs the command crossing plus the response crossing before PREADY can assert
- Additional latency if the FIFOs are full (backpressure) or the backend is slow

### 15.5.2 Timing Considerations

**Timing diagram pending.** The signals and sequence this scenario exercises:

- pclk
  - aclk (different frequency)
  - APB transaction signals
  - FIFO write/read pointers
  - Backend transaction signals
  - CDC latency
- 

## 15.6 Usage Examples

---

```
apb5_slave_cdc #(
 .ADDR_WIDTH (32),
 .DATA_WIDTH (32),
 .AUSER_WIDTH (4),
 .WUSER_WIDTH (4),
 .RUSER_WIDTH (4),
 .BUSER_WIDTH (4),
 .DEPTH (2),
 .ENABLE_PARITY (0)
) u_apb5_slave_cdc (
 // APB clock domain
 .pclk (apb_clk),
 .presetn (apb_rst_n),

 // Backend clock domain
 .aclk (backend_clk),
 .aresetn (backend_rst_n),
```

```

// APB5 slave interface (pclk domain)
.s_apb_PSEL (s_apb_psel),
.s_apb_PENABLE (s_apb_penable),
// ... other APB signals

// Backend interface (aclk domain)
.cmd_valid (backend_cmd_valid),
.cmd_ready (backend_cmd_ready),
// ... other backend signals

// Wake-up request (aclk domain, synchronized internally to pclk)
.wakeup_request (backend_wakeup)
);

```

---

## 15.7 Design Notes

---

### 15.7.1 FIFO Depth Sizing

- DEPTH sizes the apb5\_slave skid buffers; the CDC FIFOs are sized separately as  $\max(\text{DEPTH}, 4)$ , so the crossing itself is never shallower than 4 entries even at the default  $\text{DEPTH}=2$
  - Deeper FIFOs buy tolerance for a slow backend or a bursty APB master; they do not reduce the per-transfer crossing latency
  - Because APB is single-outstanding, depth beyond a handful of entries has little effect on throughput for this module
- 

## 15.8 Related Modules

---

- [APB5 Slave](#) - Base slave without CDC
  - [APB5 Slave CDC CG](#) - CDC with clock gating
  - [GAXI Async FIFO](#) - The async FIFO used for both crossings
  - [Clock Domain Crossing](#) - CDC design patterns and reset behavior
- 

## 15.9 Testing

---

`val/amba/test_apb5_slave_cdc.py` exercises this module. It collects 3 parameter cases at the default `REG_LEVEL`.

```
source env_python
pytest val/amba/test_apb5_slave_cdc.py -v
```

---

## 15.10 Navigation

---

- [← Back to APB5 Index](#)
- [← Back to rtl-amba Index](#)
- [← Back to Main Documentation Index](#)

## 16 APB5 Slave CDC (Clock-Gated)

**Module:** apb5\_slave\_cdc\_cg.sv **Location:** rtl/amba/apb5/ **Status:** Production Ready

---

### 16.1 Overview

---

The APB5 Slave CDC + Clock-Gated module combines clock domain crossing with clock gating for maximum power efficiency. It wraps apb5\_slave\_cdc with clock gating control to cut power during idle periods while keeping operation safe across asynchronous clock boundaries.

#### 16.1.1 Key Features

- Full APB5 protocol support with all extensions
  - Asynchronous clock domain crossing (APB to backend)
  - Clock gating for power reduction during idle
  - All APB5 user signals (PAUSER, PWUSER, PRUSER, PBUSER)
  - PWAKEUP signal handling across domains
  - Optional parity support for data integrity
  - Automatic wake-up on transaction activity
- 

### 16.2 Parameters

---

| Parameter   | Type | Default      | Description                     |
|-------------|------|--------------|---------------------------------|
| ADDR_WIDTH  | int  | 32           | APB address bus width           |
| DATA_WIDTH  | int  | 32           | APB data bus width              |
| PROT_WIDTH  | int  | 3            | Protection signal width         |
| AUSER_WIDTH | int  | 4            | Address user signal width       |
| WUSER_WIDTH | int  | 4            | Write user signal width         |
| RUSER_WIDTH | int  | 4            | Read user signal width          |
| BUSER_WIDTH | int  | 4            | Response user signal width      |
| STRB_WIDTH  | int  | DATA_WIDTH/8 | Write strobe width (calculated) |
| DEPTH       | int  | 2            | Skid-buffer depth of the        |

| Parameter           | Type | Default  | Description                                                                                                                                                                                                                             |
|---------------------|------|----------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| USE_JOHNSON         | int  | 0 (Gray) | wrapped slave; CDC FIFOs use $\max(\text{DEPTH}, 4)$<br>Forwarded to apb5_slave_cdc: 0 Gray, 1 Johnson, -1 auto.<br>Newly reachable from this level — it used to be unreachable, so a CG instance could not select the encoding at all. |
| ENABLE_PARITY       | bit  | 0        | Enable parity generation and checking                                                                                                                                                                                                   |
| CG_IDLE_COUNT_WIDTH | int  | 4        | Width of idle counter (max idle = $2^N - 1$ cycles)                                                                                                                                                                                     |
| USE_2_PHASE_CDC     | bit  | 1        | Deprecated and ignored — retained for source compatibility                                                                                                                                                                              |

As with [apb5\\_slave\\_cdc](#), there is no SYNC\_STAGES parameter: pointer synchronization is fixed at 2 flops inside the async FIFOs.

### 16.2.1 Derived Parameters (do not override)

These are declared as parameter so the elaborator can compute them, not so callers can set them. Each defaults to an expression over the parameters above; overriding one desynchronises it from its source and the design fails to elaborate or silently mis-sizes a bus. Set the parameters they are derived FROM and leave these alone.

| Derived parameter | Default expression  |
|-------------------|---------------------|
| DW                | DATA_WIDTH          |
| AW                | ADDR_WIDTH          |
| SW                | STRB_WIDTH          |
| PW                | PROT_WIDTH          |
| AUW               | AUSER_WIDTH         |
| WUW               | WUSER_WIDTH         |
| RUW               | RUSER_WIDTH         |
| BUW               | BUSER_WIDTH         |
| ICW               | CG_IDLE_COUNT_WIDTH |

| Derived parameter | Default expression                  |
|-------------------|-------------------------------------|
| CPW               | $AW + DW + SW + PW + AUW + WUW + 1$ |
| RPW               | $DW + RUW + BUW + 1$                |

## 16.3 Ports

### 16.3.1 Clock and Reset

| Port    | Width | Direction | Description                |
|---------|-------|-----------|----------------------------|
| pclk    | 1     | Input     | APB bus clock              |
| presetn | 1     | Input     | APB reset<br>(active low)  |
| acclk   | 1     | Input     | User/backend<br>clock      |
| aresetn | 1     | Input     | User reset<br>(active low) |

### 16.3.2 Clock Gating Configuration

| Port              | Width               | Direction | Description               |
|-------------------|---------------------|-----------|---------------------------|
| cfg_cg_enable     | 1                   | Input     | Enable clock gating       |
| cfg_cg_idle_count | CG_IDLE_COUNT_WIDTH | Input     | Idle cycles before gating |

### 16.3.3 APB5 Slave Interface

Same as [apb5\\_slave](#) - operates in pclk domain.

### 16.3.4 Backend Interface


Same command/response interface as [apb5\\_slave\\_cdc](#) - operates in acclk domain.

### 16.3.5 Status Outputs


| Port      | Width | Direction | Description                           |
|-----------|-------|-----------|---------------------------------------|
| cg_gating | 1     | Output    | Indicates clock is<br>currently gated |
| cg_idle   | 1     | Output    | Activity terms quiet                  |

| Port               | Width | Direction | Description                                               |
|--------------------|-------|-----------|-----------------------------------------------------------|
| parity_error_wdata | 1     | Output    | (registered ~wakeup).<br>Write data parity error detected |
| parity_error_ctrl  | 1     | Output    | Control signal parity error                               |

## 16.4 Functional Description



### Architecture



### Wake-up Logic

The aclk-domain activity terms are ORed together and passed through a two-flop synchronizer before being combined with the pclk-domain terms. Without that synchronizer, these signals would cross domains unsynchronized.

## 16.5 Timing Characteristics

**Timing diagram pending.** The signals and sequence this scenario exercises:

- pclk (ungated)
- gated\_pclk
- aclk (different frequency)
- s\_apb\_PSEL (wake trigger)
- cg\_gating indicator
- Transaction flow across CDC with gating
- Wake-up latency from PSEL to clock active

---

## 16.6 Usage Examples

---

```
apb5_slave_cdc_cg #(
 .ADDR_WIDTH (32),
 .DATA_WIDTH (32),
 .AUSER_WIDTH (4),
 .WUSER_WIDTH (4),
 .RUSER_WIDTH (4),
 .BUSER_WIDTH (4),
 .DEPTH (2),
 .ENABLE_PARITY (0),
 .CG_IDLE_COUNT_WIDTH(4)
) u_apb5_slave_cdc_cg (
 // APB clock domain
 .pclk (apb_clk),
 .presetn (apb_rst_n),

 // User clock domain
 .aclk (user_clk),
 .aresetn (user_rst_n),

 // Clock gating
 .cfg_cg_enable (1'b1),
 .cfg_cg_idle_count (4'd8),
 .cg_gating (slave_clk_gated),

 // APB5 slave interface (pclk domain)
 .s_apb_PSEL (s_apb_psel),
 .s_apb_PENABLE (s_apb_penable),
 .s_apb_PREADY (s_apb_pready),
 // ... other APB signals

 // Backend interface (aclk domain)
 .cmd_valid (backend_cmd_valid),
 .cmd_ready (backend_cmd_ready),
 // ... other command signals

 .rsp_valid (backend_rsp_valid),
 .rsp_ready (backend_rsp_ready),
 // ... other response signals

 // Wake-up control (aclk domain)
 .wakeup_request (backend_wakeup)
);
```

---

## 16.7 Design Notes

---

### 16.7.1 Power and Latency Trade-offs

Wake-up is **registered, not combinational**. Activity is registered once (AXI4, AXI5, AXI4-Lite, AXI4-Stream, APB) or twice (APB5, AXI5-Stream) before reaching the ICG enable, which is combinational. A pclk-domain trigger (PSEL or PENABLE) passes through this wrapper's r\_wakeup and then through the flop inside amba\_clock\_gate\_ctrl, so APB5 is a two-stage family and the first usable gated pclk rising edge arrives **3 ungated pclk cycles** after the trigger. An aclk-domain trigger (cmd\_valid, rsp\_valid, wakeup\_request) first crosses a 2-flop synchronizer into pclk, adding about two more pclk cycles.


| Configuration           | Power Savings | Wake-up Latency from PSEL                                  |
|-------------------------|---------------|------------------------------------------------------------|
| CG disabled             | None          | 0 cycles (clock always running)                            |
| CG idle=4               | Moderate      | 2 register stages; first usable edge 3 ungated pclk cycles |
| CG idle=16 <sup>1</sup> | Good          | 2 register stages; first usable edge 3 ungated pclk cycles |

The wake-up latency does not depend on cfg\_cg\_idle\_count; the idle count only controls how long the block waits before gating. Latency is absorbed as APB wait states, since the slave holds PREADY low until it is running again.

Deliberately keeping the wake-up path registered avoids driving a clock-gate enable from a combinational function of a bus input. The gating element itself is an ICG cell instantiated by clock\_gate\_ctrl, which is glitch-free by construction. ICG cells are an ASIC library primitive; FPGA targets should use a clock-enable approach instead.

---

<sup>1</sup> idle=16 requires CG\_IDLE\_COUNT\_WIDTH >= 5. At the default of 4 the counter holds 0-15, so writing 16 truncates to 0 and gates after the first idle cycle – the opposite of the intent. Raise the parameter for this row.



### *Combined Feature Hierarchy*

#### **16.7.2 Reset Considerations**

- APB domain reset (presetn) controls the APB interface, the wake-up synchronizer and the clock gate
  - User domain reset (aresetn) controls the backend side of the CDC FIFOs
  - Each reset must already be synchronized (or asynchronously asserted and synchronously deasserted) in its own domain by the integrator; this module does not synchronize either reset into the other domain
  - A one-sided reset is NOT safe: the crossed pointer copy is a live synchronizer, so the reset side returns with its own pointer at zero against a remote pointer that kept advancing. Quiesce the bus first — see [apb5\\_slave\\_cdc](#) for the mechanism and for the in-flight transaction caveat
-

## 16.8 Related Modules

---

- [APB5 Slave](#) - Base slave module
  - [APB5 Slave CDC](#) - CDC variant without clock gating
  - [APB5 Slave CG](#) - Clock gating without CDC
- 

## 16.9 Testing

---

`val/amba/test_apb5_slave_cdc_cg.py` exercises this module. It collects 3 parameter cases at the default `REG_LEVEL`.

```
source env_python
pytest val/amba/test_apb5_slave_cdc_cg.py -v
```

---

## 16.10 Navigation

---

- [← Back to APB5 Index](#)
- [← Back to rtl-amba Index](#)
- [← Back to Main Documentation Index](#)