




RTL Design Sherpa

Converters Micro-Architecture Specification 1.4

September 11, 2026

Table of Contents

1 Converter Mas Index.....	22
2 Document Information.....	23
2.1 Document Control.....	23
2.2 Purpose.....	24
2.3 Audience.....	25
2.4 Related Documents.....	26
2.5 Revision History.....	27
2.6 Conventions.....	28
2.6.1 Notation.....	28
2.6.2 Diagrams.....	28
2.6.3 Code Examples.....	28
3 Chapter 1: Introduction - Overview.....	29
3.1 1.1 Purpose.....	29
3.2 1.2 Problem Statement.....	30
3.2.1 1.2.1 Data Width Mismatch.....	30
3.2.2 1.2.2 Protocol Incompatibility.....	30
3.3 1.3 Solution Architecture.....	31
3.3.1 Figure 1.1: Converter Module Hierarchy.....	31
3.3.2 1.3.1 Data Width Converters.....	31
3.3.3 1.3.2 Protocol Converters.....	31
3.4 1.4 Key Design Decisions.....	32
3.4.1 1.4.1 Generic vs. Full Modules.....	32

3.4.2 1.4.2 Throughput vs. Area Trade-offs.....	32
3.4.3 1.4.3 Sideband Signal Handling.....	32
3.5 1.5 Performance Characteristics.....	34
3.5.1 1.5.1 Latency.....	34
3.5.2 1.5.2 Throughput.....	34
3.6 1.6 Scope.....	35
3.6.1 In Scope.....	35
3.6.2 Out of Scope.....	35
3.7 1.7 Target Applications.....	36
4 2.1 Generic Building Blocks.....	37
4.1 2.1.1 Module Hierarchy.....	38
4.2 2.1.2 Design Philosophy.....	39
4.2.1 Separation of Concerns.....	39
4.2.2 Interface Pattern.....	39
4.3 2.1.3 Width Ratio Calculation.....	40
4.4 2.1.4 Sideband Modes.....	41
4.4.1 Upsize Sideband Modes.....	41
4.4.2 Downsize Sideband Modes.....	41
4.5 2.1.5 Performance Comparison.....	43
4.5.1 Figure 2.1: Performance Comparison.....	43
4.6 2.1.6 Resource Utilization.....	44
4.6.1 Upsize Resources (NARROW=64, WIDE=512).....	44
4.6.2 Downsize Resources (WIDE=512, NARROW=64).....	44
4.7 2.1.7 Integration Guidelines.....	45

4.7.1 When to Use Generic Blocks.....	45
4.7.2 Composition Example.....	45
5 2.2 axi_data_upsize.....	46
5.1 2.2.1 Overview.....	47
5.2 2.2.2 Block Diagram.....	48
5.2.1 Figure 2.2: axi_data_upsize Architecture.....	48
5.3 2.2.3 Interface Specification.....	49
5.3.1 Parameters.....	49
5.3.2 Ports.....	49
5.4 2.2.4 Operation.....	51
5.4.1 Accumulation Cycle.....	51
5.4.2 Early LAST Handling.....	51
5.4.3 State Machine.....	51
5.5 2.2.5 Sideband Handling.....	52
5.5.1 Concatenate Mode (SB_OR_MODE=0).....	52
5.5.2 Severity-Fold Mode (SB_OR_MODE=1).....	52
5.6 2.2.6 Implementation.....	54
5.6.1 Core Logic.....	54
5.7 2.2.7 Timing.....	56
5.7.1 Latency.....	56
5.7.2 Throughput.....	56
5.7.3 Critical Paths.....	56
5.8 2.2.8 Resource Utilization.....	57
5.8.1 Typical Resources (64-bit to 512-bit).....	57

5.8.2 Scaling.....	57
5.9 2.2.9 Usage Example.....	58
6 2.3 axi_data_dnsiz.....	59
6.1 2.3.1 Overview.....	60
6.2 2.3.2 Block Diagram.....	61
6.2.1 Figure 2.3: axi_data_dnsiz Single-Buffer Architecture.....	61
6.3 2.3.3 Interface Specification.....	62
6.3.1 Parameters.....	62
6.3.2 Ports.....	62
6.4 2.3.4 Single-Buffer Mode Operation.....	64
6.4.1 Split Cycle.....	64
6.4.2 State Machine.....	64
6.4.3 Throughput Analysis.....	64
6.5 2.3.5 Sideband Handling.....	66
6.5.1 Slice Mode (SB_BROADCAST=0).....	66
6.5.2 Broadcast Mode (SB_BROADCAST=1).....	66
6.6 2.3.6 Burst Tracking.....	67
6.6.1 Purpose.....	67
6.6.2 Operation.....	67
6.7 2.3.7 Implementation.....	68
6.7.1 Single-Buffer Core Logic.....	68
6.8 2.3.8 Resource Utilization.....	69
6.8.1 Single-Buffer (512-bit to 64-bit).....	69
6.8.2 Measured Throughput.....	69

6.9 2.3.9 Usage Example.....	70
7 2.5 axi4_dwidth_converter_wr.....	71
7.1 2.5.1 Overview.....	72
7.2 2.5.2 Block Diagram.....	73
7.2.1 Figure 2.6: Write Converter Architecture.....	73
7.3 2.5.3 Interface Specification.....	74
7.3.1 Parameters.....	74
7.3.2 Ports.....	74
7.4 2.5.4 Burst Length Conversion.....	78
7.4.1 Ratio Calculation.....	78
7.4.2 Examples.....	78
7.4.3 Non-Aligned Bursts.....	78
7.5 2.5.5 Address Channel Handling.....	79
7.5.1 AW Passthrough with Adjustment.....	79
7.5.2 Skid Buffer for AW.....	80
7.6 2.5.6 Write Data Channel.....	81
7.6.1 Upsize Instance.....	81
7.7 2.5.7 Response Channel.....	82
7.7.1 B Channel.....	82
7.7.2 Response Ordering.....	82
7.7.3 Ordering Constraint (applies to both directions).....	82
7.8 2.5.8 AW/W Synchronization.....	84
7.8.1 Challenge.....	84
7.8.2 Solution.....	84

7.9 2.5.9 Resource Utilization.....	85
7.9.1 Typical Resources (64→512, ID=4, ADDR=64).....	85
7.10 2.5.10 Timing.....	86
7.10.1 Latency.....	86
7.10.2 Throughput.....	86
7.11 2.5.11 Usage Example.....	87
8 2.6 axi4_dwidth_converter_rd.....	88
8.1 2.6.1 Overview.....	89
8.2 2.6.2 Block Diagram.....	90
8.2.1 Figure 2.7: Read Converter Architecture.....	90
8.3 2.6.3 Interface Specification.....	91
8.3.1 Parameters.....	91
8.3.2 Ports.....	91
8.4 2.6.4 Burst Length Conversion.....	94
8.4.1 Ratio Calculation.....	94
8.4.2 Examples.....	94
8.5 2.6.5 Address Channel Handling.....	95
8.5.1 AR Passthrough with Adjustment.....	95
8.5.2 Burst-Length FIFO.....	96
8.6 2.6.6 Read Data Channel.....	97
8.6.1 Downsize Instance.....	97
8.7 2.6.7 RLAST Generation.....	98
8.8 2.6.8 RID Handling.....	99
8.8.1 ID Passthrough.....	99

8.8.2 Ordering Constraint (applies to both directions).....	99
8.9 2.6.9 Resource Utilization.....	100
8.9.1 Typical Resources (64→512 UPSIZE, ratio 8, ID=4).....	100
8.10 2.6.10 Timing.....	101
8.10.1 Latency.....	101
8.10.2 Throughput.....	101
8.11 2.6.11 Usage Example.....	102
9 2.7 AXIL Wide-Alignment Converters.....	103
9.1 2.7.1 Why They Exist.....	104
9.2 2.7.2 Parameters.....	105
10 3.1 Protocol Conversion Overview.....	106
10.1 3.1.1 Available Converters.....	107
10.1.1 AXI4 to AXI4-Lite (Protocol Downgrade).....	107
10.1.2 AXI4-Lite to AXI4 (Protocol Upgrade).....	107
10.1.3 Other Protocol Converters.....	107
10.2 3.1.2 Protocol Comparison.....	109
10.2.1 Figure 3.1: Protocol Feature Comparison.....	109
10.2.2 Feature Matrix.....	109
10.3 3.1.3 Conversion Strategies.....	110
10.3.1 AXI4 to AXI4-Lite.....	110
10.3.2 AXI4-Lite to AXI4.....	110
10.3.3 AXI4 to APB.....	110
10.4 3.1.4 Performance Characteristics.....	111
10.4.1 Key Observations.....	111

10.5 3.1.5 Use Case Guidelines.....	112
10.5.1 When to Use AXI4 to AXI4-Lite.....	112
10.5.2 When to Use AXI4-Lite to AXI4.....	112
10.5.3 When to Use AXI4 to APB.....	112
10.6 3.1.6 Integration Patterns.....	113
10.6.1 Pattern 1: CPU to Peripherals.....	113
10.6.2 Pattern 2: Simple IP in AXI4 Fabric.....	113
10.6.3 Pattern 3: Mixed System.....	113
11 3.2 AXI4 to AXI4-Lite Converter.....	114
11.1 3.2.1 Module Organization.....	115
11.1.1 Design Philosophy.....	115
11.2 3.2.2 Read Path (axi4_to_axil4_rd).....	116
11.2.1 Figure 3.2: AXI4 to AXI4-Lite Read Path.....	116
11.2.2 Operation.....	116
11.2.3 State Machine.....	116
11.2.4 Implementation.....	117
11.3 3.2.3 Write Path (axi4_to_axil4_wr).....	119
11.3.1 Figure 3.3: AXI4 to AXI4-Lite Write Path.....	119
11.3.2 Operation.....	119
11.3.3 AW/W Synchronization Challenge.....	120
11.3.4 Solution: W gated on the AW that owns it.....	120
11.3.5 Response Aggregation.....	121
11.4 3.2.4 Bidirectional Wrapper (axi4_to_axil4).....	122
11.4.1 Composition Pattern.....	122

11.5 3.2.5 Resource Utilization.....	123
11.6 3.2.6 Timing.....	124
11.6.1 Throughput.....	124
11.6.2 Latency.....	124
11.7 3.2.7 Testing.....	125
12 3.3 AXI4-Lite to AXI4 Converter.....	126
12.1 3.3.1 Module Organization.....	127
12.2 3.3.2 Design Philosophy.....	128
12.3 3.3.3 Signal Mapping.....	129
12.3.1 Address Channel Signals.....	129
12.3.2 Data Channel Signals.....	129
12.4 3.3.4 Read Path (axil4_to_axi4_rd).....	130
12.4.1 Figure 3.4: AXI4-Lite to AXI4 Read Path.....	130
12.4.2 Implementation.....	130
12.5 3.3.5 Write Path (axil4_to_axi4_wr).....	133
12.5.1 Figure 3.5: AXI4-Lite to AXI4 Write Path.....	133
12.5.2 Implementation.....	133
12.6 3.3.6 Bidirectional Wrapper.....	136
12.7 3.3.7 Resource Utilization.....	137
12.8 3.3.8 Timing.....	138
12.9 3.3.9 Testing.....	139
12.10 3.3.10 Usage Example.....	140
13 3.4 AXI4 to APB Converter.....	141
13.1 3.4.1 Overview.....	142

13.2 3.4.2 Block Diagram.....	143
13.2.1 Figure 3.6: AXI4 to APB Converter.....	143
13.3 3.4.3 Interface Specification.....	144
13.3.1 Parameters.....	144
13.3.2 Ports.....	145
13.4 3.4.4 Clock Domains.....	148
13.5 3.4.5 State Machine.....	149
13.5.1 Figure 3.7: AXI4 to APB FSM.....	149
13.5.2 States.....	150
13.6 3.4.6 Burst Handling.....	151
13.6.1 Burst Decomposition.....	151
13.6.2 Address Calculation.....	151
13.7 3.4.7 Address Width Adaptation.....	152
13.7.1 64-bit to 32-bit Conversion.....	152
13.8 3.4.8 Error Response Mapping.....	153
13.8.1 Error Aggregation.....	153
13.9 3.4.9 Implementation.....	154
13.10 3.4.10 Resource Utilization.....	155
13.11 3.4.11 Timing.....	156
13.11.1 Timing Analysis.....	156
13.11.2 Throughput.....	156
13.12 3.4.12 Usage Example.....	157
14 AXI4 to APB5 Shim.....	159
14.1 Overview.....	160

14.2 Parameters.....	161
14.3 Design Notes.....	162
14.4 Related Modules.....	163
14.5 Navigation.....	164
15 AXI4 to AXI5-Lite.....	165
15.1 Overview.....	166
15.2 Where each signal gets its value.....	167
15.3 The AW/AR sideband is held, not passed through.....	168
15.4 Reset style.....	169
15.5 Verification.....	170
15.6 Related Modules.....	171
15.7 Navigation.....	172
16 3.5 PeakRDL Adapter.....	173
16.1 Overview.....	174
16.2 Parameters.....	175
16.3 Ports.....	176
16.4 Functional Description.....	178
16.4.1 Write.....	178
16.4.2 Read.....	178
16.4.3 3.5.5 Implementation.....	178
16.4.4 3.5.6 Resource Utilization.....	179
16.5 Waveforms.....	180
16.5.1 Figure 3.8: PeakRDL Adapter.....	180
16.6 Usage Example.....	182

16.6.1 3.5.7 Use Cases.....	182
16.6.2 3.5.8 Integration Example.....	182
16.7 Related Modules.....	184
16.7.1 3.5.9 APB4 Front End (apb4_to_peakrdl).....	184
17 3.7 UART to AXI4-Lite Bridge.....	187
17.1 3.7.1 Module Organization.....	188
17.2 3.7.2 Command Protocol.....	189
17.3 3.7.3 Error Responses.....	190
17.4 3.7.4 Parameters.....	191
17.5 3.7.5 Ports.....	192
17.5.1 uart_axil_bridge.....	192
17.5.2 uart_rx.....	193
17.5.3 uart_tx.....	193
18 3.8 Width-plus-Protocol Chains.....	195
18.1 3.8.1 axi4_dwidth_to_axil4_wr_chain.....	196
18.2 3.8.2 axi4_dwidth_to_apb4_chain.....	197
19 3.10 AXI4-Lite to Wishbone B4 Converter.....	199
19.1 3.10.1 Module Organization.....	200
19.2 3.10.2 Design Philosophy.....	201
19.3 3.10.3 Block Diagram.....	202
19.3.1 Figure 3.12: AXI4-Lite to Wishbone B4 Converter.....	202
19.4 3.10.4 Interface Specification.....	203
19.4.1 Parameters.....	203
19.4.2 Ports.....	203

19.5 3.10.5 Command Formation.....	205
19.6 3.10.6 Response Mapping.....	206
19.7 3.10.7 Timing.....	207
19.8 3.10.8 Formal.....	208
19.9 3.10.9 Testing.....	209
19.10 3.10.10 Usage Example.....	210
20 3.11 Wishbone B4 to AXI4-Lite Converter.....	211
20.1 3.11.1 Module Organization.....	212
20.2 3.11.2 Why This Direction Is Harder.....	213
20.3 3.11.3 Interface Specification.....	214
20.3.1 Parameters.....	214
20.3.2 Ports.....	214
20.4 3.11.4 Command Formation.....	216
20.5 3.11.5 Response Mapping.....	217
20.6 3.11.6 Formal.....	218
20.7 3.11.7 Testing.....	219
20.8 3.11.8 Usage Example.....	220
20.9 Navigation.....	221
21 APB to AXI4 (Requester).....	222
21.1 Overview.....	223
21.2 Fields the requester invents.....	224
21.3 APB5 sideband.....	225
21.4 Parameters.....	226
21.5 Verification.....	227

21.6 Related Modules.....	228
21.7 Navigation.....	229
22 AXI4 to Wishbone B4.....	230
22.1 Overview.....	231
22.2 Parameters.....	232
22.3 Verification.....	233
22.4 Related Modules.....	234
22.5 Navigation.....	235
23 Wishbone B4 to AXI4.....	236
23.1 Overview.....	237
23.2 Parameters.....	238
23.3 Verification.....	239
23.4 Related Modules.....	240
23.5 Navigation.....	241
24 4.1 Width Converter FSMs.....	242
24.1 Overview.....	242
24.2 Functional Description.....	243
24.2.1 4.1.1 Upsize Structure (no FSM).....	243
24.2.2 4.1.2 Downsize Structure (single buffer, no FSM).....	243
24.2.3 4.1.3 Full Converter FSMs.....	244
24.3 Timing.....	245
24.3.1 Upsize (8:1 ratio, m_ready held high).....	245
24.3.2 Downsize (8:1 ratio, single buffer, atomic replace).....	245
25 4.2 Protocol Converter FSMs.....	246

25.1 Overview.....	246
25.2 Functional Description.....	247
25.2.1 4.2.1 AXI4 to AXI4-Lite Read FSM.....	247
25.2.2 4.2.2 AXI4 to AXI4-Lite Write FSM.....	247
25.2.3 4.2.3 AXI4 to APB FSMs.....	247
25.3 Timing.....	248
25.3.1 4.2.4 Timing Analysis.....	248
26 4.3 Burst Decomposition.....	249
26.1 Overview.....	249
26.2 Functional Description.....	250
26.2.1 4.3.1 Width Converter Burst Handling.....	250
26.2.2 4.3.2 Protocol Converter Burst Handling.....	250
26.2.3 4.3.3 Response Aggregation.....	251
26.2.4 4.3.4 Burst Tracking Registers.....	252
26.3 Timing.....	254
26.3.1 4.3.5 Timing Impact.....	254
26.4 Waveforms.....	255
26.4.1 Figure 4.5: Width Burst Conversion.....	255
26.5 Navigation.....	256
27 5.1 Test Strategy.....	257
27.1 Overview.....	257
27.2 Testing.....	258
27.2.1 5.1.1 Test Organization.....	258
27.2.2 5.1.2 Test Levels.....	258

27.2.3 5.1.3 Test Categories.....	259
27.2.4 5.1.4 Testbench Architecture.....	260
27.2.5 5.1.5 Coverage Model.....	261
27.2.6 5.1.6 Test Execution.....	262
27.3 Navigation.....	263
28 5.2 Debug Guide.....	264
28.1 Overview.....	264
28.2 Testing.....	265
28.2.1 5.2.1 Common Issues.....	265
28.2.2 5.2.2 Debug Signals.....	266
28.2.3 5.2.3 Simulation Debug.....	267
28.2.4 5.2.4 Common Mistakes.....	268
28.2.5 5.2.5 Verification Checklist.....	269
28.2.6 5.2.6 Performance Validation.....	269
28.3 Navigation.....	271

List of Figures

Figure 1.1: Converter Module Hierarchy.....	31
Figure 2.1: Performance Comparison.....	43
Figure 2.2: axi_data_upsize Architecture.....	48
Figure 2.3: axi_data_dnsiz Single-Buffer Architecture.....	61
Figure 2.6: Write Converter Architecture.....	73
Figure 2.7: Read Converter Architecture.....	90
Figure 3.1: Protocol Feature Comparison.....	109
Figure 3.2: AXI4 to AXI4-Lite Read Path.....	116
Figure 3.3: AXI4 to AXI4-Lite Write Path.....	119
Figure 3.4: AXI4-Lite to AXI4 Read Path.....	130
Figure 3.5: AXI4-Lite to AXI4 Write Path.....	133
Figure 3.6: AXI4 to APB Converter.....	143
Figure 3.7: AXI4 to APB FSM.....	149
Figure 3.8: PeakRDL Adapter.....	180
Figure 3.12: AXI4-Lite to Wishbone B4 Converter.....	202
Figure 4.5: Width Burst Conversion.....	255

List of Tables

Table 1: Table 0.1: Document Control Information.....	23
Table 2: Table 0.2: Related Documents.....	26
Table 3: Table 0.3: Revision History.....	27
Table 4: Table 1.1: Common Data Width Configurations.....	30
Table 5: Table 1.2: Protocol Mismatch Examples.....	30
Table 6: Table 1.3: Throughput vs. Area Trade-offs.....	32
Table 7: Table 1.4: Sideband Handling Modes.....	32
Table 8: Table 1.5: Latency Characteristics.....	34
Table 9: Table 1.6: Throughput Characteristics.....	34
Table 10: Table 2.1: Upsize Sideband Modes.....	41
Table 11: Table 2.2: Downsize Sideband Modes.....	41
Table 12: Table 2.3: Performance Comparison.....	43
Table 13: Table 2.4: axi_data_upsize Parameters.....	49
Table 14: Table 2.5: Upsize Latency.....	56
Table 15: Table 2.6: Upsize Resource Scaling.....	57
Table 16: Table 2.7: axi_data_dnsiz Parameters.....	62
Table 17: Table 2.8: Measured Throughput by Mode.....	65
Table 18: Table 2.9: Resources and measured rate.....	69
Table 19: Table 2.14: Write Converter Parameters.....	74
Table 20: Table 2.15: Burst Length Conversion Examples.....	78
Table 21: Table 2.16: Write Converter Latency.....	86
Table 22: Table 2.17: Read Converter Parameters.....	91
Table 23: Table 2.18: Read Burst Length Conversion.....	94
Table 24: Table 2.20: Read Converter Latency.....	101
Table 25: Wide-Alignment Converter Parameters.....	105
Table 26: Table 3.1: AXI4 to AXI4-Lite Converters.....	107
Table 27: Table 3.2: AXI4-Lite to AXI4 Converters.....	107
Table 28: Table 3.3: Other Protocol Converters.....	107
Table 29: Table 3.4: Protocol Feature Comparison.....	109
Table 30: Table 3.5: Protocol Converter Performance.....	111
Table 31: Table 3.6: AXI4 to AXI4-Lite Resources.....	123
Table 32: Table 3.7: AXI4 to AXI4-Lite Throughput.....	124
Table 33: Table 3.8: AXI4 to AXI4-Lite Latency.....	124
Table 34: Table 3.9: Test Coverage Summary.....	125
Table 35: Table 3.10: AR Channel Mapping.....	129
Table 36: Table 3.11: R Channel Mapping.....	129

Table 37: Table 3.12: AXIL4 to AXI4 Resources.....	137
Table 38: Table 3.13: AXIL4 to AXI4 Performance.....	138
Table 39: Table 3.14: Test Coverage Summary.....	139
Table 40: Table 3.15: AXI4 vs APB Comparison.....	142
Table 41: Table 3.16: AXI4 to APB Parameters.....	144
Table 42: Table 3.17: AXI4 to APB Shim Parameters.....	144
Table 43: Table 3.18: Error Mapping.....	153
Table 44: Table 3.19: APB Converter Timing.....	156
Table 45: AXI4 to APB5 Shim – Inert Parameter.....	161
Table 46: AXI4 to AXI5-Lite – Sideband Disposition.....	167
Table 47: Table 3.20: PeakRDL Adapter Parameters.....	175
Table 48: apb4_to_peakrdl Parameters.....	184
Table 49: apb4_to_peakrdl Ports.....	185
Table 50: uart_axil_bridge Parameters.....	191
Table 51: uart_axil_bridge Ports.....	192
Table 52: uart_rx Ports.....	193
Table 53: uart_tx Ports.....	193
Table 54: axi4_dwidth_to_axil4_wr_chain Parameters.....	196
Table 55: axi4_dwidth_to_apb4_chain Parameters.....	197
Table 56: Table 3.26: AXI4-Lite to Wishbone Parameters.....	203
Table 57: Table 3.27: AXI4-Lite to Wishbone Ports.....	203
Table 58: Table 3.28: Command Formation.....	205
Table 59: Table 3.29: Response Mapping.....	206
Table 60: Table 3.30: AXI4-Lite to Wishbone Timing.....	207
Table 61: Table 3.31: Wishbone to AXI4-Lite Parameters.....	214
Table 62: Table 3.32: Wishbone to AXI4-Lite Ports.....	214
Table 63: Table 3.33: Command Formation.....	216
Table 64: Table 3.34: Response Mapping.....	217
Table 65: Table 3.35: APB-to-AXI4 requester pieces.....	223
Table 66: Table 3.36: AXI4 request fields from an APB transfer.....	224
Table 67: Table 3.37: APB5 USER mapping.....	225
Table 68: Table 3.38: APB-to-AXI4 parameters.....	226
Table 69: Table 3.39: axi4_to_wb4 stages.....	231
Table 70: Table 3.40: axi4_to_wb4 parameters.....	232
Table 71: Table 3.41: AXI4 fields the wrapper invents.....	237
Table 72: Table 4.6: Decomposition Overhead.....	254
Table 73: Table 5.1: Functional Test Categories.....	259
Table 74: Table 5.2: Stress Test Categories.....	259

Table 75: Table 5.3: Error Test Categories.....	259
Table 76: Table 5.4: Coverage Targets.....	261
Table 77: Table 5.5: Expected Throughput.....	269

1 Converter Mas Index

Generated: 2026-09-11

RTL Design Sherpa · Learning Hardware Design Through Practice [GitHub](#) · [Documentation Index](#) · [MIT License](#)

2 Document Information

2.1 Document Control

Table 0.1: Document Control Information

Field	Value
Document Title	Converters Micro-Architecture Specification
Document Version	1.0
Component	Converters
Status	Active
Classification	Internal Technical
Last Updated	2026-01-03

2.2 Purpose

This Micro-Architecture Specification (MAS) describes how the Converters component is actually built — the internal detail behind the feature set. It covers:

- Internal block architectures
- State machine designs
- Signal timing and handshaking
- Resource utilization estimates
- Debug and verification strategies

2.3 Audience

This document is for:

- RTL designers implementing or modifying converter modules
- Verification engineers creating testbenches
- Integration engineers connecting converters to systems
- Performance engineers optimizing throughput and latency

2.4 Related Documents

Table 0.2: Related Documents

Document	Purpose
Bridge MAS	Related crossbar micro-architecture
Stream MAS	Related datapath micro-architecture

(The old converter spec tree was migrated into this MAS, and converters has no standalone PRD — this book is the authoritative document.)

2.5 Revision History

Table 0.3: Revision History

Version	Date	Author	Description
1.0	2026-01-03	RTL Design Sherpa	Initial MAS release
1.2	2026-09-09	RTL Design Sherpa	Added chapter 3.10, AXI4-Lite to Wishbone B4

2.6 Conventions

2.6.1 Notation

- `signal_name` - RTL signals and parameters
- **ModuleName** - Module names and key concepts
- *Figure X.X* - Figure references

2.6.2 Diagrams

All diagrams use Mermaid format rendered to PNG: - Source: `assets/mermaid/*.mmd` -
Rendered: `assets/mermaid/*.png`

2.6.3 Code Examples

Treat the SystemVerilog snippets as implementation guidance: they show the intended design pattern but may differ slightly from the actual RTL.

Next: [Chapter 1: Introduction](#)

RTL Design Sherpa · Learning Hardware Design Through Practice [GitHub](#) · [Documentation Index](#) ·
[MIT License](#)

3 Chapter 1: Introduction - Overview

3.1 1.1 Purpose

The Converters component provides data width and protocol conversion — the modules that let blocks with different data widths or communication protocols talk to each other.

3.2 1.2 Problem Statement

Sooner or later, every SoC design runs into the same two integration problems:

3.2.1 1.2.1 Data Width Mismatch

Table 1.1: Common Data Width Configurations

Component	Typical Data Width
CPU	64-bit
DDR Controller	512-bit
PCIe Endpoint	128-bit
GPU	256-bit

Challenge: You can't wire mismatched widths together directly — a width converter has to sit between them.

3.2.2 1.2.2 Protocol Incompatibility


Table 1.2: Protocol Mismatch Examples

Master Type	Protocol	Slave Type	Protocol
CPU	AXI4	DDR	AXI4
DMA	AXI4	UART	APB
CPU	AXI4	GPIO	APB
Custom IP	AXI4-Lite	Fabric	AXI4

Challenge: Different protocols can't exchange a single transfer without a bridge between them.

3.3 1.3 Solution Architecture

3.3.1 Figure 1.1: Converter Module Hierarchy



Converter Module Hierarchy

3.3.2 1.3.1 Data Width Converters

Generic Building Blocks: - **axi_data_upsize** - Accumulates N narrow beats into 1 wide beat - **axi_data_dsize** - Splits 1 wide beat into N narrow beats

Full AXI4 Converters: - **axi4_dwidth_converter_wr** - Complete write path (AW + W + B channels) - **axi4_dwidth_converter_rd** - Complete read path (AR + R channels)

3.3.3 1.3.2 Protocol Converters

AXI4 to AXI4-Lite: - **axi4_to_axil4_rd** - Read path burst decomposition - **axi4_to_axil4_wr** - Write path burst decomposition - **axi4_to_axil4** - Full bidirectional wrapper

AXI4-Lite to AXI4: - **axil4_to_axi4_rd** - Read path protocol upgrade - **axil4_to_axi4_wr** - Write path protocol upgrade - **axil4_to_axi4** - Full bidirectional wrapper

Other Converters: - **axi4_to_apb4_convert** - Full AXI4-to-APB bridge - **peakrdl_to_cmdrsp** - Register interface adapter

3.4 1.4 Key Design Decisions

3.4.1 1.4.1 Generic vs. Full Modules

The converters are built in three layers:

Layer 1: Generic Building Blocks

- `axi_data_upsize`, `axi_data_dnsz`
- Protocol-agnostic data manipulation
- Reusable across different contexts

Layer 2: Full AXI4 Converters

- `axi4_dwidth_converter_wr`, `axi4_dwidth_converter_rd`
- Compose generic blocks with AXI4 channel management
- Handle burst length adjustment, ID tracking

Layer 3: Protocol Converters

- `axi4_to_axil4`, `axi4_to_apb4`, etc.
- Full protocol translation
- State machine control

3.4.2 1.4.2 Throughput vs. Area Trade-offs

Table 1.3: Throughput vs. Area Trade-offs

Configuration	Throughput	Area	Use Case
Upsize (single buffer)	100%	1x	All narrow-to-wide
Downsize (single buffer)	0.992 beats/cycle	1x	Area-constrained

Design Decision: One buffer per direction. The downsize accepts its next wide beat during the last narrow beat of the current one — measured at 0.992 beats/cycle — so a second buffer has nothing to recover. A ping-pong mode existed for this and was removed once the single buffer was fixed.

3.4.3 1.4.3 Sideband Signal Handling

Sideband signals (WSTRB, RRESP, etc.) support three handling modes:

Table 1.4: Sideband Handling Modes

Mode	Upsize Behavior	Downsize Behavior
Concatenate	Pack narrow strobes	Slice wide strobes
Broadcast	N/A	Repeat value

Mode	Upsize Behavior	Downsize Behavior
Severity fold	Keep numeric max	N/A

3.5 1.5 Performance Characteristics

3.5.1 1.5.1 Latency

Table 1.5: Latency Characteristics

Module	Single-Beat	Burst (N beats)
axi_data_upsize	1 cycle (registered output)	RATIO narrow beats to fill a group
axi_data_dnsize (single)	1 cycle	N cycles (replacement wide beat accepted during the last narrow beat — no gap)
axi4_to_axil4	0 cycles	slave-limited (see 3.2.6)
axil4_to_axi4	0 cycles	N/A (single only)

3.5.2 1.5.2 Throughput

Table 1.6: Throughput Characteristics

Module	Configuration	Peak Throughput
axi_data_upsize	Single buffer	1 beat/cycle
axi_data_dnsize	Single buffer	0.992 beats/cycle (measured, ratio 4)
axi4_to_axil4	Burst	slave-limited; requests stream independently of responses
axil4_to_axi4	Any	1 beat/cycle

3.6 1.6 Scope

3.6.1 In Scope

- Integer width ratios (2:1, 4:1, 8:1, 16:1, etc.)
- AXI4, AXI4-Lite, and APB protocol support
- Configurable throughput vs. area trade-offs
- Generic building blocks for custom converters
- Burst-aware width conversion

3.6.2 Out of Scope

- Non-integer width ratios (e.g., 3:2 conversion)
- AXI4-Stream protocol (see Stream component)
- Buffering deeper than one beat per direction
- Clock domain crossing (use separate CDC modules)
- Address translation (handled by crossbar)

3.7 1.7 Target Applications

1. **CPU-to-DDR Integration** - 64-bit CPU to 512-bit memory controller
 2. **DMA Engines** - Variable width data movers
 3. **Mixed Protocol Systems** - AXI4 fabric with APB peripheral bus
 4. **FPGA Fabric Interfaces** - Width matching for IP integration
 5. **Register Access** - PeakRDL to custom control protocols
-

Next: [Chapter 2: Data Width Converter Blocks](#)

RTL Design Sherpa · Learning Hardware Design Through Practice [GitHub](#) · [Documentation Index](#) · [MIT License](#)

4 2.1 Generic Building Blocks

Two generic building blocks handle the data movement: **axi_data_upsize** (narrow-to-wide) and **axi_data_dnsz** (wide-to-narrow). Both are protocol-agnostic — compose them into full AXI4 converters or use them directly in custom designs.

4.1 2.1.1 Module Hierarchy

Generic Building Blocks

```
|
+-- axi_data_upsize.sv      # Narrow-to-Wide accumulator
|   +-- Accumulator buffer
|   +-- Beat counter
|   +-- Sideband packing
|
+-- axi_data_dnsize.sv      # Wide-to-Narrow splitter
|   +-- Data buffer
|   +-- Beat counter
|   +-- Sideband extraction
|   +-- Optional burst tracking
```

4.2 2.1.2 Design Philosophy

4.2.1 Separation of Concerns

The generic blocks handle **only data manipulation**: - Data packing/unpacking - Sideband signal handling - Valid/ready handshaking

They do **not** handle: - Address manipulation - Burst length adjustment - ID tracking - Protocol-specific signals (ARLEN, AWLEN, etc.)

That separation buys you three things: - Reuse in multiple contexts - Simpler verification - Cleaner composition into full converters

4.2.2 Interface Pattern

Both modules use a consistent valid/ready interface:

```
// Input (narrow for upsize, wide for dnsz)
input logic          i_valid,
output logic         o_ready,
input logic [DATA_WIDTH-1:0] i_data,
input logic [SB_WIDTH-1:0] i_sideband,
input logic          i_last,      // Optional

// Output (wide for upsize, narrow for dnsz)
output logic         o_valid,
input logic          i_ready,
output logic [DATA_WIDTH-1:0] o_data,
output logic [SB_WIDTH-1:0] o_sideband,
output logic         o_last      // Optional
```

4.3 2.1.3 Width Ratio Calculation

Both modules calculate the width ratio at elaboration time:

```
localparam int RATIO = WIDE_WIDTH / NARROW_WIDTH;  
localparam int RATIO_LOG2 = $clog2(RATIO);
```

```
// Example: 64-bit to 512-bit  
// RATIO = 512 / 64 = 8  
// RATIO_LOG2 = 3
```

Constraints: - WIDE_WIDTH must be an integer multiple of NARROW_WIDTH - Minimum ratio: 2 -
Maximum ratio: Typically 16 (limited by timing)

4.4 2.1.4 Sideband Modes

4.4.1 Upsize Sideband Modes

Table 2.1: Upsize Sideband Modes

Mode	Parameter	Behavior	Use Case
Concatenate	SB_OR_MODE=0	Pack N narrow sidebands	WSTRB packing
Severity fold	SB_OR_MODE=1	Keep the numeric max of the narrow sidebands	RRESP worst-case fold

Concatenate Example (WSTRB):

8 beats of 8-bit WSTRB → 1 beat of 64-bit WSTRB
[0]: 0xFF → output[7:0] = 0xFF
[1]: 0xF0 → output[15:8] = 0xF0
⋮
[7]: 0x0F → output[63:56] = 0x0F

OR Example (Error flags):

8 beats of error flags → 1 beat with any error
[0]: 0 → output = 0
[1]: 1 → output = 1 (error detected)
⋮
[7]: 0 → output remains 1

4.4.2 Downsize Sideband Modes

Table 2.2: Downsize Sideband Modes

Mode	Parameter	Behavior	Use Case
Slice	SB_BROADCAST=0	Extract slice per beat	WSTRB extraction
Broadcast	SB_BROADCAST=1	Repeat full value	RRESP broadcast

Slice Example (WSTRB):


1 beat of 64-bit WSTRB → 8 beats of 8-bit WSTRB
input = 0xFF_F0..._0F
[0]: output = 0xFF (bits [7:0])
[1]: output = 0xF0 (bits [15:8])
⋮
[7]: output = 0x0F (bits [63:56])

Broadcast Example (RRESP):

1 beat of 2-bit RRESP → 8 beats of 2-bit RRESP
input = OKAY (2'b00)
[0-7]: output = OKAY (all beats get same response)

4.5 2.1.5 Performance Comparison

4.5.1 Figure 2.1: Performance Comparison



Performance Comparison

Table 2.3: Performance Comparison

Module	Mode	Throughput	Latency	Area
axi_data_u psize	Single	100%	N cycles	1x
axi_data_d nsize	Single	0.992 beats/cycle	1 cycle	1x

Single-buffer downsize does not stall between wide beats.

wide_ready is asserted during the last narrow beat, so the next wide beat is accepted as the current one finishes draining. With TRACK_BURSTS=1 the replacement is suppressed at burst boundaries only.

(The former dual-buffer mode is deleted; the atomic replace above is the whole story.)

4.6 2.1.6 Resource Utilization

4.6.1 Upsize Resources (NARROW=64, WIDE=512)

Accumulator buffer: 512 bits (output register)
Beat counter: 3 bits ($\text{clog}_2(8)$)
Sideband accumulator: 64 bits (WSTRB)
Control logic: ~50 LEs

Total: ~70-100 LEs, ~580 registers

4.6.2 Downsize Resources (WIDE=512, NARROW=64)

Single Buffer:

Data buffer: 512 bits
Beat counter: 3 bits
Sideband logic: ~20 LEs
Control logic: ~50 LEs

Total: ~70-100 LEs, ~520 registers

4.7 2.1.7 Integration Guidelines

4.7.1 When to Use Generic Blocks

Use directly when: - Building custom data pipelines - Data width conversion without AXI4 protocol - Simple valid/ready streaming interfaces

Use full converters when: - Need AXI4 channel management (AW, W, B, AR, R) - Burst length adjustment required - ID tracking needed

4.7.2 Composition Example

```
// Full AXI4 write path composition
axi4_dwidth_converter_wr
#(.S_AXI_DATA_WIDTH(32), .M_AXI_DATA_WIDTH(128)) u_wr (
    // This module internally instantiates:
    // 1. Address phase skid buffer
    // 2. axi_data_upsize for write data
    // 3. Response path passthrough
    // 4. Burst length adjustment logic
);
```

Next: [axi_data_upsize Module](#)

RTL Design Sherpa · Learning Hardware Design Through Practice [GitHub](#) · [Documentation Index](#) · [MIT License](#)

5 2.2 axi_data_upsize

The **axi_data_upsize** module accumulates N narrow beats into 1 wide beat. It's the core building block for narrow-to-wide data width conversion.


5.1 2.2.1 Overview

The upsize module does four things:

1. **Data Accumulation:** Collects N narrow data beats into accumulator buffer
2. **Sideband Packing:** Concatenates or ORs sideband signals (WSTRB, etc.)
3. **Flow Control:** Manages valid/ready handshaking with single-cycle latency
4. **LAST Tracking:** Detects input LAST to generate output LAST

5.2 2.2.2 Block Diagram

5.2.1 Figure 2.2: axi_data_upsize Architecture



axi_data_upsize Architecture

5.3 2.2.3 Interface Specification

5.3.1 Parameters

Table 2.4: axi_data_upsize Parameters

Parameter	Type	Default	Description
NARROW_WIDTH	int	32	Input data width (bits)
WIDE_WIDTH	int	128	Output data width (bits)
NARROW_SB_WIDTH	int	0	Input sideband width (bits), 0 if unused
WIDE_SB_WIDTH	int	0	Output sideband width, 0 if unused
SB_OR_MODE	int	0	0=concatenate, 1=severity fold (numeric max)

5.3.2 Ports

```
module axi_data_upsize #(
    // Width Configuration
    parameter int NARROW_WIDTH    = 32,
    parameter int WIDE_WIDTH      = 128,
    parameter int NARROW_SB_WIDTH = 0,           // Sideband width (0 if
unused)
    parameter int WIDE_SB_WIDTH   = 0,           // Wide sideband width
    parameter int SB_OR_MODE      = 0,           // 0=concatenate,
1=severity fold (numeric max)

    // Calculated Parameters
    localparam int WIDTH_RATIO = WIDE_WIDTH / NARROW_WIDTH,
    localparam int PTR_WIDTH   = $clog2(WIDTH_RATIO),
    // Ensure sideband widths are at least 1 for port declarations
(unused if actual width is 0)
    localparam int NARROW_SB_PORT_WIDTH = (NARROW_SB_WIDTH > 0) ?
NARROW_SB_WIDTH : 1,
    localparam int WIDE_SB_PORT_WIDTH = (WIDE_SB_WIDTH > 0) ?
WIDE_SB_WIDTH : 1
```

```

) (
    input logic                                aclk,
    input logic                                aresetn,

    // Narrow Input (from slave or master)
    input logic                                narrow_valid,
    output logic                               narrow_ready,
    input logic [NARROW_WIDTH-1:0]            narrow_data,
    input logic [NARROW_SB_PORT_WIDTH-1:0]     narrow_sideband, // Min
width 1 to avoid [-1:0]
    input logic                                narrow_last,

    // Lane the FIRST narrow beat of a burst occupies inside the wide
    // word (addr % wide_bytes / narrow_bytes); later groups start at
    // lane 0. Tie '0 for aligned-only behavior.
    input logic [PTR_WIDTH-1:0]                start_lane,

    // Wide Output (to master or slave)
    output logic                               wide_valid,
    input logic                                wide_ready,
    output logic [WIDE_WIDTH-1:0]              wide_data,
    output logic [WIDE_SB_PORT_WIDTH-1:0]      wide_sideband, // Min
width 1 to avoid [-1:0]
    output logic                               wide_last
);

```

5.4 2.2.4 Operation

5.4.1 Accumulation Cycle

Cycle 0: s_data[0] → buffer[63:0], count = 0 → 1
Cycle 1: s_data[1] → buffer[127:64], count = 1 → 2
...
Cycle 7: s_data[7] → buffer[511:448], count = 7 → 0, m_valid = 1
Cycle 8: m_ready handshake, output complete

5.4.2 Early LAST Handling

If s_last arrives before buffer is full:

Cycle 0: s_data[0] → buffer[63:0], count = 0 → 1
Cycle 1: s_data[1] + s_last → buffer[127:64], count = 1 → 0
 m_valid = 1, m_last = 1
 Remaining bytes = don't care (masked by WSTRB)

5.4.3 State Machine

IDLE (count=0):

- s_valid=1 → load beat, increment count
- count < RATIO-1 → stay in IDLE
- count = RATIO-1 OR s_last → OUTPUT

OUTPUT (m_valid=1):

- m_ready=1 → clear buffer, → IDLE
- m_ready=0 → hold output

5.5 2.2.5 Sideband Handling

5.5.1 Concatenate Mode (SB_OR_MODE=0)

Used forWSTRB packing:

```
// Pack narrow sidebands into wide sideband
always_ff @(posedge clk) begin
    if (s_valid && s_ready) begin
        r_sideband[r_count * NARROW_SB_WIDTH +: NARROW_SB_WIDTH] <=
s_sideband;
    end
end
```

Example: 8 beats of 8-bitWSTRB to 64-bitWSTRB

```
Beat 0:WSTRB = 0xFF → output[7:0] = 0xFF
Beat 1:WSTRB = 0xF0 → output[15:8] = 0xF0
Beat 2:WSTRB = 0x0F → output[23:16] = 0x0F
...
Beat 7:WSTRB = 0xAA → output[63:56] = 0xAA
Final: output = 0xAA..._0F_F0_FF
```

5.5.2 Severity-Fold Mode (SB_OR_MODE=1)

Folds per-sub-beat responses into one wide-beat response by keeping the **numeric maximum**, not a bitwise OR. The distinction matters for the mode's primary use case, 2-bitRRESP: with the AXI encoding (OKAY=00, EXOKAY=01, SLVERR=10, DECERR=11), SLVERR | EXOKAY = DECERR — an OR inflates the error class the moment an exclusive-read beat mixes with a slave error (CONV-005). Numeric max is exactly severity order forRRESP, so the fold keeps the worst response instead:

```
// gen_or_mode: keep the numeric max across the group's sub-beats
if (narrow_valid && narrow_ready) begin
    if (r_beat_ptr == '0)
        r_sideband_accumulator <=
WIDE_SB_PORT_WIDTH'(narrow_sideband);
    else if (WIDE_SB_PORT_WIDTH'(narrow_sideband) >
r_sideband_accumulator)
        r_sideband_accumulator <=
WIDE_SB_PORT_WIDTH'(narrow_sideband);
end
```

For 1-bit sidebands (a bare error flag) max and OR coincide, so “any error in the group propagates” still holds:

```
Beat 0: RRESP = OKAY (00) → accumulator = 00
Beat 1: RRESP = EXOKAY (01) → accumulator = 01
Beat 2: RRESP = SLVERR (10) → accumulator = 10 (worst so far)
```

Beat 3: RRESP = OKAY (00) → accumulator = 10 (max retained)
Final: wide RRESP = SLVERR – not DECERR, which a bitwise OR would fabricate

5.6 2.2.6 Implementation

5.6.1 Core Logic

Wide valid is **registered**, not combinational. It is set by the narrow beat that *completes* a group, so the wide word is only offered after its final lane has actually been written:

```
// Lane the incoming narrow beat lands in (start_lane on a fresh
burst)
assign w_lane = (r_beat_ptr == '0) ? start_lane : r_beat_ptr;

// The accepted narrow beat that closes a group: last lane, or an
early LAST
logic narrow_completes_group;
assign narrow_completes_group = narrow_valid && narrow_ready &&
                                (w_lane == PTR_WIDTH'(WIDTH_RATIO-1)
||
                                narrow_last);

logic wide_accept;
assign wide_accept = r_wide_valid && wide_ready;

always_ff @(posedge aclk or negedge aresetn) begin
    if (!aresetn) begin
        r_data_accumulator <= '0;
        r_beat_ptr          <= '0;
        r_wide_valid        <= 1'b0;
        r_last_buffered     <= 1'b0;
        r_burst_fresh       <= 1'b1;
    end else begin
        if (narrow_valid && narrow_ready) begin
            // A group-opening beat ZEROES the whole accumulator and
places
            // its data in one whole-register assignment; later beats
write
            // only their own lane. Without the zeroing, an early
narrow_last
            // (or a burst starting mid-word) leaves residual data and
WSTRB
            // in the unfilled lanes and they leak into the emitted
beat.
            if (r_beat_ptr == '0)
                r_data_accumulator <= {{(WIDE_WIDTH-NARROW_WIDTH)
{1'b0}}, narrow_data}
                                << (w_lane * NARROW_WIDTH);
            else
                r_data_accumulator[w_lane*NARROW_WIDTH +:
```

```

NARROW_WIDTH] <= narrow_data;

    r_beat_ptr    <= narrow_completes_group ? '0 : (w_lane +
1'b1);
    r_burst_fresh <= narrow_last;
end

    // ONE point of control for r_wide_valid. A newly completed
group
    // dominates an in-flight wide accept -- when both land on the
same
    // cycle, two separate NBAs to this register would race.
    if (narrow_completes_group) begin
        r_wide_valid    <= 1'b1;
        r_last_buffered <= narrow_last;
    end else if (wide_accept) begin
        r_wide_valid    <= 1'b0;
        r_last_buffered <= 1'b0;
    end
end
end

assign narrow_ready = !r_wide_valid || wide_ready;
assign wide_valid   = r_wide_valid;
assign wide_last    = r_last_buffered;

```

The registered form is what makes the throughput claim in 2.2.7 true. An earlier version of this page showed a combinational `w_output_valid = (r_count == RATIO-1) || r_last` driving `m_valid`, with `s_ready` gated on `m_ready`. That machine is wrong in both directions: it asserts valid during the cycle the completing narrow beat is still being accepted, so with `m_ready` high the wide beat handshakes before the final lane's non-blocking write lands (a wide word missing its last narrow beat), and with `m_ready` low the completing narrow beat stalls (a bubble every group). It also contradicted 2.2.4, which correctly places the transition *after* the completing beat.

5.7 2.2.7 Timing

5.7.1 Latency

Table 2.5: Upsize Latency

Scenario	Latency
Full buffer (N beats)	N cycles
Early LAST (M beats)	M cycles
Output handshake	0-1 cycles

5.7.2 Throughput

100% throughput — no gaps required between input beats.

The accumulator accepts one beat per cycle, every cycle. When the output buffer completes its handshake, accumulation of the next wide beat starts immediately.

5.7.3 Critical Paths

Typical critical paths: - s_data → accumulator buffer → m_data - r_count → comparison → s_ready

Timing closure: The module is designed for single-cycle operation, with combinatorial paths only within registered stages.

5.8 2.2.8 Resource Utilization

5.8.1 Typical Resources (64-bit to 512-bit)

Accumulator buffer: 512 flip-flops
Sideband buffer: 64 flip-flops (WSTRB)
Beat counter: 3 flip-flops
Control logic: ~20 flip-flops
~50-70 LUTs

Total: ~600 flip-flops, ~50-70 LUTs

5.8.2 Scaling

Table 2.6: Upsize Resource Scaling

Configuration	Registers	LUTs
32 → 128 (4:1)	~170	~30
64 → 256 (4:1)	~330	~40
64 → 512 (8:1)	~600	~60
128 → 1024 (8:1)	~1150	~80

5.9 2.2.9 Usage Example

W-channel upsize (32 → 128) with WSTRB as the concatenated sideband:

```
axi_data_upsize #(
    .NARROW_WIDTH      (32),
    .WIDE_WIDTH        (128),
    .NARROW_SB_WIDTH   (4),    // WSTRB, narrow side
    .WIDE_SB_WIDTH     (16),   // WSTRB, wide side
    .SB_OR_MODE        (0)    // concatenate strobes
) u_w_upsize (
    .aclk               (aclk),
    .aresetn            (aresetn),
    .narrow_valid       (s_wvalid),
    .narrow_ready       (s_wready),
    .narrow_data        (s_wdata),
    .narrow_sideband    (s_wstrb),
    .narrow_last        (s_wlast),
    // '0 = aligned-only; wire addr%wide_bytes/narrow_bytes for
    // mid-word burst starts (see the wr_converter's AW-lane queue)
    .start_lane         ('0),
    .wide_valid         (m_wvalid),
    .wide_ready         (m_wready),
    .wide_data          (m_wdata),
    .wide_sideband      (m_wstrb),
    .wide_last          (m_wlast)
);
```

Next: [axi_data_dsize Module](#)

RTL Design Sherpa · Learning Hardware Design Through Practice [GitHub](#) · [Documentation Index](#) · [MIT License](#)

6 2.3 axi_data_dnsiz

The **axi_data_dnsiz** module splits 1 wide beat into N narrow beats. It accepts the next wide beat during the last narrow beat of the current one, so a steady stream costs no per-beat stall.


6.1 2.3.1 Overview

The downsize module does three things:

1. **Data Splitting:** Extracts N narrow beats from one wide beat
2. **Sideband Extraction:** Slices or broadcasts sideband signals
3. **Burst Tracking:** Optional LAST signal generation based on burst length

6.2 2.3.2 Block Diagram

6.2.1 Figure 2.3: axi_data_dnsz Single-Buffer Architecture



axi_data_dnsz Single Buffer

6.3 2.3.3 Interface Specification

6.3.1 Parameters

Table 2.7: *axi_data_dnsiz*e Parameters

Parameter	Type	Default	Description
WIDE_WIDTH	int	128	Input data width (bits)
NARROW_WIDTH	int	32	Output data width (bits)
WIDE_SB_WIDTH	int	0	Input sideband width, 0 if unused
NARROW_SB_WIDTH	int	0	Output sideband width, 0 if unused
SB_BROADCAST	int	1	0=slice, 1=broadcast sidebands
TRACK_BURSTS	int	0	Enable burst-aware LAST generation
BURST_LEN_WIDTH	int	8	Width of burst length input

6.3.2 Ports

```
module axi_data_dnsiz #(
    // Width Configuration
    parameter int WIDE_WIDTH      = 128,
    parameter int NARROW_WIDTH    = 32,
    parameter int WIDE_SB_WIDTH   = 0,           // Sideband width (0
if unused)
    parameter int NARROW_SB_WIDTH = 0,
    parameter int SB_BROADCAST    = 1,           // 1=broadcast,
0=slice
    parameter int TRACK_BURSTS    = 0,           // 1=track bursts for
LAST
    parameter int BURST_LEN_WIDTH = 8,           // Burst length
counter width
```

```

// Calculated Parameters
localparam int WIDTH_RATIO = WIDE_WIDTH / NARROW_WIDTH,
localparam int PTR_WIDTH   = $clog2(WIDTH_RATIO),
// Ensure sideband widths are at least 1 for port declarations
localparam int WIDE_SB_PORT_WIDTH = (WIDE_SB_WIDTH > 0) ?
WIDE_SB_WIDTH : 1,
localparam int NARROW_SB_PORT_WIDTH = (NARROW_SB_WIDTH > 0) ?
NARROW_SB_WIDTH : 1
) (
    input logic          aclk,
    input logic          aresetn,

    // Burst Control (only if TRACK_BURSTS=1)
    input logic [BURST_LEN_WIDTH-1:0] burst_len,          // From
address channel (ARLEN/AWLEN)
    input logic          burst_start,          // Pulse
to start new burst
    // TRACK_BURSTS only: lane the burst's first narrow beat occupies
in
    // the FIRST wide word; later wide words slice from lane 0.
    input logic [PTR_WIDTH-1:0] start_lane,

    // Wide Input (from slave or master)
    input logic          wide_valid,
    output logic         wide_ready,
    input logic [WIDE_WIDTH-1:0] wide_data,
    input logic [WIDE_SB_PORT_WIDTH-1:0] wide_sideband, // Min
width 1 to avoid [-1:0]
    input logic          wide_last,

    // Narrow Output (to master or slave)
    output logic         narrow_valid,
    input logic          narrow_ready,
    output logic [NARROW_WIDTH-1:0] narrow_data,
    output logic [NARROW_SB_PORT_WIDTH-1:0] narrow_sideband, // Min
width 1 to avoid [-1:0]
    output logic         narrow_last
);

```

6.4 2.3.4 Single-Buffer Mode Operation

6.4.1 Split Cycle

```
Cycle 0: s_data loaded → buffer, s_ready = 0
Cycle 1: m_data = buffer[63:0], count = 0, m_valid = 1
Cycle 2: m_data = buffer[127:64], count = 1
...
Cycle 8: m_data = buffer[511:448], count = 7, m_last possible
Cycle 9: wide_ready already asserted during cycle 8's last narrow beat
-- no gap
```

6.4.2 State Machine

IDLE:

- s_valid=1 → load buffer → SPLITTING

SPLITTING:

- Output beats 0 to $RATIO-1$
- m_ready=1 → increment count
- count= $RATIO-1$ AND m_ready → IDLE

6.4.3 Throughput Analysis

No per-beat load cycle

For ratio N , a wide beat costs N cycles outputting narrow beats. The load of the next wide beat is not an extra cycle: wide_ready is asserted during the N th narrow beat, so the replacement is accepted as the current beat finishes.

TRACK_BURSTS=1 is the exception. Its ready condition is mid_burst_replace, which excludes the final beat of a burst, so a cycle is given up at each burst boundary — not at each wide beat.

Measured on axi_data_dnsiz with both sides driven by the shared backtoback randomizer profile, 64 wide beats per run, timing the drain only (see measure_throughput in the dnsiz TB):

Configuration	Narrow beats	Cycles	Beats/cycle
ratio 4, single buffer	256	258	0.992
ratio 2, single buffer	128	130	0.985

Both buffering modes sustain a narrow beat every cycle, which is the most the narrow side can carry. The shortfall is a constant 2-cycle pipeline fill, not a per-beat cost — it does not grow with the run, which is how the earlier “one gap cycle per wide beat” model was ruled out.

TRACK_BURSTS=1 gives up one cycle per BURST boundary (its replace condition excludes each burst’s final beat) and measures 0.914-0.941 beats/cycle over 8 framed bursts of 32 narrow beats

— the shortfall being the same fixed fill, paid once per burst (measure_burst_throughput in the dnsz TB).

Table 2.8: Measured Throughput by Mode

Mode	Measured	Cost model
Simple, ratio 4	0.992 beats/cycle	fixed 2-cycle fill only
Simple, ratio 2	0.985 beats/cycle	fixed 2-cycle fill only
TRACK_BURSTS, ratio 4	0.914-0.941	one bubble per burst boundary

6.5 2.3.5 Sideband Handling

6.5.1 Slice Mode (SB_BROADCAST=0)

Used for WSTRB extraction:

```
// Extract sideband slice per beat  
assign m_sideband = r_sideband[r_count * NARROW_SB_WIDTH +:  
NARROW_SB_WIDTH];
```

Example: 64-bit WSTRB to 8-bit WSTRB

Input: 0xAA_BB_CC_DD_EE_FF_00_11

Beat 0: output WSTRB = 0x11 (bits [7:0])
Beat 1: output WSTRB = 0x00 (bits [15:8])
Beat 2: output WSTRB = 0xFF (bits [23:16])
...
Beat 7: output WSTRB = 0xAA (bits [63:56])

6.5.2 Broadcast Mode (SB_BROADCAST=1)

Used for RRESP:

```
// Broadcast same sideband to all beats  
assign m_sideband = r_sideband[NARROW_SB_WIDTH-1:0];
```

Example: RRESP = OKAY for all beats

Input: RRESP = 2'b00 (OKAY)

Beat 0: output RRESP = 2'b00
Beat 1: output RRESP = 2'b00
...
Beat 7: output RRESP = 2'b00

6.6 2.3.6 Burst Tracking

6.6.1 Purpose

When TRACK_BURSTS=1, the module generates narrow_last from the AXI4 burst length instead of relying on the incoming wide_last.

6.6.2 Operation

```
// from the RTL: burst_len is already in NARROW beats (the caller
frames
// the burst in output units -- ARLEN/AWLEN scaled by the
instantiator),
// so there is no xRATIO multiply here
if (TRACK_BURSTS != 0 && burst_start && !r_burst_active) begin
    r_slave_total_beats <= burst_len + 1'b1;
    r_slave_beat_count  <= '0;
    r_burst_active      <= 1'b1;
end

// count every narrow beat sent; LAST on the final one
assign narrow_last = r_wide_buffered && r_burst_active &&
    (r_slave_beat_count + 1'b1 >=
r_slave_total_beats);
```

An earlier revision showed the module multiplying burst_len by RATIO internally. It never did — the units contract is narrow beats in, and mis-framing in wide beats makes LAST fire WIDTH_RATIO times early (a real bug this exact confusion caused in the TB).

Example: ARLEN=3 (4 beats), ratio 8:1

Total narrow beats = 4 * 8 = 32
Beat 0-30: m_last = 0
Beat 31: m_last = 1

6.7 2.3.7 Implementation

6.7.1 Single-Buffer Core Logic

```
// from the RTL -- the whole point is that wide_ready re-asserts
// DURING
// the last narrow beat (atomic replace), so there is NO
// `s_ready = !r_active` style stall between wide beats:
assign wide_ready = !gen_single_buffer.r_wide_buffered ||
                    (narrow_ready && w_last_narrow_beat);

// output slice selection
assign narrow_data = gen_single_buffer.r_data_buffer
                    [r_beat_ptr * NARROW_WIDTH +: NARROW_WIDTH];
assign narrow_valid = gen_single_buffer.r_wide_buffered;

// simple mode: LAST = buffered wide_last AND final slice
assign narrow_last = gen_single_buffer.r_wide_buffered &&
                    gen_single_buffer.r_last_buffered &&
                    w_last_narrow_beat;
```

An earlier pseudocode block here used `s_ready = !r_active` — the pre-fix behaviour that stalls one cycle per wide beat, which the measured 0.992 beats/cycle disproves.

6.8 2.3.8 Resource Utilization

6.8.1 Single-Buffer (512-bit to 64-bit)

Data buffer: 512 flip-flops
Sideband buffer: 64 flip-flops
Beat counter: 3 flip-flops
Control logic: ~10 flip-flops
~30-50 LUTs

Total: ~590 flip-flops, ~30-50 LUTs

6.8.2 Measured Throughput

Table 2.9: Resources and measured rate

Registers	LUTs	Throughput
590	40	0.992 beats/cycle (ratio 4)

A narrow beat every cycle, so there is nothing left for a second buffer to recover. The ping-pong DUAL_BUFFER mode this table used to compare against was removed once the single buffer was fixed to accept its replacement during the last narrow beat — nothing instantiated it and it measured no faster.

6.9 2.3.9 Usage Example

R-channel downsize (128 → 32) with RRESP broadcast and burst tracking.

burst_start/burst_len are NOT optional with TRACK_BURSTS(1): r_burst_active only sets on a burst_start pulse, so leaving them unconnected silently produces framing with no LAST at all. burst_len is in NARROW beats minus one.

```
axi_data_dnsiz #(  
  .WIDE_WIDTH      (128),  
  .NARROW_WIDTH    (32),  
  .WIDE_SB_WIDTH   (2),      // RRESP  
  .NARROW_SB_WIDTH (2),  
  .SB_BROADCAST    (1),  
  .TRACK_BURSTS    (1),  
  .BURST_LEN_WIDTH (8)  
) u_r_dnsiz (  
  .aclk      (aclk),  
  .aresetn   (aresetn),  
  .burst_len (8'(narrow_beats - 1)), // narrow-beat units  
  .burst_start (ar_accepted),  
  // '0 = aligned-only; wire addr%wide_bytes/narrow_bytes for  
  // mid-word burst starts (see the rd converter's length+lane FIFO)  
  .start_lane ('0),  
  .wide_valid (s_rvalid),  
  .wide_ready (s_rready),  
  .wide_data  (s_rdata),  
  .wide_sideband (s_rresp),  
  .wide_last   (1'b0),      // counter drives LAST  
here  
  .narrow_valid (m_rvalid),  
  .narrow_ready (m_rready),  
  .narrow_data  (m_rdata),  
  .narrow_sideband (m_rresp),  
  .narrow_last   (m_rlast)  
);
```

Next: [axi4_dwidth_converter_wr](#)

RTL Design Sherpa · Learning Hardware Design Through Practice [GitHub](#) · [Documentation Index](#) · [MIT License](#)

7 2.5 axi4_dwidth_converter_wr

The **axi4_dwidth_converter_wr** module is the complete AXI4 write path — AW, W, and B channels with burst length adjustment for the new width.


7.1 2.5.1 Overview

The write converter combines the generic `axi_data_upsize` with AXI4 protocol handling:

1. **Address Channel (AW):** Passes through with burst length adjustment
2. **Write Data Channel (W):** Uses `axi_data_upsize` for data packing
3. **Response Channel (B):** Passes through unchanged
4. **Burst Length Adjustment:** Converts AWLEN based on width ratio

7.2 2.5.2 Block Diagram

7.2.1 Figure 2.6: Write Converter Architecture



Write Converter Architecture

7.3 2.5.3 Interface Specification

7.3.1 Parameters

Table 2.14: Write Converter Parameters

Parameter	Type	Default	Description
S_AXI_DATA_WIDTH	int	32	Slave-side data width
M_AXI_DATA_WIDTH	int	128	Master-side data width
AXI_ID_WIDTH	int	8	Transaction ID width
AXI_ADDR_WIDTH	int	32	Address width
AXI_USER_WIDTH	int	1	User-signal width
SKID_DEPTH_AW	int	2	AW skid buffer depth
SKID_DEPTH_W	int	4	W skid buffer depth
SKID_DEPTH_B	int	2	B skid buffer depth

7.3.2 Ports

```
module axi4_dwidth_converter_wr #(
    // Width Configuration
    parameter int S_AXI_DATA_WIDTH = 32,
    parameter int M_AXI_DATA_WIDTH = 128,
    parameter int AXI_ID_WIDTH     = 8,
    parameter int AXI_ADDR_WIDTH   = 32,
    parameter int AXI_USER_WIDTH   = 1,

    // Skid Buffer Depths (for timing closure)
    parameter int SKID_DEPTH_AW    = 2,
    parameter int SKID_DEPTH_W     = 4,
    parameter int SKID_DEPTH_B     = 2,

    // Calculated Parameters
    localparam int S_STRB_WIDTH = S_AXI_DATA_WIDTH / 8,
    localparam int M_STRB_WIDTH = M_AXI_DATA_WIDTH / 8,
```

```

    localparam int WIDTH_RATIO = (S_AXI_DATA_WIDTH <
M_AXI_DATA_WIDTH) ?
                                (M_AXI_DATA_WIDTH /
S_AXI_DATA_WIDTH) :
                                (S_AXI_DATA_WIDTH /
M_AXI_DATA_WIDTH),
    localparam bit UPSIZE      = (S_AXI_DATA_WIDTH <
M_AXI_DATA_WIDTH) ? 1'b1 : 1'b0,
    localparam bit DOWNSIZE    = (S_AXI_DATA_WIDTH >
M_AXI_DATA_WIDTH) ? 1'b1 : 1'b0,

    // Skid buffer packed widths
    localparam int AW_WIDTH = AXI_ID_WIDTH + AXI_ADDR_WIDTH + 8 + 3 +
2 + 1 + 4 + 3 + 4 + 4 + AXI_USER_WIDTH,
    localparam int W_WIDTH  = S_AXI_DATA_WIDTH + S_STRB_WIDTH + 1 +
AXI_USER_WIDTH,
    localparam int B_WIDTH  = AXI_ID_WIDTH + 2 + AXI_USER_WIDTH
) (
    // Clock and Reset
    input logic          aclk,
    input logic          aresetn,

    //=====
    // Slave AXI Write Interface
    //=====

    // Write Address Channel
    input logic [AXI_ID_WIDTH-1:0] s_axi_awid,
    input logic [AXI_ADDR_WIDTH-1:0] s_axi_awaddr,
    input logic [7:0] s_axi_awlen,
    input logic [2:0] s_axi_awsz,
    input logic [1:0] s_axi_awburst,
    input logic s_axi_awlock,
    input logic [3:0] s_axi_awcache,
    input logic [2:0] s_axi_awprot,
    input logic [3:0] s_axi_awqos,
    input logic [3:0] s_axi_awregion,
    input logic [AXI_USER_WIDTH-1:0] s_axi_awuser,
    input logic s_axi_awvalid,
    output logic s_axi_awready,

    // Write Data Channel
    input logic [S_AXI_DATA_WIDTH-1:0] s_axi_wdata,

```

```

input  logic [S_STRB_WIDTH-1:0]    s_axi_wstrb,
input  logic                               s_axi_wlast,
input  logic [AXI_USER_WIDTH-1:0]   s_axi_wuser,
input  logic                               s_axi_wvalid,
output logic                               s_axi_wready,

// Write Response Channel
output logic [AXI_ID_WIDTH-1:0]     s_axi_bid,
output logic [1:0]                   s_axi_bresp,
output logic [AXI_USER_WIDTH-1:0]   s_axi_buser,
output logic                               s_axi_bvalid,
input  logic                               s_axi_bready,

//=====
// Master AXI Write Interface
//=====

// Write Address Channel
output logic [AXI_ID_WIDTH-1:0]     m_axi_awid,
output logic [AXI_ADDR_WIDTH-1:0]   m_axi_awaddr,
output logic [7:0]                   m_axi_awlen,
output logic [2:0]                   m_axi_awsz,
output logic [1:0]                   m_axi_awburst,
output logic                               m_axi_awlock,
output logic [3:0]                   m_axi_awcache,
output logic [2:0]                   m_axi_awprot,
output logic [3:0]                   m_axi_awqos,
output logic [3:0]                   m_axi_awregion,
output logic [AXI_USER_WIDTH-1:0]   m_axi_awuser,
output logic                               m_axi_awvalid,
input  logic                               m_axi_awready,

// Write Data Channel
output logic [M_AXI_DATA_WIDTH-1:0] m_axi_wdata,
output logic [M_STRB_WIDTH-1:0]     m_axi_wstrb,
output logic                               m_axi_wlast,
output logic [AXI_USER_WIDTH-1:0]   m_axi_wuser,
output logic                               m_axi_wvalid,
input  logic                               m_axi_wready,

// Write Response Channel
input  logic [AXI_ID_WIDTH-1:0]     m_axi_bid,
input  logic [1:0]                   m_axi_bresp,

```

```
    input logic [AXI_USER_WIDTH-1:0] m_axi_buser,  
    input logic m_axi_bvalid,  
    output logic m_axi_bready  
);
```

7.4 2.5.4 Burst Length Conversion

7.4.1 Ratio Calculation

```
// Direction-aware: wider / narrower (a plain M/S divide evaluates to  
// 0 in DOWNSIZE mode and $clog2(0) is illegal)
```

```
localparam int RATIO = (S_AXI_DATA_WIDTH < M_AXI_DATA_WIDTH)  
    ? (M_AXI_DATA_WIDTH / S_AXI_DATA_WIDTH)  
    : (S_AXI_DATA_WIDTH / M_AXI_DATA_WIDTH);
```

```
localparam int RATIO_LOG2 = $clog2(RATIO);
```

```
// upsize: New AWLEN = ceil((AWLEN + 1) / RATIO) - 1 (round UP --  
// a floor divide underflows on non-multiples, see 2.5.5)  
// downsize: split into master bursts of <= 256 beats (see 2.5.5)
```

7.4.2 Examples

Table 2.15: Burst Length Conversion Examples

S_DATA	M_DATA	Ratio	S_AWLE	S_beats	M_AWLE	M_beats
			N		N	
64	512	8	7	8	0	1
64	512	8	15	16	1	2
64	256	4	3	4	0	1
64	256	4	7	8	1	2

7.4.3 Non-Aligned Bursts

When the burst length is not a multiple of the ratio:

S_AWLEN = 5 (6 beats), RATIO = 8

M_AWLEN = 0 (1 beat)

The 6 narrow beats pack into 1 wide beat.

Remaining 2 positions have WSTRB = 0 (no write).

7.5 2.5.5 Address Channel Handling

7.5.1 AW Passthrough with Adjustment

The rewrite differs by direction, so the RTL has two generate branches.

Narrow → wide (upsized). Beats combine, so the count divides — and it must round UP, because a burst that is not a whole multiple of the ratio still needs a final partial wide beat:

```
assign m_axi_awlen = ((int_awlen + 8'(WIDTH_RATIO)) / 8'(WIDTH_RATIO))  
- 8'd1;
```

Wide → narrow (downsize). Each wide beat becomes *RATIO* narrow ones — and the product does not fit a burst. AXI4 allows 256 beats, so a full-length slave burst needs up to $256 * \text{WIDTH_RATIO}$ narrow beats, which is neither expressible in the 8-bit AWLEN nor a legal burst. One slave burst is therefore **split** into as many master bursts of ≤ 256 beats as it takes:

```
// remaining-beat counter, 9 + $clog2(WIDTH_RATIO) bits -- 8 is  
exactly  
// what a multiply-into-AWLEN overflows  
r_split_remaining <= (CNTW'(int_awlen) + CNTW'(1)) *  
CNTW'(WIDTH_RATIO);  
...  
assign m_axi_awlen = 8'(w_this_beats - 9'd1); // w_this_beats =  
min(remaining, 256)
```

Each issued master burst is recorded in a split queue — its beat count frames that burst's WLAST on the W path, and a final-burst flag drives the B fold (several master responses collapse into one slave response, worst case wins). The address advances by $256 * \text{M_STRB_WIDTH}$ per burst for INCR and holds for FIXED; WRAP never reaches a split, since AXI4 caps it at 16 beats and $16 * \text{RATIO} \leq 256$ for every supported ratio. The slave's AW is consumed only when its final master burst has been issued.

Rounding up is the whole point of the $+ \text{WIDTH_RATIO}$ term. A floor divide underflows: $\text{AWLEN}=5$ (6 beats) at $\text{RATIO}=8$ gives $(6 \gg 3) - 1$, which is -1 — 8'hFF, a 256-beat burst — where the correct answer is 0, one wide beat.

Size is not derived from the incoming AWSIZE. Both branches drive the master side at its own full width:

```
assign m_axi_awsz = MASTER_SIZE[2:0];
```

Neither converter checks address alignment on the downsize path. The read converter aligns the address it issues on its UPSIZE path only — a wide access cannot start mid-word (see 2.6.5).

Upsize INCR bursts may start mid-wide-word. The AXI-correct behavior: the first wide beat carries the narrow data in the byte lanes the ADDRESS selects, withWSTRB covering only those lanes (leading lanes byte-disabled, not clobbered). Three pieces cooperate:

- `m_axi_awlen` counts the lane offset: `ceil((start_lane + narrow_beats) / RATIO)` wide beats. `AWADDR` passes through unchanged — unaligned is legal AXI; the first beat is partial within its size container.
- an AW-lane queue (same inline pattern as the read converter's burst-length FIFO) records each accepted AW's start lane; W beats are held off until their AW is queued (the packer's lane comes from `AWADDR`), and the entry pops on the NARROW side's last beat — the head must track the burst the CURRENT narrow beat belongs to, and a wide-side pop raced back-to-back bursts.
- `axi_data_upsize` takes a `start_lane` input: the first narrow beat of a burst lands at that lane (data AND `WSTRB` shifted, leading slots '0'); later wide groups of the burst start at lane 0.

FIXED and WRAP keep the wide-aligned requirement (asserted in simulation) — their lane semantics through the packer are not defined.

7.5.2 Skid Buffer for AW

The channel buffers are `gaxi_skid_buffer` (there is no `axi_skid_buffer`), and they carry the UNMODIFIED slave fields — the length rewrite happens after the skid, in the splitter:

```
gaxi_skid_buffer #(
    .DEPTH(SKID_DEPTH_AW),
    .DATA_WIDTH(AW_WIDTH)
) aw_skid (
    .axi_aclk      (aclk),
    .axi_aresetn   (aresetn),
    .wr_valid      (s_axi_awvalid),
    .wr_ready      (s_axi_awready),
    .wr_data       ({s_axi_awid, s_axi_awaddr, s_axi_awlen, s_axi_awsz,
                     s_axi_awburst, s_axi_awlock, s_axi_awcache,
s_axi_awprot,
                     s_axi_awqos, s_axi_awregion, s_axi_awuser}),
    .rd_valid      (int_aw_valid),
    .rd_ready      (int_aw_ready),
    .rd_data       (int_aw_data),
    .count         (),
    .rd_count      ()
);
```

7.6 2.5.6 Write Data Channel

7.6.1 Upsize Instance

```
axi_data_upsize #(
    .NARROW_WIDTH      (S_AXI_DATA_WIDTH),
    .WIDE_WIDTH        (M_AXI_DATA_WIDTH),
    .NARROW_SB_WIDTH   (S_STRB_WIDTH),    // WSTRB
    .WIDE_SB_WIDTH     (M_STRB_WIDTH),
    .SB_OR_MODE        (0)                // Concatenate WSTRB
) u_w_upsize (
    .aclk               (aclk),
    .aresetn            (aresetn),
    .narrow_valid       (int_w_valid),
    .narrow_ready       (int_w_ready),
    .narrow_data        (int_wdata),
    .narrow_sideband    (int_wstrb),
    .narrow_last        (int_wlast),
    .start_lane         (w_upsize_start_lane[$clog2(WIDTH_RATIO)-1:0]),
    .wide_valid         (m_axi_wvalid),
    .wide_ready         (m_axi_wready),
    .wide_data          (m_axi_wdata),
    .wide_sideband      (m_axi_wstrb),
    .wide_last          (m_axi_wlast)
);
```

7.7 2.5.7 Response Channel

7.7.1 B Channel

Upsize passes the response through unchanged. Downsize cannot: a split slave burst receives several master responses, and the slave expects ONE. Non-final responses are consumed immediately and folded worst-case- wins; the final one carries the folded result (see 2.5.5). The passthrough below is the UPSIZE branch:

```
// gen_b_pass (upsized)
assign s_bvalid = m_bvalid;
assign m_bready = s_bready;
assign s_bid    = m_bid;
assign s_bresp  = m_bresp;
```

7.7.2 Response Ordering

Responses return in order **provided the master-side slave returns them in issue order across all IDs** – see the constraint below, which is a real restriction and not a formality. Given that, the fold is safe because: - Upsize doesn't reorder data - B response generated after all W beats accepted

An earlier version of this list gave “single outstanding transaction per ID” as sufficient. It is not: the B fold compares no IDs at all.

7.7.3 Ordering Constraint (applies to both directions)

The master-side slave must return responses in AR/AW issue order across ALL IDs. Neither converter carries a transaction ID through its data path, so neither can attribute an out-of-order completion to the burst it belongs to.

Read side: the burst-length queue holds {narrow ARLEN, start lane} and nothing else, and s_axi_rid is a single register latched on every master R handshake. If a slave returns ID1's data before ID0's, ID1's beats are counted against ID0's queue head, s_axi_rlast fires at the wrong boundary, one aggregated wide beat can mix two transactions, and s_axi_rid is whichever ID handshaked most recently.

Write side: the split queue is popped by split_b_pop = m_axi_bvalid && m_axi_bready – any arriving B pops the head, whatever its ID – and m_axi_bready = split_b_final ?
int_b_ready : 1'b1 swallows a non-final response unconditionally. A B that arrives out of order is folded into the wrong burst's accumulator: the burst it actually belonged to never gets its response (the master hangs), the other burst's response is built from the wrong beats, and every later burst's framing is shifted by one.

AXI4 permits both behaviours. A slave may complete different-ID transactions out of order and may interleave read data across IDs, and a multi-ported DDR controller normally does. Satisfying “one outstanding transaction per ID” is **not** sufficient – two IDs with one transaction each meet that and still break the fold.

Safe configurations are: an in-order master-side slave, or a single ID outstanding at a time. Nothing in the RTL enforces or detects a violation, so a breach corrupts data silently rather than failing (tracked as CONV-001).

7.8 2.5.8 AW/W Synchronization

7.8.1 Challenge

AXI4 allows AW to arrive before, with, or after W data. The converter must handle all cases:

1. **AW before W:** Normal pipelining
2. **W before AW:** Data buffered until AW arrives
3. **Interleaved:** Multiple transactions in flight

7.8.2 Solution

The **upsized** path carries an AW-lane queue (see 2.5.5): each accepted AW pushes its start lane, and W beats are held off until their AW is queued — the packer's lane placement comes from AWADDR, so W genuinely cannot run ahead of its address. The entry pops on the narrow side's last beat.

The **downsized** path does carry one — the split queue of 2.5.5. Each issued master burst pushes {final-burst flag, beat count}; the beat count frames that burst's WLAST on the W path and the flag drives the B fold. AW runs ahead of W, which is precisely why that AW-derived information must be queued.

Compare the read converter, which needs a queue in its UPSIZE mode: its downsize block has no length queue of its own, so overlapping read bursts would lose their framing without an external ARLEN queue.

7.9 2.5.9 Resource Utilization

7.9.1 Typical Resources (64→512, ID=4, ADDR=64)

AW skid buffer: ~200 flip-flops
W upsize: ~600 flip-flops, ~60 LUTs
B skid buffer: ~20 flip-flops
Control logic: ~100 LUTs

Total: ~820 flip-flops, ~160 LUTs

Hand estimates, not synthesis results. AW-derived storage: the downsize split queue, or the upsize AW-lane queue (16 x lane-width entries + two 5-bit pointers) -- see 2.5.5 for both.

7.10 2.5.10 Timing

7.10.1 Latency

Table 2.16: Write Converter Latency

Path	Latency
AW passthrough	1-2 cycles (skid)
W upsize	N cycles (accumulation)
B passthrough	1 cycle (skid)

7.10.2 Throughput

- AW channel: 1 transaction/cycle
- W channel: 100% (upsized is 100%)
- B channel: 1 response/cycle

7.11 2.5.11 Usage Example

```
axi4_dwidth_converter_wr #(
    .S_AXI_DATA_WIDTH (128),    // slave wide
    .M_AXI_DATA_WIDTH (32),     // master narrow (downsize)
    .AXI_ID_WIDTH       (8),
    .AXI_ADDR_WIDTH     (32),
    .SKID_DEPTH_AW      (2),
    .SKID_DEPTH_W       (4),
    .SKID_DEPTH_B       (2)
) u_conv (
    .aclk                (aclk),
    .aresetn             (aresetn),
    // slave side: s_axi_aw*/... (full AXI4 channel set)
    .s_axi_awvalid       (cpu_awvalid),
    .s_axi_awready       (cpu_awready),
    // ...
    // master side: m_axi_* toward the narrow fabric
    .m_axi_awvalid       (mem_awvalid),
    .m_axi_awready       (mem_awready)
    // ...
);
```

All channel ports carry the s_axi_/m_axi_ prefix; the full list is the module header. Skid depths are per channel — there is no single SKID_DEPTH.

Next: [axi4_dwidth_converter_rd](#)

RTL Design Sherpa · Learning Hardware Design Through Practice [GitHub](#) · [Documentation Index](#) · [MIT License](#)

8 2.6 axi4_dwidth_converter_rd

The **axi4_dwidth_converter_rd** module is the complete AXI4 read path — AR and R channels with burst length adjustment and burst-aware RLAST generation.


8.1 2.6.1 Overview

The read converter combines the generic `axi_data_dsize` with AXI4 protocol handling:

1. **Address Channel (AR):** Passes through with burst length adjustment
2. **Read Data Channel (R):** Uses `axi_data_dsize` for data splitting
3. **Burst Tracking:** Generates correct RLAST based on original ARLEN
4. **Response Broadcasting:** Propagates RRESP to all narrow beats

8.2 2.6.2 Block Diagram

8.2.1 Figure 2.7: Read Converter Architecture



Read Converter Architecture

8.3 2.6.3 Interface Specification

8.3.1 Parameters

Table 2.17: Read Converter Parameters

Parameter	Type	Default	Description
S_AXI_DATA_WIDTH	int	32	Slave-side data width
M_AXI_DATA_WIDTH	int	128	Master-side data width
AXI_ID_WIDTH	int	8	Transaction ID width
AXI_ADDR_WIDTH	int	32	Address width
AXI_USER_WIDTH	int	1	User-signal width
SKID_DEPTH_AR	int	2	AR skid buffer depth
SKID_DEPTH_R	int	4	R skid buffer depth

8.3.2 Ports

```
module axi4_dwidth_converter_rd #(
    // Width Configuration
    parameter int S_AXI_DATA_WIDTH = 32,
    parameter int M_AXI_DATA_WIDTH = 128,
    parameter int AXI_ID_WIDTH = 8,
    parameter int AXI_ADDR_WIDTH = 32,
    parameter int AXI_USER_WIDTH = 1,

    // Skid Buffer Depths (for timing closure)
    parameter int SKID_DEPTH_AR = 2,
    parameter int SKID_DEPTH_R = 4,

    // Calculated Parameters
    localparam int S_STRB_WIDTH = S_AXI_DATA_WIDTH / 8,
    localparam int M_STRB_WIDTH = M_AXI_DATA_WIDTH / 8,
    localparam int WIDTH_RATIO = (S_AXI_DATA_WIDTH <
M_AXI_DATA_WIDTH) ?
(M_AXI_DATA_WIDTH /
```

```

S_AXI_DATA_WIDTH) :
                                (S_AXI_DATA_WIDTH /
M_AXI_DATA_WIDTH),
    localparam bit UPSIZE      = (S_AXI_DATA_WIDTH <
M_AXI_DATA_WIDTH) ? 1'b1 : 1'b0,
    localparam bit DOWNSIZE    = (S_AXI_DATA_WIDTH >
M_AXI_DATA_WIDTH) ? 1'b1 : 1'b0,

    // Skid buffer packed widths
    localparam int AR_WIDTH = AXI_ID_WIDTH + AXI_ADDR_WIDTH + 8 + 3 +
2 + 1 + 4 + 3 + 4 + 4 + AXI_USER_WIDTH,
    localparam int R_WIDTH  = S_AXI_DATA_WIDTH + 2 + AXI_USER_WIDTH +
1 + AXI_ID_WIDTH
) (
    // Clock and Reset
    input logic          aclk,
    input logic          aresetn,

//=====
//
    // Slave AXI Read Interface

//=====
//=====

    // Read Address Channel
    input logic [AXI_ID_WIDTH-1:0] s_axi_arid,
    input logic [AXI_ADDR_WIDTH-1:0] s_axi_araddr,
    input logic [7:0] s_axi_arlen,
    input logic [2:0] s_axi_arsize,
    input logic [1:0] s_axi_arburst,
    input logic s_axi_arlock,
    input logic [3:0] s_axi_arcache,
    input logic [2:0] s_axi_arprot,
    input logic [3:0] s_axi_arqos,
    input logic [3:0] s_axi_arregion,
    input logic [AXI_USER_WIDTH-1:0] s_axi_aruser,
    input logic s_axi_arvalid,
    output logic s_axi_arready,

    // Read Data Channel
    output logic [AXI_ID_WIDTH-1:0] s_axi_rid,
    output logic [S_AXI_DATA_WIDTH-1:0] s_axi_rdata,
    output logic [1:0] s_axi_rresp,
    output logic s_axi_rlast,
    output logic [AXI_USER_WIDTH-1:0] s_axi_ruser,

```

```

output logic                                     s_axi_rvalid,
input  logic                                     s_axi_rready,

//=====
// Master AXI Read Interface
//=====

// Read Address Channel
output logic [AXI_ID_WIDTH-1:0] m_axi_arid,
output logic [AXI_ADDR_WIDTH-1:0] m_axi_araddr,
output logic [7:0] m_axi_arlen,
output logic [2:0] m_axi_arsize,
output logic [1:0] m_axi_arburst,
output logic m_axi_arlock,
output logic [3:0] m_axi_arcache,
output logic [2:0] m_axi_arprot,
output logic [3:0] m_axi_arqos,
output logic [3:0] m_axi_arregion,
output logic [AXI_USER_WIDTH-1:0] m_axi_aruser,
output logic m_axi_arvalid,
input logic m_axi_arready,

// Read Data Channel
input logic [AXI_ID_WIDTH-1:0] m_axi_rid,
input logic [M_AXI_DATA_WIDTH-1:0] m_axi_rdata,
input logic [1:0] m_axi_rresp,
input logic m_axi_rlast,
input logic [AXI_USER_WIDTH-1:0] m_axi_ruser,
input logic m_axi_rvalid,
output logic m_axi_rready
);

```

8.4 2.6.4 Burst Length Conversion

8.4.1 Ratio Calculation

Same as write converter:

```
// Direction-aware, same as the write converter (2.5.4)
localparam int RATIO = (S_AXI_DATA_WIDTH < M_AXI_DATA_WIDTH)
                        ? (M_AXI_DATA_WIDTH / S_AXI_DATA_WIDTH)
                        : (S_AXI_DATA_WIDTH / M_AXI_DATA_WIDTH);
localparam int RATIO_LOG2 = $clog2(RATIO);

// upsize: New ARLEN = ceil((start_lane + ARLEN + 1) / RATIO) - 1
// (round UP; the start lane counts for mid-word bursts, see 2.6.5)
// downsize: split into master bursts of <= 256 beats (see 2.6.5)
```

8.4.2 Examples

Table 2.18: Read Burst Length Conversion

S_DATA	M_DATA	Ratio	S_ARLE		M_ARLE	
			N	S_beats	N	M_beats
64	512	8	7	8	0	1
64	512	8	15	16	1	2
64	512	8	31	32	3	4

8.5 2.6.5 Address Channel Handling

8.5.1 AR Passthrough with Adjustment

Same two directions as the write converter (see 2.5.4 for why the upsize divide rounds up):

```
// narrow -> wide (upsized): beats combine, round up; the start lane
// counts toward the wide total for mid-word bursts
assign m_axi_arlen = 8'(((w_ar_lane + 10'(int_arlen) +
10'(WIDTH_RATIO))
/ 10'(WIDTH_RATIO)) - 10'd1);

// wide -> narrow (downsize): each wide beat becomes RATIO narrow
beats,
// and the product can exceed both AWLEN's 8 bits and the 256-beat
legal
// maximum -- so one slave burst is SPLIT into master bursts of <= 256
// beats (same mechanism as the write converter, see 2.5.5)
assign m_axi_arlen = 8'(w_this_beats - 9'd1); // min(remaining, 256)

// size is the master's own full width, not a shift of ARSIZE
assign m_axi_arsize = MASTER_SIZE[2:0];
```

For a split read the slave must still see ONE burst: each master burst returns its own RLAST, and every one except the final master burst's is masked out of the upsize, whose accumulation simply continues across the boundary. The masking is safe by construction — 256 narrow beats is a whole number of wide beats at every ratio, so a masked boundary can never land mid-accumulation. A one-bit flag queue, pushed per issued AR and popped per master RLAST, says which burst is final.

On the UPSIZE path the issued address is aligned down to the master data width — the slave returns whole wide words — and mid-word INCR starts are handled by the R slicer: the burst-length FIFO carries the start lane alongside the narrow length, `m_axi_arlen` counts it (`ceil((start_lane + narrow_beats) / RATIO)` wide beats), and `axi_data_dnsz` slices the burst's FIRST wide word from that lane, so the narrow master receives the bytes it actually addressed. Later wide words slice from lane 0. FIXED/WRAP keep the wide-aligned requirement (asserted in simulation). The downsize path issues narrow accesses and passes the address through the splitter unmodified apart from the per-burst advance:

```
localparam int ALIGN_BITS = $clog2(M_STRB_WIDTH);
assign aligned_araddr = {int_araddr[AXI_ADDR_WIDTH-1:ALIGN_BITS],
                        {ALIGN_BITS{1'b0}}};
assign m_axi_araddr = aligned_araddr;
```

8.5.2 Burst-Length FIFO

Only the wide→narrow R data path needs one — that is the converter’s **UPSIZE** mode (S narrower than M: wide master read data sliced down to narrow slave beats through `axi_data_dnsz`). The downsize block ignores a `burst_start` pulse while a burst is active and keeps no length queue of its own, so framing only the first burst would collapse N read bursts into one — bursts 2..N would drain with `narrow_last` never asserting. A small queue holds one narrow ARLEN per outstanding burst: its head feeds the downsize, and it pops as each narrow burst completes, so every burst is framed however they overlap.

It is an inline circular buffer rather than a `fifo_sync` instance, deliberately: this converter is widely instantiated and a submodule here would add a filelist dependency to every consumer. It stores the narrow length AND the burst’s start lane (mid-word INCR starts, CONV-006) — ID is carried on the AXI channels, not through this queue — and AR is back-pressured when full, so it cannot overflow.

```
localparam int BLEN_FIFO_DEPTH = 16;
localparam int BLEN_AW          = $clog2(BLEN_FIFO_DEPTH);

localparam int R_LANE_W = $clog2(WIDTH_RATIO);
logic [7:0]      blen_mem      [BLEN_FIFO_DEPTH];
logic [R_LANE_W-1:0] blen_lane_mem [BLEN_FIFO_DEPTH]; //
start lane
for full/empty
logic [BLEN_AW:0]      blen_wptr, blen_rptr; // extra MSB
logic
w_blen_push, w_blen_pop;

// AR accepted -> enqueue its narrow length. A slave-side
(narrow)
// burst completes on its last-beat handshake -> dequeue.
assign w_blen_push = int_ar_valid && int_ar_ready;
assign w_blen_pop  = int_r_valid && int_r_ready &&
int_rlast;
```

The narrow→wide R data path (the converter’s **DOWNSIZE** mode) needs no such queue; its generate branch ties the shared handshake wires off inert.

8.6 2.6.6 Read Data Channel

8.6.1 Downsize Instance

```
axi_data_dsize #(
    .WIDE_WIDTH      (M_AXI_DATA_WIDTH),
    .NARROW_WIDTH    (S_AXI_DATA_WIDTH),
    .WIDE_SB_WIDTH   (2),           // RRESP
    .NARROW_SB_WIDTH (2),
    .SB_BROADCAST    (1),           // Broadcast RRESP
    .TRACK_BURSTS    (1),
    .BURST_LEN_WIDTH (8)
) u_r_dsize (
    .aclk             (aclk),
    .aresetn          (aresetn),
    // burst framing is MANDATORY with TRACK_BURSTS(1): burst_len is
in
    // NARROW beats - 1, and burst_start is LEVEL-HELD while the
length
    // FIFO is non-empty (the dsize samples it only when idle, which
is
    // what frames every overlapping burst); leaving them off silently
    // produces framing with no LAST (see 2.3.9)
    .burst_len        (w_blen_rd_data),
    .burst_start       (w_blen_rd_valid),
    .start_lane        (w_blen_rd_lane[$clog2(WIDTH_RATIO)-1:0]),
    .wide_valid        (m_axi_rvalid),
    .wide_ready        (m_axi_rready),
    .wide_data         (m_axi_rdata),
    .wide_sideband     (m_axi_rresp),
    .wide_last         (m_axi_rlast),
    .narrow_valid      (int_r_valid),
    .narrow_ready      (int_r_ready),
    .narrow_data       (int_rdata),
    .narrow_sideband   (int_rresp),
    .narrow_last       (int_rlast)
);
```

8.7 2.6.7 RLAST Generation

There is no local RLAST tracker in the converter. The downsize block generates `narrow_last` itself in `TRACK_BURSTS` mode, framed by the burst-length FIFO of 2.6.5: each accepted AR pushes its narrow-beat length, the FIFO head drives `burst_len/burst_start`, and the `dnsize` counts narrow beats against it — `int_rlast` comes out of the `dnsize` and passes to `s_axi_rlast` through the R skid.

(An earlier revision showed a standalone counter loading $(arlen + 1) * \text{RATIO} - 1$, the same `xRATIO` framing 2.3.4 calls out as the classic mis-framing bug. No such multiply exists anywhere in the converter: in `UPSIZE` mode the slave side is already narrow, so the `ARLEN` pushed into the FIFO is already in narrow-beat units — `blen_mem[...] <= int_arlen`; stores it unchanged.)

8.8 2.6.8 RID Handling

8.8.1 ID Passthrough

RID is sampled from the master side and HELD:

```
// from the RTL: latch rid/ruser on every master R handshake; AXI4
// keeps RID constant across a transaction's beats, so "most recent
// rid" is correct for whatever aggregated beat is being emitted
`ALWAYS_FF_RST(aclk, aresetn,
... else if (m_axi_rvalid && m_axi_rready) begin
    r_rid_held    <= m_axi_rid;
    r_ruser_held <= m_axi_ruser;
end
)
assign int_rid = r_rid_held;
```

8.8.2 Ordering Constraint (applies to both directions)

The master-side slave must return responses in AR/AW issue order across ALL IDs. Neither converter carries a transaction ID through its data path, so neither can attribute an out-of-order completion to the burst it belongs to.

Read side: the burst-length queue holds {narrow ARLEN, start lane} and nothing else, and s_axi_rid is a single register latched on every master R handshake. If a slave returns ID1's data before ID0's, ID1's beats are counted against ID0's queue head, s_axi_rlast fires at the wrong boundary, one aggregated wide beat can mix two transactions, and s_axi_rid is whichever ID handshaked most recently.

Write side: the split queue is popped by split_b_pop = m_axi_bvalid && m_axi_bready – any arriving B pops the head, whatever its ID – and m_axi_bready = split_b_final ?
int_b_ready : 1'b1 swallows a non-final response unconditionally. A B that arrives out of order is folded into the wrong burst's accumulator: the burst it actually belonged to never gets its response (the master hangs), the other burst's response is built from the wrong beats, and every later burst's framing is shifted by one.

AXI4 permits both behaviours. A slave may complete different-ID transactions out of order and may interleave read data across IDs, and a multi-ported DDR controller normally does. Satisfying “one outstanding transaction per ID” is **not** sufficient – two IDs with one transaction each meet that and still break the fold.

Safe configurations are: an in-order master-side slave, or a single ID outstanding at a time. Nothing in the RTL enforces or detects a violation, so a breach corrupts data silently rather than failing (tracked as CONV-001).

8.9 2.6.9 Resource Utilization

8.9.1 Typical Resources (64→512 UPSIZE, ratio 8, ID=4)

The burst-length FIFO exists only in the converter's UPSIZE mode (see 2.6.5), so the configuration here is upsize — an earlier revision labeled this table 512→64 (DOWNSIZE), a mode in which that FIFO is tied off and does not exist. Hand estimates, not synthesis results, except the FIFO, which is counted from its declarations.

AR skid buffer:	~150 flip-flops
R dnsz data path:	~600 flip-flops, ~50 LUTs (single-buffer slicer)
Burst-length FIFO:	186 flip-flops (16 x 8b length + 16 x 3b start lane + two 5b pointers)
Control logic:	~80 LUTs

Total: ~940 flip-flops, ~130 LUTs (sum of the lines above)

There is no separate “burst tracker” block — 2.6.7 explains RLAST comes from the dnsz's own counter, framed by the FIFO. The R data path is the single-buffer axi_data_dnsz — the only implementation; its ping-pong DUAL_BUFFER variant was removed, see 2.3.8.

8.10 2.6.10 Timing

8.10.1 Latency

Table 2.20: Read Converter Latency

Path	Latency
AR passthrough	1-2 cycles (skid)
First R beat	1 cycle (load buffer)
Subsequent R beats	1 beat/cycle

8.10.2 Throughput

- AR channel: 1 transaction/cycle
- R channel: the downsize accepts its next wide beat during the last narrow beat — 0.992 beats/cycle measured in simple mode (TRACK_BURSTS=1 pays one bubble per burst boundary and measures ~0.93; see 2.3)

8.11 2.6.11 Usage Example

```
axi4_dwidth_converter_rd #(  
    .S_AXI_DATA_WIDTH (128),  
    .M_AXI_DATA_WIDTH (32),  
    .AXI_ID_WIDTH      (8),  
    .AXI_ADDR_WIDTH    (32),  
    .SKID_DEPTH_AR     (2),  
    .SKID_DEPTH_R      (4)  
) u_conv (  
    .aclk                (aclk),  
    .aresetn              (aresetn),  
    // slave side: s_axi_ar*/... (full AXI4 channel set)  
    .s_axi_arvalid        (cpu_arvalid),  
    .s_axi_arready        (cpu_arready),  
    // ...  
    // master side: m_axi_* toward the narrow fabric  
    .m_axi_arvalid        (mem_arvalid),  
    .m_axi_arready        (mem_arready)  
    // ...  
);
```

All channel ports carry the s_axi_/m_axi_ prefix; the full list is the module header. Skid depths are per channel — there is no single SKID_DEPTH.

Next: [Protocol Conversion Overview](#)

9 2.7 AXIL Wide-Alignment Converters

`axil_to_axi4_wide_align_wr` and `axil_to_axi4_wide_align_rd` are drop-in replacements for `axi4_dwidth_converter_wr / _rd` for the one case those converters get wrong: a slave-side master that is effectively AXI4-Lite – every transfer single-beat, narrow, at a sub-row address.

9.1 2.7.1 Why They Exist

The generic converters place data by **beat counter**. `axi_data_upsize` aggregates N narrow beats into one wide beat in arrival order, and `axi_data_dnsize` emits narrow beats the same way. That is correct for a true multi-beat narrow burst.

An AXI4-Lite master never issues one. Every transfer is a single beat with `awlen=0`, and AXI4 requires the data for a narrow single-beat transfer to sit in the byte lanes selected by the address, not in lane 0. With the generic upsize, every such write lands at lane 0 whatever `awaddr` says, and every such read is taken from lane 0 rather than the requested slot.

These two modules place the beat at byte offset $(addr \ \& \ wide_mask) * 8$ instead, which is what the specification requires.

The path this matters on is the bridge generator's: it wraps an AXI4-Lite master with `axi4_slave_wr`, which presents `fub_axi_*` with `awlen=0`, a narrow `awsiz` and `awburst=01`. That is exactly the shape the generic upsize mishandles.

9.2 2.7.2 Parameters

Both modules take the same set. Defaults differ from the generic converters because the intended direction is narrow master to wide slave.

Wide-Alignment Converter Parameters

Parameter	Default	Description
S_AXI_DATA_WIDTH	32	Slave-side (narrow, AXIL-like) data width
M_AXI_DATA_WIDTH	256	Master-side (wide) data width
AXI_ID_WIDTH	8	ID width
AXI_ADDR_WIDTH	32	Address width
AXI_USER_WIDTH	1	User-signal width
SKID_DEPTH_AW	2	Accepted for API parity; unused. Tied into the module's <code>_unused</code> sink
SKID_DEPTH_W	2	Likewise unused (write module)
SKID_DEPTH_B	2	Likewise unused (write module)
SKID_DEPTH_AR	2	Likewise unused (read module)
SKID_DEPTH_R	4	Likewise unused (read module)

The `SKID_DEPTH_*` parameters exist so these modules can be swapped in for the generic converters without editing the instantiation. They buffer nothing – the RTL lists them explicitly in a verilator `lint_off UNUSED` sink rather than leaving the non-use implicit.

Several AW/AR sideband inputs are declared and deliberately unused for the same reason – `awlen`, `awsz`, `awburst`, `awlock`, `awqos`, `awregion` and `wlast` on the write side. A single-beat AXIL transfer carries no useful burst description, so the module ignores it rather than trusting it.

RTL Design Sherpa · Learning Hardware Design Through Practice [GitHub](#) · [Documentation Index](#) · [MIT License](#)

10 3.1 Protocol Conversion Overview

Protocol converters bridge components that speak different bus protocols — the glue in any SoC that mixes AXI4, AXI4-Lite, and APB.

10.1 3.1.1 Available Converters

10.1.1 AXI4 to AXI4-Lite (Protocol Downgrade)

Table 3.1: AXI4 to AXI4-Lite Converters

Module	Function	Test Status
axi4_to_axil4_rd	Read burst decomposition	5/5 configs passing
axi4_to_axil4_wr	Write burst decomposition	5/5 configs passing
axi4_to_axil4	Full bidirectional wrapper	Composed

10.1.2 AXI4-Lite to AXI4 (Protocol Upgrade)

Table 3.2: AXI4-Lite to AXI4 Converters

Module	Function	Test Status
axil4_to_axi4_rd	Read protocol upgrade	7/7 passing
axil4_to_axi4_wr	Write protocol upgrade	7/7 passing
axil4_to_axi4	Full bidirectional wrapper	Composed

10.1.3 Other Protocol Converters


Table 3.3: Other Protocol Converters

Module	Function	Status
axi4_to_apb4_convert	Full AXI4-to-APB bridge	Production
peakrdl_to_cmdrsp	Register interface adapter	Production
uart_axil_bridge	UART to AXI4-Lite	Implemented (rtl/uart_to_axil4/, dv/tests/test_uart_axil_bridge.py)
axil4_to_wb4	AXI4-Lite to Wishbone B4 (pipelined or classic)	Implemented (dv/tests/test_axil4_to_wb4.py, formal on the

Module	Function	Status
		core)
wb4_to_axil4	Wishbone B4 to AXI4-Lite; merges AXI's two response channels back into B4's in-order termination	Implemented (dv/tests/test_wb4_to_axil4.py, formal on the core)
apb4_to_axi4 / apb5_to_axi4	APB4 / APB5 completer in, single-beat AXI4 requester out (the mirror of axi4_to_apb4); APB5 USER mapped onto AXI USER	Implemented (dv/tests/test_apb{4,5}_to_axi4.py; the bridge's APB master ports, BRIDGE-014)
apb_cmdrsp_to_axi4	The requester half of the two above: APB cmd/rsp stream to one AXI4 transaction per transfer	Implemented (covered by the wrappers' tests)
axi4_to_wb4	Full AXI4 completer in, Wishbone B4 requester out: the AXI4-Lite decomposers + axil4_to_wb4	Implemented (dv/tests/test_axi4_to_wb4.py; the bridge's WB4 slave ports, BRIDGE-019)
wb4_to_axi4	Wishbone B4 completer in, single-beat AXI4 requester out: wb4_to_axil4 + the AXI4 promotion	Implemented (dv/tests/test_wb4_to_axi4.py; the bridge's WB4 master ports, BRIDGE-019)

10.2 3.1.2 Protocol Comparison

10.2.1 Figure 3.1: Protocol Feature Comparison



Protocol Comparison

10.2.2 Feature Matrix

Table 3.4: Protocol Feature Comparison

Feature	AXI4	AXI4-Lite	APB
Channels	5 (AW, W, B, AR, R)	5 (simplified)	1 (combined)
Bursts	Up to 256 beats	Single beat only	Single beat
Out-of-order	Yes (ID-based)	No	No
Pipelining	Yes	Optional	No (2-phase)
Data widths	8-1024 bits	32/64 bits	8-32 bits

10.3 3.1.3 Conversion Strategies

10.3.1 AXI4 to AXI4-Lite

Challenge: Decompose multi-beat bursts into sequential single beats.

Strategy: 1. Accept AXI4 burst transaction 2. Issue N single-beat AXI4-Lite transactions 3. Aggregate responses 4. Return combined response to AXI4 master

Complexity: Medium (FSM-based decomposition)

10.3.2 AXI4-Lite to AXI4

Challenge: Add AXI4 burst signals with appropriate defaults.

Strategy: 1. Pass through all AXI4-Lite signals 2. Add default values for burst signals (LEN=0, SIZE= $\lceil \log_2(\text{DATA_WIDTH}/8) \rceil$, BURST=INCR) 3. Add default IDs (configurable)

Complexity: Very low (combinational only)

10.3.3 AXI4 to APB

Challenge: Bridge 5-channel AXI4 to 2-phase APB protocol.

Strategy: 1. Accept AXI4 transaction 2. Execute APB setup phase 3. Execute APB access phase 4. Return AXI4 response

Complexity: High (full protocol FSM)

10.4 3.1.4 Performance Characteristics

Table 3.5: Protocol Converter Performance

Converter	Single-Beat	Burst (N)	Area
axi4_to_axil4	0 cycles	slave-limited, ~N cycles best case (see 3.2.6)	~450 LUTs
axil4_to_axi4	0 cycles	N/A	~110 LUTs
axi4_to_apb4	3-5 cycles	~2N+1 APB cycles sustained (see 3.4.11)	~150 LUTs (hand estimate; matches 3.4.10)

10.4.1 Key Observations

1. **Zero-overhead upgrade:** AXI4-Lite to AXI4 is purely combinational
2. **Burst rate:** AXI4 to AXI4-Lite adds no wait state between beats — decomposed requests stream at whatever rate the AXIL4 slave sustains (a zero-wait-state slave is served one beat per cycle)
3. **APB overhead:** 3-5 cycle minimum per APB transaction

10.5 3.1.5 Use Case Guidelines

10.5.1 When to Use AXI4 to AXI4-Lite

Use when: - CPU/DMA with burst support needs simple peripheral access - Want to simplify peripheral design (no burst handling) - Data widths match (no width conversion needed) - Burst performance is not critical

Avoid when: - High-bandwidth streaming data - Latency-critical paths - Many back-to-back bursts

10.5.2 When to Use AXI4-Lite to AXI4

Use when: - Legacy AXI4-Lite IP connects to AXI4 fabric - Designing simple peripheral for AXI4 system - Want zero-overhead protocol upgrade - Don't need burst capability

Avoid when: - Need actual burst support (use native AXI4) - Width conversion needed (use width converters)

10.5.3 When to Use AXI4 to APB

Use when: - AXI4 masters need APB peripheral access - Building CPU-to-peripheral bridges - Integrating legacy APB devices

Avoid when: - High-performance paths (APB is slow) - Streaming data (APB is sequential)

10.6 3.1.6 Integration Patterns

10.6.1 Pattern 1: CPU to Peripherals

CPU (AXI4) → AXI4-to-AXIL4 → AXIL4 Peripherals
→ AXI4-to-APB → APB Peripherals

10.6.2 Pattern 2: Simple IP in AXI4 Fabric

Simple IP (AXIL4) → AXIL4-to-AXI4 → AXI4 Crossbar

10.6.3 Pattern 3: Mixed System

CPU (AXI4) → AXI4 Crossbar → DDR (AXI4)
→ AXI4-to-AXIL4 → Config Regs (AXIL4)
→ AXI4-to-APB → UART, GPIO (APB)

Next: [AXI4 to AXI4-Lite](#)

RTL Design Sherpa · Learning Hardware Design Through Practice [GitHub](#) · [Documentation Index](#) · [MIT License](#)

11 3.2 AXI4 to AXI4-Lite Converter

The **axi4_to_axil4** converter family decomposes AXI4 burst transactions into sequential AXI4-Lite single-beat transactions.

11.1 3.2.1 Module Organization


```
axi4_to_axil4.sv      # Full bidirectional wrapper
├─ axi4_to_axil4_rd.sv # Read path converter
└─ axi4_to_axil4_wr.sv # Write path converter
```

11.1.1 Design Philosophy

Read and write are separate modules, which buys: - Independent optimization - Selective instantiation (read-only, write-only, or both) - Simpler verification (test paths independently)

11.2 3.2.2 Read Path (axi4_to_axil4_rd)

11.2.1 Figure 3.2: AXI4 to AXI4-Lite Read Path



AXI4 to AXI4 Read

11.2.2 Operation

Single-Beat (ARLEN=0):

Cycle 0: AR accepted (passthrough)
Cycle 1: R returned (passthrough)
Total: 0 extra cycles (pure passthrough)

Multi-Beat (ARLEN=N-1):

Cycle 0: AR[0] issued to AXI4
Cycle 1: R[0] received, AR[1] issued
Cycle 2: R[1] received, AR[2] issued
...
R[N-1] received (RLAST)

AR issue and R return are fully independent: the address FSM streams one AXI4 AR per beat as fast as m_axil_arready accepts, with no reference to R at all — nothing in the AR path looks at m_axil_rvalid. Against a slave with multi-cycle response latency the requests pipeline ahead of the data, and end-to-end duration is set by the slave, not the converter.

Bursts do not overlap each other: AR is held off until the burst in flight has delivered its last beat, because the burst-tracking registers hold one burst at a time.

11.2.3 State Machine

```
typedef enum logic [1:0] {  
    RD_IDLE      = 2'b00, // No burst in flight  
    RD_BURST     = 2'b01, // Issuing the AXI4 reads of a burst  
    RD_LAST_BEAT = 2'b10, // Issuing the burst's final read  
} rd_state_t;
```

A single-beat read (ARLEN == 0) never leaves RD_IDLE: it takes the passthrough leg rather than the decomposition FSM. That is also why the one-outstanding-burst guard cannot rely on the FSM state alone.

11.2.4 Implementation

```
// Read state machine
typedef enum logic [1:0] {
    RD_IDLE      = 2'b00,
    RD_BURST     = 2'b01,
    RD_LAST_BEAT = 2'b10
} rd_state_t;

rd_state_t r_rd_state, w_rd_next_state;


always_ff @(posedge aclk or negedge aresetn) begin
    if (!aresetn) begin
        r_rd_state <= RD_IDLE;
    end else begin
        r_rd_state <= w_rd_next_state;
    end
end

// Read next state logic
always_comb begin
    w_rd_next_state = r_rd_state;
    case (r_rd_state)
        RD_IDLE: begin
            if (s_axi_arvalid && s_axi_arready) begin
                if (s_axi_arlen == 0)
                    w_rd_next_state = RD_IDLE; // Single beat,
stay idle
                else
                    w_rd_next_state = RD_BURST; // Start burst
decomposition
            end
        end
        RD_BURST: begin
            if (m_axil_arvalid && m_axil_arready) begin
                if (r_ar_beat_count == r_ar_len - 1)
                    w_rd_next_state = RD_LAST_BEAT;
            end
        end
        RD_LAST_BEAT: begin
            if (m_axil_arvalid && m_axil_arready)
                w_rd_next_state = RD_IDLE;
        end
        default: w_rd_next_state = RD_IDLE;
    end
end
```

```
endcase
end
```

11.3 3.2.3 Write Path (axi4_to_axil4_wr)

11.3.1 Figure 3.3: AXI4 to AXI4-Lite Write Path



AXI4 to AXIL4 Write

11.3.2 Operation

Single-Beat (AWLEN=0):

Cycle 0: AW+W accepted (passthrough)
Cycle 1: B returned (passthrough)
Total: 0 extra cycles

Multi-Beat (AWLEN=N-1):

Cycle 0: AW[0] + W[0] issued
Cycle 1: B[0] received, AW[1] + W[1] issued
...

B[N-1] received (timing set by the AXIL4 slave)
 B[N-1] received

As on the read path, AW/W issue is independent of B return — the write FSM streams beats as fast as the AXIL4 slave accepts them, and mid-burst B responses are consumed immediately (`m_axil_bready = s_axi_bready || !w_b_all_beats_done`). The single B the master receives is emitted once every decomposed write has responded. Bursts do not overlap each other: the next AW is held off until the burst in flight has its B accepted.

11.3.3 AW/W Synchronization Challenge

AXI4 allows AW and W to arrive in any order: - AW before W - W before AW - Interleaved

11.3.4 Solution: W gated on the AW that owns it

W is gated directly against the state of the AW it belongs to, rather than through pending-arrival flags, so a data beat cannot reach the AXIL4 slave ahead of its own address:

- **Burst capture cycle.** When a burst's AW is being taken into the registers, W is held off for that cycle (`w_burst_capture`). Without it, the first W beat of the next burst slips out while the previous burst is still finishing, and the AXIL4 slave pairs an address from burst N with data from burst N+1.
 - **During a burst.** W passes only once this burst's AW has actually been sent (`r_aw_sent`), or is being sent in this very cycle.
 - **Single beats.** Straight passthrough; AW and W are presented together.
- ```
// s_axi_awready restricts the block to the ACCEPT cycle. Without
it
// a burst AW merely PARKED by the one-outstanding guard
(awvalid=1,
// awready=0) deadlocked the write already in flight (CONV-007).
wire w_burst_capture = !r_aw_active && s_axi_awvalid &&
s_axi_awready &&
(s_axi_awlen > 0);

assign m_axil_wvalid = w_burst_capture ? 1'b0 :
r_aw_active ? (s_axi_wvalid &&
(r_aw_sent ||
(m_axil_awvalid &&
m_axil_awready))) :
s_axi_wvalid;
assign s_axi_wready = w_burst_capture ? 1'b0 :
r_aw_active ? (m_axil_wready &&
(r_aw_sent ||
(m_axil_awvalid &&
m_axil_awready))) :
m_axil_wready;
```

The failure this guards against only appears back to back — the next AW arriving the cycle after WR\_LAST\_BEAT completes. A sequential test with a cooldown between bursts never opens that window, which is why it survived the FUB tests and was caught by a bridge-level probe. Classic bring-up bug, in other words: it hides from every directed test and fires the moment real traffic shows up.

### 11.3.5 Response Aggregation

```
// from the RTL: accumulate per response, reset on AW acceptance
always_ff @(posedge aclk or negedge aresetn) begin
 ...
 if (s_axi_awvalid && s_axi_awready)
 r_b_resp_accum <= 2'b00;
 if (m_axil_bvalid && m_axil_bready)
 if (m_axil_bresp > r_b_resp_accum)
 r_b_resp_accum <= m_axil_bresp;
end

// live-beat fold at emission (see 4.3.3 for why the registered-only
// form was a shipped bug)
assign w_b_resp_worst = (m_axil_bresp > r_b_resp_accum) ? m_axil_bresp
 :
r_b_resp_accum;
assign s_axi_bresp = w_b_resp_worst;
```

## 11.4 3.2.4 Bidirectional Wrapper (axi4\_to\_axil4)

---

### 11.4.1 Composition Pattern

```
module axi4_to_axil4 #(
 parameter int AXI_ID_WIDTH = 8,
 parameter int AXI_ADDR_WIDTH = 32,
 parameter int AXI_DATA_WIDTH = 32,
 parameter int AXI_USER_WIDTH = 1
) (
 // aclk/aresetn + full s_axi*/m_axil* channel set
);

 axi4_to_axil4_rd #(
 .AXI_ID_WIDTH (AXI_ID_WIDTH),
 .AXI_ADDR_WIDTH (AXI_ADDR_WIDTH),
 .AXI_DATA_WIDTH (AXI_DATA_WIDTH),
 .AXI_USER_WIDTH (AXI_USER_WIDTH)
) u_rd_converter (/* read channels */);

 axi4_to_axil4_wr #(
 .AXI_ID_WIDTH (AXI_ID_WIDTH),
 .AXI_ADDR_WIDTH (AXI_ADDR_WIDTH),
 .AXI_DATA_WIDTH (AXI_DATA_WIDTH),
 .AXI_USER_WIDTH (AXI_USER_WIDTH)
) u_wr_converter (/* write channels */);

endmodule
```

## 11.5 3.2.5 Resource Utilization

---

*Table 3.6: AXI4 to AXIL4 Resources*

| Module                      | Registers | LUTs | BRAM |
|-----------------------------|-----------|------|------|
| axi4_to_axil4_r<br>d        | ~120      | ~180 | 0    |
| axi4_to_axil4_<br>wr        | ~150      | ~220 | 0    |
| axi4_to_axil4<br>(combined) | ~270      | ~400 | 0    |

## 11.6 3.2.6 Timing

### 11.6.1 Throughput

*Table 3.7: AXI4 to AXIL4 Throughput*

| Transaction Type    | Behaviour                                                                                       |
|---------------------|-------------------------------------------------------------------------------------------------|
| Single-beat         | Passthrough; no converter-inserted wait state                                                   |
| N-beat burst        | One AXIL4 access per beat; requests stream at the slave's accept rate, independent of responses |
| Back-to-back bursts | Serialized: the next AR/AW waits for the previous burst's last beat                             |

The previous “50% for any burst” figure was not measured and does not follow from the design: the converter adds no wait state between beats, so the achievable rate is whatever the AXIL4 slave sustains. A measured characterization per slave latency has not been done and is not claimed here.

### 11.6.2 Latency

*Table 3.8: AXI4 to AXIL4 Latency*

| Transaction Type | Latency                                                               |
|------------------|-----------------------------------------------------------------------|
| Single-beat      | 0 extra cycles                                                        |
| N-beat burst     | slave-limited; requests stream independently of responses (see 3.2.6) |

## 11.7 3.2.7 Testing

---

**Test Suite:** both directions green — `test_axi4_to_axil4_{wr,rd}.py`, 5 parametrized configurations each, every scenario suite asserted (the counts below are scenario groups within those runs)

*Table 3.9: Test Coverage Summary*

| Test Category     | Tests | Status |
|-------------------|-------|--------|
| Single-beat read  | 4     | Pass   |
| Multi-beat read   | 6     | Pass   |
| Single-beat write | 4     | Pass   |
| Multi-beat write  | 6     | Pass   |
| Mixed traffic     | 8     | Pass   |
| Error injection   | 6     | Pass   |
| Edge cases        | 8     | Pass   |

Next: [AXI4-Lite to AXI4](#)

RTL Design Sherpa · Learning Hardware Design Through Practice [GitHub](#) · [Documentation Index](#) · [MIT License](#)

---

## 12 3.3 AXI4-Lite to AXI4 Converter

The **axil4\_to\_axi4** converter family upgrades AXI4-Lite single-beat transactions to full AXI4 protocol by adding default burst signals.

## 12.1 3.3.1 Module Organization

---

```
axil4_to_axi4.sv # Full bidirectional wrapper
├─ axil4_to_axi4_rd.sv # Read path converter
└─ axil4_to_axi4_wr.sv # Write path converter
```

## 12.2 3.3.2 Design Philosophy

---

**Zero-Overhead Upgrade:** - Purely combinational logic - No state machines or buffers - Adds default values for missing AXI4 signals

**Why This Works:** - AXI4-Lite is a subset of AXI4 - All AXI4-Lite transactions are single-beat - Missing AXI4 signals have well-defined defaults

## 12.3 3.3.3 Signal Mapping

### 12.3.1 Address Channel Signals

Table 3.10: AR Channel Mapping

| AXI4-Lite Signal | AXI4 Signal | Default/Mapping                       |
|------------------|-------------|---------------------------------------|
| ARADDR           | ARADDR      | Passthrough                           |
| ARPROT           | ARPROT      | Passthrough                           |
| ARVALID          | ARVALID     | Passthrough                           |
| ARREADY          | ARREADY     | Passthrough                           |
| -                | ARLEN       | 8'h00 (single beat)                   |
| -                | ARSIZE      | $\$clog2(DATA\_WIDTH/8)$              |
| -                | ARBURST     | 2'b01 (INCR)                          |
| -                | ARLOCK      | 1'b0 (normal)                         |
| -                | ARCACHE     | 4'b0011 (bufferable, modifiable)      |
| -                | ARQOS       | 4'(DEFAULT_QOS), parameter, default 0 |
| -                | ARID        | Configurable default                  |


### 12.3.2 Data Channel Signals

Table 3.11: R Channel Mapping

| AXI4-Lite Signal | AXI4 Signal | Default/Mapping    |
|------------------|-------------|--------------------|
| RDATA            | RDATA       | Passthrough        |
| RRESP            | RRESP       | Passthrough        |
| RVALID           | RVALID      | Passthrough        |
| RREADY           | RREADY      | Passthrough        |
| -                | RLAST       | 1'b1 (always last) |
| -                | RID         | Matches ARID       |

## 12.4 3.3.4 Read Path (axil4\_to\_axi4\_rd)

### 12.4.1 Figure 3.4: AXI4-Lite to AXI4 Read Path



### AXI4 to AXI4 Read

### 12.4.2 Implementation

```

module axil4_to_axi4_rd #(
 // Width Configuration
 parameter int AXI_ID_WIDTH = 8,
 parameter int AXI_ADDR_WIDTH = 32,
 parameter int AXI_DATA_WIDTH = 32,
 parameter int AXI_USER_WIDTH = 1,

 // Default Values for AXI4-only Fields
 parameter int DEFAULT_ID = 0,
 parameter int DEFAULT_REGION = 0,
 parameter int DEFAULT_QOS = 0,

 // Skid Buffer Depths (for timing closure)

```

```

// Calculated Parameters
localparam int STRB_WIDTH = AXI_DATA_WIDTH / 8,
localparam int SIZE_VAL = $clog2(STRB_WIDTH) // ARSIZE for full
width
) (
 // Clock and Reset
 input logic aclk,
 input logic aresetn,

 //=====
 // Slave AXI4-Lite Read Interface (Input - Simplified Protocol)
 //=====

 // Read Address Channel
 input logic [AXI_ADDR_WIDTH-1:0] s_axil_araddr,
 input logic [2:0] s_axil_arprot,
 input logic s_axil_arvalid,
 output logic s_axil_arready,

 // Read Data Channel
 output logic [AXI_DATA_WIDTH-1:0] s_axil_rdata,
 output logic [1:0] s_axil_rresp,
 output logic s_axil_rvalid,
 input logic s_axil_rready,

 //=====
 // Master AXI4 Read Interface (Output - Full Protocol)
 //=====

 // Read Address Channel
 output logic [AXI_ID_WIDTH-1:0] m_axi_arid,
 output logic [AXI_ADDR_WIDTH-1:0] m_axi_araddr,
 output logic [7:0] m_axi_arlen,
 output logic [2:0] m_axi_arsize,
 output logic [1:0] m_axi_arburst,
 output logic m_axi_arlock,
 output logic [3:0] m_axi_arcache,
 output logic [2:0] m_axi_arprot,
 output logic [3:0] m_axi_arqos,

```

```

output logic [3:0] m_axi_arregion,
output logic [AXI_USER_WIDTH-1:0] m_axi_aruser,
output logic m_axi_arvalid,
input logic m_axi_arready,


// Read Data Channel
input logic [AXI_ID_WIDTH-1:0] m_axi_rid,
input logic [AXI_DATA_WIDTH-1:0] m_axi_rdata,
input logic [1:0] m_axi_rresp,
input logic m_axi_rlast,
input logic [AXI_USER_WIDTH-1:0] m_axi_ruser,
input logic m_axi_rvalid,
output logic m_axi_rready
);

```

**Key Points:** - No registers - purely combinational - RLAST from AXI4 is ignored (always 1 for AXI4) - RID from AXI4 is ignored (no ID tracking in AXI4)

## 12.5 3.3.5 Write Path (axil4\_to\_axi4\_wr)

### 12.5.1 Figure 3.5: AXI4-Lite to AXI4 Write Path



### AXI4 to AXI4 Write

### 12.5.2 Implementation

```
module axil4_to_axi4_wr #(
 // Width Configuration
 parameter int AXI_ID_WIDTH = 8,
 parameter int AXI_ADDR_WIDTH = 32,
 parameter int AXI_DATA_WIDTH = 32,
 parameter int AXI_USER_WIDTH = 1,
```

```

// Default Values for AXI4-only Fields
parameter int DEFAULT_ID = 0,
parameter int DEFAULT_REGION = 0,
parameter int DEFAULT_QOS = 0,

// Skid Buffer Depths (for timing closure)

// Calculated Parameters
localparam int STRB_WIDTH = AXI_DATA_WIDTH / 8,
localparam int SIZE_VAL = $clog2(STRB_WIDTH) // AWSIZE for full
width
) (
 // Clock and Reset
 input logic aclk,
 input logic aresetn,

//=====
// Slave AXI4-Lite Write Interface (Input - Simplified Protocol)
//=====

// Write Address Channel
input logic [AXI_ADDR_WIDTH-1:0] s_axil_awaddr,
input logic [2:0] s_axil_awprot,
input logic s_axil_awvalid,
output logic s_axil_awready,

// Write Data Channel
input logic [AXI_DATA_WIDTH-1:0] s_axil_wdata,
input logic [STRB_WIDTH-1:0] s_axil_wstrb,
input logic s_axil_wvalid,
output logic s_axil_wready,

// Write Response Channel
output logic [1:0] s_axil_bresp,
output logic s_axil_bvalid,
input logic s_axil_bready,

//=====

```

```

=====
// Master AXI4 Write Interface (Output - Full Protocol)

//=====

// Write Address Channel
output logic [AXI_ID_WIDTH-1:0] m_axi_awid,
output logic [AXI_ADDR_WIDTH-1:0] m_axi_awaddr,
output logic [7:0] m_axi_awlen,
output logic [2:0] m_axi_awsz,
output logic [1:0] m_axi_awburst,
output logic m_axi_awlock,
output logic [3:0] m_axi_awcache,
output logic [2:0] m_axi_awprot,
output logic [3:0] m_axi_awqos,
output logic [3:0] m_axi_awregion,
output logic [AXI_USER_WIDTH-1:0] m_axi_awuser,
output logic m_axi_awvalid,
input logic m_axi_awready,

// Write Data Channel
output logic [AXI_DATA_WIDTH-1:0] m_axi_wdata,
output logic [STRB_WIDTH-1:0] m_axi_wstrb,
output logic m_axi_wlast,
output logic [AXI_USER_WIDTH-1:0] m_axi_wuser,
output logic m_axi_wvalid,
input logic m_axi_wready,

// Write Response Channel
input logic [AXI_ID_WIDTH-1:0] m_axi_bid,
input logic [1:0] m_axi_bresp,
input logic [AXI_USER_WIDTH-1:0] m_axi_buser,
input logic m_axi_bvalid,
output logic m_axi_bready
);

```

## 12.6 3.3.6 Bidirectional Wrapper

---

```
module axil4_to_axi4 #(
 parameter int AXI_ID_WIDTH = 8,
 parameter int AXI_ADDR_WIDTH = 32,
 parameter int AXI_DATA_WIDTH = 32,
 parameter int AXI_USER_WIDTH = 1,
 parameter int DEFAULT_ID = 0,
 parameter int DEFAULT_REGION = 0,
 parameter int DEFAULT_QOS = 0
) (
 // ... all port declarations
);

axil4_to_axi4_rd #(
 .AXI_ID_WIDTH (AXI_ID_WIDTH),
 .AXI_ADDR_WIDTH (AXI_ADDR_WIDTH),
 .AXI_DATA_WIDTH (AXI_DATA_WIDTH),
 .AXI_USER_WIDTH (AXI_USER_WIDTH),
 .DEFAULT_ID (DEFAULT_ID),
 .DEFAULT_REGION (DEFAULT_REGION),
 .DEFAULT_QOS (DEFAULT_QOS)
) u_rd_converter (/* connections */);

axil4_to_axi4_wr #(
 .AXI_ID_WIDTH (AXI_ID_WIDTH),
 .AXI_ADDR_WIDTH (AXI_ADDR_WIDTH),
 .AXI_DATA_WIDTH (AXI_DATA_WIDTH),
 .AXI_USER_WIDTH (AXI_USER_WIDTH),
 .DEFAULT_ID (DEFAULT_ID),
 .DEFAULT_REGION (DEFAULT_REGION),
 .DEFAULT_QOS (DEFAULT_QOS)
) u_wr_converter (/* connections */);

endmodule
```

## 12.7 3.3.7 Resource Utilization

---

*Table 3.12: AXIL4 to AXI4 Resources*

| Module                      | Registers | LUTs |
|-----------------------------|-----------|------|
| axil4_to_axi4_rd            | 0         | ~50  |
| axil4_to_axi4_wr            | 0         | ~60  |
| axil4_to_axi4<br>(combined) | 0         | ~110 |

**Note:** Zero registers — purely combinational logic.

## 12.8 3.3.8 Timing

---

*Table 3.13: AXIL4 to AXI4 Performance*

| Metric        | Value      |
|---------------|------------|
| Latency       | 0 cycles   |
| Throughput    | 100%       |
| Max frequency | Wire speed |

## 12.9 3.3.9 Testing

---

**Test Suite:** 14 tests passing

*Table 3.14: Test Coverage Summary*

| Test Category           | Tests | Status |
|-------------------------|-------|--------|
| Single-beat read        | 3     | Pass   |
| Single-beat write       | 3     | Pass   |
| Mixed traffic           | 4     | Pass   |
| Default ID verification | 2     | Pass   |
| Edge cases              | 2     | Pass   |

## 12.10 3.3.10 Usage Example

---

```
axil4_to_axi4 #(
 .AXI_DATA_WIDTH (32),
 .AXI_ADDR_WIDTH (32),
 .AXI_ID_WIDTH (8),
 .DEFAULT_ID (8'h05),
 .DEFAULT_QOS (0)
) u_axil2axi (
 .aclk (aclk),
 .aresetn (aresetn),
 // AXI4-Lite slave side: s_axil_*
 .s_axil_arvalid (lite_arvalid),
 .s_axil_arready (lite_arready),
 .s_axil_araddr (lite_araddr),
 // ...
 // AXI4 master side: m_axi_* (ID/len/size driven from parameters;
 // the conversion is combinational, no internal pipelining)
 .m_axi_arvalid (fab_arvalid),
 .m_axi_arready (fab_arready)
 // ...
);
```

---

Next: [AXI4 to APB](#)

RTL Design Sherpa · Learning Hardware Design Through Practice [GitHub](#) · [Documentation Index](#) · [MIT License](#)

---

## 13 3.4 AXI4 to APB Converter

The **axi4\_to\_apb4\_convert** module provides full protocol translation from AXI4 to APB, enabling AXI4 masters to access APB peripherals.

## 13.1 3.4.1 Overview

---


AXI4 and APB differ on just about every axis:

*Table 3.15: AXI4 vs APB Comparison*

| Aspect        | AXI4                | APB                     |
|---------------|---------------------|-------------------------|
| Channels      | 5 (AW, W, B, AR, R) | 1 (combined)            |
| Phases        | Pipelined           | 2-phase (setup, access) |
| Bursts        | Up to 256 beats     | Single transfer         |
| Address width | Up to 64 bits       | Typically 32 bits       |
| Data width    | 8-1024 bits         | 8-32 bits               |

## 13.2 3.4.2 Block Diagram

13.2.1 Figure 3.6: AXI4 to APB Converter



*AXI4 to APB*

## 13.3 3.4.3 Interface Specification

### 13.3.1 Parameters

**axi4\_to\_apb4\_convert** — the conversion core:

*Table 3.16: AXI4 to APB Parameters*

| Parameter          | Type | Default | Description                                                |
|--------------------|------|---------|------------------------------------------------------------|
| AXI_ADDR_WIDT<br>H | int  | 32      | AXI4 address width                                         |
| AXI_DATA_WIDT<br>H | int  | 32      | AXI4 data width                                            |
| AXI_ID_WIDTH       | int  | 8       | AXI4 ID width                                              |
| AXI_USER_WIDTH     | int  | 1       | AXI4 user-signal width                                     |
| APB_ADDR_WIDT<br>H | int  | 32      | APB address width                                          |
| APB_DATA_WIDT<br>H | int  | 32      | APB data width                                             |
| SIDE_DEPTH         | int  | 6       | Side-channel FIFO depth (ID/user carried past the APB hop) |

**axi4\_to\_apb4\_shim** — the wrapper integrators actually instantiate. It adds the channel FIFOs around the core, and is what the bridge generator emits:

*Table 3.17: AXI4 to APB Shim Parameters*

| Parameter              | Type | Default | Description                           |
|------------------------|------|---------|---------------------------------------|
| DEPTH_AW /<br>DEPTH_AR | int  | 2       | Write/read address channel FIFO depth |
| DEPTH_W /<br>DEPTH_R   | int  | 4       | Write/read data channel FIFO depth    |
| DEPTH_B                | int  | 2       | Write response channel FIFO           |

| Parameter        | Type | Default  | Description                                                                                                   |
|------------------|------|----------|---------------------------------------------------------------------------------------------------------------|
|                  |      |          | depth                                                                                                         |
| SIDE_DEPTH       | int  | 4        | Side-channel FIFO depth                                                                                       |
| APB_CMD_DEPTH    | int  | 4        | APB command FIFO depth                                                                                        |
| APB_RSP_DEPTH    | int  | 4        | APB response FIFO depth                                                                                       |
| USE_JOHNSON      | int  | 0        | CDC-FIFO pointer encoding: 0 = Gray (power-of-2 depth only), 1 = Johnson                                      |
| USE_2_PHASE_CD C | bit  | 1        | <b>Deprecated and ignored.</b> Declared for compatibility and never referenced, so setting it changes nothing |
| AXI_/APB_ widths | int  | as above | Same width parameters as the core                                                                             |

The shim's defaults differ from the core's SIDE\_DEPTH (4 vs 6); it sets its own rather than inheriting.

### 13.3.2 Ports

```
// The conversion CORE. It does not sit on the AXI4 channels directly:
// the shim packs each channel into a packet plus a count and hands
those
// in, which is why the ports below are r_s_axi*_pkt rather than the
raw
// AXI signals. Integrators instantiate axi4_to_apb4_shim instead; its
// parameters are in the shim table above.
```

```
module axi4_to_apb4_convert #(
 parameter int SIDE_DEPTH = 6,
 parameter int AXI_ID_WIDTH = 8,
```

```

parameter int AXI_ADDR_WIDTH = 32,
parameter int AXI_DATA_WIDTH = 32,
parameter int AXI_USER_WIDTH = 1,
parameter int APB_ADDR_WIDTH = 32,
parameter int APB_DATA_WIDTH = 32,
parameter int AXI_WSTRB_WIDTH = AXI_DATA_WIDTH / 8,
parameter int APB_WSTRB_WIDTH = APB_DATA_WIDTH / 8,
// short and calculated params
parameter int AW = AXI_ADDR_WIDTH,
parameter int DW = AXI_DATA_WIDTH,
parameter int IW = AXI_ID_WIDTH,
parameter int UW = AXI_USER_WIDTH,
parameter int SW = AXI_DATA_WIDTH / 8,
parameter int APBAW = APB_ADDR_WIDTH,
parameter int APBDW = APB_DATA_WIDTH,
parameter int APBSW = APB_DATA_WIDTH / 8,
parameter int AXI2APBRATIO = DW / APBDW,
parameter int ASize = IW + AW + 8 + 3 + 2 + 1 + 4 + 3
+ 4 + 4 + UW,
parameter int WSize = DW + SW + 1 + UW,
parameter int BSize = IW + 2 + UW,
parameter int ARSize = IW + AW + 8 + 3 + 2 + 1 + 4 + 3
+ 4 + 4 + UW,
parameter int RSize = IW + DW + 2 + 1 + UW,
parameter int APBCmdWidth = APBAW + APBDW + APBSW + 3 + 1 +
1 + 1,
parameter int APBRspWidth = APBDW + 1 + 1 + 1,
parameter int SideSize = 1 + IW + 1 + UW
) (
// Clock and Reset
input logic aclk,
input logic aresetn,

// Inputs from axi_slave_stub
input logic [ASize-1:0] r_s_axi_aw_pkt,
input logic [3:0] r_s_axi_aw_count,
input logic r_s_axi_awvalid,
output logic w_s_axi_awready,

input logic [WSize-1:0] r_s_axi_w_pkt,
input logic r_s_axi_wvalid,
output logic w_s_axi_wready,

output logic [BSize-1:0] r_s_axi_b_pkt,
output logic w_s_axi_bvalid,
input logic r_s_axi_bready,

```

```

input logic [ARSize-1:0] r_s_axi_ar_pkt,
input logic [3:0] r_s_axi_ar_count,
input logic r_s_axi_arvalid,
output logic w_s_axi_arready,

output logic [RSize-1:0] r_s_axi_r_pkt,
output logic w_s_axi_rvalid,
input logic r_s_axi_rready,

// APB Master Interface
output logic w_cmd_valid,
input logic r_cmd_ready,
output logic [APBCmdWidth-1:0] r_cmd_data,

input logic r_rsp_valid,
output logic w_rsp_ready,
input logic [APBRspWidth-1:0] r_rsp_data
);

```

## 13.4 3.4.4 Clock Domains

---

The shim is a **two-clock** block. `aclk` runs the AXI side and `pclk` the APB side, with independent resets (`aresetn`, `presetn`). Commands cross `aclk`→`pclk` and responses `pclk`→`aclk` through gray-pointer asynchronous FIFOs (`gaxi_fifo_async`).

The reset behaviour is the part worth knowing. Each domain resets its own pointer and its crossed copy of the remote pointer from its LOCAL reset, so resetting one side alone leaves that side self-consistent — both pointers at 0, meaning empty. Because the pointers are absolute positions rather than toggle parity, an independent reset of one side cannot fabricate or swallow a transfer. The earlier 2-phase handshake could, and the failure was permanent rather than transient: the response stream ends up offset by one and every read returns the previous read's data.


CDC FIFO depths follow the APB cmd/rsp queue depths, floored at 4:

```
localparam int CDC_CMD_DEPTH = (APB_CMD_DEPTH < 4) ? 4 :
APB_CMD_DEPTH;
localparam int CDC_RSP_DEPTH = (APB_RSP_DEPTH < 4) ? 4 :
APB_RSP_DEPTH;
```

A power-of-2 depth is preferred: the default `USE_JOHNSON=0` Gray encoding requires it (see the parameter table above).

## 13.5 3.4.5 State Machine

13.5.1 Figure 3.7: AXI4 to APB FSM



*AXI4 to APB FSM*

## 13.5.2 States

The convert core runs TWO state machines, and neither of them touches PSEL/PENABLE — the APB setup/access phases belong to `apb4_master`, on the far side of the shim's CDC. The core's job is packetizing.

**Command FSM** — walks the accepted AXI burst and emits one APB command packet per APB-width beat into the cmd stream:

```
typedef enum logic [2:0] {
 IDLE = 3'b001, // wait for an accepted AW+W pair or AR
 READ = 3'b010, // emit one read command per beat
 WRITE = 3'b100 // emit one write command per beat
} apb_state_t;
```

Writes are preferred over reads when both are pending. Each command carries first/last flags; last is set on the final beat, which is also when the AXI address channel is released (`awready/arready`). A data pointer sub-divides each AXI beat into `AXI2APBRATIO` APB beats when the widths differ, and the next address comes from the shared `axi_gen_addr` (`INCR/FIXED/WRAP` per the AXI burst type).

**Response FSM** — assembles returning rsp packets into AXI responses:

```
typedef enum logic [1:0] {
 RSP_IDLE = 2'b01, // no response in flight
 RSP_ACTIVE = 2'b10 // consuming rsp packets for one transaction
} rsp_state_t;
```

Read data re-packs APB-width beats into AXI-width R beats; write responses collapse into the single B, worst response winning. The first flag in the command stream is what arms this FSM — stamping it from transaction progress rather than the previous command state is the fix for a hang recorded in the RTL history (a `first=0` first command froze the response FSM in `RSP_IDLE`).

## 13.6 3.4.6 Burst Handling

---

### 13.6.1 Burst Decomposition

AXI4 bursts are decomposed into sequential APB transfers:

AXI4: AWADDR=0x1000, AWLEN=3 (4 beats)

APB sequence:

Transfer 0: PADDR=0x1000

Transfer 1: PADDR=0x1004

Transfer 2: PADDR=0x1008

Transfer 3: PADDR=0x100C

### 13.6.2 Address Calculation

The per-beat address comes from the shared `axi_gen_addr` block, keyed on the AXI burst type (INCR/FIXED/WRAP) and size — the convert core does not roll its own increment, and there is no APB-phase state here (PREADY pacing happens in `apb4_master` beyond the CDC; see 3.4.5).

## 13.7 3.4.7 Address Width Adaptation

---

### 13.7.1 64-bit to 32-bit Conversion

*// Truncate upper address bits*

**assign** paddr = s\_awaddr[APB\_ADDR\_WIDTH-1:0];

*// Optional: Check for out-of-range access*

**wire** w\_addr\_oor = |s\_awaddr[AXI\_ADDR\_WIDTH-1:APB\_ADDR\_WIDTH];

## 13.8 3.4.8 Error Response Mapping

---

*Table 3.18: Error Mapping*

| APB Signal  | AXI4 Response  |
|-------------|----------------|
| PSLVERR = 0 | OKAY (2'b00)   |
| PSLVERR = 1 | SLVERR (2'b10) |

### 13.8.1 Error Aggregation

Each `rsp` packet carries the APB `pslverr` for its beat. The response FSM accumulates across the transaction's packets and maps any error to SLVERR on the AXI side — B for writes, per-beat RRESP for reads (the error lands on the beats it belongs to, not smeared across the burst). There is no APB-phase sampling here; `pslverr` is captured by `apb4_master` and travels back in the packet.

## 13.9 3.4.9 Implementation

---

The implementation is the two FSMs of 3.4.5 plus the packet plumbing: the shim skid-buffers each AXI channel, packs AW/W/AR into per-channel packets for the core, and carries the core's cmd/rsp streams through gray-pointer async FIFOs to `apb4_master`, which owns the actual APB setup/access phases. The full source is `projects/components/converters/rtl/axi4_to_apb4_convert.sv` (core) and `axi4_to_apb4_shim.sv` (integration); their headers document the packet formats. No separate simplified listing is maintained here — an earlier one drifted into describing states the RTL never had.

## 13.10      3.4.10 Resource Utilization

---

|                 |                    |
|-----------------|--------------------|
| State machine:  | ~50 LUTs, ~20 regs |
| Address logic:  | ~30 LUTs, ~40 regs |
| Data buffering: | ~10 LUTs, ~70 regs |
| Control:        | ~60 LUTs, ~20 regs |

Total: ~150 LUTs, ~150 regs

## 13.11 3.4.11 Timing

### 13.11.1 Timing Analysis

*Table 3.19: APB Converter Timing*

| Operation          | Cycles                                                                    |
|--------------------|---------------------------------------------------------------------------|
| Single write       | 3-4 APB-side + ~2 pclk in + ~2 aclk out for the CDC crossings (see 3.4.4) |
| Single read        | 3-4 (setup + access + R)                                                  |
| N-beat write burst | ~2N+1 APB-side + the same two CDC crossings, paid once per command run    |
| N-beat read burst  | ~2N+1 APB-side                                                            |

The 2N+1 shape falls out of the FSM: the first transfer pays IDLE→SETUP→ACCESS (3 pclk), but with commands queued in the skid the ACCESS state hands straight back to SETUP — every subsequent transfer is SETUP→ACCESS, 2 pclk, PREADY permitting.

### 13.11.2 Throughput

**Best case (sustained):** 1 transfer per 2 pclk cycles once the first transfer's extra IDLE cycle is paid

**With slow PREADY:** Additional cycles per transfer

## 13.12 3.4.12 Usage Example

---

Integrators instantiate the SHIM (the core's cmd/rsp packet ports are internal plumbing):

```
axi4_to_apb4_shim #(
 .AXI_ID_WIDTH (8),
 .AXI_ADDR_WIDTH (32),
 .AXI_DATA_WIDTH (32),
 .APB_ADDR_WIDTH (32),
 .APB_DATA_WIDTH (32),
 .APB_CMD_DEPTH (4),
 .APB_RSP_DEPTH (4)
) u_axi2apb (
 // AXI side clock/reset
 .aclk (aclk),
 .aresetn (aresetn),
 // APB side clock/reset -- this is a TWO-CLOCK block (see 3.4.4)
 .pclk (pclk),
 .presetn (presetn),

 // AXI4 slave interface: s_axi_aw*/w*/b*/ar*/r* (full signal list
in // the module header)
 .s_axi_awvalid (cpu_awvalid),
 .s_axi_awready (cpu_awready),
 // ...

 // APB master interface
 .m_apb_PSEL (periph_psel),
 .m_apb_PENABLE (periph_penable),
 .m_apb_PADDR (periph_paddr),
 .m_apb_PWRITE (periph_pwrite),
 .m_apb_PWDATA (periph_pwdata),
 .m_apb_PSTRB (periph_pstrb),
 .m_apb_PPROT (periph_pprot),
 .m_apb_PRDATA (periph_prdata),
 .m_apb_PSLVERR (periph_pslverr),
 .m_apb_PREADY (periph_pready)
);
```

---

**Next:** [AXI4 to APB5 Shim](#)



## 14 AXI4 to APB5 Shim

**Module:** axi4\_to\_apb5\_shim.sv **Filelist:** rtl/filelists/axi4\_to\_apb5\_shim.f (-f's the APB4 shim's closure)

## 14.1 Overview

APB5 changes less than the version bump suggests. The transfer protocol is pure APB4 — PSEL/PENABLE phases, PREADY, PSLVERR, all unchanged — and everything new lives in the sideband: requester-driven user signals on one side, completer-driven wakeup/user signals on the other. So this block is a thin **pin-superset wrapper** over `axi4_to_apb4_shim`. The entire protocol engine — CDC, cmd/rsp FIFOs, burst decomposition, response assembly — is the APB4 shim instantiated unchanged, with all 58 ports and 18 parameters forwarded 1:1.

The APB5 additions on the requester surface:

| Signal                                         | Dir | Handling                                                                                 |
|------------------------------------------------|-----|------------------------------------------------------------------------------------------|
| <code>m_apb_PAUSER[APB_AUSER_WIDTH-1:0]</code> | out | tied '0 — nothing upstream sources it (AXI USER bits do not map onto APB user semantics) |
| <code>m_apb_PWUSER[APB_WUSER_WIDTH-1:0]</code> | out | tied '0                                                                                  |
| <code>m_apb_PWAKEUP</code>                     | in  | accepted and terminated                                                                  |
| <code>m_apb_PRUSER[APB_RUSER_WIDTH-1:0]</code> | in  | accepted and terminated                                                                  |
| <code>m_apb_PBUSER[APB_BUSER_WIDTH-1:0]</code> | in  | accepted and terminated                                                                  |

One thing to get right before you wire this up. User-signal widths default to 1 here, but `rtl/amba/apb5/apb5_slave.sv` defaults its own to 4, so the two do NOT line up out of the box — set `APB_{A,W,R,B}USER_WIDTH` to match whatever completer you attach. The signal set is otherwise pin-for-pin with that slave, including the repo's convention that PWAKEUP rides completer→requester.

## 14.2 Parameters

---

The width and depth parameters are the same as the APB4 shim's – see [Table 3.17](#) – with the APB5 user-signal widths (APB\_AUSER\_WIDTH, APB\_WUSER\_WIDTH, APB\_RUSER\_WIDTH, APB\_BUSER\_WIDTH) added for the sideband above.

One parameter is inert, and I'd rather call it out here than watch someone wire a knob to it:

### *AXI4 to APB5 Shim – Inert Parameter*

| Parameter       | Type | Default | Description                                                                                                                                                                                                 |
|-----------------|------|---------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| USE_2_PHASE_CDC | bit  | 1       | <b>Deprecated and ignored.</b> This shim wraps axi4_to_apb4_shim and forwards the parameter to it; that module declares it // deprecated, ignored and never references it, so it has no effect at any level |

## 14.3 Design Notes

---

Deriving the wrapper from the APB4 shim's actual port surface — rather than reimplementing the conversion — means the APB4 engine's fixes and its dependency closure are inherited automatically. The closure filelist simply - f's the APB4 shim's. When a future consumer needs real PAUSER/PWUSER sourcing or PWAKEUP-gated clocking, those grow here without touching the conversion core.

## 14.4 Related Modules

---

The bridge generator instantiates this shim for `protocol = "apb5"` slaves (BRIDGE-002 A5-3c) through the same component path as the APB4 shim — the `Axi4ToApbShim` component takes `protocol='apb5'` and wires the five extra pairs. Verified by `projects/components/bridge/dv/tests/test_bridge_1x2_rw_apb5.py` (APB4 BFM legally drives the port: same transfer protocol).

## 14.5 Navigation

---

Next: [AXI4 to AXI5-Lite](#)

[RTL Design Sherpa](#) · [Learning Hardware Design Through Practice](#) [GitHub](#) · [Documentation Index](#) · [MIT License](#)

---

## 15 AXI4 to AXI5-Lite

**Modules:** `axi4_to_axil5.sv`, `axi4_to_axil5_wr.sv`, `axi4_to_axil5_rd.sv` **Filelists:**  
`rtl/filelists/axi4_to_axil5{,_wr,_rd}.f` (each -f's the AXI4-Lite converter's closure)

## 15.1 Overview

---

AXI5-Lite changes less than the version bump suggests, and in exactly the way APB5 does: the transfer protocol is pure AXI4-Lite – single-beat, no burst, no ID, same handshakes and response codes – and everything new is sideband. So these blocks are thin wrappers over `axi4_to_axil4_{wr,rd}`. Burst decomposition, address incrementing, response folding and the one-outstanding guards are that module, instantiated unchanged. What lives here is the sideband, and one timing decision it forces.

`axi4_to_axil5` pairs the two halves exactly as `axi4_to_axil4` pairs its own, and contains nothing but the two instantiations.

The AXI4 slave surface is the **full** AXI4 interface, not a subset. A master upstream of this converter may be feeding a fabric whose other branches reach AXI4 slaves, so nothing is dropped from the boundary on the grounds that this particular downstream path cannot use it.

## 15.2 Where each signal gets its value

Every AXI5-Lite signal is in one of three groups. This is the whole content of the block, so it is worth reading as a table rather than as prose.

### *AXI4 to AXI5-Lite – Sideband Disposition*

| Group                           | Signals                                                                              | Handling                                                            |
|---------------------------------|--------------------------------------------------------------------------------------|---------------------------------------------------------------------|
| <b>FORWARDED</b>                | awlock/arlock,<br>awuser/aruser, wuser                                               | carries the AXI4 value,<br>gated by ENABLE_LOCK /<br>ENABLE_USER    |
| <b>FORWARDED<br/>(response)</b> | buser -> s_axi_buser,<br>ruser -> s_axi_ruser                                        | returned onto the AXI4<br>response channel, gated<br>by ENABLE_USER |
| <b>TIED</b>                     | awloop, awmpam, awmecid,<br>awnsaid, awtrace,<br>wpoison (and the AR<br>equivalents) | driven '0; AXI4 has no<br>source for them                           |
| <b>TERMINATED</b>               | bloop, btrace, rloop,<br>rtrace, rpoison                                             | completer-driven,<br>nothing on the AXI4 side<br>to return them to  |

Two consequences of that table are worth stating outright.

**The tied group has no ENABLE\_ parameter, and that is deliberate.** A knob that cannot change the design's behaviour is worse than no knob: a reader sets it and believes something happened. ENABLE\_LOCK and ENABLE\_USER exist because they gate a real choice; ENABLE\_TRACE, ENABLE\_MPAM and the rest do not exist at all. The width parameters (USER\_WIDTH, LOOP\_WIDTH, MPAM\_WIDTH, MECID\_WIDTH, NSAID\_WIDTH) do remain, because they set the port shape.

**The tied group's ports still exist.** An AXI5-Lite boundary whose shape changes with a config knob cannot be wired to a fixed external completer, so the ports are always present and always driven – zero rather than floating.

rpoison deserves its own note. AXI4 has no poison bit, so a poisoned read beat arriving from an AXI5-Lite completer cannot be signalled to the AXI4 master through any protocol field. It is terminated rather than folded into RRESP, because turning poison into SLVERR would invent an error the completer never reported. A design that must observe poison needs an AXI5 master port, not an AXI4 one.

## 15.3 The AW/AR sideband is held, not passed through

---

This is the one place the wrapper is not trivial, and it is worth understanding before modifying it.

The core **decomposes**: one AXI4 AW handshake becomes N AXI5-Lite AW handshakes, one per beat. And axi4\_to\_axil4\_wr drops s\_axi\_awready as soon as it accepts the AW (s\_axi\_awready = !r\_aw\_active && !r\_wr\_outstanding && ...). So from the cycle after acceptance, an AXI4 master is free to present the **next** transaction's AW signals – new address, new USER, new LOCK – while beats 2..N of the current burst are still going out.

A combinational passthrough of the sideband therefore stamps those later beats with the wrong values. The core has the identical problem with the address and solves it with a capture-and-hold mux:

```
assign m_axil_awaddr = r_aw_active ? r_aw_addr : s_axi_awaddr;
```

r\_aw\_active is registered, so it is still 0 during the accept cycle itself and 1 for every beat after. The wrapper cannot see that signal, so it reconstructs the same selector from outside:

```
wire w_aw_accept = s_axi_awvalid && s_axi_awready;
// ... captured into r_held_aw{lock,user} on w_aw_accept ...
wire [AXI_USER_WIDTH-1:0] w_awuser_sel = w_aw_accept ? s_axi_awuser
 : r_held_awuser;
```

Accept cycle takes the live value; every beat after takes the captured one – identical to the core's mux, so the sideband travels with the address it was issued against.

The first version of this RTL passed it through combinationaly. The failure is not subtle once you can see it: with two overlapping writes, burst A's beats carry burst B's USER **from beat 0**, because B's AW is already on the bus. It is invisible to sequential traffic, which is why the test that catches it (test\_axi4\_to\_axil5.py, full level) issues two writes concurrently rather than one after another.

W beats need no hold. The core assigns m\_axil\_wdata = s\_axi\_wdata directly – W is 1:1 with the AXI4 side – so wuser is a plain passthrough.

## 15.4 Reset style

---

These modules use `always_ff @(posedge aclk or negedge aresetn)` rather than the components-area `ALWAYS_FF_RST` macro, matching the AXI4-Lite converters they wrap. Using the macro here while the core stays manual would make one compiled design half synchronous-reset and half asynchronous – Verilator says so with `SYNCASYNENET`. Converting the whole area is tracked as CONV-009, with what it actually costs recorded there.

## 15.5 Verification

---

`projects/components/converters/dv/tests/test_axi4_to_axil5.py` covers only what the wrapper adds; the transfer protocol is already covered by the AXI4-Lite suites, and duplicating it would test code this module does not contain. Eight configurations across three levels check that forwarded values arrive and stay with their burst, that tied signals read 0 (never X) at every handshake, that the response sideband returns on the AXI4 channels, and that an `ENABLE_*=0` build produces zeros from non-zero AXI4 traffic. The overlapping-burst case was confirmed RED against the unfixed RTL and GREEN after.

## 15.6 Related Modules

---

The bridge generator instantiates these converters for `protocol = "axil5"` slaves through the same component path as the AXI4-Lite ones – the `Axi4ToAxi5Shim` component takes `protocol='axil5'`. Which sideband groups are live comes from the port's `axi5_features`, which on an `axil5` port accepts only `user` and `exclusive`: naming a tied group there would suggest it changes something, so the validator rejects it rather than ignoring it. Verified by `projects/components/bridge/dv/tests/test_bridge_1x2_rw_axil5.py`.

## 15.7 Navigation

---

Next: [PeakRDL Adapter](#)

[RTL Design Sherpa](#) · [Learning Hardware Design Through Practice](#) [GitHub](#) · [Documentation Index](#) · [MIT License](#)

---

## 16 3.5 PeakRDL Adapter

The **peakrdl\_to\_cmdrsp** module drives a PeakRDL-generated register block from a command/response stream. Data flows cmd/rsp -> PeakRDL: cmd\_\* are inputs and regblk\_\* request signals are outputs, with the register block's acks and read data coming back in. Fair warning — the name reads in the opposite order to the dataflow. The ports below are authoritative.

## 16.1 Overview

---

PeakRDL generates register blocks with a selectable cpuif, and this adapter mates with the **passthrough** cpuif (regblk\_req / req\_is\_wr / wr\_biten / stall / ack / err — NOT the APB cpuif's PSEL/PENABLE pins, so if you came here looking for PSEL, back up). It:

1. Decouples the register interface from the implementation
2. Provides a clean handshake protocol
3. Enables pipelined register access
4. Supports custom control logic integration

## 16.2 Parameters

---

*Table 3.20: PeakRDL Adapter Parameters*

| Parameter  | Type | Default | Description                                |
|------------|------|---------|--------------------------------------------|
| ADDR_WIDTH | int  | 12      | Address width<br>(register-block<br>sized) |
| DATA_WIDTH | int  | 32      | Data width                                 |

## 16.3 Ports

---

The full port list, straight from the RTL:

```
module peakrdl_to_cmdrsp #(
 parameter int ADDR_WIDTH = 12, // Address width for cmd/rsp
 interface
 parameter int DATA_WIDTH = 32 // Must match PeakRDL generation
 (typically 32)
) (
 // Clock and Reset
 input logic aclk,
 input logic aresetn,

 //
 =====
 // CMD/RSP Interface (rtldesignsherpa standard)
 //
 =====
 // Command Channel
 input logic cmd_valid,
 output logic cmd_ready,
 input logic cmd_pwrite, // 1=write,
0=read
 input logic [ADDR_WIDTH-1:0] cmd_paddr, // Byte
address
 input logic [DATA_WIDTH-1:0] cmd_pwdata, // Write data
 input logic [DATA_WIDTH/8-1:0] cmd_pstrb, // Byte
strokes

 // Response Channel
 output logic rsp_valid,
 input logic rsp_ready,
 output logic [DATA_WIDTH-1:0] rsp_prdata, // Read data
 output logic rsp_pslverr, // Error flag

 //
 =====
 // PeakRDL Passthrough Interface
 //
 =====
 output logic regblk_req, // Request --
```

```

HELD until ack
 output logic
 output logic [ADDR_WIDTH-1:0]
 output logic [DATA_WIDTH-1:0]
 output logic [DATA_WIDTH-1:0]
enables
 input logic
stall
 input logic
 input logic
acknowledge
 input logic
 input logic [DATA_WIDTH-1:0]
 input logic
acknowledge
 input logic
);

regblk_req_is_wr, // Write flag
regblk_addr, // Address
regblk_wr_data, // Write data
regblk_wr_biten, // Write bit

regblk_req_stall_wr, // Write

regblk_req_stall_rd, // Read stall
regblk_rd_ack, // Read

regblk_rd_err, // Read error
regblk_rd_data, // Read data
regblk_wr_ack, // Write

regblk_wr_err // Write error

```

## 16.4 Functional Description

Direction: cmd/rsp -> PeakRDL. cmd\_\* arrive from upstream (e.g. the APB shim's packet stream); the adapter drives the register block's regblk\_\* request signals and returns its acks as rsp packets.

### 16.4.1 Write

Cycle 0: cmd\_valid, cmd\_pwrite=1 accepted (cmd\_ready high in CMD\_IDLE);  
regblk\_req driven the same cycle, strobes converted to bit enables

Cycle N: regblk\_wr\_ack -> rsp queued (RSP\_VALID), pslverr from regblk\_wr\_err  
(if regblk\_req\_stall\_wr: state parks in CMD\_STALLED and retries when the stall clears)

Cycle M: rsp\_valid && rsp\_ready -> back to idle

`regblk\_req` is HELD asserted from acceptance through the ack cycle (`(cmd_state == CMD_WAIT_ACK) || (CMD_IDLE && cmd_valid)``) -- it is not a one-cycle strobe. The generated passthrough regblock samples the request level and needs the hold; a consumer that starts a new access every cycle req is high would double-fire non-idempotent registers (TASK-064 records the earlier strobe mis-documentation).

### 16.4.2 Read

Same shape via regblk\_rd\_ack/regblk\_rd\_data/regblk\_rd\_err, with rsp\_prdata carried in the response.

### 16.4.33.5 Implementation

Two small FSMs do all the work, from the RTL:

```
typedef enum logic [1:0] {
 CMD_IDLE = 2'b00, // ready to accept a command
 CMD_WAIT_ACK = 2'b01, // register block has the request
 CMD_STALLED = 2'b10 // req_stall_* asserted; retry
} cmd_state_t;
```

```
typedef enum logic {
 RSP_IDLE = 1'b0,
 RSP_VALID = 1'b1 // response held until rsp_ready
} rsp_state_t;
```

The command is registered on acceptance so the request can be replayed out of CMD\_STALLED; in CMD\_IDLE the request muxes straight from the live cmd inputs, so an unstalled single-cycle ack

costs no extra latency. `cmd_ready = (cmd_state == CMD_IDLE)` gives one outstanding command, which is what a register block wants.

#### 16.4.43.5.6 Resource Utilization


FSM state: 3 regs (`cmd_state` 2b + `rsp_state` 1b)  
Command capture: 77 regs (`pwrite` 1 + `paddr` 12 + `pdata` 32 +  
`wr_biten` 32)  
Response capture: 33 regs (`prdata` 32 + `pslverr` 1)  
Request replay muxes: ~77 LUTs (four 2:1 selects, 1+12+32+32 bits --  
IDLE-vs-registered request presentation)  
FSM + control decode: ~25 LUTs

Total: ~113 regs, ~100 LUTs (counted from the declarations at the documented defaults `ADDR_WIDTH=12`, `DATA_WIDTH=32`; all of it is live state -- the command capture replays the request out of `CMD_STALLED/CMD_WAIT_ACK`, and the response capture holds the `rsp` packet)

## 16.5 Waveforms

---

### 16.5.1 Figure 3.8: PeakRDL Adapter



## PeakRDL Adapter

## 16.6 Usage Example

---

### 16.6.13.5.7 Use Cases

The adapter sits between a command/response packet stream and a PeakRDL-generated register block:

```
peakrdl_to_cmdrsp #(
 .ADDR_WIDTH (12),
 .DATA_WIDTH (32)
) u_adapter (
 .aclk (aclk),
 .aresetn (aresetn),

 // command/response stream (e.g. from apb4_slave_cdc's unpacker)
 .cmd_valid (cmd_valid),
 .cmd_ready (cmd_ready),
 .cmd_pwrite (cmd_pwrite),
 .cmd_paddr (cmd_paddr),
 .cmd_pwdata (cmd_pwdata),
 .cmd_pstrb (cmd_pstrb),
 .rsp_valid (rsp_valid),
 .rsp_ready (rsp_ready),
 .rsp_prdata (rsp_prdata),
 .rsp_pslverr (rsp_pslverr),

 // PeakRDL register block hookup
 .regblk_req (hwif_req),
 .regblk_req_is_wr (hwif_req_is_wr),
 .regblk_addr (hwif_addr),
 .regblk_wr_data (hwif_wr_data),
 .regblk_wr_biten (hwif_wr_biten),
 .regblk_req_stall_wr (hwif_req_stall_wr),
 .regblk_req_stall_rd (hwif_req_stall_rd),
 .regblk_rd_ack (hwif_rd_ack),
 .regblk_rd_err (hwif_rd_err),
 .regblk_rd_data (hwif_rd_data),
 .regblk_wr_ack (hwif_wr_ack),
 .regblk_wr_err (hwif_wr_err)
);
```

Typical deployments: CSR blocks behind the APB shim (see 3.4), or any fabric whose endpoint speaks the cmd/rsp packet convention.

### 16.6.23.5.8 Integration Example

The flow a real deployment uses (CSRs behind the APB shim):

```

AXI4 master → axi4_to_apb4_shim → APB bus → apb4_to_peakrdl
 └─ apb4_slave_cdc (APB
-> cmd/rsp)
 └─ peakrdl_to_cmdrsp
 │
 └─ regblk_* request/ack
 ↓
 PeakRDL-generated register
block

```

The shim's own cmd/rsp stream is internal — its external pins are APB. The block that turns a bus back into the cmd/rsp surface this adapter consumes is `apb4_to_peakrdl` (which wraps `apb4_slave_cdc + peakrdl_to_cmdrsp`); wire the adapter directly only when your fabric already speaks cmd/rsp packets.

The wiring is exactly the instantiation in 3.5.7; the earlier example here showed the adapter hanging off a register block's APB port through `reg_*` signals that do not exist — the reversed-direction reading this page used to make.

## 16.7 Related Modules

### 16.7.1 3.5.9 APB4 Front End (apb4\_to\_peakrdl)

peakrdl\_to\_cmdrsp above takes a cmd/rsp stream. apb4\_to\_peakrdl is the module that produces one from an APB4 slave port, and it crosses clock domains while doing it – APB on pclk, the register block on aclk. It is a two-stage composition and adds no register logic of its own:

| Stage | Module            | Domain            |
|-------|-------------------|-------------------|
| 1     | apb4_slave_cdc    | pclk in, aclk out |
| 2     | peakrdl_to_cmdrsp | aclk              |

#### *apb4\_to\_peakrdl Parameters*

| Parameter       | Type | Default      | Description                                                                     |
|-----------------|------|--------------|---------------------------------------------------------------------------------|
| ADDR_WIDTH      | int  | 12           | APB and cpuif byte address width                                                |
| DATA_WIDTH      | int  | 32           | Must match the PeakRDL generation                                               |
| PROT_WIDTH      | int  | 3            | PPROT width                                                                     |
| STRB_WIDTH      | int  | DATA_WIDTH/8 | Derived; do not override                                                        |
| CDC_DEPTH       | int  | 2            | Command/response CDC FIFO depth ( $\geq 2$ )                                    |
| USE_JOHNSON     | int  | 0            | CDC pointer encoding: 0 = Gray (power-of-2 depth only), 1 = Johnson (any depth) |
| USE_2_PHASE_CDC | bit  | 1            | <b>Deprecated and ignored.</b> Accepted for compatibility; apb4_slave_cdc       |

| Parameter | Type | Default | Description                                          |
|-----------|------|---------|------------------------------------------------------|
|           |      |         | does not reference it, so setting it changes nothing |

### *apb4\_to\_peakrdl Ports*

| Port            | Width      | Direction | Description                                   |
|-----------------|------------|-----------|-----------------------------------------------|
| aclk            | 1          | Input     | Register-block clock; drives the cpuif master |
| aresetn         | 1          | Input     | Active-low reset, aclk domain                 |
| pclk            | 1          | Input     | APB clock                                     |
| presetn         | 1          | Input     | Active-low reset, pclk domain                 |
| s_apb_PSEL      | 1          | Input     | APB select                                    |
| s_apb_PENABLE   | 1          | Input     | APB enable                                    |
| s_apb_PREADY    | 1          | Output    | APB ready                                     |
| s_apb_PADDR     | ADDR_WIDTH | Input     | APB address                                   |
| s_apb_PWRITE    | 1          | Input     | Write/read direction                          |
| s_apb_PWDATA    | DATA_WIDTH | Input     | Write data                                    |
| s_apb_PSTRB     | STRB_WIDTH | Input     | Write strobes                                 |
| s_apb_PPROT     | PROT_WIDTH | Input     | Protection attributes                         |
| s_apb_PRDATA    | DATA_WIDTH | Output    | Read data                                     |
| s_apb_PSLVERR   | 1          | Output    | Slave error                                   |
| cpuif_req       | 1          | Output    | Passthrough request                           |
| cpuif_req_is_wr | 1          | Output    | Request is a write                            |
| cpuif_addr      | ADDR_WIDTH | Output    | Request address                               |
| cpuif_wr_data   | DATA_WIDTH | Output    | Write data                                    |

| Port               | Width      | Direction | Description                                |
|--------------------|------------|-----------|--------------------------------------------|
| cpuif_wr_biten     | DATA_WIDTH | Output    | Per-bit write enables, expanded from PSTRB |
| cpuif_req_stall_wr | 1          | Input     | Register block stalls a write              |
| cpuif_req_stall_rd | 1          | Input     | Register block stalls a read               |
| cpuif_rd_ack       | 1          | Input     | Read acknowledge                           |
| cpuif_rd_err       | 1          | Input     | Read error, returned as PSLVERR            |
| cpuif_rd_data      | DATA_WIDTH | Input     | Read data                                  |
| cpuif_wr_ack       | 1          | Input     | Write acknowledge                          |
| cpuif_wr_err       | 1          | Input     | Write error, returned as PSLVERR           |

## 17 3.7 UART to AXI4-Lite Bridge

The `uart_to_axil4` block turns a two-wire serial link into an AXI4-Lite master, so a host can read and write a design's register space over a USB serial adapter with no JTAG, no debug core and no vendor tooling. It is the access path the characterization harnesses use for board bring-up.

Three modules ship in `rtl/uart_to_axil4/`: the byte-level `uart_rx` and `uart_tx`, and `uart_axil_bridge`, which parses commands and drives the bus.

## 17.1 3.7.1 Module Organization

---

uart\_axil\_bridge instantiates four blocks:

| Instance         | Module          | Role                                               |
|------------------|-----------------|----------------------------------------------------|
| u_uart_rx        | uart_rx         | Serial to byte,<br>with a ready/valid<br>handshake |
| u_uart_tx        | uart_tx         | Byte to serial                                     |
| u_axil_wr_timing | axil4_master_wr | Write-channel<br>skid isolation                    |
| u_axil_rd_timing | axil4_master_rd | Read-channel skid<br>isolation                     |

The two axil4\_master\_\* instances are the same AMBA library blocks documented in the RTL AMBA AXI4-Lite book; the bridge does not hand-roll a bus master. The SKID\_DEPTH\_\* parameters below are passed straight through to them.

uart\_rx and uart\_tx use the i\_clk/i\_rst\_n common-block naming while the bridge presents aclk/aresetn. The bridge maps them directly – there is one clock domain and no CDC anywhere in this block.

## 17.2 3.7.2 Command Protocol

---

Commands are ASCII, terminated by newline:

| Command | Form                         | Response       |
|---------|------------------------------|----------------|
| Write   | W <addr_hex><br><data_hex>\n | OK\n           |
| Read    | R <addr_hex>\n               | 0x<data_hex>\n |

The command letter is **case-insensitive**: w/W and r/R are both accepted.

Hex digits are consumed as nibbles and shifted in, with ' ', '\n' and '\r' skipped as separators. The field widths follow the bus parameters –  $\text{ceil}(\text{AXIL\_ADDR\_WIDTH}/4)$  address digits and  $\text{ceil}(\text{AXIL\_DATA\_WIDTH}/4)$  data digits, so 8 of each at the 32-bit defaults.

Because digits shift rather than fill a fixed field, **the parser does not validate length**. Supplying too few digits left-aligns nothing and yields a small address; supplying too many silently discards the oldest nibbles. The host is responsible for sending exactly the digit count the bus width implies.

## 17.3 3.7.3 Error Responses

---

A transfer that returns a non-OKAY response answers ERR\n:

| Outcome                        | Response   |
|--------------------------------|------------|
| Write accepted (bresp == OKAY) | OK\n       |
| Write rejected (SLVERR/DECERR) | ERR\n      |
| Read accepted (rresp == OKAY)  | 0x<data>\n |
| Read rejected (SLVERR/DECERR)  | ERR\n      |

A failed read returns no data at all rather than 0x followed by whatever sat on rdata, because a value that looks well-formed is worse than no value.

Until 2026-09-03 the bridge did not do this. It brought bresp and rresp back from the two axil4\_master\_\* instances and never examined either, so a write that returned SLVERR still answered OK and a failed read returned stale bus data formatted as a normal result. A host had no way to tell a failure from a success.

The host side needed no change: uart\_axi\_bridge.py already returns response.strip() == b'OK' from write(), and read() already requires a 0x prefix. Both were written expecting error signalling the RTL was not producing.

## 17.4 3.7.4 Parameters

---

uart\_axil\_bridge:

*uart\_axil\_bridge Parameters*

| Parameter       | Default | Description                                                             |
|-----------------|---------|-------------------------------------------------------------------------|
| AXIL_ADDR_WIDTH | 32      | AXI4-Lite address width; sets the address digit count                   |
| AXIL_DATA_WIDTH | 32      | AXI4-Lite data width; sets the data digit count and m_axil_wstrb width  |
| CLKS_PER_BIT    | 868     | Baud divisor, aclk cycles per bit. The default is 100 MHz / 115200 baud |
| SKID_DEPTH_AR   | 2       | Passed to axil4_master_rd                                               |
| SKID_DEPTH_R    | 4       | Passed to axil4_master_rd                                               |
| SKID_DEPTH_AW   | 2       | Passed to axil4_master_wr                                               |
| SKID_DEPTH_W    | 4       | Passed to axil4_master_wr                                               |
| SKID_DEPTH_B    | 2       | Passed to axil4_master_wr                                               |

uart\_rx and uart\_tx each take one parameter, CLKS\_PER\_BIT (default 868), with the same meaning.

## 17.5 3.7.5 Ports

### 17.5.1 uart\_axil\_bridge

#### *uart\_axil\_bridge Ports*

| Port            | Width             | Direction | Description                                          |
|-----------------|-------------------|-----------|------------------------------------------------------|
| aclk            | 1                 | Input     | Bridge clock; also clocks both UART blocks           |
| aresetn         | 1                 | Input     | Active-low reset                                     |
| i_uart_rx       | 1                 | Input     | Serial receive line                                  |
| o_uart_tx       | 1                 | Output    | Serial transmit line                                 |
| m_axil_aw_addr  | AXIL_ADDR_WIDTH   | Output    | Write address                                        |
| m_axil_aw_prot  | 3                 | Output    | Write protection attributes                          |
| m_axil_aw_valid | 1                 | Output    | Write address valid                                  |
| m_axil_aw_ready | 1                 | Input     | Write address ready                                  |
| m_axil_wdata    | AXIL_DATA_WIDTH   | Output    | Write data                                           |
| m_axil_ws_trb   | AXIL_DATA_WIDTH/8 | Output    | Write byte strobes                                   |
| m_axil_wvalid   | 1                 | Output    | Write data valid                                     |
| m_axil_wr_ready | 1                 | Input     | Write data ready                                     |
| m_axil_bresp    | 2                 | Input     | Write response; a non-OKAY value yields ERR\n(3.7.3) |
| m_axil_bvalid   | 1                 | Input     | Write response valid                                 |
| m_axil_br_ready | 1                 | Output    | Write response ready                                 |
| m_axil_ar_addr  | AXIL_ADDR_WIDTH   | Output    | Read address                                         |

| Port            | Width           | Direction | Description                                         |
|-----------------|-----------------|-----------|-----------------------------------------------------|
| m_axil_ar_prot  | 3               | Output    | Read protection attributes                          |
| m_axil_ar_valid | 1               | Output    | Read address valid                                  |
| m_axil_ar_ready | 1               | Input     | Read address ready                                  |
| m_axil_rdata    | AXIL_DATA_WIDTH | Input     | Read data, returned in the 0x... response           |
| m_axil_resp     | 2               | Input     | Read response; a non-OKAY value yields ERR\n(3.7.3) |
| m_axil_rvalid   | 1               | Input     | Read data valid                                     |
| m_axil_rready   | 1               | Output    | Read data ready                                     |

### 17.5.2 uart\_rx

#### uart\_rx Ports

| Port       | Width | Direction | Description         |
|------------|-------|-----------|---------------------|
| i_clk      | 1     | Input     | Clock               |
| i_rst_n    | 1     | Input     | Active-low reset    |
| i_rx       | 1     | Input     | Serial receive line |
| o_rx_data  | 8     | Output    | Received byte       |
| o_rx_valid | 1     | Output    | Byte valid          |
| i_rx_ready | 1     | Input     | Consumer ready      |

### 17.5.3 uart\_tx

#### uart\_tx Ports

| Port    | Width | Direction | Description      |
|---------|-------|-----------|------------------|
| i_clk   | 1     | Input     | Clock            |
| i_rst_n | 1     | Input     | Active-low reset |

| Port       | Width | Direction | Description                         |
|------------|-------|-----------|-------------------------------------|
| o_tx       | 1     | Output    | Serial transmit line                |
| i_tx_data  | 8     | Input     | Byte to send                        |
| i_tx_valid | 1     | Input     | Byte valid                          |
| o_tx_ready | 1     | Output    | Transmitter ready for the next byte |

## 18 3.8 Width-plus-Protocol Chains

A chain is a FUB-level module that wires one width converter directly to one protocol converter. Neither chain adds logic of its own: they exist because that pairing is what a real bridge instantiates, and because the bugs worth catching live at the seam between the two, not inside either one.

Both are downsizing chains – a wide AXI4 master talking to a narrow peripheral – and both were built to reproduce the back-to-back page-boundary class of defect at FUB level, where it can be driven directly, instead of only at the bridge level where it was first seen.

## 18.1 3.8.1 axi4\_dwidth\_to\_axil4\_wr\_chain

Write path. The slave port is wide AXI4 (S\_AXI\_DATA\_WIDTH); the master port is AXI4-Lite at M\_AXIL\_DATA\_WIDTH.

| Stage | Module                   | What it does                                                                                                                |
|-------|--------------------------|-----------------------------------------------------------------------------------------------------------------------------|
| 1     | axi4_dwidth_converter_wr | Downsizes the W beats and rewrites awlen/awsize, so the next stage sees a narrow AXI4 burst of M_AXIL_DATA_WIDTH-wide beats |
| 2     | axi4_to_axil4_wr         | Decomposes that narrow burst into single-beat AXI4-Lite writes                                                              |

### *axi4\_dwidth\_to\_axil4\_wr\_chain Parameters*

| Parameter         | Default | Description                                                        |
|-------------------|---------|--------------------------------------------------------------------|
| S_AXI_DATA_WIDTH  | 64      | Slave-side (wide) AXI4 data width                                  |
| M_AXIL_DATA_WIDTH | 32      | Master-side AXI4-Lite data width; the width beats are downsized to |
| AXI_ID_WIDTH      | 8       | Slave-side ID width                                                |
| AXI_ADDR_WIDTH    | 32      | Address width, both sides                                          |
| AXI_USER_WIDTH    | 1       | User-signal width                                                  |

Ports are the slave-side AXI4 write channels and the master-side AXI4-Lite write channels, unchanged from the two modules it wraps – see [AXI4 to AXI4-Lite](#) for the AXI4-Lite side and the width block chapters for the AXI4 side.

## 18.2 3.8.2 axi4\_dwidth\_to\_apb4\_chain

Read path only. **The write channels on the slave port are tied off inside the chain** – this is a read chain, not a general bridge.

| Stage | Module                   | What it does                                                                                                               |
|-------|--------------------------|----------------------------------------------------------------------------------------------------------------------------|
| 1     | axi4_dwidth_converter_rd | Rewrites arlen/arsize down to APB_DATA_WIDTH-wide narrow beats, and upsizes the returning R beats back to S_AXI_DATA_WIDTH |
| 2     | axi4_to_apb4_shim        | Decomposes the narrow burst into single-beat APB reads                                                                     |

### *axi4\_dwidth\_to\_apb4\_chain Parameters*

| Parameter        | Default | Description                                                                                    |
|------------------|---------|------------------------------------------------------------------------------------------------|
| S_AXI_DATA_WIDTH | 64      | Slave-side (wide) AXI4 data width                                                              |
| APB_DATA_WIDTH   | 32      | APB data width on the master side                                                              |
| AXI_ID_WIDTH     | 8       | Slave-side ID width                                                                            |
| AXI_ADDR_WIDTH   | 32      | AXI address width                                                                              |
| APB_ADDR_WIDTH   | 32      | APB address width                                                                              |
| AXI_USER_WIDTH   | 1       | User-signal width                                                                              |
| USE_JOHNSON      | 0       | CDC-FIFO pointer encoding forwarded to the shim: 0 = Gray (power-of-2 depth only), 1 = Johnson |

The direction is worth stating plainly, because the two readings of “downsize” point opposite ways. The wide master issues a wide-beat AR burst; the read converter rewrites it into narrow APB\_DATA\_WIDTH beats for the shim; the R data coming back is reassembled into wide beats for the master.



## 19 3.10 AXI4-Lite to Wishbone B4 Converter

The **axil4\_to\_wb4** converter presents an AXI4-Lite slave and drives a Wishbone B4 master. It is the bridge to drop in when an AXI4-Lite interconnect has to reach a Wishbone peripheral, and it is built from parts the library already has: the AXI4-Lite slave skids, a small conversion core, and **wb4\_master** from `rtl/amba/wb4/`.

## 19.1 3.10.1 Module Organization

---

```
axil4_to_wb4.sv # Wrapper: skids + core + wb4_master
├─ axil4_slave_wr.sv # AW / W / B skid buffers
(rtl/amba/axil4)
├─ axil4_slave_rd.sv # AR / R skid buffers (rtl/amba/axil4)
├─ axil4_to_wb4_core.sv # Five channels -> one command and one
response queue
└─ wb4_master.sv # Queues -> Wishbone B4 bus
(rtl/amba/wb4)
```

## 19.2 3.10.2 Design Philosophy

---

**One queue, in order.** AXI4-Lite has separate write and read paths; Wishbone has one bus. The core merges the two paths into the single command queue of `wb4_master` and relies on a B4 rule to route the answers back: terminations return in issue order. A one-bit side queue, written at issue with the direction of each command, tells the core whether the oldest response belongs on B or on R. There is no ID, no reorder buffer, and no state machine.


**No starvation.** A write is eligible when both AW and W are present, a read when AR is. When both are eligible in the same clock the side not served last is picked, so a stream of writes cannot hold reads off the bus and vice versa.

**Wishbone status becomes an AXI response, nothing more.** ACK maps to OKAY, ERR to SLVERR, and RTY to the RTY\_RESP parameter (SLVERR by default). AXI has no retry, so a retrying Wishbone slave is reported to the requester as an error; a retry wrapper on the FUB side of `wb4_master` is the place to absorb RTY transparently.

**No width conversion.** The AXI4-Lite and Wishbone address and data widths are the same. Put `axi4_dwidth_*` or the wide-alignment blocks in front when the interconnect width differs.

## 19.3 3.10.3 Block Diagram

### 19.3.1 Figure 3.12: AXI4-Lite to Wishbone B4 Converter



## 19.4 3.10.4 Interface Specification

### 19.4.1 Parameters

Table 3.26: AXI4-Lite to Wishbone Parameters

| Parameter              | Default | Description                                                                                |
|------------------------|---------|--------------------------------------------------------------------------------------------|
| ADDR_WIDTH             | 32      | Address width, both sides                                                                  |
| DATA_WIDTH             | 32      | Data width, both sides                                                                     |
| SKID_DEPTH_AW/W/B/AR/R | 2       | AXI4-Lite channel skid depths (2..8)                                                       |
| CMD_DEPTH              | 4       | wb4_master command queue depth                                                             |
| RSP_DEPTH              | 4       | wb4_master response queue depth; bounds the transfers on the bus                           |
| SIDE_DEPTH             | 8       | Core direction queue; at least CMD_DEPTH + RSP_DEPTH so the master sets the pipeline depth |
| CLASSIC                | 0       | 0 = B4 pipelined, 1 = B4 standard (classic) mode; match the slave                          |
| RTY_RESP               | 2'b10   | AXI response returned for a Wishbone RTY                                                   |

The core alone (axil4\_to\_wb4\_core) takes ADDR\_WIDTH, DATA\_WIDTH, SIDE\_DEPTH and RTY\_RESP.

### 19.4.2 Ports

Table 3.27: AXI4-Lite to Wishbone Ports

| Port group                       | Direction | Description                                             |
|----------------------------------|-----------|---------------------------------------------------------|
| s_axil_aw*, s_axil_w*, s_axil_b* | slave     | AXI4-Lite write channels (AWPROT accepted, not carried) |
| s_axil_ar*, s_axil_r*            | slave     | AXI4-Lite read channels                                 |

| Port group                        | Direction | Description                                                     |
|-----------------------------------|-----------|-----------------------------------------------------------------|
| m_wb_CYC/STB/WE/ADR/<br>DAT_W/SEL | out       | Wishbone request                                                |
| m_wb_STALL/ACK/ERR/<br>RTY/DAT_R  | in        | Wishbone flow control<br>and termination                        |
| busy                              | out       | Anything buffered in the<br>skids or a cycle open on<br>the bus |

The core's ports are the five fub\_\* channels of the two slave blocks on one side and the cmd\_\* / rsp\_\* queues of wb4\_master on the other.

## 19.5 3.10.5 Command Formation

---

*Table 3.28: Command Formation*

| AXI4-Lite           | Wishbone command                              |
|---------------------|-----------------------------------------------|
| AWADDR, WDATA,WSTRB | we=1, adr, dat, sel =WSTRB                    |
| ARADDR              | we=0, adr, sel = all ones                     |
| AWPROT / ARPROT     | not carried (Wishbone has no protection bits) |

A write consumes AW and W in the same clock (fub\_awready and fub\_wready are the same signal), so a W beat that arrives before its AW simply waits in its skid. A command is only issued when the side queue has room; when the side queue is full the core stalls, which with  $SIDE\_DEPTH \geq CMD\_DEPTH + RSP\_DEPTH$  never happens before the master's own queues are full.

## 19.6 3.10.6 Response Mapping

---

*Table 3.29: Response Mapping*

| Wishbone status | B / R response            |
|-----------------|---------------------------|
| ACK             | OKAY (2'b00)              |
| ERR             | SLVERR (2'b10)            |
| RTY             | RTY_RESP (default SLVERR) |

RDATA is the Wishbone DAT\_R of the termination. A response that arrives with nothing recorded in the side queue is a master protocol violation (`wb4_master` reports it in simulation); the core drops it so the queue cannot wedge on it.

## 19.7 3.10.7 Timing

*Table 3.30: AXI4-Lite to Wishbone Timing*

| Path                                              | Clocks                                                               |
|---------------------------------------------------|----------------------------------------------------------------------|
| AW+W or AR accepted at s_axil_* to STB on the bus | 2 (skid, master command queue)                                       |
| Termination on the bus to BVALID / RVALID         | 2 (master response queue, skid)                                      |
| Sustained rate, pipelined slave                   | one transfer per clock per direction pair, up to RSP_DEPTH in flight |
| Sustained rate, CLASSIC=1                         | one transfer at a time (B4 standard mode)                            |

## 19.8 3.10.8 Formal

---

formal/converters/axil4\_to\_wb4\_core/ (sv2v flatten, then SymbiYosys) proves on the core: a command is issued only with a request present and carries the payload of the channel it took; B and R never fire together and each consumes exactly one response; the response steered to B or R is the oldest open transfer's; the open count never exceeds SIDE\_DEPTH; OKAY / SLVERR / RTY\_RESP follow the status; and a write is never issued straight after a write while a read is pending. Covers reach both directions, the alternation, an SLVERR on B, a retry on R, and a full side queue.

## 19.9 3.10.9 Testing

---

`dv/tests/test_axil4_to_wb4.py` drives `s_axil_*` with the AXI4-Lite master BFM and answers `m_wb_*` with the framework's Wishbone slave BFM over a memory model (ERR and RTY windows by address) while the Wishbone monitor checks the B4 rules on the wire. The scoreboard mirrors every OKAY write byte-by-byte (strobed writes included) and checks every read against the mirror, and every response code against the window its address is in. Phases: sequential; several transfers in flight against a stalling, a slow-terminating and a mixed slave; drain with one monitor transfer per transaction and no violations. The classic build runs the same phases against a classic-mode slave and returns DECERR for a retry to prove the RTY\_RESP path. Four RTL mutations (direction bit inverted, RTY mapped to OKAY, write strobes dropped, read address taken from AW) all fail the test.

## 19.10 3.10.10 Usage Example

---

```
axil4_to_wb4 #(
 .ADDR_WIDTH (32),
 .DATA_WIDTH (32),
 .CMD_DEPTH (4),
 .RSP_DEPTH (4),
 .SIDE_DEPTH (8),
 .CLASSIC (0)
) u_axil_to_wb (
 .aclk (clk),
 .aresetn (rst_n),
 .s_axil_awaddr (awaddr), .s_axil_awprot
 (awprot), .s_axil_awvalid (awvalid), .s_axil_awready (awready),
 .s_axil_wdata (wdata), .s_axil_wstrb (wstrb), .s_axil_wvalid
 (wvalid), .s_axil_wready (wready),
 .s_axil_bresp (bresp), .s_axil_bvalid (bvalid), .s_axil_bready
 (bready),
 .s_axil_araddr (araddr), .s_axil_arprot
 (arprot), .s_axil_arvalid (arvalid), .s_axil_arready (arready),
 .s_axil_rdata (rdata), .s_axil_rresp (rresp), .s_axil_rvalid
 (rvalid), .s_axil_rready (rready),
 .m_wb_CYC (wb_cyc), .m_wb_STB (wb_stb), .m_wb_WE
 (wb_we), .m_wb_ADR (wb_adr),
 .m_wb_DAT_W (wb_dat_w), .m_wb_SEL (wb_sel),
 .m_wb_STALL (wb_stall), .m_wb_ACK (wb_ack), .m_wb_ERR
 (wb_err), .m_wb_RTY (wb_rty),
 .m_wb_DAT_R (wb_dat_r),
 .busy ()
);
```

RTL Design Sherpa · Learning Hardware Design Through Practice [GitHub](#) · [Documentation Index](#) · [MIT License](#)

---

## 20 3.11 Wishbone B4 to AXI4-Lite Converter

The **wb4\_to\_axil4** converter presents a Wishbone B4 slave and drives an AXI4-Lite master. It is the mirror of [3.10](#) and the harder direction, for one reason given below.

## 20.1 3.11.1 Module Organization

---

```
wb4_to_axil4.sv # Wrapper: slave + core + the two masters
├─ wb4_slave.sv # Wishbone bus -> cmd/rsp queues
(rtl/amba/wb4)
├─ wb4_to_axil4_core.sv # Queues -> five AXI4-Lite channels,
responses merged
├─ axil4_master_wr.sv # AW / W / B channels (rtl/amba/axil4)
└─ axil4_master_rd.sv # AR / R channels (rtl/amba/axil4)
```

## 20.2 3.11.2 Why This Direction Is Harder

---

Going AXI4-Lite to Wishbone, the two request paths merge into one bus and the answers come back in issue order for free, because that is what Wishbone does.

Coming back the other way the ordering runs against you. **Wishbone B4 terminates in issue order. AXI4-Lite's B and R channels are independent and a slave may answer them in any order.** A read issued behind a slow write will routinely finish first. If the converter passed each response straight through, the Wishbone master would see its transfers terminate out of order, which the protocol forbids and its own bookkeeping would mispair.

So the core records the direction of every command it issues in an in-order side queue and releases only the response at the head. A response that completed early waits in its channel's skid until its predecessors have retired.

**Ordering cost.** A slow write at the head holds back a read that finished behind it. `OUTSTANDING = 1` avoids the question entirely and is the default; raise it only when the AXI slave answers roughly in order, or head-of-line waiting eats the benefit.

## 20.3 3.11.3 Interface Specification

### 20.3.1 Parameters

Table 3.31: Wishbone to AXI4-Lite Parameters

| Parameter                  | Default | Description                                                  |
|----------------------------|---------|--------------------------------------------------------------|
| ADDR_WIDTH                 | 32      | Address width, both sides                                    |
| DATA_WIDTH                 | 32      | Data width, both sides                                       |
| CMD_DEPTH /<br>RSP_DEPTH   | 2       | The wb4_slave's queue depths                                 |
| MAX_OUTSTANDING            | 16      | The wb4_slave's bus-side bound                               |
| CLASSIC                    | 0       | 0 = B4 pipelined, 1 = B4 standard; match the Wishbone master |
| OUTSTANDING                | 1       | Commands in flight through the core. 1 = strictly serialised |
| SKID_DEPTH_AW/W/B/<br>AR/R | 2       | The AXI4-Lite masters' channel skids                         |
| AXIL_PROT                  | 3'b000  | AWPROT/ARPROT driven on every command                        |

### 20.3.2 Ports

Table 3.32: Wishbone to AXI4-Lite Ports

| Port group                                | Direction | Description                               |
|-------------------------------------------|-----------|-------------------------------------------|
| s_wb_CYC/STB/WE/ADR/<br>DAT_W/SEL/CTI/BTE | in        | Wishbone request                          |
| s_wb_STALL/ACK/ERR/<br>RTY/DAT_R          | out       | Flow control and termination              |
| m_axil_aw*, m_axil_w*,<br>m_axil_b*       | master    | AXI4-Lite write channels                  |
| m_axil_ar*, m_axil_r*                     | master    | AXI4-Lite read channels                   |
| busy                                      | out       | Anything buffered, or a cycle open on the |

| Port group | Direction | Description   |
|------------|-----------|---------------|
|            |           | Wishbone side |

## 20.4 3.11.4 Command Formation

---

Table 3.33: Command Formation

| Wishbone              | AXI4-Lite                   |
|-----------------------|-----------------------------|
| WE=1, ADR, DAT_W, SEL | AWADDR + WDATA, WSTRB = SEL |
| WE=0, ADR             | ARADDR                      |
| CTI / BTE             | <b>dropped</b> (see below)  |
| -                     | AWPROT/ARPROT = AXIL_PROT   |

A write only launches when **both** AWREADY and WREADY are high in the same clock. Letting AW go while W stalled would leave a half-issued transfer that the outstanding count could not describe honestly.

**The burst hints stop here.** CTI and BTE describe a registered-feedback burst, and AXI4-Lite is single-beat with no field to carry them. A burst-aware target behind an AXI4-Lite bus is a contradiction, so the converter drops them rather than inventing a mapping.

## 20.5 3.11.5 Response Mapping

---

*Table 3.34: Response Mapping*

| AXI4-Lite BRESP / RRESP | Wishbone termination  |
|-------------------------|-----------------------|
| OKAY                    | ACK                   |
| SLVERR                  | ERR                   |
| DECERR                  | ERR                   |
| -                       | RTY is never produced |

B4 has one error termination and no way to say which kind, so SLVERR and DECERR both become ERR. RTY never appears because an AXI slave has no way to ask for a retry; a Wishbone master that wants retries needs `wb4_retry` on the other side of the bus, not here.

## 20.6 3.11.6 Formal

---

formal/converters/wb4\_to\_axil4\_core/ proves: a write takes AW and W in the same clock and never one alone; a command becomes a write or a read, never both; the response consumed is the head's channel **and is backed by a real response on that channel**; the open count never exceeds OUTSTANDING; OKAY becomes ACK, any error becomes ERR, and RTY is never produced.

The backing property earned its place. An earlier property set proved every ordering rule above and still passed a mutation that asserted `rsp_valid` whenever *either* channel had a response, returning a status read off a channel that had nothing. Simulation caught it as a hang; formal did not, until the property was added.

## 20.7 3.11.7 Testing

---

`dv/tests/test_wb4_to_axil4.py` drives `s_wb_*` with the framework's Wishbone master BFM and answers `m_axil_*` with the AXI4-Lite slave BFMs over one shared memory model, so a write really lands where a later read finds it. Addresses beyond the memory model answer SLVERR and must arrive as ERR.

The phases are sequential traffic, pipelined traffic under gappy and sparse request pacing, and a directed **ordering probe**: write-then-read pairs where the testbench answers the write channel slowly and the read channel quickly, so the AXI side completes the read first every time. The scoreboard pairs terminations with issue order, so a response that overtook its predecessor shows up as a status or data mismatch rather than as silence.

Three RTL mutations fail the test: completing from whichever channel is ready, corrupting read data, and launching AW without W.

## 20.8 3.11.8 Usage Example

---

```
wb4_to_axil4 #(
 .ADDR_WIDTH (32),
 .DATA_WIDTH (32),
 .OUTSTANDING (1), // raise only if the AXI slave answers in
order
 .CLASSIC (0)
) u_wb_to_axil (
 .aclk (clk), .aresetn (rst_n),
 .s_wb_CYC (wb_cyc), .s_wb_STB (wb_stb), .s_wb_WE
(wb_we), .s_wb_ADR (wb_adr),
 .s_wb_DAT_W (wb_dat_w), .s_wb_SEL (wb_sel), .s_wb_CTI
(wb_cti), .s_wb_BTE (wb_bte),
 .s_wb_STALL (wb_stall), .s_wb_ACK (wb_ack), .s_wb_ERR
(wb_err), .s_wb_RTY (wb_rty),
 .s_wb_DAT_R (wb_dat_r),
 .m_axil_awaddr (awaddr), .m_axil_awprot (awprot), .m_axil_awvalid
(awvalid), .m_axil_awready (awready),
 .m_axil_wdata (wdata), .m_axil_wstrb (wstrb), .m_axil_wvalid
(wvalid), .m_axil_wready (wready),
 .m_axil_bresp (bresp), .m_axil_bvalid (bvalid), .m_axil_bready
(bready),
 .m_axil_araddr (araddr), .m_axil_arprot (arprot), .m_axil_arvalid
(arvalid), .m_axil_arready (arready),
 .m_axil_rdata (rdata), .m_axil_rresp (rresp), .m_axil_rvalid
(rvalid), .m_axil_rready (rready),
 .busy ()
);
```

## 20.9 Navigation

---

Next: [APB to AXI4 \(Requester\)](#)

[RTL Design Sherpa](#) · [Learning Hardware Design Through Practice](#) [GitHub](#) · [Documentation Index](#) · [MIT License](#)

---

## 21 APB to AXI4 (Requester)

**Modules:** apb4\_to\_axi4.sv, apb5\_to\_axi4.sv, apb\_cmdrsp\_to\_axi4.sv **Filelists:** rtl/filelists/apb4\_to\_axi4.f, rtl/filelists/apb5\_to\_axi4.f (each -f's its rtl/amba APB completer and the shared requester core)

## 21.1 Overview

The mirror of [AXI4 to APB](#). There an AXI4 requester is turned into an APB requester; here an APB requester – a small CPU, a debug port, a register sequencer – becomes an AXI4 requester. The bridge generator puts one of these in front of a master port declared `protocol = "apb" or "apb5"` (BRIDGE-014).

Two proven pieces and one small new one:

*Table 3.35: APB-to-AXI4 requester pieces*

| Piece                                           | Job                                                                                                                      |
|-------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------|
| <code>apb4_slave / apb5_slave (rtl/amba)</code> | The APB completer surface:<br>PSEL/PENABLE/PREADY handshake,<br>skid-buffered cmd/rsp pair, the<br>orphan-response guard |
| <code>apb_cmdrsp_to_axi4</code>                 | One command -> one single-beat AXI4<br>transaction -> one response                                                       |
| <code>apb4_to_axi4 / apb5_to_axi4</code>        | The wiring, and for APB5 the USER<br>mapping                                                                             |

APB is strictly one-outstanding, so there is never more than one AXI4 transaction in flight and the response needs no ID matching: every request carries the constant `DEFAULT_ID`. AW and W are presented together and each is held until its own handshake, so the requester is legal against a completer that takes W before AW or the other way round.

```
pwrite=1 : AW + W (awlen=0, wlast=1) -> B -> rsp (PSLVERR =
bresp[1])
pwrite=0 : AR (arlen=0) -> R -> rsp (PRDATA, PSLVERR
= rresp[1])
```

Both SLVERR and DECERR fold to PSLVERR: APB has one error bit, so a requester cannot tell an unmapped address from a failing completer. R is drained to RLAST even though every request asks for one beat, so a completer that answers with more cannot wedge the port; only the first beat's data is returned.

## 21.2 Fields the requester invents

Table 3.36: AXI4 request fields from an APB transfer

| Field                   | Value                                 | Why                                                                 |
|-------------------------|---------------------------------------|---------------------------------------------------------------------|
| AxID                    | DEFAULT_ID                            | APB has no ID; one value is enough with one transaction outstanding |
| AxLEN                   | 0                                     | one beat per transfer                                               |
| AxSIZE                  | $\$c\log_2(\text{DATA\_WIDTH}/8)$     | the full data width; PSTRB becomesWSTRB unchanged                   |
| AxBURST                 | INCR                                  |                                                                     |
| AxCACHE                 | DEFAULT_CACHE (device non-bufferable) | the natural attribute for a peripheral-bus requester                |
| AxPROT                  | PPROT                                 | forwarded                                                           |
| AxLOCK, AxQOS, AxREGION | 0                                     |                                                                     |
| WLAST                   | 1                                     |                                                                     |

## 21.3 APB5 sideband

---

apb5\_to\_axi4 maps the APB5 USER signals onto the AXI USER fields with a plain size cast – the low bits travel, anything wider is zero-filled or dropped:

*Table 3.37: APB5 USER mapping*

| APB5   | AXI4            |
|--------|-----------------|
| PAUSER | awuser / aruser |
| PWUSER | wuser           |
| buser  | PBUSER          |
| ruser  | PRUSER          |

PWAKEUP is requester-driven (AMBA APB5, IHI 0024E) and is an **input** here, accepted and terminated – there is nothing on the AXI side for it to become. apb5\_slave's own PWAKEUP output (its wake-up *request* toward the requester) is left unconnected with wakeup\_request tied low. Parity is not part of this surface: apb5\_slave is instantiated with ENABLE\_PARITY=0 and its parity pins tied off, matching axi4\_to\_apb5\_shim.

Inside the bridge the AXI USER width is 1, so one bit of PAUSER reaches the fabric – the same width the AXI5-Lite slave side exposes, which is how an APB5 requester's user bit can be observed at an AXI5-Lite completer. The return direction is not symmetrical: the bridge's master adapters tie response-side USER to zero (see the bridge PRD), so PBUSER/PRUSER read 0 on a bridge port even though the converter itself carries them.

## 21.4 Parameters

Table 3.38: APB-to-AXI4 parameters

| Parameter               | Default | Description                                                  |
|-------------------------|---------|--------------------------------------------------------------|
| APB_ADDR_WIDTH          | 32      | PADDR and AxADDR width (one address space; no windowing)     |
| APB_DATA_WIDTH          | 32      | PWDATA/PRDATA and the AXI data width – equal by construction |
| AXI_ID_WIDTH            | 1       | width of the constant ID                                     |
| AXI_USER_WIDTH          | 1       | AXI USER width (APB4 wrapper drives 0)                       |
| DEFAULT_ID              | '0      | the ID on every request                                      |
| DEPTH                   | 2       | cmd/rsp skid depth inside the APB completer                  |
| APB_{A,W,R,B}USER_WIDTH | 1       | APB5 only; independent of AXI_USER_WIDTH                     |

DEFAULT\_CACHE on the core is not exposed by the wrappers; a consumer that needs a different cache attribute instantiates `apb_cmdrsp_to_axi4` directly.

## 21.5 Verification

---

`dv/tests/test_apb4_to_axi4.py` and `dv/tests/test_apb5_to_axi4.py` share `dv/tbclasses/apb_to_axi4_tb.py`: an APB master BFM drives the completer surface, memory-backed AXI4 slave BFMs answer the requester surface, and a monitor samples every AW/W/AR handshake. Checked: the data round trip (the completer's memory is read directly as well as through a second transfer, so a converter echoing its own write data is caught), the single-beat shape of every request, PPR0T->AxPR0T, PSTRB->WSTRB against a byte-accurate shadow, out-of-range SLVERR folding to PSLVERR with nothing written and the port working afterwards, the APB5 USER mapping in both directions, and a slow completer holding AW/W/AR to their handshakes. Twelve cells at FULL (data widths 16/32/64, AXI USER 1/2/4, APB USER 2/4).

## 21.6 Related Modules

---

The bridge generator's master adapter instantiates `apb4_to_axi4` / `apb5_to_axi4` for protocol = "apb" / "apb5" master ports and feeds the result to the same `axi4_slave_{wr,rd}` timing wrapper an AXI4 master port gets, so decode, width adaptation (including the wide-slave aligner) and the response mux are untouched. Sign-off:  
`projects/components/bridge/dv/tests/test_bridge_2x3_apb_req_paths.py`.

## 21.7 Navigation

---

**Previous:** [Wishbone B4 to AXI4-Lite](#) **Next:** [AXI4 to Wishbone B4](#)

[RTL Design Sherpa](#) · [Learning Hardware Design Through Practice](#) [GitHub](#) · [Documentation Index](#) · [MIT License](#)

---

## 22 AXI4 to Wishbone B4

**Module:** `axi4_to_wb4.sv` **Filelist:** `rtl/filelists/axi4_to_wb4.f` (-f's the two AXI4-Lite decomposers and `axi4_to_wb4`)

## 22.1 Overview

Full AXI4 completer in, Wishbone B4 requester out – the Wishbone counterpart of `axi4_to_apb4_shim`, built entirely from blocks this component already has:

*Table 3.39: axi4\_to\_wb4 stages*

| Stage | Block                                                            | Job                                                                                                      |
|-------|------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------|
| 1     | <code>axi4_to_axil4_wr</code> ,<br><code>axi4_to_axil4_rd</code> | burst decomposition:<br>one AXI4-Lite beat per<br>AXI4 beat, response<br>folding, ID/LAST<br>bookkeeping |
| 2     | <code>axil4_to_wb4</code>                                        | five AXI4-Lite channels<br>to one Wishbone<br>command stream and<br>back, in-order<br>termination        |

Every AXI4 beat becomes one Wishbone transfer: WE from the channel it came on, SEL from WSTRB, ADR the beat's address. Termination maps ACK to OKAY, ERR to SLVERR, RTY to RTY\_RESP (SLVERR by default; AXI has no retry, so a completer that wants retries handled needs a `wb4_retry` between it and this block). Same address and data width on both sides; the bridge's width converters sit in front when they differ.

The Wishbone side is B4 pipelined unless `CLASSIC = 1` (B4 standard mode); match the completer – the two do not mix (see the `rtl/amba/wb4` README). CTI/BTE are driven CLASSIC/LINEAR: the decomposer issues one transfer at a time and burst hints are advisory in B4.

## 22.2 Parameters

Table 3.40: axi4\_to\_wb4 parameters

| Parameter                                                             | Default      | Description                                                        |
|-----------------------------------------------------------------------|--------------|--------------------------------------------------------------------|
| AXI_ID_WIDTH,<br>AXI_ADDR_WIDTH,<br>AXI_DATA_WIDTH,<br>AXI_USER_WIDTH | 8, 32, 32, 1 | the AXI4 face; data width is also the Wishbone width               |
| SKID_DEPTH_AW/W/B/<br>AR/R                                            | 2            | axi4_to_wb4's channel skids                                        |
| CMD_DEPTH, RSP_DEPTH                                                  | 4, 4         | wb4_master queues; RSP_DEPTH bounds transfers in flight on the bus |
| SIDE_DEPTH                                                            | 8            | direction queue in axi4_to_wb4_core ( $\geq$ CMD + RSP)            |
| CLASSIC                                                               | 0            | 0 = B4 pipelined, 1 = B4 standard                                  |
| RTY_RESP                                                              | 2'b10        | AXI response for a Wishbone RTY                                    |

## 22.3 Verification

---

`dv/tests/test_axi4_to_wb4.py` with `dv/tbclasses/axi4_to_wb4_tb.py`: AXI4 master BFM's drive bursts of 1..16 beats into a `WB4Slave` whose `status_hook` turns two address windows into `ERR` and `RTY`. Checked: an N-beat burst is exactly N Wishbone transfers in address order, each landing in the completer's memory, and the read burst returns it beat by beat; `WSTRB` becomes `SEL` against a byte shadow; `ERR` folds to `SLVERR` and `RTY` to `RTY_RESP` on both B and R with the port working after; fixed, stalling and slow completer profiles; pipelined and classic builds.

## 22.4 Related Modules

---

The bridge generator's slave adapter instantiates this for a slave port declared `protocol = "wb4"` (BRIDGE-019), through `Axi4ToWb4Shim`, with the same monitor-wrapper sandwich and response intercepts the APB shim gets. Sign-off:

`projects/components/bridge/dv/tests/test_bridge_2x2_wb4_paths.py`.

## 22.5 Navigation

---

**Previous:** [APB to AXI4 \(Requester\)](#) **Next:** [Wishbone B4 to AXI4](#)

[RTL Design Sherpa](#) · [Learning Hardware Design Through Practice](#) [GitHub](#) · [Documentation Index](#) · [MIT License](#)

---

## 23 Wishbone B4 to AXI4

**Module:** wb4\_to\_axi4.sv **Filelist:** rtl/filelists/wb4\_to\_axi4.f (-f's wb4\_to\_axi4's closure)

## 23.1 Overview

---

Wishbone B4 completer in, full AXI4 requester out – the Wishbone counterpart of `apb4_to_axi4.wb4_to_axil4` does the work (`wb4_slave`, the in-order response merge, the AXI4-Lite masters); this wrapper promotes its AXI4-Lite face to AXI4 the way any AXI4-Lite requester is promoted in the fabric:

*Table 3.41: AXI4 fields the wrapper invents*

| Field                                     | Value                                                                                  |
|-------------------------------------------|----------------------------------------------------------------------------------------|
| AxID                                      | DEFAULT_ID (constant; Wishbone has no ID)                                              |
| AxLEN, AxBURST, WLAST<br>AxSIZE           | 0, INCR, 1 – every transfer is one beat the full data width; SEL rides WSTRB unchanged |
| AxPROT                                    | AXIL_PROT (Wishbone carries no protection bits)                                        |
| AxLOCK, AxCACHE, AxQOS, AxREGION,<br>USER | 0 (AxCACHE = DEFAULT_CACHE)                                                            |
| BID/RID, RLAST, response USER             | dropped: one ID in use, one beat per read                                              |

ERR to the requester when the AXI response is SLVERR or DECERR (Wishbone has one way to say either); RTY is never generated. The core's OUTSTANDING (default 1) is how many commands it lets into the AXI side at once – B4 terminates in issue order, and above 1 a slow write at the head holds back a read that finished behind it.

## 23.2 Parameters

---

ADDR\_WIDTH, DATA\_WIDTH, CMD\_DEPTH, RSP\_DEPTH, MAX\_OUTSTANDING, CLASSIC, OUTSTANDING, SKID\_DEPTH\_\*, AXIL\_PROT are wb4\_to\_axil4's (Table 3.31); added here: AXI\_ID\_WIDTH (1), AXI\_USER\_WIDTH (1), DEFAULT\_ID ('0'), DEFAULT\_CACHE (4'b0000).

## 23.3 Verification

---

`dv/tests/test_wb4_to_axi4.py` with `dv/tbclasses/wb4_to_axi4_tb.py`: a WB4Master drives the completer surface, memory-backed AXI4 slave BFM's answer, a monitor samples every AW/W/AR handshake. Checked: the single-beat shape of every request, SEL toWSTRB against a byte shadow, the data round trip (memory read directly and through a second transfer), the completer's out-of-range SLVERR terminating ERR with nothing written and the port working after, a slow completer, pipelined and classic builds.

## 23.4 Related Modules

---

The bridge generator's master adapter instantiates this for a master port declared `protocol = "wb4"` (BRIDGE-019) and feeds the ordinary `axi4_slave_{wr,rd}` timing wrapper with its `m_axi` face, exactly as it does `apb4_to_axi4` for an APB requester.

## 23.5 Navigation

---

**Previous:** [AXI4 to Wishbone B4](#)

[RTL Design Sherpa](#) · [Learning Hardware Design Through Practice](#) [GitHub](#) · [Documentation Index](#) · [MIT License](#)

---

## 24 4.1 Width Converter FSMs

### 24.1 Overview

---

Here's the twist with the width converters: the upsize and downsize primitives don't have a state machine at all, and that's deliberate — what they have instead is tighter. The full AXI4 dwidth converters do run explicit FSMs, and those are small.

## 24.2 Functional Description

---

### 24.2.1 4.1.1 Upsize Structure (no FSM)

The **axi\_data\_upsize** module has NO state machine. Earlier revisions of this chapter documented an ACCUMULATE/OUTPUT FSM; the RTL is a register-and-gating structure built deliberately so that output and the next accumulation OVERLAP — an explicit OUTPUT state would insert the very bubble the design avoids.

The real structure:

- `r_beat_ptr` walks slots as narrow beats land in the accumulator. A fresh BURST's first beat lands at `start_lane` (mid-word INCR starts, CONV-006), so the effective lane is `w_lane = (ptr=='0') ? (burst_fresh ? start_lane : 0) : ptr;`
- a group completes when the LANE reaches `WIDTH_RATIO-1` or on an early `narrow_last` — a mid-word burst's first wide group holds `WIDTH_RATIO - start_lane` beats, not `WIDTH_RATIO`; later groups start at lane 0 and hold the full count;
- `r_wide_valid` presents the completed group on the wide side while the NEXT group starts accumulating — when a wide handshake and a completing narrow beat coincide, a single priority selection keeps the freshly completed group. (Last-write-wins ordering across two NBA blocks was the BUG here, not the fix — the wide-accept block's clear clobbered the fresh group's valid; the `dnsz` is the one that legitimately relies on last-write-wins, for its atomic replace.)
- `narrow_ready = !r_wide_valid || wide_ready`, so the narrow side stalls only when the wide side is holding a beat hostage.

On an early `narrow_last`, unwritten accumulator slots are zeroed in the same single write that deposits the first beat of a group — residual data from the previous group must not leak into a partial wide beat.

### 24.2.2 4.1.2 Downsize Structure (single buffer, no FSM)

Same story: no IDLE/LOAD/OUTPUT machine. The single buffer performs an **atomic replace** — `wide_ready` asserts during the LAST narrow beat of the current wide beat, so the replacement lands as the drain finishes and a steady stream pays no per-beat bubble (measured 0.992 beats/cycle, see 2.3):

```
assign wide_ready = !r_wide_buffered || (narrow_ready && w_last_narrow_beat);
```

With TRACK\_BURSTS=1 the condition narrows to mid\_burst\_replace, which excludes each burst's final beat — one bubble per burst boundary, not per beat. A beat pointer (r\_beat\_ptr) selects the slice driven onto the narrow side; on the burst's FIRST wide word it starts at start\_lane (mid-word INCR starts, CONV-006 — the narrow master receives the bytes it addressed, not the aligned-down word's bytes), and at lane 0 for every later wide word. narrow\_last in tracked mode comes only from the beat counter reaching burst\_len + 1.

### 24.2.34.1.3 Full Converter FSMs

#### 24.2.3.1 Write Converter (axi4\_dwidth\_converter\_wr)

IDLE:

- AW valid → accept AW, store info
- W valid → buffer W data

AW\_ACCEPT:

- downstream AW ready → forward adjusted AW

W\_CONVERT:

- upsize accumulating narrow W beats
- on output → forward wide W beat

B\_FORWARD:

- B from downstream → forward to master

#### 24.2.3.2 Read Converter (axi4\_dwidth\_converter\_rd)

IDLE:

- AR valid → accept AR, store info, forward adjusted AR

AR\_FORWARD:

- downstream AR ready → wait for R

R\_CONVERT:

- downsize splitting wide R into narrow beats
- track burst count for RLAST


R\_FORWARD:

- narrow R beat ready → forward to master
- on RLAST → IDLE

## 24.3 Timing

### 24.3.1 Upsize (8:1 ratio, m\_ready held high)


No s\_ready bubble: the wide handshake overlaps the next group's first narrow beat.



s\_ready only drops if the wide side stalls with a completed group waiting (!r\_wide\_valid || wide\_ready fails).

### 24.3.2 Downsize (8:1 ratio, single buffer, atomic replace)

s\_ready re-asserts DURING the last narrow beat, so WIDE1 is accepted as D7 drains — back-to-back wide beats with no gap:



## 25 4.2 Protocol Converter FSMs

### 25.1 Overview

---

The state machines inside the protocol converters: every one of them is smaller than the earlier documentation made it look.

## 25.2 Functional Description

---

### 25.2.1 4.2.1 AXI4 to AXI4-Lite Read FSM

The real machine has THREE states, and none of them waits for R:

```
typedef enum logic [1:0] {
 RD_IDLE = 2'b00, // no burst in flight
 RD_BURST = 2'b01, // issuing the AXI4 reads of a burst
 RD_LAST_BEAT = 2'b10 // issuing the burst's final read
} rd_state_t;
```

Earlier revisions documented IDLE/SINGLE/DECOMPOSE/WAIT\_R. Two of those never existed: a single-beat read (ARLEN==0) takes a passthrough leg and never leaves RD\_IDLE, and there is no wait-for-R state at all — the address and data paths are decoupled, with ARs issuing at the downstream accept rate regardless of R (see 3.2). What serializes bursts is not an FSM state but the one-outstanding-burst guard: the burst-tracking registers (r\_r\_id/r\_r\_len/r\_r\_beat\_count) hold exactly one burst, so s\_axi\_arready is held off until the in-flight burst delivers its last beat. Response aggregation is the live-beat worst-case fold of 4.3.3.

### 25.2.2 4.2.2 AXI4 to AXI4-Lite Write FSM

Also three states (WR\_IDLE/WR\_BURST/WR\_LAST\_BEAT), same passthrough leg for AWLEN==0, same one-outstanding guard on the B side. There are no r\_aw\_pending/r\_w\_pending flags: W is gated directly against the state of the AW that owns it —

- during the burst-capture cycle W is held off (w\_burst\_capture), or the first W beat of the next burst would slip out while the previous burst is finishing and pair with the wrong address;
- during a burst, W passes only once this burst's AW has been sent (r\_aw\_sent) or is being sent in the same cycle;
- single beats pass straight through.

See 3.2's "W gated on the AW that owns it" for the exact equations.

### 25.2.3 4.2.3 AXI4 to APB FSMs

The convert core runs a command FSM (IDLE/READ/WRITE, writes preferred) that packetizes each burst into per-APB-beat commands, and a response FSM (RSP\_IDLE/RSP\_ACTIVE) that reassembles rsp packets into AXI responses. PSEL/PENABLE setup/access phases are NOT here — they belong to apb4\_master on the far side of the shim's CDC. Full description in 3.4.5.

## 25.3 Timing

---

### 25.3.1 4.2.4 Timing Analysis

**AXI4 to AXIL4.** A single-beat transfer is a combinational passthrough — no converter-inserted wait state. In a burst the decomposed requests issue at the downstream ACCEPT rate, independent of responses: `m_axil_arvalid` in `RD_BURST` holds every cycle and the beat tracker advances on `m_axil_arvalid && m_axil_arready` — no signal in the request path samples `m_axil_rvalid/m_axil_bvalid`. Against a pipelining slave, `N` requests issue in `N` consecutive cycles and the burst finishes one response latency later; the old “2 cycles overhead /  $2N$  cycles per burst” figures described a serialization the RTL does not have. Bursts do not overlap each other (the one-outstanding guard above); that is the only converter-imposed serialization.

**AXI4 to APB.** Per APB beat: command packet -> CDC (2-flop gray crossing each way) -> APB setup + access (2 `pclk` minimum, plus `PREADY` stalls) -> response packet -> CDC back. Latency is dominated by the two clock crossings and the APB protocol itself; no measured characterization is published here.

RTL Design Sherpa · Learning Hardware Design Through Practice [GitHub](#) · [Documentation Index](#) · [MIT License](#)

---

## 26 4.3 Burst Decomposition

### 26.1 Overview

---

Burst handling is where the two converter families diverge: width converters rescale the burst length, protocol converters split the burst into single beats. Same AXI burst in — two very different kinds of bookkeeping out.

## 26.2 Functional Description

---

### 26.2.1 4.3.1 Width Converter Burst Handling

#### 26.2.1.1 Burst Length Adjustment

When converting widths, burst length changes inversely with data width:

$$M\_AWLEN = \text{ceil}((S\_AWLEN + 1) / \text{RATIO}) - 1 = ((S\_AWLEN + 1) + \text{RATIO} - 1) / \text{RATIO} - 1$$

Example (64-bit to 512-bit, RATIO=8):

S\_AWLEN = 7 (8 beats × 64 bits = 512 bits)

M\_AWLEN = (7 + 1) / 8 - 1 = 0 (1 beat × 512 bits)

S\_AWLEN = 15 (16 beats × 64 bits = 1024 bits)

M\_AWLEN = (15 + 1) / 8 - 1 = 1 (2 beats × 512 bits)

#### 26.2.1.2 Non-Aligned Bursts

When the burst length is not a multiple of the ratio:

S\_AWLEN = 5 (6 beats), RATIO = 8

M\_AWLEN = ((5 + 1) + 7) / 8 - 1 = 0 (1 wide beat -- the ceiling term is what prevents the -1 a floor divide gives)

The 6 narrow beats pack into 1 wide beat.  
Last 2 positions have WSTRB = 0 (no write).

### 26.2.2 4.3.2 Protocol Converter Burst Handling

#### 26.2.2.1 AXI4 to AXI4-Lite Decomposition

AXI4-Lite only supports single-beat transactions, so all bursts must be decomposed:

AXI4 Burst:

ARADDR = 0x1000

ARLEN = 3 (4 beats)

ARSIZE = 2 (4 bytes)

AXI4 Sequence:

Transaction 0: ARADDR = 0x1000

Transaction 1: ARADDR = 0x1004

Transaction 2: ARADDR = 0x1008

Transaction 3: ARADDR = 0x100C

### 26.2.2.2 Address Increment Calculation

```
// Calculate address increment based on burst type and size
function automatic [ADDR_WIDTH-1:0] next_address(
 input [ADDR_WIDTH-1:0] current_addr,
 input [2:0] size,
 input [1:0] burst,
 input [7:0] len,
 input [7:0] beat
);
 logic [ADDR_WIDTH-1:0] increment;
 logic [ADDR_WIDTH-1:0] wrap_mask;

 increment = 1 << size;

 case (burst)
 2'b00: // FIXED
 return current_addr; // No increment

 2'b01: // INCR
 return current_addr + increment;

 2'b10: // WRAP
 wrap_mask = ((len + 1) << size) - 1;
 return (current_addr & ~wrap_mask) |
 ((current_addr + increment) & wrap_mask);

 default:
 return current_addr + increment;
 endcase
endfunction
```

### 26.2.34.3.3 Response Aggregation

Several downstream responses collapse into one upstream response, worst case winning (OKAY < EXOKAY < SLVERR < DECERR by numeric compare). The subtlety is WHERE the final beat's own response enters the result, and getting it wrong is not hypothetical: the converters shipped with the registered-only version below, and every single-beat error was reported upstream as OKAY.

**Wrong (the shipped bug).** Accumulate in a register and emit the register:

```
always_ff @(posedge aclk or negedge aresetn) begin
 ...
 else if (m_axil_bvalid && m_axil_bready)
 if (m_axil_bresp > r_b_resp_accum)
 r_b_resp_accum <= m_axil_bresp;
end
```

```
assign s_axi_bresp = r_b_resp_accum; // WRONG
```

The accumulator is a non-blocking assign: on the handshake that emits the response it does not yet contain THAT beat's bresp. For a single-beat transfer the only beat is the final one, so its error is exactly the one dropped.

**Right (the current RTL).** Fold the live beat combinationally at the point of emission:

```
assign w_b_resp_worst = (m_axil_bresp > r_b_resp_accum) ? m_axil_bresp
 :
r_b_resp_accum;
assign s_axi_bresp = w_b_resp_worst;
```

The read path is the same shape, gated to the final beat:

```
assign s_axi_rresp = (r_r_beat_count == r_r_len) ? w_r_resp_worst
 : m_axil_rresp;
```

Both were fixed with tests that inject SLVERR from the downstream slave and assert it reaches the upstream master — including the last-beat-only case, where earlier beats cannot mask the loss. See test\_error\_response in the axi4\_to\_axil4 testbenches.

### 26.2.44.3.4 Burst Tracking Registers

#### 26.2.4.1 Required State

The listing below is the generic PATTERN, not the RTL's register names: the real modules split this state between separate read and write converters and count UP (r\_ar\_beat\_count/r\_aw\_beat\_count against a stored length) rather than down — grep for those names, not these.

```
// Burst tracking registers (generic pattern)
logic [ADDR_WIDTH-1:0] r_base_addr;
logic [ADDR_WIDTH-1:0] r_current_addr;
logic [7:0] r_original_len;
logic [7:0] r_remaining_beats;
logic [2:0] r_size;
logic [1:0] r_burst;
logic [ID_WIDTH-1:0] r_id;
logic r_is_write;
```

#### 26.2.4.2 Initialization

```
always_ff @(posedge clk) begin
 if (accept_new_transaction) begin
 r_base_addr <= s_axaddr;
 r_current_addr <= s_axaddr;
 r_original_len <= s_axlen;
 r_remaining_beats <= s_axlen;
 end
end
```

```

 r_size <= s_axsize;
 r_burst <= s_axburst;
 r_id <= s_axid;
 r_is_write <= is_write_transaction;
 end else if (beat_complete) begin
 r_current_addr <= next_address(...);
 r_remaining_beats <= r_remaining_beats - 1;
 end
end

```

## 26.3 Timing

### 26.3.1 4.3.5 Timing Impact

#### 26.3.1.1 *Decomposition Overhead*

*Table 4.6: Decomposition Overhead*

| Transaction Type    | Overhead                                                                                            |
|---------------------|-----------------------------------------------------------------------------------------------------|
| Single-beat         | 0 cycles (passthrough)                                                                              |
| N-beat burst        | none per beat – requests stream at the downstream accept rate, independent of responses (see 4.2.4) |
| Back-to-back bursts | serialized by the one-outstanding-burst guard                                                       |

The old “2 cycles per beat” figure described a request-response lockstep the RTL does not have; nothing in the address path waits on responses.


#### 26.3.1.2 *Pipeline Considerations*

Decomposed requests are independent of responses (Table 4.6 above): - ARs issue at the downstream accept rate — nothing in the AR path samples `m_axil_rvalid` - AW/W beats advance on their own handshakes; intermediate Bs are absorbed (`m_axil_bready = s_axi_bready || !w_b_all_beats_done`) without stalling issuance

Against a pipelining AXIL4 slave, an N-beat burst costs N issue cycles plus one response latency — not N round trips. The only serialization the converter imposes is between whole bursts (the one-outstanding guard).

## 26.4 Waveforms

### 26.4.1 Figure 4.5: Width Burst Conversion



#### Width Burst Conversion

## 26.5 Navigation

---

**Next:** [Chapter 5: Verification](#)

[RTL Design Sherpa](#) · [Learning Hardware Design Through Practice](#) [GitHub](#) · [Documentation Index](#) · [MIT License](#)

---

## 27 5.1 Test Strategy

### 27.1 Overview

---

How the converter modules get verified, from unit tests to system stress: three test levels, a coverage model, and the exact commands to run it all.

## 27.2 Testing

---

### 27.2.15.1.1 Test Organization

```
projects/components/converters/dv/tests/
├── width_converters/
│ ├── test_axi_data_upsize.py
│ ├── test_axi_data_dsize.py
│ ├── test_axi4_dwidth_converter_wr.py
│ └── test_axi4_dwidth_converter_rd.py
├── protocol_converters/
│ ├── test_axi4_to_axil4_rd.py
│ ├── test_axi4_to_axil4_wr.py
│ ├── test_axil4_to_axi4_rd.py
│ ├── test_axil4_to_axi4_wr.py
│ ├── test_axi4_to_apb4.py
│ └── test_peakrdl_adapter.py
└── integration/
 ├── test_width_protocol_chain.py
 └── test_full_system.py
```

### 27.2.25.1.2 Test Levels

#### 27.2.2.1 *Level 1: Unit Tests*

**Purpose:** Verify individual module functionality in isolation.

**Coverage:** - All parameters combinations - Edge cases (min/max values) - Error injection

**Example:**

```
@pytest.mark.parametrize("narrow_width,wide_width", [
 (32, 64),
 (32, 128),
 (64, 256),
 (64, 512),
 (128, 1024),
])
async def test_upsize_ratios(dut, narrow_width, wide_width):
 """Test various width ratios."""
 tb = UpsizeTB(dut, narrow_width, wide_width)
 await tb.reset()
 await tb.run_basic_transfer(count=100)
 assert tb.scoreboard.check_passed()
```

#### 27.2.2.2 *Level 2: Integration Tests*

**Purpose:** Verify module combinations work together.

**Coverage:** - Width converter + protocol converter chains - Back-to-back converters - Mixed traffic patterns

### 27.2.2.3 *Level 3: System Tests*

**Purpose:** Verify in realistic system context.

**Coverage:** - CPU-to-DDR paths - Peripheral access paths - Full bandwidth stress

### 27.2.35.1.3 Test Categories

#### 27.2.3.1 *Functional Tests*

*Table 5.1: Functional Test Categories*

| Category | Description             | Example                     |
|----------|-------------------------|-----------------------------|
| Basic    | Single transactions     | Single read, single write   |
| Burst    | Multi-beat transactions | INCR burst, WRAP burst      |
| Mixed    | Interleaved R/W         | Read-modify-write sequences |
| Edge     | Boundary conditions     | Max burst length, min width |

#### 27.2.3.2 *Stress Tests*

*Table 5.2: Stress Test Categories*

| Category     | Description             | Duration              |
|--------------|-------------------------|-----------------------|
| Throughput   | Maximum bandwidth       | 10,000 transactions   |
| Backpressure | Ready signal variations | Random delays         |
| Reset        | Reset during operation  | Mid-transaction reset |

#### 27.2.3.3 *Error Tests*

*Table 5.3: Error Test Categories*


| Category | Description           | Expected Behavior |
|----------|-----------------------|-------------------|
| SLVERR   | Slave error injection | Error propagation |
| DECERR   | Decode error          | Error response    |

| Category | Description      | Expected Behavior |
|----------|------------------|-------------------|
| Timeout  | Response timeout | Error handling    |

## 27.2.45.1.4 Testbench Architecture

### 27.2.4.1 Components

#### Testbench



### 27.2.4.2 Driver Implementation

```

class AXI4MasterDriver:
 def __init__(self, dut, clock, prefix="s_axi"):
 self.dut = dut
 self.clock = clock
 self.prefix = prefix

 async def write(self, addr, data, burst_len=0):
 """Issue AXI4 write transaction."""
 # AW phase
 # the DUT ports are prefixed (s_axi_awvalid, ...) -- resolve
 # through the stored prefix or every access AttributeError
 sig = lambda name: getattr(self.dut, f"{self.prefix}_{name}")
 sig("awvalid").value = 1
 sig("awaddr").value = addr
 sig("awlen").value = burst_len
 await self._wait_ready("awready")

 # W phase
 for i, d in enumerate(data):
 sig("wvalid").value = 1
 sig("wdata").value = d
 sig("wlast").value = (i == len(data) - 1)

```

```

 await self._wait_ready("wready")

 # B phase
 await self._wait_valid("bvalid")
 return sig("bresp").value

```

### 27.2.55.1.5 Coverage Model

#### 27.2.5.1 Functional Coverage

```

@coverage
class ConverterCoverage:
 # Width ratio coverage
 ratio_cp = coverpoint(
 lambda: self.width_ratio,
 bins=[2, 4, 8, 16]
)

 # Burst length coverage
 burst_len_cp = coverpoint(
 lambda: self.burst_len,
 bins=list(range(0, 256, 16)) + [255]
)

 # Burst type coverage
 burst_type_cp = coverpoint(
 lambda: self.burst_type,
 bins=["FIXED", "INCR", "WRAP"]
)

 # Cross coverage
 ratio_x_burst = cross(ratio_cp, burst_len_cp)

```

#### 27.2.5.2 Code Coverage

**Target:** >95% line coverage, >90% branch coverage

Table 5.4: Coverage Targets

| Module            | Line | Branch | FSM                      |
|-------------------|------|--------|--------------------------|
| axi_data_upsizer  | 98%  | 95%    | n/a (no FSM – see 4.1.1) |
| axi_data_dnsizer  | 97%  | 93%    | n/a (no FSM – see 4.1.2) |
| axi4_to_axil4_rid | 96%  | 91%    | 100%                     |

| Module               | Line | Branch | FSM  |
|----------------------|------|--------|------|
| axi4_to_axil4_<br>wr | 95%  | 90%    | 100% |

## 27.2.65.1.6 Test Execution

### 27.2.6.1 *Running Tests*

*# Run all converter tests*

```
cd projects/components/converters/dv/tests
pytest -v
```

*# Run specific module tests*

```
pytest test_axi_data_upsize.py -v
```

*# Run with coverage*

```
pytest --cov=projects/components/converters/rtl -v
```

*# Run with waveform dump*

```
pytest test_axi_data_upsize.py -v --waves
```

### 27.2.6.2 *CI/CD Integration*

```
converter_tests:
 stage: test
 script:
 - cd projects/components/converters/dv/tests
 - pytest -v --junitxml=results.xml
 - pytest --cov=rtl --cov-report=html
 artifacts:
 reports:
 junit: results.xml
 paths:
 - htmlcov/
```

## 27.3 Navigation

---

Next: [Debug Guide](#)

[RTL Design Sherpa](#) · [Learning Hardware Design Through Practice](#) [GitHub](#) · [Documentation Index](#) · [MIT License](#)

---

## 28 5.2 Debug Guide

### 28.1 Overview

---

Start here when a converter misbehaves. Each entry runs symptoms, probe points, then fix — the order you want them in at the bench.

## 28.2 Testing

---

### 28.2.1 5.2.1 Common Issues

#### 28.2.1.1 Width Converter Issues

##### 28.2.1.1.1 Issue: Data Corruption

**Symptoms:** - Output data doesn't match expected - Random bits flipped or missing

**Debug Steps:** 1. Check width ratio calculation:  $RATIO = WIDE\_WIDTH / NARROW\_WIDTH$  2. Verify beat counter is correct:  $\$clog2(RATIO)$  bits 3. Check data packing/unpacking slice indices 4. Verify sideband mode matches use case

**Waveform Checkpoints:**

|                                 |                                                |
|---------------------------------|------------------------------------------------|
| <code>r_beat_ptr</code>         | - slot/slice pointer, cycles 0 to $RATIO-1$    |
| <code>r_data_accumulator</code> | - (usize) each slot fills as narrow beats land |
| <code>r_data_buffer</code>      | - (usize) captured wide beat being sliced      |
| <code>narrow_/wide_ data</code> | - compare input/output patterns                |

##### 28.2.1.1.2 Issue: LAST Signal Incorrect

**Symptoms:** - Transaction ends early or late - Master receives wrong beat count

**Debug Steps:** 1. Check `TRACK_BURSTS` parameter (`axi_data_dsize`); there is no `USE_LAST` 2. Verify burst tracker logic if enabled 3. Check `s_last` input timing

**Solution:**

```
// LAST is the buffered last flag AND the final beat -- from the RTL:
// usize: wide_last = r_last_buffered && r_wide_valid;
// dsize: narrow_last = r_wide_buffered && r_last_buffered
// && w_last_narrow_beat;
// An OR of (final beat, saw-last) asserts LAST at the end of EVERY
// group and truncates any burst spanning more than one wide beat.
```

##### 28.2.1.1.3 Issue: Throughput Lower Than Expected

**Symptoms:** - Gaps between output beats - Single-buffer stalling between wide beats (it should replace during the last narrow beat)

**Debug Steps:** 1. Check downstream ready signal behavior 2. Look for backpressure stalls

#### 28.2.1.2 Protocol Converter Issues

##### 28.2.1.2.1 Issue: Burst Decomposition Incorrect

**Symptoms:** - Wrong number of single transactions - Address increment wrong

**Debug Steps:** 1. Check burst type (FIXED, INCR, WRAP) 2. Verify size calculation 3. Check address increment logic

**Waveform Checkpoints:**

r\_ar\_addr / r\_aw\_addr - walk by (1 << size) per issued beat  
r\_ar\_beat\_count / r\_aw\_beat\_count - count to arlen/awlen  
m\_axil\_arvalid - asserts for each decomposed beat

### 28.2.1.2.2 Issue: Response Aggregation Wrong

**Symptoms:** - Wrong BRESP/RRESP value - Error not propagated correctly

**Debug Steps:** 1. Check worst-case response tracking 2. Verify response counter 3. Check RLAST generation

**Solution:**

```
// Accumulate per beat...
always_ff @(posedge aclk or negedge aresetn) begin
 ...
 else if (m_axil_rvalid && m_axil_rready)
 if (m_axil_rresp > r_r_resp_accum)
 r_r_resp_accum <= m_axil_rresp;
end

// ...and ALWAYS fold the live beat in at emission. The accumulator
// alone is the shipped bug (see 4.3.3): it does not yet contain the
// final beat's own response, so a single-beat error reports OKAY.
assign w_r_resp_worst = (m_axil_rresp > r_r_resp_accum) ? m_axil_rresp
 :
 r_r_resp_accum;
assign s_axi_rresp = (r_r_beat_count == r_r_len) ? w_r_resp_worst
 : m_axil_rresp;
```

## 28.2.25.2 Debug Signals

### 28.2.2.1 Recommended Internal Signals

For width converters:

```
// Add debug outputs
output logic [$clog2(RATIO)-1:0] dbg_beat_count,
output logic dbg_buffer_valid,
output logic [1:0] dbg_state
```

For protocol converters:

```
// Add debug outputs
output logic [7:0] dbg_remaining_beats,
```

```

output logic [2:0] dbg_state,
output logic [1:0] dbg_worst_resp,
output logic dbg_in_burst

```

### 28.2.2.2 ILA Configuration

*# Create ILA for converter debug*

```
create_debug_core u_ila ila
```

*# Add probes*

```
set_property probe_count 10 [get_debug_cores u_ila]
connect_debug_port u_ila/clk [get_nets aclk]
```

*# Key signals*

```
connect_debug_port u_ila/probe0 [get_nets r_beat_ptr]
connect_debug_port u_ila/probe1 [get_nets r_ar_beat_count]
connect_debug_port u_ila/probe2 [get_nets s_axi_arvalid]
connect_debug_port u_ila/probe3 [get_nets m_axil_arvalid]
```

## 28.2.35.2.3 Simulation Debug

### 28.2.3.1 Waveform Analysis

**Key Signal Groups** (real port names — the width primitives are named by WIDTH, not direction: narrow\_\* is the INPUT of the upsize but the OUTPUT of the dnsize):

#### 1. Narrow side:

- narrow\_valid, narrow\_ready, narrow\_data, narrow\_sideband, narrow\_last (upsized input / dnsized output)

#### 2. Wide side:

- wide\_valid, wide\_ready, wide\_data, wide\_sideband, wide\_last (upsized output / dnsized input)


#### 3. Control:

- r\_beat\_ptr, r\_slave\_beat\_count, r\_burst\_active (tracked mode), start\_lane, burst\_len/burst\_start (dnsized tracked mode)

#### 4. Sideband contents:

- WSTRB rides narrow\_/wide\_sideband on the write path
- RRESP rides narrow\_/wide\_sideband on the read path

### 28.2.3.2 Timing Diagram Template



```

s_data
s_ready _____

m_valid __|_____|____
 [WIDE_DATA]
m_data

```

Check:

1. s\_ready stays high during accumulation
2. m\_valid asserts after RATIO beats
3. Data packing is correct

## 28.2.45.2.4 Common Mistakes

### 28.2.4.1 Mistake 1: Wrong Width Ratio

```

// WRONG: Manual ratio
localparam RATIO = 8; // May not match actual widths

// CORRECT: Calculated ratio
localparam RATIO = WIDE_WIDTH / NARROW_WIDTH;

```

### 28.2.4.2 Mistake 2: Missing Sideband Handling

```

// WRONG: Forgetting sideband
assign m_data = r_data;
// Missing: assign m_wstrb = ...

```

```

// CORRECT: Handle both
assign m_data = r_data;
assign m_wstrb = r_sideband;

```

### 28.2.4.3 Mistake 3: Incorrect LAST Timing

```

// WRONG: LAST on wrong beat
assign narrow_last = (r_beat_ptr == 0); // First slice!

// CORRECT: LAST on final beat
// AND of final-beat with the buffered last flag (see 5.2.1); the OR
// form truncates multi-wide-beat bursts
assign narrow_last = r_wide_buffered && r_last_buffered &&
w_last_narrow_beat;

```

### 28.2.4.4 Mistake 4: Burst Length Calculation Error

```

// WRONG: Off-by-one
assign m_awlen = s_awlen / RATIO; // Wrong!

```

```
// CORRECT: Account for LEN encoding
assign m_axi_awlen = ((int_awlen + 8'(WIDTH_RATIO)) / 8'(WIDTH_RATIO))
- 8'd1; // round UP
```

### 28.2.55.2.5 Verification Checklist

Before signoff, verify:

- ☐ All parameter combinations tested
- ☐ Single-beat transactions work
- ☐ Multi-beat bursts work (INCR, WRAP, FIXED)
- ☐ Backpressure handling correct
- ☐ Error propagation correct
- ☐ LAST signal timing correct
- ☐ Sideband signals handled correctly
- ☐ Reset behavior verified
- ☐ Coverage targets met

### 28.2.65.2.6 Performance Validation

#### 28.2.6.1 *Throughput Measurement*

```
async def measure_throughput(tb, transaction_count=1000):
 start_time = get_sim_time()

 for _ in range(transaction_count):
 await tb.send_transaction()

 end_time = get_sim_time()
 elapsed_cycles = (end_time - start_time) / clock_period

 throughput = transaction_count / elapsed_cycles
 print(f"Throughput: {throughput:.2f} transactions/cycle")

 return throughput
```

#### 28.2.6.2 *Expected Throughput*

Table 5.5: Expected Throughput

| Module          | Mode   | Expected                               |
|-----------------|--------|----------------------------------------|
| axi_data_upsize | Single | 1.0 trans/cycle                        |
| axi_data_dnsz   | Single | 0.992 narrow beats/cycle<br>(measured) |

| Module        | Mode        | Expected                                                                                                                                           |
|---------------|-------------|----------------------------------------------------------------------------------------------------------------------------------------------------|
| axi4_to_axil4 | Single-beat | ~1/(2+slave latency) trans/cycle (one-outstanding guard)                                                                                           |
| axi4_to_axil4 | Burst       | ~N/(N+2) beats/cycle vs a pipelining slave (N issue cycles + one response latency, see 4.2.4); whole bursts serialize on the one-outstanding guard |

## 28.3 Navigation

---

**End of Micro-Architecture Specification**