

Address Increment Patterns for DMA and Descriptor-Based Data Movement

Document Version: 1.0 **Last Updated:** 2026-01-13 **Scope:** Comprehensive analysis of address increment patterns for descriptor-based systems **Reference Implementation:** RAPIDS (Rapid AXI Programmable In-band Descriptor System)

Table of Contents

1. Executive Summary
 2. RAPIDS Current Implementation
 3. AXI Protocol Burst Modes
 4. Linear Address Increment Patterns
 5. 2D Transfer Patterns (Stride-Based)
 6. 3D Transfer Patterns
 7. Scatter-Gather Patterns
 8. Circular and Ring Buffer Patterns
 9. Tensor Memory Access Patterns
 10. Implementation Comparison Matrix
 11. RAPIDS Extension Recommendations
 12. References
-

1. Executive Summary

Address increment patterns determine how DMA engines and descriptor-based data movement systems calculate memory addresses during transfers. The choice of pattern significantly impacts:

- **Bandwidth efficiency:** Alignment and burst optimization
- **Application fit:** Different workloads require different patterns
- **Hardware complexity:** More patterns = more control logic
- **Software flexibility:** Richer patterns = simpler programming model

Pattern Categories

Category	Description	Use Cases
Linear/Contiguous	Sequential address increment	Block copy, streaming data
2D Stride	Row-major with inter-row gaps	Image processing, video frames
3D Volume	Multi-plane with inter-plane gaps	3D imaging, tensor operations
Scatter-Gather	Non-contiguous fragment list	Protocol buffers, fragmented memory
Circular	Wrap-around at boundary	Audio streaming, network buffers
Tensor	Multi-dimensional with strides	Neural network accelerators

2. RAPIDS Current Implementation

2.1 Descriptor Format (256-bit)

RAPIDS uses a fixed 256-bit descriptor structure defined in `rapids_pkg.sv`:

Bit Field	Width	Description
[63:0]	64-bit	src_addr - Source address (byte-aligned)
[127:64]	64-bit	dst_addr - Destination address (byte-aligned)
[159:128]	32-bit	length - Transfer length in BEATS (not bytes)
[191:160]	32-bit	next_descriptor_ptr - Next descriptor address
[207:200]	8-bit	desc_priority - Transfer priority
[199:196]	4-bit	channel_id - Channel ID
[195]	1-bit	error - Error flag
[194]	1-bit	last - Last descriptor in chain
[193]	1-bit	gen_irq - Generate interrupt
[192]	1-bit	valid - Valid descriptor
[255:208]	48-bit	reserved - Future use

2.2 Current Address Increment Modes

Mode 1: Beat-Based Linear (Default)

Address Calculation:

$$\text{next_addr} = \text{current_addr} + (\text{beat_count} \times 64 \text{ bytes})$$

Where: RAPIDS_BYTES_PER_BEAT = 64 bytes (512-bit data width)

- All transfers aligned to 64-byte boundaries for optimal throughput
- Descriptor length field specifies count in beats
- Software converts byte count to beats: beats = (bytes + 63)
 $\gg 6$

Mode 2: Chunk-Based Alignment (Sub-Beat Transfers)

For alignment/final phases handling partial beats:

Address Calculation for Alignment Phase:

```
offset = address[5:0]           // 6-bit offset (0-63)
bytes_to_boundary = (offset == 0) ? 0 : 64 - offset
```

Transfer Phases:

Phase 1 (Alignment): Transfer bytes_to_boundary to reach 64B boundary
Phase 2 (Streaming): Transfer full 64B beats
Phase 3 (Final): Transfer remaining < 64 bytes

2.3 Alignment Information Structure

```
typedef struct packed {
    logic [15:0] first_chunk_enables;      // Phase 1 chunk mask
    logic [15:0] streaming_chunk_enables; // Phase 2 (always 0xFFFF)
    logic [15:0] final_chunk_enables;       // Phase 3 chunk mask
    logic [31:0] first_transfer_bytes;     // Phase 1 size
    logic [31:0] streaming_bytes;          // Phase 2 size
    logic [31:0] final_transfer_bytes;     // Phase 3 size
    logic [3:0]  first_start_chunk;        // Starting chunk (0-15)
    logic [3:0]  first_valid_chunks;       // Chunks in phase 1
    logic [3:0]  final_valid_chunks;        // Chunks in phase 3
} alignment_info_t;
```

2.4 AXI Translation

Transfer Phase	AxSIZE	AxBURST	Address Increment
Alignment	3'b010	INCR	4 bytes × (AxLEN+1)
Streaming	3'b110	INCR	64 bytes × (AxLEN+1)
Final	3'b010	INCR	4 bytes × (AxLEN+1)

3. AXI Protocol Burst Modes

The AXI protocol defines three burst types that constrain address generation:

3.1 FIXED Burst (AxBURST = 2'b00)

Address Formula:

```
beat[n].addr = AxADDR    (constant for all beats)
```

Use Cases: - FIFO access (single address register) - Memory-mapped I/O ports - Audio codec sample registers

Constraints: - All beats transfer to same address - Cannot cross 4KB boundary (trivially satisfied)

3.2 INCREMENT Burst (AxBURST = 2'b01)

Address Formula:

```
beat[n].addr = AxADDR + n × (1 << AxSIZE)
```

Example (AxSIZE=6, 64-byte):

```
beat[0] = 0x1000  
beat[1] = 0x1040  
beat[2] = 0x1080  
...
```

Use Cases: - Linear memory access - Block copy operations - Stream data to/from memory

Constraints: - Cannot cross 4KB boundary - Maximum burst length varies by protocol version

3.3 WRAP Burst (AxBURST = 2'b10)

Address Formula:

```
wrap_boundary = AxADDR & ~((AxLEN+1) × (1<<AxSIZE) - 1)  
wrap_size = (AxLEN+1) × (1 << AxSIZE)
```

```
beat[n].addr = wrap_boundary + ((AxADDR + n×(1<<AxSIZE)) mod wrap_size)
```

Example (AxADDR=0x1020, AxLEN=3, AxSIZE=4):

```
Wrap boundary: 0x1000  
Wrap size: 64 bytes (4 × 16)  
beat[0] = 0x1020  
beat[1] = 0x1030  
beat[2] = 0x1000  (wrapped)  
beat[3] = 0x1010
```

Use Cases: - Cache line fills (critical-word-first) - Circular buffer access

Constraints: - AxLEN must be 1, 3, 7, or 15 - Starting address must be aligned to transfer size

4. Linear Address Increment Patterns

4.1 Simple Contiguous

The most basic pattern - sequential addresses:

Descriptor Fields:

```
src_addr: Starting source address  
dst_addr: Starting destination address  
length: Total transfer size
```

Address Calculation:

```
for each beat b in [0, length/beat_size):  
    src = src_addr + b × beat_size  
    dst = dst_addr + b × beat_size
```

RAPIDS Implementation: Native support via beat-based transfers.

4.2 Unaligned Start/End

Handling transfers that don't start or end on natural boundaries:

Example: Transfer 200 bytes from 0x1040

Phase 1 (Align): Transfer 64 bytes (0x1040-0x107F) to reach 0x1080
Phase 2 (Stream): Transfer 128 bytes (0x1080-0x10FF) in 2 full beats
Phase 3 (Final): Transfer 8 bytes (0x1100-0x1107)

RAPIDS Implementation: Native support via alignment_info_t structure.

4.3 Size-Constrained Bursts

Breaking large transfers into AXI-compliant bursts:

Constraints:

- AXI4: Max 256 beats per burst
- No 4KB boundary crossing

Address Calculation:

```
remaining = total_length  
current_addr = start_addr
```

```

while (remaining > 0):
    burst_len = min(remaining, 256×beat_size)
        burst_len = min(burst_len, next_4kb_boundary - current_addr)
    issue_burst(current_addr, burst_len)
    current_addr += burst_len
    remaining -= burst_len

```

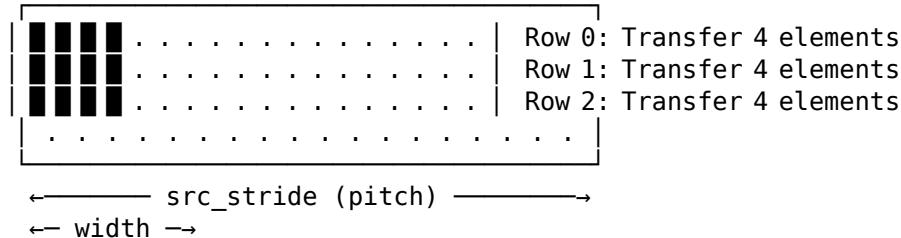
RAPIDS Implementation: Handled by AXI engine burst splitting.

5. 2D Transfer Patterns (Stride-Based)

5.1 Concept

2D transfers move rectangular regions where source and destination may have different memory layouts (pitches/strides).

Memory Layout:



5.2 Descriptor Fields (Extended)

2D Descriptor Format:

src_addr:	Starting source address (top-left corner)
dst_addr:	Starting destination address
x_length:	Bytes per row (width of transfer region)
y_length:	Number of rows
src_stride:	Source bytes between row starts
dst_stride:	Destination bytes between row starts

5.3 Address Calculation

```

for row in [0, y_length):
    row_src = src_addr + row × src_stride
    row_dst = dst_addr + row × dst_stride
    transfer(row_src, row_dst, x_length)

```

5.4 Use Cases

Image Processing:

Example: Extract 640x480 ROI from 1920x1080 frame (32-bit pixels)

```
src_addr: frame_base + (y_offset × 1920 + x_offset) × 4
dst_addr: roi_buffer
x_length: 640 × 4 = 2560 bytes
y_length: 480
src_stride: 1920 × 4 = 7680 bytes
dst_stride: 640 × 4 = 2560 bytes (packed destination)
```

Video Frame Tiling:

Example: Copy 16×16 macroblock

```
x_length: 16 × bytes_per_pixel
y_length: 16
src_stride: frame_width × bytes_per_pixel
dst_stride: 16 × bytes_per_pixel
```

5.5 Analog Devices AXI DMAC Implementation

Reference implementation from ADI's high-speed DMA controller:

Registers:

```
X_LENGTH (0x418): Row width - 1
Y_LENGTH (0x41C): Row count - 1
SRC_STRIDE (0x424): Source row spacing
DEST_STRIDE (0x420): Destination row spacing
```

Address Formulas:

```
ROW_SRC_ADDRESS = SRC_ADDRESS + SRC_STRIDE × N
ROW_DEST_ADDRESS = DEST_ADDRESS + DEST_STRIDE × N
```

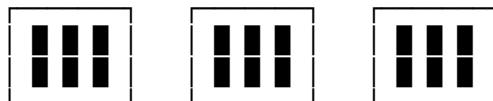
6. 3D Transfer Patterns

6.1 Concept

3D transfers extend 2D with an additional depth/plane dimension:

Memory Layout (Z planes stacked):

Plane 0: Plane 1: Plane 2:



6.2 Descriptor Fields (Extended)

3D Descriptor Format:

```
src_addr: Starting address (origin corner)
dst_addr: Destination address
x_length: Width in bytes
y_length: Height in rows
z_length: Depth in planes
src_stride_x: Source row stride (unused, implicit x_length)
src_stride_y: Source row-to-row stride
src_stride_z: Source plane-to-plane stride
dst_stride_y: Destination row stride
dst_stride_z: Destination plane stride
```

6.3 Address Calculation

```
for plane in [0, z_length):
    for row in [0, y_length):
        src = src_addr + planexsrc_stride_z + row×src_stride_y
        dst = dst_addr + planedst_stride_z + row×dst_stride_y
        transfer(src, dst, x_length)
```

6.4 Use Cases

3D Medical Imaging:

Example: Extract 64×64×64 VOI from 512×512×256 CT volume

```
src_addr: volume_base + z_off×512×512 + y_off×512 + x_off
x_length: 64 bytes
y_length: 64
z_length: 64
src_stride_y: 512
src_stride_z: 512 × 512
```

Neural Network Tensors:

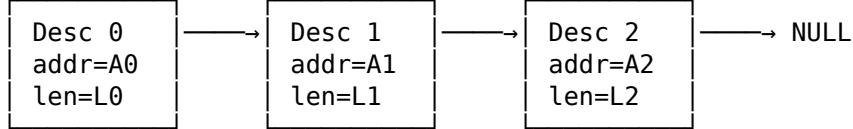
Example: Move 4D tensor slice [batch, channel, H, W]
Requires nested 3D operations or generalized tensor addressing

7. Scatter-Gather Patterns

7.1 Concept

Scatter-Gather (SG) handles non-contiguous memory regions using linked lists of descriptors:

Descriptor Chain:



7.2 SG Descriptor Fields

Scatter-Gather Descriptor:

buffer_addr:	Address of this fragment
buffer_length:	Size of this fragment
next_desc_ptr:	Address of next descriptor (0 = last)
control_flags:	Completion interrupt, error handling, etc.
status:	Completion status (written by DMA)

7.3 Operation Modes

Gather (Multiple Source -> Single Destination):

Source fragments: [A0:L0], [A1:L1], [A2:L2]

Destination: Contiguous buffer B

Result: B = concat(mem[A0:A0+L0], mem[A1:A1+L1], mem[A2:A2+L2])

Scatter (Single Source -> Multiple Destinations):

Source: Contiguous buffer B of length L0+L1+L2

Destination fragments: [A0:L0], [A1:L1], [A2:L2]

Result: mem[A0:A0+L0] = B[0:L0]
mem[A1:A1+L1] = B[L0:L0+L1]
mem[A2:A2+L2] = B[L0+L1:end]

7.4 AMD/Xilinx AXI DMA SG Descriptor Format

Standard Descriptor (32-bit addressing):

Offset	Field	Description
0x00	NXTDESC	Next descriptor pointer
0x04	Reserved	
0x08	BUFFER_ADDRESS	Data buffer pointer
0x0C	Reserved	
0x10	Reserved	
0x14	Reserved	
0x18	CONTROL	Transfer length, SOF, EOF
0x1C	STATUS	Completion status

7.5 Use Cases

Network Packet Assembly:

Example: Assemble TCP packet from scattered headers and payload

```
Desc 0: Ethernet header (14 bytes at addr_eth)
Desc 1: IP header (20 bytes at addr_ip)
Desc 2: TCP header (20 bytes at addr_tcp)
Desc 3: Payload (1460 bytes at addr_payload)
```

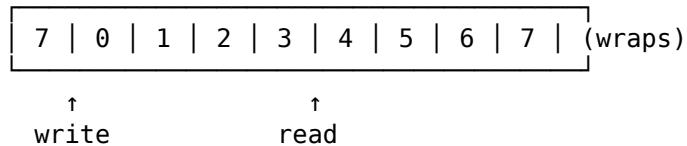
RAPIDS Implementation: Native support via next_descriptor_ptr chaining.

8. Circular and Ring Buffer Patterns

8.1 Concept

Circular buffers wrap addresses at a boundary, enabling continuous streaming without descriptor reload:

Memory Layout:



8.2 Descriptor Fields (Extended)

Circular Buffer Descriptor:

```
base_addr: Buffer base address
buffer_size: Total buffer size (must be power of 2 for efficient wrap)
current_ptr: Current read/write position
wrap_enable: Enable address wrapping
```

8.3 Address Calculation

Efficient wrap (power-of-2 size):

```
next_addr = base_addr + ((current_ptr + increment) & (buffer_size - 1))
```

General wrap:

```
offset = (current_ptr + increment) % buffer_size
next_addr = base_addr + offset
```

8.4 Ping-Pong Double Buffering

Common implementation using two descriptors pointing to each other:

Descriptor A:	Descriptor B:
buffer_addr = BUF_A	buffer_addr = BUF_B
next_desc = &Desc_B	next_desc = &Desc_A

Operation:

```
While DMA processes BUF_A, CPU fills BUF_B
On completion, DMA auto-loads Desc_B, starts BUF_B
CPU processes BUF_A while DMA runs BUF_B
(repeat)
```

8.5 Use Cases

Audio Streaming:

Example: 48kHz stereo audio, 20ms buffers
buffer_size: $48000 \times 2 \text{ channels} \times 2 \text{ bytes} \times 0.02\text{s} = 3840 \text{ bytes}$
Ring with 4 buffers: 15360 bytes total
Wrap at: base + 15360

Network RX Ring:

Example: Ethernet receive ring
descriptor_count: 256 (power of 2)
Each descriptor points to packet buffer
Last descriptor.next points to first descriptor
Hardware auto-advances through ring

9. Tensor Accelerator Memory Access Patterns (PRIORITY 0)

Tensor accelerators for neural network inference represent the most demanding and sophisticated address generation requirements. This section provides comprehensive coverage of industry-standard approaches.

9.1 Systolic Array Dataflow Fundamentals

Systolic arrays (used in Google TPU, NVIDIA Tensor Cores, etc.) require coordinated data movement across three matrices: inputs (activations), weights, and outputs (partial sums).

Three Primary Dataflows:

Dataflow	Stationary Data	Moving Data	Best For
Weight Stationary (WS)	Weights preloaded in PEs	Inputs + partial sums flow	Weight reuse (many inputs)
Input Stationary (IS)	Inputs held in PEs	Weights + partial sums flow	Input reuse (many kernels)
Output Stationary (OS)	Partial sums accumulate in PEs	Inputs + weights flow	Output accumulation

Address Pattern Implications:

Weight Stationary:

- Weights: Load once per layer, stream through array
- Inputs: Stream row-by-row, reuse across weight tiles
- Address pattern: 2D tiled with weight-tile-aligned boundaries

Output Stationary:

- Outputs: Accumulate in place
- Inputs/Weights: Double-buffered streaming
- Address pattern: Output-tile-centric addressing

9.2 Data Cube Model (NVDLA Reference)

NVIDIA's Deep Learning Accelerator (NVDLA) defines tensor addressing using a "data cube" model with three key stride parameters:

Data Cube Structure:

Dimensions: Width (W) × Height (H) × Channels (C)

Memory Layout:

- Data divided into 1×1×32byte "atom cubes"
- Scanning order: C'(32B) → W → H → C (surfaces)
- C' changes fastest (innermost loop)

Address Calculation:

```
element_addr = base_addr
              + surface_index × surface_stride
              + line_index × line_stride
              + element_offset_within_line
```

Stride Parameters:

Parameter	Description	Alignment
line_stride	Bytes from line N to line N+1	32 bytes
surface_stride	Bytes from surface N to surface N+1	32 bytes
base_addr	Starting address	32 bytes

NVDLA Register Configuration:

```
D_DATAIN_FORMAT      : Data format selection
D_DATAIN_SIZE_0     : Width x Height
D_DATAIN_SIZE_1     : Channels
D_LINE_STRIDE       : Line-to-line spacing (0x5040)
D_SURF_STRIDE       : Surface-to-surface spacing
```

Example:

```
Width=26, Height=32, Channels=3
line_stride = 0x1A0 (aligned width)
surface_stride = 0x1520 (line_stride x height)
```

9.3 Tensor Memory Accelerator (TMA) - NVIDIA Hopper

NVIDIA's Hopper architecture introduces hardware TMA for efficient multi-dimensional tensor transfers:

TMA Descriptor Format:

CUtensorMap structure (64-bit packed descriptor):

- Global memory base address
- Tensor rank (1D to 5D supported)
- Dimension sizes and strides
- Swizzle pattern for bank conflict avoidance
- Box dimensions (tile size for transfer)

Key Constraints:

Alignment Requirements:

- Contiguous dimension must have stride=1
- All other strides must be multiples of 16 bytes
- For float tensors [M,N] with stride [N,1]: $N \% 4 == 0$

Tile Quantization:

- TPU/GPU systolic arrays process 128x128 or 128x8 tiles
- Tensors must align to tile boundaries to avoid padding waste
- Performance penalty for partial tiles (idle MACs)

Address Generation (Hardware):

TMA eliminates per-thread address calculation:

Traditional: Each thread computes $\text{addr} = \text{base} + \text{thread_id} \times \text{stride}$

TMA: Single descriptor, hardware generates all addresses

Coordinate-based access:

$\text{ArithTuple}(\text{row}, \text{col}) \rightarrow \text{Hardware translates to physical address}$

Automatic out-of-bounds predication (no manual boundary checks)

9.4 Convolution Address Patterns

9.4.1 Direct Convolution

Input Feature Map: $[N, C_{in}, H_{in}, W_{in}]$
Weight Kernel: $[C_{out}, C_{in}, K_h, K_w]$
Output Feature Map: $[N, C_{out}, H_{out}, W_{out}]$

Output position ($n, c_{out}, h_{out}, w_{out}$):

```
For each (c_in, k_h, k_w):
    input_addr = base_in + n×(C_in×H_in×W_in)
                  + c_in×(H_in×W_in)
                  + (h_out×stride_h + k_h)×W_in
                  + (w_out×stride_w + k_w)

    weight_addr = base_w + c_out×(C_in×K_h×K_w)
                  + c_in×(K_h×K_w)
                  + k_h×K_w
                  + k_w
```

9.4.2 Im2Col Transformation Converts convolution to GEMM for systolic array efficiency:

Original Convolution:

```
Output[h,w] = sum(Input[h+kh, w+kw] × Weight[kh, kw])
```

Im2Col Transformation:

- Unroll each receptive field into a column
- Weight kernels become rows
- Convolution becomes matrix multiplication

Address Pattern for Im2Col:

```
For output position (h_out, w_out):
    col_index = h_out × W_out + w_out
    For kernel position (c_in, k_h, k_w):
        row_index = c_in × K_h × K_w + k_h × K_w + k_w
        src_addr = base + c_in×(H×W) + (h_out×s_h+k_h)×W + (w_out×s_w+k_w)
        dst_addr = im2col_base + row_index × num_output_positions + col_index
```

Data Expansion Factor:

Naive: $K_h \times K_w \times$ data replication

Hardware Im2Col: Generate addresses on-the-fly, no physical expansion

9.4.3 Strided and Dilated Convolution

Strided Convolution ($\text{stride} > 1$):

$\text{input}_h = \text{output}_h \times \text{stride} + \text{kernel}_h$

$\text{input}_w = \text{output}_w \times \text{stride} + \text{kernel}_w$

Reduces output spatial dimensions

Dilated/Atrous Convolution:

$\text{input}_h = \text{output}_h + \text{kernel}_h \times \text{dilation}$

$\text{input}_w = \text{output}_w + \text{kernel}_w \times \text{dilation}$

Expands receptive field without increasing parameters

Address Modification:

$\text{base_addr} + (\text{h} \times \text{stride} + \text{kh} \times \text{dilation}) \times \text{line_stride}$
 $+ (\text{w} \times \text{stride} + \text{kw} \times \text{dilation}) \times \text{element_size}$

9.4.4 Depthwise Separable Convolution

Standard Convolution: $C_{in} \times C_{out} \times K \times K$ multiplications

Depthwise Separable: $C_{in} \times K \times K + C_{in} \times C_{out}$ multiplications

Phase 1 - Depthwise (per-channel):

For each channel c :

$\text{output}[c, h, w] = \text{conv2d}(\text{input}[c], \text{kernel}[c])$

Address Pattern:

Each channel processes independently

No cross-channel data movement

$\text{input_addr} = \text{base} + c \times H \times W + h \times W + w$

Phase 2 - Pointwise (1×1 convolution):

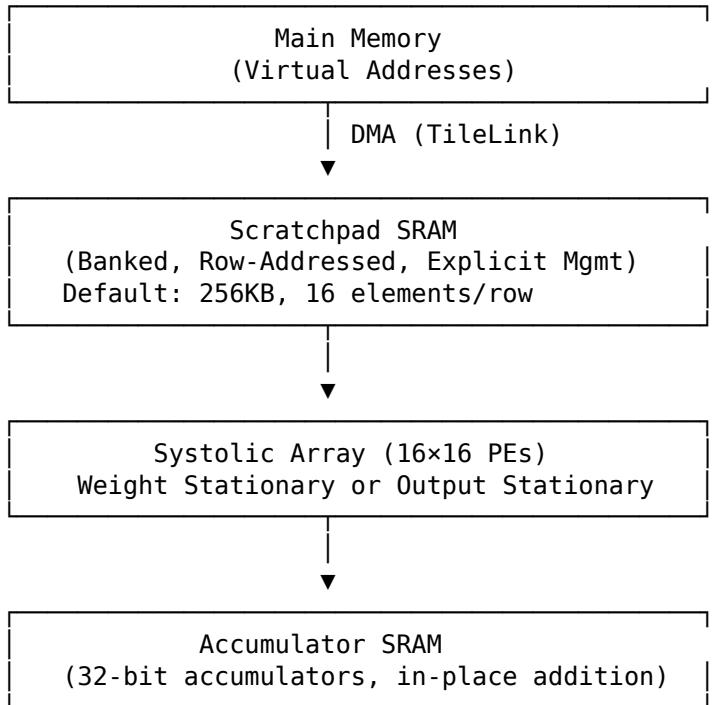
Standard GEMM across channels

$\text{input_addr} = \text{base} + c \times (H \times W) + \text{position}$

9.5 Gemmini Accelerator Architecture (UC Berkeley Reference)

Open-source systolic array generator with well-documented DMA and addressing:

Memory Hierarchy:



Scratchpad Address Format (32-bit):

Bit 31: 0=Scratchpad, 1=Accumulator
 Bit 30: Accumulator mode (0=overwrite, 1=accumulate)
 Bit 29: Read format (0=scaled inputType, 1=raw accType)
 Bits 28:0: Row address within selected memory

Row Width:

Scratchpad: DIM × inputType bits (default: $16 \times 8 = 128$ bits)
 Accumulator: DIM × accType bits (default: $16 \times 32 = 512$ bits)

DMA Operations:

```

// mvin: Move data from DRAM to scratchpad
// Configuration registers set stride
mvin(
    dram_addr,           // Virtual address (64-bit)
    sp_addr,             // Scratchpad row address
    rows,                // Number of rows to transfer
    cols                 // Elements per row (<= DIM)
);

// mvout: Move data from accumulator to DRAM
// Supports integrated max-pooling

```

```

mvout(
    dram_addr,           // Destination address
    acc_addr,            // Accumulator row address
    rows, cols,
    pool_size,           // Optional pooling window
    pool_stride
);

```

Stride Configuration:

```

config_mvin:
    stride[0]: Source memory stride for mvin_A
    stride[1]: Source memory stride for mvin_B
    stride[2]: Source memory stride for mvin_D

config_mvout:
    stride: Destination memory stride
    pool_params: Pooling window and stride

```

9.6 Tiling for Large Matrices

When matrices exceed on-chip memory, tiling is essential:

Hierarchical Tiling (CUTLASS Model):

Level 1 - Thread Block Tile:

Load tile from global memory to shared memory
 Tile size: $M_{\text{tile}} \times K_{\text{tile}}$ (A), $K_{\text{tile}} \times N_{\text{tile}}$ (B)

Level 2 - Warp Tile:

Load from shared memory to registers
 Tile size: $Warp_M \times Warp_K$ (A), $Warp_K \times Warp_N$ (B)

Level 3 - Thread Tile:

Individual thread's computation
 Register-to-register operations

Address Generation per Level:

Global: $\text{base} + \text{block_id} \times \text{block_tile_size}$
 Shared: $\text{smem_base} + \text{warp_id} \times \text{warp_tile_size}$
 Register: Direct indexing within tile

Double-Buffering for Latency Hiding:

Buffer Structure:
 $\text{tile_buffer}[2][\text{TILE_ROWS}][\text{TILE_COLS}]$

Pipeline:

Cycle N: Compute on buffer[0], Load to buffer[1]
Cycle N+1: Compute on buffer[1], Load to buffer[0]

Address Pattern:

```
compute_addr = base + (cycle % 2) × tile_size + element_offset  
load_addr   = base + ((cycle + 1) % 2) × tile_size + element_offset
```

9.7 Sparse Tensor Patterns

Sparse accelerators require index-based indirect addressing:

Compressed Sparse Row (CSR) Format:

Dense: [a 0 b]	CSR: values = [a, b, c, d]
[0 c 0]	col_idx = [0, 2, 1, 0]
[d 0 0]	row_ptr = [0, 2, 3, 4]

Address Calculation:

```
For row r, iterate col_idx[row_ptr[r]] to col_idx[row_ptr[r+1]-1]  
value_addr = values_base + idx  
col_addr = col_idx_base + idx
```

Indirect Access:

```
actual_col = memory[col_addr]  
dense_addr = base + r × stride + actual_col × element_size
```

Hardware Implications:

Dense Accelerator (TPU): Initiation interval = 1 (fully pipelined)
Sparse Accelerator (EIE): Initiation interval = 17 (metadata overhead)

Trade-off:

- Sparse saves memory bandwidth (skip zeros)
- But adds index indirection latency
- Effective for >90% sparsity

9.8 Transformer/Attention Patterns

Modern LLM accelerators require specialized patterns:

Q, K, V Matrix Generation:

Input: X [seq_len, d_model]
Weights: W_Q, W_K, W_V [d_model, d_k]

```
Q = X × W_Q → [seq_len, d_k]  
K = X × W_K → [seq_len, d_k]
```

$V = X \times W_V \rightarrow [seq_len, d_k]$

Address Pattern: Standard GEMM with different weight matrices

```
q_addr = q_base + seq_idx * d_k + head_idx  
k_addr = k_base + seq_idx * d_k + head_idx  
v_addr = v_base + seq_idx * d_k + head_idx
```

Attention Score Computation:

$\text{Attention} = \text{softmax}(Q \times K^T / \sqrt{d_k}) \times V$

Memory Challenge:

```
 $Q \times K^T$  produces  $[seq\_len \times seq\_len]$  matrix  
 $O(n^2)$  memory for sequence length n
```

FlashAttention Tiling:

```
Tile Q, K, V into blocks that fit in SRAM  
Compute attention incrementally without materializing full matrix
```

Tile Address Pattern:

```
q_tile_addr = q_base + tile_row * tile_size * d_k  
k_tile_addr = k_base + tile_col * tile_size * d_k
```

Online Softmax:

```
Track running max and sum across tiles  
No need to store intermediate attention matrix
```

9.9 Programmable Address Generation Unit (PAGU)

Academic reference for flexible tensor addressing:

PAGU Instruction Set:

Operations Supported:

- Standard Convolution
- Padded Convolution
- Strided Convolution
- Dilated (Atrous) Convolution
- Pooling (Max, Average)
- Upsampled Convolution
- Transposed Convolution

Performance: 1.7 cycles per address for convolution

Area Overhead: ~4.6x vs fixed datapath (justified by flexibility)

Configuration Registers:

`tensor_config_t`:

```

base_addr[63:0]           // Tensor base address
dim_sizes[4][31:0]         // Size of each dimension
dim_strides[4][31:0]       // Stride of each dimension
kernel_size[15:0]          // Convolution kernel size
conv_stride[7:0]           // Convolution stride
dilation[7:0]              // Dilation factor
padding[7:0]               // Padding amount
operation_mode[3:0]         // Select operation type

```

Address Generation FSM:

```

always_ff @(posedge clk) begin
    case (state)
        IDLE: begin
            if (start) begin
                dim_counters <= '0;
                state <= GENERATE;
            end
        end

        GENERATE: begin
            // Multi-dimensional address calculation
            addr <= base_addr;
            for (int d = 0; d < num_dims; d++) begin
                case (operation_mode)
                    CONV_STANDARD:
                        addr <= addr + dim_counters[d] * dim_strides[d];
                    CONV_STRIDED:
                        addr <= addr + (dim_counters[d] * conv_stride) * dim_strides[d];
                    CONV_DILATED:
                        addr <= addr + (dim_counters[d] * dilation) * dim_strides[d];
                endcase
            end

            // Increment counters (innermost first)
            increment_counters();

            if (done) state <= IDLE;
        end
    endcase
end

```

10. Implementation Comparison Matrix

10.1 Pattern Complexity vs. Capability

Pattern	Descriptor Fields	Address Logic	Hardware Cost	SW Flexibility
Linear	3 (src, dst, len)	Simple add	Low	Limited
2D Stride	6 (+x_len, y_len, strides)	Nested loop	Medium	Good
3D Vol- ume	9 (+z_len, z_strides)	Triple nested	Medium-High	Very Good
Scatter-Gather	4 per fragment	List traversal	Medium	Excellent
Circular	4 (+size, wrap)	Modulo/mask	Low-Medium	Good
Tensor	N (variable stride)	Sum-of-products	High	Excellent

10.2 Use Case Mapping

Application Domain	Primary Pattern	Secondary Pattern
Block memory copy	Linear	-
Network packets	Scatter-Gather	Linear
Audio streaming	Circular	Linear
Image processing	2D Stride	Linear
Video encode/decode	2D Stride	Scatter-Gather
3D graphics	3D Volume	2D Stride
Neural networks	Tensor	2D Stride
Database operations	Scatter-Gather	Linear

10.3 RAPIDS Current vs. Extended Capability

Capability	Current RAPIDS	Industry Standard
Linear contiguous	YES	YES
Unaligned transfers	YES (3-phase)	Varies
Descriptor chaining	YES	YES
2D stride	NO	Common
3D volume	NO	Rare
Scatter-Gather	Partial (chain)	Full SG
Circular wrap	NO	Common
Tensor addressing	NO	Emerging

11. RAPIDS Extension Recommendations

11.0 PRIORITY 0: Tensor Accelerator Addressing (CRITICAL)

Rationale: Essential for AI/ML workloads. Tensor accelerators represent the highest-growth market for DMA engines. Without tensor addressing, RAPIDS cannot serve neural network inference applications.

Implementation Approach: Data Cube Model (NVDLA-style)

This approach provides the best balance of flexibility and hardware complexity, supporting convolution, pooling, and GEMM operations.

Extended Descriptor Format:

```
typedef struct packed {
    // === Base Fields (existing, 256 bits) ===
    logic [63:0] src_addr;           // Base source address
    logic [63:0] dst_addr;           // Base destination address
    logic [31:0] length;             // For linear mode: beats
    logic [31:0] next_descriptor_ptr;
    logic [7:0] desc_priority;
    logic [3:0] channel_id;
    logic error;
    logic last;
    logic gen_irq;
    logic valid;

    // === Tensor Extension (repurpose reserved bits + add 2nd descriptor word) ===
    // Mode selection
    logic [3:0] transfer_mode;      // 0=linear, 1=2D, 2=3D/tensor, 3=conv, 4=pool

    // Data cube dimensions
    logic [15:0] width;              // W dimension (elements)
    logic [15:0] height;             // H dimension (lines)
    logic [15:0] channels;           // C dimension (surfaces)

    // Stride parameters (32-byte aligned)
    logic [31:0] line_stride;        // Bytes between lines
    logic [31:0] surface_stride;     // Bytes between surfaces

    // Convolution parameters (when transfer_mode == CONV)
    logic [7:0] kernel_width;        // Kw
    logic [7:0] kernel_height;       // Kh
    logic [7:0] conv_stride_x;       // Horizontal stride
    logic [7:0] conv_stride_y;       // Vertical stride
    logic [7:0] dilation_x;          // Horizontal dilation
```

```

    logic [7:0] dilation_y;           // Vertical dilation
    logic [7:0] pad_left;            // Left padding
    logic [7:0] pad_right;           // Right padding
    logic [7:0] pad_top;             // Top padding
    logic [7:0] pad_bottom;          // Bottom padding

    // Pooling parameters (when transfer_mode == POOL)
    logic [7:0] pool_width;
    logic [7:0] pool_height;
    logic [7:0] pool_stride_x;
    logic [7:0] pool_stride_y;
    logic [1:0] pool_type;           // 0=max, 1=avg, 2=min

    // Data type and precision
    logic [3:0] data_type;           // 0=int8, 1=int16, 2=fp16, 3=bf16, 4=fp32
    logic [3:0] element_size;        // Bytes per element (1, 2, 4)
} descriptor_tensor_t;

```

Address Generation Unit (AGU) Architecture:

```

module tensor_address_generator #(
    parameter int PIPE_STAGES = 3      // Pipeline for multiply-accumulate
) (
    input  logic         clk,
    input  logic         rst_n,

    // Configuration from descriptor
    input  descriptor_tensor_t desc,
    input  logic         start,

    // Address output stream
    output logic [63:0] addr,
    output logic         addr_valid,
    input  logic         addr_ready,

    // Status
    output logic         done
);

    // Dimension counters
    logic [15:0] w_cnt, h_cnt, c_cnt;
    logic [7:0]  kw_cnt, kh_cnt;        // Kernel position (for conv)

    // Pipeline registers for address calculation
    logic [63:0] addr_stage[PIPE_STAGES];

    // FSM states

```

```

typedef enum logic [2:0] {
    IDLE,
    CALC_BASE,
    CALC_LINE,
    CALC_SURFACE,
    OUTPUT
} state_t;
state_t state;

// Address calculation based on mode
always_ff @(posedge clk or negedge rst_n) begin
    if (!rst_n) begin
        state <= IDLE;
        w_cnt <= '0;
        h_cnt <= '0;
        c_cnt <= '0;
    end else begin
        case (state)
            IDLE: begin
                if (start) begin
                    w_cnt <= '0;
                    h_cnt <= '0;
                    c_cnt <= '0;
                    state <= CALC_BASE;
                end
            end
    end

    CALC_BASE: begin
        // Stage 1: Base address
        addr_stage[0] <= desc.src_addr;
        state <= CALC_LINE;
    end

    CALC_LINE: begin
        // Stage 2: Add line offset
        case (desc.transfer_mode)
            MODE_LINEAR:
                addr_stage[1] <= addr_stage[0] + (w_cnt * desc.element_s

            MODE_2D, MODE_3D:
                addr_stage[1] <= addr_stage[0]
                    + (h_cnt * desc.line_stride)
                    + (w_cnt * desc.element_size);

            MODE_CONV:
                // Convolution: account for stride and kernel position

```

```

        addr_stage[1] <= addr_stage[0]
        + ((h_cnt * desc.conv_stride_y + kh_cnt *
        + ((w_cnt * desc.conv_stride_x + kw_cnt *

    MODE_POOL:
        addr_stage[1] <= addr_stage[0]
        + ((h_cnt * desc.pool_stride_y) * desc.li
        + ((w_cnt * desc.pool_stride_x) * desc.el

    endcase
    state <= CALC_SURFACE;
end

CALC_SURFACE: begin
    // Stage 3: Add surface offset (for 3D/tensor)
    if (desc.transfer_mode inside {MODE_3D, MODE_CONV}) begin
        addr_stage[2] <= addr_stage[1] + (c_cnt * desc.surface_strid
    end else begin
        addr_stage[2] <= addr_stage[1];
    end
    state <= OUTPUT;
end

OUTPUT: begin
    if (addr_ready) begin
        // Increment counters (innermost first)
        if (desc.transfer_mode == MODE_CONV) begin
            // Convolution: iterate kernel, then spatial, then channel
            if (kw_cnt < desc.kernel_width - 1) begin
                kw_cnt <= kw_cnt + 1;
            end else begin
                kw_cnt <= '0;
                if (kh_cnt < desc.kernel_height - 1) begin
                    kh_cnt <= kh_cnt + 1;
                end else begin
                    kh_cnt <= '0;
                    // Continue to spatial/channel iteration
                    increment_spatial_counters();
                end
            end
        end else begin
            increment_spatial_counters();
        end
    end
    if (transfer_complete())
        state <= IDLE;
    else

```

```

                state <= CALC_BASE;
            end
        end
    endcase
end
end

assign addr = addr_stage[PIPE_STAGES-1];
assign addr_valid = (state == OUTPUT);
assign done = (state == IDLE) && !start;

endmodule

```

Scratchpad Integration (Gemmini-style):

```

// Memory address format for tensor scratchpad
typedef struct packed {
    logic      mem_select;          // 0=scratchpad, 1=accumulator
    logic      acc_mode;            // 0=overwrite, 1=accumulate
    logic      read_format;         // 0=scaled, 1=raw
    logic [28:0] row_addr;          // Row address within memory
} scratchpad_addr_t;

// DMA configuration for tensor loads
typedef struct packed {
    logic [63:0] dram_addr;          // Virtual address in main memory
    scratchpad_addr_t sp_addr;        // Scratchpad destination
    logic [15:0] rows;               // Number of rows to transfer
    logic [15:0] cols;               // Elements per row
    logic [31:0] dram_stride;         // Source memory stride
    logic [15:0] sp_stride;           // Scratchpad row stride (usually 1)
} tensor_dma_config_t;

```

Alignment Requirements:

All tensor addresses must satisfy:

- base_addr aligned to 32 bytes (NVDLA requirement)
- line_stride aligned to 32 bytes
- surface_stride aligned to 32 bytes
- For optimal performance: align to 64 bytes (RAPIDS beat size)

Verification:

```

assert((desc.src_addr & 32'h1F) == 0);
assert((desc.line_stride & 32'h1F) == 0);
assert((desc.surface_stride & 32'h1F) == 0);

```

Performance Considerations:

Address Generation Rate:

- Target: 1 address per cycle (pipelined)
- Convolution: 1.7 cycles/address typical (due to kernel iteration)

Memory Bandwidth:

- TPU-class: 600+ GB/s (HBM)
- FPGA-class: 20-80 GB/s (DDR4/5)
- RAPIDS target: Match AXI interface bandwidth

Tiling Support:

- Hardware computes tile boundaries
- Double-buffering via ping-pong descriptors
- Software provides tile dimensions, hardware handles addressing

11.1 Priority 1: 2D Stride Support

Rationale: High value for image/video processing with moderate complexity. Also serves as foundation for tensor addressing.

Descriptor Extension:

```
typedef struct packed {
    // Existing fields (256 bits)...

    // 2D extension (96 bits in reserved space)
    logic [31:0] x_length;           // Bytes per row
    logic [15:0] y_length;           // Number of rows
    logic [31:0] src_stride;         // Source row stride
    logic [31:0] dst_stride;         // Destination row stride
    logic          mode_2d;           // Enable 2D mode
} descriptor_2d_t;
```

Address Generation:

```
always_comb begin
    if (mode_2d) begin
        row_src_addr = src_addr + current_row * src_stride;
        row_dst_addr = dst_addr + current_row * dst_stride;
        row_length   = x_length;
    end else begin
        row_src_addr = src_addr;
        row_dst_addr = dst_addr;
        row_length   = length * BYTES_PER_BEAT;
    end
end
```

11.2 Priority 2: Circular Buffer Mode

Rationale: Essential for streaming applications, low complexity addition.

Descriptor Extension:

```
typedef struct packed {
    // Existing fields...

    // Circular buffer extension
    logic [31:0] buffer_size;      // Total buffer size
    logic [31:0] buffer_mask;     // Size-1 for power-of-2
    logic         circular_mode;   // Enable wrap-around
    logic         src_circular;    // Apply to source
    logic         dst_circular;    // Apply to destination
} descriptor_circular_t;
```

Address Generation:

```
function automatic logic [63:0] wrap_address(
    logic [63:0] base,
    logic [63:0] offset,
    logic [31:0] mask,
    logic         enable
);
    if (enable)
        return base + (offset & {32'b0, mask});
    else
        return base + offset;
endfunction
```

11.3 Priority 3: Enhanced Scatter-Gather

Rationale: Full SG capability for network and fragmented memory use cases.

Current Limitation: Chain-only SG requires one descriptor per fragment.

Enhancement: Add fragment list support within single descriptor:

```
typedef struct packed {
    logic [63:0] fragment_list_ptr; // Pointer to fragment array
    logic [15:0] fragment_count;    // Number of fragments
    logic         sg_mode;          // Enable SG mode

    // Fragment entry format (in memory):
    // [63:0] address
```

```
// [31:0] length  
// [31:0] flags  
} descriptor_sg_t;
```

12. References

Industry Documentation

1. [AMD AXI DMA Product Guide \(PG021\)](#) - Scatter-Gather descriptor format
2. [Analog Devices AXI DMAC](#) - 2D transfer implementation
3. [Understanding AXI Addressing \(ZipCPU\)](#) - Burst mode analysis
4. [WRAP Address Calculation \(Verification Guide\)](#) - Wrap burst formulas

GPU and Accelerator Patterns

5. [NVIDIA CUDA Memory Coalescing](#) - Coalesced access patterns
6. [CUDA Pitch Linear Memory](#) - 2D array allocation
7. [NVIDIA TMA Tutorial \(CUTLASS\)](#) - Tensor Memory Accelerator
8. [Efficient GEMM in CUDA \(CUTLASS\)](#) - Hierarchical tiling

DMA Architecture

9. [Intel DMA Descriptors](#) - PCIe DMA linked lists
10. [DMA330 Microcode](#) - Linked-list implementation

Neural Network Accelerators

11. [NVDLA Hardware Architecture](#) - Data cube model, CDMA
12. [NVDLA In-Memory Data Formats](#) - Line stride, surface stride
13. [NVDLA Unit Description](#) - Convolution DMA details
14. [Gemmini Accelerator \(UC Berkeley\)](#) - Open-source systolic array
15. [Gemmini Documentation \(Chipyard\)](#) - DMA and scratchpad
16. [Google TPU Architecture](#) - Systolic array overview
17. [TPU System Architecture](#) - Memory hierarchy
18. [Programmable AGU for DNNs](#) - Flexible address generation

Convolution and Tensor Operations

19. [Im2Col Characterization](#) - Convolution memory patterns
20. [ST Neural-ART Programming Model](#) - NPU descriptor concepts

21. [Systolic Array Dataflows Survey](#) - Weight/input/output stationary
22. [FlashAttention Analysis](#) - Attention tiling patterns
23. [Depthwise Separable Convolutions](#) - Channel-wise addressing

Sparse Tensor Formats

24. [FLAASH Sparse Tensor Accelerator](#) - CSR/CSF formats
 25. [SparTen CNN Accelerator](#) - Sparse tensor addressing
 26. [Indirection Stream Architecture](#) - Hardware indirect addressing
-

Document History

Version	Date	Author	Changes
1.0	2026-01-13	RTL Design Sherpa	Initial comprehensive analysis

This document is part of the RAPIDS (Rapid AXI Programmable In-band Descriptor System) design documentation.