



RTL Design Sherpa

**STREAM Subsystem Hardware Architecture
Specification 0.90**

January 3, 2026

Table of Contents

| | |
|---|----|
| 1 Stream Has Index..... | 22 |
| 2 Document Information..... | 22 |
| 2.1 STREAM Hardware Architecture Specification..... | 22 |
| 2.2 Document Purpose..... | 22 |
| 2.3 References..... | 22 |
| 2.4 Terminology..... | 23 |
| 2.5 Revision History..... | 24 |
| 3 Purpose and Scope..... | 24 |
| 3.1 Document Purpose..... | 24 |
| 3.2 Scope..... | 25 |
| 3.2.1 In Scope..... | 25 |
| 3.2.2 Out of Scope..... | 25 |
| 3.3 Design Philosophy..... | 26 |
| 3.3.1 Tutorial-Focused Simplicity..... | 26 |
| 3.3.2 Relationship to RAPIDS..... | 26 |
| 3.4 Intended Audience..... | 26 |
| 4 Document Conventions..... | 27 |
| 4.1 Notation Conventions..... | 27 |
| 4.1.1 Signal Naming..... | 27 |
| 4.1.2 Interface Prefixes..... | 27 |
| 4.2 Diagram Conventions..... | 28 |
| 4.2.1 Block Diagrams..... | 28 |

| | |
|---|----|
| 4.2.2 Timing Diagrams..... | 28 |
| 4.3 Numeric Conventions..... | 28 |
| 4.3.1 Units..... | 28 |
| 4.3.2 Address Notation..... | 28 |
| 4.3.3 Size Notation..... | 28 |
| 4.4 Requirement Notation..... | 29 |
| 4.5 Reference Format..... | 29 |
| 5 Definitions and Acronyms..... | 29 |
| 5.1 Acronyms..... | 29 |
| 5.2 Definitions..... | 30 |
| 5.2.1 Architectural Terms..... | 30 |
| 5.2.2 Interface Terms..... | 31 |
| 5.2.3 Transfer Terms..... | 31 |
| 5.3 Units and Measurements..... | 32 |
| 5.4 Document-Specific Terms..... | 32 |
| 6 Use Cases..... | 33 |
| 6.1 Primary Use Cases..... | 33 |
| 6.1.1 Memory-to-Memory Data Movement..... | 33 |
| 6.1.2 Scatter-Gather Operations..... | 33 |
| 6.1.3 Multi-Channel Concurrent Transfers..... | 33 |
| 6.2 Secondary Use Cases..... | 34 |
| 6.2.1 Buffer Management..... | 34 |
| 6.2.2 Memory Initialization..... | 34 |
| 6.3 Target Applications..... | 34 |

| | |
|--|----|
| 7 Key Features..... | 34 |
| 7.1 Feature Summary..... | 34 |
| 7.2 Descriptor-Based Operation..... | 35 |
| 7.2.1 Single-Write Kick-off..... | 35 |
| 7.2.2 Autonomous Chaining..... | 35 |
| 7.3 Multi-Channel Architecture..... | 35 |
| 7.3.1 Channel Independence..... | 35 |
| 7.3.2 Resource Sharing..... | 36 |
| 7.3.3 Priority-Based Arbitration..... | 36 |
| 7.4 Data Path Features..... | 36 |
| 7.4.1 Parameterizable Data Width..... | 36 |
| 7.4.2 Aligned Address Requirement..... | 36 |
| 7.5 Monitoring and Debug..... | 37 |
| 7.5.1 MonBus Integration..... | 37 |
| 7.5.2 Status Visibility..... | 37 |
| 7.6 Error Handling..... | 37 |
| 7.6.1 Error Detection..... | 37 |
| 7.6.2 Error Isolation..... | 38 |
| 8 System Context..... | 38 |
| 8.1 System Integration Overview..... | 38 |
| 8.2 External Connections..... | 40 |
| 8.2.1 Configuration Path..... | 40 |
| 8.2.2 Data Paths..... | 40 |
| 8.2.3 Monitoring Path..... | 40 |

| | |
|--------------------------------------|----|
| 8.3 Clock and Reset..... | 41 |
| 8.3.1 Clock Domain..... | 41 |
| 8.3.2 Reset Requirements..... | 41 |
| 8.4 Memory Map Requirements..... | 41 |
| 8.4.1 Descriptor Memory..... | 41 |
| 8.4.2 Data Memory..... | 41 |
| 8.5 Interrupt Integration..... | 42 |
| 8.5.1 Interrupt Signals..... | 42 |
| 8.5.2 Interrupt Conditions..... | 42 |
| 8.6 Power Considerations..... | 42 |
| 8.6.1 Clock Gating..... | 42 |
| 8.6.2 Power States..... | 42 |
| 9 Block Diagram..... | 44 |
| 9.1 Top-Level Architecture..... | 44 |
| 9.2 Functional Block Summary..... | 45 |
| 9.3 Block Interactions..... | 45 |
| 9.3.1 Control Flow..... | 45 |
| 9.3.2 Data Flow..... | 45 |
| 9.3.3 Monitoring Flow..... | 45 |
| 9.4 Parameterization..... | 46 |
| 10 Data Flow..... | 46 |
| 10.1 Transfer Sequence Overview..... | 46 |
| 10.2 Phase Details..... | 48 |
| 10.2.1 Phase 1: Kick-off..... | 48 |

| | |
|---|----|
| 10.2.2 Phase 2: Descriptor Fetch..... | 48 |
| 10.2.3 Phase 3: Descriptor Parse..... | 48 |
| 10.2.4 Phase 4: Arbitration..... | 48 |
| 10.2.5 Phase 5: Read Phase..... | 49 |
| 10.2.6 Phase 6: Write Phase..... | 49 |
| 10.2.7 Phase 7: Chain Check..... | 49 |
| 10.3 Concurrent Operation..... | 49 |
| 10.3.1 Multi-Channel Concurrency..... | 49 |
| 10.3.2 Pipeline Overlap..... | 49 |
| 10.4 Data Path Bandwidth..... | 50 |
| 10.4.1 Theoretical Maximum..... | 50 |
| 10.4.2 Practical Considerations..... | 50 |
| 11 Multi-Channel Architecture..... | 50 |
| 11.1 Channel Organization..... | 50 |
| 11.2 Per-Channel Resources..... | 51 |
| 11.3 Shared Resources..... | 51 |
| 11.4 Arbitration Scheme..... | 52 |
| 11.4.1 Priority-Based Arbitration..... | 52 |
| 11.4.2 Fairness Mechanism..... | 52 |
| 11.4.3 Arbitration Points..... | 52 |
| 11.5 Channel States..... | 52 |
| 11.6 SRAM Partitioning..... | 54 |
| 11.6.1 Static Partitioning (Default)..... | 54 |
| 11.6.2 Dynamic Allocation (Optional)..... | 55 |

| | |
|--|----|
| 11.7 Error Isolation..... | 55 |
| 12 AXI Master Interfaces..... | 55 |
| 12.1 Overview..... | 55 |
| 12.2 Descriptor Fetch Master..... | 56 |
| 12.2.1 Signal Summary..... | 56 |
| 12.2.2 Transaction Characteristics..... | 56 |
| 12.2.3 Address Alignment..... | 57 |
| 12.3 Data Read Master..... | 57 |
| 12.3.1 Signal Summary..... | 57 |
| 12.3.2 Transaction Characteristics..... | 58 |
| 12.3.3 Burst Generation..... | 58 |
| 12.4 Data Write Master..... | 58 |
| 12.4.1 Signal Summary..... | 58 |
| 12.4.2 Transaction Characteristics..... | 59 |
| 12.5 Common AXI Attributes..... | 59 |
| 12.5.1 Supported Features..... | 59 |
| 12.5.2 Error Handling..... | 60 |
| 13 APB Configuration Interface..... | 60 |
| 13.1 Overview..... | 60 |
| 13.2 Signal Summary..... | 60 |
| 13.3 Register Map Overview..... | 61 |
| 13.4 Key Registers..... | 62 |
| 13.4.1 Global Control (CTRL) - 0x000..... | 62 |
| 13.4.2 Global Status (STATUS) - 0x004..... | 62 |

| | |
|--|----|
| 13.4.3 Channel Control (CHn_CTRL) - 0x040 + n*0x10..... | 63 |
| 13.4.4 Channel Status (CHn_STATUS) - 0x044 + n*0x10..... | 63 |
| 13.5 Access Timing..... | 64 |
| 13.5.1 Write Access..... | 64 |
| 13.5.2 Read Access..... | 64 |
| 13.6 Kick-off Blocking..... | 65 |
| 14 Monitoring Interface..... | 65 |
| 14.1 Overview..... | 65 |
| 14.2 Signal Summary..... | 65 |
| 14.3 Packet Format..... | 65 |
| 14.4 Event Types..... | 66 |
| 14.4.1 Transfer Events (EVENT_TYPE = 0x01)..... | 66 |
| 14.4.2 Error Events (EVENT_TYPE = 0x02)..... | 67 |
| 14.4.3 Performance Events (EVENT_TYPE = 0x03)..... | 67 |
| 14.5 Event Generation..... | 67 |
| 14.5.1 Configurable Events..... | 67 |
| 14.5.2 Event Priority..... | 67 |
| 14.6 Backpressure Handling..... | 68 |
| 14.6.1 Flow Control..... | 68 |
| 14.6.2 Recommendations..... | 68 |
| 14.7 Typical MonBus Downstream..... | 68 |
| 14.8 Example Event Sequence..... | 68 |
| 15 Throughput Targets..... | 69 |
| 15.1 Theoretical Maximum..... | 69 |

| | |
|--|----|
| 15.1.1 Per-Interface Bandwidth..... | 69 |
| 15.1.2 Aggregate Bandwidth..... | 69 |
| 15.2 Practical Throughput..... | 69 |
| 15.2.1 Single-Channel Performance..... | 69 |
| 15.2.2 Multi-Channel Performance..... | 70 |
| 15.3 Throughput Limiting Factors..... | 70 |
| 15.3.1 Memory System Factors..... | 70 |
| 15.3.2 STREAM Internal Factors..... | 70 |
| 15.4 Performance Targets..... | 71 |
| 15.4.1 Design Targets..... | 71 |
| 15.4.2 Verification Criteria..... | 71 |
| 15.5 Descriptor Chain Efficiency..... | 71 |
| 15.5.1 Descriptor Overhead..... | 71 |
| 15.5.2 Chaining Efficiency..... | 71 |
| 16 Latency Characteristics..... | 72 |
| 16.1 Latency Components..... | 72 |
| 16.1.1 End-to-End Latency Breakdown..... | 72 |
| 16.1.2 Total Latency Formula..... | 72 |
| 16.2 First-Byte Latency..... | 73 |
| 16.3 Latency by Transfer Size..... | 73 |
| 16.3.1 Small Transfers..... | 73 |
| 16.3.2 Large Transfers..... | 74 |
| 16.4 Latency Variability..... | 74 |
| 16.4.1 Sources of Variability..... | 74 |

| | |
|---------------------------------------|----|
| 16.4.2 Worst-Case Latency..... | 74 |
| 16.5 Latency Optimization..... | 74 |
| 16.5.1 Software Recommendations..... | 74 |
| 16.5.2 Hardware Configuration..... | 75 |
| 16.6 Interrupt Latency..... | 75 |
| 17 Resource Estimates..... | 75 |
| 17.1 FPGA Resource Summary..... | 75 |
| 17.2 Resource Breakdown by Block..... | 76 |
| 17.2.1 APB Configuration Slave..... | 76 |
| 17.2.2 Descriptor Engine..... | 76 |
| 17.2.3 Scheduler (8 channels)..... | 76 |
| 17.2.4 Channel Arbiter..... | 76 |
| 17.2.5 AXI Read Engine..... | 77 |
| 17.2.6 SRAM Buffer..... | 77 |
| 17.2.7 AXI Write Engine..... | 77 |
| 17.2.8 MonBus Reporter..... | 77 |
| 17.3 SRAM Sizing..... | 77 |
| 17.4 Scaling with Parameters..... | 78 |
| 17.4.1 Channel Count Impact..... | 78 |
| 17.4.2 Data Width Impact..... | 78 |
| 17.5 ASIC Estimates..... | 78 |
| 17.5.1 Power Estimates..... | 78 |
| 17.6 Timing Estimates..... | 79 |
| 17.6.1 Target Frequencies..... | 79 |

| | |
|---|----|
| 17.6.2 Critical Paths..... | 79 |
| 18 System Requirements..... | 79 |
| 18.1 Interface Requirements..... | 79 |
| 18.1.1 APB Interface..... | 79 |
| 18.1.2 AXI Interfaces..... | 80 |
| 18.1.3 MonBus Interface..... | 80 |
| 18.2 Memory Requirements..... | 80 |
| 18.2.1 Descriptor Memory..... | 80 |
| 18.2.2 Data Memory..... | 80 |
| 18.3 System Bandwidth Allocation..... | 81 |
| 18.3.1 AXI Interconnect Requirements..... | 81 |
| 18.3.2 Interconnect Recommendations..... | 81 |
| 18.4 Interrupt Requirements..... | 81 |
| 18.4.1 Interrupt Controller..... | 81 |
| 18.4.2 Interrupt Handling..... | 81 |
| 18.5 Address Space Requirements..... | 82 |
| 18.5.1 STREAM Register Space..... | 82 |
| 18.5.2 Memory Map Constraints..... | 82 |
| 18.6 Software Requirements..... | 82 |
| 18.6.1 Driver Requirements..... | 82 |
| 18.6.2 Programming Sequence..... | 82 |
| 19 Clocking and Reset..... | 83 |
| 19.1 Clock Requirements..... | 83 |
| 19.1.1 Primary Clock (aclk)..... | 83 |

| | |
|--|----|
| 19.1.2 Clock Relationship..... | 83 |
| 19.2 Reset Requirements..... | 84 |
| 19.2.1 Reset Signal (aresetn)..... | 84 |
| 19.2.2 Reset Sequence..... | 84 |
| 19.2.3 Reset Behavior..... | 85 |
| 19.3 Power Management..... | 85 |
| 19.3.1 Clock Gating Support..... | 85 |
| 19.3.2 Low Power Recommendations..... | 85 |
| 19.4 Timing Constraints..... | 85 |
| 19.4.1 Clock Period Constraints..... | 85 |
| 19.4.2 False Path Constraints..... | 86 |
| 19.5 CDC Considerations..... | 86 |
| 19.5.1 Single Clock Domain..... | 86 |
| 19.5.2 CDC Recommendations..... | 86 |
| 20 Verification Strategy..... | 87 |
| 20.1 Verification Approach..... | 87 |
| 20.1.1 Multi-Level Verification..... | 87 |
| 20.2 Block-Level Verification..... | 87 |
| 20.2.1 Test Categories..... | 87 |
| 20.2.2 Coverage Metrics..... | 88 |
| 20.3 Integration-Level Verification..... | 88 |
| 20.3.1 Interface Verification..... | 88 |
| 20.3.2 Multi-Channel Scenarios..... | 89 |
| 20.4 System-Level Verification..... | 89 |

| | |
|--|----|
| 20.4.1 End-to-End Test Cases..... | 89 |
| 20.4.2 Performance Verification..... | 90 |
| 20.5 Testbench Architecture..... | 90 |
| 20.5.1 CocoTB-Based Verification..... | 90 |
| 20.5.2 Key Testbench Components..... | 90 |
| 20.6 Verification Environment..... | 91 |
| 20.6.1 Simulator Support..... | 91 |
| 20.6.2 Test Execution..... | 91 |
| 20.7 Formal Verification..... | 91 |
| 20.7.1 Properties for Formal..... | 91 |
| 20.7.2 Formal Tool Usage..... | 92 |
| 20.8 Sign-Off Criteria..... | 92 |
| 20.8.1 Block-Level Sign-Off..... | 92 |
| 20.8.2 Integration Sign-Off..... | 92 |
| 20.8.3 System Sign-Off..... | 92 |
| 21 STREAM Specification Index..... | 93 |
| 21.1 Document Organization..... | 93 |
| 21.1.1 Front Matter..... | 93 |
| 21.1.2 Chapter 1: Overview..... | 93 |
| 21.1.3 Chapter 2: Functional Blocks..... | 93 |
| 21.1.4 Chapter 3: Interfaces..... | 93 |
| 21.1.5 Chapter 4: Registers..... | 93 |
| 21.1.6 Chapter 5: Programming..... | 94 |
| 21.1.7 Chapter 6: Configuration Reference..... | 94 |

| | |
|---|-----|
| 21.2 Quick Reference..... | 94 |
| 21.2.1 Functional Unit Blocks (FUB)..... | 94 |
| 21.2.2 Integration Blocks (MAC)..... | 95 |
| 21.3 Performance Modes (AXI Engines)..... | 95 |
| 21.3.1 Holistic Overview: V1/V2/V3 Working Together..... | 95 |
| 21.3.2 V1 - Low Performance (Single Outstanding Transaction)..... | 96 |
| 21.3.3 V2 - Medium Performance (Command Pipelined)..... | 96 |
| 21.3.4 V3 - High Performance (Out-of-Order Completion)..... | 96 |
| 21.3.5 Performance Comparison Summary..... | 96 |
| 21.4 Clock and Reset Summary..... | 97 |
| 21.4.1 Clock Domains..... | 97 |
| 21.4.2 Reset Signals..... | 97 |
| 21.5 Interface Summary..... | 98 |
| 21.5.1 External Interfaces..... | 98 |
| 21.5.2 Internal Buses..... | 98 |
| 21.6 Area Estimates..... | 98 |
| 21.6.1 By Performance Mode..... | 98 |
| 21.6.2 Breakdown (Low Performance)..... | 99 |
| 21.7 Related Documentation..... | 99 |
| 21.8 Specification Conventions..... | 100 |
| 21.8.1 Signal Naming..... | 100 |
| 21.8.2 Parameter Naming..... | 100 |
| 21.8.3 State Machine Naming..... | 100 |
| 22 Product Requirements Document (PRD)..... | 100 |

| | |
|--|-----|
| 22.1 STREAM - Scatter-gather Transfer Rapid Engine for AXI Memory..... | 100 |
| 22.2 1. Executive Summary..... | 101 |
| 22.2.1 1.1 Quick Stats..... | 101 |
| 22.2.2 1.2 Project Goals..... | 101 |
| 22.3 2. Key Design Principles..... | 101 |
| 22.3.1 2.1 Simplifications from RAPIDS..... | 101 |
| 22.3.2 2.2 Tutorial-Friendly Features..... | 102 |
| 22.4 3. Architecture Overview..... | 102 |
| 22.4.1 3.1 Top-Level Block Diagram..... | 102 |
| 22.4.2 3.2 Data Flow..... | 103 |
| 22.5 4. Interfaces..... | 103 |
| 22.5.1 4.1 External Interfaces..... | 103 |
| 22.5.2 4.2 Descriptor Format..... | 104 |
| 22.6 5. Key Components..... | 105 |
| 22.6.1 5.1 Descriptor Engine (APB-Only for STREAM)..... | 105 |
| 22.6.2 5.2 Scheduler Group (Integration Wrapper)..... | 105 |
| 22.6.3 5.3 Scheduler (Simplified from RAPIDS)..... | 107 |
| 22.6.4 5.3 AXI Read Engine (Streaming Pipeline - NO FSM)..... | 107 |
| 22.6.5 5.4 AXI Write Engine (Streaming Pipeline - NO FSM)..... | 108 |
| 22.6.6 5.5 Simple SRAM..... | 109 |
| 22.7 6. Configuration and Control..... | 109 |
| 22.7.1 6.1 APB Register Map..... | 109 |
| 22.7.2 6.2 Channel Configuration..... | 110 |
| 22.8 7. Resource Sharing and Arbitration..... | 110 |

| | |
|--|-----|
| 22.8.1 7.1 Shared Resources..... | 110 |
| 22.8.2 7.2 Arbitration Strategy..... | 110 |
| 22.9 8. Error Detection and Recovery..... | 111 |
| 22.9.1 8.1 Error Types..... | 111 |
| 22.9.2 8.2 Error Recovery..... | 111 |
| 22.10 9. MonBus Integration..... | 111 |
| 22.10.1 9.1 Standard MonBus Format..... | 111 |
| 22.10.2 9.2 STREAM Event Codes..... | 112 |
| 22.10.3 9.3 Default Configuration..... | 112 |
| 22.11 10. Design Constraints..... | 112 |
| 22.11.1 10.1 Tutorial Constraints (Intentional Simplifications)..... | 112 |
| 22.11.2 10.2 Implementation Constraints..... | 113 |
| 22.12 11. Verification Strategy..... | 113 |
| 22.12.1 11.1 Test Organization..... | 113 |
| 22.12.2 11.2 Test Levels..... | 113 |
| 22.13 12. Performance Characteristics..... | 114 |
| 22.13.1 12.1 Throughput by Engine Version..... | 114 |
| 22.13.2 12.2 Latency..... | 114 |
| 22.13.3 12.3 Resource Utilization by Engine Version..... | 114 |
| 22.14 13. Development Roadmap..... | 115 |
| 22.14.1 13.1 Phase 1: Foundation (Current)..... | 115 |
| 22.14.2 13.2 Phase 2: Core Blocks..... | 115 |
| 22.14.3 13.3 Phase 3: Data Path..... | 115 |
| 22.14.4 13.4 Phase 4: Integration..... | 115 |

| | |
|--|-----|
| 22.14.5 13.5 Phase 5: Multi-Channel..... | 116 |
| 22.14.6 13.6 Phase 6: Advanced Engines (Future - V2/V3)..... | 116 |
| 22.15 14. Educational Value..... | 116 |
| 22.15.1 14.1 Learning Objectives..... | 116 |
| 22.15.2 14.2 Progression Path..... | 117 |
| 22.16 15. Success Criteria..... | 117 |
| 22.16.1 15.1 Functional..... | 117 |
| 22.16.2 15.2 Quality..... | 117 |
| 22.16.3 15.3 Performance..... | 117 |
| 22.17 16. Open Questions (For Review)..... | 118 |
| 22.17.1 16.1 Descriptor Engine Adaptation..... | 118 |
| 22.17.2 16.2 AXI Descriptor Master..... | 118 |
| 22.17.3 16.3 Channel Arbitration..... | 118 |
| 22.17.4 16.4 SRAM Partitioning..... | 118 |
| 22.18 16. Attribution and Contribution Guidelines..... | 118 |
| 22.18.1 16.1 Git Commit Attribution..... | 118 |
| 22.19 16.2 PDF Generation Location..... | 119 |
| 22.20 17. References..... | 119 |
| 22.20.1 17.1 Internal Documentation..... | 119 |
| 22.20.2 17.2 External References..... | 119 |
| 22.21 Navigation..... | 120 |
| 23 Claude Code Guide: STREAM Subsystem..... | 120 |
| 23.1 Quick Context..... | 120 |
| 23.2 Global Requirements Reference..... | 120 |

| | |
|---|-----|
| 23.3 Critical Rules for This Subsystem..... | 121 |
| 23.3.1 Rule #0: Attribution Format for Git Commits..... | 121 |
| 23.3.2 Rule #0.1: TUTORIAL FOCUS - Intentional Simplifications..... | 121 |
| 23.3.3 Rule #0.1: Testbench Location and Test Structure (MANDATORY).... | 121 |
| 23.3.4 Rule #0.2: Three Mandatory TB Methods (MANDATORY)..... | 122 |
| 23.3.5 Rule #1: REUSE from RAPIDS Where Appropriate..... | 122 |
| 23.3.6 Rule #2: Descriptor Format is DIFFERENT from RAPIDS..... | 123 |
| 23.3.7 Rule #3: Know the Shared Resources..... | 123 |
| 23.4 Architecture Quick Reference..... | 123 |
| 23.4.1 Block Organization..... | 123 |
| 23.4.2 Module Status..... | 123 |
| 23.5 Common User Questions and Responses..... | 124 |
| 23.5.1 Q: “How is STREAM different from RAPIDS?”..... | 124 |
| 23.5.2 Q: “How do I kick off a transfer?”..... | 125 |
| 23.5.3 Q: “How many channels can I use?”..... | 125 |
| 23.5.4 Q: “What’s the descriptor format?”..... | 126 |
| 23.5.5 Q: “How do I run STREAM tests?”..... | 126 |
| 23.6 Integration Patterns..... | 127 |
| 23.6.1 Pattern 1: Basic STREAM Instantiation..... | 127 |
| 23.6.2 Pattern 2: Descriptor Creation (Software Model)..... | 128 |
| 23.6.3 Pattern 3: MonBus Integration..... | 129 |
| 23.7 Anti-Patterns to Catch..... | 129 |
| 23.7.1 ✗ Anti-Pattern 1: Adding Alignment Fixup..... | 129 |
| 23.7.2 ✗ Anti-Pattern 2: Using Length in Bytes..... | 129 |

| | |
|---|-----|
| 23.7.3 ✗ Anti-Pattern 3: Circular Buffer Descriptors..... | 129 |
| 23.7.4 ✗ Anti-Pattern 4: Assuming Exclusive Channel Access..... | 130 |
| 23.8 Debugging Workflow..... | 130 |
| 23.8.1 Issue: Descriptor Not Fetched..... | 130 |
| 23.8.2 Issue: Data Transfer Stalls..... | 130 |
| 23.8.3 Issue: Chained Descriptors Not Following..... | 131 |
| 23.9 Testing Guidance..... | 131 |
| 23.9.1 Test Organization..... | 131 |
| 23.9.2 Test Levels..... | 131 |
| 23.10 Key Documentation Links..... | 131 |
| 23.10.1 Always Reference These..... | 131 |
| 23.11 Quick Commands..... | 132 |
| 23.12 Remember..... | 132 |
| 23.13 PDF Generation Location..... | 133 |

List of Figures

No figures in this document.

List of Tables

| | |
|---|----|
| Table 1: Reference Documents..... | 22 |
| Table 2: Terminology and Definitions..... | 23 |
| Table 3: Document Revision History..... | 24 |

1 Stream Has Index

Generated: 2026-01-04

RTL Design Sherpa · Learning Hardware Design Through Practice GitHub · Documentation Index · MIT License

2 Document Information

2.1 STREAM Hardware Architecture Specification

Document Number: STREAM-HAS-001 **Version:** 0.90 **Status:** Draft **Classification:** Open Source - Apache 2.0 License

2.2 Document Purpose

This Hardware Architecture Specification (HAS) provides a high-level architectural overview of the STREAM (Scatter-gather Transfer Rapid Engine for AXI Memory) subsystem. It describes the system-level design, external interfaces, performance characteristics, and integration requirements without detailing internal implementation specifics.

Target Audience: - System architects evaluating STREAM for integration - Hardware engineers planning system-level integration - Software engineers developing drivers and firmware - Verification engineers planning system-level testing

Companion Documents: - STREAM Micro-Architecture Specification (MAS) - Detailed block-level implementation - STREAM Product Requirements Document (PRD) - Requirements and rationale

2.3 References

Reference Documents

| ID | Document | Description |
|---------|------------------|----------------------------------|
| [REF-1] | STREAM MAS v0.90 | Micro-Architecture Specification |
| [REF-2] | STREAM PRD | Product |

| ID | Document | Description |
|---------|---------------|-----------------------------|
| | | Requirements Document |
| [REF-3] | ARM AMBA AXI4 | AXI4 Protocol Specification |
| [REF-4] | ARM AMBA APB | APB Protocol Specification |

2.4 Terminology

Terminology and Definitions

| Term | Definition |
|----------------|---|
| AXI | Advanced eXtensible Interface - ARM AMBA high-performance bus |
| APB | Advanced Peripheral Bus - ARM AMBA low-power configuration bus |
| Beat | Single data transfer on AXI bus (one clock cycle of valid data) |
| Burst | Sequence of consecutive beats forming a single AXI transaction |
| Channel | Independent DMA transfer context (STREAM supports 8 channels) |
| Descriptor | 256-bit data structure defining a single DMA transfer operation |
| DMA | Direct Memory Access - data transfer without CPU involvement |
| HAS | Hardware Architecture Specification |
| MAS | Micro-Architecture Specification |
| MonBus | Monitor Bus - internal debug/trace event streaming interface |
| Scatter-Gather | DMA mode using linked descriptors for non-contiguous transfers |

2.5 Revision History

Document Revision History

| Version | Date | Author | Description |
|---------|------------|--------|---------------------|
| 0.90 | 2026-01-03 | seang | Initial HAS release |

Last Updated: 2026-01-03

RTL Design Sherpa · Learning Hardware Design Through Practice GitHub · Documentation Index · MIT License

3 Purpose and Scope

3.1 Document Purpose

This Hardware Architecture Specification (HAS) defines the high-level architecture of the STREAM (Scatter-gather Transfer Rapid Engine for AXI Memory) subsystem. It serves as the primary reference for:

- **System Integration** - Understanding STREAM's external interfaces and system-level requirements
- **Performance Planning** - Evaluating throughput, latency, and resource characteristics
- **Driver Development** - Programming model and software interface requirements
- **Verification Planning** - System-level test scenarios and coverage requirements

This document complements the STREAM Micro-Architecture Specification (MAS), which provides detailed block-level implementation specifics.

3.2 Scope

3.2.1 In Scope

This specification covers:

1. System Architecture

- Block-level functional partitioning
- Data flow through the subsystem
- Multi-channel operation model

2. External Interfaces

- APB configuration slave interface
- AXI4 master interfaces (descriptor fetch, data read, data write)
- Monitor bus (MonBus) output interface

3. Performance Characteristics

- Throughput targets and limiting factors
- Latency breakdown for typical operations
- Resource estimates for FPGA/ASIC targets

4. Integration Requirements

- Clock and reset requirements
- System-level constraints
- Verification strategy overview

3.2.2 Out of Scope

This specification does not cover:

- **Detailed Micro-Architecture** - See STREAM MAS for block-level implementation
 - **RTL Implementation Details** - Register-level timing, pipeline stages
 - **Testbench Architecture** - See verification documentation
 - **Software Driver Implementation** - See software programming guide
-

3.3 Design Philosophy

3.3.1 Tutorial-Focused Simplicity

STREAM is intentionally designed as a beginner-friendly DMA engine tutorial. Key simplifications include:

| Aspect | Simplification | Rationale |
|----------------------------|---------------------------|---------------------------------|
| Address Alignment | Aligned addresses only | Eliminates complex fixup logic |
| Transfer Length | Specified in beats | Simplifies datapath math |
| Descriptor Chaining | Linear chains only | No circular buffer management |
| Flow Control | Simple transaction limits | No credit management complexity |

These simplifications make STREAM an ideal learning platform while maintaining production-quality RTL practices.

3.3.2 Relationship to RAPIDS

STREAM is derived from the RAPIDS (Rapid AXI Programmable In-band Descriptor System) architecture but removes:

- Network interfaces (source/sink paths)
- Address alignment fixup logic
- Credit-based flow control
- Control read/write engines

This creates a pure memory-to-memory DMA engine suitable for understanding core DMA concepts.

3.4 Intended Audience

| Audience | Primary Use |
|--------------------|--|
| System Architects | Evaluate STREAM for system integration |
| Hardware Engineers | Plan physical integration and timing closure |

| Audience | Primary Use |
|------------------------|---------------------------------|
| Software Engineers | Develop drivers and firmware |
| Verification Engineers | Plan system-level test coverage |

Last Updated: 2026-01-03

RTL Design Sherpa · Learning Hardware Design Through Practice GitHub · Documentation Index · MIT License

4 Document Conventions

4.1 Notation Conventions

4.1.1 Signal Naming

| Prefix/Suffix | Meaning | Example |
|---------------|--------------------|------------------|
| i_ | Input signal | i_clk, i_data |
| o_ | Output signal | o_valid, o_ready |
| r_ | Registered signal | r_state, r_count |
| w_ | Combinational wire | w_next_state |
| _n | Active-low signal | aresetn, rst_n |

4.1.2 Interface Prefixes

| Prefix | Protocol | Direction |
|-------------|------------------------|-----------|
| s_apb_ | APB slave | Inbound |
| m_axi_desc_ | AXI4 descriptor master | Outbound |
| m_axi_rd_ | AXI4 read master | Outbound |
| m_axi_wr_ | AXI4 write master | Outbound |
| monbus_ | Monitor bus | Outbound |

4.2 Diagram Conventions

4.2.1 Block Diagrams

- **Rectangles** - Functional blocks
- **Arrows** - Data flow direction
- **Dashed Lines** - Configuration/control paths
- **Double Lines** - High-bandwidth data paths

4.2.2 Timing Diagrams

- **High** - Logic 1
- **Low** - Logic 0
- **X** - Don't care / undefined
- **Hatched** - Valid data window

4.3 Numeric Conventions

4.3.1 Units

| Unit | Meaning |
|-------------|--|
| Beat | Single data transfer (one clock cycle of valid data) |
| Burst | Sequence of consecutive beats |
| Transaction | Complete AXI transaction (address + all data beats) |

4.3.2 Address Notation

- All addresses are specified in bytes unless otherwise noted
- Address alignment requirements are specified relative to data width
- Example: 512-bit data width requires 64-byte (512/8) alignment

4.3.3 Size Notation

| Notation | Value |
|----------|-----------------|
| KB | 1,024 bytes |
| MB | 1,048,576 bytes |
| Kbit | 1,024 bits |

4.4 Requirement Notation

Requirements and constraints are identified using the following format:

| Prefix | Category |
|-----------|-------------------------|
| [REQ-x] | Functional requirement |
| [PERF-x] | Performance requirement |
| [INTEG-x] | Integration requirement |
| [CONST-x] | Design constraint |

4.5 Reference Format

Cross-references to other documents use the format:

- [MAS-x.y] - Reference to MAS section x.y
 - [REF-n] - Reference to document in reference table
-

Last Updated: 2026-01-03

RTL Design Sherpa · Learning Hardware Design Through Practice GitHub · Documentation Index · MIT License

5 Definitions and Acronyms

5.1 Acronyms

| Acronym | Definition |
|---------|---|
| APB | Advanced Peripheral Bus |
| AXI | Advanced eXtensible Interface |
| AMBA | Advanced Microcontroller Bus Architecture |
| BFM | Bus Functional Model |
| CDC | Clock Domain Crossing |

| Acronym | Definition |
|---------|---|
| DMA | Direct Memory Access |
| FIFO | First-In First-Out buffer |
| FSM | Finite State Machine |
| HAS | Hardware Architecture Specification |
| MAS | Micro-Architecture Specification |
| MonBus | Monitor Bus (internal debug/trace interface) |
| PRD | Product Requirements Document |
| RTL | Register Transfer Level |
| SG | Scatter-Gather |
| SRAM | Static Random Access Memory |
| STREAM | Scatter-gather Transfer Rapid Engine for AXI Memory |

5.2 Definitions

5.2.1 Architectural Terms

| Term | Definition |
|-------------------|--|
| Beat | A single data transfer on an AXI bus, corresponding to one clock cycle where valid data is transferred. For a 512-bit data bus, one beat transfers 64 bytes. |
| Burst | A sequence of consecutive beats forming a single AXI transaction. STREAM supports burst lengths up to 256 beats (AXI4 maximum). |
| Channel | An independent DMA transfer context. STREAM supports 8 channels, each capable of processing its own descriptor chain. |
| Descriptor | A 256-bit data structure in memory that defines a single DMA transfer operation, including source address, destination address, length, and chaining |

| Term | Definition |
|-------------------------|---|
| Descriptor Chain | information. |
| Scatter-Gather | A linked list of descriptors where each descriptor points to the next, enabling complex scatter-gather transfer patterns. |
| Scatter-Gather | A DMA technique where data is transferred to/from non-contiguous memory locations using a list of descriptors. |

5.2.2 Interface Terms

| Term | Definition |
|---------------------|---|
| APB Slave | STREAM's configuration interface, receiving programming commands from the system processor. |
| AXI Master | STREAM's memory interfaces for descriptor fetches and data transfers. Three masters: descriptor fetch, data read, data write. |
| MonBus | Monitor Bus - a 64-bit streaming interface for debug and performance monitoring packets. |
| Backpressure | Flow control mechanism where a receiver signals inability to accept data, causing the sender to pause. |

5.2.3 Transfer Terms

| Term | Definition |
|------------------------|--|
| Aligned Address | An address that is a multiple of the data width in bytes. For 512-bit (64-byte) data width, aligned addresses are multiples of 64. |
| Kick-off | Initiating a DMA transfer by writing the first descriptor address to a channel's control register. |
| Completion | The successful end of a descriptor chain, optionally generating an interrupt. |
| Chaining | Following the <code>next_descriptor_ptr</code> to fetch and process subsequent descriptors. |

5.3 Units and Measurements

| Term | Definition |
|--------------------|--|
| Clock Cycle | One period of the primary system clock (aclk). |
| Latency | Time from initiating an operation to its completion, typically measured in clock cycles. |
| Throughput | Data transfer rate, typically measured in bytes per second or beats per cycle. |
| Utilization | Percentage of available bandwidth actually used for data transfer. |

5.4 Document-Specific Terms

| Term | Definition |
|-------------------------------|---|
| External Interface | A signal or bus that crosses the STREAM subsystem boundary. |
| Internal Interface | A signal or bus between blocks within STREAM, not visible externally. |
| Configuration Register | An APB-accessible register for controlling STREAM behavior. |
| Status Register | An APB-readable register reporting STREAM operational state. |

Last Updated: 2026-01-03

RTL Design Sherpa · Learning Hardware Design Through Practice GitHub · Documentation Index · MIT License

6 Use Cases

6.1 Primary Use Cases

6.1.1 Memory-to-Memory Data Movement

Description: Transfer data between two memory regions without CPU involvement.

Operation: 1. Software prepares descriptor(s) in memory with source/destination addresses 2. Software writes first descriptor address to channel control register 3. STREAM autonomously fetches descriptor and executes transfer 4. Optional interrupt notifies software of completion

Key Benefits: - CPU offloading for bulk data transfers - Deterministic transfer timing - Support for large transfer sizes via descriptor chaining

6.1.2 Scatter-Gather Operations

Description: Collect data from multiple non-contiguous source regions into a contiguous destination, or distribute contiguous source data to multiple destinations.

Gather Example:

```
Descriptor 0: src=0x1000, dst=0x8000, len=64 beats
Descriptor 1: src=0x3000, dst=0x8100, len=32 beats (chained)
Descriptor 2: src=0x5000, dst=0x8180, len=16 beats (chained, last)
```

Scatter Example:

```
Descriptor 0: src=0x8000, dst=0x1000, len=64 beats
Descriptor 1: src=0x8100, dst=0x3000, len=32 beats (chained)
Descriptor 2: src=0x8180, dst=0x5000, len=16 beats (chained, last)
```

6.1.3 Multi-Channel Concurrent Transfers

Description: Execute multiple independent transfer operations simultaneously.

Operation: - Up to 8 channels operate concurrently - Each channel processes its own descriptor chain - Shared resources (SRAM, AXI masters) arbitrated by priority

Use Cases: - Parallel DMA for multi-stream applications - Priority-based bandwidth allocation - Independent channel error isolation

6.2 Secondary Use Cases

6.2.1 Buffer Management

Description: Move data between system memory and local buffers for processing pipelines.

Pattern:

System Memory → STREAM → Processing Buffer → STREAM → System Memory

6.2.2 Memory Initialization

Description: Fill memory regions with pattern data using chained descriptors.

Note: Requires source buffer containing pattern data.

6.3 Target Applications

| Application | Channel Usage | Typical Transfer Size |
|-------------------------|---------------|-----------------------|
| Video Frame Copy | 1-2 channels | 1-8 MB per frame |
| Audio Buffer Management | 1 channel | 4-64 KB per buffer |
| Data Collection | 4-8 channels | Variable |
| Memory Test Patterns | 1-4 channels | System-dependent |

Last Updated: 2026-01-03

RTL Design Sherpa · Learning Hardware Design Through Practice GitHub · Documentation Index · MIT License

7 Key Features

7.1 Feature Summary

| Feature | Specification |
|-----------------|------------------------|
| Channels | 8 independent channels |

| Feature | Specification |
|------------------------|------------------------------------|
| Descriptor Size | 256 bits (32 bytes) |
| Data Width | Parameterizable (default 512 bits) |
| Address Width | 64 bits |
| Maximum Burst | 256 beats (AXI4) |
| Chaining | Unlimited descriptor chain depth |
| Arbitration | Priority-based with round-robin |

7.2 Descriptor-Based Operation

7.2.1 Single-Write Kick-off

A single APB write to a channel's control register initiates an entire transfer sequence:

1. Write descriptor address to CHn_CTRL register
2. STREAM automatically fetches and processes descriptor
3. Chained descriptors processed without software intervention
4. Completion indicated via status register and optional interrupt

7.2.2 Autonomous Chaining

STREAM follows next_descriptor_ptr automatically:

- **Non-zero pointer:** Fetch next descriptor, continue processing
 - **Zero pointer:** Chain complete, transition to idle
 - **Last flag:** Explicit chain termination regardless of pointer
-

7.3 Multi-Channel Architecture

7.3.1 Channel Independence

Each channel maintains:

- Independent FSM state
- Separate descriptor chain pointer
- Individual error status

- Private completion interrupt

7.3.2 Resource Sharing

All channels share:

- Descriptor fetch AXI master
- Data read AXI master
- Data write AXI master
- Internal SRAM buffer
- MonBus reporter

7.3.3 Priority-Based Arbitration

- 8-bit priority field in each descriptor
 - Higher priority descriptors serviced first
 - Round-robin within same priority level
 - Anti-starvation timeout mechanism
-

7.4 Data Path Features

7.4.1 Parameterizable Data Width

Compile-time configurable data width:

| Parameter | Description | Typical Values |
|------------|-----------------------|--------------------------|
| DATA_WIDTH | AXI data bus width | 128, 256, 512 bits |
| SRAM_DEPTH | Internal buffer depth | 1024, 2048, 4096 entries |

7.4.2 Aligned Address Requirement

[CONST-1] All source and destination addresses must be aligned to data width.

| Data Width | Alignment Requirement |
|------------|-----------------------|
| 128 bits | 16-byte aligned |
| 256 bits | 32-byte aligned |
| 512 bits | 64-byte aligned |

This simplification eliminates complex alignment fixup logic for tutorial clarity.

7.5 Monitoring and Debug

7.5.1 MonBus Integration

64-bit monitor packets for:

- Transfer initiation events
- Descriptor fetch events
- Transfer completion events
- Error condition reporting
- Performance metrics

7.5.2 Status Visibility

APB-readable status registers provide:

- Channel state (idle, active, error)
- Descriptor count per channel
- Error codes and flags
- Performance counters (optional)

7.6 Error Handling

7.6.1 Error Detection

| Error Type | Detection | Response |
|----------------------|--------------------|----------------------------|
| AXI read error | RRESP != OKAY | Stop chain, set error flag |
| AXI write error | BRESP != OKAY | Stop chain, set error flag |
| Invalid descriptor | Validation failure | Stop chain, set error flag |
| Address out of range | Limit check | Stop chain, set error flag |

7.6.2 Error Isolation

Errors in one channel do not affect other channels. Each channel has independent error status and recovery.

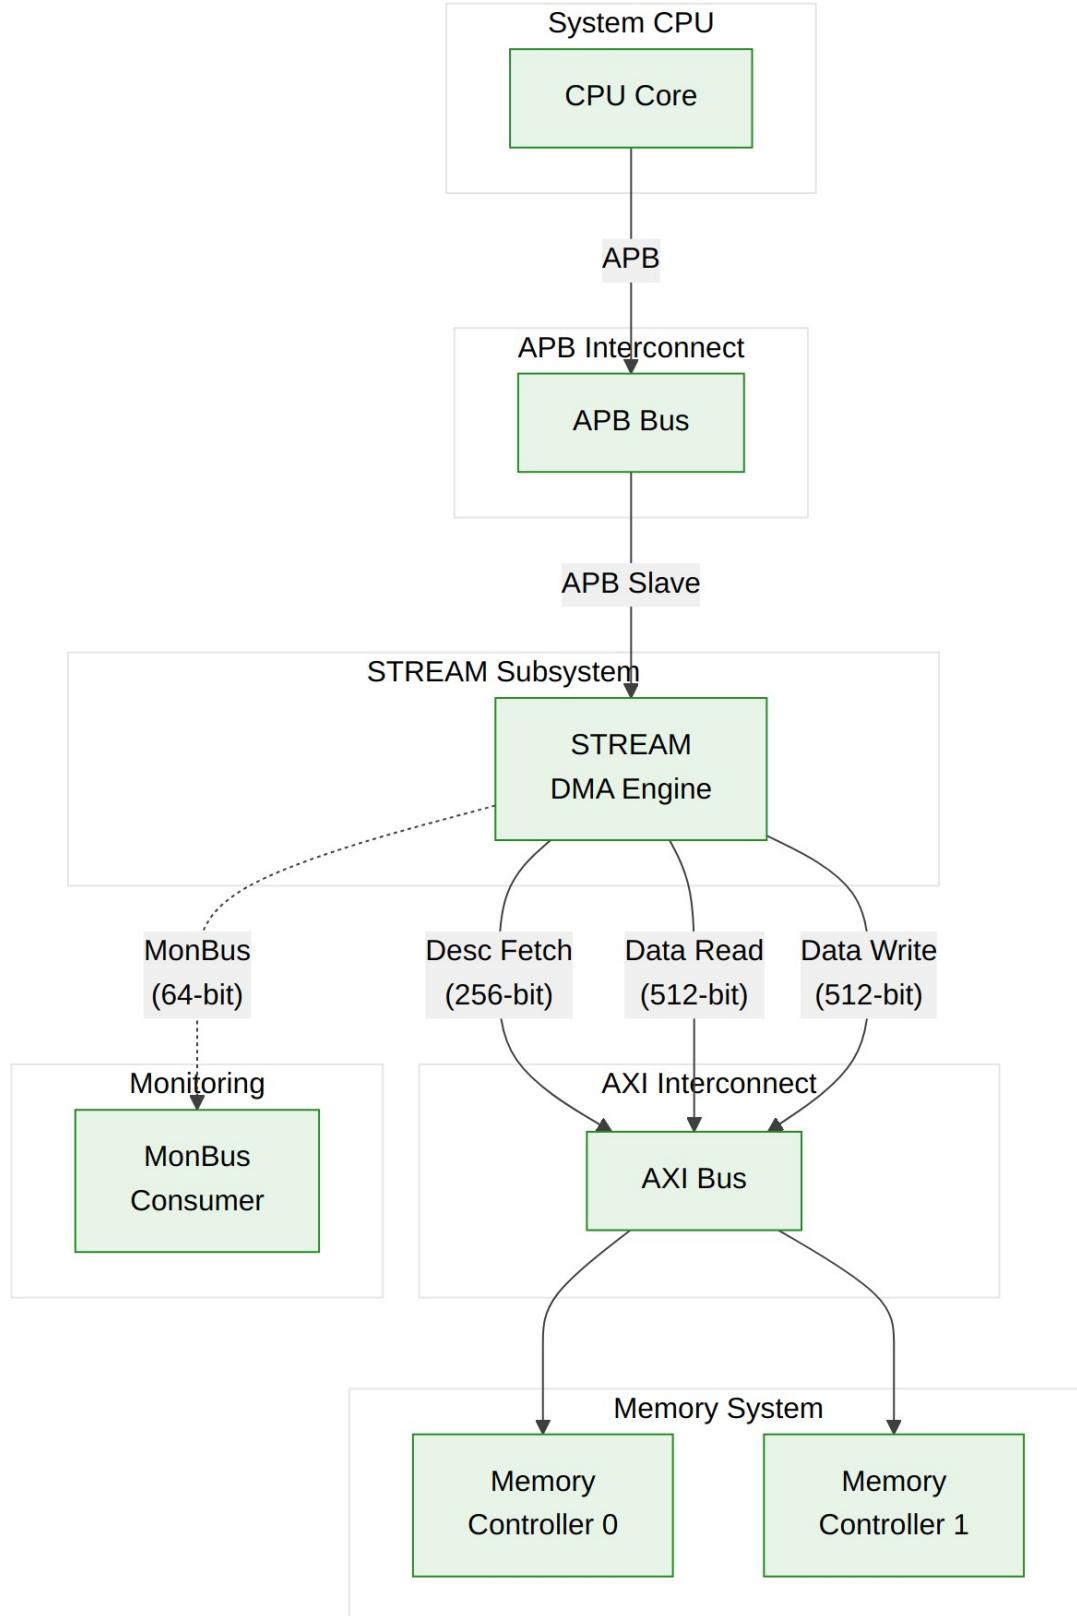
Last Updated: 2026-01-03

RTL Design Sherpa · Learning Hardware Design Through Practice GitHub · Documentation Index · MIT License

8 System Context

8.1 System Integration Overview

STREAM integrates into a typical SoC as a peripheral subsystem with the following connections:



System Context Diagram

Source: [01_system_context.mmd](#)

8.2 External Connections

8.2.1 Configuration Path

| Connection | Protocol | Description |
|------------|-----------|-------------------------------------|
| APB Slave | APB3/APB4 | Configuration and control interface |

Functions: - Channel kick-off (write descriptor address) - Status monitoring - Interrupt control - Error handling

8.2.2 Data Paths

| Connection | Protocol | Width | Description |
|-------------------|----------|---------|----------------------------------|
| Descriptor Master | AXI4 | 256-bit | Fetch descriptors from memory |
| Read Data Master | AXI4 | Param | Read source data from memory |
| Write Data Master | AXI4 | Param | Write destination data to memory |

8.2.3 Monitoring Path

| Connection | Protocol | Width | Description |
|---------------|----------|--------|---------------------------|
| MonBus Master | MonBus | 64-bit | Debug/trace packet output |

8.3 Clock and Reset

8.3.1 Clock Domain

STREAM operates in a single clock domain:

| Signal | Description |
|---------|--|
| aclk | Primary system clock, all interfaces synchronous |
| aresetn | Active-low asynchronous reset |

[INTEG-1] All external interfaces must be synchronous to aclk.

8.3.2 Reset Requirements

- Asynchronous assertion, synchronous de-assertion
 - Minimum reset pulse width: 2 clock cycles
 - All state machines return to idle on reset
 - SRAM contents undefined after reset
-

8.4 Memory Map Requirements

8.4.1 Descriptor Memory

Descriptors must be placed in memory accessible via the descriptor fetch AXI master:

- **Alignment:** 32-byte aligned (256-bit descriptor size)
- **Accessibility:** Read access required
- **Coherency:** Software must ensure descriptors are coherent before kick-off

8.4.2 Data Memory

Source and destination regions must be accessible via the appropriate AXI master:

- **Source:** Read access via read data master
 - **Destination:** Write access via write data master
 - **Alignment:** Per data width (see Key Features)
-

8.5 Interrupt Integration

8.5.1 Interrupt Signals

| Signal | Description |
|--------------|--|
| irq[7:0] | Per-channel completion/error interrupt |
| irq_combined | OR of all channel interrupts |

8.5.2 Interrupt Conditions

- Descriptor with interrupt flag completes successfully
 - Channel encounters error condition
 - Software-triggered interrupt (via control register)
-

8.6 Power Considerations

8.6.1 Clock Gating

STREAM supports clock gating when all channels are idle:

- Monitor `all_channels_idle` status bit
- External clock gate can disable `aclk` when idle
- Re-enable clock before writing kick-off register

8.6.2 Power States

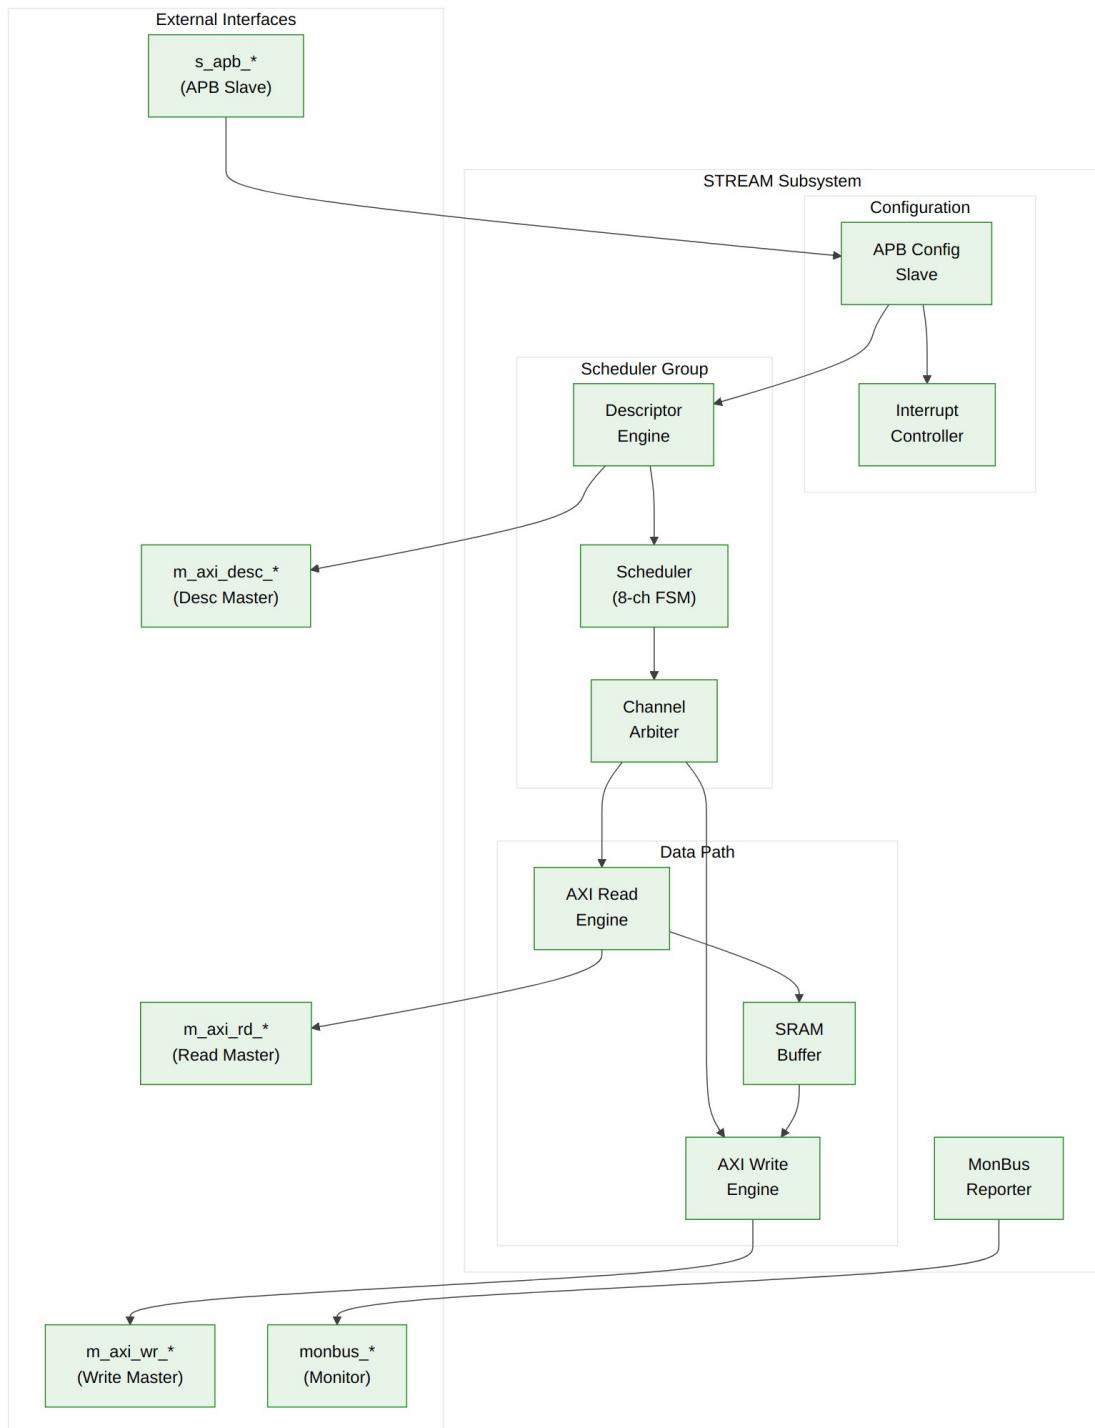
| State | Channels | Clock | Power |
|--------|---------------|---------|------------------------|
| Active | ≥ 1 busy | Running | Full |
| Idle | All idle | Running | Reduced (no switching) |
| Gated | All idle | Stopped | Minimal (leakage only) |

Last Updated: 2026-01-03

RTL Design Sherpa · Learning Hardware Design Through Practice GitHub · Documentation Index · MIT License

9 Block Diagram

9.1 Top-Level Architecture



STREAM Block Diagram

Source: [02_block_diagram.mmd](#)

9.2 Functional Block Summary

| Block | Function | Key Interfaces |
|--------------------------|------------------------------------|-------------------------------|
| APB Config Slave | Channel control registers, status | APB slave, kick-off signals |
| Descriptor Engine | Fetch and parse descriptors | AXI read, descriptor output |
| Scheduler | Coordinate transfer phases | Control signals to all blocks |
| Channel Arbiter | Multi-channel priority arbitration | Grant signals |
| AXI Read Engine | Read source data to SRAM | AXI read master |
| SRAM Buffer | Temporary data storage | Read/write ports |
| AXI Write Engine | Write SRAM data to destination | AXI write master |
| MonBus Reporter | Generate monitoring packets | MonBus output |

9.3 Block Interactions

9.3.1 Control Flow

APB Write -> APB Config -> Descriptor Engine -> Scheduler -> Data Path, with AXI Desc Master fetching descriptors.

9.3.2 Data Flow

Memory -> AXI Read Engine -> SRAM Buffer -> AXI Write Engine -> Memory

9.3.3 Monitoring Flow

All Blocks -> Event Signals -> MonBus Reporter -> MonBus Output

9.4 Parameterization

| Parameter | Description | Default | Range |
|-----------------|-------------------------|---------|---------------|
| NUM_CHANNELS | Number of DMA channels | 8 | 1-8 |
| DATA_WIDTH | AXI data bus width | 512 | 128, 256, 512 |
| ADDR_WIDTH | Address width | 64 | 32, 64 |
| SRAM_DEPTH | Buffer depth in entries | 4096 | 256-8192 |
| DESC_FIFO_DEPTH | Descriptor FIFO depth | 4 | 2-8 |

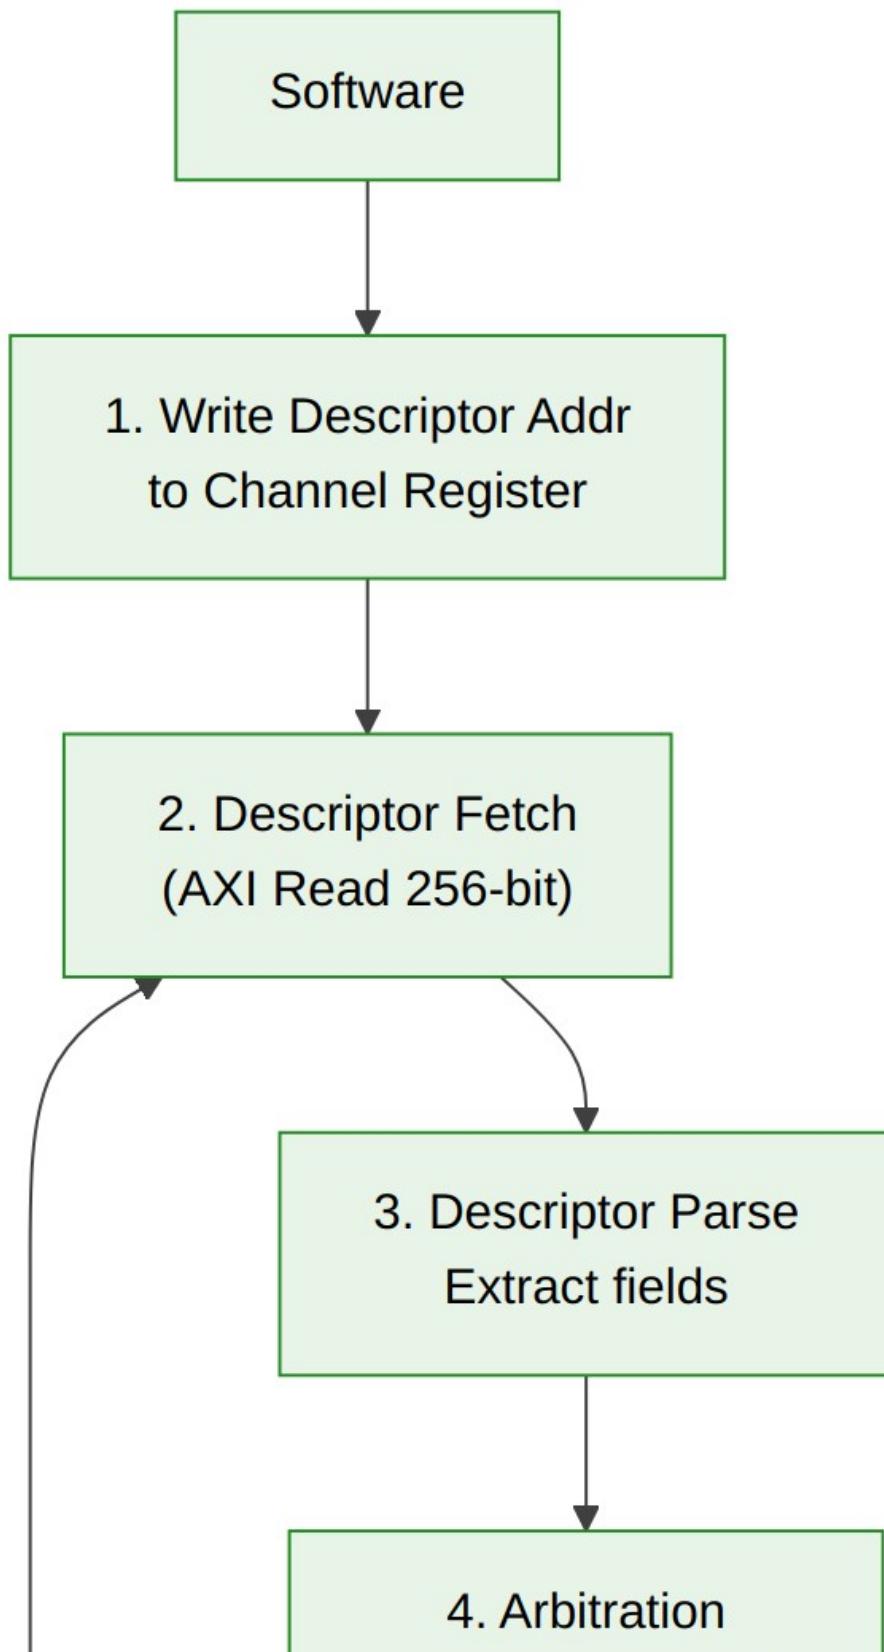
Last Updated: 2026-01-03

RTL Design Sherpa · Learning Hardware Design Through Practice GitHub · Documentation Index · MIT License

10 Data Flow

10.1 Transfer Sequence Overview

A complete STREAM transfer follows this sequence:



Transfer Sequence Flowchart

Source: [03_data_flow.mmd](#)

10.2 Phase Details

10.2.1 Phase 1: Kick-off

Software writes the first descriptor address to the channel control register:

- APB write triggers channel state transition from IDLE to ACTIVE
- Descriptor address stored in channel's descriptor pointer register
- Kick-off signal sent to Descriptor Engine

10.2.2 Phase 2: Descriptor Fetch

Descriptor Engine issues AXI read transaction:

- Address: From kick-off or `next_descriptor_ptr`
- Size: 256 bits (single beat or 2-beat depending on bus width)
- Response: Descriptor data captured on valid read data

10.2.3 Phase 3: Descriptor Parse

Descriptor fields extracted and validated:

| Field | Validation |
|-----------------------|------------------------------|
| <code>valid</code> | Must be 1 |
| <code>src_addr</code> | Alignment check, range check |
| <code>dst_addr</code> | Alignment check, range check |
| <code>length</code> | Non-zero, within limits |

10.2.4 Phase 4: Arbitration

Channel requests access to shared data path:

- Multiple channels may be active simultaneously
- Arbiter grants based on priority and fairness
- Granted channel proceeds to read phase

10.2.5 Phase 5: Read Phase

AXI Read Engine fetches source data:

- Issues AXI read bursts to source address
- Data written to SRAM buffer
- Continues until length beats transferred
- Read/write to SRAM can overlap (ping-pong or streaming)

10.2.6 Phase 6: Write Phase

AXI Write Engine sends data to destination:

- Reads data from SRAM buffer
- Issues AXI write bursts to destination address
- Awaits write responses
- Continues until all data written

10.2.7 Phase 7: Chain Check

Scheduler determines next action:

- If `last == 1` or `next_descriptor_ptr == 0`: Chain complete
 - Otherwise: Fetch next descriptor (return to Phase 2)
-

10.3 Concurrent Operation

10.3.1 Multi-Channel Concurrency

Multiple channels can be at different phases simultaneously. For example: - Ch0: Kick -> Fetch -> Read -> Write -> Complete - Ch1: (offset) Kick -> Fetch -> Arb Wait -> Read -> Write - Ch2: (offset) Kick -> Fetch -> Arb Wait -> Read

10.3.2 Pipeline Overlap

Within a single channel, read and write phases can overlap. As read beats 0-63 complete, write beats 0-63 can begin while read continues with beats 64-127.

The SRAM buffer enables this overlap by storing in-flight data.

10.4 Data Path Bandwidth

10.4.1 Theoretical Maximum

With 512-bit data width at 200 MHz:

- Per-direction: $512 \text{ bits} \times 200 \text{ MHz} = 102.4 \text{ Gbps} = 12.8 \text{ GB/s}$
- Bidirectional: 25.6 GB/s (simultaneous read + write)

10.4.2 Practical Considerations

Actual throughput affected by:

- AXI arbitration latency
 - Memory access latency
 - SRAM buffer depth (limits outstanding data)
 - Channel arbitration overhead
-

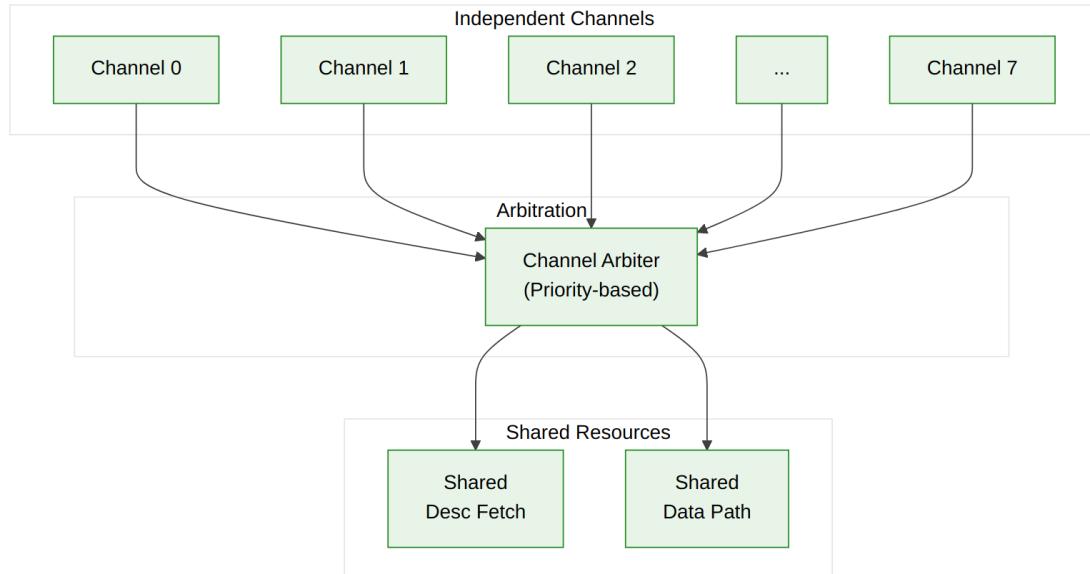
Last Updated: 2026-01-03

RTL Design Sherpa · Learning Hardware Design Through Practice GitHub · Documentation Index · MIT License

11 Multi-Channel Architecture

11.1 Channel Organization

STREAM supports 8 independent DMA channels, each capable of processing its own descriptor chain:



Channel Architecture

Source: [04_channel_architecture.mmd](#)

11.2 Per-Channel Resources

Each channel maintains independent state:

| Resource | Description | Size |
|---------------------------|-------------------------------------|---------|
| FSM State | Current operation phase | 4 bits |
| Descriptor Pointer | Current/next descriptor address | 64 bits |
| Transfer Counter | Beats remaining in current transfer | 32 bits |
| Error Status | Error flags and codes | 8 bits |
| Priority | Current descriptor priority | 8 bits |

11.3 Shared Resources

All channels share the following resources:

| Resource | Arbitration | Capacity |
|------------------------------|------------------------|---------------------------|
| Descriptor AXI Master | Round-robin | 1 outstanding transaction |
| Data Read AXI Master | Priority-based | Configurable outstanding |
| Data Write AXI Master | Priority-based | Configurable outstanding |
| SRAM Buffer | Partitioned by channel | SRAM_DEPTH entries |
| MonBus Reporter | FIFO-based | 64-entry FIFO |

11.4 Arbitration Scheme

11.4.1 Priority-Based Arbitration

Descriptor priority field (8 bits) determines service order:

- **Priority 255:** Highest priority
- **Priority 0:** Lowest priority
- **Same Priority:** Round-robin among equal priority

11.4.2 Fairness Mechanism

To prevent starvation:

- Maximum service time per grant (configurable)
- Anti-starvation counter promotes low-priority after timeout
- Emergency preemption for high-priority arrivals (optional)

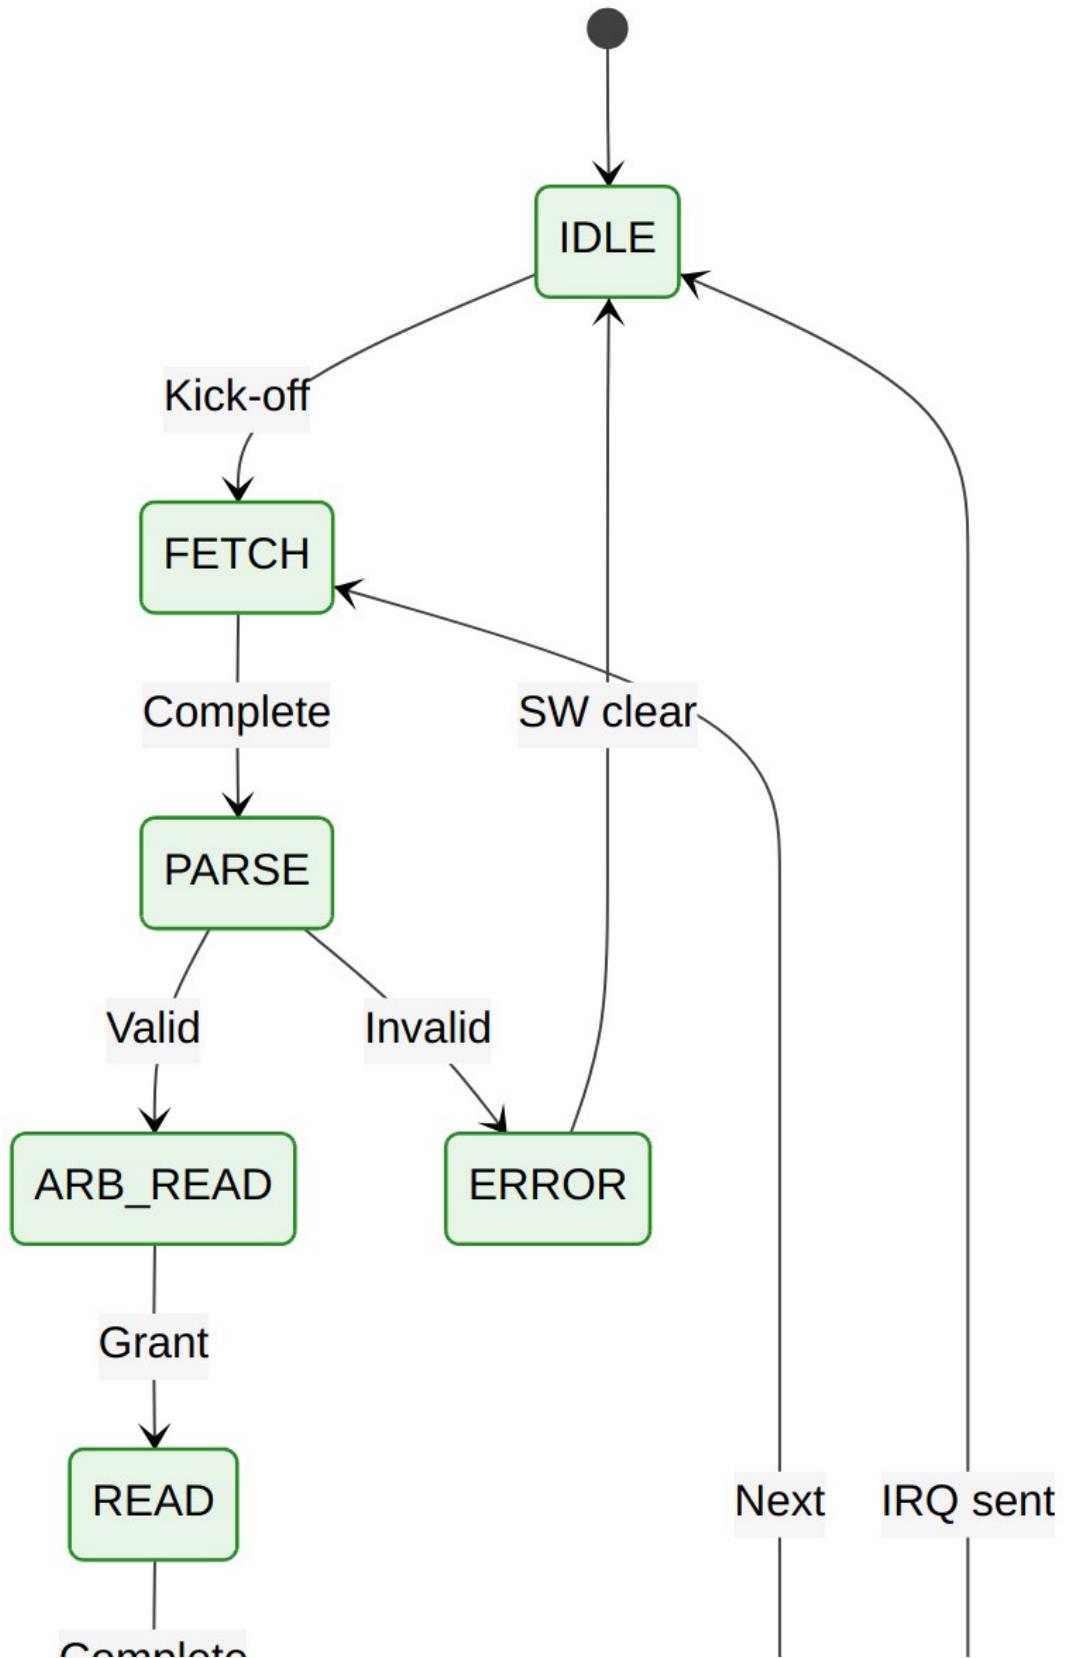
11.4.3 Arbitration Points

Arbitration occurs at:

1. **Descriptor Fetch:** When channel needs to fetch descriptor
 2. **Read Phase Start:** When channel ready to read source data
 3. **Write Phase Start:** When channel ready to write destination
-

11.5 Channel States

Each channel FSM traverses these states:



Channel FSM State Diagram

Source: [05_channel_fsm.mmd](#)

| State | Description | Transitions |
|------------------|-----------------------|-------------------------------------|
| IDLE | No active transfer | Kick-off -> FETCH |
| FETCH | Fetching descriptor | Complete -> PARSE |
| PARSE | Validating descriptor | Valid -> ARB_READ, Invalid -> ERROR |
| ARB_READ | Awaiting read grant | Grant -> READ |
| READ | Reading source data | Complete -> ARB_WRITE |
| ARB_WRITE | Awaiting write grant | Grant -> WRITE |
| WRITE | Writing destination | Complete -> CHAIN |
| CHAIN | Checking for next | Next -> FETCH, Last -> COMPLETE |
| COMPLETE | Transfer done | IRQ sent -> IDLE |
| ERROR | Error condition | SW clear -> IDLE |

11.6 SRAM Partitioning

The shared SRAM buffer is partitioned among active channels:

11.6.1 Static Partitioning (Default)

SRAM Buffer (4096 entries) is divided equally among 8 channels:

| Partition | Entries |
|-----------|---------|
| Channel 0 | 512 |
| Channel 1 | 512 |
| Channel 2 | 512 |
| Channel 3 | 512 |
| Channel 4 | 512 |
| Channel 5 | 512 |
| Channel 6 | 512 |
| Channel 7 | 512 |

11.6.2 Dynamic Allocation (Optional)

When fewer than 8 channels are active, partitions can be expanded:

| Active Channels | Entries Per Channel |
|-----------------|---------------------|
| 1 | 4096 |
| 2 | 2048 |
| 4 | 1024 |
| 8 | 512 |

11.7 Error Isolation

Errors are isolated to the affected channel:

- Error in Channel N does not affect Channels 0..(N-1), (N+1)..7
- Shared resource errors logged but don't block other channels
- Software must clear error and re-kick to resume

Last Updated: 2026-01-03

RTL Design Sherpa · Learning Hardware Design Through Practice GitHub · Documentation Index · MIT License

12 AXI Master Interfaces

12.1 Overview

STREAM includes three AXI4 master interfaces for memory access:

| Interface | Purpose | Data Width | Address Width |
|---------------------------|------------------|-----------------|---------------|
| <code>m_axi_desc_*</code> | Descriptor fetch | 256 bits | 64 bits |
| <code>m_axi_rd_*</code> | Source data read | Parameterizable | 64 bits |
| <code>m_axi_wr_*</code> | Destination | Parameterizable | 64 bits |

| Interface | Purpose | Data Width | Address Width |
|-----------|---------|------------|---------------|
| | write | le | |

12.2 Descriptor Fetch Master

12.2.1 Signal Summary

| Signal | Direction | Width | Description |
|---------------------|-----------|----------|---------------------------|
| m_axi_desc_ar_id | Output | ID_WIDTH | Read transaction ID |
| m_axi_desc_ar_addr | Output | 64 | Read address |
| m_axi_desc_ar_len | Output | 8 | Burst length (0 = 1 beat) |
| m_axi_desc_ar_size | Output | 3 | Beat size (5 = 32 bytes) |
| m_axi_desc_ar_burst | Output | 2 | Burst type (INCR) |
| m_axi_desc_ar_valid | Output | 1 | Address valid |
| m_axi_desc_ar_ready | Input | 1 | Address ready |
| m_axi_desc_rd_id | Input | ID_WIDTH | Read data ID |
| m_axi_desc_rd_ata | Input | 256 | Read data |
| m_axi_desc_rr_esp | Input | 2 | Read response |
| m_axi_desc_rl_ast | Input | 1 | Last beat |
| m_axi_desc_rv_alid | Input | 1 | Read data valid |
| m_axi_desc_rr_ready | Output | 1 | Read data ready |

12.2.2 Transaction Characteristics

| Characteristic | Value | Notes |
|---------------------|--------|----------------|
| Burst Length | 1 beat | Single 256-bit |

| Characteristic | Value | Notes |
|--------------------|----------|-----------------|
| | | descriptor |
| Burst Type | INCR | Fixed increment |
| Outstanding | 1-2 | Configurable |
| Ordering | In-order | Single ID used |

12.2.3 Address Alignment

Descriptor addresses must be 32-byte aligned (256-bit boundary).

12.3 Data Read Master

12.3.1 Signal Summary

| Signal | Direction | Width | Description |
|------------------------------|-----------|------------|---------------------|
| m_axi_rd_arid | Output | ID_WIDTH | Read transaction ID |
| m_axi_rd_arad _{dr} | Output | 64 | Read address |
| m_axi_rd_arle _n | Output | 8 | Burst length |
| m_axi_rd_arsize | Output | 3 | Beat size |
| m_axi_rd_arbu _{rst} | Output | 2 | Burst type (INCR) |
| m_axi_rd_arva _{lid} | Output | 1 | Address valid |
| m_axi_rd_arre _{ady} | Input | 1 | Address ready |
| m_axi_rd_rid | Input | ID_WIDTH | Read data ID |
| m_axi_rd_rdat _a | Input | DATA_WIDTH | Read data |
| m_axi_rd_rres _p | Input | 2 | Read response |
| m_axi_rd_rlas _t | Input | 1 | Last beat |
| m_axi_rd_rval _{id} | Input | 1 | Read data valid |

| Signal | Direction | Width | Description |
|------------------|-----------|-------|-----------------|
| m_axi_rd_rrea_dy | Output | 1 | Read data ready |

12.3.2 Transaction Characteristics

| Characteristic | Value | Notes |
|---------------------|--------------|--------------------------|
| Burst Length | 1-256 beats | AXI4 maximum |
| Burst Type | INCR | Fixed increment |
| Outstanding | Configurable | 4-16 typical |
| Ordering | In-order | Channel-based ID tagging |

12.3.3 Burst Generation

Large transfers are broken into maximum-length bursts:

| Transfer Size | Burst Count | Burst Lengths |
|---------------|-------------|---------------|
| 256 beats | 1 | 256 |
| 512 beats | 2 | 256, 256 |
| 300 beats | 2 | 256, 44 |

12.4 Data Write Master

12.4.1 Signal Summary

| Signal | Direction | Width | Description |
|-------------------|-----------|----------|----------------------|
| m_axi_wr_awid | Output | ID_WIDTH | Write transaction ID |
| m_axi_wr_awad_dr | Output | 64 | Write address |
| m_axi_wr_awle_n | Output | 8 | Burst length |
| m_axi_wr_awsi_ze | Output | 3 | Beat size |
| m_axi_wr_awbu_RST | Output | 2 | Burst type (INCR) |
| m_axi_wr_awva_lid | Output | 1 | Address valid |

| Signal | Direction | Width | Description |
|------------------------------|------------------|--------------|--------------------|
| m_axi_wr_awre _{ady} | Input | 1 | Address ready |
| m_axi_wr_wdat _a | Output | DATA_WIDTH | Write data |
| m_axi_wr_wstr _b | Output | DATA_WIDTH/8 | Byte strobes |
| m_axi_wr_wlas _t | Output | 1 | Last beat |
| m_axi_wr_wval _{id} | Output | 1 | Write data valid |
| m_axi_wr_wrea _{dy} | Input | 1 | Write data ready |
| m_axi_wr_bid | Input | ID_WIDTH | Response ID |
| m_axi_wr_bres _p | Input | 2 | Write response |
| m_axi_wr_bval _{id} | Input | 1 | Response valid |
| m_axi_wr_brea _{dy} | Output | 1 | Response ready |

12.4.2 Transaction Characteristics

| Characteristic | Value | Notes |
|-----------------------|--------------|------------------------|
| Burst Length | 1-256 beats | Matches read bursts |
| Burst Type | INCR | Fixed increment |
| Outstanding | Configurable | 4-16 typical |
| Write Strobes | All-ones | Aligned transfers only |

12.5 Common AXI Attributes

12.5.1 Supported Features

| Feature | Descriptor | Read | Write |
|----------------|-------------------|-------------|--------------|
| INCR Burst | Yes | Yes | Yes |
| WRAP Burst | No | No | No |

| Feature | Descriptor | Read | Write |
|-------------|------------|------|-------|
| FIXED Burst | No | No | No |
| Exclusive | No | No | No |
| Locked | No | No | No |

12.5.2 Error Handling

- SLVERR/DECERR on any channel stops the current transfer
 - Error logged in channel status register
 - Subsequent descriptors in chain not processed
 - Software intervention required to clear and resume
-

Last Updated: 2026-01-03

RTL Design Sherpa · Learning Hardware Design Through Practice GitHub · Documentation Index · MIT License

13 APB Configuration Interface

13.1 Overview

STREAM provides an APB slave interface for software configuration and control. The interface supports:

- Channel kick-off (write descriptor address to start transfer)
 - Status monitoring
 - Interrupt control
 - Error handling
-

13.2 Signal Summary

| Signal | Direction | Width | Description |
|-------------|-----------|-------|-------------|
| s_apb_paddr | Input | 12 | Address bus |
| s_apb_psel | Input | 1 | Peripheral |

| Signal | Direction | Width | Description |
|---------------|------------------|--------------|--------------------|
| | | | select |
| s_apb_penable | Input | 1 | Enable phase |
| s_apb_pwrite | Input | 1 | Write enable |
| s_apb_pwdata | Input | 32 | Write data |
| s_apb_pstrb | Input | 4 | Byte strobes |
| s_apb_pready | Output | 1 | Ready signal |
| s_apb_prdata | Output | 32 | Read data |
| s_apb_pslverr | Output | 1 | Slave error |

13.3 Register Map Overview

| Offset | Register | Access | Description |
|---------------|-----------------|---------------|------------------------------|
| 0x000 | CTRL | RW | Global control register |
| 0x004 | STATUS | RO | Global status register |
| 0x008 | IRQ_EN | RW | Interrupt enable mask |
| 0x00C | IRQ_STATUS | RW1C | Interrupt status |
| 0x010 | ERR_STATUS | RO | Error status |
| 0x014 | ERR_ADDR | RO | Error address capture |
| 0x020-0x03C | Reserved | - | Reserved |
| 0x040 | CH0_CTRL | RW | Channel 0 control |
| 0x044 | CH0_STATUS | RO | Channel 0 status |
| 0x048 | CH0_DESC_PTR | RO | Channel 0 current descriptor |
| 0x04C | Reserved | - | Reserved |

| Offset | Register | Access | Description |
|--------|--------------|--------|------------------------------|
| 0x050 | CH1_CTRL | RW | Channel 1 control |
| ... | ... | ... | ... |
| 0x0B0 | CH7_CTRL | RW | Channel 7 control |
| 0x0B4 | CH7_STATUS | RO | Channel 7 status |
| 0x0B8 | CH7_DESC_PTR | RO | Channel 7 current descriptor |

13.4 Key Registers

13.4.1 Global Control (CTRL) - 0x000

| Bits | Field | Access | Description |
|--------|------------|--------|----------------------------|
| [0] | ENABLE | RW | Global enable |
| [1] | SOFT_RESET | RW | Soft reset (self-clearing) |
| [31:2] | Reserved | - | Reserved |

13.4.2 Global Status (STATUS) - 0x004

| Bits | Field | Access | Description |
|---------|-----------|--------|-------------------------|
| [0] | BUSY | RO | Any channel active |
| [7:1] | Reserved | - | Reserved |
| [15:8] | CH_ACTIVE | RO | Per-channel active mask |
| [23:16] | CH_ERROR | RO | Per-channel error mask |
| [31:24] | Reserved | - | Reserved |

13.4.3 Channel Control (CHn_CTRL) - 0x040 + n*0x10

| Bits | Field | Access | Description |
|--------|-----------|--------|--|
| [31:0] | DESC_ADDR | RW | Descriptor address (kick-off on write) |

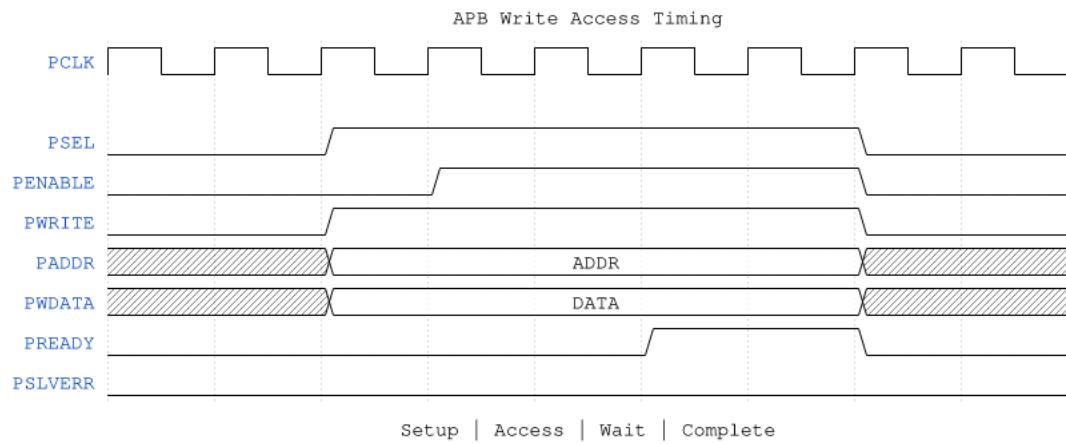
Writing to CHn_CTRL initiates a transfer: 1. Descriptor address stored 2. Channel transitions from IDLE to FETCH 3. PREADY de-asserted until channel accepts kick-off

13.4.4 Channel Status (CHn_STATUS) - 0x044 + n*0x10

| Bits | Field | Access | Description |
|---------|------------|--------|-------------------------|
| [3:0] | STATE | RO | Channel FSM state |
| [4] | IDLE | RO | Channel is idle |
| [5] | BUSY | RO | Channel is processing |
| [6] | ERROR | RO | Error condition |
| [7] | COMPLETE | RO | Last transfer complete |
| [15:8] | DESC_COUNT | RO | Descriptors processed |
| [23:16] | ERR_CODE | RO | Error code if ERROR set |
| [31:24] | Reserved | - | Reserved |

13.5 Access Timing

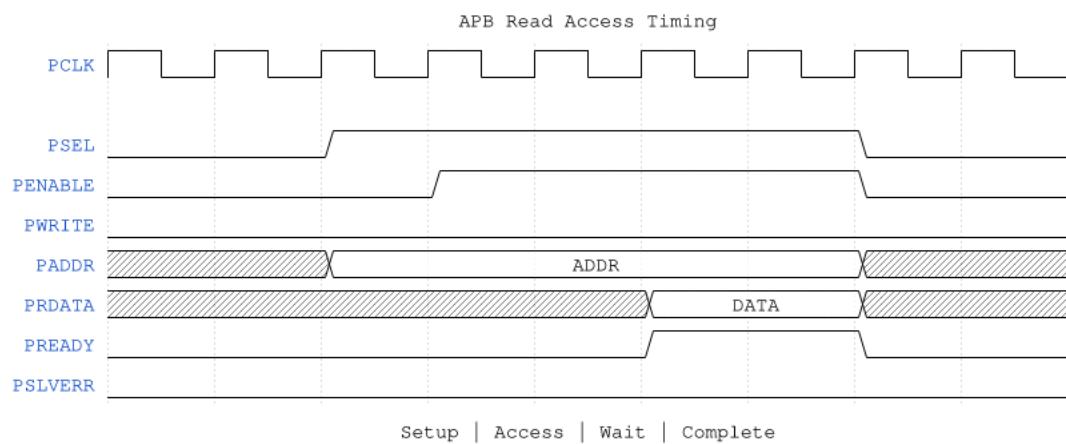
13.5.1 Write Access



APB Write Access Timing

Source: [apb_write_access.json](#)

13.5.2 Read Access



APB Read Access Timing

Source: [apb_read_access.json](#)

13.6 Kick-off Blocking

When software writes to CHn_CTRL:

- If channel is IDLE: Immediate accept, PREADY asserted
 - If channel is BUSY: PREADY held low until channel completes
 - Prevents descriptor pointer corruption during active transfer
-

Last Updated: 2026-01-03

RTL Design Sherpa · Learning Hardware Design Through Practice GitHub · Documentation Index · MIT License

14 Monitoring Interface

14.1 Overview

STREAM provides a 64-bit Monitor Bus (MonBus) interface for debug and performance monitoring. This interface streams event packets capturing key operational events.

14.2 Signal Summary

| Signal | Direction | Width | Description |
|----------------------|-----------|-------|------------------|
| monbus_pkt_va lid | Output | 1 | Packet valid |
| monbus_pkt_re ady | Input | 1 | Downstream ready |
| monbus_pkt_da ta | Output | 64 | Packet data |

14.3 Packet Format

Each 64-bit MonBus packet follows this format:

| Bits | Field | Description |
|-------------|--------------|--------------------------------|
| [63:56] | SUBSYSTEM_ID | STREAM subsystem identifier |
| [55:48] | EVENT_TYPE | Event category |
| [47:40] | EVENT_CODE | Specific event within category |
| [39:32] | CHANNEL_ID | Associated channel (0-7) |
| [31:0] | PAYLOAD | Event-specific data |

14.4 Event Types

14.4.1 Transfer Events (EVENT_TYPE = 0x01)

| EVENT_CODE | Event | Payload |
|-------------------|--------------|-----------------------------|
| 0x01 | KICKOFF | Descriptor address [31:0] |
| 0x02 | DESC_FETCH | Descriptor address [31:0] |
| 0x03 | DESC_VALID | Descriptor fields summary |
| 0x04 | READ_START | Source address [31:0] |
| 0x05 | READ_DONE | Beat count |
| 0x06 | WRITE_START | Destination address [31:0] |
| 0x07 | WRITE_DONE | Beat count |
| 0x08 | CHAIN_NEXT | Next descriptor address |
| 0x09 | COMPLETE | Total descriptors processed |

14.4.2 Error Events (EVENT_TYPE = 0x02)

| EVENT_CODE | Event | Payload |
|------------|---------------|--------------------------|
| 0x01 | AXI_READ_ERR | AXI response code |
| 0x02 | AXI_WRITE_ERR | AXI response code |
| 0x03 | DESC_INVALID | Validation failure code |
| 0x04 | ADDR_RANGE | Offending address [31:0] |
| 0x05 | TIMEOUT | Timeout counter value |

14.4.3 Performance Events (EVENT_TYPE = 0x03)

| EVENT_CODE | Event | Payload |
|------------|---------------|------------------------|
| 0x01 | ARB_WAIT | Wait cycles |
| 0x02 | READ_LATENCY | First beat latency |
| 0x03 | WRITE_LATENCY | First response latency |
| 0x04 | THROUGHPUT | Bytes transferred |

14.5 Event Generation

14.5.1 Configurable Events

Events can be enabled/disabled via APB registers:

| Register | Bit | Event Category |
|------------|-------|--------------------|
| MONBUS_CFG | [0] | Transfer events |
| MONBUS_CFG | [1] | Error events |
| MONBUS_CFG | [2] | Performance events |
| MONBUS_CFG | [7:4] | Channel mask |

14.5.2 Event Priority

When multiple events occur simultaneously:

1. Error events (highest priority)

2. Transfer completion events
 3. Transfer start events
 4. Performance events (lowest priority)
-

14.6 Backpressure Handling

14.6.1 Flow Control

- Events generated when `monbus_pkt_ready` is asserted
- If `monbus_pkt_ready` is low, events are queued in internal FIFO
- FIFO depth: 64 entries (configurable)
- Overflow behavior: Oldest events discarded, overflow counter incremented

14.6.2 Recommendations

- Connect MonBus to FIFO with sufficient depth
 - Monitor overflow counter in status register
 - Size downstream FIFO based on expected event rate
-

14.7 Typical MonBus Downstream

STREAM MonBus → MonBus FIFO → MonBus Arbiter → System Trace
(64 deep) (multi-source) Interface

14.8 Example Event Sequence

A single-descriptor transfer generates these events:

| Time | Event | - |
|------|-------------|---|
| T0 | KICKOFF | - Software kicks off channel |
| T1 | DESC_FETCH | - Descriptor fetch initiated |
| T10 | DESC_VALID | - Descriptor parsed successfully |
| T12 | ARB_WAIT | - Waiting for data path grant (if applicable) |
| T15 | READ_START | - AXI read burst initiated |
| T75 | READ_DONE | - All source data read |
| T80 | WRITE_START | - AXI write burst initiated |
| T140 | WRITE_DONE | - All data written, responses received |
| T141 | COMPLETE | - Channel complete, back to IDLE |

Last Updated: 2026-01-03

RTL Design Sherpa · Learning Hardware Design Through Practice GitHub · Documentation Index · MIT License

15 Throughput Targets

15.1 Theoretical Maximum

15.1.1 Per-Interface Bandwidth

| Interface | Data Width | Frequency | Bandwidth |
|------------------|------------|-----------|-----------|
| Descriptor Fetch | 256 bits | 200 MHz | 6.4 GB/s |
| Data Read | 512 bits | 200 MHz | 12.8 GB/s |
| Data Write | 512 bits | 200 MHz | 12.8 GB/s |

15.1.2 Aggregate Bandwidth

- **Bidirectional:** 25.6 GB/s (simultaneous read + write)
 - **Descriptor Overhead:** ~0.1% for large transfers
-

15.2 Practical Throughput

15.2.1 Single-Channel Performance

| Scenario | Expected Efficiency | Effective Bandwidth |
|------------------------------|---------------------|---------------------|
| Large sequential (>1MB) | >95% | >12.0 GB/s |
| Medium sequential (64KB-1MB) | 85-95% | 10.8-12.0 GB/s |
| Small sequential (4KB-64KB) | 70-85% | 9.0-10.8 GB/s |
| Very small (<4KB) | 40-70% | 5.0-9.0 GB/s |

15.2.2 Multi-Channel Performance

| Active Channels | Per-Channel BW | Aggregate BW |
|-----------------|----------------|--------------|
| 1 | 12.8 GB/s | 12.8 GB/s |
| 2 | 6.4 GB/s | 12.8 GB/s |
| 4 | 3.2 GB/s | 12.8 GB/s |
| 8 | 1.6 GB/s | 12.8 GB/s |

Note: Aggregate limited by shared AXI masters. Individual channel throughput scales inversely with active channel count.

15.3 Throughput Limiting Factors

15.3.1 Memory System Factors

| Factor | Impact | Mitigation |
|-------------------------|--|-----------------------------------|
| Memory latency | Reduces efficiency for small transfers | Use deeper outstanding |
| Memory bandwidth | Hard limit | Match STREAM to memory capability |
| Interconnect contention | Variable impact | Priority configuration |

15.3.2 STREAM Internal Factors

| Factor | Impact | Mitigation |
|---------------------------|------------------------|--------------------------|
| Descriptor fetch overhead | Fixed per descriptor | Use longer transfers |
| SRAM depth | Limits outstanding | Configure adequate depth |
| Arbitration overhead | Multi-channel penalty | Reduce active channels |
| Burst splitting | 4KB boundary crossings | Align transfers |

15.4 Performance Targets

15.4.1 Design Targets

| Metric | Target | Condition |
|----------------------------------|------------|---------------------------------|
| Peak throughput | 12.8 GB/s | Single direction, ideal |
| Sustained throughput | >11.0 GB/s | Large transfers, single channel |
| Multi-channel aggregate | >10.0 GB/s | 8 channels active |
| Small transfer efficiency | >50% | 4KB transfers |

15.4.2 Verification Criteria

| Scenario | Pass Criteria |
|-----------------------|-----------------|
| 1MB sequential read | >95% efficiency |
| 1MB sequential write | >95% efficiency |
| 64KB scatter (4x16KB) | >80% efficiency |
| 4KB × 100 descriptors | >60% efficiency |

15.5 Descriptor Chain Efficiency

15.5.1 Descriptor Overhead

| Transfer Size | Descriptor Fetches | Overhead |
|--------------------|--------------------|----------|
| 64 KB | 1 | 0.05% |
| 256 KB | 1 | 0.01% |
| 1 MB | 1 | <0.01% |
| 4 KB (100 chained) | 100 | 0.5% |

15.5.2 Chaining Efficiency

For chained descriptors, next descriptor fetch can overlap with current transfer:

```
Desc 0: [Fetch][-----Transfer 0-----]
Desc 1:           [Fetch][-----Transfer 1-----]
Desc 2:           [Fetch][-----Transfer 2-----]
```

This pipelining minimizes chaining overhead for continuous transfers.

16 Latency Characteristics

16.1 Latency Components

16.1.1 End-to-End Latency Breakdown

For a single-descriptor transfer (kick-off to completion):

| Phase | Typical Cycles | Notes |
|----------------------|----------------|----------------------------|
| APB write processing | 2-3 | APB transaction |
| Descriptor fetch | 10-50 | Memory latency dependent |
| Descriptor parse | 2-4 | Fixed pipeline |
| Arbitration wait | 0-100+ | Depends on contention |
| Read phase setup | 2-3 | AXI AR channel |
| Read data transfer | N + 10-50 | N beats + memory latency |
| Write phase setup | 2-3 | AXI AW channel |
| Write data transfer | N + 5-20 | N beats + response latency |
| Completion handling | 2-4 | Status update, IRQ |

16.1.2 Total Latency Formula

$$\text{Latency} = T_{\text{apb}} + T_{\text{desc_fetch}} + T_{\text{parse}} + T_{\text{arb}} + T_{\text{read}} + T_{\text{write}} + T_{\text{complete}}$$

Where:

```

T_app          = 3 cycles (fixed)
T_desc_fetch  = Memory_latency + 10 cycles
T_parse        = 3 cycles (fixed)
T_arb          = 0 to N (contention dependent)
T_read         = Memory_latency + (Transfer_length / Beats_per_burst)
* Burst_overhead + Transfer_length
T_write        = Memory_latency + Transfer_length + Response_latency
T_complete     = 3 cycles (fixed)

```

16.2 First-Byte Latency

Time from kick-off to first data beat arriving at destination:

| Component | Cycles | Cumulative |
|--------------------|------------|------------|
| APB processing | 3 | 3 |
| Descriptor fetch | 30 | 33 |
| Parse + setup | 5 | 38 |
| Read AR accepted | 10 | 48 |
| First read data | 20 | 68 |
| Write AW + first W | 10 | 78 |
| Total | ~80 | - |

At 200 MHz: ~400 ns first-byte latency (typical)

16.3 Latency by Transfer Size

16.3.1 Small Transfers

| Transfer Size | Total Latency | Efficiency |
|---------------------|---------------|------------|
| 64 bytes (1 beat) | ~100 cycles | 1% |
| 256 bytes (4 beats) | ~110 cycles | 4% |
| 1 KB (16 beats) | ~130 cycles | 12% |
| 4 KB (64 beats) | ~180 cycles | 35% |

16.3.2 Large Transfers

| Transfer Size | Total Latency | Efficiency |
|---------------------|---------------|------------|
| 16 KB (256 beats) | ~400 cycles | 64% |
| 64 KB (1024 beats) | ~1200 cycles | 85% |
| 256 KB (4096 beats) | ~4400 cycles | 93% |
| 1 MB (16384 beats) | ~17000 cycles | 96% |

16.4 Latency Variability

16.4.1 Sources of Variability

| Source | Impact | Range |
|---------------------|--------|----------------|
| Memory system load | High | 10-100+ cycles |
| AXI interconnect | Medium | 5-50 cycles |
| Channel arbitration | Medium | 0-100+ cycles |
| Descriptor chaining | Low | 2-5 cycles |

16.4.2 Worst-Case Latency

Conditions for worst-case latency:

1. All 8 channels active (maximum arbitration delay)
2. Memory system fully loaded (maximum memory latency)
3. Small transfer size (maximum overhead ratio)
4. Low priority descriptor (maximum wait time)

Under these conditions, first-byte latency could exceed 1000 cycles.

16.5 Latency Optimization

16.5.1 Software Recommendations

| Recommendation | Benefit |
|----------------------------------|--------------------|
| Use larger transfers | Better efficiency |
| Minimize active channels | Reduce arbitration |
| Use higher priority for latency- | Earlier service |

| Recommendation | Benefit |
|--------------------------------|--------------|
| sensitive | |
| Pre-stage descriptors in cache | Faster fetch |

16.5.2 Hardware Configuration

| Setting | Effect |
|----------------------------|--------------------------|
| Increase outstanding limit | Hide memory latency |
| Enable descriptor prefetch | Reduce chaining overhead |
| Configure larger SRAM | Support longer bursts |

16.6 Interrupt Latency

Time from transfer completion to interrupt assertion:

| Phase | Cycles |
|------------------------|-----------------|
| Last write response | 0 |
| Completion FSM | 2 |
| Status register update | 1 |
| Interrupt assertion | 1 |
| Total | 4 cycles |

At 200 MHz: 20 ns from completion to interrupt.

Last Updated: 2026-01-03

RTL Design Sherpa · Learning Hardware Design Through Practice GitHub · Documentation Index · MIT License

17 Resource Estimates

17.1 FPGA Resource Summary

Estimates for typical configuration (8 channels, 512-bit data width, 4096-entry SRAM):

| Resource | Estimate | Notes |
|-----------|-----------------|--------------------|
| LUTs | 15,000 - 25,000 | Logic and control |
| Registers | 20,000 - 30,000 | Pipeline and state |
| Block RAM | 40 - 60 | SRAM buffer |
| DSP | 0 | No DSP usage |

17.2 Resource Breakdown by Block

17.2.1 APB Configuration Slave

| Resource | Estimate |
|-----------|-----------|
| LUTs | 500 - 800 |
| Registers | 400 - 600 |
| BRAM | 0 |

17.2.2 Descriptor Engine

| Resource | Estimate |
|-----------|---------------------|
| LUTs | 1,500 - 2,500 |
| Registers | 1,000 - 1,500 |
| BRAM | 1 (descriptor FIFO) |

17.2.3 Scheduler (8 channels)

| Resource | Estimate |
|-----------|---------------|
| LUTs | 3,000 - 5,000 |
| Registers | 4,000 - 6,000 |
| BRAM | 0 |

17.2.4 Channel Arbiter

| Resource | Estimate |
|-----------|-------------|
| LUTs | 500 - 1,000 |
| Registers | 200 - 400 |
| BRAM | 0 |

17.2.5 AXI Read Engine

| Resource | Estimate |
|-----------|---------------|
| LUTs | 2,000 - 3,500 |
| Registers | 3,000 - 4,500 |
| BRAM | 0 |

17.2.6 SRAM Buffer

| Resource | Estimate |
|-----------|---------------------------------|
| LUTs | 100 - 200 |
| Registers | 100 - 200 |
| BRAM | 32 - 64 (depth/width dependent) |

17.2.7 AXI Write Engine

| Resource | Estimate |
|-----------|---------------|
| LUTs | 2,000 - 3,500 |
| Registers | 3,000 - 4,500 |
| BRAM | 0 |

17.2.8 MonBus Reporter

| Resource | Estimate |
|-----------|----------------|
| LUTs | 800 - 1,200 |
| Registers | 600 - 900 |
| BRAM | 1 (event FIFO) |

17.3 SRAM Sizing

SRAM buffer size depends on configuration:

| Configuration | BRAM (18Kb) | BRAM (36Kb) |
|---------------|-------------|-------------|
| 512b x 1024 | 16 | 8 |
| 512b x 2048 | 32 | 16 |
| 512b x 4096 | 64 | 32 |
| 256b x 4096 | 32 | 16 |
| 128b x 4096 | 16 | 8 |

17.4 Scaling with Parameters

17.4.1 Channel Count Impact

| Channels | LUT Delta | Register Delta |
|----------|-----------|----------------|
| 1 | Base | Base |
| 2 | +500 | +800 |
| 4 | +1,500 | +2,400 |
| 8 | +3,500 | +5,600 |

17.4.2 Data Width Impact

| Data Width | LUT Delta | Register Delta | BRAM Delta |
|------------|-----------|----------------|------------|
| 128 bits | -30% | -40% | -75% |
| 256 bits | -15% | -20% | -50% |
| 512 bits | Base | Base | Base |

17.5 ASIC Estimates

For 28nm technology node (typical):

| Metric | Estimate |
|------------|--------------------------------------|
| Gate Count | 200K - 350K gates |
| Area | 0.2 - 0.4 mm ² |
| SRAM Area | 0.1 - 0.2 mm ² (4096x512) |
| Total Area | 0.3 - 0.6 mm ² |

17.5.1 Power Estimates

| Condition | Power |
|-----------------------|--------------|
| Idle (clock gated) | <1 mW |
| Idle (clock running) | 10 - 20 mW |
| Single channel active | 50 - 100 mW |
| All channels active | 150 - 300 mW |

17.6 Timing Estimates

17.6.1 Target Frequencies

| Technology | Target Fmax |
|--------------------|-------------|
| Xilinx Ultrascale+ | 250 MHz |
| Intel Agilex | 250 MHz |
| 28nm ASIC | 400 MHz |
| 16nm ASIC | 600 MHz |

17.6.2 Critical Paths

| Path | Concern Level | Mitigation |
|------------------|---------------|------------------------|
| AXI data path | Low | Fully pipelined |
| SRAM read/write | Low | Registered outputs |
| Arbiter decision | Medium | Priority encoder depth |
| Descriptor parse | Low | Simple combinational |

Last Updated: 2026-01-03

RTL Design Sherpa · Learning Hardware Design Through Practice GitHub · Documentation Index · MIT License

18 System Requirements

18.1 Interface Requirements

18.1.1 APB Interface

| Requirement | Specification |
|---------------|-----------------|
| Protocol | APB3 or APB4 |
| Data Width | 32 bits |
| Address Width | 12 bits minimum |

| Requirement | Specification |
|---------------|---------------------|
| Timing | Synchronous to aclk |

18.1.2 AXI Interfaces

| Requirement | Descriptor | Read Data | Write Data |
|----------------------|------------|-----------------|-----------------|
| Protocol | AXI4 | AXI4 | AXI4 |
| Data Width | 256 bits | Parameterizable | Parameterizable |
| Address Width | 64 bits | 64 bits | 64 bits |
| ID Width | 4 bits min | 4 bits min | 4 bits min |
| Outstanding | 2 | Configurable | Configurable |

18.1.3 MonBus Interface

| Requirement | Specification |
|---------------------|--|
| Data Width | 64 bits |
| Protocol | Valid/Ready streaming |
| Backpressure | Required (must handle ready deassertion) |

18.2 Memory Requirements

18.2.1 Descriptor Memory

| Requirement | Specification |
|------------------|---------------------------------------|
| Alignment | 32-byte aligned |
| Access | Read access via descriptor AXI master |
| Coherency | Software responsibility |
| Latency | <100 cycles recommended |

18.2.2 Data Memory

| Requirement | Specification |
|------------------|--------------------------------------|
| Alignment | Per data width (64-byte for 512-bit) |

| Requirement | Specification |
|------------------|--|
| Access | Read via read master, Write via write master |
| Bandwidth | Sufficient for target throughput |

18.3 System Bandwidth Allocation

18.3.1 AXI Interconnect Requirements

To achieve maximum STREAM throughput:

| AXI Master | Required Bandwidth | Priority |
|------------------|--------------------|----------|
| Descriptor Fetch | 0.1 GB/s | Medium |
| Data Read | 12.8 GB/s | High |
| Data Write | 12.8 GB/s | High |

18.3.2 Interconnect Recommendations

- Configure STREAM AXI masters with high priority
 - Ensure sufficient outstanding transaction support
 - Consider dedicated paths for high-bandwidth DMA
-

18.4 Interrupt Requirements

18.4.1 Interrupt Controller

| Requirement | Specification |
|------------------------|---|
| Interrupt Lines | 8 (per-channel) or 1 (combined) |
| Type | Level-sensitive, active-high |
| Latency | <10 cycles from assertion to controller |

18.4.2 Interrupt Handling

- Software must read `IRQ_STATUS` to identify source
 - Write-1-to-clear to acknowledge interrupts
 - Channel remains blocked until error cleared (if applicable)
-

18.5 Address Space Requirements

18.5.1 STREAM Register Space

| Requirement | Specification |
|---------------------|----------------|
| Base Address | System-defined |
| Size | 4 KB (0x1000) |
| Alignment | 4 KB aligned |

18.5.2 Memory Map Constraints

- Descriptor addresses must be accessible via descriptor AXI master
 - Source addresses must be accessible via read AXI master
 - Destination addresses must be accessible via write AXI master
 - All addresses must respect configured address range limits
-

18.6 Software Requirements

18.6.1 Driver Requirements

| Requirement | Description |
|-------------------------------|--|
| Descriptor Preparation | Build descriptors in coherent memory |
| Cache Management | Flush descriptor cache before kick-off |
| Interrupt Handling | Service interrupts promptly |
| Error Recovery | Clear error status, optionally retry |

18.6.2 Programming Sequence

1. Verify STREAM is enabled (CTRL.ENABLE = 1)
 2. Check target channel is idle (CHn_STATUS.IDLE = 1)
 3. Prepare descriptor chain in memory
 4. Flush cache (if applicable)
 5. Write first descriptor address to CHn_CTRL
 6. Wait for interrupt or poll CHn_STATUS
-

19 Clocking and Reset

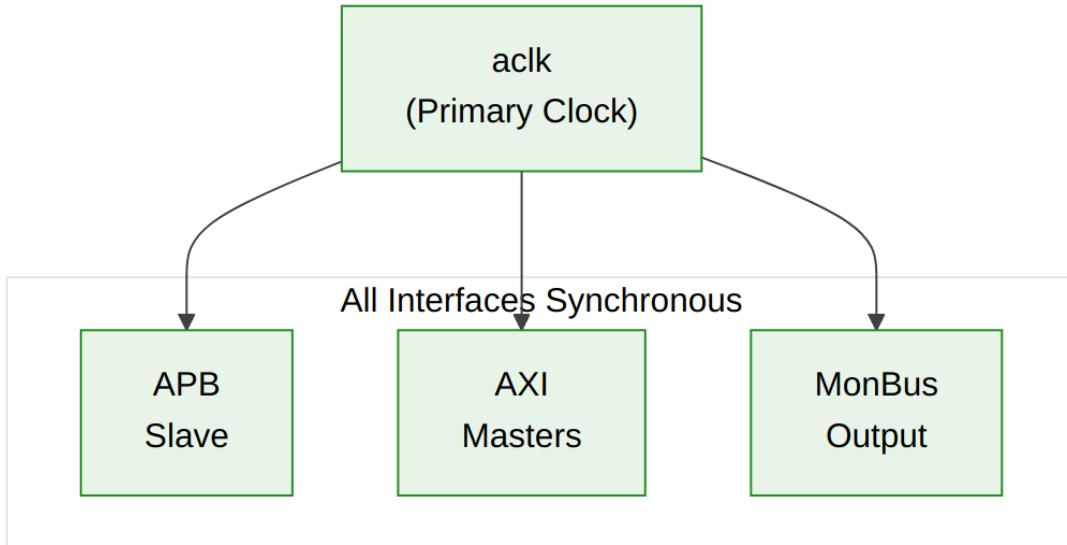
19.1 Clock Requirements

19.1.1 Primary Clock (aclk)

| Parameter | Requirement |
|-------------------|--|
| Frequency | 100 - 400 MHz (target dependent) |
| Duty Cycle | 40% - 60% |
| Jitter | <100 ps peak-to-peak |
| Domain | Single clock domain for all interfaces |

19.1.2 Clock Relationship

All STREAM interfaces are synchronous to aclk:



Clock Distribution

Source: [06_clock_distribution.mmd](#)

[INTEG-1] External interfaces must be synchronous to `aclk`. Clock domain crossing is the system integrator's responsibility.

19.2 Reset Requirements

19.2.1 Reset Signal (`aresetn`)

| Parameter | Requirement |
|----------------------|---|
| Polarity | Active-low |
| Type | Asynchronous assert, synchronous deassert |
| Minimum Pulse | 2 clock cycles |
| Glitch Filter | Not required (external synchronizer assumed) |

19.2.2 Reset Sequence

Reset is asserted asynchronously and de-asserted synchronously:

1. `aresetn` goes low (asynchronous assertion)
2. Reset remains active for minimum 2 clock cycles

3. aresetn goes high on rising edge of aclk (synchronous release)
4. All blocks begin normal operation

19.2.3 Reset Behavior

| Block | Reset State |
|-------------------|-----------------------------|
| APB Config | All registers to default |
| Scheduler | All channels IDLE |
| Descriptor Engine | FSM to IDLE, FIFOs empty |
| AXI Engines | No outstanding transactions |
| SRAM | Contents undefined |
| MonBus | FIFO empty, no output |

19.3 Power Management

19.3.1 Clock Gating Support

STREAM supports external clock gating when idle:

| Signal | Purpose |
|-------------------|------------------------------|
| all_channels_idle | Indicates safe to gate clock |

Clock Gating Sequence: 1. Check STATUS.BUSY = 0 2. Disable clock via external gate 3. Re-enable clock before any APB access

19.3.2 Low Power Recommendations

| Technique | Implementation |
|-----------------|----------------------------------|
| Clock gating | Gate aclk when all_channels_idle |
| Power gating | Not recommended (state loss) |
| Voltage scaling | Supported if timing met |

19.4 Timing Constraints

19.4.1 Clock Period Constraints

```
# Example Vivado XDC constraints
create_clock -period 5.000 -name aclk [get_ports aclk]
```

```

# APB interface (same clock)
set_input_delay -clock aclk 1.0 [get_ports s_apb_*]
set_output_delay -clock aclk 1.0 [get_ports s_apb_pready]
set_output_delay -clock aclk 1.0 [get_ports s_apb_prdata*]
set_output_delay -clock aclk 1.0 [get_ports s_apb_pslverr]

# AXI interfaces (same clock)
set_input_delay -clock aclk 1.0 [get_ports m_axi_*ready]
set_input_delay -clock aclk 1.0 [get_ports m_axi_*rdata*]
set_input_delay -clock aclk 1.0 [get_ports m_axi_*rresp*]
set_input_delay -clock aclk 1.0 [get_ports m_axi_*bresp*]
set_output_delay -clock aclk 1.0 [get_ports m_axi_*valid]
set_output_delay -clock aclk 1.0 [get_ports m_axi_*addr*]
set_output_delay -clock aclk 1.0 [get_ports m_axi_*data*]

```

19.4.2 False Path Constraints

```

# Reset is asynchronous
set_false_path -from [get_ports aresetn]

```

```

# Configuration registers to datapath (multi-cycle)
set_multicycle_path 2 -setup -from [get_cells *cfg_*]
set_multicycle_path 1 -hold -from [get_cells *cfg_*]

```

19.5 CDC Considerations

19.5.1 Single Clock Domain

STREAM operates in a single clock domain. All clock domain crossing must be handled externally:

| Interface | CDC Responsibility |
|-------------|---|
| APB Slave | External if APB on different clock |
| AXI Masters | External if AXI interconnect on different clock |
| MonBus | External if consumer on different clock |

19.5.2 CDC Recommendations

For systems requiring CDC:

1. Use async FIFO for MonBus output
 2. Use AXI clock converter IP for AXI interfaces
 3. Use APB bridge with clock crossing for APB
-

Last Updated: 2026-01-03

RTL Design Sherpa · Learning Hardware Design Through Practice GitHub · Documentation Index · MIT License

20 Verification Strategy

20.1 Verification Approach

20.1.1 Multi-Level Verification

STREAM verification follows a hierarchical approach:

| Level | Focus | Coverage Target |
|--------------------|----------------------------|------------------|
| Block | Individual module function | 100% line/toggle |
| Integration | Block-to-block interfaces | 95% functional |
| System | End-to-end operation | 90% use cases |

20.2 Block-Level Verification

20.2.1 Test Categories

| Block | Key Test Scenarios |
|--------------------------|--|
| APB Config | Register access, kick-off timing, error response |
| Descriptor Engine | Fetch, parse, chain, error handling |

| Block | Key Test Scenarios |
|-------------------------|--|
| Scheduler | State transitions, multi-channel coordination |
| Channel Arbiter | Priority, fairness, starvation prevention |
| AXI Read Engine | Burst generation, error handling, backpressure |
| AXI Write Engine | Burst generation, response handling |
| SRAM Buffer | Read/write concurrent, full/empty |
| MonBus Reporter | Event generation, backpressure, overflow |

20.2.2 Coverage Metrics

| Metric | Target |
|-------------------------|----------------------|
| Line coverage | 100% |
| Toggle coverage | 95% |
| FSM state coverage | 100% |
| FSM transition coverage | 100% |
| Parameter coverage | All supported values |

20.3 Integration-Level Verification

20.3.1 Interface Verification

| Interface | Test Focus |
|--------------------------------|---------------------------------------|
| APB to Descriptor Engine | Kick-off handshake, blocking behavior |
| Descriptor Engine to Scheduler | Descriptor passing, backpressure |
| Scheduler to AXI Engines | Grant/request, data flow |
| AXI Engines to SRAM | Concurrent access, pointer management |

| Interface | Test Focus |
|----------------------|-------------------------------|
| All blocks to MonBus | Event generation, arbitration |

20.3.2 Multi-Channel Scenarios

| Scenario | Description |
|-----------------|---------------------------------------|
| Sequential | Channels operate one at a time |
| Concurrent | All 8 channels active simultaneously |
| Priority | High priority preempts low priority |
| Fairness | Equal priority channels share fairly |
| Error isolation | Error in one channel, others continue |

20.4 System-Level Verification

20.4.1 End-to-End Test Cases

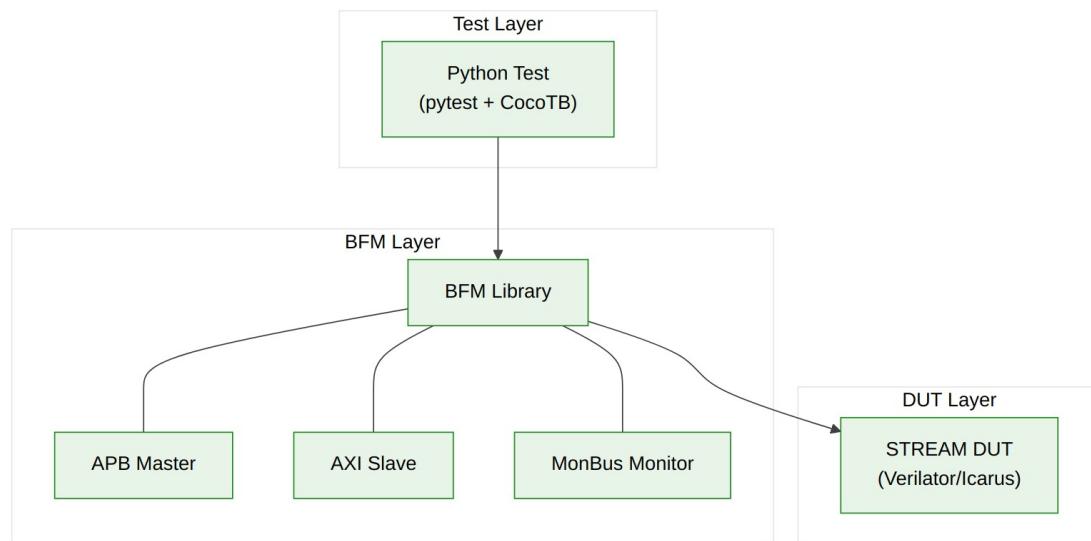
| Test Case | Description |
|---------------------------|--|
| Basic Transfer | Single descriptor, single channel |
| Chained Transfer | Multiple descriptors, single channel |
| Multi-Channel | Multiple channels, independent transfers |
| Scatter | Single source to multiple destinations |
| Gather | Multiple sources to single destination |
| Maximum Throughput | Sustained bandwidth test |
| Error Injection | AXI errors, invalid descriptors |
| Stress | Random traffic, long duration |

20.4.2 Performance Verification

| Metric | Verification Method |
|------------|---|
| Throughput | Measure bytes/cycle over large transfer |
| Latency | Measure kick-off to completion cycles |
| Efficiency | Compare actual vs theoretical maximum |

20.5 Testbench Architecture

20.5.1 CocoTB-Based Verification



Testbench Architecture

Source: [07_testbench_stack.mmd](#)

20.5.2 Key Testbench Components

| Component | Purpose |
|-----------------------|-------------------------------|
| APB Master BFM | Drive configuration interface |
| AXI Slave BFM | Model memory system |
| MonBus Monitor | Capture and verify events |
| Scoreboard | Verify data integrity |

| Component | Purpose |
|---------------------------|---------------------------|
| Coverage Collector | Track functional coverage |

20.6 Verification Environment

20.6.1 Simulator Support

| Simulator | Status |
|--------------------------|-----------------------|
| Verilator | Primary (open-source) |
| Icarus Verilog | Secondary |
| Commercial (VCS, Questa) | Supported |

20.6.2 Test Execution

Run all STREAM tests

```
pytest projects/components/stream/dv/tests/ -v
```

Run specific test category

```
pytest projects/components/stream/dv/tests/fub_tests/ -v
```

```
pytest projects/components/stream/dv/tests/integration_tests/ -v
```

Generate coverage report

```
pytest projects/components/stream/dv/tests/ --cov-report=html
```

20.7 Formal Verification

20.7.1 Properties for Formal

| Category | Properties |
|-------------------|---|
| Protocol | AXI handshake rules, APB timing |
| Liveness | No deadlock in arbitration |
| Safety | No data corruption, proper error handling |
| Functional | Descriptor chaining correctness |

20.7.2 Formal Tool Usage

Recommended for: - AXI protocol compliance - Arbiter deadlock freedom - FSM reachability analysis

20.8 Sign-Off Criteria

20.8.1 Block-Level Sign-Off

- 100% line coverage
- 95% toggle coverage
- All FSM states exercised
- All FSM transitions exercised
- No simulation errors
- No lint warnings

20.8.2 Integration Sign-Off

- All interface scenarios passed
- Multi-channel concurrency verified
- Error handling verified
- Performance targets met

20.8.3 System Sign-Off

- All end-to-end test cases passed
 - Throughput targets achieved
 - Latency targets achieved
 - 72-hour stress test passed
-

Last Updated: 2026-01-03

RTL Design Sherpa · Learning Hardware Design Through Practice GitHub · Documentation Index · MIT License

21 STREAM Specification Index

Version: 0.90 **Date:** 2025-11-22 **Purpose:** Complete technical specification for STREAM subsystem

21.1 Document Organization

Note: All chapters linked below for automated document generation.

21.1.1 Front Matter

- [Document Information](#)

21.1.2 Chapter 1: Overview

- [Architecture](#)
- [Top-Level Port List](#)
- [Clocks and Reset](#)

21.1.3 Chapter 2: Functional Blocks

Macro/Integration Level: - Stream Core - Scheduler Group Array - Scheduler Group

Control Path: - Scheduler - Descriptor Engine

Read Datapath: - AXI Read Engine - Stream Alloc Control

SRAM Buffering: - SRAM Controller - SRAM Controller Unit - Stream Latency Bridge

Write Datapath: - Stream Drain Control - AXI Write Engine

Configuration & Monitoring: - APB to Descriptor Router - APB Config - Performance Profiler - MonBus AXIL Group

21.1.4 Chapter 3: Interfaces

- [Interface Overview](#)
- [AXI4 Interface Specification](#)
- [AXIL4 Interface Specification](#)
- [APB Programming Interface](#)
- [MonBus Interface Specification](#)

21.1.5 Chapter 4: Registers

- [Register Map](#)

21.1.6 Chapter 5: Programming

- [Programming Guide](#)

21.1.7 Chapter 6: Configuration Reference

- [Complete Configuration Guide](#)
-

21.2 Quick Reference

21.2.1 Functional Unit Blocks (FUB)

| Module | File | Purpose | Lines | Status |
|-------------------|--|----------------------------------|-------|---|
| descriptor_engine | stream_fub/ descriptor_engine. .sv | Descriptor fetch/parse (256-bit) | ~300 | [Done] |
| scheduler | stream_fub/ scheduler .sv | Transfer coordinator | ~400 | [Done] Created (corrected) |
| axi_read_engine | stream_fub/ axi_read_engine.sv | AXI read master | ~350 | [Done] Created (no FSM - streaming pipeline) |
| axi_write_engine | stream_fub/ axi_write_engine.sv | AXI write master | ~400 | [Done] Created (no FSM - streaming pipeline) |
| sram_controller | stream_fub/ sram_controller.sv | Per-channel buffer management | ~350 | [Done] Created (no FSM - pointer arithmetic + pre-allocation) |
| simple_sram | stream_fub/ simple_sram.sv | Dual-port SRAM primitive | ~150 | [Done] |
| perf_profiler | stream_fub/ perf_prof | Channel performance | ~400 | [Done] Dual-mode |

| Module | File | Purpose | Lines | Status |
|--------|---------|-----------|-------|----------------------------|
| | iler.sv | profiling | | timestamp/elapsed tracking |

21.2.2 Integration Blocks (MAC)

| Module | File | Purpose | Lines | Status |
|-------------------|-------------------------------------|----------------------|-------|-------------------------|
| channel_arbiter | stream_maestro/channel_arbiter.sv | Priority arbitration | ~200 | [Pending] To be created |
| apb_config | regs/stream_registers.rdl + wrapper | Config registers | ~350 | [Future] PeakRDL-based |
| monbus_axil_group | stream_maestro/monbus_axil_group.sv | MonBus + AXIL | ~800 | [Done] |
| stream_top | stream_maestro/stream_top.sv | Top-level | ~500 | [Pending] To be created |

21.3 Performance Modes (AXI Engines)

STREAM AXI engines support three performance modes via compile-time parameters, offering area/performance trade-offs from tutorial simplicity to datacenter throughput.

21.3.1 Holistic Overview: V1/V2/V3 Working Together

Design Philosophy: - **V1 (Low):** Tutorial-focused, simple architecture, minimal area - **V2 (Medium):** Command pipelining hides memory latency, 6.7x throughput improvement - **V3 (High):** Out-of-order completion maximizes memory controller efficiency, 7.0x throughput

Key Insight: The benefit scales with memory latency. For low-latency SRAM (2-3 cycles), V1 achieves 40% throughput. For high-latency DDR4 (70-100 cycles), V1 drops to 14% but V2/V3 maintain 94-98% throughput.

Parameterization Strategy: - `ENABLE_CMD_PIPELINE = 0`: V1 (default, tutorial mode) - `ENABLE_CMD_PIPELINE = 1`: V2 (command pipelined, best area efficiency) - `ENABLE_CMD_PIPELINE = 1, ENABLE_000_DRAIN = 1` (write) or `ENABLE_000_READ = 1` (read): V3 (out-of-order, maximum throughput)

21.3.2 V1 - Low Performance (Single Outstanding Transaction)

Architecture: Single-burst-per-request, blocking on completion - **Area:** ~1,250 LUTs per engine - **Throughput:** 0.14 beats/cycle (DDR4), 0.40 beats/cycle (SRAM) - **Outstanding Txns:** 1 - **Use Case:** Tutorial examples, embedded systems, low-latency SRAM - **Key Feature:** Zero-bubble streaming pipeline (NO FSM!)

Blocking Behavior: - Write: Blocks on B response before accepting next request - Read: Blocks on R last before accepting next request - Simple control: 3 flags (r_ar_inflight, r_ar_valid, completion tracking)

21.3.3 V2 - Medium Performance (Command Pipelined)

Architecture: Command queue decouples command acceptance from data completion - **Area:** ~2,000 LUTs per engine (1.6x increase) - **Throughput:** 0.94 beats/cycle (DDR4), 0.85 beats/cycle (SRAM) - **Outstanding Txns:** 4-8 (command queue depth) - **Use Case:** General-purpose FPGA, DDR3/DDR4, balanced area/performance - **Key Feature:** Hides memory latency via pipelining (6.7x improvement)

Command Pipelining: - Write: AW queue + W drain FSM + B scoreboard (async response tracking) - Read: AR queue + R in-order reception (simpler than write, no B channel) - Per-command SRAM pointers enable multiple bursts from same channel

Best Area Efficiency: 4.19x throughput per unit area (6.7x throughput / 1.6x area)

21.3.4 V3 - High Performance (Out-of-Order Completion)

Architecture: OOO command selection maximizes memory controller efficiency - **Area:** ~3,500-4,000 LUTs per engine (2.8-3.2x increase) - **Throughput:** 0.98 beats/cycle (DDR4), 0.92 beats/cycle (SRAM) - **Outstanding Txns:** 8-16 (larger command queue) - **Use Case:** Datacenter, ASIC, HBM2, high-performance memory controllers - **Key Feature:** Flexible drain order based on SRAM data availability (7.0x improvement)

Out-of-Order Mechanisms: - Write: OOO W drain (select command with SRAM data ready), AXI ID matching for B responses - Read: OOO R reception (match m_axi_rid to queue entry), independent SRAM write per command - Transaction ID structure: {counter, channel_id} preserves channel routing for MonBus

Why OOO is Naturally Supported: 1. Per-channel SRAM partitioning (no cross-channel hazards) 2. Per-command pointer tracking (no pointer collision) 3. Transaction ID matching (channel ID in lower bits for routing)

21.3.5 Performance Comparison Summary

| Configuration | Area | DDR4 Throughput | SRAM Throughput | Area Efficiency | Use Case |
|---------------|------|------------------|------------------|-----------------|------------|
| V1 | 1.0x | 0.14 beats/cycle | 0.40 beats/cycle | 1.00 | Tutorial , |

| Configuration | Area (1.0x) | DDR4 Throughput (1.0x) | SRAM Throughput (1.0x) | Area Efficiency | Use Case |
|---------------|-------------|----------------------------|----------------------------|-----------------|---------------------|
| V2 | 1.6x | 0.94 beats/cycle (6.7x) | 0.85 beats/cycle (2.1x) | 4.19 | embedded |
| V3 | 2.8x | 0.98 beats/cycle (7.0x) | 0.92 beats/cycle (2.3x) | 2.50 | Datacenter, ASIC |

Key Takeaway: V2 offers best area efficiency (4.19x) for typical FPGA use cases. V3 provides marginal throughput improvement (7.0x vs 6.7x) at higher area cost, justified only for high-performance memory controllers that support out-of-order responses.

21.4 Clock and Reset Summary

21.4.1 Clock Domains

| Clock | Frequency | Usage |
|-------|-------------|--|
| aclk | 100-500 MHz | Primary - all STREAM logic, AXI/AXIL interfaces |
| pclk | 50-200 MHz | APB configuration interface (may be async to aclk) |

21.4.2 Reset Signals

| Reset | Polarity | Type | Usage |
|---------|------------|-----------------------------|-----------------------------|
| aresetn | Active-low | Async assert, sync deassert | Primary - all STREAM logic |
| presetn | Active-low | Async assert, sync deassert | APB configuration interface |

See: [Clocks and Reset](#) for complete timing specifications

21.5 Interface Summary

21.5.1 External Interfaces

| Interface | Type | Width | Purpose |
|------------------|--------|-----------------|---------------------------------|
| APB | Slave | 32-bit | Configuration registers |
| AXI (Descriptor) | Master | 256-bit | Descriptor fetch |
| AXI (Read) | Master | 512-bit (param) | Source data read |
| AXI (Write) | Master | 512-bit (param) | Destination data write |
| AXIL (Slave) | Slave | 32-bit | Error/interrupt FIFO access |
| AXIL (Master) | Master | 32-bit | MonBus packet logging to memory |
| IRQ | Output | 1-bit | Error interrupt |

21.5.2 Internal Buses

| Interface | Width | Purpose |
|-----------|--------|---|
| MonBus | 64-bit | Internal monitoring bus (channels -> monbus_axil_group) |

21.6 Area Estimates

21.6.1 By Performance Mode

| Configuration | Total LUTs | SRAM | Use Case |
|------------------|------------|-------|-------------------------------|
| Low (Tutorial) | ~9,500 | 64 KB | Educational, area-constrained |
| Medium (Typical) | ~11,200 | 64 KB | Balanced FPGA implementation |

| Configuration | Total LUTs | SRAM | Use Case |
|-----------------------|------------|-------|---------------------------|
| High (Performance) | ~13,700 | 64 KB | ns |
| | | | High-throughput ASIC/FPGA |

21.6.2 Breakdown (Low Performance)

| Component | Instances | Area/Instance | Total |
|------------------------|-----------|----------------|---------------------------|
| Descriptor Engine | 8 | ~300 LUTs | ~2,400 LUTs |
| Scheduler | 8 | ~400 LUTs | ~3,200 LUTs |
| AXI Read Engine (Low) | 1 | ~250 LUTs | ~250 LUTs |
| AXI Write Engine (Low) | 1 | ~250 LUTs | ~250 LUTs |
| SRAM Controller | 1 | ~1,600 LUTs | ~1,600 LUTs |
| Simple SRAM (internal) | 1-8 | 1024x64B total | 64 KB |
| Channel Arbiter | 3 | ~150 LUTs | ~450 LUTs |
| APB Config | 1 | ~350 LUTs | ~350 LUTs |
| MonBus AXIL Group | 1 | ~1,000 LUTs | ~1,000 LUTs |
| Total | - | - | ~9,500 LUTs + 64KB |

21.7 Related Documentation

- [PRD.md](#) - Product requirements and overview
- [ARCHITECTURAL_NOTES.md](#) - Critical design decisions
- [CLAUDE.md](#) - AI development guide
- [Register Generation](#) - PeakRDL workflow

21.8 Specification Conventions

21.8.1 Signal Naming

- **Clock:** aclk, pclk
- **Reset:** aresetn, presetn (active-low)
- **Valid/Ready:** Standard AXI/custom handshake
- **Registers:** r_ prefix (e.g., r_state, r_counter)
- **Wires:** w_ prefix (e.g., w_next_state, w_grant)

21.8.2 Parameter Naming

- **Uppercase:** NUM_CHANNELS, DATA_WIDTH, ADDR_WIDTH
- **String parameters:** PERFORMANCE (“LOW”, “MEDIUM”, “HIGH”)

21.8.3 State Machine Naming

```
typedef enum logic [3:0] {
    IDLE    = 4'h0,
    ACTIVE  = 4'h1,
    // ...
} state_t;

state_t r_state, w_next_state; // Current and next state
```

Last Updated: 2025-10-17 Maintained By: STREAM Architecture Team

RTL Design Sherpa · Learning Hardware Design Through Practice GitHub · Documentation Index · MIT License

22 Product Requirements Document (PRD)

22.1 STREAM - Scatter-gather Transfer Rapid Engine for AXI Memory

Version: 1.0 **Date:** 2025-10-17 **Status:** Nearly Complete - Final Integration Pending **Owner:** RTL Design Sherpa Project **Parent Document:** /PRD.md

22.2 1. Executive Summary

The **STREAM** (Scatter-gather Transfer Rapid Engine for AXI Memory) is a simplified DMA-style accelerator designed for efficient memory-to-memory data movement with descriptor-based scatter-gather support. STREAM serves as a beginner-friendly tutorial demonstrating descriptor-based DMA engine design patterns while maintaining production-quality RTL practices.

22.2.1 1.1 Quick Stats

- **Modules:** ~8-10 SystemVerilog files (estimated)
- **Channels:** Maximum 8 independent channels
- **Interfaces:** APB (config), AXI4 (descriptor fetch + data read/write), MonBus (monitoring)
- **Architecture:** Simplified from RAPIDS - pure memory-to-memory
- **Tutorial Focus:** Aligned addresses only, straightforward data flow
- **Status:** Nearly complete - config interface and top-level wrapper pending

22.2.2 1.2 Project Goals

- **Primary:** Educational DMA engine demonstrating scatter-gather descriptor chains
 - **Secondary:** Production-quality RTL suitable for FPGA/ASIC implementation
 - **Tertiary:** Foundation for understanding more complex DMA architectures (e.g., RAPIDS)
-

22.3 2. Key Design Principles

22.3.1 2.1 Simplifications from RAPIDS

STREAM is intentionally simplified for tutorial purposes:

| Feature | RAPIDS | STREAM |
|---------------------------|----------------------------|---------------------------------|
| Network Interfaces | Yes (Network master/slave) | ✗ No (pure memory-to-memory) |
| Address Alignment | Complex fixup logic | ✓ Aligned addresses only |
| Credit Management | Exponential encoding | ✓ Simple transaction limits |
| Control Engines | Control read/write | ✗ No (direct APB config) |

| Feature | RAPIDS | STREAM |
|--------------------------|--------------------------|-----------------------------|
| | engines | |
| Descriptor Length | Chunks (4-byte) | ✓ Beats (data width) |
| Program Engine | Complex alignment FSM | ✓ Simplified coordination |

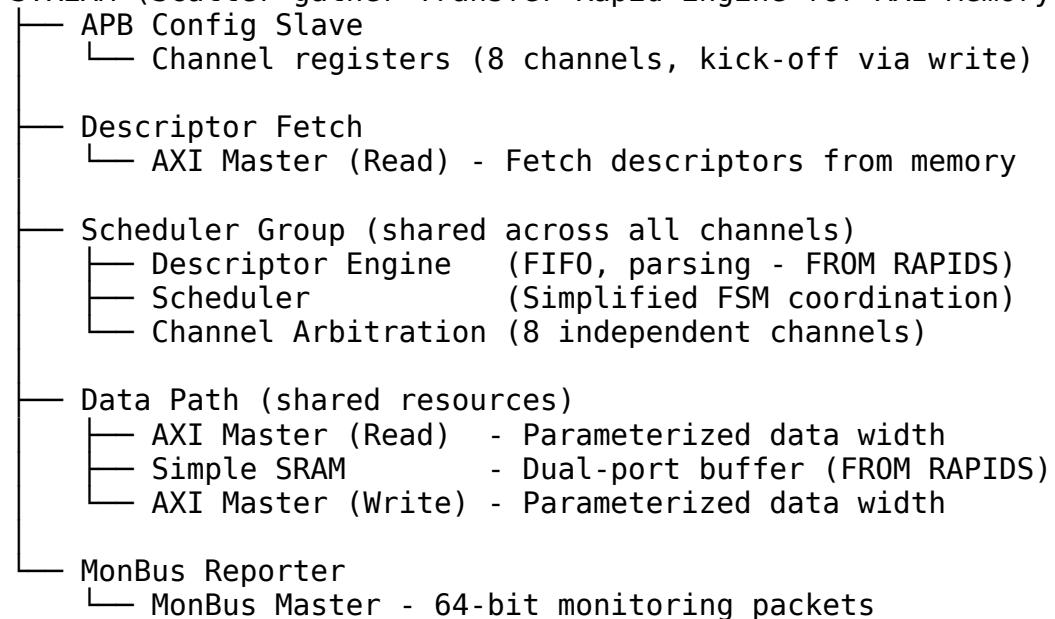
22.3.2 2.2 Tutorial-Friendly Features

- **Aligned Addresses:** Source and destination addresses must be aligned to data width
 - **Length in Beats:** Descriptor length specified in data beats (not bytes/chunks)
 - **Single APB Write:** One APB register write kicks off entire descriptor chain
 - **No Circular Buffers:** Chained descriptors with explicit termination
 - **Parameterized Engines:** Multiple AXI engine versions (compile-time selection)
-

22.4 3. Architecture Overview

22.4.1 3.1 Top-Level Block Diagram

STREAM (Scatter-gather Transfer Rapid Engine for AXI Memory)



22.4.23.2 Data Flow

Descriptor-Based Transfer Sequence:

1. Software writes to APB channel register
↓
2. Descriptor fetch via AXI descriptor master
↓
3. Descriptor Engine parses descriptor fields
↓
4. Scheduler coordinates data transfer:
 - a. AXI Read Engine fetches source data → SRAM buffer
 - b. AXI Write Engine writes SRAM → destination
↓
5. Check for chained descriptor (next_descriptor_ptr != 0)
↓ (if chained)
6. Fetch next descriptor, repeat from step 3
↓ (if last)
7. Generate MonBus completion packet

Channel Independence: - 8 channels operate independently - All channels share: SRAM, AXI data masters, descriptor fetch master - Arbitration required for shared resources

22.5 4. Interfaces

22.5.14.1 External Interfaces

| Interface | Type | Width | Purpose | Notes |
|--------------------------------|--------|-----------------|---------------------------------|---|
| APB Slave | Slave | 32-bit | Configuration, channel kick-off | Write to channel register starts transfer |
| AXI Master (Descriptor) | Master | 256-bit | Fetch descriptors from memory | Dedicated descriptor fetch path |
| AXI Master (Data Read) | Master | Parameterizable | Read source data | Multiple engine versions (compile-time) |

| Interface | Type | Width | Purpose | Notes |
|--------------------------------|--------|-----------------|------------------------|---|
| AXI Master (Data Write) | Master | Parameterizable | Write destination data | Multiple engine versions (compile-time) |
| MonBus Master | Master | 64-bit | Monitor packet output | Standard AMBA format |

22.5.24.2 Descriptor Format

256-bit Descriptor Structure:

| Bits | Field | Description |
|-----------|---------------------|---|
| [63:0] | src_addr | Source address (64-bit, must be aligned to data width) |
| [127:64] | dst_addr | Destination address (64-bit, must be aligned to data width) |
| [159:128] | length | Transfer length in BEATS (not bytes!) |
| [191:160] | next_descriptor_ptr | Address of next descriptor (0 = last in chain) |
| [192] | valid | Descriptor is valid |
| [193] | interrupt | Generate interrupt on completion |
| [194] | last | Last descriptor in chain (explicit flag) |
| [195] | error | Error status (used for reporting) |
| [199:196] | channel_id | Channel ID (0-7) |
| [207:200] | priority | Transfer priority (for arbitration) |
| [255:208] | reserved | Reserved for future use |

Key Descriptor Features: - ✓ **Chained descriptors:** next_descriptor_ptr links to next descriptor - ✗ **No circular buffers:** Explicit termination (last flag or ptr=0) - ✓ **Length in beats:** Simplified for tutorial (no byte/chunk conversion) - ✓ **Aligned addresses:** Tutorial constraint (performance hidden for now)

22.6 5. Key Components

22.6.15.1 Descriptor Engine (APB-Only for STREAM)

Source: Adapted from RAPIDS descriptor_engine.sv

Purpose: - Autonomous descriptor fetch and chaining - APB interface for initial descriptor address - AXI read interface for descriptor memory fetches - Descriptor FIFO storage and distribution

Key Features: - ✓ **Autonomous chaining:** Automatically fetches next descriptor if `next_descriptor_ptr != 0` AND `last == 0` - ✓ **Address validation:** Validates next descriptor addresses against `cfg_addr0/1_base/limit` - ✓ **APB blocking:** APB blocked until `channel_idle == 1` (channel fully idle) - ✓ **Error handling:** AXI errors stop chaining, set `descriptor_error`, `block_descriptor_valid`

Adaptations from RAPIDS: - ✗ **RDA removed:** STREAM is memory-to-memory only (no network interfaces) - ✓ **APB-only:** Single APB write kicks off entire descriptor chain - ✓ **Descriptor Read Address FIFO:** 2-deep FIFO stores addresses for AXI fetch (APB + chaining) - ✓ **Chaining logic:** Descriptor engine autonomously manages `next_descriptor_ptr` chaining

Idle Signal: - `descriptor_engine_idle` asserted when: - FSM in RD_IDLE state - No pending descriptor fetches (address FIFO empty) - No active AXI transactions

22.6.25.2 Scheduler Group (Integration Wrapper)

Purpose: Wraps descriptor engine and scheduler into a single channel processing unit

Architecture:

```
scheduler_group (
    // APB interface (from APB config slave)
    .apb_valid      (apb_valid),
    .apb_ready      (apb_ready),  // Blocked when channel not idle
    .apb_addr       (descriptor_addr),

    // Channel idle signal composition (CRITICAL!)
    .channel_idle   (channel_idle),

    // Descriptor → Scheduler flow
    .descriptor_valid (desc_valid),
    .descriptor_ready (desc_ready),
    .descriptor_packet (desc_packet),

    // Data engine interfaces
    .datard_*        (datard_*),    // Read engine
    .datawr_*        (datawr_*),    // Write engine
```

```

    // Status
    .scheduler_idle      (sched_idle),
    .descriptor_idle     (desc_idle)
);

```

Channel Idle Signal Composition:

```

// Channel is idle ONLY when BOTH sub-blocks are idle
assign channel_idle = scheduler_idle && descriptor_engine_idle;

```

Why Both Signals Matter:

| Signal | Indicates | Used For |
|----------------------------|---|---|
| scheduler_idle | No active data transfers, all descriptors processed | Prevents new APB request during active transfer |
| descriptor_engine _idle | No pending descriptor fetches (FIFO empty) | Prevents new APB request during chaining |
| channel_idle (AND of both) | Channel fully quiescent | Gates APB interface |

APB Blocking Logic:

```

// Descriptor engine blocks APB when channel not idle
assign apb_ready = apb_skid_ready_in &&
                   !r_channel_reset_active &&
                   w_desc_addr_fifo_empty &&           // No pending fetches
                   channel_idle;                         // Scheduler +
                                         descriptor_idle

```

Example Scenario:

1. Software writes APB → descriptor_addr = 0x1000
 - channel_idle = 1 (both idle)
 - APB accepted
2. Descriptor engine fetches descriptor @ 0x1000
 - descriptor_engine_idle = 0 (fetch in progress)
 - channel_idle = 0
 - APB BLOCKED
3. Descriptor pushed to scheduler
 - descriptor_engine_idle = 1 (fetch complete)
 - scheduler_idle = 0 (transfer starting)
 - channel_idle = 0
 - APB BLOCKED
4. Scheduler completes data transfer

- Descriptor has `next_descriptor_ptr = 0x1100` (chained!)
 - Descriptor engine autonomously fetches @ `0x1100`
 - `descriptor_engine_idle = 0` (autonomous fetch)
 - `channel_idle = 0`
 - APB BLOCKED
5. Final descriptor completes (last = 1 OR `next_ptr = 0`)
- `scheduler_idle = 1` (transfer done)
 - `descriptor_engine_idle = 1` (no more fetches)
 - `channel_idle = 1`
 - APB UNBLOCKED (ready for next transfer!)

Key Insight: The AND gate ensures software cannot interrupt a descriptor chain in progress!

22.6.35.3 Scheduler (Simplified from RAPIDS)

Purpose: - Coordinate descriptor-to-data-transfer flow - Manage 8 independent channels - Arbitrate shared resources (SRAM, AXI masters)

FSM States:

```
typedef enum logic [7:0] {
    SCHED_IDLE          = 8'b00000001, // Idle, waiting for
    channel activation
    SCHED_FETCH_DESCRIPTOR = 8'b00000010, // Fetch descriptor via
    AXI master
    SCHED_PARSE_DESCRIPTOR = 8'b00000100, // Parse descriptor fields
    SCHED_READ_PHASE     = 8'b00001000, // Coordinate read engine
    SCHED_WRITE_PHASE    = 8'b00010000, // Coordinate write engine
    SCHED_CHAIN_CHECK    = 8'b00100000, // Check for next
    descriptor
    SCHED_COMPLETE       = 8'b01000000, // Transfer complete,
    report status
    SCHED_ERROR          = 8'b10000000 // Error state
} scheduler_state_t;
```

Key Differences from RAPIDS: - ✗ No credit management (just simple transaction limits) - ✗ No program engine coordination (no alignment fixup) - ✓ Simplified FSM (no control read/write phases)

22.6.45.3 AXI Read Engine (Streaming Pipeline - NO FSM)

Purpose: High-performance streaming reads from memory to SRAM buffer

Architecture: Pipelined streaming design (NO FSM for performance)

Key Insight: FSMs are horrible for performance! Instead, use: - **Arbiter** selects which channel's `datard_*` interface gets access - **Streaming pipeline** continuously moves data when granted - **Data interface** (`datard_valid`, `datard_ready`, `datard_beats_remaining`) controls flow

Data Read Interface (per channel):

```
// Channel requests read access
 datard_valid;           // Channel has read request
 datard_ready;          // Engine ready for request
 [63:0] datard_addr;        // Source address (aligned)
 [31:0] datard_beats_remaining; // Beats left to read
 [7:0] datard_burst_len;      // Preferred burst length
 [3:0] datard_channel_id;    // Channel ID for tracking
```

Multiple Versions (Compile-Time Selection): 1. **Version 1 - Basic:** Single outstanding read, fixed burst length 2. **Version 2 - Pipelined:** Multiple outstanding reads, configurable bursts 3. **Version 3 - Adaptive:** Dynamic burst sizing based on remaining beats

Operation:

1. Arbiter grants channel access based on priority
2. Engine accepts datard_* request (continuous streaming)
3. AXI AR channel issues read burst
4. AXI R channel streams data → SRAM (no FSM stalls!)
5. Engine updates beats_remaining, accepts next request
6. Different channels can have different burst lengths (e.g., CH0: 8 beats, CH1: 16 beats)

Example: Different Burst Lengths per Channel

```
// Channel 0 prefers 8-beat bursts
datard_burst_len[0] = 8'd8;

// Channel 1 prefers 16-beat bursts
datard_burst_len[1] = 8'd16;

// Engine adapts to requested burst length (within MAX_BURST_LEN)
```

22.6.5.4 AXI Write Engine (Streaming Pipeline - NO FSM)

Purpose: High-performance streaming writes from SRAM buffer to memory

Architecture: Pipelined streaming design (NO FSM for performance)

Key Insight: Same as read engine - no FSMs! Use streaming pipeline with arbiter.

Data Write Interface (per channel):

```
// Channel requests write access
 datawr_valid;           // Channel has write request
 datawr_ready;          // Engine ready for request
 [63:0] datawr_addr;        // Destination address
```

```

(aligned)
input logic [31:0] datawr_beats_remaining; // Beats left to write
input logic [7:0] datawr_burst_len; // Preferred burst length
input logic [3:0] datawr_channel_id; // Channel ID for tracking

```

Multiple Versions (Compile-Time Selection): 1. **Version 1 - Basic:** Single outstanding write, fixed burst length 2. **Version 2 - Pipelined:** Multiple outstanding writes, configurable bursts 3. **Version 3 - Adaptive:** Dynamic burst sizing based on remaining beats

Operation:

1. Arbiter grants channel access based on priority
2. Engine accepts datawr_* request (continuous streaming)
3. Engine reads data from SRAM
4. AXI AW channel issues write address
5. AXI W channel streams data (no FSM stalls!)
6. AXI B channel receives response, updates beats_remaining
7. Different channels can have different burst lengths (e.g., CH0: 8 beats, CH2: 32 beats)

Read/Write Asymmetry Example:

```

// Channel can use different burst lengths for read vs write
// Example: Read in small bursts, write in large bursts
datard_burst_len[0] = 8'd8; // Read: 8 beats
datawr_burst_len[0] = 8'd16; // Write: 16 beats

// Engine handles asymmetry via SRAM buffering

```

22.6.6.5.5 Simple SRAM

Source: Direct copy from RAPIDS simple_sram.sv

Purpose: - Dual-port SRAM buffer - Decouples read and write engines - Shared across all channels (arbitration required)

Why Reuse: - Standard dual-port SRAM design - Proven in RAPIDS integration tests - Parameterizable depth and width

22.7 6. Configuration and Control

22.7.16.1 APB Register Map

| Offset | Register | Access | Description |
|--------|---------------|--------|----------------------------|
| 0x0000 | GLOBAL_CTRL | RW | Global enable, reset |
| 0x0004 | GLOBAL_STATUS | RO | Global status, error flags |

| Offset | Register | Access | Description |
|--------|----------------|--------|---|
| 0x0100 | CH0_CTRL | WO | Channel 0 kick-off (write descriptor address) |
| 0x0104 | CH0_STATUS | RO | Channel 0 status |
| 0x0108 | CH0_DESC_ADDR | RO | Channel 0 current descriptor address |
| 0x010C | CH0_BYTES_XFER | RO | Channel 0 bytes transferred |
| ... | ... | ... | ... (repeat for channels 1-7) |
| 0x0200 | CH1_CTRL | WO | Channel 1 kick-off |
| ... | | | |
| 0x0700 | CH7_CTRL | WO | Channel 7 kick-off |

Kick-Off Sequence: 1. Software writes descriptor address to CHx_CTRL register 2. STREAM fetches descriptor from memory 3. Transfer begins automatically 4. If chained, STREAM follows next_descriptor_ptr automatically 5. Completion reported via MonBus packet

22.7.26.2 Channel Configuration

Each channel independently configurable: - **Descriptor start address:** Written to CHx_CTRL - **Priority:** Encoded in descriptor (arbitration) - **Interrupt enable:** Per-descriptor flag - **Status monitoring:** Read CHx_STATUS

22.8 7. Resource Sharing and Arbitration

22.8.1 7.1 Shared Resources

All channels share: 1. **Descriptor Fetch AXI Master** - Fetches descriptors for all channels 2. **Data Read AXI Master** - Reads source data for all channels 3. **Data Write AXI Master** - Writes destination data for all channels 4. **SRAM Buffer** - Shared buffer (dual-port, but still arbitrated) 5. **MonBus Reporter** - Single monitor output

22.8.2 7.2 Arbitration Strategy

Priority-Based Round-Robin: - Channels have priority field in descriptor - Higher priority = serviced first - Within same priority: round-robin - Prevents starvation with timeout

Example Arbitration:

Channel 0: Priority 7 (highest)
Channel 1: Priority 5
Channel 2: Priority 5
Channel 3: Priority 3

Service order: CH0 → CH1 → CH2 (round-robin) → CH0 → CH1 → CH2 → CH3 ...

22.9 8. Error Detection and Recovery

22.9.18.1 Error Types

| Error Type | Detection | Response |
|---------------------------|---------------------|------------------------------|
| Invalid descriptor | Valid bit = 0 | Skip, move to next |
| Alignment error | Address not aligned | Set error flag, halt channel |
| AXI SLVERR | AXI response | Set error flag, halt channel |
| AXI DECERR | AXI response | Set error flag, halt channel |
| Timeout | Transaction timeout | Set error flag, halt channel |
| SRAM overflow | Buffer full | Backpressure, wait |

22.9.28.2 Error Recovery

Per-Channel Error Handling: - Error sets channel to CH_ERROR state - Channel halts, does not affect other channels - Software must: 1. Read CHx_STATUS to identify error 2. Clear error condition 3. Re-kick channel with new descriptor

No Automatic Retry: - Tutorial design keeps error handling simple - Software responsible for retry logic

22.10 9. MonBus Integration

22.10.1 9.1 Standard MonBus Format

Uses standard 64-bit MonBus packet format: - [63:60] Packet Type (0=ERROR, 1=COMPL, etc.) - [59:57] Protocol (custom STREAM protocol) - [56:53] Event Code (STREAM-specific events) - [52:47]

Channel ID (0-7) - [46:43] Unit ID (unused for STREAM) - [42:35] Agent ID (unused for STREAM) - [34:0] Event Data (address, byte count, etc.)

22.10.2 9.2 STREAM Event Codes

| Code | Event | Description |
|------|----------------|--------------------------|
| 0x0 | DESC_START | Descriptor started |
| 0x1 | DESC_COMPLETE | Descriptor completed |
| 0x2 | READ_START | Read phase started |
| 0x3 | READ_COMPLETE | Read phase completed |
| 0x4 | WRITE_START | Write phase started |
| 0x5 | WRITE_COMPLETE | Write phase completed |
| 0x6 | CHAIN_FETCH | Chained descriptor fetch |
| 0xF | ERROR | Error occurred |

22.10.3 9.3 Default Configuration

Tutorial-Friendly Defaults: - Errors only: `cfg_error_enable = 1`, all others = 0 - **Interrupts:** Descriptor flag controls per-transfer interrupt - **Reduces MonBus traffic** for beginner understanding

22.11 10. Design Constraints

22.11.1 10.1 Tutorial Constraints (Intentional Simplifications)

| Constraint | Rationale |
|-------------------------------|---|
| Aligned addresses only | Simplify data path, hide alignment complexity |
| Length in beats | Avoid byte/chunk conversion math |
| No circular buffers | Explicit termination easier to understand |
| Single APB kick-off | Simple software model |
| No credit management | Avoid exponential encoding |

| Constraint | Rationale |
|------------|------------|
| | complexity |

22.11.2 10.2 Implementation Constraints

| Parameter | Value | Notes |
|------------------|-----------------|----------------------------|
| Max channels | 8 | Compile-time parameter |
| Max burst length | 256 | AXI4 spec limit |
| Descriptor size | 256 bits | 4 × 64-bit words |
| SRAM depth | Parameterizable | Typical: 1024-4096 entries |
| Data width | Parameterizable | Typical: 512-bit |

22.12 11. Verification Strategy

22.12.1 11.1 Test Organization

```
projects/components/stream/dv/tests/
└── fub_tests/
    ├── descriptor_engine/          # Functional Unit Block tests
    ├── scheduler/                 # Copy from RAPIDS (adapt imports)
    ├── axi_engines/               # Simplified scheduler tests
    └── sram/                      # Read/write engine tests

└── integration_tests/
    ├── single_channel/           # Multi-block scenarios
    ├── multi_channel/            # Single channel transfers
    ├── chained_descriptors/      # 8-channel concurrent
    └── error_handling/           # Descriptor chain tests
                                # Error recovery tests
```

22.12.2 11.2 Test Levels

Basic (Quick Smoke Tests): - Single descriptor, single channel - Aligned addresses, simple transfers - ~30 seconds runtime

Medium (Typical Scenarios): - Multiple descriptors, 2-4 channels - Chained descriptors (2-3 deep) - ~90 seconds runtime

Full (Comprehensive Validation): - All 8 channels concurrent - Long descriptor chains (10+ deep) - Error injection, stress testing - ~180 seconds runtime

22.13 12. Performance Characteristics

22.13.1 12.1 Throughput by Engine Version

V1 (Low Performance - Tutorial Mode): - **Throughput:** 0.14 beats/cycle (DDR4), 0.40 beats/cycle (SRAM) - **Architecture:** Single outstanding transaction, blocks on completion - **Use Case:** Tutorial examples, embedded systems, low-latency SRAM

V2 (Medium Performance - Command Pipelined): - **Throughput:** 0.94 beats/cycle (DDR4), 0.85 beats/cycle (SRAM) - **Improvement:** 6.7x over V1 (DDR4), 2.1x over V1 (SRAM) - **Architecture:** Command queue (4-8 deep), hides memory latency - **Use Case:** General-purpose FPGA, DDR3/DDR4, best area efficiency

V3 (High Performance - Out-of-Order): - **Throughput:** 0.98 beats/cycle (DDR4), 0.92 beats/cycle (SRAM) - **Improvement:** 7.0x over V1 (DDR4), 2.3x over V1 (SRAM) - **Architecture:** OOO command selection, maximizes memory controller efficiency - **Use Case:** Datacenter, ASIC, HBM2, high-performance memory

Key Insight: Benefit scales with memory latency. V1 throughput degrades from 40% (SRAM) to 14% (DDR4), while V2/V3 maintain 94-98% throughput regardless of latency.

Configuration Parameters: - `ENABLE_CMD_PIPELINE = 0`: V1 (default, tutorial mode) - `ENABLE_CMD_PIPELINE = 1`: V2 (command pipelined) - `ENABLE_CMD_PIPELINE = 1`, `ENABLE_000_DRAIN = 1` (write) or `ENABLE_000_READ = 1` (read): V3

Factors Affecting Throughput: - Memory latency (V2/V3 hide latency via pipelining) - SRAM buffer size (limits burst pipelining) - Channel arbitration overhead - Descriptor fetch latency - Engine version (V1/V2/V3 configuration)

22.13.2 12.2 Latency

Transfer Latency Breakdown: - Descriptor fetch: ~10-50 cycles (depends on memory) - Read phase: $(\text{length} / \text{burst_len}) \times \text{burst_latency}$ - Write phase: $(\text{length} / \text{burst_len}) \times \text{burst_latency}$ - End-to-end: Typically <200 cycles for small transfers (V1), <100 cycles (V2/V3 pipelined)

V2/V3 Latency Hiding: - Command pipelining overlaps descriptor fetch, read, write operations - Multiple outstanding transactions hide memory latency - OOO completion (V3) reduces head-of-line blocking

22.13.3 12.3 Resource Utilization by Engine Version

V1 Configuration (Tutorial - Minimum Area): - Total: ~9,500 LUTs + 64 KB SRAM - Descriptor Engine (8×): ~2,400 LUTs - Scheduler (8×): ~3,200 LUTs - AXI Read Engine: ~1,250 LUTs - AXI Write Engine: ~1,250 LUTs - SRAM Controller: ~1,600 LUTs - APB Config: ~350 LUTs - MonBus AXIL Group: ~1,000 LUTs

V2 Configuration (Balanced - Best Area Efficiency): - Total: ~11,000 LUTs + 64 KB SRAM (1.16x area) - AXI Read Engine: ~2,000 LUTs (1.6x increase, 6.7x throughput) - AXI Write Engine: ~2,500 LUTs (2.0x increase, 6.7x throughput) - Other blocks: Same as V1

V3 Configuration (Maximum Performance): - Total: ~14,000 LUTs + 64 KB SRAM (1.47x area) - AXI Read Engine: ~3,500 LUTs (2.8x increase, 7.0x throughput) - AXI Write Engine: ~4,000 LUTs (3.2x increase, 7.0x throughput) - Other blocks: Same as V1

Area Efficiency Comparison: - V1: 1.00 throughput / 1.00 area = 1.00 - V2: 6.70 throughput / 1.16 area = 5.78 (best efficiency) - V3: 7.00 throughput / 1.47 area = 4.76

Recommendation: V2 provides best area efficiency for most use cases. V3 justified only for high-performance memory controllers that support OOO responses.

22.14 13. Development Roadmap

22.14.1 13.1 Phase 1: Foundation (Current)

- ✓ Directory structure
- ✓ Package definitions (`stream_pkg.sv`)
- ✓ Imports header (`stream_imports.svh`)
- ⏳ Documentation (this PRD)

22.14.2 13.2 Phase 2: Core Blocks

- Adapt descriptor engine from RAPIDS
- Design simplified scheduler
- Create APB config interface
- Copy simple SRAM from RAPIDS

22.14.3 13.3 Phase 3: Data Path

- AXI read engine (version 1 - basic)
- AXI write engine (version 1 - basic)
- SRAM integration
- Channel arbitration

22.14.4 13.4 Phase 4: Integration

- Top-level module
- MonBus reporter
- Integration testbench

- Single-channel validation

22.14.5 13.5 Phase 5: Multi-Channel

- 8-channel support
- Arbiter implementation
- Multi-channel tests
- Performance tuning

22.14.6 13.6 Phase 6: Advanced Engines (Future - V2/V3)

Goal: Add parameterized high-performance engine variants

V2 - Command Pipelined (Medium Performance): - Command queue implementation (4-8 deep) - W drain FSM for write engine - B response scoreboard (write) or in-order R reception (read) - Per-command SRAM pointer tracking - Parameter: ENABLE_CMD_PIPELINE = 1 - Expected: 6.7x throughput improvement over V1

V3 - Out-of-Order Completion (High Performance): - OOO command selection logic - Transaction ID matching (AXI ID to queue entry) - SRAM data availability checking (write engine) - R beat matching to queue entry (read engine) - Parameters: ENABLE_CMD_PIPELINE = 1, ENABLE_000_DRAIN = 1 (write) or ENABLE_000_READ = 1 (read) - Expected: 7.0x throughput improvement over V1

Deliverables: - Updated RTL with parameterization - Performance comparison tests (V1 vs V2 vs V3) - Tutorial documentation explaining trade-offs - Area/throughput measurements on target FPGA

22.15 14. Educational Value

22.15.1 14.1 Learning Objectives

STREAM teaches: 1. **Descriptor-based DMA design patterns** - Descriptor structure and parsing - Chained descriptors (scatter-gather) - Descriptor fetch via AXI

2. AXI4 Memory Interface Usage

- Burst transactions
- Address/data/response channels
- Outstanding transactions

3. Resource Sharing and Arbitration

- Multiple channels sharing AXI masters
- Priority-based arbitration

- Conflict resolution
- 4. FSM Design and Coordination**
- Multiple interconnected FSMs
 - State machine composition
 - Error handling
- 5. Buffer Management**
- SRAM-based buffering
 - Rate matching
 - Flow control

22.15.2 14.2 Progression Path

Learning Sequence: 1. **STREAM (this project):** Memory-to-memory DMA, aligned addresses 2. **STREAM Extended:** Add alignment fixup, more complex scenarios 3. **RAPIDS:** Add network interfaces, credit management, full complexity

22.16 15. Success Criteria

22.16.1 15.1 Functional

- ✓ Single descriptor transfer working
- ✓ Chained descriptors (2-3 deep)
- ✓ Multi-channel operation (8 channels concurrent)
- ✓ Error detection and reporting
- ✓ MonBus packet generation
- ✓ >90% functional test coverage

22.16.2 15.2 Quality

- ✓ Verilator compiles with 0 warnings
- ✓ All tests passing (100% success rate)
- ✓ Comprehensive documentation
- ✓ Tutorial-quality code comments

22.16.3 15.3 Performance

- ✓ Achieves >80% of theoretical AXI bandwidth
- ✓ <5K LUT utilization (excluding SRAM)
- ✓ <200 cycle end-to-end latency for small transfers

22.17 16. Open Questions (For Review)

22.17.1 16.1 Descriptor Engine Adaptation

- **Q:** Should descriptor engine use APB-only, RDA-only, or mixed mode?
- **A (pending):** TBD - depends on software use case preference

22.17.2 16.2 AXI Descriptor Master

- **Q:** Fixed 256-bit width, or parameterizable?
- **A (pending):** Propose fixed 256-bit for simplicity

22.17.3 16.3 Channel Arbitration

- **Q:** Fixed priority, or dynamic priority based on age/fairness?
- **A (pending):** Propose fixed priority with round-robin for tutorial simplicity

22.17.4 16.4 SRAM Partitioning

- **Q:** Single shared SRAM, or per-channel SRAMs?
 - **A (pending):** Propose single shared SRAM with arbitration (matches RAPIDS pattern)
-

22.18 16. Attribution and Contribution Guidelines

22.18.1 16.1 Git Commit Attribution

When creating git commits for STREAM documentation or implementation:

Use:

Documentation and implementation support by Claude.

Do NOT use:

Co-Authored-By: Claude <noreply@anthropic.com>

Rationale: STREAM documentation and organization receives AI assistance for structure and clarity, while design concepts and architectural decisions remain human-authored.

22.19 16.2 PDF Generation Location

IMPORTANT: PDF files should be generated in the docs directory:

```
/mnt/data/github/rtldesignsherpa/projects/components/stream/docs/
```

Quick Command: Use the provided shell script:

```
cd /mnt/data/github/rtldesignsherpa/projects/components/stream/docs  
./generate_pdf.sh
```

The shell script will automatically: 1. Use the md_to_docx.py tool from bin/ 2. Process the stream_spec index file 3. Generate both DOCX and PDF files in the docs/ directory 4. Create table of contents and title page

 See: bin/md_to_docx.py for complete implementation details

22.20 17. References

22.20.1 17.1 Internal Documentation

- **RAPIDS PRD:** projects/components/rapids/PRD.md - Parent architecture
- **RAPIDS Descriptor Engine:**
projects/components/rapids/rtl/rapids_fub/descriptor_engine.sv
- **RAPIDS Simple SRAM:**
projects/components/rapids/rtl/rapids_fub/simple_sram.sv
- **AMBA PRD:** rtl/amba/PRD.md - MonBus integration
- **Repository Guide:** /CLAUDE.md - Design patterns and conventions

22.20.2 17.2 External References

- **AXI4 Specification:** ARM IHI0022E
 - **APB Specification:** ARM IHI0024C
 - CocoTB Documentation: <https://docs.cocotb.org/>
 - Verilator Manual: <https://verilator.org/guide/latest/>
-

Document Version: 1.0 **Last Updated:** 2025-10-17 **Review Cycle:** Weekly during initial design

Next Review: TBD (after user feedback) **Owner:** RTL Design Sherpa Project

22.21 Navigation

- ← Back to Root: /PRD.md
- Parent Architecture: projects/components/rapids/PRD.md
- AI Guidance: CLAUDE.md (to be created)
- Quick Start: README.md (to be created)

RTL Design Sherpa · Learning Hardware Design Through Practice GitHub · Documentation Index · MIT License

23 Claude Code Guide: STREAM Subsystem

Version: 1.0 **Last Updated:** 2025-10-17 **Purpose:** AI-specific guidance for working with STREAM subsystem

23.1 Quick Context

What: STREAM - Scatter-gather Transfer Rapid Engine for AXI Memory **Status:** 🟡 Initial design - tutorial-focused DMA engine **Your Role:** Help users build a beginner-friendly descriptor-based DMA engine

📖 **Complete Specification:** projects/components/stream/PRD.md ← **Always reference this for technical details**

23.2 📖 Global Requirements Reference

IMPORTANT: Review /GLOBAL_REQUIREMENTS.md for all mandatory requirements

All mandatory requirements are consolidated in the global requirements document: - **See:** /GLOBAL_REQUIREMENTS.md - Repository-wide mandatory requirements - **STREAM-Specific:** Attribution format, tutorial focus (intentional simplifications) - **Universal:** TB location, three methods, TBBBase inheritance, 100% success

This CLAUDE.md provides STREAM-specific guidance. Also review: - Root /CLAUDE.md - Repository-wide patterns - projects/components/CLAUDE.md - Project area standards (reset macros, FPGA attributes) - bin/CocoTBFramework/CLAUDE.md - Framework usage patterns

23.3 Critical Rules for This Subsystem

23.3.1 Rule #0: Attribution Format for Git Commits

IMPORTANT: When creating git commit messages for STREAM documentation or code:

Use:

Documentation and implementation support by Claude.

Do NOT use:

Co-Authored-By: Claude <noreply@anthropic.com>

Rationale: STREAM receives AI assistance for structure and clarity, while design concepts and architectural decisions remain human-authored.

23.3.2 Rule #0.1: TUTORIAL FOCUS - Intentional Simplifications

⚠ STREAM is INTENTIONALLY SIMPLIFIED for educational purposes ⚠

Key Simplifications (DO NOT “fix” these): 1. ✓ **Aligned addresses only** - No alignment fixup logic (kept for RAPIDS) 2. ✓ **Length in beats** - Not bytes or chunks (simplifies math) 3. ✓ **No circular buffers** - Explicit chain termination only 4. ✓ **No credit management** - Simple transaction limits 5. ✓ **Pure memory-to-memory** - No network interfaces

When users ask “Can we add alignment fixup?”: - ✓ **Correct answer:** “STREAM intentionally keeps addresses aligned for tutorial simplicity. For complex alignment, see RAPIDS.” - ✗ **Wrong answer:** “Sure, let me add alignment logic...” (defeats tutorial purpose!)

23.3.3 Rule #0.1: Testbench Location and Test Structure (MANDATORY)

 See: /GLOBAL REQUIREMENTS.md Section 2.1 for complete requirement

STREAM-Specific Directory Structure:

STREAM Import Pattern:

```

# Import STREAM TB from project area
from projects.components.stream.dv.tbclasses.scheduler_tb import
StreamSchedulerTB

# Shared utilities from framework
from CocoTBFramework.tbclasses.shared.tbbase import TBBBase

```

Complete Pattern: [projects/components/rapids/CLAUDE.md](#) Rule #0.1

23.3.4 Rule #0.2: Three Mandatory TB Methods (MANDATORY)

See: [/GLOBAL_REQUIREMENTS.md](#) Section 2.2 for complete requirement

STREAM-Specific Context:

STREAM has simpler config requirements than RAPIDS - most config can be set after reset.

Standard STREAM Pattern:

```

class StreamSchedulerTB(TBBBase):
    async def setup_clocks_and_reset(self):
        """Standard STREAM initialization"""
        await self.start_clock('clk', freq=10, units='ns')
        await self.assert_reset()
        await self.wait_clocks('clk', 10)
        await self.deassert_reset()
        await self.wait_clocks('clk', 5)

    async def assert_reset(self):
        self.dut.rst_n.value = 0

    async def deassert_reset(self):
        self.dut.rst_n.value = 1

```

Note: Unlike RAPIDS, STREAM typically doesn't need config set before reset.

23.3.5 Rule #1: REUSE from RAPIDS Where Appropriate

Direct Reuse (No Modification): - ✓ descriptor_engine.sv - Works with STREAM descriptors - ✓ simple_sram.sv - Standard dual-port SRAM - ✓ Descriptor engine tests - Adapt imports only

Adapt from RAPIDS: - ⚠ scheduler.sv - **Simplify FSM** (no credit management, no control engines) - ⚠ AXI engines - **Create simplified versions** (no alignment fixup)

Create New for STREAM: - APB config interface - PeakRDL-generated (like HPET), 8 channels, kick-off registers - Top-level integration - Different interface set

Always Ask Yourself: “Can I reuse from RAPIDS instead of creating new?”

23.3.6 Rule #2: Descriptor Format is DIFFERENT from RAPIDS

STREAM Descriptor (256-bit): - `src_addr` (64-bit), `dst_addr` (64-bit), `length (beats)`, `next_descriptor_ptr` (32-bit) - **Length is in BEATS, not chunks!**

RAPIDS Descriptor: - Uses chunks (4-byte units) - Has alignment metadata

When comparing/referencing RAPIDS: - ✓ “RAPIDS uses chunks, STREAM uses beats for tutorial simplicity” - ✗ Don’t assume RAPIDS descriptor format applies to STREAM

23.3.7 Rule #3: Know the Shared Resources

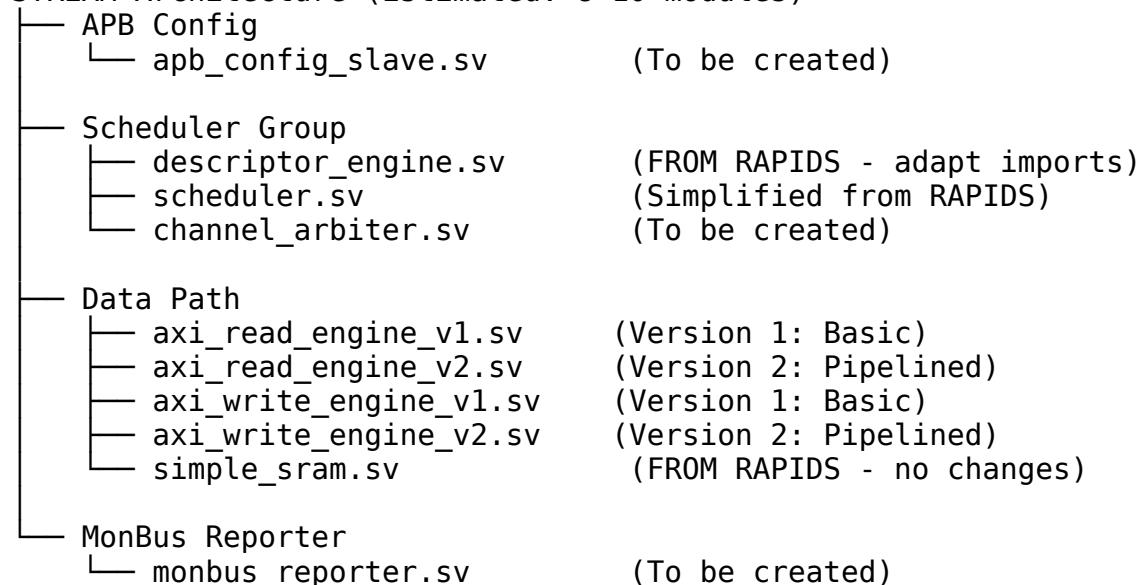
All 8 channels share: 1. Descriptor fetch AXI master 2. Data read AXI master 3. Data write AXI master 4. SRAM buffer 5. MonBus reporter

Arbitration is required! - Never assume exclusive access - All shared resources need arbiter logic - Priority-based round-robin (descriptor priority field)

23.4 Architecture Quick Reference

23.4.1 Block Organization

STREAM Architecture (Estimated: 8-10 modules)



23.4.2 Module Status

| Module | Source | Status | Notes |
|-----------------------------------|---------------|----------|---------------------|
| <code>descriptor_engine.sv</code> | RAPIDS (copy) | ✓ Copied | Adapt #include only |

| Module | Source | Status | Notes |
|------------------------|-------------------|-----------|-------------------------------------|
| simple_sram.sv | RAPIDS (copy) | ✓ Copied | No changes needed |
| stream_pkg.sv | New | ✓ Created | Descriptor format defined |
| stream_imports.svh | New | ✓ Created | Package imports |
| scheduler.sv | RAPIDS (simplify) | ⌚ Pending | Remove credit mgmt, control engines |
| apb_config_slave.sv | New | ⌚ Pending | 8 channel registers |
| axi_read_engine_v1.sv | New | ⌚ Pending | Basic version |
| axi_write_engine_v1.sv | New | ⌚ Pending | Basic version |
| channel_arbiter.sv | New | ⌚ Pending | Priority-based round-robin |
| monbus_reporter.sv | New | ⌚ Pending | STREAM event codes |
| stream_top.sv | New | ⌚ Pending | Top-level integration |

23.5 Common User Questions and Responses

23.5.1 Q: "How is STREAM different from RAPIDS?"

A: STREAM is intentionally simplified for tutorial purposes:

Simplifications: | Feature | RAPIDS | STREAM | |——|——|——| | **Interfaces** | APB + AXI + Network | APB + AXI only | | **Address Alignment** | Complex fixup | Aligned only | | **Credit Management** | Exponential encoding | Simple limits | | **Descriptor Length** | Chunks (4-byte) | Beats (data width) | | **Control Engines** | Control read/write | None (direct APB) |

Learning Path: 1. **STREAM** - Basic DMA with scatter-gather 2. **STREAM Extended** - Add alignment fixup 3. **RAPIDS** - Full complexity with network + credit management

 See: PRD.md Section 2.1 for complete comparison table

23.5.2 Q: "How do I kick off a transfer?"

A: Single APB write starts descriptor chain:

```
// Software writes descriptor address to channel register
write_apb(ADDR_CH0_CTRL, 0x1000_0000); // Start address of descriptor

// STREAM automatically:
// 1. Fetches descriptor from 0x1000_0000
// 2. Parses src_addr, dst_addr, length
// 3. Reads source data → SRAM
// 4. Writes SRAM → destination
// 5. Checks next_descriptor_ptr
//     - If != 0: Fetch next descriptor, repeat
//     - If == 0 or last flag set: Complete
```

Chained Descriptors:

Descriptor 0 @ 0x1000_0000:

```
src_addr = 0x2000_0000
dst_addr = 0x3000_0000
length = 64 beats
next_descriptor_ptr = 0x1000_0100 ← Points to next descriptor
```

Descriptor 1 @ 0x1000_0100:

```
src_addr = 0x2000_1000
dst_addr = 0x3000_1000
length = 32 beats
next_descriptor_ptr = 0x0000_0000 ← Last descriptor (0 = stop)
```

 See: PRD.md Section 3.2 for complete data flow

23.5.3 Q: "How many channels can I use?"

A: Maximum 8 independent channels:

Channel Operation: - Each channel has own FSM state - Each channel can have separate descriptor chain - All channels share: AXI masters, SRAM, MonBus

Arbitration: - Priority-based (descriptor priority field) - Round-robin within same priority - Prevents starvation with timeout

Example:

```
// Kick off 3 channels concurrently
write_apb(ADDR_CH0_CTRL, desc0_addr); // Channel 0: Priority 7
write_apb(ADDR_CH1_CTRL, desc1_addr); // Channel 1: Priority 5
write_apb(ADDR_CH2_CTRL, desc2_addr); // Channel 2: Priority 5
```

// Service order: CH0 → CH1 → CH2 → CH0 → CH1 → CH2 ...

📖 See: PRD.md Section 7 for arbitration details

23.5.4 Q: "What's the descriptor format?"

A: 256-bit descriptor (4 × 64-bit words):

```
typedef struct packed {
    logic [63:0] reserved;           // [255:192] Reserved
    logic [7:0] priority;           // [207:200] Priority
    logic [3:0] channel_id;         // [199:196] Channel ID
    logic       error;              // [195] Error flag
    logic       last;               // [194] Last in chain
    logic       interrupt;          // [193] Generate interrupt
    logic       valid;              // [192] Valid descriptor
    logic [31:0] next_descriptor_ptr; // [191:160] Next descriptor
address
    logic [31:0] length;            // [159:128] Length in BEATS
*
    logic [63:0] dst_addr;          // [127:64] Destination
address
    logic [63:0] src_addr;          // [63:0] Source address
} descriptor_t;
```

★ CRITICAL: length is in BEATS, not bytes or chunks!

Example:

Data width = 512 bits = 64 bytes

Length = 16 beats

Total transfer = 16 × 64 = 1024 bytes

📖 See: PRD.md Section 4.2 for complete descriptor specification

23.5.5 Q: "How do I run STREAM tests?"

A: Multi-layered test approach (same as RAPIDS):

```
# 1. FUB (Functional Unit Block) Tests - Individual blocks
pytest projects/components/stream/dv/tests/fub_tests/scheduler/ -v
pytest
projects/components/stream/dv/tests/fub_tests(descriptor_engine/ -v
pytest projects/components/stream/dv/tests/fub_tests/ -v # All FUB
tests

# 2. Integration Tests - Multi-block scenarios
pytest projects/components/stream/dv/tests/integration_tests/ -v
```

```
# Run with waveforms for debugging
pytest projects/components/stream/dv/tests/fub_tests/scheduler/ --
vcd=debug.vcd
gtkwave debug.vcd
```

Test Organization: - **FUB tests:** Focus on individual block functionality - **Integration tests:** Verify block-to-block interfaces and complete data flows

23.6 Integration Patterns

23.6.1 Pattern 1: Basic STREAM Instantiation

```
stream_top #(
    .NUM_CHANNELS(8),
    .DATA_WIDTH(512),
    .ADDR_WIDTH(64),
    .SRAM_DEPTH(4096)
) u_stream (
    // Clock and Reset
    .aclk                (system_clk),
    .aresetn             (system_rst_n),

    // APB Configuration Interface
    .s_apb_paddr         (apb_paddr),
    .s_apb_psel          (apb_psel),
    .s_apb_penable       (apb_penable),
    .s_apb_pwrite        (apb_pwrite),
    .s_apb_pwdata        (apb_pwdata),
    .s_apb_pready        (apb_pready),
    .s_apb_prdata        (apb_prdata),
    .s_apb_pslverr       (apb_pslverr),

    // AXI Master - Descriptor Fetch (256-bit)
    .m_axi_desc_araddr   (desc_araddr),
    .m_axi_desc_arlen    (desc_arlen),
    .m_axi_desc_arsize   (desc_arsize),
    .m_axi_desc_arvalid  (desc_arvalid),
    .m_axi_desc_arready  (desc_arready),
    .m_axi_desc_rdata    (desc_rdata),
    .m_axi_desc_rresp    (desc_rresp),
    .m_axi_desc_rlast    (desc_rlast),
    .m_axi_desc_rvalid   (desc_rvalid),
    .m_axi_desc_rready   (desc_rready),
```

```

// AXI Master - Data Read (parameterizable width)
.m_axi_rd_araddr      (rd_araddr),
// ... (full AXI4 AR + R channels)

// AXI Master - Data Write (parameterizable width)
.m_axi_wr_awaddr      (wr_awaddr),
// ... (full AXI4 AW + W + B channels)

// MonBus Output
.monbus_pkt_valid     (stream_mon_valid),
.monbus_pkt_ready     (stream_mon_ready),
.monbus_pkt_data      (stream_mon_data)
);

```

23.6.2 Pattern 2: Descriptor Creation (Software Model)

```

// C/C++ software model for descriptor creation
typedef struct {
    uint64_t src_addr;           // Source address (must be aligned!)
    uint64_t dst_addr;           // Destination address (must be
aligned!)
    uint32_t length;             // Transfer length in BEATS
    uint32_t next_descriptor_ptr; // Next descriptor address (0 =
last)
    uint8_t control;             // valid | interrupt | last | error |
channel_id | priority
} stream_descriptor_t;

// Create descriptor chain
stream_descriptor_t desc[2];

// Descriptor 0
desc[0].src_addr = 0x80000000ULL; // Aligned to 64B (512-bit data)
desc[0].dst_addr = 0x90000000ULL;
desc[0].length = 64; // 64 beats x 64 bytes = 4KB transfer
desc[0].next_descriptor_ptr = (uint32_t)&desc[1]; // Chain to next
desc[0].control = 0x01; // valid = 1

// Descriptor 1 (last)
desc[1].src_addr = 0x80001000ULL;
desc[1].dst_addr = 0x90001000ULL;
desc[1].length = 32; // 32 beats x 64 bytes = 2KB transfer
desc[1].next_descriptor_ptr = 0; // Last descriptor
desc[1].control = 0x45; // valid | last | interrupt

// Kick off transfer
write_apb_reg(CH0_CTRL, (uint32_t)&desc[0]);

```

23.6.3 Pattern 3: MonBus Integration

```
// Always add downstream FIFO for MonBus
gaxi_fifo_sync #(
    .DATA_WIDTH(64),
    .DEPTH(256)
) u_stream_mon_fifo (
    .i_clk      (aclk),
    .i_rst_n   (aresetn),
    .i_data     (monbus_pkt_data),
    .i_valid    (monbus_pkt_valid),
    .o_ready    (monbus_pkt_ready),
    .o_data     (fifo_mon_data),
    .o_valid    (fifo_mon_valid),
    .i_ready    (consumer_ready)
);
```

23.7 Anti-Patterns to Catch

23.7.1 X Anti-Pattern 1: Adding Alignment Fixup

✗ WRONG:

"Let me add alignment fixup logic to handle unaligned addresses..."

✓ CORRECTED:

"STREAM intentionally requires aligned addresses for tutorial simplicity.

If you need unaligned transfers, that's a great learning exercise **for** extending STREAM, **or use** RAPIDS which has full alignment support."

23.7.2 X Anti-Pattern 2: Using Length in Bytes

✗ WRONG:

descriptor.length = 4096; // Thinking this is 4096 bytes

✓ CORRECTED:

```
// Length is in BEATS, not bytes!
// For 512-bit data width (64 bytes per beat):
descriptor.length = 4096 / 64; // = 64 beats for 4096 bytes
```

23.7.3 X Anti-Pattern 3: Circular Buffer Descriptors

✗ WRONG:

```
// Descriptor chain that loops back
desc[9].next_descriptor_ptr = &desc[0]; // Circular!
```

✓ CORRECTED:

```
"STREAM does not support circular buffers (no stop condition).
Always terminate chains explicitly:
desc[last].next_descriptor_ptr = 0; // Stop
desc[last].last = 1; // Explicit last flag
```

23.7.4 X Anti-Pattern 4: Assuming Exclusive Channel Access

x WRONG:

```
// Assume channel has exclusive AXI master access
assign m_axi_arvalid = channel_request; // No arbitration!
```

✓ CORRECTED:

```
// All channels share AXI masters - arbitration required
channel_arbiter u_arbiter (
    .ch_requests (channel_requests[7:0]),
    .ch_grant    (channel_grant[7:0]),
    .axi_master_available (axi_master_idle)
);
```

23.8 Debugging Workflow

23.8.1 Issue: Descriptor Not Fetched

Check in order: 1. ✓ APB write to CHx_CTRL register successful? 2. ✓ Descriptor address valid? 3. ✓ AXI descriptor master not stalled? 4. ✓ Descriptor memory accessible? 5. ✓ Arbiter granting descriptor fetch?

Debug commands:

```
pytest projects/components/stream/dv/tests/fub_tests/scheduler/ -v -s
# Verbose test
pytest projects/components/stream/dv/tests/fub_tests/scheduler/ --
vcd=debug.vcd
gtkwave debug.vcd # Inspect FSM state transitions
```

23.8.2 Issue: Data Transfer Stalls

Check in order: 1. ✓ SRAM buffer depth sufficient? 2. ✓ Source/destination addresses aligned? 3. ✓ AXI read/write engines getting grants? 4. ✓ Downstream AXI ready signals asserted? 5. ✓ Channel not in error state?

Waveform Analysis: - Check SRAM read/write pointers - Verify AXI AR/AW/R/W/B handshakes - Inspect scheduler FSM state - Check arbiter grant signals

23.8.3 Issue: Chained Descriptors Not Following

Check in order: 1. ✓ next_descriptor_ptr != 0? 2. ✓ last flag not set prematurely? 3. ✓ Descriptor fetch completing successfully? 4. ✓ Scheduler transitioning to CHAIN_CHECK state? 5. ✓ MonBus showing CHAIN_FETCH event?

23.9 Testing Guidance

23.9.1 Test Organization

```
projects/components/stream/dv/tests/
  └── fub_tests/                      # Individual block tests
      └── descriptor_engine/          # Adapt from RAPIDS tests
          └── scheduler/              # Simplified scheduler tests
          └── axi_engines/             # Read/write engine tests
          └── sram/                   # SRAM tests (from RAPIDS)

  └── integration_tests/            # Multi-block scenarios
      └── single_channel/           # Single channel transfers
      └── multi_channel/            # 8-channel concurrent
      └── chained_descriptors/     # Descriptor chain tests
          └── error_handling/        # Error recovery tests
```

23.9.2 Test Levels

Basic (Quick Smoke Tests): - Single descriptor, single channel - Aligned addresses, simple transfers - ~30 seconds runtime

Medium (Typical Scenarios): - Multiple descriptors, 2-4 channels - Chained descriptors (2-3 deep) - ~90 seconds runtime

Full (Comprehensive Validation): - All 8 channels concurrent - Long descriptor chains (10+ deep) - Error injection, stress testing - ~180 seconds runtime

23.10 Key Documentation Links

23.10.1 Always Reference These

This Subsystem: - projects/components/stream/PRD.md - **Complete specification** - projects/components/stream/README.md - Quick start guide (to be created) - projects/components/stream/known_issues/ - Bug tracking

Related: - projects/components/rapids/PRD.md - Parent architecture (for comparison) - rtl/amba/PRD.md - MonBus integration - /PRD.md - Master requirements - /CLAUDE.md - Repository guide

23.11 Quick Commands

```
# View complete specification  
cat projects/components/stream/PRD.md  
  
# Check package definition  
cat projects/components/stream/rtl/includes/stream_pkg.sv  
  
# Run tests (once created)  
pytest projects/components/stream/dv/tests/fub_tests/ -v  
pytest projects/components/stream/dv/tests/integration_tests/ -v  
  
# Lint (once RTL created)  
verilator --lint-only  
projects/components/stream/rtl/stream_fub/scheduler.sv  
  
# Search for modules  
find projects/components/stream/rtl/ -name "*.sv" -exec grep -H  
"^module" {} \;
```

23.12 Remember

1. **Tutorial focus** - Intentional simplifications for learning
 2. **Reuse from RAPIDS** - Descriptor engine, SRAM, patterns
 3. **Length in beats** - Not bytes or chunks!
 4. **Aligned addresses** - No fixup logic
 5. **Chained descriptors** - No circular buffers
 6. **8 channels** - Shared resources, arbitration required
 7. **MonBus standard** - Same format as AMBA/RAPIDS
 8. **Testbench reuse** - Always create TB classes in
bin/CocoTBFramework/tbclasses/stream/
-

23.13 PDF Generation Location

IMPORTANT: PDF files should be generated in the docs directory:

```
/mnt/data/github/rtldesignsherpa/projects/components/stream/docs/
```

Quick Command: Use the provided shell script:

```
cd /mnt/data/github/rtldesignsherpa/projects/components/stream/docs  
./generate_pdf.sh
```

The shell script will automatically: 1. Use the md_to_docx.py tool from bin/ 2. Process the stream_spec index file 3. Generate both DOCX and PDF files in the docs/ directory 4. Create table of contents and title page

 See: bin/md_to_docx.py for complete implementation details

Version: 1.0 **Last Updated:** 2025-10-17 **Maintained By:** RTL Design Sherpa Project