

# Contents

<b>Stream Index</b>	<b>1</b>
<b>STREAM Specification Index</b>	<b>1</b>
Document Organization . . . . .	2
Chapter 1: Overview . . . . .	2
Chapter 2: Functional Blocks . . . . .	2
Quick Reference . . . . .	2
Functional Unit Blocks (FUB) . . . . .	2
Integration Blocks (MAC) . . . . .	2
Performance Modes (AXI Engines) . . . . .	3
Holistic Overview: V1/V2/V3 Working Together . . . . .	3
V1 - Low Performance (Single Outstanding Transaction) . . . . .	3
V2 - Medium Performance (Command Pipelined) . . . . .	4
V3 - High Performance (Out-of-Order Completion) . . . . .	4
Performance Comparison Summary . . . . .	5
Clock and Reset Summary . . . . .	5
Clock Domains . . . . .	5
Reset Signals . . . . .	5
Interface Summary . . . . .	6
External Interfaces . . . . .	6
Internal Buses . . . . .	6
Area Estimates . . . . .	6
By Performance Mode . . . . .	6
Breakdown (Low Performance) . . . . .	6
Related Documentation . . . . .	7
Specification Conventions . . . . .	7
Signal Naming . . . . .	7
Parameter Naming . . . . .	7
State Machine Naming . . . . .	7

## Stream Index

**Generated:** 2025-11-16

## STREAM Specification Index

**Version:** 0.26 **Date:** 2025-10-17 **Purpose:** Complete technical specification for STREAM subsystem

---

## Document Organization

**Note:** All chapters linked below for automated document generation.

### Chapter 1: Overview

- Clocks and Reset

### Chapter 2: Functional Blocks

- Descriptor Engine
  - Scheduler
  - AXI Read Engine
  - AXI Write Engine
  - SRAM Controller
  - Simple SRAM
  - Channel Arbiter
  - APB Config
  - MonBus AXIL Group
  - Top-Level Integration
  - Performance Profiler
- 

## Quick Reference

### Functional Unit Blocks (FUB)

Module	File	Purpose	Lines	Status
descriptor_engine	stream_fub/descriptor_engine.sv	Descriptor fetch/parse (256-bit)	~300	[Pen]
scheduler	stream_fub/scheduler.sv	Transfer coordinator	~40	[Pen]
axi_read_engine	stream_fub/axi_read_engine.sv	AXI read master	~30	[Pen]
axi_write_engine	stream_fub/axi_write_engine.sv	AXI write master	~40	[Pen]
sram_controller	stream_fub/sram_controller.sv	Per-channel buffer management	~30	[Pen]
simple_sram	stream_fub/simple_sram.sv	Dual-port SRAM primitive	~10	[Pen]
perf_profiler	stream_fub/perf_profiler.sv	Channel performance profiling	~40	[Pen]

### Integration Blocks (MAC)

Module	File	Purpose	Lines	Status
channel_arbiter	stream_macro/channel_arbiter.sv	Priority arbitration	~200	[Pen]
apb_config	regs/stream_regs.rdl + wrapper	Config registers	~350	[Future]

Module	File	Purpose	Lines	Stat
monbus_axil_group	stream_macro/monbus_axil_group.sv	MonBus + AXIL	~800	[Don]
stream_top	stream_macro/stream_top.sv	Top-level	~500	[Pen]

---

## Performance Modes (AXI Engines)

STREAM AXI engines support three performance modes via compile-time parameters, offering area/performance trade-offs from tutorial simplicity to datacenter throughput.

### Holistic Overview: V1/V2/V3 Working Together

#### Design Philosophy:

- **V1 (Low):** Tutorial-focused, simple architecture, minimal area
- **V2 (Medium):** Command pipelining hides memory latency, 6.7x throughput improvement
- **V3 (High):** Out-of-order completion maximizes memory controller efficiency, 7.0x throughput

**Key Insight:** The benefit scales with memory latency. For low-latency SRAM (2-3 cycles), V1 achieves 40% throughput. For high-latency DDR4 (70-100 cycles), V1 drops to 14% but V2/V3 maintain 94-98% throughput.

#### Parameterization Strategy:

- `ENABLE_CMD_PIPELINE = 0`: V1 (default, tutorial mode)
- `ENABLE_CMD_PIPELINE = 1`: V2 (command pipelined, best area efficiency)
- `ENABLE_CMD_PIPELINE = 1, ENABLE_000_DRAIN = 1` (write) or `ENABLE_000_READ = 1` (read): V3 (out-of-order, maximum throughput)

### V1 - Low Performance (Single Outstanding Transaction)

**Architecture:** Single-burst-per-request, blocking on completion

- **Area:** ~1,250 LUTs per engine
- **Throughput:** 0.14 beats/cycle (DDR4), 0.40 beats/cycle (SRAM)
- **Outstanding Txns:** 1
- **Use Case:** Tutorial examples, embedded systems, low-latency SRAM
- **Key Feature:** Zero-bubble streaming pipeline (NO FSM!)

### **Blocking Behavior:**

- Write: Blocks on B response before accepting next request
- Read: Blocks on R last before accepting next request
- Simple control: 3 flags (`r_ar_inflight`, `r_ar_valid`, completion tracking)

## **V2 - Medium Performance (Command Pipelined)**

**Architecture:** Command queue decouples command acceptance from data completion

- **Area:** ~2,000 LUTs per engine (1.6x increase)
- **Throughput:** 0.94 beats/cycle (DDR4), 0.85 beats/cycle (SRAM)
- **Outstanding Txns:** 4-8 (command queue depth)
- **Use Case:** General-purpose FPGA, DDR3/DDR4, balanced area/performance
- **Key Feature:** Hides memory latency via pipelining (6.7x improvement)

### **Command Pipelining:**

- Write: AW queue + W drain FSM + B scoreboard (async response tracking)
- Read: AR queue + R in-order reception (simpler than write, no B channel)
- Per-command SRAM pointers enable multiple bursts from same channel

**Best Area Efficiency:** 4.19x throughput per unit area (6.7x throughput / 1.6x area)

## **V3 - High Performance (Out-of-Order Completion)**

**Architecture:** OOO command selection maximizes memory controller efficiency

- **Area:** ~3,500-4,000 LUTs per engine (2.8-3.2x increase)
- **Throughput:** 0.98 beats/cycle (DDR4), 0.92 beats/cycle (SRAM)
- **Outstanding Txns:** 8-16 (larger command queue)
- **Use Case:** Datacenter, ASIC, HBM2, high-performance memory controllers
- **Key Feature:** Flexible drain order based on SRAM data availability (7.0x improvement)

### **Out-of-Order Mechanisms:**

- Write: OOO W drain (select command with SRAM data ready), AXI ID matching for B responses
- Read: OOO R reception (match m\_axi\_rid to queue entry), independent SRAM write per command
- Transaction ID structure: {counter, channel\_id} preserves channel routing for MonBus

### Why OOO is Naturally Supported:

1. Per-channel SRAM partitioning (no cross-channel hazards)
2. Per-command pointer tracking (no pointer collision)
3. Transaction ID matching (channel ID in lower bits for routing)

### Performance Comparison Summary

Configuration	Area	DDR4 Throughput	SRAM Throughput	Area Efficiency	Usage
<b>V1</b>	1.0x	0.14 beats/cycle (1.0x)	0.40 beats/cycle (1.0x)	1.00	Tutorial
<b>V2</b>	1.6x	0.94 beats/cycle (6.7x)	0.85 beats/cycle (2.1x)	4.19	General
<b>V3</b>	2.8x	0.98 beats/cycle (7.0x)	0.92 beats/cycle (2.3x)	2.50	Data

**Key Takeaway:** V2 offers best area efficiency (4.19x) for typical FPGA use cases. V3 provides marginal throughput improvement (7.0x vs 6.7x) at higher area cost, justified only for high-performance memory controllers that support out-of-order responses.

---

### Clock and Reset Summary

#### Clock Domains

Clock	Frequency	Usage
aclk	100-500 MHz	Primary - all STREAM logic, AXI/AXIL interfaces
pclk	50-200 MHz	APB configuration interface (may be async to aclk)

#### Reset Signals

Reset	Polarity	Type	Usage
aresetn	Active-low	Async assert, sync deassert	Primary - all STREAM logic
presetn	Active-low	Async assert, sync deassert	APB configuration interface

**See:** Clocks and Reset for complete timing specifications

---

## Interface Summary

### External Interfaces

Interface	Type	Width	Purpose
APB	Slave	32-bit	Configuration registers
AXI (Descriptor)	Master	256-bit	Descriptor fetch
AXI (Read)	Master	512-bit (param)	Source data read
AXI (Write)	Master	512-bit (param)	Destination data write
AXIL (Slave)	Slave	32-bit	Error/interrupt FIFO access
AXIL (Master)	Master	32-bit	MonBus packet logging to memory
IRQ	Output	1-bit	Error interrupt

### Internal Buses

Interface	Width	Purpose
MonBus	64-bit	Internal monitoring bus (channels -> monbus_axil_group)

---

## Area Estimates

### By Performance Mode

Configuration	Total LUTs	SRAM	Use Case
Low (Tutorial)	~9,500	64 KB	Educational, area-constrained
Medium (Typical)	~11,200	64 KB	Balanced FPGA implementations
High (Performance)	~13,700	64 KB	High-throughput ASIC/FPGA

### Breakdown (Low Performance)

Component	Instances	Area/Instance	Total
Descriptor Engine	8	~300 LUTs	~2,400 LUTs
Scheduler	8	~400 LUTs	~3,200 LUTs
AXI Read Engine (Low)	1	~250 LUTs	~250 LUTs
AXI Write Engine (Low)	1	~250 LUTs	~250 LUTs
SRAM Controller	1	~1,600 LUTs	~1,600 LUTs

Component	Instances	Area/Instance	Total
Simple SRAM (internal)	1-8	1024x64B total	64 KB
Channel Arbiter	3	~150 LUTs	~450 LUTs
APB Config	1	~350 LUTs	~350 LUTs
MonBus AXIL Group	1	~1,000 LUTs	~1,000 LUTs
<b>Total</b>	-	-	<b>~9,500 LUTs + 64KB</b>

---

## Related Documentation

- **PRD.md** - Product requirements and overview
  - **ARCHITECTURAL NOTES.md** - Critical design decisions
  - **CLAUDE.md** - AI development guide
  - **Register Generation** - PeakRDL workflow
- 

## Specification Conventions

### Signal Naming

- **Clock:** aclk, pclk
- **Reset:** aresetn, presetn (active-low)
- **Valid/Ready:** Standard AXI/custom handshake
- **Registers:** r\_ prefix (e.g., r\_state, r\_counter)
- **Wires:** w\_ prefix (e.g., w\_next\_state, w\_grant)

### Parameter Naming

- **Uppercase:** NUM\_CHANNELS, DATA\_WIDTH, ADDR\_WIDTH
- **String parameters:** PERFORMANCE ("LOW", "MEDIUM", "HIGH")

### State Machine Naming

```
typedef enum logic [3:0] {
    IDLE    = 4'h0,
    ACTIVE  = 4'h1,
    // ...
} state_t;

state_t r_state, w_next_state; // Current and next state
```

---

**Last Updated:** 2025-10-17 **Maintained By:** STREAM Architecture Team