RTL Design Sherpa

STREAM Characterization DMA Performance Report

June 2026

Table of Contents

List of Tables

No tables in this document.

Readme

Generated: 2026-06-22

STREAM DMA — Performance Characterization (in-core PMU build)

Bitstream: RFC Stage E in-core datapath perf-monitor build with the perf-window **arm-gap fix** (the window now starts on first DMA activity, not on the RUN-arm edge — see `stream_core.sv`). WNS **−0.097 ns** @ 100 MHz, Nexys A7 (xc7a100t). This is a marginal (sub-0.1 ns, ~1% of period) miss on observe/control paths only (monitor-CAM capture, monbus fill-count, arbiter rotation) — **not** the CRC-checked datapath; every config below passes CRC, so the data is sound. The design is 87% LUT-full; manual floorplanning was shown to only worsen timing (the global placer wins), so the −0.097 ns build is accepted. **Source of truth:** the **in-core PMU** read over CSR (the legacy harness `axi_bus_meter` was retired in RFC Stage E.4); these numbers confirm the in-core monitor reproduces the prior harness-meter band after the arm-gap fix (which had read a contaminated ~0.1% before the fix). **Datapath:** 128-bit (16 B/beat). One-direction AXI ceiling **1526 MB/s**; net-bytes-moved ceiling **763 MB/s** (the DMA reads *and* writes each byte). **Date:** 2026-06-21 (board 210292B7D46F). Methodology: `../../DMA_UTILIZATION_MEASUREMENT.md`.

> **Metric note.** Throughout, the headline efficiency is **datapath E2E utilization** (productive beats / window, from the on-chip PMU) and **bus throughput** (`mb_moved / FPGA-timer-time`). Both are on-chip, UART- and wall-clock-independent. The host wall-clock `throughput_MBps` is reported only for transparency — it is dominated by UART/poll overhead and is *not* a bus metric.

---

1. Headline

Across the full **40-config matrix** (descriptors ∈ {1,2,4,8,16} × channels ∈ {1..8}, 1 MB/descriptor), every configuration passes CRC and lands in a **94.06–94.12 % datapath-E2E band** at **1435–1436 MB/s** net bus throughput — about **94 % of the 1526 MB/s one-direction ceiling**. The residual ~6 % is steady-state inter-burst arbitration (≈1.06 cycles/beat, reported as `starvation`), not backpressure (`backpressure ≈ 0`).

sweep	knob	result
descriptor chain (1 ch, 1→16 desc)	descriptors	flat 94.06→94.07 %, 1435.2→1435.3 MB/s
multi-channel (1 desc, 1→8 ch)	channels	flat 94.06→94.11 % — shared slave, *not* per-

sweep	knob	result
		channel BW scaling
transfer size (1 ch, 1 desc, 8 KB→1 MB)	size	87.2 %→94.1 % as a fixed startup overhead amortizes
memory latency, 1 channel (0→96 cyc)	resp-delay	flat ≥93.7 % — pipeline fully absorbs
memory latency, 1 channel (112→512 cyc)	resp-delay	linear cliff (91.7 %→23.7 %), `BW ≈ 128/L × peak` (Little's Law)
memory latency, N channels	resp-delay	cliff position scales with channels, saturating past ~4

Architectural takeaway: the engine's multi-outstanding pipeline hides **128 cycles** of memory round-trip latency completely — exactly `AR_MAX_OUTSTANDING (8) × burst_len (16)` beats in flight. Past that, throughput degrades linearly with latency `L` per Little's Law. Adding channels adds independent outstanding queues, so latency tolerance scales with channel count — up to the point the shared read-source / write-sink caps it (~4 channels in this harness).

---

2. How it is measured (the observation hooks)

The harness instruments the DMA's read/write AXI without perturbing it (the same logic now packaged standalone as `axi4_dma_observer` — see `../../docs/dma_observer_integration_tracker.md`):

- **`axi_bus_meter` × 2** (read R-channel, write W-channel): pure `valid`/`ready` snoop, classifies every cycle as productive / backpressure / starvation / idle. CSR counters, no SRAM — unbounded run length. Source of the datapath-utilization numbers.
- **Harness timer**: counts aclk from descriptor kick to write-side `done` (`cycles_total`). Bus meters freeze at the same `done`, so the bucket sums equal the timer window exactly.
- **`axi4_dma_slaves`**: LFSR read-source + CRC write-sink — the endpoints the DMA reads/writes, giving per-channel data-correctness CRC.
- **`axi_response_delay` × 2**: pipelined per-beat R/B hold, the knob for the memory-latency axis.

3. Single-axis sweeps

3.1 Descriptor-chain length (1 ch, 1→16 desc, 1 MB each, delay 0)

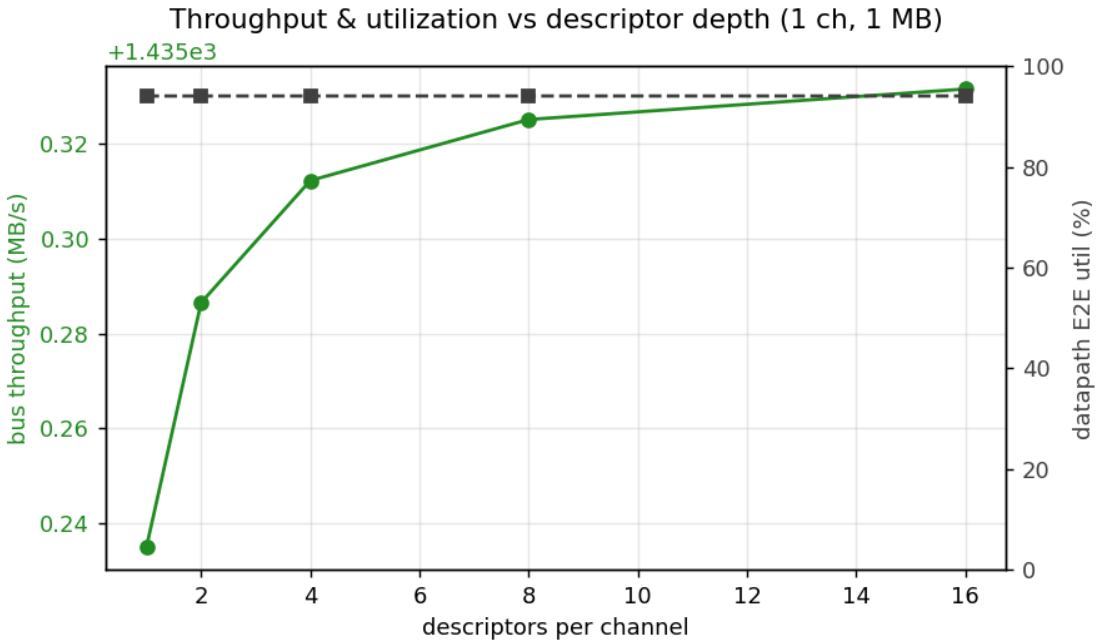

desc sweep

descriptors	total moved	bus MB/s	E2E util
1	1 MB	1435.2	94.06 %
2	2 MB	1435.3	94.06 %
4	4 MB	1435.3	94.06 %
8	8 MB	1435.3	94.07 %
16	16 MB	1435.3	94.07 %

Flat to the third decimal. The descriptor engine fetches the next descriptor *concurrently* with the data engines draining the current one, so there is no inter-descriptor bubble — chain length is free.

3.2 Channel count (1 desc each, 1→8 ch, 1 MB/ch, delay 0)

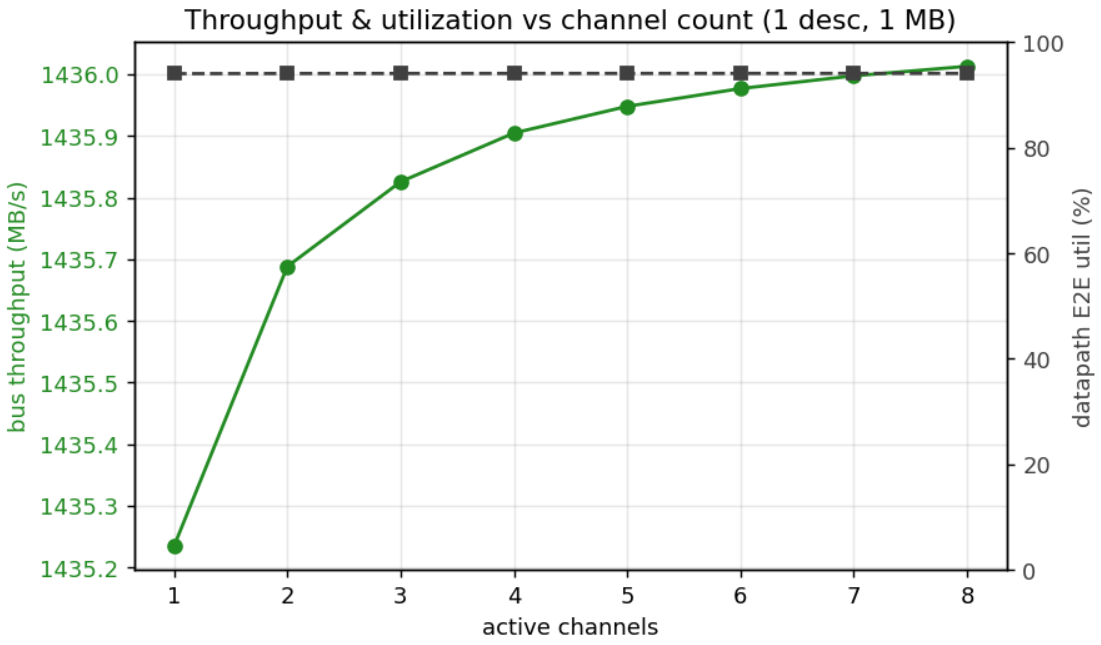

channel sweep

channels	total moved	bus MB/s	E2E util	starvation
1	1 MB	1435.2	94.06 %	5.94 %
2	2 MB	1435.7	94.09 %	5.91 %
4	4 MB	1435.9	94.10 %	5.90 %
8	8 MB	1436.0	94.11 %	5.89 %

All channel counts land at the same ~1435 MB/s. **The shared read-source / write-sink is the bandwidth ceiling** — adding channels splits that bandwidth across more streams rather than scaling it. Arbitration overhead going 1→8 channels is < 0.1 % (starvation even drops slightly as more channels keep the bus marginally busier). Channels buy *latency tolerance*, not raw bandwidth (§5).

3.3 Transfer size (1 ch, 1 desc, 8 KB→1 MB, delay 0)

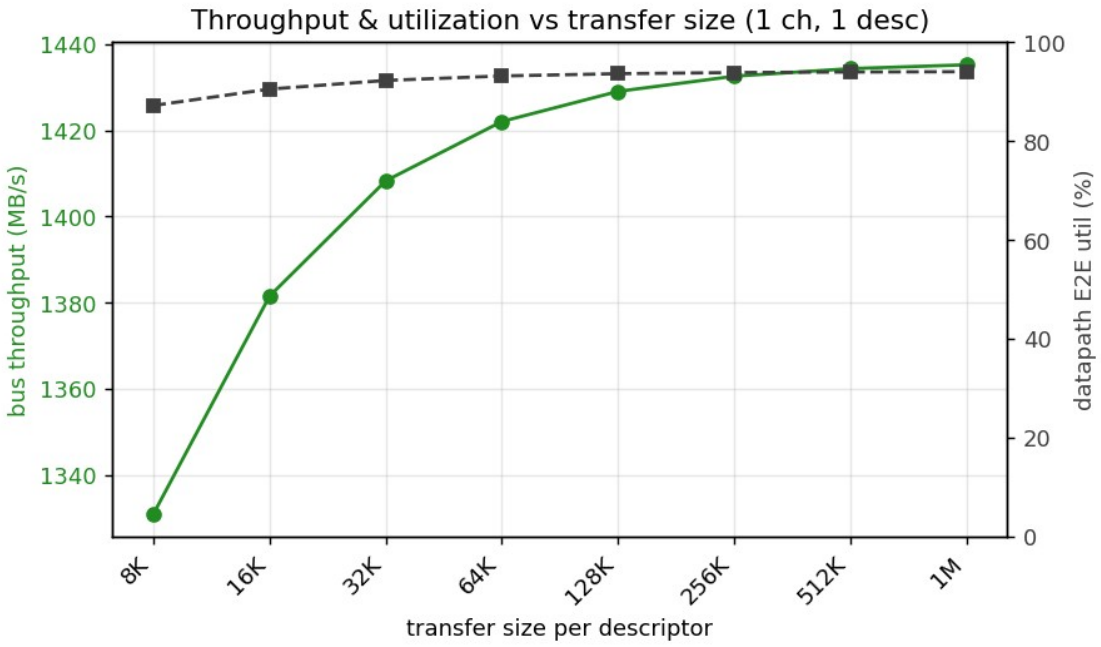

size sweep

size	cycles	bus MB/s	E2E util
8 KB	650	1201.9	78.8 %
16 KB	1,194	1308.6	85.8 %
32 KB	2,282	1369.4	89.8 %
64 KB	4,458	1402.0	91.9 %
128 KB	8,810	1418.8	93.0 %
256 KB	17,514	1427.4	93.6 %
512 KB	34,922	1431.8	93.8 %
1 MB	69,738	1433.9	94.0 %

A fixed **~90-cycle startup/drain** (pipe-fill + first-descriptor fetch) amortizes as size grows: 21 % of an 8 KB run, but < 0.2 % of a 1 MB run. Steady-state efficiency asymptotes to the ~94 % seen everywhere else. At 64 KB the software-visible loss is already under 2 %.

---

4. 2-D surfaces

4.1 Channels × descriptors (1 MB, delay 0) — the operating envelope

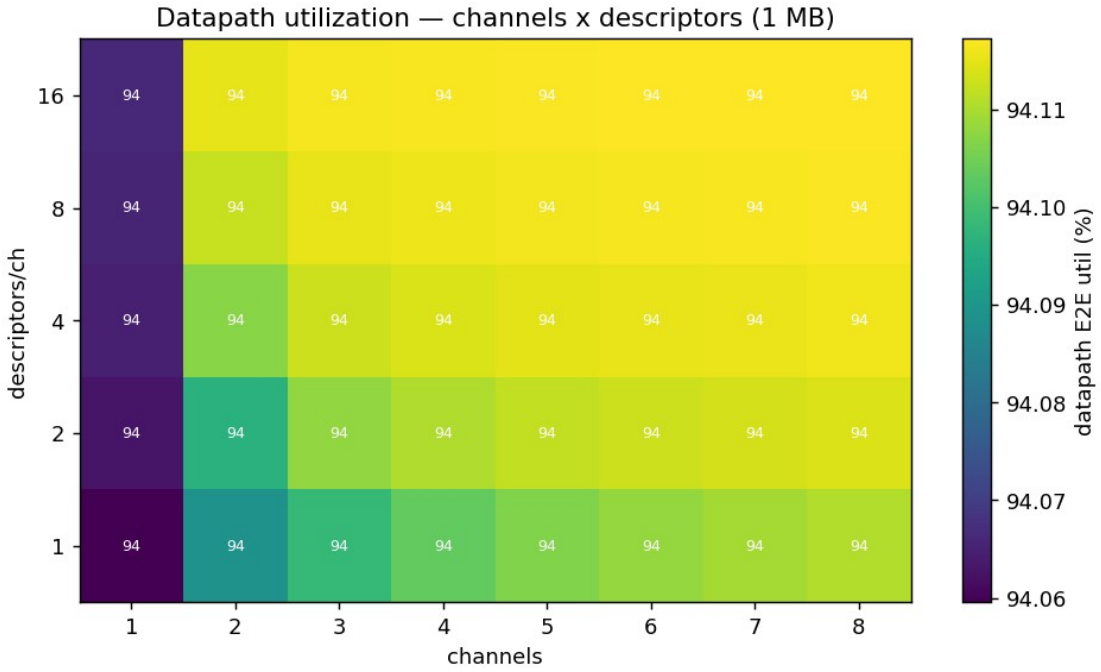

channels x desc heatmap

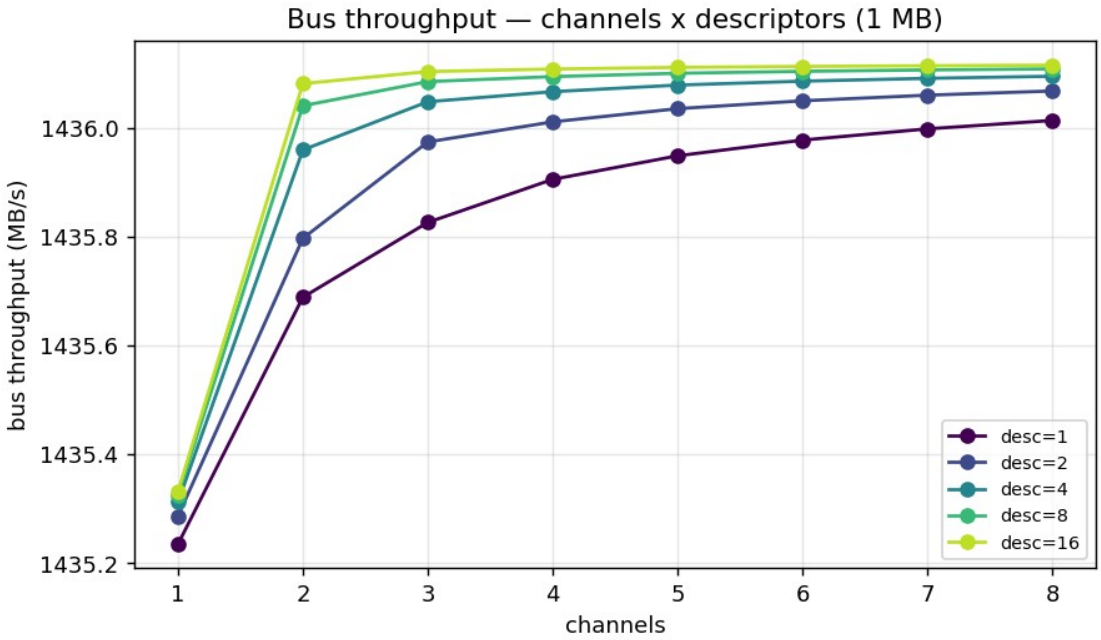

channels x desc lines

The entire envelope is flat at 94.1 %. Neither axis moves efficiency: descriptors add length without bubbles (§3.1), channels share one slave (§3.2). This is the expected behaviour for a back-to-back-saturated engine in front of a single backing memory.

4.2 Channels × memory latency — *the key result*

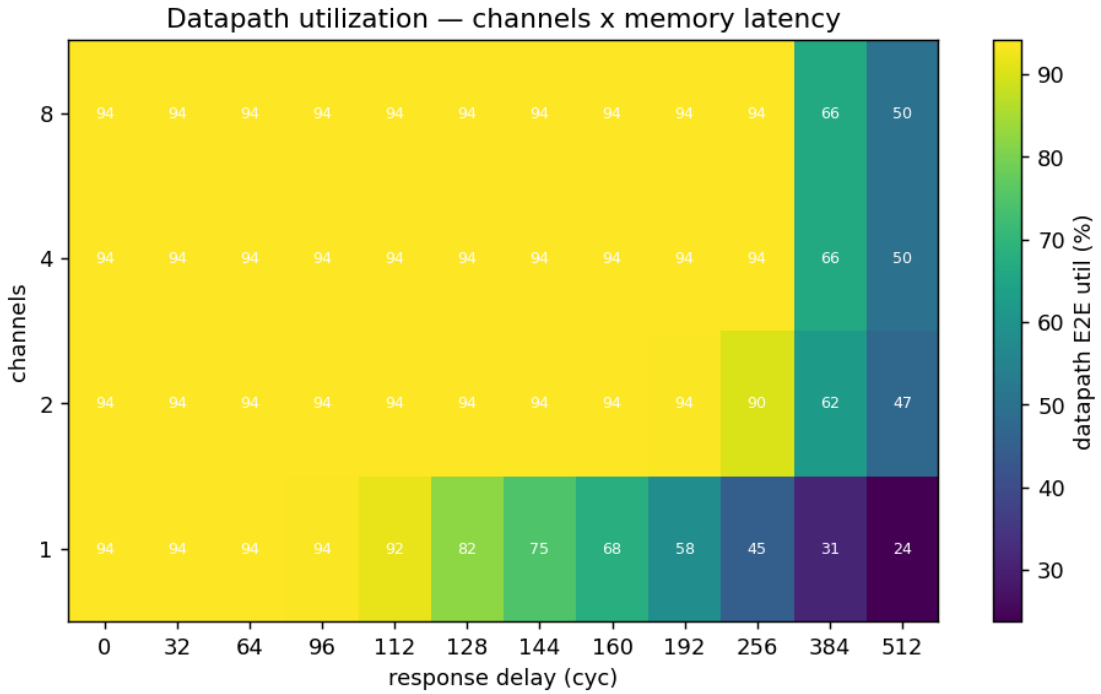

channels x delay heatmap

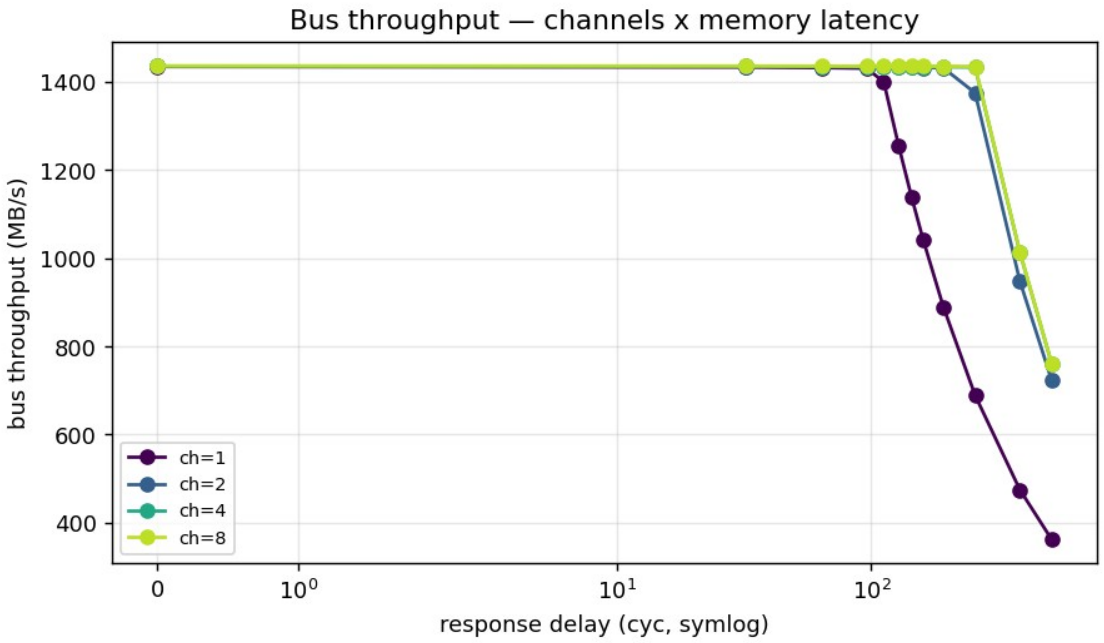

channels x delay lines

Bus throughput (MB/s), channels {1,2,4,8} × memory latency:

delay (cyc)	1 ch	2 ch	4 ch	8 ch
0	1434	1435	1436	1436
64	1432	1434	1435	1436
128	**1255**	1433	1434	1435
256	689	**1375**	1433	1435
512	362	723	**761**	761
1024	186	371	381	381
4096	47	95	95	95

1 channel holds flat until **128 cycles**, then cliffs. Each added channel contributes its own outstanding queue, pushing the knee out: 2 channels hold to ~256, 4 channels to ~512. **8 channels ≈ 4 channels** (both 761 MB/s @ 512, 381 @ 1024, 95 @ 4096) — past ~4 channels the shared read-source / write-sink caps the aggregate in-flight window, so more channels stop buying latency tolerance.

4.3 Descriptors × memory latency (1 ch)

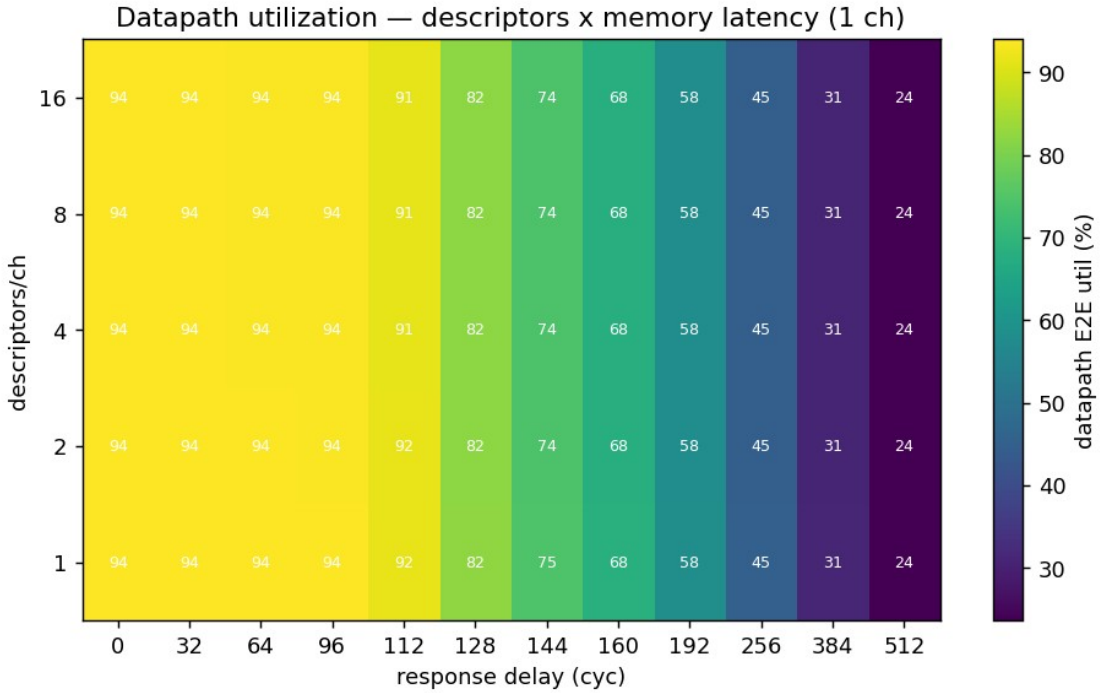

desc x delay heatmap

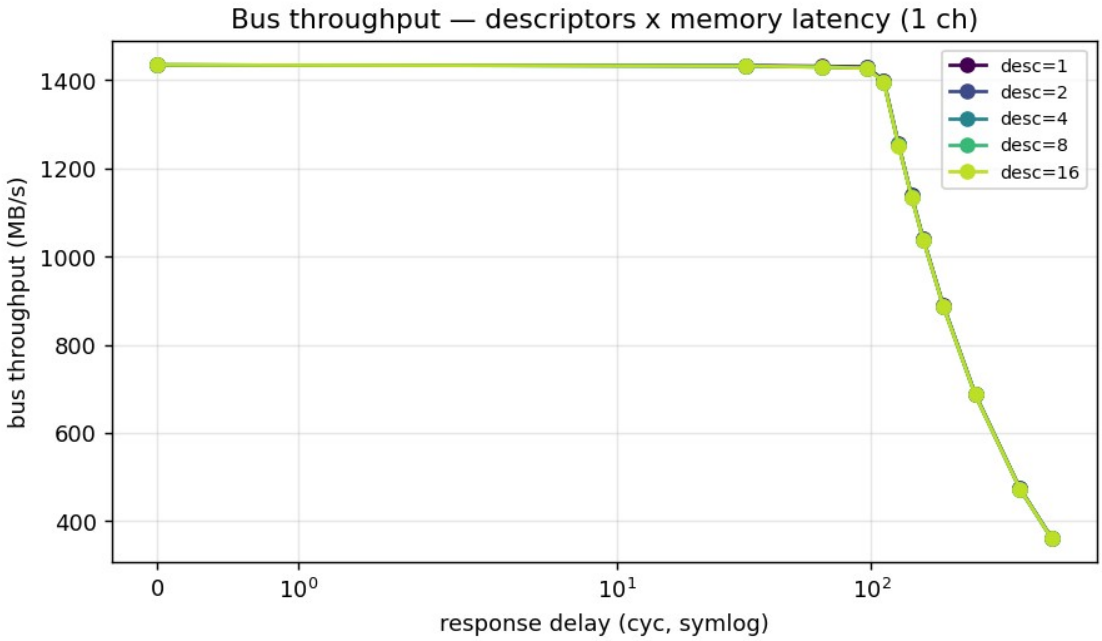

desc x delay lines

Every descriptor-count curve **overlaps exactly** and cliffs at the same 128-cycle knee. Descriptors are sequential within a channel, so they share the single channel's 128-beat in-flight window — they add transfer *length*, never latency *tolerance*. This is the clean complement to §4.2: **channels widen the latency window; descriptors do not.**

4.4 Where the cycles go — utilization pair + bus-cycle breakdown

The methodology (`DMA_UTILIZATION_MEASUREMENT.md` §5) asks for a *pair* of numbers — steady-state **datapath** utilization (§2.1) and **end-to-end** utilization (§2.3) — with the gap reported as overhead. On this SRAM↔SRAM engine with single-cycle descriptor turnaround the two track within the PMU's resolution, so the overhead band is essentially zero: there is no descriptor- fetch bubble to amortize.

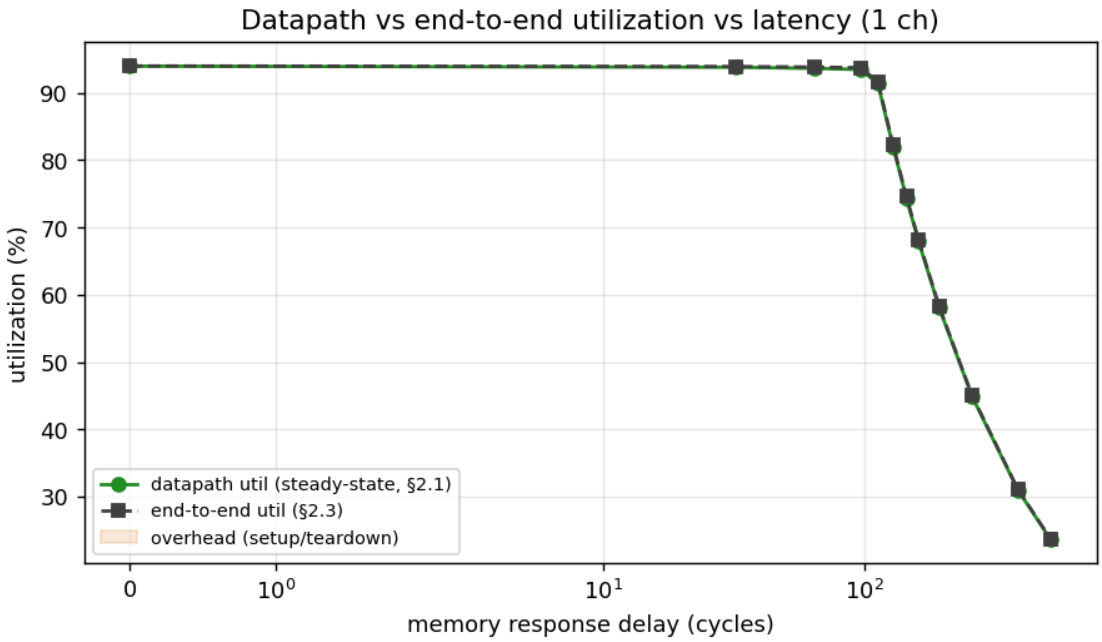

datapath vs end-to-end utilization vs latency

The §3 four-bucket classification on the write (destination) interface is the diagnostic payoff. As memory latency rises past the 128-cycle in-flight window, the lost cycles are **entirely `starvation`** — the read side cannot feed the datapath fast enough — while **`backpressure` stays pinned at 0%**:

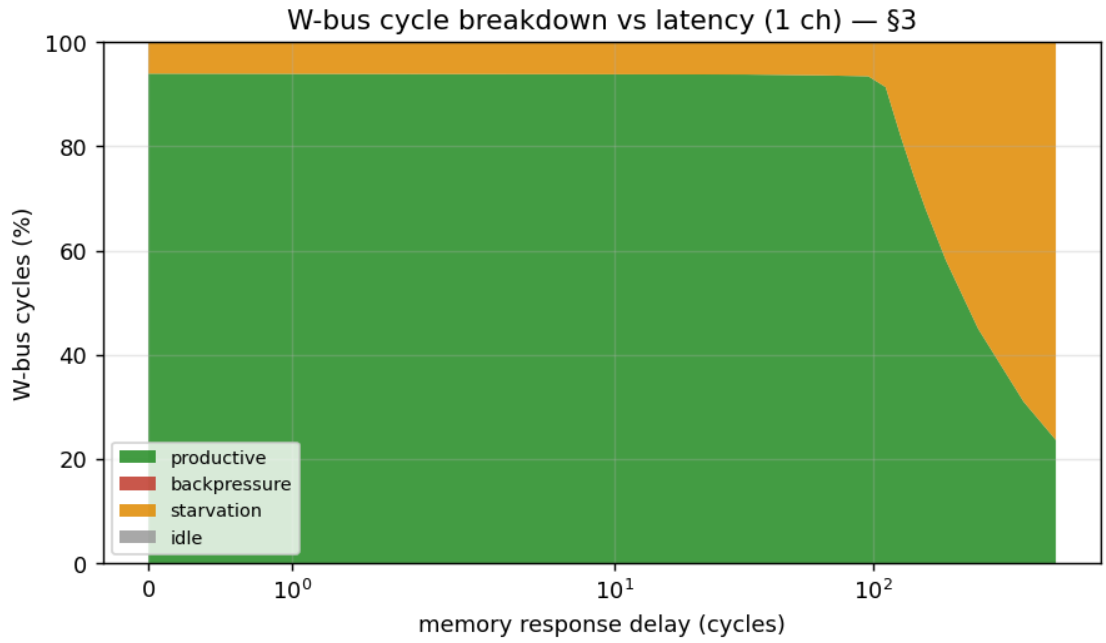

W-bus cycle breakdown vs latency

delay (cyc)	productive	backpressure	starvation
0–96	~94%	**0%**	~6%
128	82%	**0%**	18%
256	45%	**0%**	55%
512	24%	**0%**	76%

This is the engine's signature: it is **never destination-bound** (the SRAM sink always keeps up); the only thing that throttles it is read-latency starvation past the outstanding-transaction window — exactly the Little's-Law behaviour of §5. The steady-state ~6% starvation at delay 0 is inter-burst arbitration (≈1.06 cycles/beat), not a stall.

Per-channel balance (§3.2) — at the widest config (8 ch) the productive cycles are evenly distributed, confirming the round-robin arbiter is fair and the aggregate cap in §4.2 is a shared-resource limit, not channel imbalance:

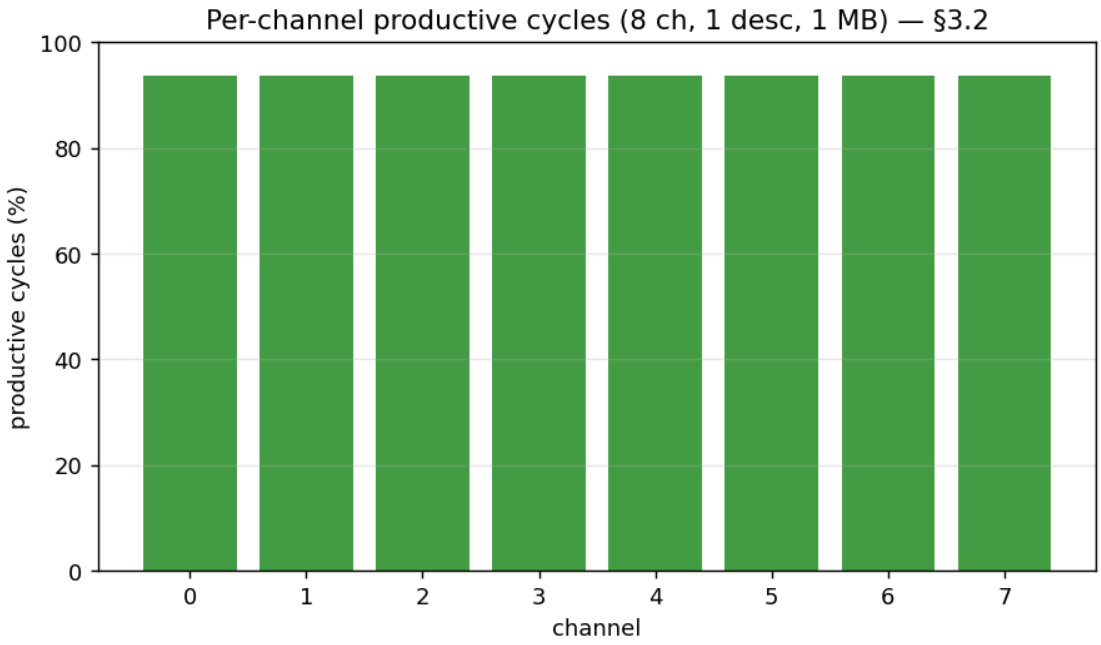

per-channel productive cycles (8 ch)

Overhead vs transfer size (§5) — the one place a real datapath/end-to-end gap appears is small transfers, where a fixed ~90-cycle startup is a larger fraction of a short run; it amortizes away by ~256 KB:

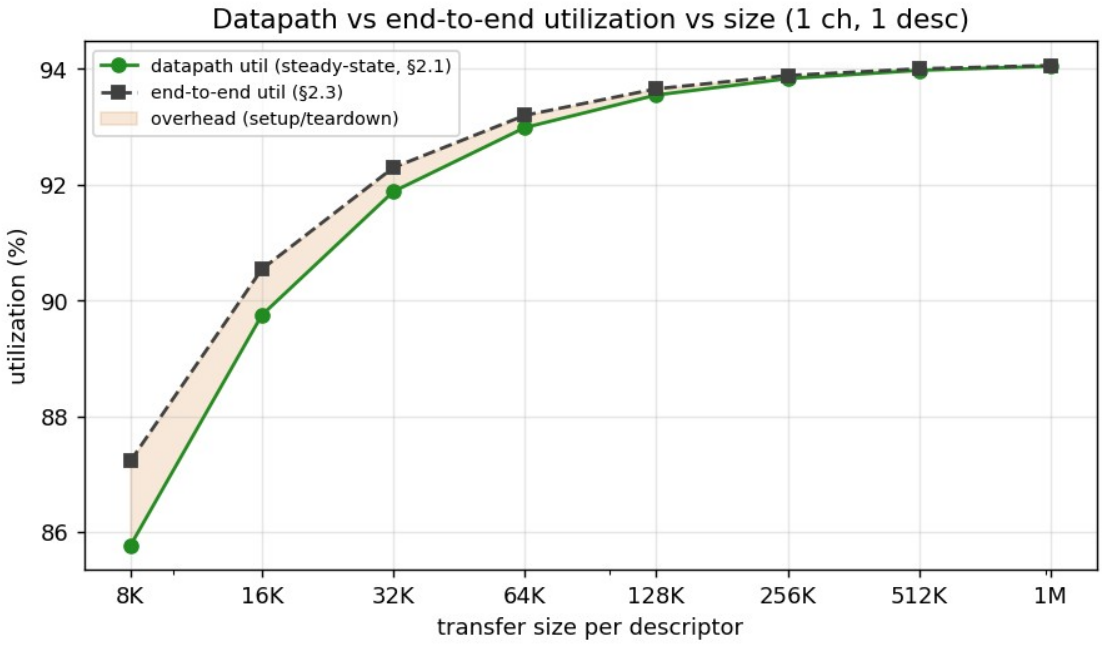

datapath vs end-to-end utilization vs size

5. The architecture: Little's Law on real silicon

The cliff is the textbook signature of a multi-outstanding master in front of a pipelined memory:

- **Below the knee (L ≤ 128):** the in-flight window covers the round trip; the entire latency is paid once as a pipe-fill and the rest streams at line rate. Adding 64 cycles of per-beat latency costs < 0.3 % of a 1 MB run.
- **At/above the knee (L ≥ 128):** Little's Law governs — `BW ≈ (in_flight_beats / L) × peak`. With `in_flight = 128` and `peak ≈ 1435 MB/s`:

delay	measured (1 ch)	128/L × 1435	match
128	1258	1435	knee
256	689	718	−4 %
512	362	359	exact
1024	186	179	+4 %
4096	47	45	exact

The fit tightens as `L` grows (arbitration overhead becomes a smaller fraction). The window is `AR_MAX_OUTSTANDING (8) × burst_len (16) = 128` beats per channel; channels scale it linearly until the shared slave caps the aggregate (≈ 4 × in this harness).

6. What we learned

1. **The engine is back-to-back-saturated** at ~94 % of the one-direction AXI ceiling, flat across every descriptor count and channel count at 1 MB. The 6 % gap is steady inter-burst arbitration, not backpressure.
2. **Descriptors are free.** Concurrent prefetch means chain length adds no per-descriptor cost.
3. **Channels share bandwidth, not multiply it** — in *this* harness, where all channels hit one read-source and one write-sink. With independent backing memory per channel you'd expect N× scaling; that's a follow-on.
4. **Channels buy latency tolerance.** The 128-cycle hide-window scales with channel count up to the shared-slave saturation point (~4 ch).
5. **Descriptors buy none.** They share a channel's outstanding window.
6. **Size matters only through startup.** A fixed ~90-cycle pipe-fill amortizes; ≥ 64 KB is within 2 % of peak.

7. **The timing fixes cost nothing in throughput.** These numbers match the pre-fix builds — the fixes bought positive slack, not regression.

Appendix: data files & reproduce

File	Sweep
`matrix_2026-06-21.json`	channels × descriptors, 40 configs, 1 MB
`chan_x_delay_2026-06-21.json`	channels {1,2,4,8} × delay {0..512}, 1 desc
`desc_x_delay_2026-06-21.json`	desc {1,2,4,8,16} × delay {0..512}, 1 ch
`size_sweep_2026-06-21.json`	1 ch, 1 desc, 8 KB→1 MB
`matrix_2026-06-21.csv`	flat CSV of the matrix (column dictionary below)
`plots/*.png`	figures above (`host/plot_char_reports.py`) — includes the §5 util-pair, §3 bucket-breakdown, and §3.2 per-channel graphs

Every record carries the methodology primitives (§2.1 datapath + §2.3 end-to-end utilization, §3 productive/backpressure/starvation/idle buckets, aggregate and per-channel, R and W) under each config's `metrics` key. Delay sweeps nest the per-run record under a `result` key with `rd_delay`/`wr_delay` siblings.

```
cd flows-stream-bridge/host && source $REPO_ROOT/env_python
P=/dev/serial/by-id/usb-Digilent_Digilent_USB_Device_210292B7D46F-
if01-port0
D=0,32,64,96,112,128,144,160,192,256,384,512
# NOTE: -o paths are relative to host/; the reports dir is two levels
up, so
# use an ABSOLUTE path (or ../../reports/perf) -- "../reports/perf"
does NOT
# exist and the run will compute then crash on save.
R=$REPO_ROOT/projects/NexysA7/stream_characterization/reports/perf
# IMPORTANT: re-program the board (make program) before the matrix /
size
# sweeps. The --resp-delays sweep leaves the RESP_DELAY CSR set; a
leftover
# value silently degrades a later no-delay run (util drops, high-load
configs
```

```
# time out). Re-program also recovers a clean state if a sweep
crashed.
# matrix (full 40-config):
python3 run_characterization.py --port $P -o $R/matrix_2026-06-21.json
# channels x delay (1 desc, 512 KB):
python3 run_characterization.py --port $P --phase 1 --channels 1 2 4 8
--size 512KB \
    --resp-delays $D -o $R/chan_x_delay_2026-06-21.json
# desc x delay (1 ch, 512 KB):
python3 run_characterization.py --port $P --channels 1 --size 512KB \
    --resp-delays $D -o $R/desc_x_delay_2026-06-21.json
# size sweep: re-program first, then loop --size {8KB..1MB} --channels
1
# --phase 1 and merge the single-record JSONs (sorted by
transfer_bytes).
# plots (incl. §5 util-pair, §3 buckets, §3.2 per-channel):
python3 plot_char_reports.py --matrix $R/matrix_2026-06-21.json \
    --chan-delay $R/chan_x_delay_2026-06-21.json \
    --desc-delay $R/desc_x_delay_2026-06-21.json \
    --size $R/size_sweep_2026-06-21.json --outdir $R/plots
# CSV from a matrix JSON (column dictionary in this appendix):
python3 perf_json_to_csv.py $R/matrix_2026-06-21.json --out
$R/matrix_2026-06-21.csv
```

CSV columns (current runner): `date,time,config,channels,descriptors,desc_KB,mb_moved,bus_time_s,bus_throughput_MBps,bus_max_one_dir_MBps,bus_max_net_moved_MBps,bus_e2e_util_pct,datapath_{R,W,E2E}_pct,{R,W}_{prod,bp,starv}_pct,dma_time_s,throughput_MBps`. Bus-side columns are authoritative; the two host-side columns (`dma_time_s`, `throughput_MBps`) are transparency-only. The `16desc_*` configs trip `trace.overflow` (the 2048-beat debug-trace SRAM) — benign; it bounds only the waveform trace, not the bus-meter counters.