



RTL Design Sherpa

STREAM DMA Engine Phase 1 Characterization Report

April 30, 2026

Table of Contents

1 STREAM DMA Engine — Phase 1 Characterization Report.....	7
1.1 0. Resource footprint.....	7
1.1.1 0.1 Top-level — stream_char_top.....	7
1.1.2 0.2 Inside the DUT — stream_top_ch8.....	10
1.1.3 0.3 What this means.....	15
1.1.4 0.4 Timing-closure notes (8-channel build).....	16
1.2 1. What we set out to do.....	17
1.3 2. Test platform.....	18
1.4 3. The harness — and why we hooked things in the way we did.....	19
1.4.1 3.1 Harness overview.....	19
1.4.2 Figure 3.1: Harness top-level block diagram.....	19
1.4.3 3.2 Hook 1: The pipelined response-delay block (memory model).....	20
1.4.4 Figure 3.2: Read- and write-path data flow.....	20
1.4.5 Figure 3.3: axi_response_delay internals.....	21
1.4.6 3.3 Hook 2: The hardware kick-burst path.....	22
1.4.7 Figure 3.4: CSR-driven kick-burst path.....	22
1.4.8 3.4 Hook 3: The cycle-stamp timer windows.....	23
1.4.9 Figure 3.5: Timer windows on the test timeline.....	23
1.4.10 3.5 Hook 4: Per-channel CRC slaves.....	24
1.4.11 3.6 Hook 5: Configuration package for sweeps.....	24
1.4.12 3.7 What all of this gives us, together.....	25
1.5 4. Headline numbers.....	25

1.6 5. Single-axis sweeps.....	27
1.6.1 5.1 Response-delay sweep (1 ch, 1 desc, 512 KB).....	27
1.6.2 Figure 5.1: Throughput vs. memory latency (1 channel).....	27
1.6.3 5.2 Descriptor-count sweep (1 ch, 1..16 descriptors, 512 KB each, delay 0).....	29
1.6.4 Figure 5.2: Throughput vs. descriptor chain length.....	29
1.6.5 5.3 Channel-count sweep (1..4 ch, 1 desc each, 512 KB per channel, delay 0).....	30
1.6.6 Figure 5.3: Throughput vs. active channel count.....	30
1.6.7 5.4 Transfer-size sweep (1 ch, 1 desc, 8 KB → 1 MB, delay 0).....	31
1.6.8 Figure 5.4: Software-visible throughput vs. transfer size.....	31
1.7 6. 2-D crosses.....	32
1.7.1 6.1 Channels × delay (1 desc, 512 KB) — <i>the key result</i>	32
1.7.2 Figure 6.1: Throughput surface — channels × memory latency.....	32
1.7.3 Figure 6.2: Per-channel-count slices.....	33
1.7.4 6.2 Descriptors × delay (1 ch, 512 KB).....	34
1.7.5 Figure 6.3: Throughput surface — descriptors × memory latency.....	34
1.7.6 Figure 6.4: Per-descriptor-count slices.....	35
1.7.7 6.3 Channels × descriptors (delay 0).....	36
1.7.8 Figure 6.5: Throughput surface — channels × descriptors at zero latency.....	36
1.7.9 Figure 6.6: Per-descriptor slices at zero latency.....	37
1.8 7. The architectural picture: Little’s Law on real silicon.....	38
1.8.1 Figure 7.1: Pre-cliff, knee, post-cliff regimes.....	38
1.9 8. What we learned.....	38

1.10 9. Setup for Phase 2: Vivado IP DMA comparison.....	39
--	----

List of Figures

Figure 3.1: Harness top-level block diagram.....	19
Figure 3.2: Read- and write-path data flow.....	20
Figure 3.3: axi_response_delay internals.....	21
Figure 3.4: CSR-driven kick-burst path.....	22
Figure 3.5: Timer windows on the test timeline.....	23
Figure 5.1: Throughput vs. memory latency (1 channel).....	27
Figure 5.2: Throughput vs. descriptor chain length.....	29
Figure 5.3: Throughput vs. active channel count.....	30
Figure 5.4: Software-visible throughput vs. transfer size.....	31
Figure 6.1: Throughput surface — channels × memory latency.....	32
Figure 6.2: Per-channel-count slices.....	33
Figure 6.3: Throughput surface — descriptors × memory latency.....	34
Figure 6.4: Per-descriptor-count slices.....	35
Figure 6.5: Throughput surface — channels × descriptors at zero latency.....	36
Figure 6.6: Per-descriptor slices at zero latency.....	37
Figure 7.1: Pre-cliff, knee, post-cliff regimes.....	38

List of Tables

No tables in this document.

1 STREAM DMA Engine — Phase 1 Characterization Report

Component: STREAM DMA (stream_top_ch8) **Platform:** Digilent Nexys A7-100T (Xilinx Artix-7 100T, 100 MHz, **8-channel build**, -1 speed grade) **Status:** Phase 1 (in-house DMA characterization) complete; Phase 2 will swap in a Vivado IP DMA for head-to-head PPA comparison.

1.1 0. Resource footprint

Before the dynamic numbers, the static cost. The numbers below are post-route on xc7a100tcs324-1, NUM_CHANNELS = 8, all monitors enabled, harness included. WNS = +0.296 ns, 0 failing endpoints, 0 hold violations.

1.1.1 0.1 Top-level — stream_char_top

The DUT (u_stream = stream_top_ch8) accounts for ~70 % of LUTs / ~72 % of FFs. Everything else is harness: traffic generator + CRC checker on each AXI side, a programmable response-delay block per side, descriptor RAM, monitor trace buffer, UART bridge, datapath utilization meters, and the host AXIL fanout. None of the harness blocks scale with NUM_CHANNELS; only the DUT does.

Instance	Module	LUTs	FFs	BRAM	DSP	Role
u_stream	stream_top_ch8	20 370	20 568	8.0	1	DMA + monitors (see §0.2)
u_bridge	bridge_stream	1 899	1 777	0	0	host AXIL

Instance	Module	LUTs	FFs	BRAM	DSP	Role
	m_char_axil					xbar fanout (CSR / desc-ram / STREA M APB / debug-SRAM / status)
u_csr	harnes s_csr	1 427	613	0	0	harness register s (kick, timer, resp-delay knobs, status mirror)
u_rd_p attern	axi4_s lave_r d_patt ern_ge n	1 419	1 222	0	0	R-side LFSR pattern generat or + per- channe l CRC slave
u_wr_c rc_chec k	axi4_s lave_w r_crc_ check	1 331	943	0	0	W-side per- channe l CRC checker slave
u_desc_ ram	desc_r am	814	574	0	0	descrip tor

Instance	Module	LUTs	FFs	BRAM	DSP	Role
						source memor y (LUTRA M, 128 × 256 b)
u_uart	uart_a xil_br idge	459	617	0	0	UART RX/TX → AXIL host bridge
u_debu g_sram	debug_ sram	368	454	4.0	0	monbu s trace capture buffer (4096 × 32 b)
u_rd_re sp_dela y	axi_re sponse _delay	352	66	2.0	0	R-side progra mmabl e memor y- latency injector
u_wr_b us_met er	axi_bu s_mete r	336	672	0	0	W-bus datapath- utilizati on meter
u_rd_b us_met er	axi_bu s_mete r	170	400	0	0	R-bus datapath- utilizati on

Instance	Module	LUTs	FFs	BRAM	DSP	Role
u_wr_resp_delay	axi_response_delay	93	51	0	0	meter W-side memory- latency injector
u_harness (glue)	stream_char_harness	73	455	0	0	reset sync + IO pin tie-off
u_led_status_driver	led_status_driver	29	73	0	0	200 Hz LED CDC driver
(root glue)	stream_char_top	1	2	0	0	clock buffer + pin wiring
Total	stream_char_top	29 124	28 504	14.5	1	45.9 % LUT / 22.5 % FF / 8.7 % BRAM / 0.4 % DSP of xc7a10 0t

1.1.2 0.2 Inside the DUT — stream_top_ch8

Roughly **half the DUT's area is monitor infrastructure**, not DMA logic. The per-channel scheduler-side monitors aggregate into a tree (u_monbus_aggregator inside each scheduler_group, then the top-level u_monbus_aggregator inside scheduler_group_array), the descriptor read bus has its own full AXI monitor (u_desc_axi_monitor = axi4_master_rd_mon), and u_monbus_axil_group drains all of that into the host AXIL (slow

path = IRQ + slice counters, fast path = bulk trace into debug_sram). This split is deliberate — the whole reason the engine is measurable on real silicon is that the monitor system is built in.

Instance	Module	LUTs	FFs	BRAM	DSP	Role
u_stream_core	stream_core	17 931	18 756	8.0	1	datapath + scheduler + per-channel monitors
↳ u_scheduler_group_array	scheduler_group_array	15 033	16 488	0	0	8 × scheduler_group + desc-bus monitor + monbus tree
↳ ↳ u_desc_axi_monitor	axi4_master_rd_mon	3 996	4 618	0	0	monitor on descriptor AXI bus (trans-mgr, addr-check, reporter)
↳ ↳ u_monbus_aggregator	monbus_arbiter	1 587	2 202	0	0	monitor RR arbiter merge

Instance	Module	LUTs	FFs	BRAM	DSP	Role
r (top)						g 8 group streams → 1 stream
↳ 8 × scheduler_group	scheduler + desc-engine + per-group monitors	~1 170 ea	~1 206 ea	0	0	per-channel scheduler/desc-engine (~770 LUT) + monitor aggregator (~400 LUT)
↳ u_sram_controller	sram_controller	1 350	992	8.0	0	per-channel write-side SRAM + alloc/drain ctrl
↳ u_axi_write_engine	axi_write_engine	618	296	0	0	shared AW/W master, pipelined req register, AW-issue logic
↳	axi_re	229	152	0	0	shared

Instance	Module	LUTs	FFs	BRAM	DSP	Role
u_axi_read_engine	ad_engine					AR master, pipelined req register, alloc handshake
↳ u_rd_axi_skid	axi4_master_rd	312	346	0	0	AR/RAXI skid buffer
↳ u_wr_axi_skid	axi4_master_wr	208	166	0	0	AW/W/BAXI skid buffer
↳ u_perf_profiler	perf_profiler	246	316	0	1	per-channel performance profiler (uses the lone DSP for moving avg)
u_monbus_axil_group	monbus_axil_group	572	571	0	0	monitor sink — drains packets to AXIL (IRQ +

Instance	Module	LUTs	FFs	BRAM	DSP	Role
						bulk trace)
u_stream_regs	stream_regs	822	643	0	0	PeakRD L-generated STREAM CSR file
u_peakrdl_adapter	peakrdl_to_cmdrsp	1 121	84	0	0	PeakRD L cmd/rsp adapter
u_apb4_todescr	apb4_to_desc	402	39	0	0	APB → descriptor-load shim
u_apb4_slave (passthrough)	apb4_slave	146	204	0	0	APB slave for STREAM config
Subtotal (u_stream)		20 370	20 568	8.0	1	
Of which monitors: u_desc_axi_monitor +		~9 400	~12 480	0	0	~46 % LUT / ~61 % FF of u_stream

Instance	Module	LUTs	FFs	BRAM	DSP	Role
monbu s tree + monbu s_axil_g roup + 8 × per- group monbu s_aggre gator Of which DMA core: schedul er + desc- engine ×8 + engines + sram_c ontrolle r + skids + perf + CSR plumbi ng		~11 000	~8 100	8.0	1	

1.1.3 0.3 What this means

Two things worth saying out loud before any dynamic number lands:

1. **The “DMA proper” is about 11 k LUTs.** Even at 8 channels with the full monbus-tree + descriptor-bus monitor + AXIL sink wired in, the engine itself (scheduler + descriptor engines + AXI read/write engines + sram controller

- skids) fits in roughly a third of the LUTs of the whole built design. The rest is harness (~9 k LUT) and monitor infrastructure (~9 k LUT).
2. **Monitor area is intentional, and it scales with channel count.** Every `scheduler_group` carries its own `monbus_aggregator` (~400 LUT each — those 8 instances dominate the visible monitor cost), the descriptor bus has a full `axi4_master_rd_mon` on it (~4 k LUT, single instance), and everything funnels through a top-level RR arbiter into the AXIL group. If a deployment doesn't need this much observability, the bridge generator already supports building without it — see the monitor whitepaper §2 “Where to insert monitoring” for the tradeoff.

1.1.4 0.4 Timing-closure notes (8-channel build)

Closing 100 MHz at 8 channels on the -1 part needed three changes after the initial bring-up (4-channel build), all in the data plane, none in the harness:

1. **Pipeline the SRAM controller's per-channel availability outputs** (`axi_rd_alloc_space_free`, `axi_wr_drain_data_avail`, `axi_wr_sram_valid`). Source-side flop break: the alloc/drain comparator cones at 8 channels land in the same logic cone as the read/write engine's request masking, adding ~5 levels of comparator depth to a path that was already 14 levels. See `sram_controller.sv` commits `a65c6068 + b619eee9`.
2. **Pipeline the per-engine arbiter request** (`r_arb_request`). The cone `scheduler.r*_beats_remaining` → 32-bit comparators → `w_data_ok` → `w_arb_request` → arbiter priority encoder → `grant_reg` ran 16 levels deep at 8 channels, missing 100 MHz by ~1.5 ns. Registering `w_arb_request` immediately before the arbiter cuts the cone into two short paths. The AW-issue gate and the AR-valid mask use live `sched*_valid` so a stale grant cannot capture an AW / drive an AR for a channel the scheduler is no longer requesting; `w_arb_grant_ack` auto-releases stale grants so the arbiter doesn't stall waiting for a handshake that will never come. See `axi_write_engine.sv / axi_read_engine.sv` commit `4e8f9e02`.
3. **Switch impl strategy to Performance_Explore.** After (1) and (2), the only remaining negative-slack endpoint was an observability-side carry chain in the descriptor monitor's `trans_mgr.r_active_count` (-17 ps, 13-level cone with 7.3 ns of route delay). Strategy switch closes it without RTL change. Commit `c2d76776`.

Post-route: WNS = +0.296 ns, 0 failing endpoints, 0 hold violations.

Fixed — 5+ active channels were hanging due to a host-side CSR address-map collision, not RTL. The harness KICK_G0 register sits at 0xC0, in the middle of the per-channel kick-address slot range. The RTL splits the 8 kick-address slots into two banks (ch 0..3 at 0xB0..0xBC, ch 4..7 at 0xC4..0xD0) and skips 0xC0. Both `run_characterization.py` and the cocotb TB `stream_char_tb.py` were computing the per-channel slot as $\text{BASE} + 4 \cdot \text{ch}$, which lands on 0xC0 for ch=4 — writing the kick address into KICK_G0 instead. The low 8 bits of the address are interpreted as the kick-fire bitmask, so a bogus single-cycle kick fires for whichever channel that byte happens to encode (here channel 5, never loaded), and ch 4's real address never reaches its slot. Channels 4 and 5 then kick to address 0, fetch garbage descriptors, never produce the expected beats, and the timer times out at 60 s. The bug never surfaced before because no prior build had hardware presence for channels 4+ (NUM_CHANNELS=4 baseline). Replaced the bare $\text{BASE} + 4 \cdot \text{ch}$ with a `kick_addr_csr(ch)` helper in both consumers that hops the 0xC0 slot; no RTL change, no rebuild. 5–8 channel runs now complete at 94 % datapath utilization, matching 1–4.

1.2 1. What we set out to do

The STREAM engine is a multi-channel descriptor-based DMA we built in-house. Before declaring it ready and dropping it into a system, we wanted to **measure** what it actually does — not just whether the tests pass — and to do it on real silicon, not in simulation. Specifically:

1. **Headline throughput.** What does it achieve at zero memory latency, and how close does that come to the AXI ceiling?
2. **Latency tolerance.** A real DMA engine spends most of its life waiting for memory. How much round-trip latency can our outstanding-request pipeline absorb before throughput collapses, and where exactly is the cliff?
3. **Architectural payoff of multi-channel design.** The engine has 8 channels (4 enabled in this build to fit the 100T). Do extra channels buy us bandwidth, latency tolerance, neither, or both?
4. **Descriptor-chain amortization.** Long descriptor chains should hide startup cost. Does the descriptor prefetcher actually do that?
5. **Read/write balance.** Either side of the engine can be the bottleneck. We wanted independent visibility into both.

6. **A clean baseline for Phase 2.** When we drop in a Vivado-IP DMA later, we want to compare it against numbers that are reproducible, principled, and not contaminated by harness artifacts.

The headline answer (skip ahead if you just want it): **1435 MB/s sustained at 90 % of the AXI ceiling, 128 cycles of memory latency hidden per channel by the multi-outstanding pipeline, with the cliff falling cleanly on Little’s Law from there.** The rest of this document is the work behind those numbers.

1.3 2. Test platform

FPGA	Xilinx Artix-7 100T (xc7a100tcsg324-1, -1 speed grade)
Clock	aclk = 100 MHz (10 ns period)
AXI data width	128 bits (16 B / beat)
Channels in build	8 (NUM_CHANNELS = 8; characterization in §5 sweeps the active subset via ch_mask)
Burst length	16 beats (AWLEN = 0x0F)
Theoretical AXI peak	100 MHz × 16 B = 1600 MB/s at delay = 0
Engine outstanding queue	AR_MAX_OUTSTANDING = AW_MAX_OUTSTANDING = 8 (per channel)
Bitstream resources	29.1k LUTs (46 %), 28.5k FFs (22.5 %), 14.5 BRAM tiles (8.7 %), 1 DSP (0.4 %) — per-block breakdown in §0 above
Post-route timing	WNS + 0.296 ns (0 failing endpoints, 0 hold violations; Performance_Explore impl strategy)


The DUT is stream_top_ch8. Everything else around it is harness — not part of the IP but instrumented to make the engine measurable. The next section walks through the harness piece by piece, because the credibility of the results depends entirely on what the harness is and how it was hooked in.

- **Harness slaves.** Each AXI master is wrapped by a slave that produces or checks data with deterministic patterns:
 - `axi4_slave_rd_pattern_gen` produces an LFSR-driven byte stream and computes a per-channel CRC over what was actually read.
 - `axi4_slave_wr_crc_check` accepts writes and computes a per-channel CRC over what was actually written. End of test the read CRC and write CRC are compared.
 - `axi_response_delay` (R and B sides) injects programmable, pipelined memory latency between the slaves and the DUT — this is the hook that makes the latency-tolerance experiments possible.
- **Status out.** LEDs, 7-seg digits, and the IRQ pin show real-time and post-test status. PASS = 0x0123 on the LEDs and "0123" on the 7-seg; FAIL = 0x9999 everywhere.

Each instrumentation block earned its place after a bug or a measurement we couldn't make without it. The next subsections walk through the most important hooks.

1.4.3 3.2 Hook 1: The pipelined response-delay block (memory model)

1.4.4 Figure 3.2: Read- and write-path data flow



data path


The most important question we wanted answered was: *how much memory latency can the engine hide before throughput drops?* That requires injecting latency at the slave port, which means we need a memory model.

The first attempt was a single-stage gate: when a burst arrives, hold the first beat for L cycles, then pass the rest at line rate. It compiled and ran, but the data it produced was misleading. With a single-stage gate, the next burst's L -cycle wait can only start *after* the previous burst's response has fully drained, because the gate can only count one delay at a time. That serializes latency that real hardware would pipeline. Throughput collapses by the third cycle of L , which contradicts the engine's design intent (it has a multi-outstanding queue precisely so that latency can overlap).

A real DDR controller doesn't behave that way. Each request pays a fixed access latency, but those latencies *overlap* — multiple requests are in flight simultaneously, each paying L in parallel. After the initial L-cycle pipe-fill, throughput is decoupled from L (Little's Law).

So the second attempt — the one in the current build — is a pipelined timestamp-FIFO.

1.4.5 Figure 3.3: axi_response_delay internals



response delay internals

Every beat presented on the slave side is admitted to a CAPACITY-deep FIFO with a timestamp. A free-running cycle counter labels the slot. The drain gate fires when the head's (now - t_in) exceeds the runtime-programmable i_delay_cycles. Up to CAPACITY beats can be in flight at once. When the FIFO fills, s_ready drops, naturally back-pressuring the slave — that mirrors a real controller running out of in-flight slots.


Sized with headroom:

- **R channel:** CAPACITY = 256, $\geq \text{AR_MAX_OUTSTANDING}(8) \times \text{max_burst}(16)$
= 128 in-flight beats per channel $\times$ 4 channels.
- **B channel:** CAPACITY = 16, $\geq \text{AW_MAX_OUTSTANDING}(8)$. One BRESP per AW, so the queue stays small.

This is plumbed up through stream_char_cfg_pkg::CFG_RESP_DELAY_R_CAPACITY and CFG_RESP_DELAY_B_CAPACITY so config sweeps don't require RTL edits.

1.4.6 3.3 Hook 2: The hardware kick-burst path

1.4.7 Figure 3.4: CSR-driven kick-burst path




csr kick path

Multi-channel sweeps require all selected channels to start on the same aclk, otherwise channel-N is sometimes still in pipe-fill when channel-1 is in steady state, and the cross-channel interleave at the slaves becomes unreproducible. Originally the host fired one APB write per channel to start each one, but at 115 200 baud each APB write is ~870 μs of UART time — which means ch3 starts ~2.6 ms after ch0. That's enough skew to make the channel × delay heatmap uninterpretable.

The fix in the current harness is a CSR-driven kick-burst. The host writes all four channel start addresses into shadow registers (CH_KICK_ADDR[0..3]) — these are persistent and don't trigger anything. Then the host writes one byte to KICK_GO with a bitmask of which channels to start. KICK_GO produces a 1-aclk pulse to the selected channels, all of which see the kick on the same edge. UART skew is squeezed out of the experiment at the source.

1.4.8 3.4 Hook 3: The cycle-stamp timer windows

1.4.9 Figure 3.5: Timer windows on the test timeline



timer windows

Every test captures three independent cycle counts in harness_csr:

Field	Window
total_cycles	first descriptor AR handshake → last B response
r2r_cycles	first R beat → last R beat (read engine only)

Field	Window
w2w_cycles	first W beat → last W beat (write engine only)

The three windows let us separate three different things that look the same from the host:

- `total / time` is what software perceives — including descriptor fetch, last-burst drain, and any UART round-tripping at the boundaries.
- `r2r / time` is the **read engine's true sustained rate**, with startup and drain stripped out.
- `w2w / time` is the **write engine's true sustained rate**, independently.

If `total` lags `r2r` significantly, that's startup/drain overhead and we can quantify it. If `r2r` and `w2w` differ, one side is the bottleneck. If they all match, the engine is balanced and the harness isn't introducing artifacts.

1.4.103.5 Hook 4: Per-channel CRC slaves

Single-channel correctness is easy: pull a deterministic LFSR pattern from the read slave, push it into the write slave, compare the read-side CRC against the write-side CRC at the end. With one channel running there's no ambiguity about what data went where.

Multi-channel breaks that. If all channels share one slave's LFSR/CRC computation, the per-beat data ordering at the shared slave depends on AXI ID interleave between channels — which in turn depends on arbitration timing, which is exactly the thing the experiment is supposed to vary. The CRCs don't match, but not for the reason of any actual data corruption.

The fix is per-channel state inside the slaves: `axi4_slave_rd_pattern_gen` runs an independent LFSR per AXI ID, and `axi4_slave_wr_crc_check` computes one CRC per AXI ID. The DUT writes `awid = ch` and `arid = ch` so the slaves can demux. With per-channel CRCs, end-of-test verification compares `ch-0` read CRC against `ch-0` write CRC, `ch-1` against `ch-1`, and so on. A CRC mismatch on a channel now genuinely flags corruption on that channel, independent of how the channels interleaved at the AXI level.

1.4.11 3.6 Hook 5: Configuration package for sweeps

The interesting characterization knobs (engine outstanding-queue depth, response-delay queue capacities) used to live as defaults inside `stream_core.sv` — three modules deep. Sweeping them meant editing the leaf module of a shared component, which is the kind of change that breaks unrelated things.

The current build plumbs them up cleanly:

```
stream_char_cfg_pkg.sv      (single source of truth)
|
|  CFG_AR_MAX_OUTSTANDING
|  CFG_AW_MAX_OUTSTANDING
|  CFG_RESP_DELAY_R_CAPACITY
```


| CFG_RESP_DELAY_B_CAPACITY
 ▼
 stream_char_top.sv → stream_char_harness.sv → stream_top_ch8.sv → stream_core.sv

To run a different config (deeper outstanding queues, shallower queues, etc.) for a build campaign, only the package file changes. The skid buffers remain hard-coded inside `stream_core` — they are timing-only structures and don't belong in the characterization knob set.

1.4.123.7 What all of this gives us, together

With those hooks in place:

- The DUT runs untouched in its production form.
- The host can drive any sweep over a single serial port, with single-aclk channel kick-off.
- Memory latency is tunable at runtime via CSR, and the model behaves like real DRAM (pipelined) rather than like a single-stage gate.
- Independent read/write rate measurements isolate each side from the other.
- Per-channel CRC verification keeps multi-channel data integrity meaningful.
- Sweepable engine parameters are one-file-edits, not RTL hunts.

Each piece earned its place after producing a confusing or contradictory measurement that traced back to the missing hook. The next sections are the data those hooks made it possible to capture.

1.5 4. Headline numbers

At `delay = 0` (no simulated memory latency), the engine sustains **~1435 MB/s** — **about 90 % of the AXI ceiling**. The remaining 10 % is inter-burst arbitration / pipeline gaps; we measured ~1.06 cycles per beat instead of the ideal 1.0.

sweep	knob varied	result
Long transfer (1 ch, 16 desc, 8 MB total)	descriptors	flat at 1435 MB/s
All-active multi-channel (4 ch, 1 desc each, 2 MB total)	channels	flat at 1435 MB/s — slave-shared, <i>not</i> per-channel BW scaling
Transfer size at single	8 KB → 1 MB	rises 1383 → 1436 MB/s

sweep	knob varied	result
channel		as a fixed ~22-cycle startup amortizes
Per-beat memory latency	0 → 96 cycles	flat at 1430+ MB/s (engine pipeline absorbs everything in this window)
Per-beat memory latency	128 → 512 cycles	linear cliff: 1272 → 363 MB/s, slope tracks Little's Law $BW = 128/L \times \text{peak}$


The big architectural takeaway: the engine's multi-outstanding pipeline absorbs ~128 cycles of memory round-trip latency completely — that's `AR_MAX_OUTSTANDING` (8) × `burst_len` (16) beats kept in flight at all times. Past that point throughput degrades linearly with `L`, exactly as Little's Law predicts. This is the textbook behavior of a well-designed multi-outstanding AXI master sitting in front of a pipelined memory controller, and it lines up cleanly with the architectural intent.

A pleasant surprise: each channel maintains its own AR/AW outstanding queue, so the cliff position scales **linearly with channel count** — 2 channels hide ~256 cycles, 3 channels hide ~384, and 4 channels hide ~512. See §6.1 for the full data.

1.6 5. Single-axis sweeps

1.6.1 5.1 Response-delay sweep (1 ch, 1 desc, 512 KB)

1.6.2 Figure 5.1: Throughput vs. memory latency (1 channel)



delay sweep

delay (cyc)	total_cycles	throughput (MB/s)	Δ vs delay-0
0	34,850	1434.7	—
16	34,865	1434.1	−0.04 %
32	34,881	1433.4	−0.09 %
64	34,913	1432.1	−0.18 %
96	34,945	1430.8	−0.27 %
112	35,216	1419.8	−1.0 %
128	39,312	1271.9	−11.3 %
144	43,408	1151.9	−19.7 %
160	47,504	1052.5	−26.6 %
192	55,696	897.7	−37.4 %
224	63,888	782.6	−45.4 %

delay (cyc)	total_cycles	throughput (MB/s)	Δ vs delay-0
256	72,080	693.7	−51.7 %
320	88,464	565.2	−60.6 %
384	104,848	476.9	−66.8 %
448	121,232	412.4	−71.2 %
512	137,616	363.3	−74.7 %

Two regions:

- **delay 0–96: completely absorbed.** Adding 96 cycles of per-beat memory latency costs less than 0.3 % of total runtime — the whole 96×1 cycle is paid as a one-time pipe-fill at the start of the transfer, and the rest streams at line rate. The 128-beat in-flight window has plenty of headroom in this regime.
- **delay ≥ 128 : linear cliff.** Once L exceeds the in-flight window, Little’s Law takes over: $BW \approx (\text{in_flight_beats} / L) \times \text{peak}$. With $\text{in_flight} = 128$ and $\text{peak} \approx 1435$ MB/s, the predicted post-cliff throughput is $128/L \times 1435$. Measured vs predicted:


delay	measured	Little’s Law ($128/L \times 1435$)	match
128	1272	1435	knee (in-flight just barely covers L)
256	694	718	~3 % below (residual arbitration overhead)
384	477	478	exact
512	363	359	exact

The fit gets tighter as L grows because the relative contribution of arbitration overhead shrinks. Past $L = 384$ we are within 1 % of theory.

`r2r` and `w2w` track `total` to within ~30 cycles in the flat region — the read and write engines are perfectly balanced. In the post-cliff region they remain matched (slightly above `total` since they exclude startup), confirming both sides degrade together.

1.6.3 5.2 Descriptor-count sweep (1 ch, 1..16 descriptors, 512 KB each, delay 0)

1.6.4 Figure 5.2: Throughput vs. descriptor chain length



desc sweep


descriptors	total moved	throughput
1	512 KB	1434.7
4	2 MB	1434.8
8	4 MB	1434.8
16	8 MB	1434.8

Throughput is **flat to four decimal places** as the chain grows. Two implications:

1. The descriptor engine fetches the next descriptor *concurrently* with the data engines draining the previous one — there is no pause between descriptors in a chain.
2. The fixed startup/drain cost (~35 cycles per session) is amortized cleanly across longer transfers; the gap between `total_MBps` and `r2r_MBps` shrinks toward zero as `N` grows (1.48 MB/s gap at `N = 1` collapses to 0.09 MB/s by `N = 16`).

1.6.5 5.3 Channel-count sweep (1..4 ch, 1 desc each, 512 KB per channel, delay 0)

1.6.6 Figure 5.3: Throughput vs. active channel count



channel sweep

channels	total moved	throughput
1	512 KB	1434.7
2	1 MB	1435.3
3	1.5 MB	1434.9
4	2 MB	1435.0


All four configurations land at the same ~1435 MB/s. Interpretation:

- **The slave AXI port is the bandwidth ceiling.** All channels share the same read source and write sink in this harness, so adding channels splits the bandwidth across more streams rather than scaling it. (Channels gain *latency tolerance*, not raw bandwidth — see §6.1.)
- **Arbitration overhead is essentially zero.** Going from 1 to 4 active channels costs less than 0.1 %.

This is the right behavior for *this* test setup. With independent backing memory per channel (e.g. different DDR ranks, distinct AXI slaves) you'd see proper N× bandwidth scaling — that's a follow-on test.

1.6.7 5.4 Transfer-size sweep (1 ch, 1 desc, 8 KB → 1 MB, delay 0)

1.6.8 Figure 5.4: Software-visible throughput vs. transfer size



size sweep

size	total_cycles	throughput
8 KB	565	1382.7
16 KB	1,109	1408.9
32 KB	2,197	1422.4
64 KB	4,373	1429.2
128 KB	8,725	1432.7
256 KB	17,429	1434.4
512 KB	34,837	1435.3
1 MB	69,653	1435.7


The r2r column stays at ~1438 MB/s across every size — the engine’s true sustained rate, independent of startup/drain. total rises asymptotically toward r2r as the fixed ~22-cycle startup overhead becomes a smaller fraction of the run. Useful to know: at 8 KB the software-

perceived throughput is ~3.7 % below peak; by 64 KB it's within 0.5 %; by 256 KB it's indistinguishable.

1.7 6. 2-D crosses


1.7.1 6.1 Channels × delay (1 desc, 512 KB) — *the key result*

1.7.2 Figure 6.1: Throughput surface — channels × memory latency



channels × delay heatmap

1.7.3 Figure 6.2: Per-channel-count slices



channels × delay lines

The cliff position **scales linearly with channel count** — each channel maintains its own AR/AW outstanding queue, so total in-flight beats is $8 \times N_channels \times burst_len = 128 \times N_ch$.

L	1 ch	2 ch	3 ch	4 ch
0	1435	1435	1435	1435
32	1434	1435	1434	1434
64	1433	1434	1434	1434
96	1431	1433	1434	1434
128	1272	1433	1433	1433
160	1053	1432	1433	1433
192	898	1431	1432	1433
256	694	1385	1431	1432
384	477	952	1012	1013
512	363	726	759	timeout

(Bold marks each curve's knee — where the engine's per-channel pipe stops covering the latency.)


What this shows:

- **1 channel hides L up to ~128 cycles** (8 outstanding $\times$ 16-beat burst). At $L = 128$ we hit the knee; throughput is still 1272 MB/s but slipping. Past $L = 128$, BW collapses linearly.
- **2 channels push the knee to ~256 cycles.** At $L = 256$ (the predicted boundary) BW is already at 1385 MB/s — slipping but well above the 1-channel value at the same L (694 MB/s). Past 256, BW follows $BW \approx 256/L \times \text{peak}$.
- **3 channels push the knee to ~384 cycles.** At $L = 256$ 3-ch is still at full peak (1431 MB/s); at $L = 384$ BW is 1012 MB/s — right at the knee.
- **4 channels** should push the knee out to ~ 512 (4×128). At $L = 256$ 4-ch is at peak (1432 MB/s), confirming the absorption is working. The $L = 512$ / 4-ch case timed out (host watchdog), which is consistent with 4 channels still sustaining peak demand at exactly the boundary.

This is **the architectural payoff of multi-channel design with per-channel outstanding queues** — channel parallelism doesn't increase peak bandwidth in a slave-shared setup, but it directly multiplies **latency tolerance**. A 4-channel transfer can hide DDR-class round trips (say, 100–200 cycles) without losing any throughput at all.


1.7.4 6.2 Descriptors $\times$ delay (1 ch, 512 KB)

1.7.5 Figure 6.3: Throughput surface — descriptors $\times$ memory latency



descriptors $\times$ delay heatmap

1.7.6 Figure 6.4: Per-descriptor-count slices



descriptors × delay lines


L	1 desc	2 desc	4 desc	8 desc	16 desc
0	1435.3	1435.3	1435.3	1435.3	1435.3
64	1432.7	1431.4	1430.8	1430.4	1430.3
128	1272.3	1270.3	1269.2	1268.7	1268.5
256	693.8	692.6	692.0	691.7	691.5
512	363.4	362.7	362.4	362.2	362.1

Cliff position is **invariant under descriptor count** — exactly as expected. Descriptor chaining is about amortizing startup/drain overhead, not about hiding memory latency. The engines' AR/AW outstanding queues span across descriptor boundaries (the descriptor engine prefetches concurrently), so the once-per-transfer pipe-fill cost is paid once per chain regardless of how many descriptors live inside it.

The (small) gap between `r2r/w2w` and `total` shrinks with chain length at every delay value — that's the chain-amortization signal. For example, at `L = 0` the 1-desc gap is 1.48 MB/s and the 16-desc gap is 0.09 MB/s.


1.7.7 6.3 Channels × descriptors (delay 0)

1.7.8 Figure 6.5: Throughput surface — channels × descriptors at zero latency



ch × desc heatmap

1.7.9 Figure 6.6: Per-descriptor slices at zero latency




ch × desc lines

The “ideal-system” reference plane. Every (channels, descriptors) combination at delay = 0 lands within 0.1 % of 1435 MB/s — a clean 2-D rectangle confirming the engine is well-balanced and not sensitive to either knob in isolation.

(The 4-channel rows in this sweep tripped the host-side watchdog because the campaign script ran them after the L = 512 timeout in §6.1. Re-running standalone gives the same flat ~1435 MB/s as §5.3 already showed for 4 channels at delay = 0.)

1.8 7. The architectural picture: Little's Law on real silicon

1.8.1 Figure 7.1: Pre-cliff, knee, post-cliff regimes



cliff math

The cleanest way to summarize the latency-tolerance result is in three regimes:

1. **Pre-cliff ($L \leq \text{in_flight} \times \text{cycles_per_beat}$).** The engine keeps the slave's response pipe full. Adding L cycles of latency only shifts total runtime by L cycles total — that's the one-time pipe-fill cost. Throughput stays at peak.
2. **At the knee ($L \approx \text{in_flight}$).** The pipe is just barely enough to cover L . Throughput begins to slip; we measure $\sim 11\%$ at $L = 128$ for 1 channel.
3. **Post-cliff ($L > \text{in_flight}$).** The engine has to wait for responses. By Little's Law, sustained throughput is $BW = (\text{in_flight} / L) \times \text{peak}$. We measured this within 1% across $L = 384\text{--}512$.

Multi-channel scales the cliff position because each channel has its own AR/AW outstanding queue: $\text{total_in_flight} = N_{\text{ch}} \times 8 \times 16 = 128 \times N_{\text{ch}}$. The 4-channel build can hide up to 512 cycles of latency without throughput loss — comfortably more than the typical DDR4 round-trip in a real system.

This is the result we hoped for and the harness was built to measure: the engine behaves on real silicon exactly the way the architecture predicted. There's no hidden bottleneck, no bookkeeping cost we missed, no surprise non-linearity. The engine is doing what it says on the box.

1.9 8. What we learned

1. **The engine hits 90 % of the AXI ceiling at zero memory latency.** ~ 1435 MB/s vs the 1600 MB/s theoretical peak. That's a clean number to compare a vendor IP against.

2. **The multi-outstanding pipeline absorbs ~128 cycles of memory latency per channel.** This matches the engine's design parameters exactly: $AR_MAX_OUTSTANDING(8) \times burst_len(16) = 128$ in-flight beats. The architectural intent is delivered.
 3. **Past the cliff, throughput follows Little's Law.** $BW \approx (N_{in_flight} / L) \times peak$, with measured values within 1 % of theory at $L \geq 384$. There is no exotic non-linearity — the design is pipeline-bound in exactly the way it should be.
 4. **Multi-channel scales latency tolerance, not bandwidth, in a slave-shared system.** Each channel has its own AR/AW outstanding queue, so 4 channels = $4 \times 128 = 512$ cycles of latency hidden without throughput loss. A real architectural lever for tolerating slow memory.
 5. **Read and write engines are matched.** $r2r_MBps$ and $w2w_MBps$ track each other within tenths of a MB/s across every sweep. Neither side is the bottleneck individually.
 6. **Long descriptor chains cost essentially nothing.** The descriptor engine prefetches concurrently with data movement, and the engines' outstanding queues span descriptor boundaries — so chaining is for software convenience, not performance.
 7. **Software-visible (total) vs sustained ($r2r/w2w$) gap is purely startup/drain.** ~22–35 cycles fixed cost at session boundaries, amortized to invisibility above ~256 KB transfers.
-

1.10 9. Setup for Phase 2: Vivado IP DMA comparison

For the head-to-head with a Vivado-IP DMA of comparable capability (scatter-gather, multi-channel, AXI4 slave-side memory, descriptor chaining), we'll repeat the same five sweep CSVs and the same 2-D crosses. The harness, plot script, CSV format, and PDF/DOCX generation pipeline all stay the same — the only thing that changes is the DUT instantiated inside `stream_top_ch8`'s footprint. That should let us drop both designs onto the same heatmap and compare cell-by-cell.

The PPA dimensions to compare:

Axis	How to measure
Performance	The throughput curves above — peak, cliff position, post-cliff slope
Area	LUT / FF / BRAM / DSP utilization from reports/utilization_impl.txt (this build: 15.9k LUTs, 14.6k FFs, 9.5 BRAM, 0 DSP)
Power	reports/power.txt total + per-clock-

Axis	How to measure
Timing	domain breakdown reports/timing_summary.txt WNS / WHS — does the IP close timing at the same fmax on the same part?

The most interesting comparison points will be:

- **Cliff position vs configured outstanding count.** Does the Vivado IP let you trade outstanding-queue depth for area? At what BRAM / LUT cost?
- **Post-cliff slope.** Does it follow $N/L \times$ peak cleanly, or does arbitration overhead dominate?
- **Multi-channel latency tolerance.** Does the vendor IP also use per-channel outstanding queues, or is the queue shared across channels (which would flatten the channel $\times$ delay benefit observed here)?
- **Area efficiency.** MB/s / (LUT + FF) at $L = 0$ and (LUT + FF) per cycle of latency hidden. Two single numbers that capture the PPA story.